

Subqueries

Boolean Operators IN, EXISTS,
ANY, ALL



Announcement

- MT Tutorial Session

Monday 3/7 - 1⁰⁰pm - 1304 SC

Jargon

Beer → Book

Bar → Lib

Drinker → Scholar

green → red

Joe Bar ~ Job

Green St Bar →

Red St
Lib,

Subqueries

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- Example: in place of a relation in the FROM clause, we can place another query, and then query its result.
 - Better use a tuple-variable to name tuples of the result.

*select name, price
from belts, (select from ...)
where ...*

Rel. ↑

45

select price from Sells → 1. holds ⇒

price
3.5

 { (3.5) }

Subqueries that Return Scalar

where bar = 'joe bar'
^ beer = 'bud'

- If a subquery is guaranteed to produce one tuple with one component, then the subquery can be used as a value.
 - “Single” tuple often guaranteed by key constraint.
 - A run-time error occurs if there is no tuple or more than one tuple.

Sells (bar, beer, price)

when price > 3.5

Example

- From Sells(bar, beer, price), find the bars that serve Miller for the same price Joe charges for Bud.
- Two queries would surely work:
 1. Find the price Joe charges for Bud.
 2. Find the bars that serve Miller at that price.

Query + Subquery Solution

```
SELECT bar
FROM Sells
WHERE beer = 'Miller' AND
      price = (SELECT price
                FROM Sells
                WHERE bar = 'Joe Bar'
                AND beer = 'Bud');
```

∈ The IN Operator

- $\langle \text{tuple} \rangle$ IN $\langle \text{relation} \rangle$ is true if and only if the tuple is a member of the relation.
 - $\langle \text{tuple} \rangle$ NOT IN $\langle \text{relation} \rangle$ means the opposite.
- IN-expressions can appear in WHERE clauses.
- The $\langle \text{relation} \rangle$ is often a subquery.

price in { }

Example

- From Beers(name, manf) and Likes(drinker, beer), find the name and manufacturer of each beer that Fred likes.

SELECT *

FROM Beers

WHERE name IN (SELECT beer

The set of
beers Fred
likes

FROM Likes

WHERE drinker = 'Fred');

The Exists Operator *not empty*

- EXISTS(<relation>) is true if and only if the <relation> is not empty.
- Being a boolean-valued operator, EXISTS can appear in WHERE clauses.
- Example: From Beers(name, manf), find those beers that are the only beer by their manufacturer.

Example Query with EXISTS

SELECT name

FROM Beers b1

WHERE NOT EXISTS(

SELECT *

FROM Beers

WHERE manf = b1.manf AND

name <> b1.name);

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

Notice the SQL “not equals” operator

Set of beers with the same manf as b1, but not the same beer

The Operator ANY

All

→ one attr

- $x = \text{ANY}(\text{<relation>})$ is a boolean condition meaning that x equals at least one tuple in the relation.
- Similarly, $=$ can be replaced by any of the comparison operators.
- Example: $x \geq \text{ANY}(\text{<relation>})$ means x is not smaller than all tuples in the relation.
 - Note tuples must have one component only.

The Operator ALL

- Similarly, $x \neq \text{ALL}(\langle \text{relation} \rangle)$ is true if and only if for every tuple t in the relation, x is not equal to t .
 - That is, x is not a member of the relation.
- The \neq can be replaced by any comparison operator.
- Example: $x \geq \text{ALL}(\langle \text{relation} \rangle)$ means there is no tuple larger than x in the relation.

Example

- From Sells(bar, beer, price), find the beer(s) sold for the highest price.

SELECT beer

FROM Sells

WHERE price >= ALL(
SELECT price
FROM Sells);

price from the outer
Sells must not be
less than any price.

CS411

Database Systems

06b: SQL-2

Grouping and Aggregation

Why Do We Learn This?

- organize (grouping)
- Gen. Report/Stat.

Q: What is “aggregate”?

Avg, Max, Sum, ...

Rel.

$\{v_1, \dots, v_n\} \rightarrow V$

Avg(price ...) $\rightarrow 35$

Aggregations

- SUM, AVG, COUNT, MIN, and MAX can be applied to a column in a SELECT clause to produce that aggregation on the column.
- Also, COUNT(*) counts the number of tuples.

freq.

Example: Aggregation

- From Sells(bar, beer, price), find the average price of Bud:

```
SELECT AVG(price)  
FROM Sells  
WHERE beer = 'Bud';
```

Eliminating Duplicates in an Aggregation

- DISTINCT inside an aggregation causes duplicates to be eliminated before the aggregation.
- Example: find the number of different prices charged for Bud:

```
SELECT COUNT(DISTINCT price)
FROM Sells
WHERE beer = 'Bud';
```

NULL's Ignored in Aggregation

- NULL never contributes to a sum, average, or count, and can never be the minimum or maximum of a column.
- But if there are no non-NULL values in a column, then the result of the aggregation is NULL.

Example: Effect of NULL's

```
SELECT count(*)  
FROM Sells  
WHERE beer = 'Bud';
```

The number of bars
that sell Bud.

```
SELECT count(price)  
FROM Sells  
WHERE beer = 'Bud';
```

The number of bars
that sell Bud at a
known price.

Grouping

- We may follow a SELECT-FROM-WHERE expression by GROUP BY and a list of attributes.
- The relation that results from the SELECT-FROM-WHERE is grouped according to the values of all those attributes, and any aggregation is applied only within each group.

Example: Grouping

- From Sells(bar, beer, price), find the average price for each beer:

```
SELECT beer, AVG(price)
FROM Sells
GROUP BY beer;
```

Example: Grouping

- From Sells(bar, beer, price) and Frequents(drinker, bar), find for each drinker the average price of Bud at the bars they frequent:

```
SELECT drinker, AVG(price)
FROM Frequents, Sells
```

```
WHERE beer = 'Bud' AND
      Frequents.bar = Sells.bar
GROUP BY drinker;
```

Compute
drinker-bar-
price of Bud
tuples first,
then group
by drinker.

Restriction on SELECT Lists With

Aggregation

- If any aggregation is used, then each element of the SELECT list must be either:

1. Aggregated, or s.t. $Avg(price)$
2. An attribute on the GROUP BY list.

Gr. w/ one value

a) select beer, avg(price)
from sells
group by beer

beer avg(-)

bud	
s.a	

b) select avg(price)
from sells
where beer = 'bud'
⇒ 1 group
No gr. attr.

Q: How about this query?

SQ L

~~SELECT bar, MIN(price)
FROM Sells
WHERE beer = 'Bud';~~


one group
no gr. attr

⇒ (bar with min price, the price)

~~select bar, avg(price)
from sells~~

sel name
avg(GPA)
~
~ From Stu.

Q: How to do it right, then?

 **SELECT bar, MIN(price)**
FROM Sells
WHERE beer = 'Bud';

only combined stuff

 price < min(price)

Q: How to do it right?

① where price \leq all (select min(price) from ...)
(select price from sells ...)