

# Updating Views

How can I insert a tuple into a table that doesn't exist?

```
CREATE VIEW JoeBarSells AS  
  SELECT beer, price  
  FROM Sells  
  WHERE bar = 'joe bar';
```

If we make the  
following insertion:

```
INSERT INTO JoeBarSells  
VALUES("bud special", 3.5)
```

It becomes:

```
INSERT INTO Sells  
VALUES(NULL, 'bud special', 3.5)
```

Q: Is the new tuple in table Sells? In the view JoeBarSells?

# Non-Updatable Views

```
CREATE VIEW Champaign-view AS  
  
  SELECT name, product, store  
  FROM   Person, Purchase  
  WHERE  Person.city = "Champaign" AND  
         Person.name = Purchase.buyer
```

How can we add the following tuple to the view?

("Joe", "Shoe Model 12345", "Nine West")

We need to add "Joe" to Person first. One copy ? More copies ?

# CS411

## Database Systems

06d: SQL-4  
Constraints and Triggers

view



# Why Do We Learn This?

Expected behavior of program  
"Assertion"



Expected behavior of data  
"health"

part of schema

# Constraints and Triggers

①      ②      View updates



①

- A constraint is a relationship among data elements that the DBMS is required to enforce.

rel.  
tables"  
attr.  
values


②

- Example: key constraints.
- Triggers are only executed when a specified condition occurs, e.g., insertion of a tuple.
  - Easier to implement than many constraints.

## Kinds of Constraints

- Keys. = *unique null*
- Foreign-key, or referential-integrity.
- a) • Value-based constraints.
  - Constrain values of a particular attribute.
- b) • Tuple-based constraints.
  - Relationship among components.
- c) • Assertions: any SQL boolean expression.  
*table-based,*

ER



## Foreign Keys

"beer" = key at  
at Sells Beers.

a key that is  
at another table,



- Consider Relation Sells(bar, beer, price).



We might expect that a beer value is a real  
beer --- something appearing in Beers.name

- A constraint that requires a beer in Sells to  
be some key in Beers is called a foreign-key  
constraint.

constraint?

- must be in beers  
table.
- has to be  
a key.

# Expressing Foreign Keys

- Use the keyword REFERENCES, either:
  1. Within the declaration of an attribute, when only one attribute is involved.
  2. As an element of the schema, as:

FOREIGN KEY ( <list of attributes> )

REFERENCES <relation> ( <attributes> )

- Referenced attributes must be declared PRIMARY KEY or UNIQUE.



## Example: With Attribute

```
CREATE TABLE Beers (  
  name CHAR(20) PRIMARY KEY,  
  manf CHAR(20) );  
  
CREATE TABLE Sells (  
  bar CHAR(20),  
  beer CHAR(20) REFERENCES Beers(name),  
  price REAL );
```

*Handwritten annotations:*

- Beers (circled in red)
- name (circled in red)
- PRIMARY KEY (underlined in red)
- Unique* (written in red, with an arrow pointing to the PRIMARY KEY)
- Sells (circled in blue)
- REFERENCES (circled in blue)
- Beers(name) (underlined in blue)
- A blue arrow points from the REFERENCES circle to the name attribute in the Beers table.


## Example: As Element

```
CREATE TABLE Beers (  
  name CHAR(20) PRIMARY KEY,  
  manf  CHAR(20) );
```

```
CREATE TABLE Sells (  
  bar    CHAR(20),  
  beer   CHAR(20),  
  price  REAL,  
  FOREIGN KEY (beer) REFERENCES  
    Beers (name) );
```

# Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from attributes of relation  $R$  to the primary key of relation  $S$ , two violations are possible:
  1. An insert or update to  $R$  introduces values not found in  $S$ .
  2. A deletion or update to  $S$  causes some tuples of  $R$  to “dangle.”



## Actions Taken -- 1

Sells                      Beers?

joe bar	Strag Bar	10
---------	-----------	----

Str?	?
------	---

X                      X

- Suppose  $R = \text{Sells}$ ,  $S = \text{Beers}$ .
- An insert or update to Sells that introduces a nonexistent beer must be rejected X
- A deletion or update to Beers that removes a beer value found in some tuples of Sells can be handled in three ways.

① cascade  
 ② update  
 ③ delete,  
 ④ set Null

Sells                      Beers

J.B.	Budweiser
G.B.	Budweiser
	:

Bud	Buseh

① Budweiser                      ②

## Actions Taken -- 2

- The three possible ways to handle beers that suddenly cease to exist are:

1. Default : Reject the modification.
2. Cascade : Make the same changes in Sells.
  - ♦ Deleted beer: delete Sells tuple.
  - ♦ Updated beer: change value in Sells.
- X 3. Set NULL : Change the beer to NULL.  
beer@ Sells if beer is key

## Example: Cascade

- Suppose we delete the Bud tuple from Beers.
  - Then delete all tuples from Sells that have beer = 'Bud'.
- Suppose we update the Bud tuple by changing 'Bud' to 'Budweiser'.
  - Then change all Sells tuples with beer = 'Bud' so that beer = 'Budweiser'.

## Example: Set NULL

- Suppose we delete the Bud tuple from Beers.
  - Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Suppose we update the Bud tuple by changing 'Bud' to 'Budweiser'.
  - Same change.



## Choosing a Policy

- When we declare a foreign key, we may choose policies **SET NULL** or **CASCADE** independently for deletions and updates.
- Follow the foreign-key declaration by:  
**ON [UPDATE, DELETE][SET NULL  
CASCADE]**
- Two such clauses may be used.
- Otherwise, the default (reject) is used.



## Example

```
CREATE TABLE Sells (  
  bar    CHAR(20),  
  beer   CHAR(20),  
  price  REAL,  
  FOREIGN KEY (beer)  
    REFERENCES Beers(name)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE );
```

Check (not in current MySQL)

- Attribute-based
- Tuple-based

# Attribute-Based Checks

- Put a constraint on the value of a particular attribute.  
*a condition that must hold.*
- CHECK( <condition> ) must be added to the declaration for the attribute.
- The condition may use the name of the attribute, but any other relation or attribute name must be in a subquery.

## Example

① ~ attr check

② ~ any condition

```
CREATE TABLE Sells (
```

```
  bar CHAR(20),
```

```
  beer CHAR(20)
```

```
  CHECK ( beer IN
```

```
    (SELECT name FROM Beers)),
```

```
  price REAL CHECK ( price <= 5.00 )
```

) ; ③ when to "check"  
⇒ [insert this attr.  
update

i) name in Beers  
X key  
ii) "one-side"  
check.

# Foreign  
key?

## Q: How is Check different from Foreign Key?

```
... beer CHAR(20) CHECK ( beer IN  
                        (SELECT name FROM Beers))  
... price REAL CHECK (price <= 5.00)
```

1. The kind of conditions to enforce?
2. The timing/actions of enforcing?

## Timing of Checks

- An attribute-based check is checked only when a value for that attribute is inserted or updated.
  - Example: CHECK (price <= 5.00) checks every new price and rejects it if it is more than \$5.
  - Example: CHECK (beer IN (SELECT name FROM Beers)) not checked if a beer is deleted from Beers (unlike foreign-keys).

## Tuple-Based Checks

- CHECK ( <condition> ) may be added as another element of a schema definition.
- The condition may refer to any attribute of the relation, but any other attributes or relations require a subquery.
- Checked on insert or update only.

## Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5

```
CREATE TABLE Sells (
```

```
  bar① CHAR(20),
```

```
  beer CHAR(20),
```

```
  price② REAL,
```

```
  CHECK (bar = 'Joe's Bar' OR  
         price <= 5.00)
```

```
);
```

select avg(price)  
from ...

sub query to  
ref any  
attr/table.

only Joe Bar can be exp



## Q: Why do we need tuple-level check?

- We can do attribute-based check, why tuple level?
- Give examples that you need to use tuple check.

attr-level  
tuple-level

Table

## Assertions (not in current MySQL)

- These are database-schema elements, like relations or views.

CREATE Table sells (

- Defined by:

CREATE ASSERTION <name>  
CHECK ( <condition> );

- Condition may refer to any relation or attribute in the database schema.

schema + ~~error~~ assertion  
= rel. assertion

## Example: Assertion

? why use assert  
can I use ?

[ - attr check  
- tuple check

- In Sells(bar, beer, price), no bar may charge an average of more than \$5.

CREATE ASSERTION NoRipoffBars CHECK

(

NOT EXISTS (

SELECT bar FROM Sells

GROUP BY bar

HAVING 5.00 < AVG(price)

));

Bars with an  
average price  
above \$5

arg of price by bar

hold.  
if  
all  
bars  
cheap

## Example: Assertion

- In Drinkers(name, addr, phone) and Bars(name, addr, license), there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
  { (SELECT COUNT(*) FROM Bars) <=  
    (SELECT COUNT(*) FROM Drinkers)  
);
```

# Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.
  - Example: No change to Beers can affect FewBar. Neither can an insertion to Drinkers.

## Triggers. Motivation

- Attribute- and tuple-based checks have limited capabilities.
- Assertions are sufficiently general for most constraint applications, but they are hard to implement efficiently.
  - The DBMS must have real intelligence to avoid checking assertions that couldn't possibly have been violated.

## Triggers: Solution

- A trigger allows <sup>programmer.</sup> the user to specify when the check occurs.
- Like an assertion, a trigger has a general-purpose condition and also can perform any sequence of SQL database modifications.

## Event-Condition-Action Rules

Why more eff. than assertion?

- Another name for “trigger” is ECA rule or event-condition-action rule.
- Event: typically a type of database modification, e.g., “insert on Sells.” *→ insert on sells*
- Condition: Any SQL boolean-valued expression. *new value & Beers.name*
- Action: Any SQL statements. *→ insert new val to Beers.*



## Example: A Trigger

- There are many details to learn about triggers.
- Here is an example to set the stage.
- Instead of using a foreign-key constraint and rejecting insertions into Sells(bar, beer, price) with unknown beers, a trigger can add that beer to Beers, with a NULL manufacturer.

schema element  $\Rightarrow$  create table R (...)  
**Example: Trigger Definition**

CREATE TRIGGER BeerTrig

The event

AFTER INSERT ON Sells

REFERENCING NEW ROW AS NewTuple

FOR EACH ROW

WHEN (NewTuple.beer NOT IN  
(SELECT name FROM Beers))

The condition

INSERT INTO Beers(name)  
VALUES(NewTuple.beer);

The action

*Sells*

*"bud"*


*Beers*

<del>Bud</del>	<del>NULL</del>

What happen?  
if no,  $\Rightarrow$  reject.

# Options: The Event

- AFTER can be BEFORE.
  - Also, INSTEAD OF, if the relation is a view.
    - A great way to execute view modifications: have triggers translate them to appropriate modifications on the base tables.
- INSERT can be DELETE or UPDATE.
  - And UPDATE can be UPDATE . . . ON a particular attribute.

## Options: FOR EACH ROW

- Triggers are either *row-level* or *statement-level*.
- FOR EACH ROW indicates row-level; its absence indicates statement-level.
- Row level triggers are executed once for each modified tuple.
- Statement-level triggers execute once for an SQL statement, regardless of how many tuples are modified.

## Options: REFERENCING

- INSERT statements imply a new tuple (for row-level) or new set of tuples (for statement-level).
- DELETE implies an old tuple or table.
- UPDATE implies both.
- Refer to these by

[NEW OLD][TUPLE TABLE] AS <name>

# Options: The Condition

- Any boolean-valued condition is appropriate.
- It is evaluated before or after the triggering event, depending on whether BEFORE or AFTER is used in the event.
- Access the new/old tuple or set of tuples through the names declared in the REFERENCING clause. (or fixed by “OLD”, “NEW” in MySQL.)

## Options: The Action

- There can be more than one SQL statement in the action.
  - Surround by BEGIN . . . END if there is more than one.
- But queries make no sense in an action, so we are really limited to modifications.

## Another Example

- Using Sells(bar, beer, price) and a unary relation RipoffBars(bar) created for the purpose, maintain a list of bars that raise the price of any beer by more than \$1.



# The Trigger

CREATE TRIGGER PriceTrig

AFTER UPDATE OF price ON Sells

REFERENCING

OLD ROW as old

NEW ROW as new

FOR EACH ROW

WHEN(new.price > old.price + 1.00)

INSERT INTO RipoffBars

VALUES(new.bar);

The event –  
only changes  
to prices

Updates let us  
talk about old  
and new tuples

We need to consider  
each price change

Condition:  
a raise in  
price > \$1

When the price change  
is great enough, add  
the bar to RipoffBars

## Behind the Scene: Why Trigger was invented?

Aspects of a trigger subsystem in an integrated database system. Proceedings of the 2nd international conference on Software engineering.  
1976.

1. Extended assertions. (why?)
2. ??

# Behind the Scene: This is why...

## Aspects of a Trigger Subsystem in an Integrated Database System

by

Kapali P. Eswaran  
IBM Research Laboratory  
San Jose

view  
↑

ABSTRACT. This paper considers the specifications and design of a trigger subsystem in a database management system. The use of triggers as extended assertions and as a means to materialize virtual data objects are discussed. The functional requirements of a trigger subsystem and different implementation issues are studied. We also examine the relationships between a trigger subsystem and the rest of the database system, in particular the authorization and locking subsystems.