

## B+ Trees

+ Seq scan  
is fast.

What's wrong with sequential index?

\ - costly to  
maintain

# B-Trees/B+Trees: B     ?    ? Trees

Balanced,

- Intuition:

- Give up on sequentiality of index
- Try to get “balance” by dynamic reorganization

- B+trees:

- Textbook refers to B+trees (a popular variant) as B-trees (as most people do)
- Distinction will be clear later (ok to confuse now)

B-tree 1972

✓ B+trees

# Behind the Scene: UIUC (Alumni) Contribution!



*Prof. Rudolf Bayer*

*Rudolf Bayer studied Mathematics in Munich and at the University of Illinois, where he received his Ph.D. in 1966. After working at Boeing Research Labs he became an Associate Professor at Purdue University. He is a Professor of Informatics at the Technische Universität München since 1972 and ... ..*


*Urbana  
→ Champaign*

*The 2001 SIGMOD Innovations Award goes to Prof. Rudolf Bayer of the Technical University of Munich, for his invention of the B-Tree (with Edward M. McCreight), of B-Tree prefix compression, and of lock coupling (a.k.a. crabbing) for concurrent access to B-Trees (with Mario Schkolnick). All of these techniques are widely used in commercial database products. ....*

## The Original Publication

*Rudolf Bayer, Edward M. McCreight: Organization and Maintenance of Large Ordered Indices. Acta Informatica 1: 173-189 (1972)*

# Behind the Scene: And he said Hello!





# B+ Trees Basics (B-tree) a.w. $n = 2d$

- Parameter  $d$  = the degree
- Each node has  $[d, 2d]$  keys (except root)

at most 4  
 $v_1 | v_2 | v_3 | v_4$   
 2 values at least

$$d=2 \Rightarrow 2 \leq v \leq 4$$

Why define capacity?  
 store index

$\Rightarrow$  size of each node

= 1 "page"  
 "block"  
 (4k, 8k)

Internal node:

30	120	240

[30, 120)

[120, 240)

[240, Y)

[30, 120)

[240, 300)

3 key values

Leaf:

40	50	60

40

50

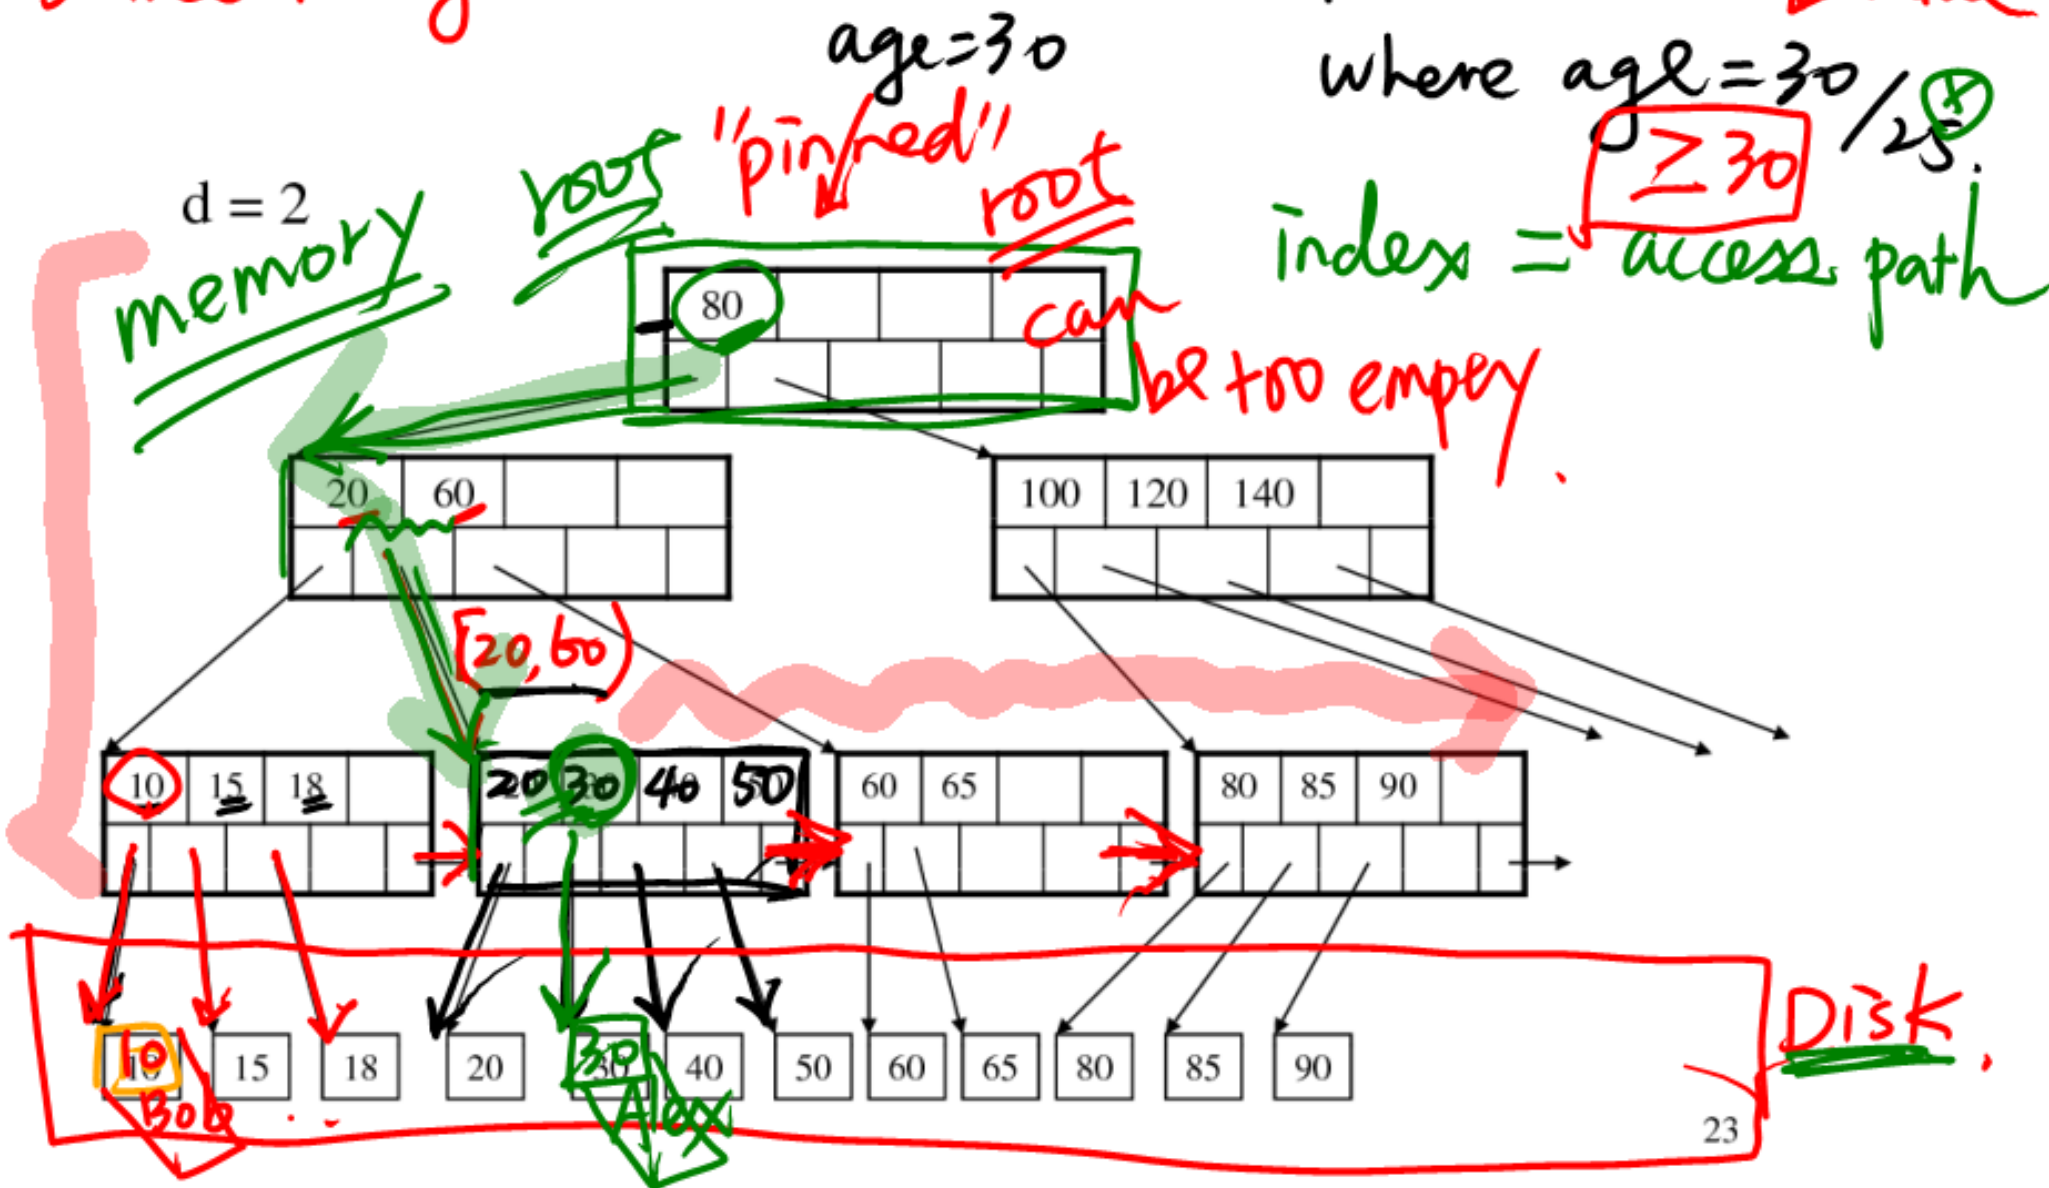
60

next leaf

Student (name, age)

# B+ Tree Example

Select name dec. by  
from stu  
where age = 30 / 25.  
index = access path



# B+ Tree Design




$d, 2d$  max # of keys  
 $\Rightarrow$  as large as u can fit on one disk block.

- How large  $d$ ?

- Example:

(integer)

- Key size = 4 bytes
- Pointer size = 8 bytes
- Block size = 4096 bytes



- $\underline{2d} \times \underline{4} + (\underline{2d+1}) \times \underline{8} \leq \underline{4096}$

- $d = \underline{170} \rightarrow$  at most  $2d = 340$  keys,  $2d+1 = 341$  ptrs.

# Searching a B+ Tree

- Exact key values:
  - Start at the root
  - Proceed down, to the leaf
- Range queries:
  - As above
  - Then sequential traversal

```
Select name  
From people  
Where age = 25
```

```
Select name  
From people  
Where 20 <= age  
and age <= 30
```



# B+ Trees in Practice

$d=100$   $100 \leq \leq 200$   
 Typical order: 100. Typical fill-factor: 67%.

- Typical order: 100. Typical fill-factor: 67%.

– average fanout = 133

- Typical capacities:

– Height 4:  $133^4 = 312,900,700$  records

– Height 3:  $133^3 = 2,352,637$  records

- Can often hold top levels in buffer pool:


– Level 1 = 1 page = 8 Kbytes

– Level 2 = 133 pages = 1 Mbyte

– Level 3 = 17,689 pages = 133 MBytes

~ top-level in mem

– tree short

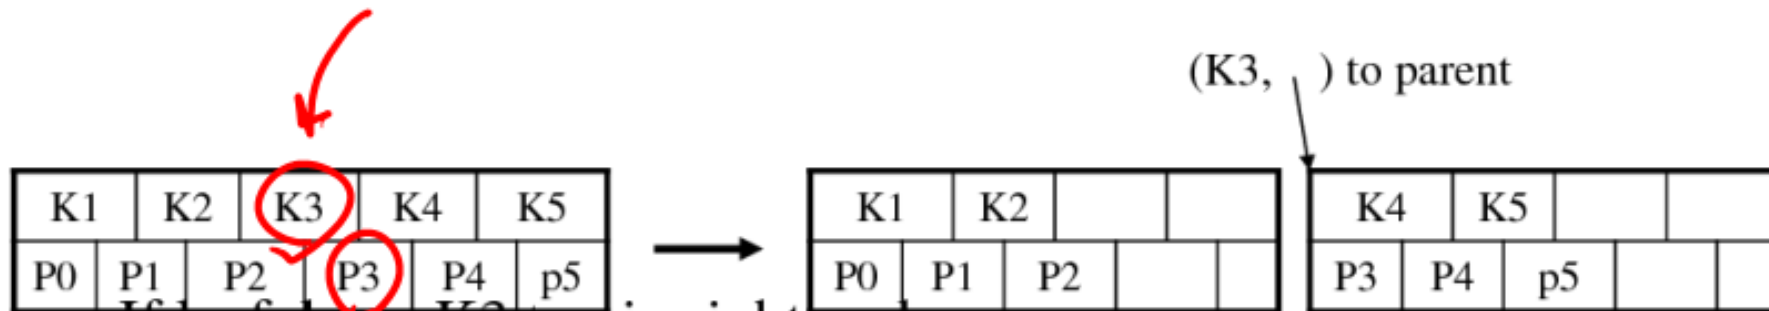


# Insertion in a B+ Tree

~~Deletion~~. [update (search)]

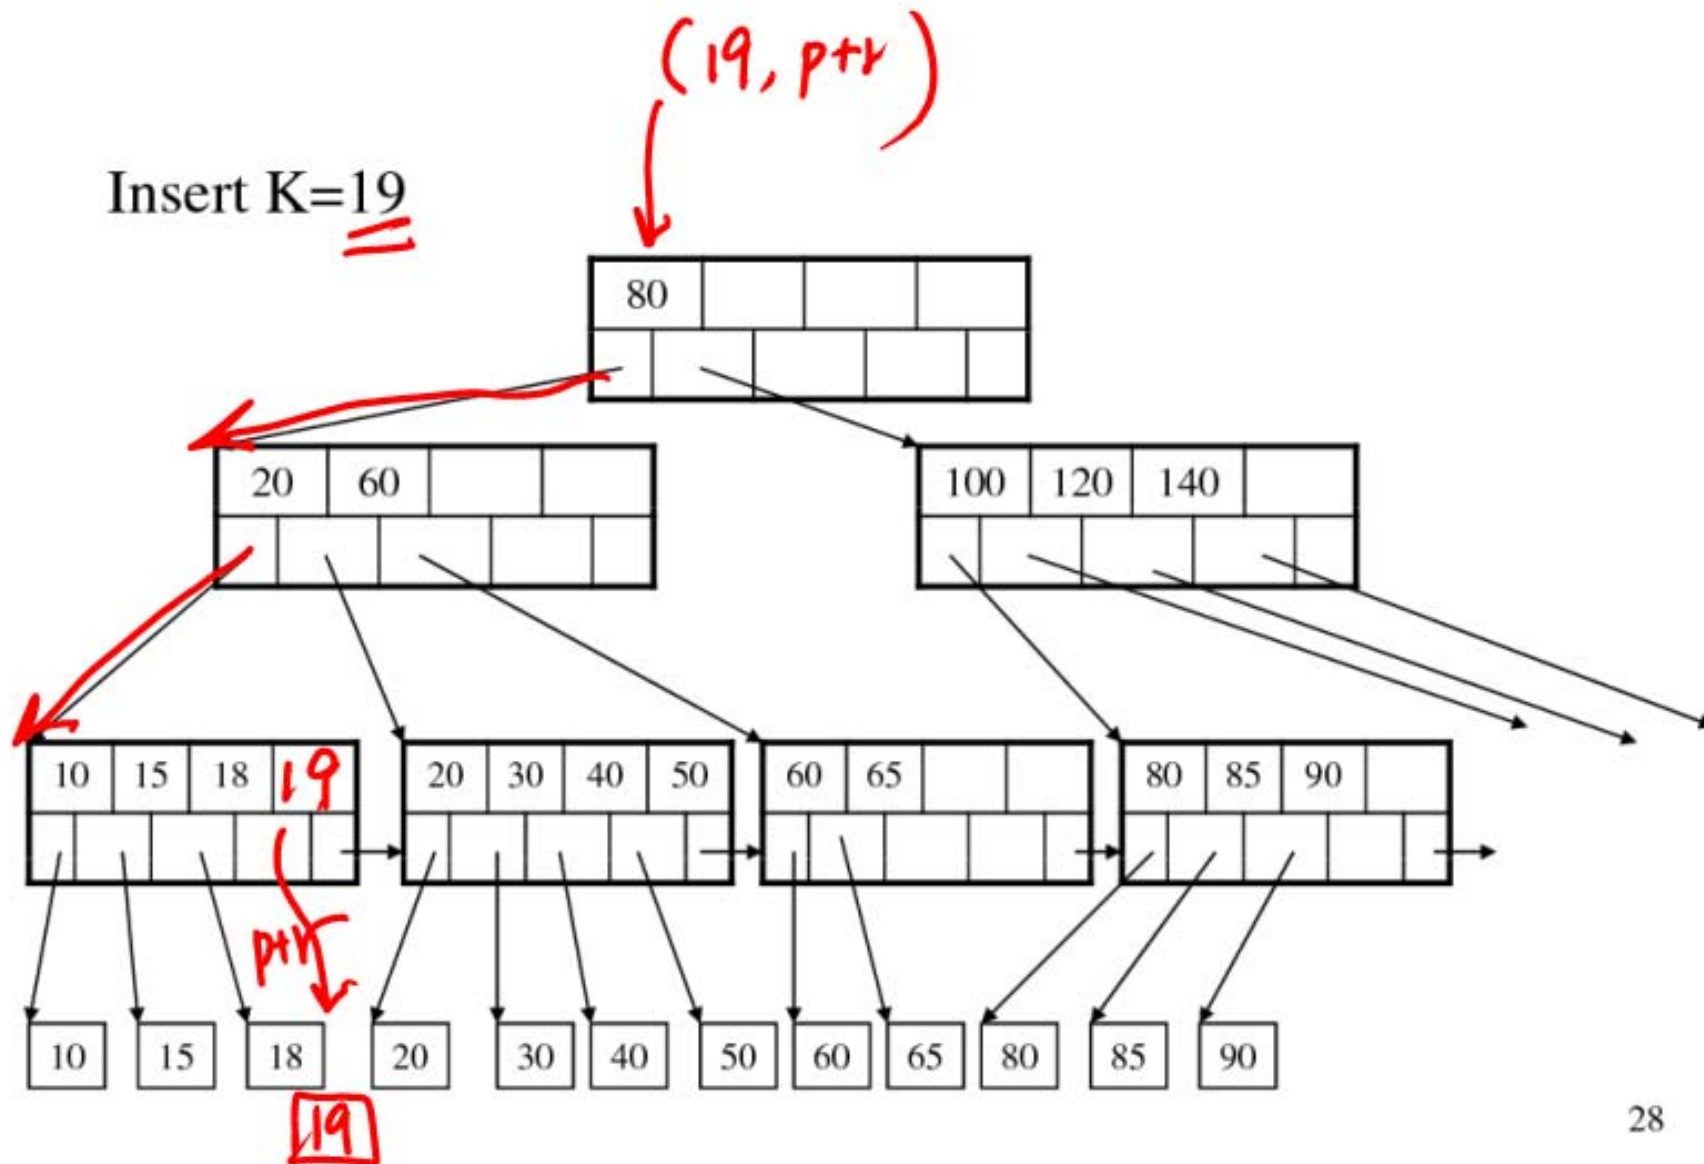
Insert (K, P)

- Find leaf where K belongs, insert
- If no overflow ( $2d$  keys or less), halt
- If overflow ( $2d+1$  keys), split node, insert in parent:



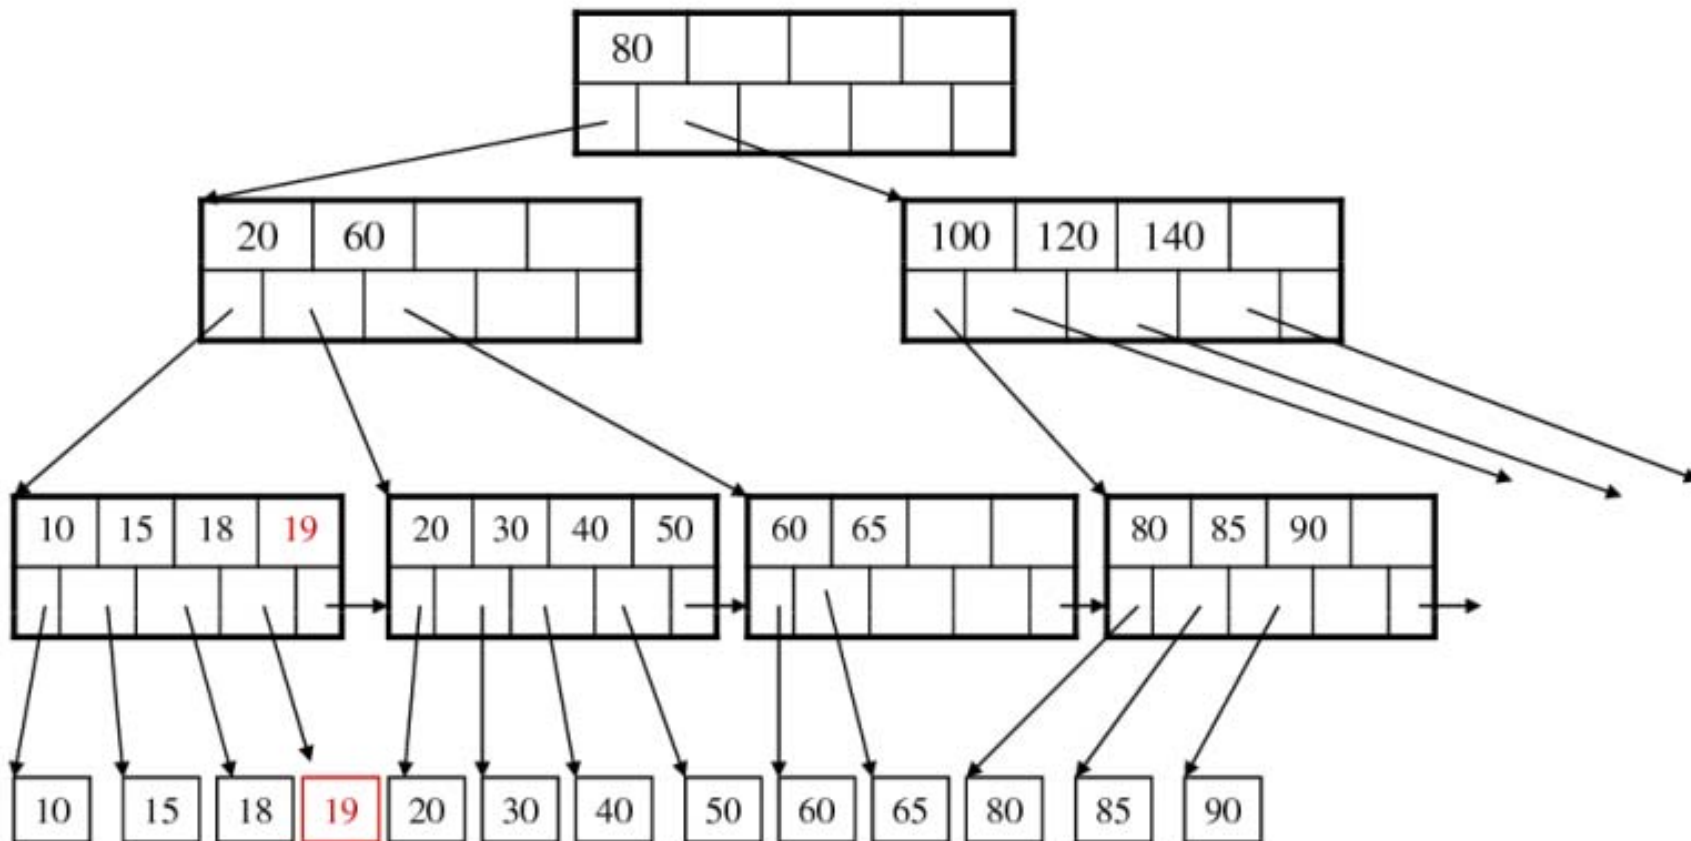
- If leaf, keep K3 too in right node
- When root splits, new root has 1 key only
  - that's why root is special for degree satisfaction

# Insertion in a B+ Tree



# Insertion in a B+ Tree

After insertion






Tree

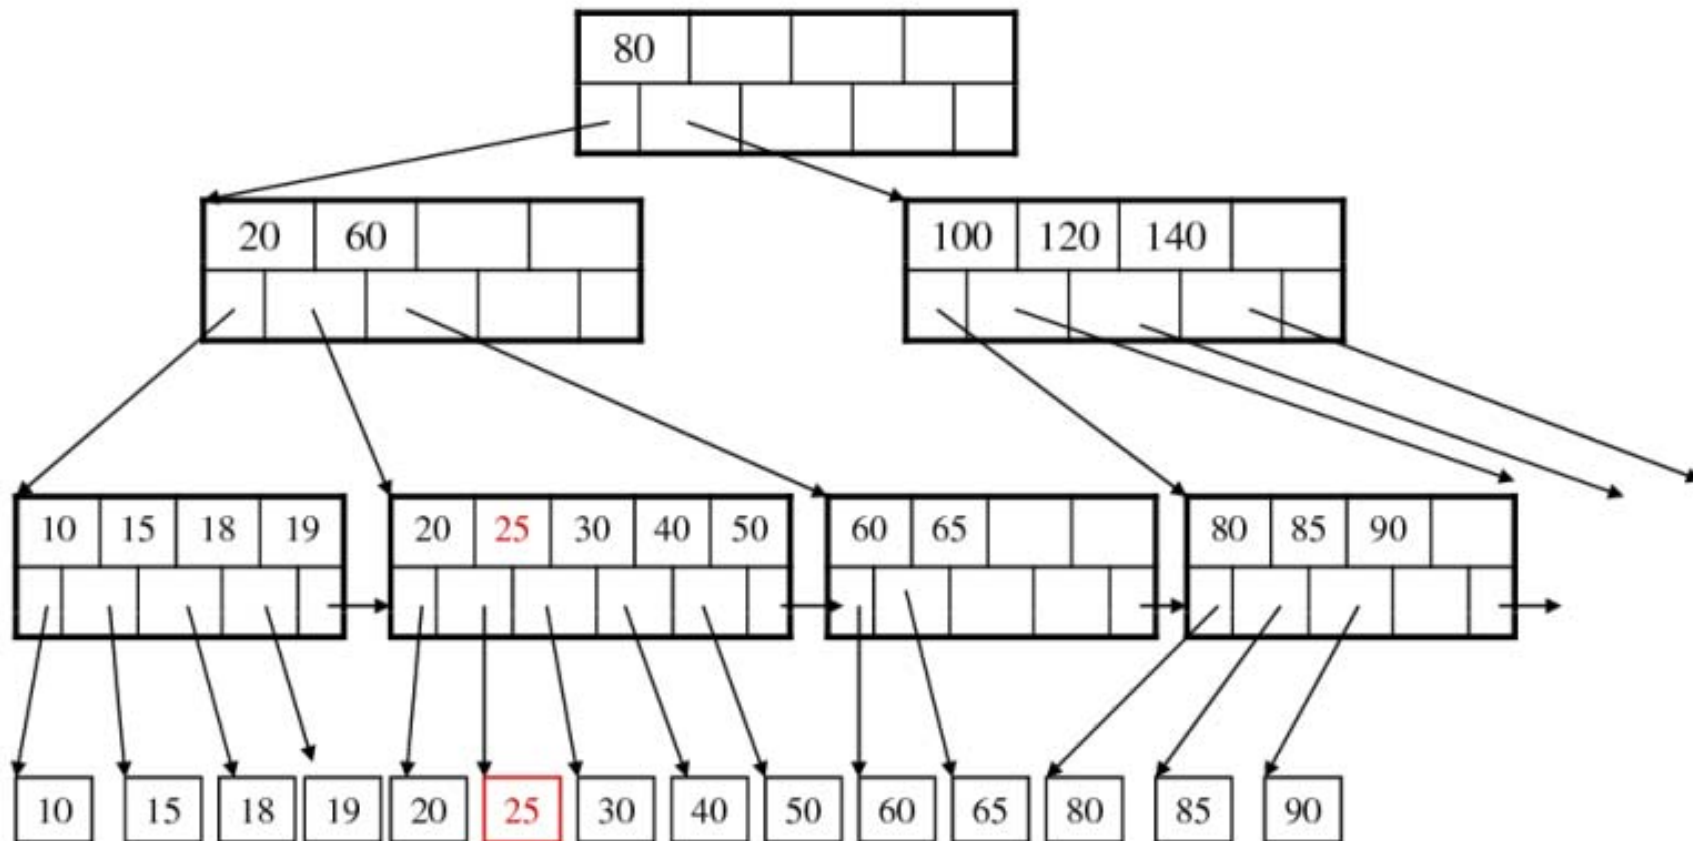
20	25	30	40	50
----	----	----	----	----

↓



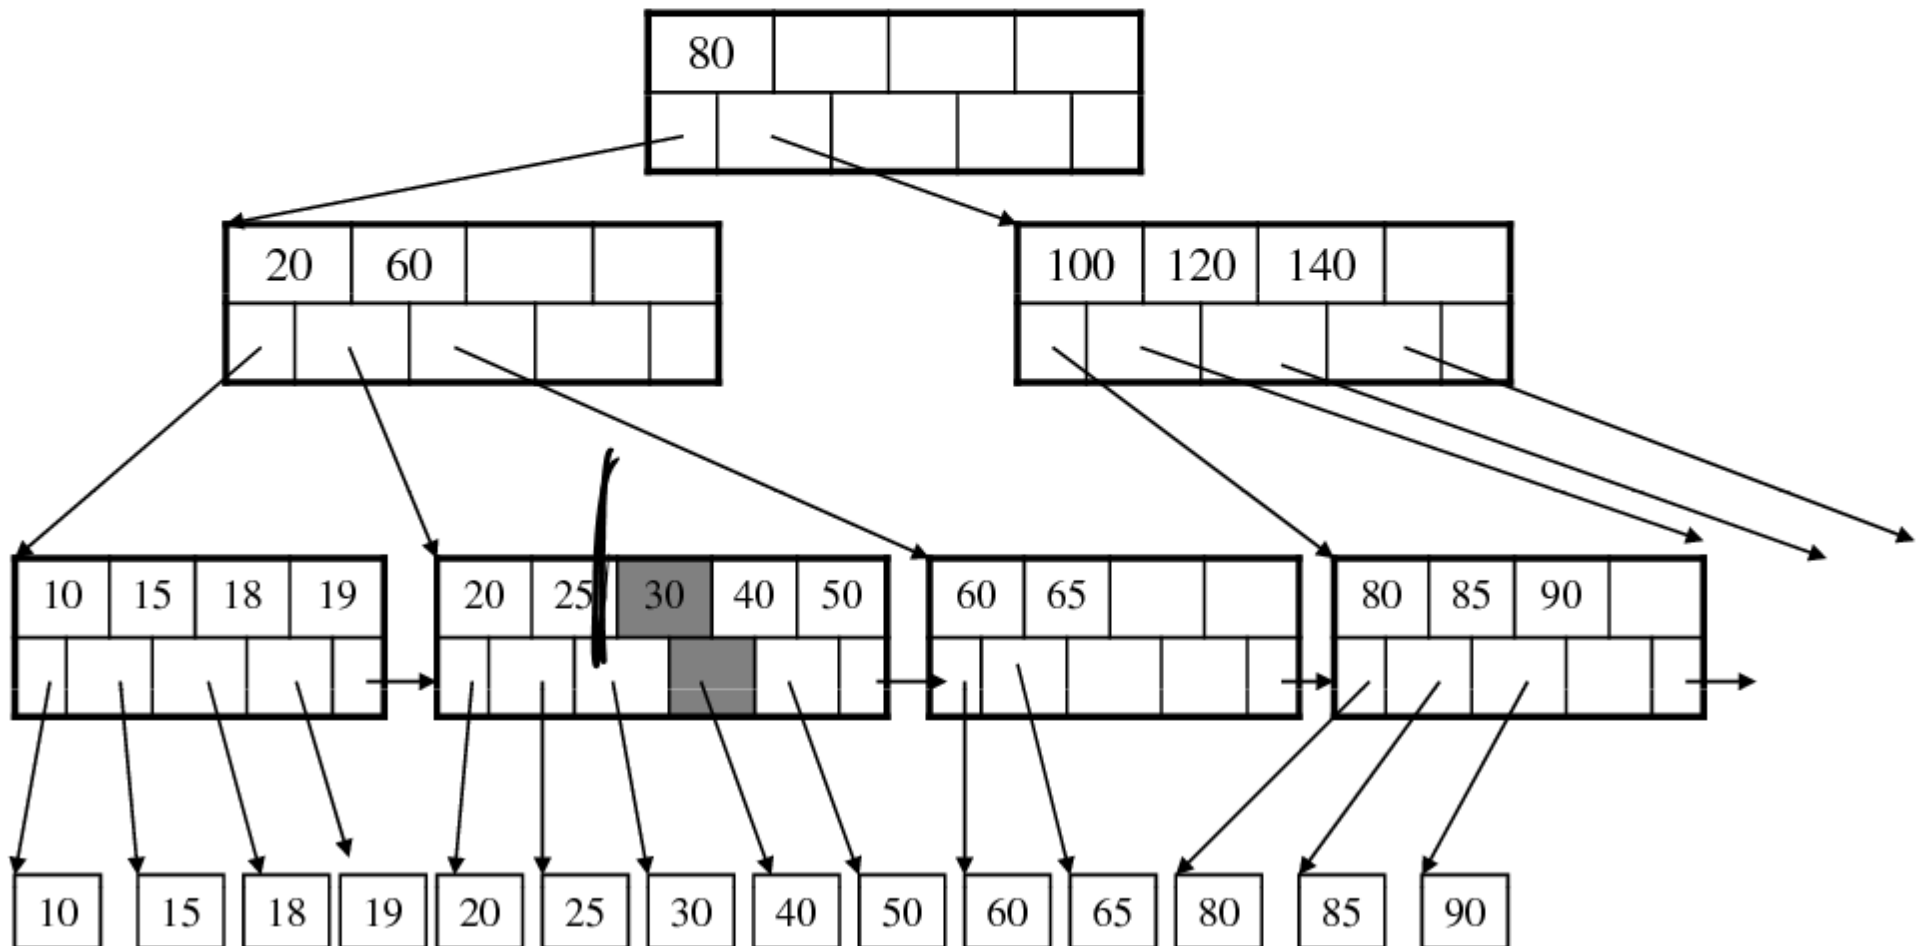
# Insertion in a B+ Tree

After insertion




# Insertion in a B+ Tree

But now have to split !



# Insertion in a B+ Tree


After the split <sup>②</sup> insert 30 at parent






# Deletion from a B+ Tree

Delete 30



# Deletion from a B+ Tree


After deleting 30



# Deletion from a B+ Tree

Rotate if u can

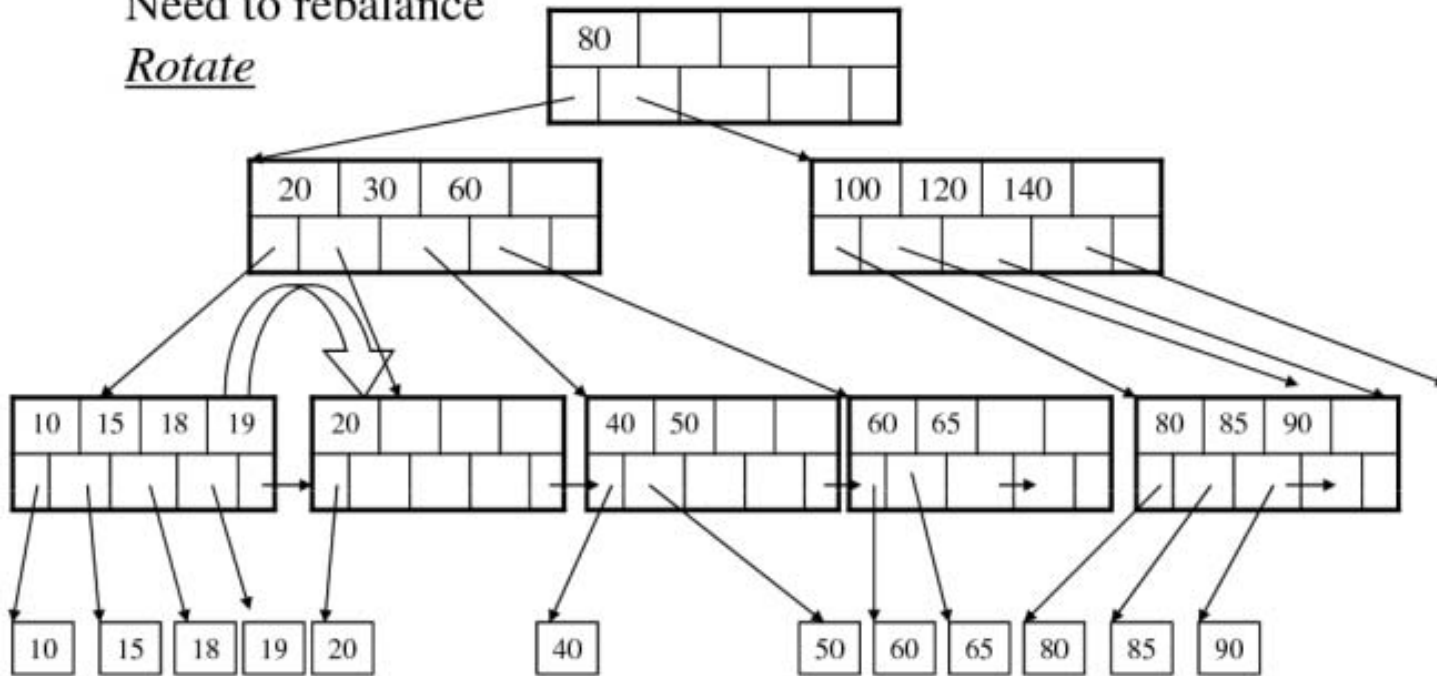
Now delete 25



# Deletion from a B+ Tree

After deleting 25  
Need to rebalance

Rotate






# Deletion from a B+ Tree

*Merge if you have to.*

19 | 20 | 50

Now delete 40

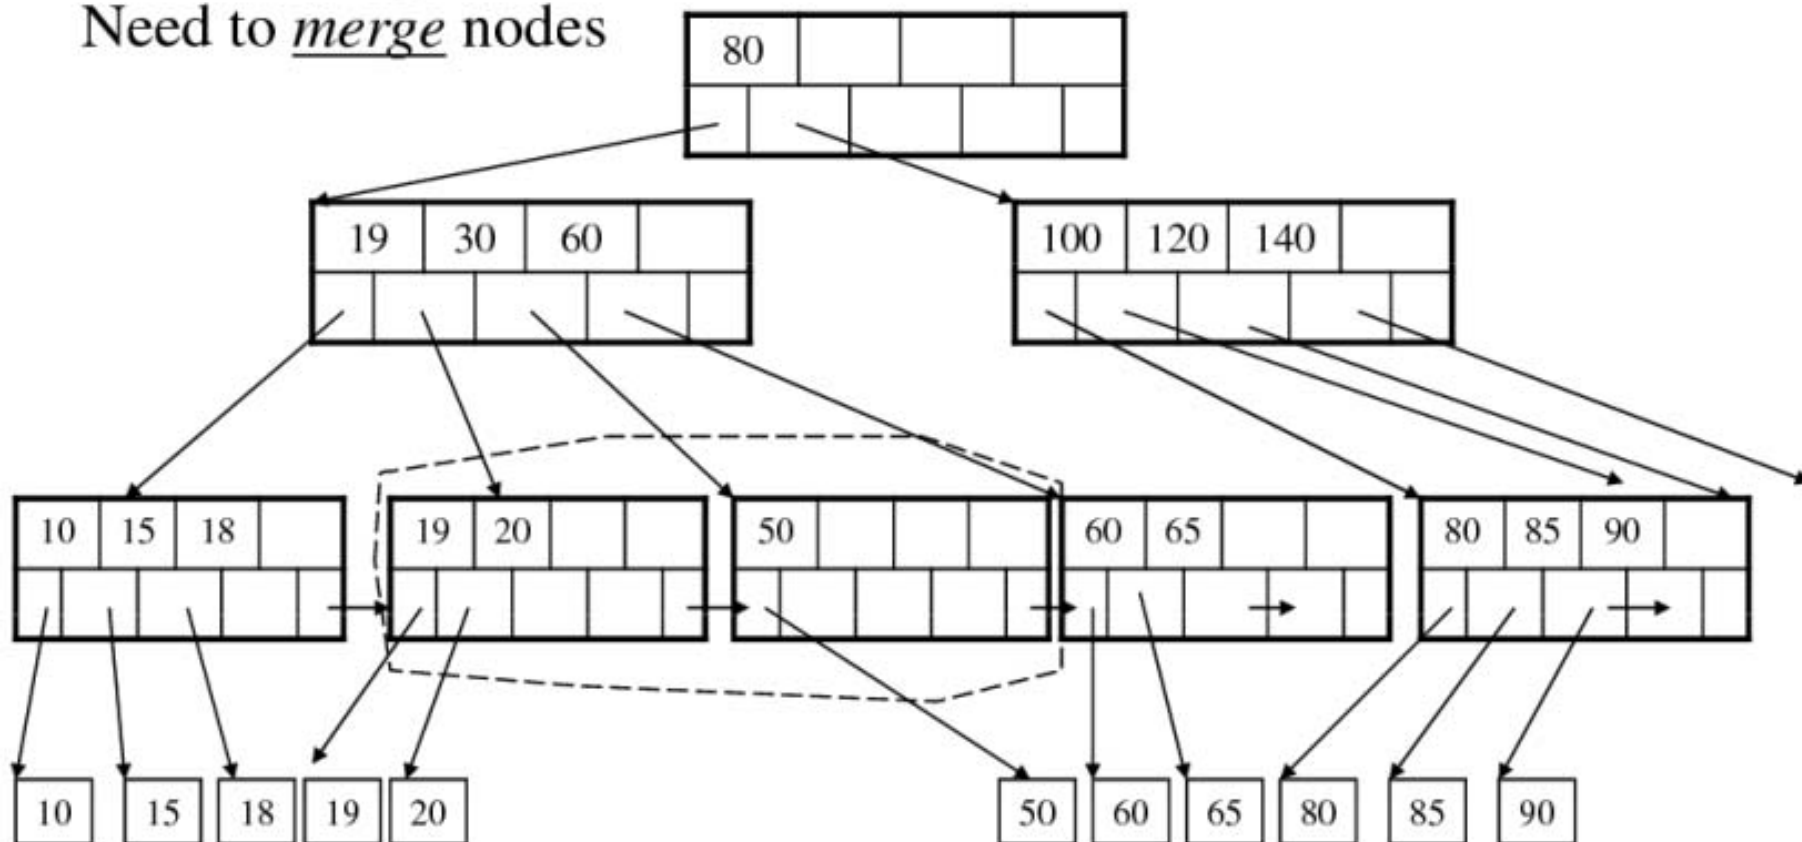


# Deletion from a B+ Tree

After deleting 40


Rotation not possible

Need to merge nodes



# Deletion from a B+ Tree


Final tree



## Variation on B+tree: B-tree (no +)

- Idea:
  - Avoid duplicate keys
  - Have record pointers in non-leaf nodes
- Note: Textbook's B-Tree means B+-tree!

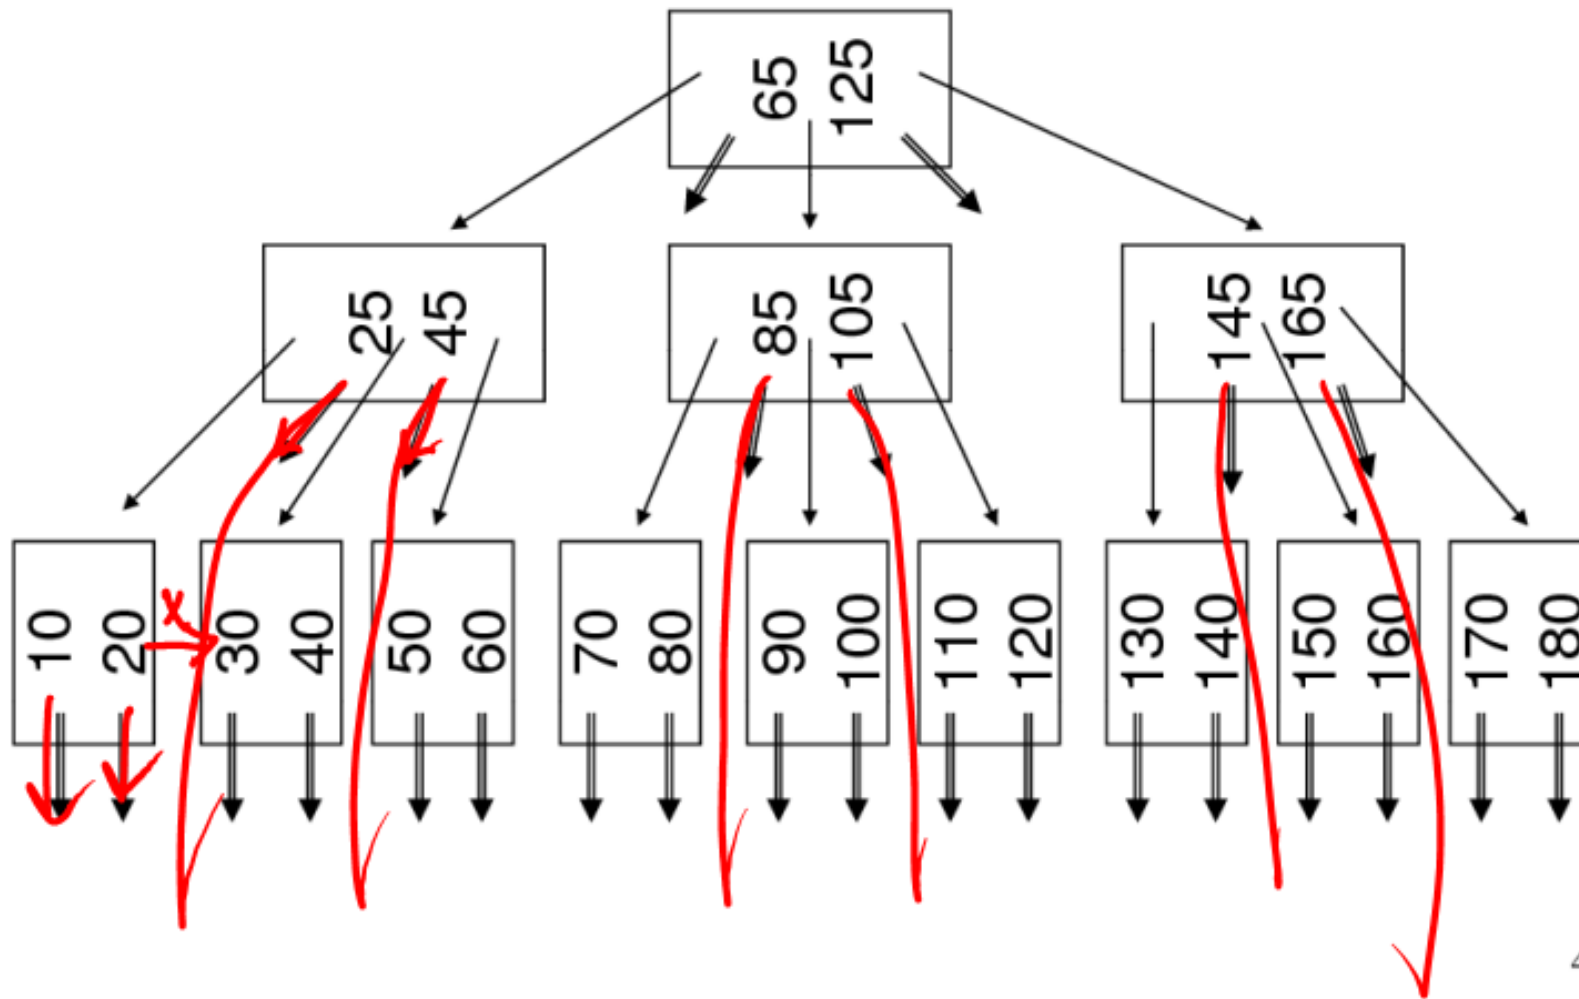




## B-tree example

B-Tree (X B+-tree)  
n=2

- Sequence pointers not useful now!



# Quiz

Name: \_\_\_\_\_

NETID: \_\_\_\_\_

Q. How do u like special topics?

SQLite : 1 - 3 - 5

SQL Tuing dislike 1 - - - - like 5

Q. Suggestions for Special Topics?

e.g. stop it,  
do it, but ...