

CS411
Database Systems

09: Query Optimization

answers
all the same
"performance"

$Q \rightsquigarrow P(Q)$

that
is optimal

Why this?

i) We really want to speed.

TPC measure perf of DBMS.
on Hardware
Why not "programmer"?
auto

— difficult
— vary every time \Rightarrow

programming

ii) Engineering Science

Optimization


- At the heart of the database engine
- Step 1: convert the SQL query to some logical plan (Rel Alg)
- Step 2: find a better logical plan, find an associated physical plan
→ concrete w/ all details

SQL \rightarrow Logical Query Plans

initial
to Logical

```
Select a1, ..., an
From R1, ..., Rk
Where C
```

```
Select a1, ..., an
From R1, ..., Rk
Where C
Group by b1, ..., bl
```

$$\prod_{a_1, \dots, a_n} \sigma_c(R_1 \times R_2 \times \dots \times R_k)$$

$$\Pi_{a1, \dots, an}(\gamma_{b1, \dots, bm, aggs}(\sigma_C(R1 \bowtie R2 \bowtie \dots \bowtie Rk)))$$

Skip

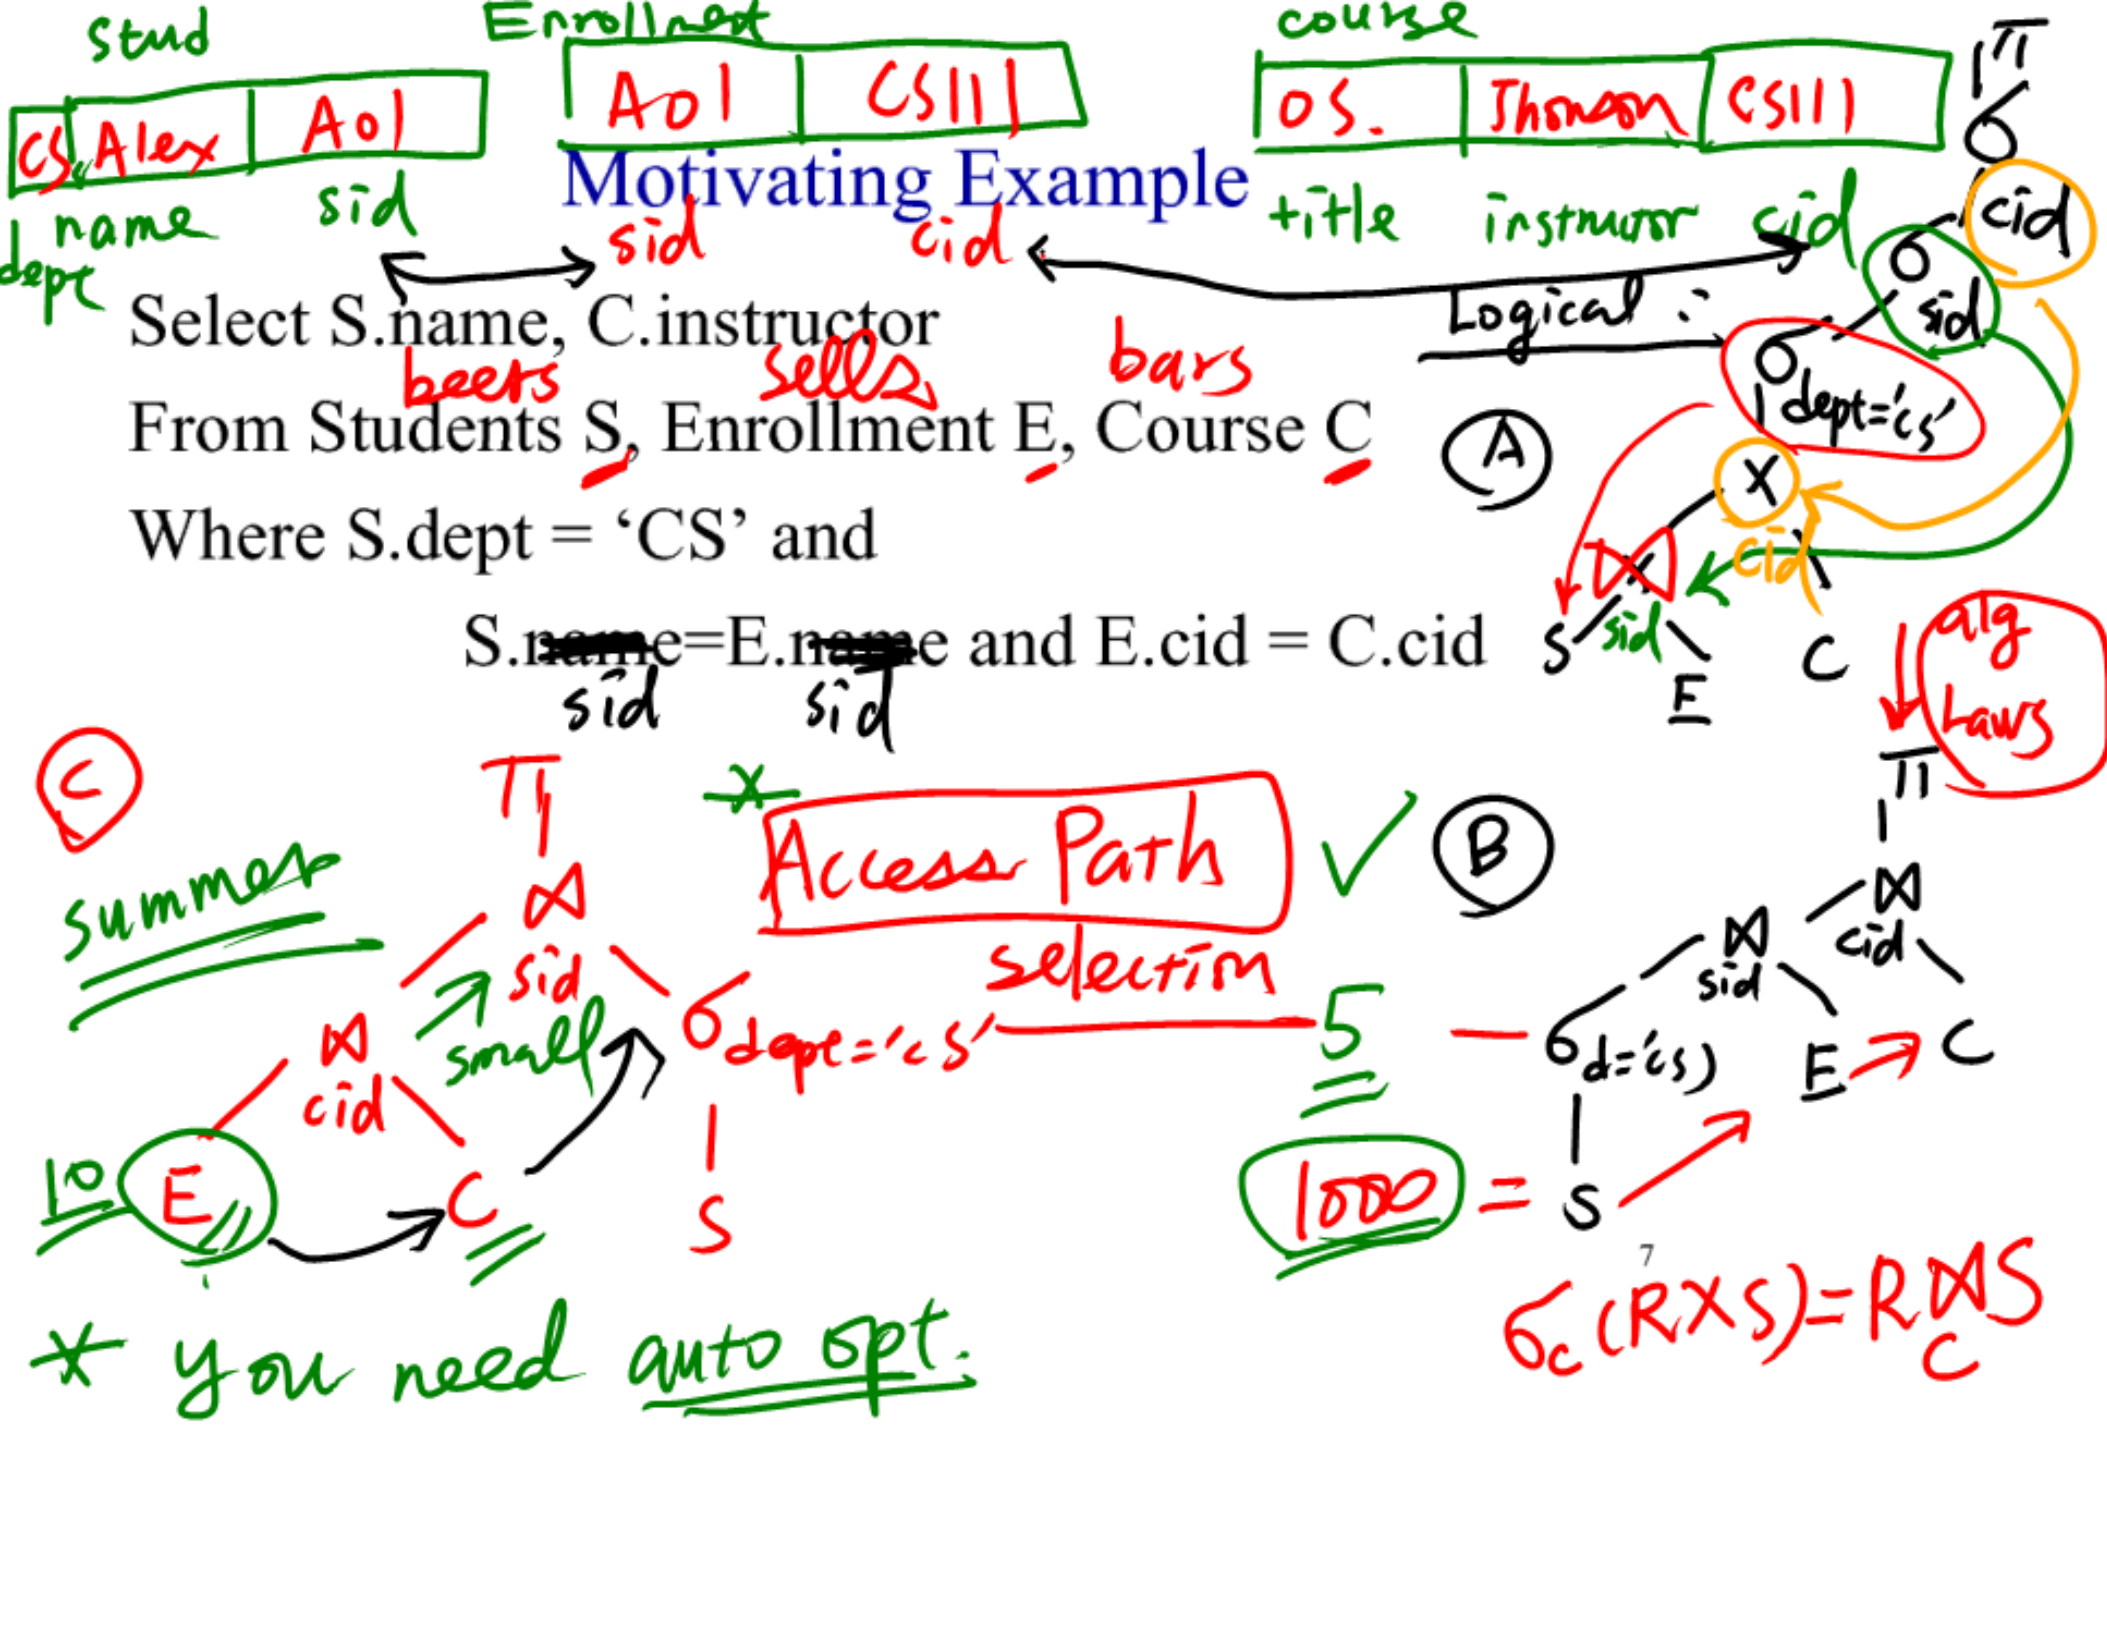
Converting Nested Queries

```
Select distinct product.name  
From product  
Where product.maker in (Select company.name  
                        From company  
                        where company.city="Urbana")
```

```
Select distinct product.name  
From product, company  
Where product.maker = company.name AND  
      company.city="Urbana"
```


Optimization: Logical Query Plan

- Now we have one logical plan
- Algebraic laws: *get to diff logical plan*
 - foundation for every optimization
- Two approaches to optimizations:
 - Rule-based (heuristics): apply laws that seem to result in cheaper plans
 - Cost-based: estimate size and cost of intermediate results, ~~search~~ systematically for best plan




The three components of an optimizer

We need three things in an optimizer:

- Algebraic laws
- An optimization algorithm
- A cost estimator

$P_1 \neq P_2$



Algebraic Laws

Algebraic Laws

- Commutative and Associative Laws

- $R \cup S = S \cup R$, $R \cup (S \cup T) = (R \cup S) \cup T$

- $R \cap S = S \cap R$, $R \cap (S \cap T) = (R \cap S) \cap T$


- $R \bowtie S = S \bowtie R$, $R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$

- Distributive Laws

- $R \bowtie (S \cup T) = (R \bowtie S) \cup (R \bowtie T)$

if $t \in LHS \Leftrightarrow t \in RHS$

Q: How to prove these laws? Make sense?



Algebraic Laws

• Laws involving selection: σ

$$- \sigma_{C \text{ AND } C'}(R) = \sigma_C(\sigma_{C'}(R)) = \sigma_C(R) \cap \sigma_{C'}(R)$$

$$- \sigma_{C \text{ OR } C'}(R) = \sigma_C(R) \cup \sigma_{C'}(R)$$

$$- \sigma_C(R \bowtie S) = \sigma_C(R) \bowtie S$$

- When C involves only attributes of R


$$- \sigma_C(R - S) = \sigma_C(R) - S$$

$$- \sigma_C(R \cup S) = \sigma_C(R) \cup \sigma_C(S)$$

$$- \sigma_C(R \cap S) = \sigma_C(R) \cap S$$

almost always
true

X Expensive prediction



Q: What do they mean? Make sense?

Algebraic Laws

- Example: $R(A, B, C, D)$, $S(E, F, G)$
 - $\sigma_{F=3}(R \bowtie_{D=E} S) \stackrel{?}{=} \sigma_{A=5 \text{ AND } G=9}(R \bowtie_{D=E} S) = ?$

Algebraic Laws

- Laws involving projections


- $$\Pi_M(R \bowtie S) = \Pi_N(\Pi_P(R) \bowtie \Pi_Q(S))$$

- Where N, P, Q are appropriate subsets of attributes of M

- $$\Pi_M(\Pi_N(R)) = \Pi_{M \cap N}(R)$$

- Example R(A,B,C,D), S(E,F,G)

- $$\Pi_{A,B,G}(R \bowtie_{D=E} S) = \Pi_?(\Pi_?(R) \bowtie \Pi_?(S))$$



Q: Again, what do they mean? Make sense?

Behind the Scene: Oracle RBO and CBO

TECHNOLOGY: Talking Tuning

As Published In

Understanding Optimization

By Kimberly Floss

ORACLE
MAGAZINE
January/February
2005

Improvements in the Oracle Database 10g Optimizer make it even more valuable for tuning.

Since its introduction in Oracle7, the cost-based optimizer (CBO) has become more valuable and relevant with each new release of the Oracle database while its counterpart, the rule-based optimizer (RBO), has become increasingly less so. The difference between the two optimizers is relatively clear: The CBO chooses the best path for your queries, based on what it knows about your data and by leveraging Oracle database features such as bitmap indexes, function-based indexes, hash joins, index-organized tables, and partitioning, whereas the RBO just follows established rules (heuristics). With the release of Oracle Database 10g, the RBO's obsolescence is official and the CBO has been significantly improved yet again.

- Oracle 7 (1992) prior (since 1979) RBO.
- Oracle 7-10: RBO + CBO.
- Oracle 10g (2003) CBO.

Behind the Scene: Oracle RBO and CBO

Rule-based optimization sometimes provided better performance than the early versions of Oracle's cost-based optimizer for specific situations. The rule-based optimizer had several weaknesses, including offering only a simplistic set of rules. The Oracle rule-based optimizer had about 20 rules and assigned a weight to each one of them. In a complex database, a query can easily involve several tables, each with several indexes and complex selection conditions and ordering. This complexity means that there were a lot of options, and the simple set of rules used by the rule-based optimizer might not differentiate the choices well enough to make the best choice.



The rule-based optimizer assigned an optimization score to each potential execution path and then took the path with the best optimization score. Another weakness in the rule-based optimizer was resolution of optimization choices made in the event of a "tie" score. When two paths presented the same optimization score, the rule-based optimizer looked to the syntax of the SQL statement to resolve the tie. The winning execution path was based on the order in which the tables occur in the SQL statement.

Rule-based Optimization

Heuristic Based Optimizations

- Query rewriting based on algebraic laws
- Result in better queries most of the time
- Heuristics number 1:
 - Push selections down
- Heuristics number 2:
 - Sometimes push selections up, then down

Predicate Pushdown



The earlier we process selections, less tuples we need to manipulate higher up in the tree (but may cause us to lose an important ordering of the tuples, if we use indexes).

Predicate Pushdown ~~min~~ X

if a group has max(price) > 100 \Rightarrow then it must come from some prod.

```

Select y.name, Max(x.price)
From   product x, company y
Where  x.maker = y.name
GroupBy y.name
Having Max(x.price) > 100

```

each company

min

```

Select y.name, Max(x.price)
From   product x, company y
Where  x.maker=y.name and
       x.price > 100
GroupBy y.name
Having Max(x.price) > 100

```

price
v
100

- For each company, find the maximal price of its products.
- Advantage: the size of the join will be smaller.
- Requires transformation rules specific to the grouping/aggregation operators.
- **Won't work if we replace Max by Min.**

Ex: push selection into views
(subqueries)

Pushing predicates up

Bargain view V1: categories with some price < 20, and the cheapest price

```
Select V2.name, V2.price
From   V1, V2
Where  V1.category = V2.category and
       V1.p = V2.price
```

```
Create View V1 AS
Select  x.category,
        Min(x.price) AS p
From    product x
Where   x.price < 20
GroupBy x.category
```

```
Create View V2 AS
Select  y.name, x.category, x.price
From    product x, company y
Where   x.maker=y.name
```


Query Rewrite: Pushing predicates up

Bargain view V1: categories with some price<20, and the cheapest price

```
Select  V2.name, V2.price
From    V1, V2
Where   V1.category = V2.category and
        V1.p = V2.price AND V1.p < 20
```


```
Create View V1 AS
Select  x.category,
        Min(x.price) AS p
From    product x
Where   x.price < 20
GroupBy x.category
```

```
Create View V2 AS
Select  y.name, x.category, x.price
From    product x, company y
Where   x.maker=y.name
```


Query Rewrite: Pushing predicates up

Bargain view V1: categories with some price<20, and the cheapest price

```
Select  V2.name, V2.price  
From    V1, V2  
Where   V1.category = V2.category and  
        V1.p = V2.price AND V1.p < 20
```



```
Create View V1 AS  
Select  x.category,  
        Min(x.price) AS p  
From    product x  
Where   x.price < 20  
GroupBy x.category
```

```
Create View V2 AS  
Select  y.name, x.category, x.price  
From    product x, company y  
Where   x.maker=y.name  
        AND x.price < 20
```


Rule-based Optimization



✓ Cost-based Optimization

Behind the Scene: The Selinger Style!

Patricia Selinger

from Wikipedia, the free encyclopedia

Patricia Selinger is an American computer scientist and IBM Fellow, best known for her work on relational database management systems. She played a fundamental role in the development of System R, a pioneering relational database implementation, and wrote the canonical paper on relational query optimization.^[1] The dynamic programming algorithm for determining join order proposed in that paper still forms the basis for most of the query optimizers used in modern relational systems.

She was made an IBM Fellow in 1994, was elected to the National Academy of Engineering in 1999, and won the SIGMOD Edgar F. Codd Innovations Award in 2002. Before her retirement, she was the Vice President of Data Management Architecture and Technology at IBM. She received A.B., S.M., and Ph.D. degrees in applied mathematics from Harvard University.

References

- ¹ Selinger, P. G.; Astrahan, M. M.; Chamberlin, D. D.; Lorie, R. A.; Price, T. (1979). "Access Path Selection in a Relational Database Management System", *Proceedings of the 1979 ACM SIGMOD International Conference on Management of Data*, pp. 23-34, doi:10.1145/582095.582099

ABSTRACT: In a high level query and data manipulation language such as SQL, requests are stated non-procedurally, without reference to access paths. This paper describes how System R chooses access paths for both simple (single relation) and complex queries (such as joins), given a user specification of desired data as a boolean expression of predicates. System R is an experimental database management system developed to carry out research on the relational model of data. System R was designed and built by members of the IBM San Jose Research Laboratory.



Behind the Scene: The Selinger Style!

In my view, the query optimizer was the first attempt at what we call autonomic computing or self-managing, self-tuning technology. Query optimizers have been 25 years in development, with enhancements of the cost-based query model and the optimization that goes with it, and a richer and richer variety of execution techniques that the optimizer chooses from. We just have to keep working on this. It's a never-ending quest for an increasingly better model and repertoire of optimization and execution techniques. So the more the model can predict what's really happening in the data and how the data is really organized, the closer and closer we will come [to the ideal system].

In hindsight, I didn't really have enough experience as a system designer and developer to realize that System R was going to be something that people [within IBM] were going to take almost as is, and make a productized copy of it and [sell] it. So designing control blocks, designing interfaces and APIs to programming languages---we just sort of put together something that we thought would work. The error control block, the way that you pass variables and parameters into the system---those were not designed with thought and care and elegance. They were [just]

Cost-based Optimizations



- Main idea: apply algebraic laws, until estimated cost is minimal *search \Rightarrow construction*
- Practically: start from partial plans, introduce operators one by one
 - Will see in a few slides
- Problem: there are too many ways to apply the laws, hence too many (partial) plans

Cost-based Optimizations

Approaches:

- **Top-down:** the partial plan is a top fragment of the logical plan

- **Bottom up:** the partial plan is a bottom fragment of the logical plan




Search Strategies

- **Branch-and-bound:**
 - Remember the cheapest complete plan P seen so far and its cost C
 - Stop generating partial plans whose cost is $> C$
 - If a cheaper complete plan is found, replace P, C
- **Hill climbing:**
 - Remember only the cheapest partial plan seen so far
- **Dynamic programming:** bottom-up
 - Remember the all cheapest partial plans

Dynamic Programming

Unit of Optimization: select-project-join




$R_1 \bowtie R_2 \bowtie R_3 \bowtie R_4$

Join Trees

Assumption :=

join is
plan binary op

- $R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$
- Join tree:




- A plan = a join tree
- A partial plan = a subtree of a join tree

Types of Join Trees

- Left deep:

iterator



Assump :

only care about
left - deep


benefits :

- ① Easily pipelined
- ② mem req small

Types of Join Trees

- Bushy:


x pipeline not
easy



x mem req
high

Types of Join Trees

- Right deep: \leftrightarrow *Left deep*



Problem

- Given: a query $R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$
- Assume we have a function $\text{cost}()$ that gives us the cost of every join tree
- Find the best join tree for the query

Dynamic Programming

- Idea: for each subset of $\{R_1, \dots, R_n\}$, compute the best plan for that subset
- In increasing order of set cardinality:
 - Step 1: for $\{R_1\}, \{R_2\}, \dots, \{R_n\}$
 - Step 2: for $\{R_1, R_2\}, \{R_1, R_3\}, \dots, \{R_{n-1}, R_n\}$
 - ...
 - Step n: for $\{R_1, \dots, R_n\}$
- It is a bottom-up strategy
- A subset of $\{R_1, \dots, R_n\}$ is also called a *subquery*

Dynamic Programming

- For each subquery $Q \subseteq \{R_1, \dots, R_n\}$ compute the following:
 - $\text{Size}(Q)$
 - A best plan for Q : $\text{Plan}(Q)$
 - The cost of that plan: $\text{Cost}(Q)$

Dynamic Programming

- **Step 1:** For each $\{R_i\}$ do:
 - $\text{Size}(\{R_i\}) = B(R_i)$
 - $\text{Plan}(\{R_i\}) = R_i$
 - $\text{Cost}(\{R_i\}) = (\text{cost of scanning } R_i)$

Dynamic Programming

- **Step i:** For each $Q \subseteq \{R_1, \dots, R_n\}$ of cardinality i do:
 - Compute $\text{Size}(Q)$ (later...)
 - For every pair of subqueries Q', Q''
s.t. $Q = Q' \cup Q''$
compute $\text{cost}(\text{Plan}(Q') \bowtie \text{Plan}(Q''))$
 - $\text{Cost}(Q)$ = the smallest such cost
 - $\text{Plan}(Q)$ = the corresponding plan

Dynamic Programming

- Return $\text{Plan}(\{R_1, \dots, R_n\})$

Dynamic Programming

To illustrate, we will make the following simplifications:

- $\text{Cost}(P1 \bowtie P2) = \text{Cost}(P1) + \text{Cost}(P2) + \text{size}(\text{intermediate result})$
- Intermediate results:
 - If $P1$ = a join, then the size of the intermediate result is $\text{size}(P1)$, otherwise the size is 0
 - Similarly for $P2$
- Cost of a scan = 0

Dynamic Programming

- Example:
- $\text{Cost}(R5 \bowtie R7) = 0$ (no intermediate results)
- $\text{Cost}((R2 \bowtie R1) \bowtie R7)$
 $= \text{Cost}(R2 \bowtie R1) + \text{Cost}(R7) + \text{size}(R2 \bowtie R1)$
 $= \text{size}(R2 \bowtie R1)$

Dynamic Programming

- Relations: R, S, T, U
- Number of tuples: 2000, 5000, 3000, 1000
- Size estimation: $T(A \bowtie B) = 0.01 * T(A) * T(B)$