

Dynamic Programming

- Idea: for each subset of $\{R1, \dots, Rn\}$, compute the best plan for that subset

Drinker

Beers \bowtie Sells \bowtie Bar

- In increasing order of set cardinality:

$k=1$ – Step 1: for $\{R1\}, \{R2\}, \dots, \{Rn\}$ $\{Beers\}, \{Sells\}, \{Drinker\}$

$k=2$ – Step 2: for $\{R1, R2\}, \{R1, R3\}, \dots, \{Rn-1, Rn\}$

...

$k=n$ – Step n: for $\{R1, \dots, Rn\}$ $B \bowtie S \bowtie D$ $\{B \bowtie S\}, \{S \bowtie D\}$

- It is a bottom-up strategy
- A subset of $\{R1, \dots, Rn\}$ is also called a *subquery*


Dynamic Programming

- For each subquery $Q \subseteq \{R_1, \dots, R_n\}$ compute the following:

- ① Size(Q) = ?
- ② A best plan for Q: Plan(Q)
- ③ The cost of that plan: Cost(Q)

? for subsequent query operations

R_1 . R_2 . R_3 . R_4
 1000 2000 3000 4000



Dynamic Programming

- **Step 1:** For each $\{R_i\}$ do:
 - $\text{Size}(\{R_i\}) = \underline{B(R_i)}$ *# blocks.*
 - $\text{Plan}(\{R_i\}) = \underline{R_i}$
 - $\text{Cost}(\{R_i\}) = (\text{cost of scanning } R_i)$

$R_1 \bowtie R_2 \bowtie R_3 \bowtie R_4 \bowtie R_5$

Dynamic Programming

$i=4$

- **Step i:** For each $Q \subseteq \{R_1, \dots, R_n\}$ of cardinality i do:

- Compute Size(Q) (later...)
- For every pair of subqueries Q', Q'' s.t. $Q = Q' \cup Q''$ compute $\text{cost}(\text{Plan}(Q') \bowtie \text{Plan}(Q''))$
- $\text{Cost}(Q)$ = the smallest such cost
- $\text{Plan}(Q)$ = the corresponding plan

join is binary,

$R_1 \cdot R_2 \cdot R_3 \cdot R_4$

R_1

$R_2 R_3 R_4$

best from $i=3$

R_2

$R_1 R_3 R_4$

$R_1 R_2$

$R_3 R_4$

best for $i=4$

$i = \underline{n}$

Dynamic Programming

- Return Plan($\{R1, \dots, Rn\}$)

Dynamic Programming

this is why we need to remember cost of subq.

To illustrate, we will make the following simplifications:

- $\text{Cost}(\underbrace{P_1}_{R_1} \bowtie \underbrace{P_2}_{R_2 R_3 R_4}) = \underbrace{\text{Cost}(P_1)}_{\text{size(intermediate result)}} + \underbrace{\text{Cost}(P_2)}$
- Intermediate results:
 - If P_1 = a join, then the size of the intermediate result is $\text{size}(P_1)$, otherwise the size is 0
 - Similarly for P_2
- Cost of a scan = 0

$\text{size}(P_1) + \text{size}(P_2)$ e.g. $P_2 = R_2 R_3 R_4$

$P_1 = R_1$ X

if P_1, P_2 are joins.

Dynamic Programming

- Example:
- $\text{Cost}(R_1 \bowtie R_2) = 0$ (no intermediate results)
- $\text{Cost}((R_2 \bowtie R_1) \bowtie R_7)$
 $= \text{Cost}(R_2 \bowtie R_1) + \text{Cost}(R_7) + \text{size}(R_2 \bowtie R_1)$
 $= \text{size}(R_2 \bowtie R_1)$

$$\textcircled{1} \text{ cost}(R_1 \bowtie R_2) = \text{cost}(R_1) + \text{cost}(R_2) + \text{size}(R_1) = 0 + 0 + \text{size}(R_2) = 0$$

$$= 0$$

$$\textcircled{2} \text{ cost}(\overset{P_1}{(R_1 \bowtie R_2)} \overset{P_2}{\bowtie} R_3) = C(\overset{0}{(R_1 \bowtie R_2)}) + \overset{0}{C(R_3)} + \underbrace{S(R_1 \bowtie R_2)} + \cancel{S(R_3)} = 0$$

$$= \underbrace{S(R_1 \bowtie R_2)}$$

$$\textcircled{3} C((R_1 \bowtie R_2) \bowtie (R_3 \bowtie R_4)) = S(R_1 \bowtie R_2) + S(R_3 \bowtie R_4)$$

$$\textcircled{4} C((R_1 \bowtie R_2 \bowtie R_3) \bowtie R_4) = C(R_1 \bowtie R_2 \bowtie R_3) + S(R_1 \bowtie R_2 \bowtie R_3)$$

Dynamic Programming

- Relations: R, S, T, U
- Number of tuples: 2000, 5000, 3000, 1000
- Size estimation: $T(A \bowtie B) = 0.01 * T(A) * T(B)$

A: 1000 tuples

B: 50 tuples

$$A \bowtie B = ? \quad \underline{0.01} \times 1000 \times 50 = \underline{\underline{500}}$$

R ⋈ S ⋈ T ⋈ U ? ① ③ ②

$i=2$ rel
 $C_2^4 = \frac{4 \cdot 3}{2 \cdot 1} = 6$


$i=3$
 $C_3^4 = 4$

$i=4$
 $C_4^4 = 1$

Subquery	<u>Size</u>	<u>Cost</u>	<u>Plan</u>
RS 1	$1 \cdot 1 \times R \times S = 100k$	0	$R \text{ --- } S$
RT 2			
RU 3			
ST 4			
SU 5			
TU 6			
RST			
RSU			
RTU			
STU			
RSTU			

$(RT) \bowtie (SU)$ = 43


Bushy



RSNT $p=?$

How to assemble?

then \bowtie



size:
 $100k \times 3k \times 0.01 = \underline{\underline{3M}}$

cost
 $C(RS \bowtie T) = s(RS) = \underline{\underline{100k}}$

Subquery	Size	Cost	Plan
RS	<u>100k</u>	0	<u>RS</u>
RT	60k	0	RT
RU	20k	0	RU
ST	150k	0	ST
SU	50k	0	SU
TU	30k	0	TU
<u>RST</u>	<u>3M</u>	60k	<u>(RT)S</u>
<u>RSU</u>	1M	20k	(RU)S
<u>RTU</u>	0.6M	20k	(RU)T
<u>STU</u>	1.5M	30k	(TU)S
RSTU	30M	60k+50k=110k	(RT)(SU)

$i=2$

$i=3$

$C_3^4 = 4$

$i=4$

RSTU

$\bar{v}=4$ (FINALLY!) RSTU

$$4 = \begin{matrix} \times 0+4 \\ 1+3 \\ \hline 2+2 \\ \hline 3+1 \\ \times 4+0 \end{matrix}$$

$$\begin{matrix} 1+3 \\ \hline 2+2 \\ \hline 3+1 \\ \times 4+0 \end{matrix}$$

our cost model is insensitive to le/right

L.D.

$$\times 3+1$$

$$\times 4+0$$

$\bar{v}=3!$

$$\frac{1+3}{2}$$


$$C^4 = 4 = 4$$



$$\frac{C^4}{2} = 2$$

\Rightarrow

Bushy



Dynamic Programming

- Summary: computes optimal plans for subqueries:
 - Step 1: $\{R_1\}, \{R_2\}, \dots, \{R_n\}$
 - Step 2: $\{R_1, R_2\}, \{R_1, R_3\}, \dots, \{R_{n-1}, R_n\}$
 - ...
 - Step n: $\{R_1, \dots, R_n\}$
- We used naïve size/cost estimations
- In practice:
 - more realistic size/cost estimations (next time)
 - heuristics for Reducing the Search Space
 - Restrict to left linear trees
 - Restrict to trees “without cartesian product”: $R(A,B), S(B,C), T(C,D)$
 $(R \text{ join } T) \text{ join } S$ has a cartesian product

Completing Physical Query Plan

How can we get order for free?

Completing the Physical Query Plan



- Choose algorithm to implement each operator

– Need to account for more than cost:

- How much memory do we have?
- Are the input operand(s) sorted?

Hash join

- Decide for each intermediate result:

To materialize

To pipeline

Join index join

S.m. join

req: S.g has index


mem: 1 block (R) X

Cost: $\underline{B(R)} + \underline{T(R)} \times \frac{B(S)}{v(S,a)}$

$$B(R) + B(S) \leq M^2$$

$$3B(R) + 3B(S)$$

Materialize Intermediate Results Between Operators



```

HashTable  $\leftarrow S$ 
repeat
  read(R, x)
  y  $\leftarrow$  join(HashTable, x)
  write(V1, y)
  
```

```

HashTable  $\leftarrow T$ 
repeat
  read(V1, y)
  z  $\leftarrow$  join(HashTable, y)
  write(V2, z)
  
```

```


HashTable  $\leftarrow U$ 
repeat
  read(V2, z)
  u  $\leftarrow$  join(HashTable, z)
  write(Answer, u)
  
```


Materialize Intermediate Results Between Operators

Given $B(R)$, $B(S)$, $B(T)$, $B(U)$

- What is the total cost of the plan ? $\approx B(R \bowtie S)$
 - Cost = $B(R) + B(S) + B(T) + B(U) + \dots \approx B(R \bowtie S \bowtie T)$
- How much main memory do we need ?
 - $M = \max(B(S), B(T), B(U))$

Pipeline Between Operators



How much main memory do we need? $M =$ 50

\downarrow Cost $= B(R) + B(S) + B(T) + B(u) +$
 \uparrow Mem $= \text{Sum of } B(S), B(T), B(u)$

~~X~~
~~X~~ z B()
~~X~~ z B()

Pipeline Between Operators

Given $B(R)$, $B(S)$, $B(T)$, $B(U)$

- What is the total cost of the plan ?
 - Cost =
- How much main memory do we need ?
 - M =


Completing the Physical Query Plan

- Choose algorithm to implement each operator
 - Need to account for more than cost:
 - How much memory do we have ?
 - Are the input operand(s) sorted ?
- Decide for each intermediate result:
 - To materialize
 - To pipeline

skip


Example

- Logical plan is:



- Main memory $M = 101$ buffers


Example



Naïve evaluation:

- 2 partitioned hash-joins
- Cost $3B(R) + 3B(S) + 4k + 3B(U) = 75000 + 4k$


Example



Smarter:

- Step 1: hash R on x into 100 buckets, each of 50 blocks; to disk
- Step 2: hash S on x into 100 buckets; to disk
- Step 3: read each R_i in memory (50 buffer) join with S_i (1 buffer); hash result on y into 50 buckets (50 buffers) -- here we pipeline
- Cost so far: $3B(R) + 3B(S)$


Example



Continuing:

- How large are the 50 buckets on y ? Answer: $k/50$.
- If $k \leq 50$ then keep all 50 buckets in Step 3 in memory, then:
- Step 4: read U from disk, hash on y and join with memory
- Total cost: $3B(R) + 3B(S) + B(U) = 55,000$


Example



Continuing:

- If $50 < k \leq 5000$ then send the 50 buckets in Step 3 to disk
 - Each bucket has size $k/50 \leq 100$
- Step 4: partition U into 50 buckets
- Step 5: read each partition and join in memory
- Total cost: $3B(R) + 3B(S) + 2k + 3B(U) = 75,000 + 2k$

Example



Continuing:

- If $k > 5000$ then materialize instead of pipeline
- 2 partitioned hash-joins
- Cost $3B(R) + 3B(S) + 4k + 3B(U) = 75000 + 4k$

Example

Summary:

- If $k \leq 50$, $\text{cost} = 55,000$
- If $50 < k \leq 5000$, $\text{cost} = 75,000 + 2k$
- If $k > 5000$, $\text{cost} = 75,000 + 4k$

Estimating Sizes

$$0.01 \times \zeta(R) \times S(S)$$

Estimating Sizes

- Need size in order to estimate cost
- Example:
 - Cost of partitioned hash-join $E1 \bowtie E2$ is $3B(E1) + 3B(E2)$
 - $B(E1) = T(E1) / \text{block size}$
 - $B(E2) = T(E2) / \text{block size}$
 - So, we need to estimate $T(E1)$, $T(E2)$

Estimating Sizes

Estimating the size of a projection

- Easy: $T(\Pi_L(R)) = T(R)$
- This is because a projection doesn't eliminate duplicates

$$\begin{aligned} R &= 1000 \\ \Pi_{dept}(R) &= 1000 \\ &\approx \underline{4} \cdot V(R, dept) \end{aligned}$$

Estimating Sizes

- Estimating the size of a selection = 1000 tuple. $V(R,A)=4$
- $S = \sigma_{A=c}(R)$ $\sigma_{dept="cs"}(stud)$ $\{"cs", "EE", "bio", "ME"\}$
 - $T(S)$ can be anything from 0 to $T(R) - V(R,A) + 1$
 - Mean value: $T(S) = T(R)/V(R,A) = \frac{1000}{4} = 250$
 - $S = \sigma_{A < c}(R)$ $age < 38$
 - $T(S)$ can be anything from 0 to $T(R)$
 - Heuristics: $T(S) = T(R)/3$ The Lady's Sellinger magic #.
- $1000 = \frac{1}{3} \times 1000$