

CS411

Database Systems

01: Introduction

Kevin C. Chang

Welcome to CS411

- Web site:

<https://wiki.engr.illinois.edu/display/cs411fa13/>

- Announcements, syllabus, policies, schedule, ...
- Please read the class syllabus, policies, and lecture schedule; ask now if you have questions.

What is this?



Teaching Staff: The Front-End



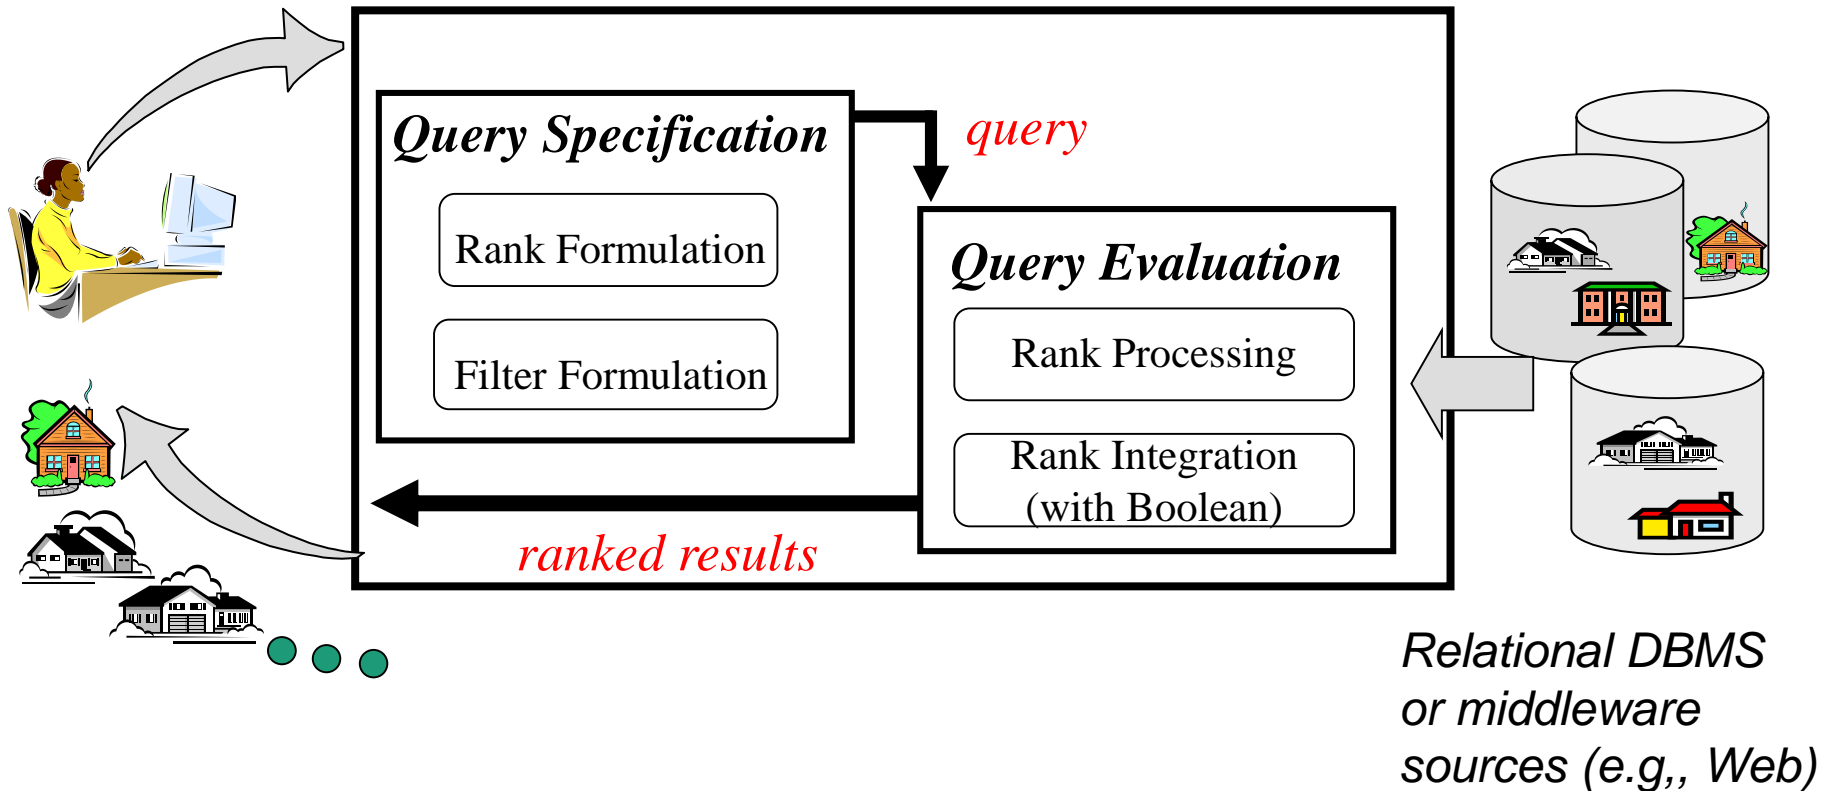
Miawei, Tanvi, Mangesh, Phuong

Teaching Staff: The Back-End

- Kevin C. Chang
- Research interest:
 - Database systems, Web integration and mining
 - Research projects:
 - AIM, MetaQuerier, WISDM, ARISE, and we are recruiting!
- Hobbies:
 - Ocean diving, mountain climbing.
- Brief history
 - Taiwan (BS in EE from National Taiwan University)
 - California (MS in CS, PhD in EE from Stanford)
 - Illinois (associate professor in CS, UIUC)
 - “Data mining”: what can you predict? East or west? CS or EE?

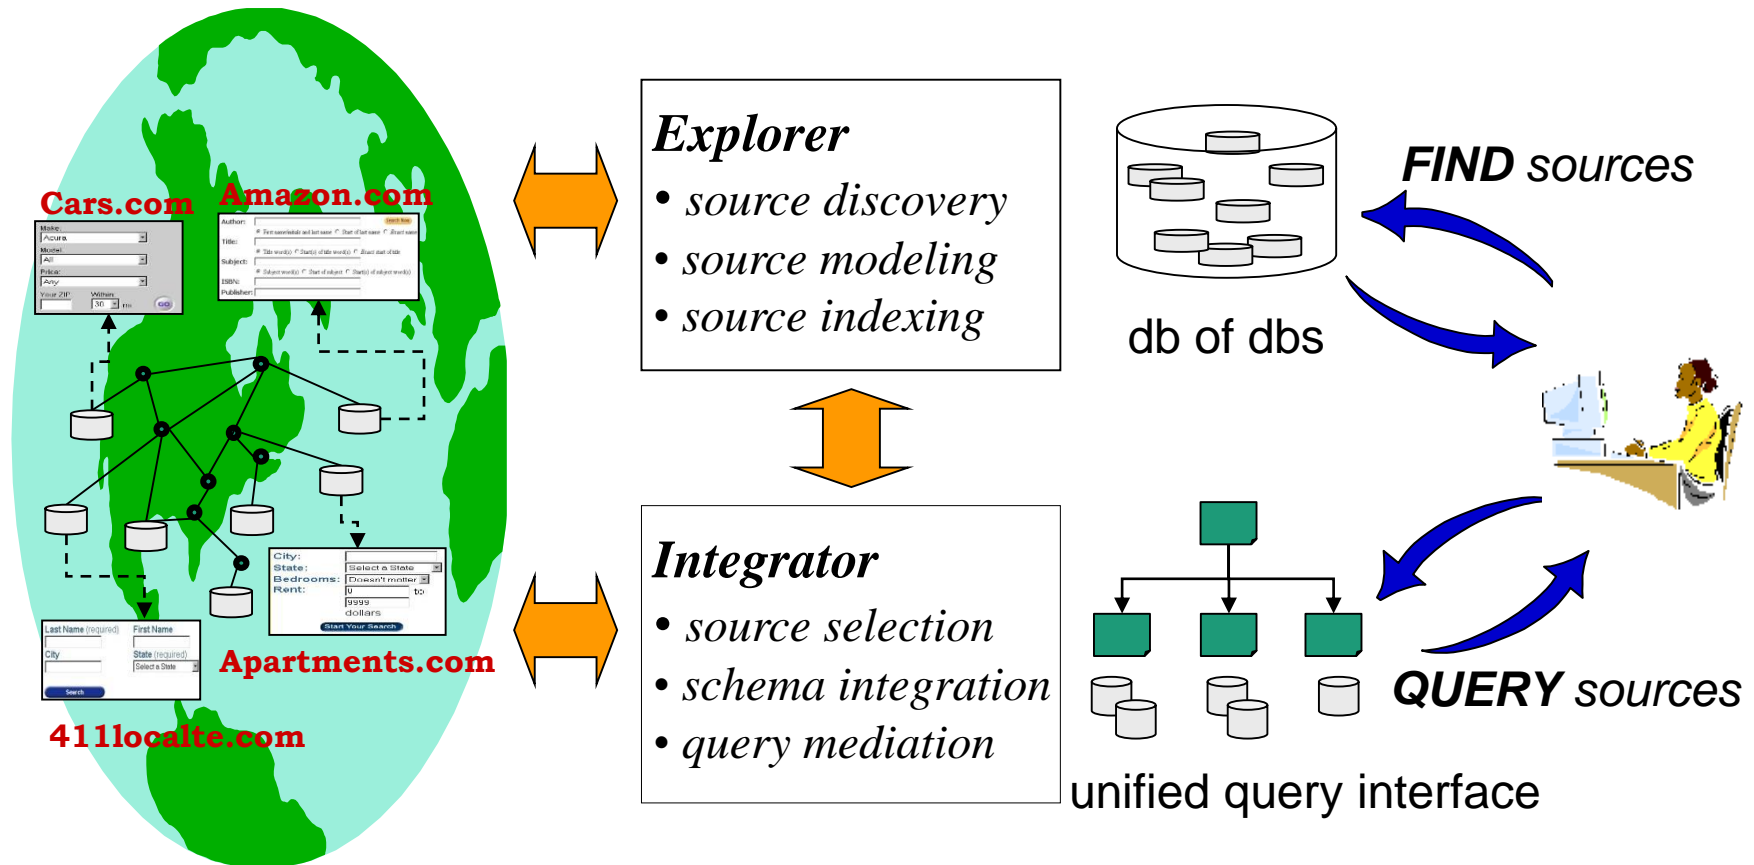
The AIM Project: AIMing to the top

Supporting ranking in data retrieval



The MetaQuerier Project

Exploring and integrating deep Web




Data Aware Search over the Web



Data Aware Search?



The ARISE Project: Augmented Reality Information Search Engine



Augmented Reality Information Search Engine



You Tell Me --

- Why are you taking this course?
- It enables many possibilities:
 - An IT Guru at Goldman-Sachs or Boeing.
 - A System Developer at Oracle or Google.
 - A CEO + receptionist + janitor (AKA startup cofounder).
 - None of the above— Please name it.

CS411 Goal: Two Perspectives of DBMS

- User perspective
 - how to use a database system?
 - conceptual data modeling, the relational and other data models, database schema design, relational algebra, and the SQL query language.
- System perspective
 - how to design and implement a database system?
 - data representation, indexing, query optimization and processing, transaction processing, concurrency control, and crash recovery

Perquisite

- Must have data structure and algorithm background
 - CS 225 or 400 equivalent
- Good programming skill
 - project will require lot of programming
 - need C++, Java, or PHP ... to do a good job at talking with DB
 - you or your project group picks the language

Textbook



- Textbook:
Database Systems: The Complete Book, 2/e, by Hector Garcia-Molina, Jeffrey D. Ullman, and Jennifer D. Widom
- Good references:
 - *Database Management Systems*, by Raghu Ramakrishnan and Johannes Gehrke, McGraw-Hill
 - *Database System Concepts*, by Abraham Silberschatz, Henry F. Korth, and S. Sudarshan, McGraw Hill
 - *Fundamentals of Database Systems*, by Ramez Elmasri and Shamkant Navathe, Addison Wesley
 - *An Introduction to Database Systems*, by C. J. Date, Addison Wesley

Course Format

- For all students
 - two 75-min lectures / week
 - 5-6 assignments planned ← “MP Inside!”
 - project (significant)
 - a midterm and a final exam
- Graduate students: 4 credits option
 - do an extra project

Lectures

- Lecture slides will be posted shortly before or after the lecture
 - are to complement the lectures
- Lectures are important for guiding your reading of textbook
- Without your participation, a lecture is impossible
 - so please attend lectures regularly

Homework Assignments

- Mostly paper-based, some may involve machine exercise/ programming
- Due by noon of the due date
- No late homework will be accepted

Project:

The 411 Incubator, Where Ideas are Hatched!

- General Track: Database-Driven Web/Mobile-based Applications
 - select an “useful” and “innovative” application that needs a database
 - design and build it from start to finish
 - your choice of topic:
useful, innovative, database-driven
- System Track: Take an open source DBMS (SQLite) and hack it.
 - strongly encouraged if you already have app experience
- Team work
- Significant amount of programming (we will provide tutorials)
- Will be done in stages
 - you will submit some work at the end of each stage
- Start from your PITCH and end with your DEMO.

Project Groups

- Project will be done in group of 3-4 students
 - learn how to work in a group: valuable skills
 - also use project group as study partners
- Try to form groups as soon as possible
 - can start by posting requests on the class newsgroup
- There will be a deadline soon for forming groups
 - if you have not formed groups by then
 - we will help assign you to groups
- Grading:
 - all members receive same grading
 - if someone drops out, the rest pick up the work

Exams

- Midterm and final
- There will be some brief review before each exam
- Check dates and make sure no conflict!
 - generally no makeup exams unless exceptional cases (see policy page)

In-Class Quiz

To encourage attending and participating in class:

- Three brief in class quiz.
- Each tests your participation in class.
- Each accounts for 1.5% of class grading.

Tutorials: 5-10 Sessions by Needs

- Homework Tutorial:
 - One for each homework assignment, to teach problem solving.
- Exam Tutorial:
 - One for each exam, to review and practice old exams.
- Project Tutorials:
 - Web/mobile programming, DB programming.
- Scheduled a few days before due. Likely at Tu or Th 5pm.
- Will be recorded.

Tentative Grading Breakdown

- Homework: 15%
- Project: 35%
- Midterm: 20%
- Final: 30%

How Do We Work Together?

Contacting the Staff --

Office Hours

- The best way for asking questions and clarifications
- Will have office hours every day
- See course Web for schedule


Communications

- Website: “Announcements” page
- Q/A forum: <https://piazza.com/class/hj3y02os5no6zy>
 - vitally important!
 - make sure to check it regularly for questions/clarifications
- If you have a question/problem
 1. talk to people in your group first
 2. post your question on newsgroup
 3. email TA
 4. go to office hours to talk to TA or instructor

Q/A Forum: Piazza

- Designed for you and your peer
 - to communicate and help one another
 - please do not post solutions/admin-requests
- TAs will monitor and try their best to help with your questions
- There can be many questions
 - may not be able to answer all of them timely manner
 - not good for more complex questions
 - hence should come to office hours or email TA

Data Management Evolution



Jim Gray: *Evolution of Data Management*.
IEEE Computer 29(10): 38-46 (1996):

- Manual processing: -- 1900
- Mechanical punched-cards: 1900-1955
- Stored-program computer-- sequential record processing: 1955-1970
- Online navigational network DBs: 1965-1980
 - many applications still run today!
- Relational DB: 1980-1995
- Post-relational and the Internet: 1995-

Database Management System (DBMS)?

- System for providing EFFICIENT, CONVENIENT, and SAFE MULTI-USER storage of and access to MASSIVE amounts of PERSISTENT data
- Really?? Let's contrast with a **File System**:
 - Persistent?
 - Efficient? Convenient? Safe? Multi-user?
 - Massive?
- Homework: Contrast with main memory system.

DBMS Examples

- Most familiar use: many Web sites rely heavily on DBMS's. Examples?
- And many non-Web examples

Example: Banking system

- Data = information on accounts, customers, balances, current interest rates, transaction histories, etc.
- MASSIVE: many gigabytes at a minimum for big banks, more if keep history of all transactions, even more if keep images of checks -> Far too big for memory
- PERSISTENT: data outlives programs that operate on it

MULTI-USER Access

- MULTI-USER: many people/programs accessing same database, or even same data, simultaneously -> Need careful controls
- Alex @ ATM1: withdraw \$100 from account #002
- Bob @ ATM2: withdraw \$50 from account #002

- What should happen, then?

Why Direct Implementation Won't Work

- Storing data: file system is limited
 - size limit by disk or address space
 - when system crashes we may loose data
 - Password/file-based authorization insufficient
- Query/update:
 - need to write a new C++/Java program for every new query
 - need to worry about performance

- Concurrency: limited protection
 - need to worry about interfering with other users
 - need to offer different views to different users (e.g. registrar, students, professors)
- Schema change:
 - entails changing file formats
 - need to rewrite virtually all applications
- That's why the notion of DBMS was motivated!

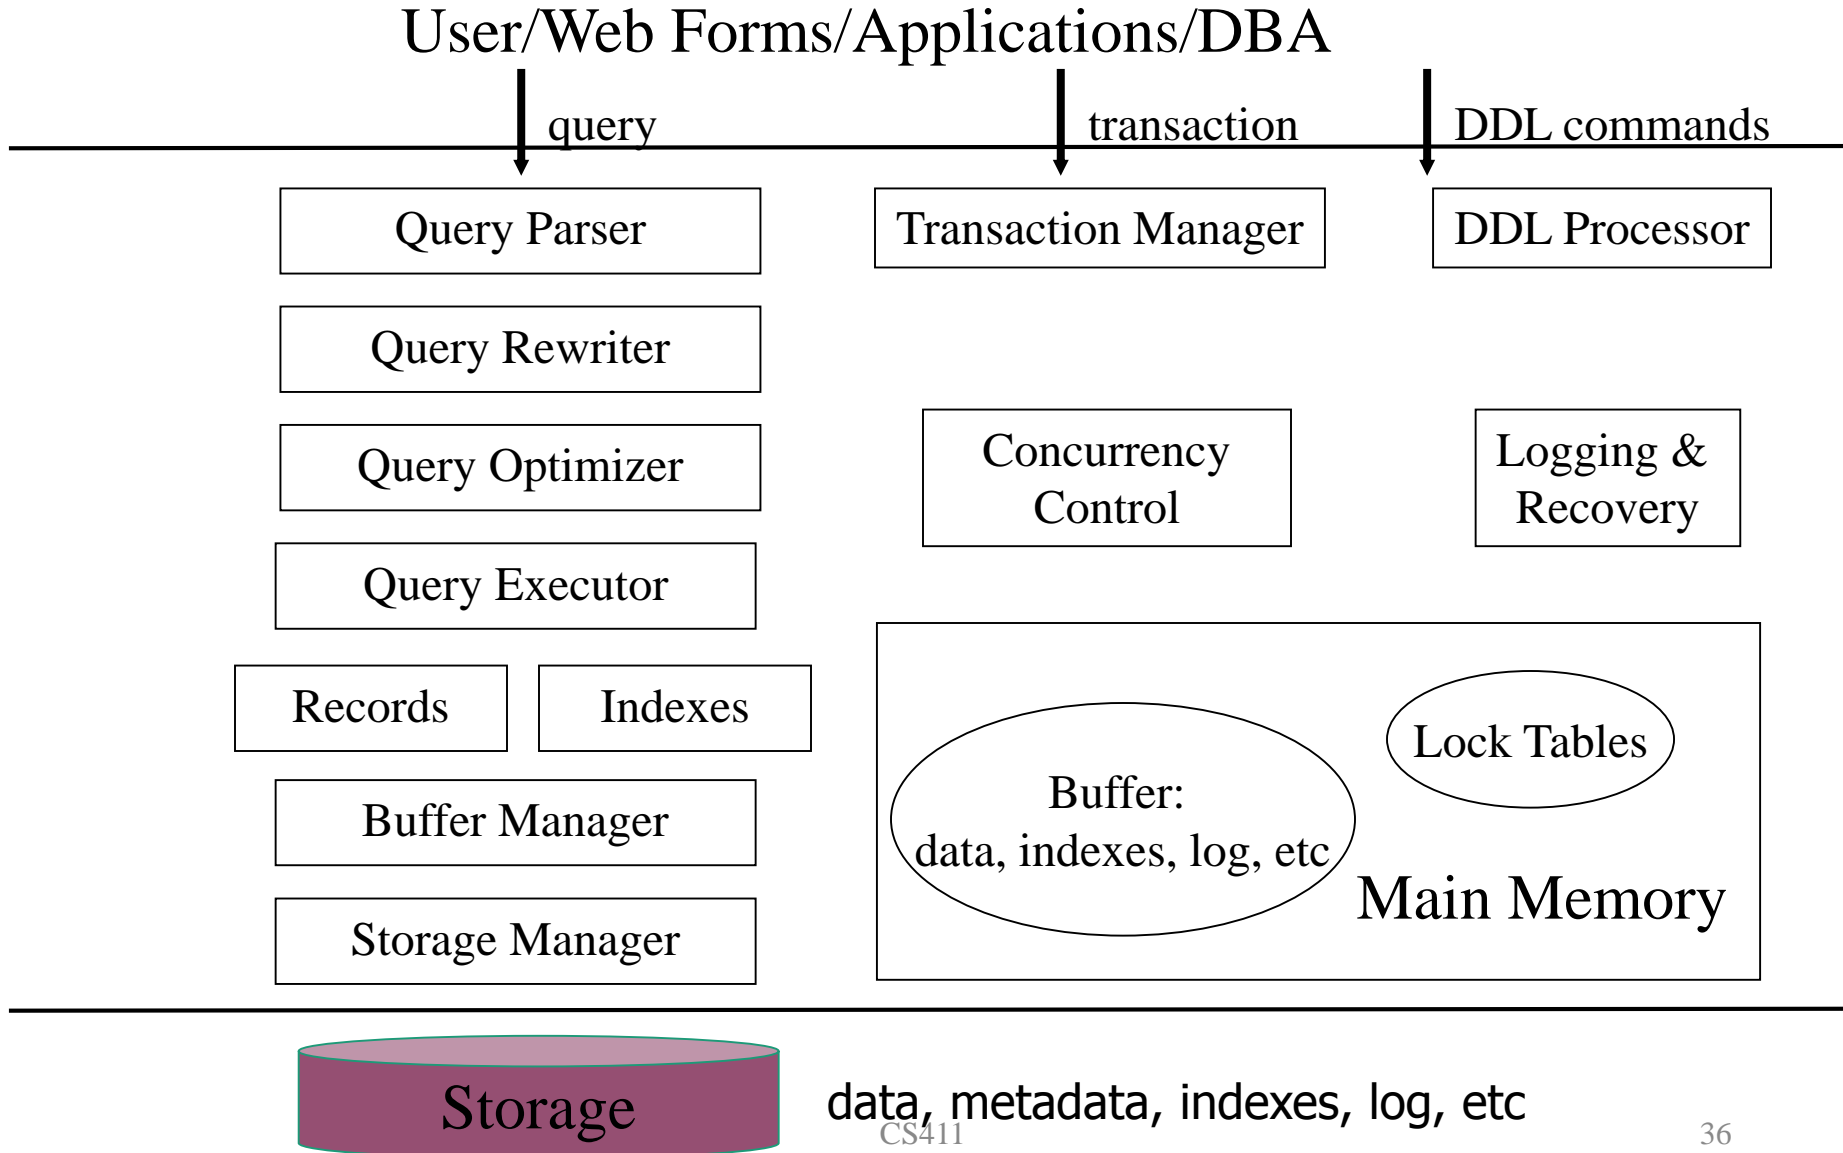
DBMS: More Requirements

- **SAFE:**
 - from system failures
 - from malicious users
- **CONVENIENT:**
 - simple commands to - debit account, get balance, write statement, transfer funds, etc. ->
 - also unpredicted queries should be easy
- **EFFICIENT:**
 - don't search all files in order to - get balance of one account, get all accounts with low balances, get large transactions, etc.
 - massive data! -> DBMS's carefully tuned for performance

DBMS: A Software System

- Buy, install, set up for particular application
- Available for PC's, workstations, mainframes, supercomputers
- Major vendors:
 - Oracle
 - IBM (DB2)
 - Microsoft (SQL Server, Access)
 - Sybase
- all are "relational" (or "object-relational") DBMS

DBMS Architecture



People

- DBMS user: queries/modifies data
- DBMS application designer
 - set up schema, loads data, ...
- DBMS administrator
 - user management, performance tuning, ...
- DBMS implementer: builds systems

Questions we must ask...

- What is data? How to view our data?
- How to “organize” data?
- How to “compute”/”query” on data?
- How to categorize (index) data?
- How to process data?
- How to acquire data?
- How to further scale up?

1/3 Topics: User Perspective

- Entity-Relationship Model
- Relational Model
- Relational Database Design
- Relational Algebra
- SQL and DBMS Functionalities:
 - SQL Programming
 - Queries and Updates
 - Indexes and Views
 - Constraints and Triggers

1/3 Topics: System Perspective

- Storage and Representation
- Indexing
- Query Execution and Optimization
- Transaction Management

1/3 Topics:

From **DB** (*relational*) to **BD** (*everything?*)

- Flexible (non-relational) data representation
 - XML/SOAP
 - JSON
 - NoSQL Databases
- Big data acquisition – Information extraction
- Big data processing – Map-Reduce

Special Topics: If You Ask!

- Databases in the Real World
- Advanced Database Research

Please nominate topics you like to hear about!

How to Get the Most out of CS411? – The classic wisdom:

- **Read and think before class**
 - welcome to ask questions before class!
- Study and discuss with your peers
 - discuss readings to enhance understanding
 - discuss assignments but write your own solution!
- Use lectures to guide your study
 - use it as a roadmap for what's important
 - lectures are **starting** points– they do not cover everything you should read

How to Get the Most out of CS411?

- Knowing why you are here.

You can work at Goldman Sachs if you want.

You can work at Google if you want.

You are qualified to be a receptionist or a janitor, or a CEO, if you want.

Questions?

- Any questions? Please come talk to me.