

SQL1: The Basics

CS411 Database Systems

Kevin C. Chang

SQL Introduction

Standard language for querying and manipulating data

Structured Query Language

Many standards out there: SQL92, SQL2, SQL3, SQL99

Vendors support various subsets of these, but all of what we'll be talking about.

Behind the Scene: The Birth of SQL

- Chamberlin, D. D. and Boyce, R. F. 1974. SEQUEL: A structured English query language. In Proceedings of the 1974 ACM SIGFIDET (Now Sigmod) Workshop on Data Description, Access and Control (Ann Arbor, Michigan, May 01 - 03, 1974). FIDET '74. ACM, New York, NY, 249-264.

SEQUEL: A STRUCTURED ENGLISH QUERY LANGUAGE

by

Donald D. Chamberlin
Raymond F. Boyce

IBM Research Laboratory
San Jose, California

SQL

~~A~~

X

ABSTRACT: In this paper we present the data manipulation facility for a (X) which can be used for accessing data in an integrated relational database. Without resorting to the concepts of relational algebra, SEQUEL identifies a set of simple operations on tabular structures, which can be shown to be of equivalent power to the first order predicate calculus. A SEQUEL user is presented with a consistent set of templates which reflect how people use tables to obtain information. Moreover, the SEQUEL user is able to compose these basic templates in a structured manner in order to form more complex queries. SEQUEL is intended as a data base sublanguage for both sequential and random access data bases.

Why SQL?

- SQL is a very-high-level language, in which the programmer is able to avoid specifying a lot of data-manipulation details that would be necessary in languages like C++.
- What makes SQL viable is that its queries are “optimized” quite well, yielding efficient query executions.

Select-From-Where Statements

- The principal form of a query is:

<code>SELECT</code>	desired attributes
<code>FROM</code>	one or more tables
<code>WHERE</code>	condition about tuples of the tables

Running Example



Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

Example

- Using Beers(name, manf), what beers are made by Anheuser-Busch?

```
SELECT name  
FROM Beers  
WHERE manf = 'Busch';
```

Meaning of Single-Relation Query

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

“Everything” In SELECT Clause: STAR

- When there is one relation in the FROM clause, * in the SELECT clause stands for “all attributes of this relation.”
- Example using Beers(name, manf):

```
SELECT *
```

```
FROM Beers
```

```
WHERE manf = 'busch';
```

Renaming Attributes

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.
- Example based on Beers(name, manf):

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Busch'
```

Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- Example: from Sells(bar, beer, price):

```
SELECT bar, beer, price * 120 AS priceInYen  
FROM Sells;
```

Complex Conditions in WHERE Clause

- Composing conditions:
 - Use AND, OR, NOT, and parentheses.
- From Sells(bar, beer, price), find the price Joe's Bar charges for Bud:

```
SELECT bar, beer, price
FROM Sells
WHERE bar = 'joe bar' AND price < 5.0;
```

```
SELECT bar, beer, price
FROM Sells
WHERE (bar = 'joe bar' AND price < 5.0) or (beer = 'sam adam');
```

Selection Conditions

- Each condition
 - Attribute names of the relation(s) used in the FROM.
 - Comparison operators: =, <>, <, >, <=, >=
 - Apply arithmetic operations: stockprice*2
 - Operations on strings (e.g., "||" for concatenation).
 - Pattern matching: s LIKE p
Pattern: % = "any string"; _ = "any character."

```
SELECT name, phone FROM Drinkers  
WHERE phone LIKE '%333-__ __';
```

- Lexicographic order on strings.
SELECT name, city FROM Drinkers
WHERE name < 'c';
- Special stuff for comparing dates and times.
- SQL is *case-insensitive*. In general, upper and lower case characters are the same, except inside quoted strings.

What Should This Query Return?

```
SELECT bar, beer  
FROM Sells  
WHERE price < 5.00 OR price >= 5.00;
```

vs.

```
SELECT bar, beer FROM Sells;
```

NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
 - *Missing value* : e.g., we know Joe's Bar has some address, but we don't know
 - *Inapplicable* e.g., the value of attribute *spouse* for an unmarried person.

Comparing NULL's to Values

- The truth of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- When any value is compared with NULL, the truth value is UNKNOWN.
- But a query only produces a tuple in the answer if its truth value for the WHERE clause is TRUE (not FALSE or UNKNOWN).

Three-Valued Logic

To understand how AND, OR, and NOT work in 3-valued logic:

- TRUE = 1.
- FALSE = 0.
- UNKNOWN = $\frac{1}{2}$.

Operators:

- AND = MIN.
- OR = MAX.
- NOT(x) = 1-x.

- Ex: TRUE AND (FALSE OR NOT(UNKNOWN))

$$= \text{MIN}(1, \text{MAX}(0, (1 - \frac{1}{2})))$$

$$= \text{MIN}(1, \text{MAX}(0, \frac{1}{2}))$$

$$= \text{MIN}(1, \frac{1}{2}) = \frac{1}{2}.$$

$$\text{price} \geq 5 \cup \text{price} < 5$$

when $\text{pr} = \text{Null}$.

$$\rightarrow \frac{1}{2} \cup \frac{1}{2} = \max(\frac{1}{2}, \frac{1}{2}) = \frac{1}{2}$$

$$\frac{1}{2} = \text{Min}(x, y) = \text{Min}(1, \frac{1}{2})$$

$$T \wedge T = 1. \quad T \wedge F = 0. \quad T \wedge U = U \\ = \text{Min}(1, 0) = 0 \quad = \frac{1}{2}$$

Reason: 2-Valued Laws \neq 3-Valued Laws

- Some common laws, like the commutativity of AND, hold in 3-valued logic.
- But others do not:
 - Ex: „ p OR NOT p = TRUE (“law of excluded middle”)
 - When p = UNKNOWN, the left side is $\text{MAX}(\frac{1}{2}, (1 - \frac{1}{2})) = \frac{1}{2}$ which is not 1 or TRUE.

Testing for Null

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

SELECT bar, beer

FROM Sells

WHERE price < 5.00 OR price >= 5.00 OR price IS NULL

Multi-relation Queries

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>”

Example

- Find the beers liked by at least one person who frequents Joe's Bar.
- Tip: Always prefix with relation name to make it clear/easier to read.

```
SELECT Likes.beer  
FROM Likes, Frequents  
WHERE Frequents.bar = 'joe bar' AND  
       Frequents.drinker = Likes.drinker;
```