

UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Indexing



Announcements

- MP2 will be posted tonight
 - About 10 SQL statements to write
 - One week to finish
- Exam scores should all be posted



Review

- We’re studying how databases are implemented
- What are the levels of the memory hierarchy?
- Which level are we optimizing for databases?
- What is meant by “volatile” memory?



Review

- What steps are required to access a block?
- What is the advantage sequential reads?
- What is a spanned record?
- What are BLOBS?
- What is a tombstone?



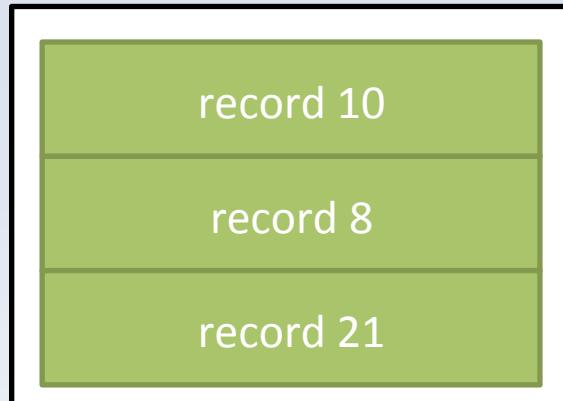
So far...

- We've studied how to arrange data on the disk
- Now we're going to study how to lay out the data so it can be queried more efficiently

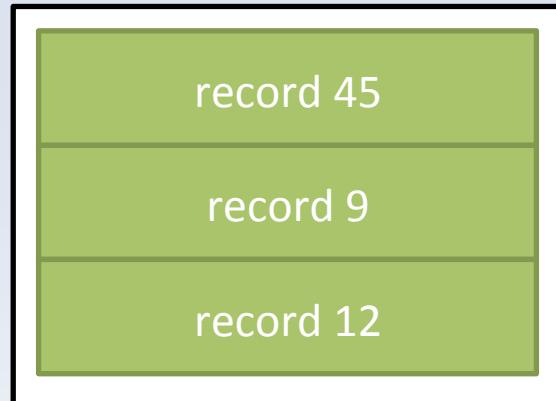


A relation on disk

block 1



block 2



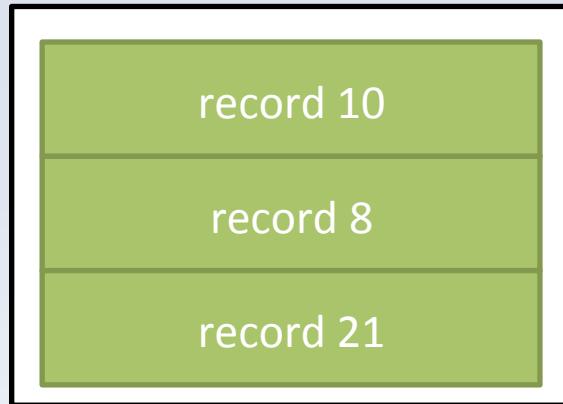
block 3



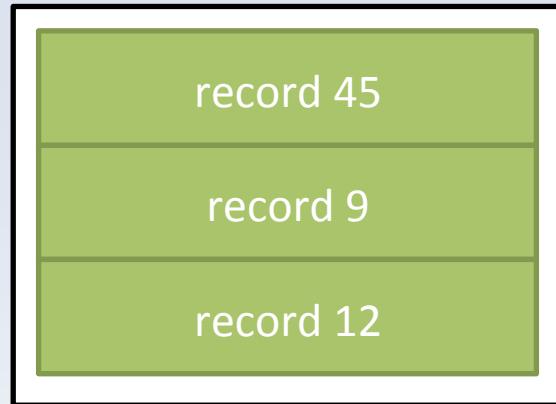
block 4

A relation on disk

block 1



block 2



```
SELECT *  
FROM R;
```

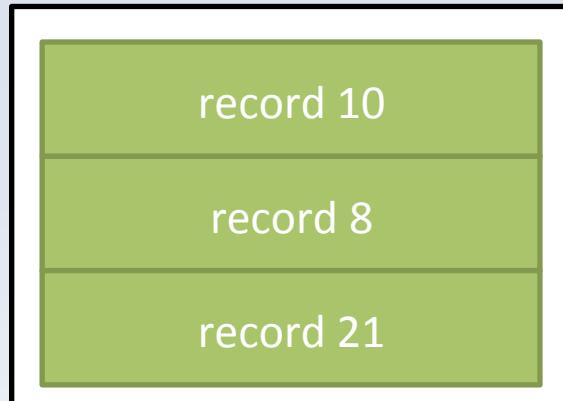


block 3

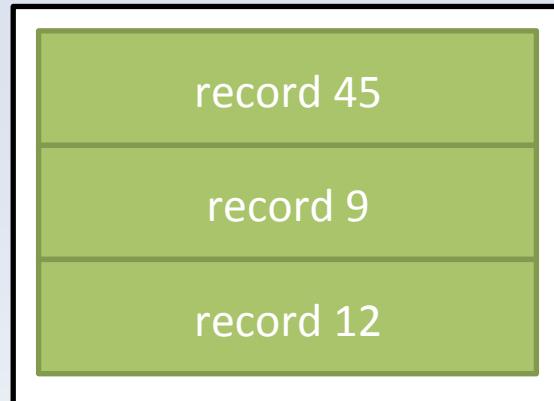
block 4

A relation on disk

block 1



block 2



```
SELECT *  
FROM R  
WHERE id=33;
```



block 3

block 4

Indexes

- An index on a relation speeds up selection on the search key fields
- Search key - any subset of the attributes of a relation
 - ***not*** to be confused with a relation's key



Indexes

- Entries in an index: (k,r)
 - k = the key
 - r = the record OR a pointer to a record OR a pointer to a collection of records
- Stored in an “index file”
 - “data file” is what we call a relation on disk



Index terminology

- Clustered/unclustered
 - Clustered - records sorted in key order
 - Unclustered - not sorted by search key
- Dense/sparse
 - Dense - each record has an entry in index file
 - Sparse - only some records have entries



Primary/secondary

- Primary - index is on the primary key of the relation
 - Records are sorted on disk by primary key
 - Clustered index
- Secondary - index is on some other set of attributes



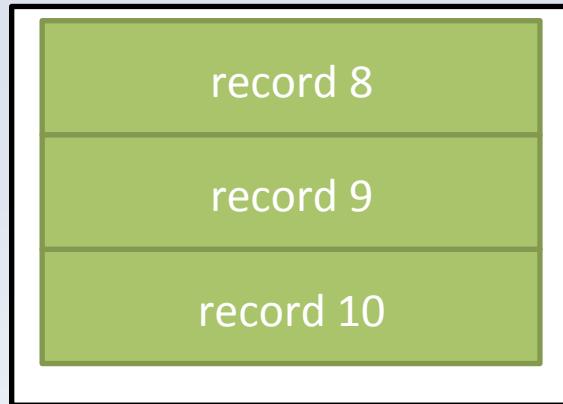
Sequential File

- A data file sorted by the primary key is called a *sequential file*

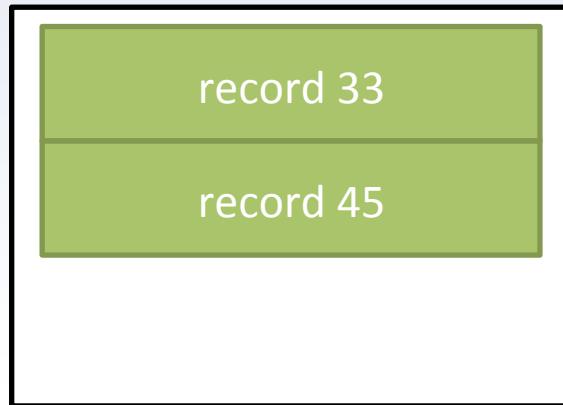
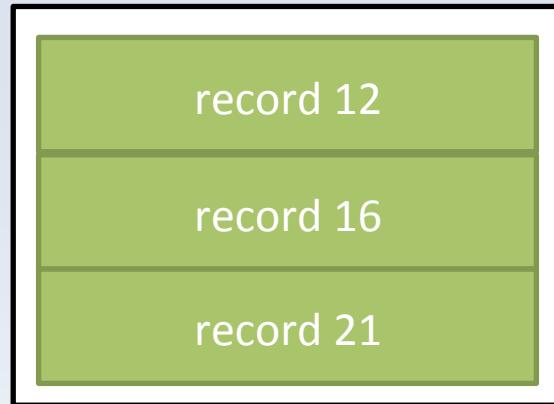


Sequential File

block 1



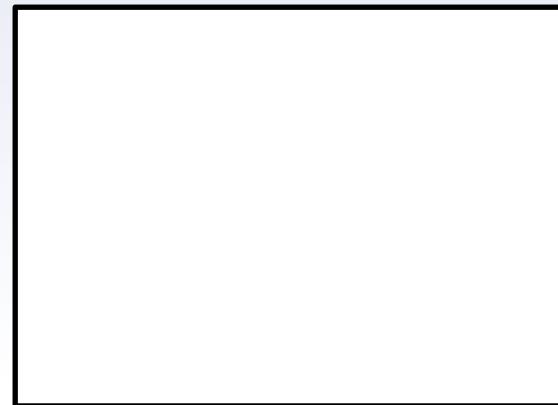
block 2



block 3



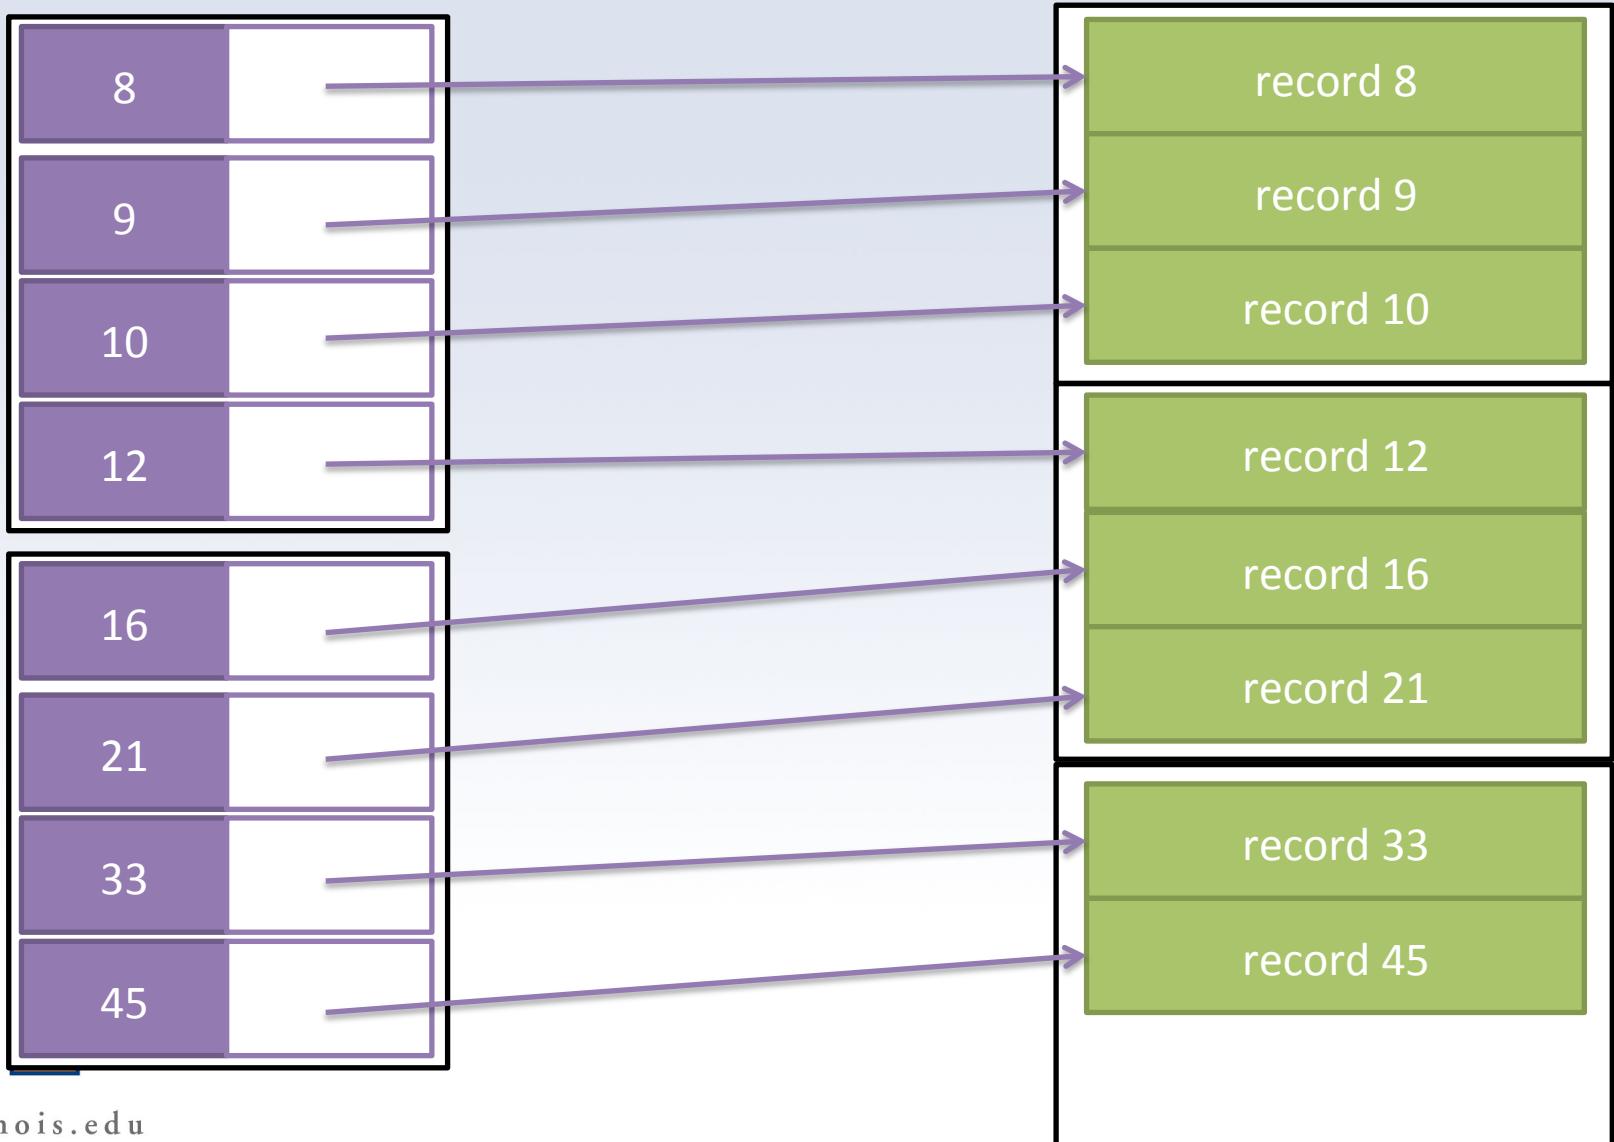
block 4



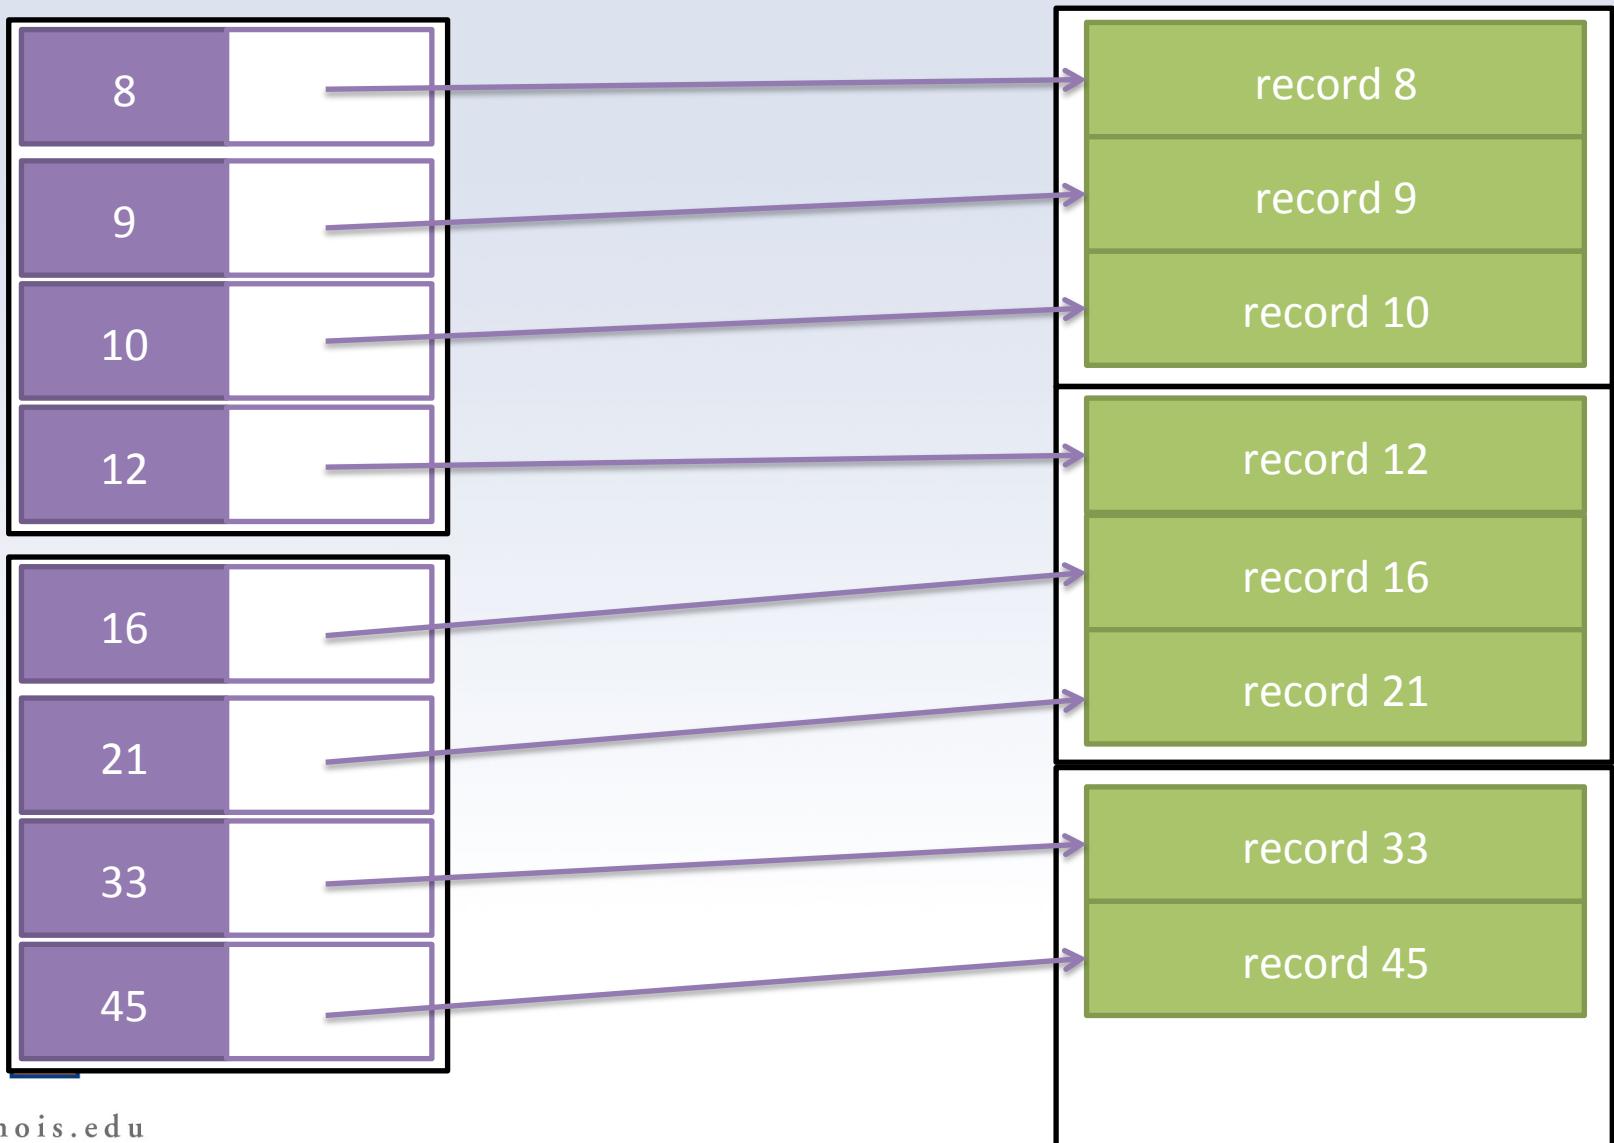
Sequential File



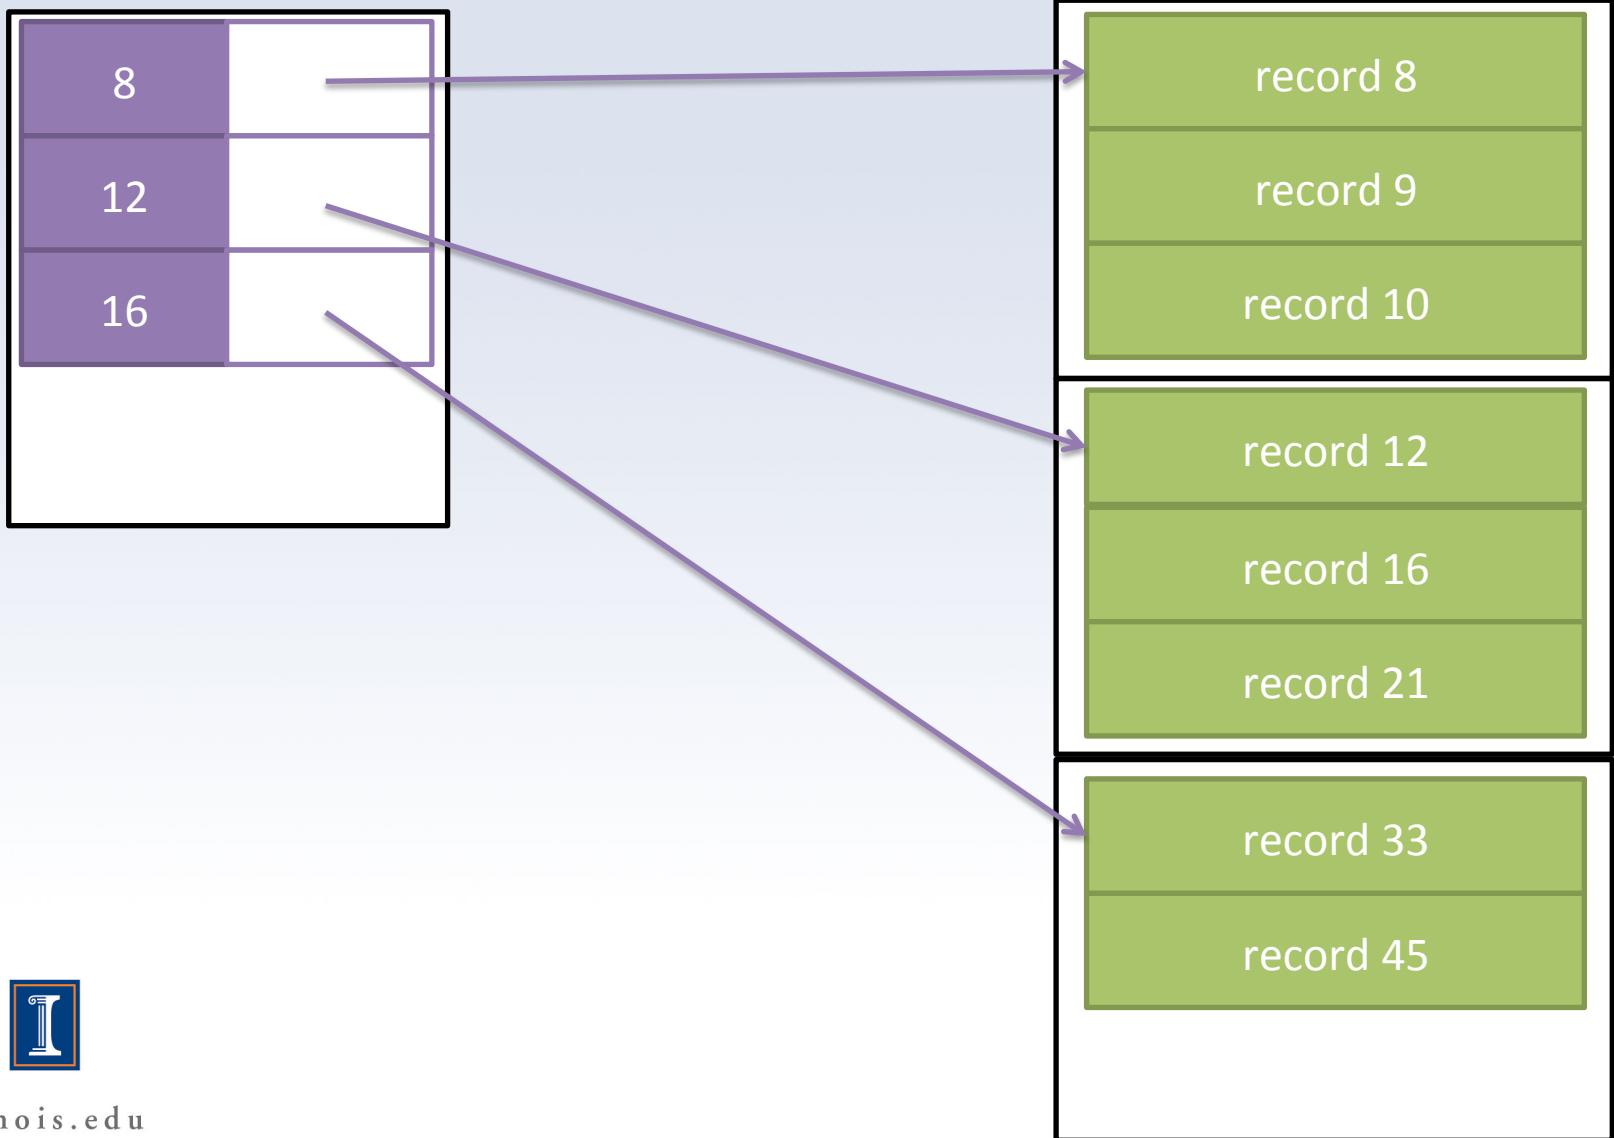
Dense Index



Dense Index



Sparse Index



Sparse index

- Note: can only use sparse index if data file is sorted
 - if index is not for primary key, ***must*** use a dense index

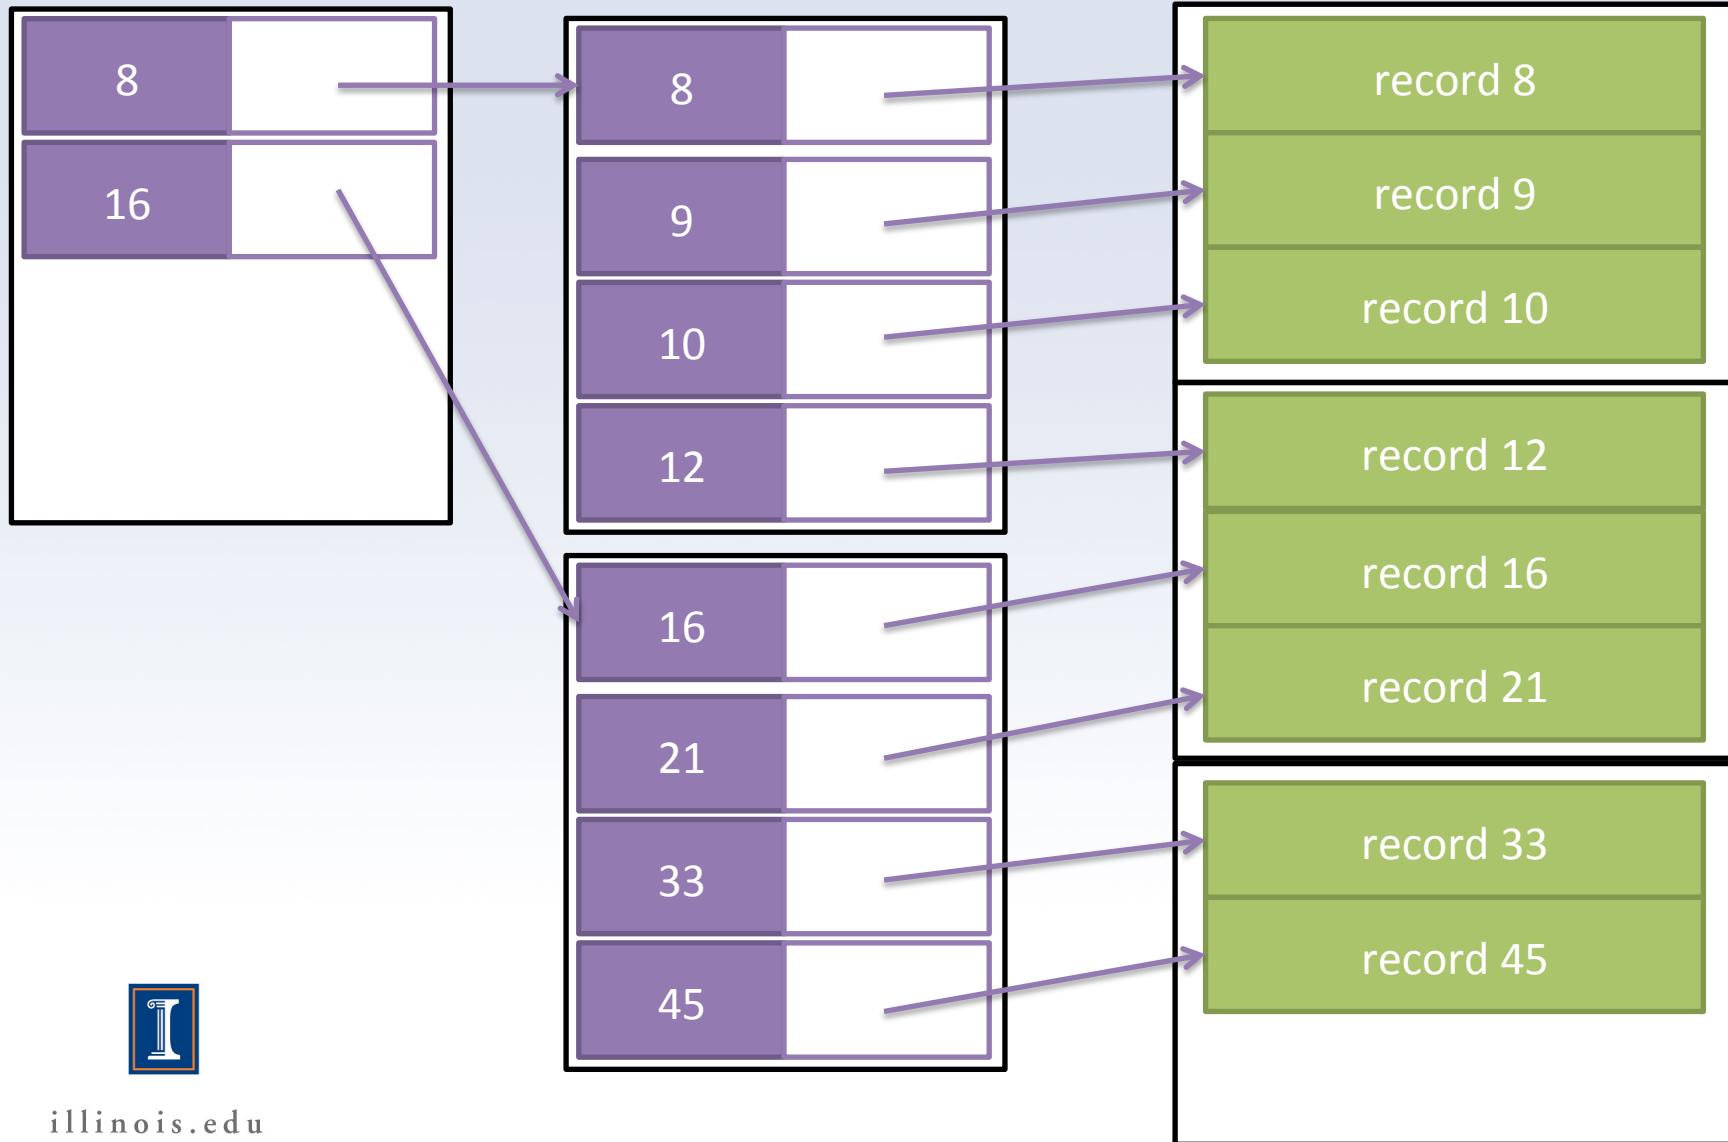


Multilevel index

- If index file is large, binary search for key can require many disk operations
- Why?
- To mitigate this, we can put an index on the index
- Obviously, this can be repeated
- B-trees do this in an optimal way



Multilevel Index

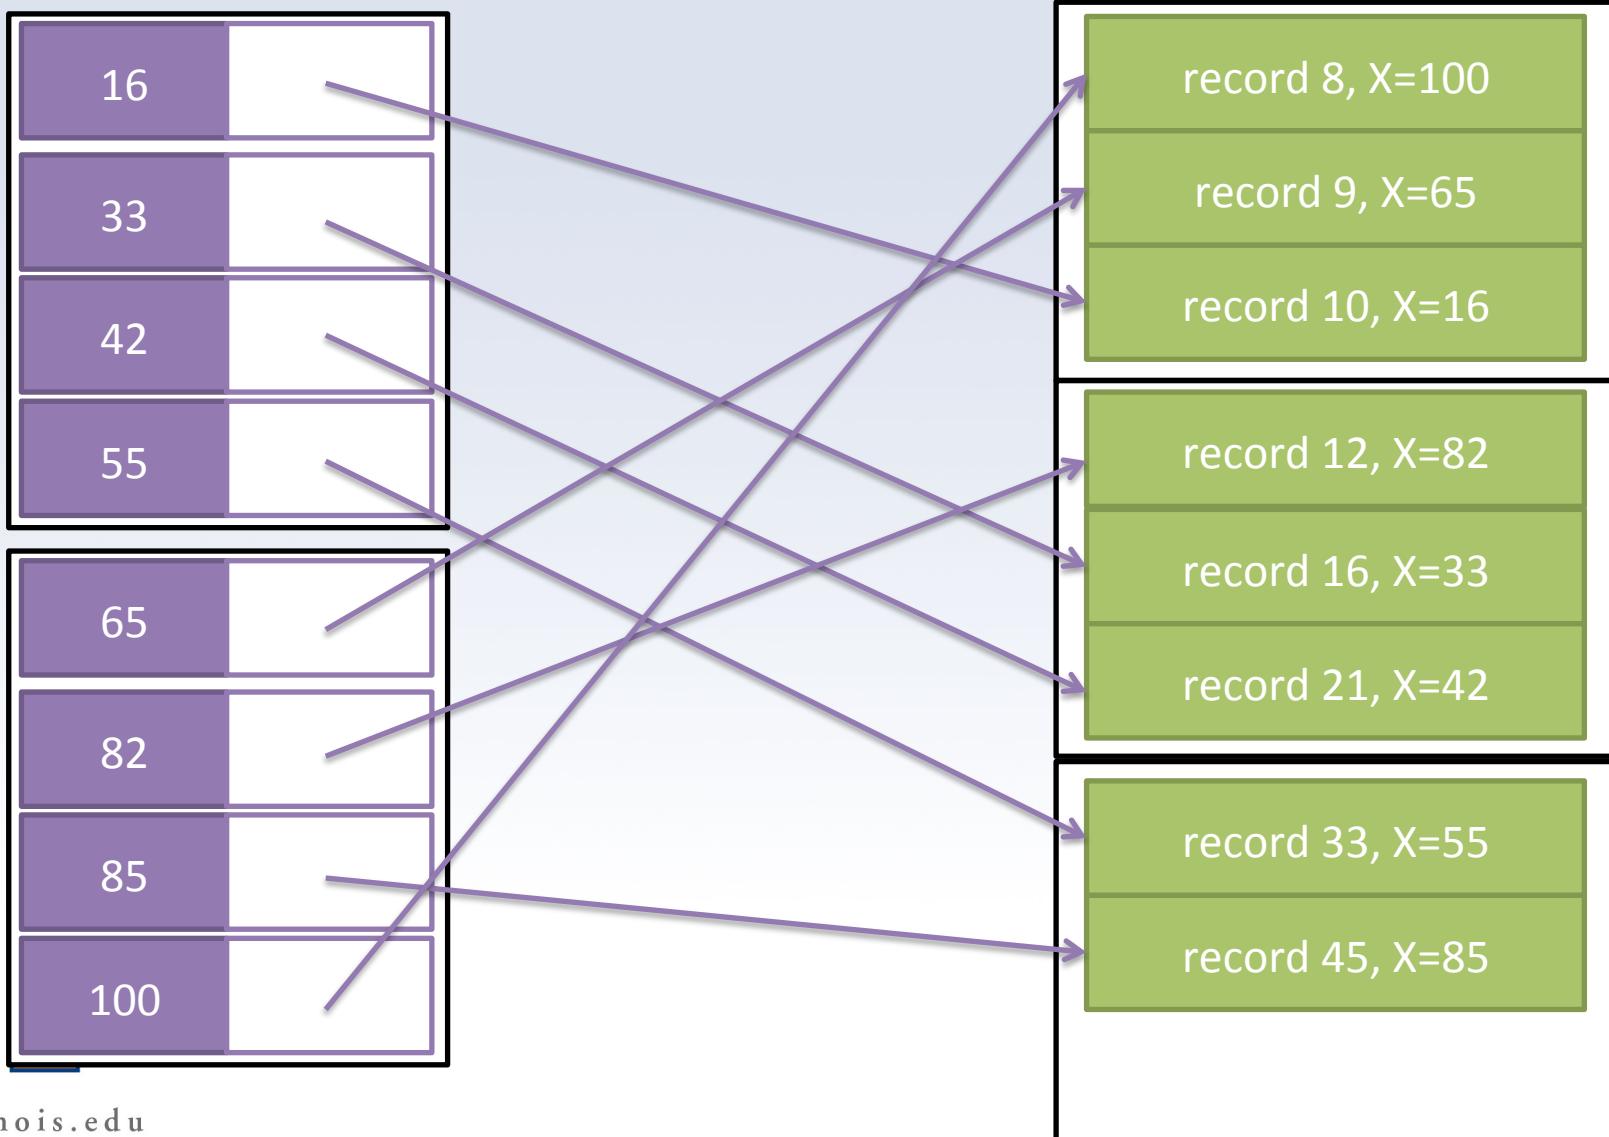


Secondary indexes

- To speed up search for non-primary key fields, we create a ***secondary index***
 - Must be unclustered
 - Must be dense



Secondary Index

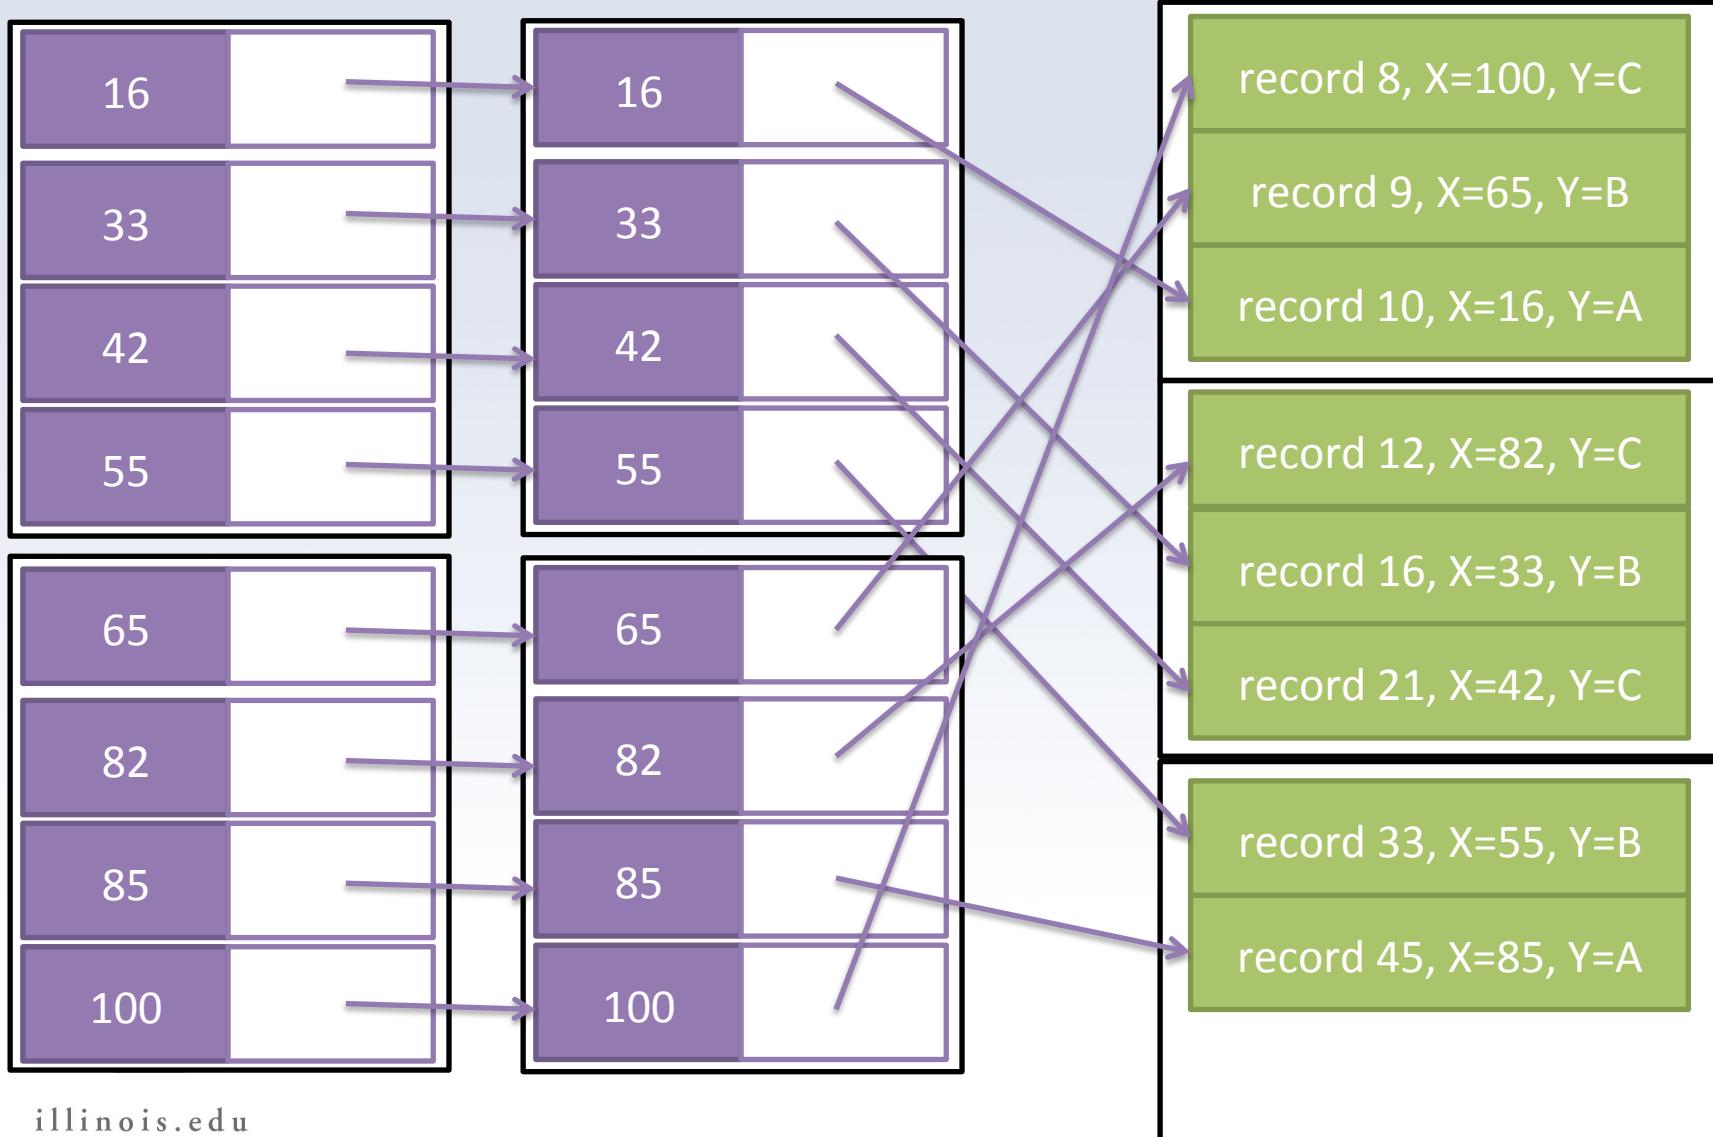


Secondary Indexes

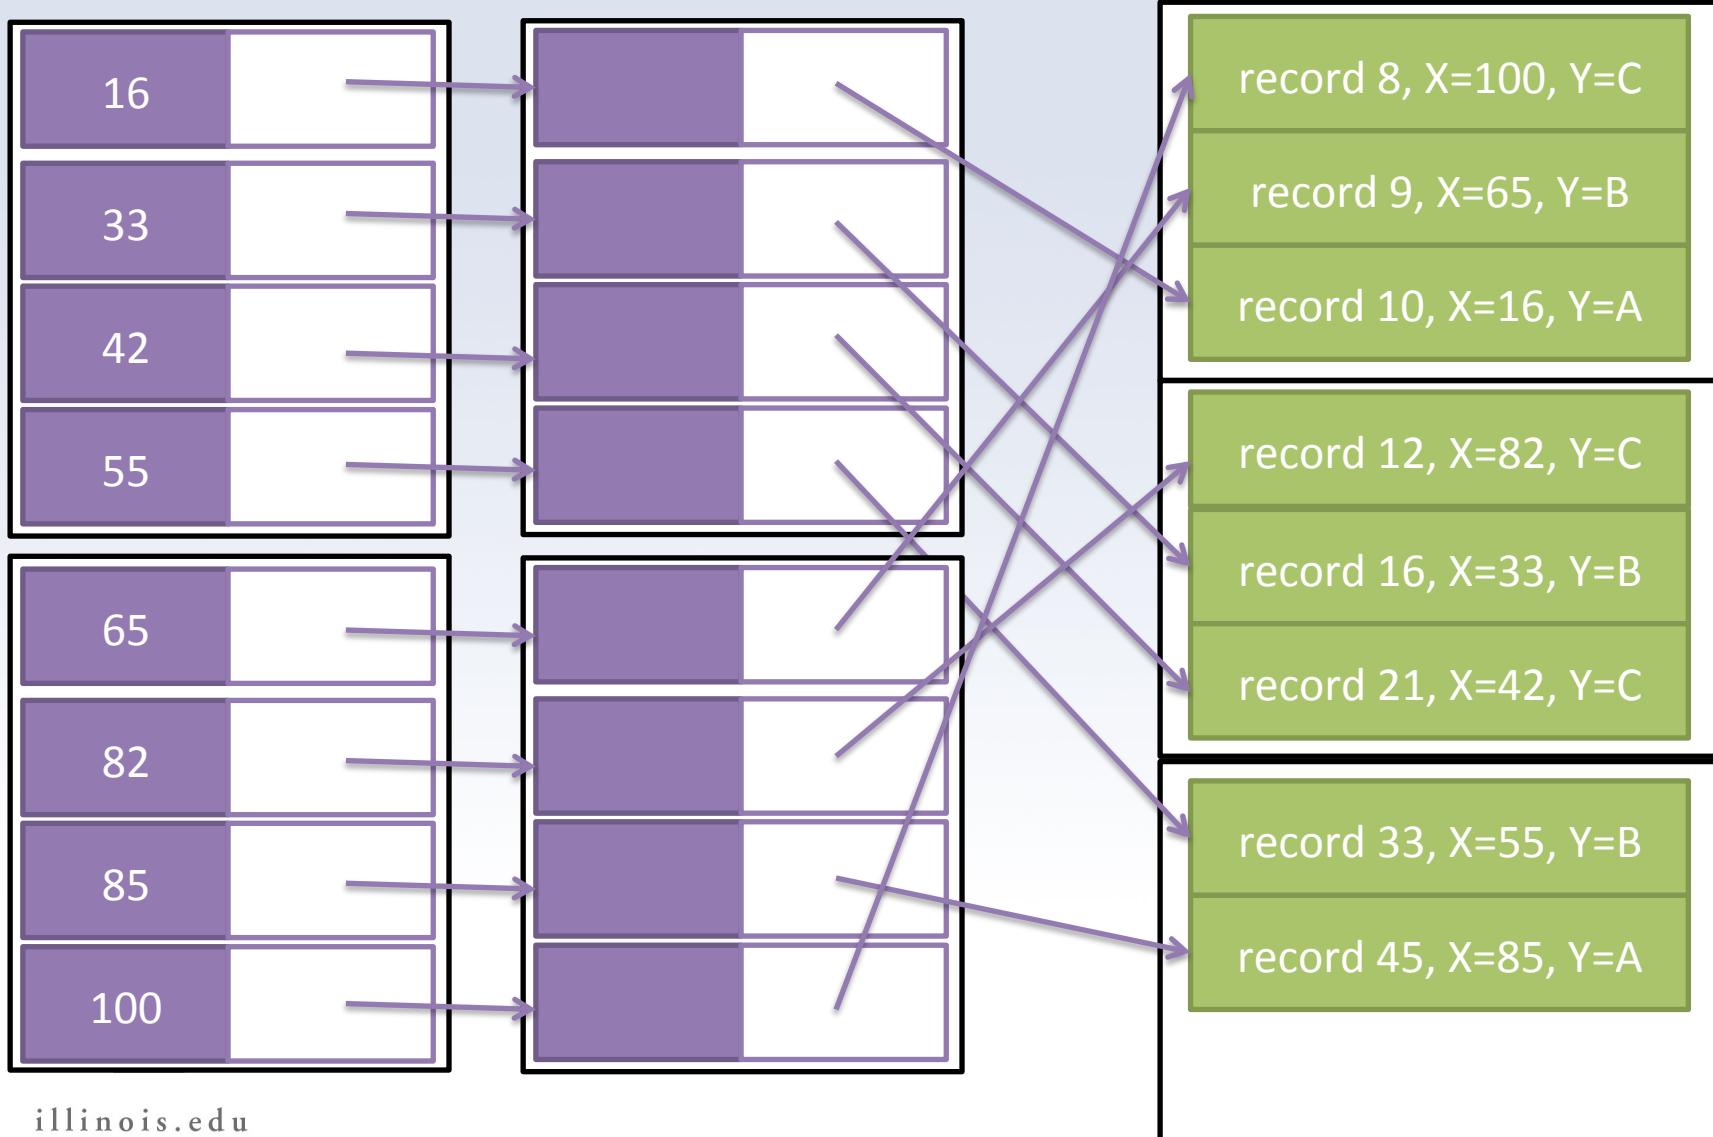
- A layer of indirection can improve a secondary index
 - Eliminates replication of search key values
 - Can be used to perform intersection
- We call this layer of indirection a ***bucket***
 - We call intersection ***bucket intersection***



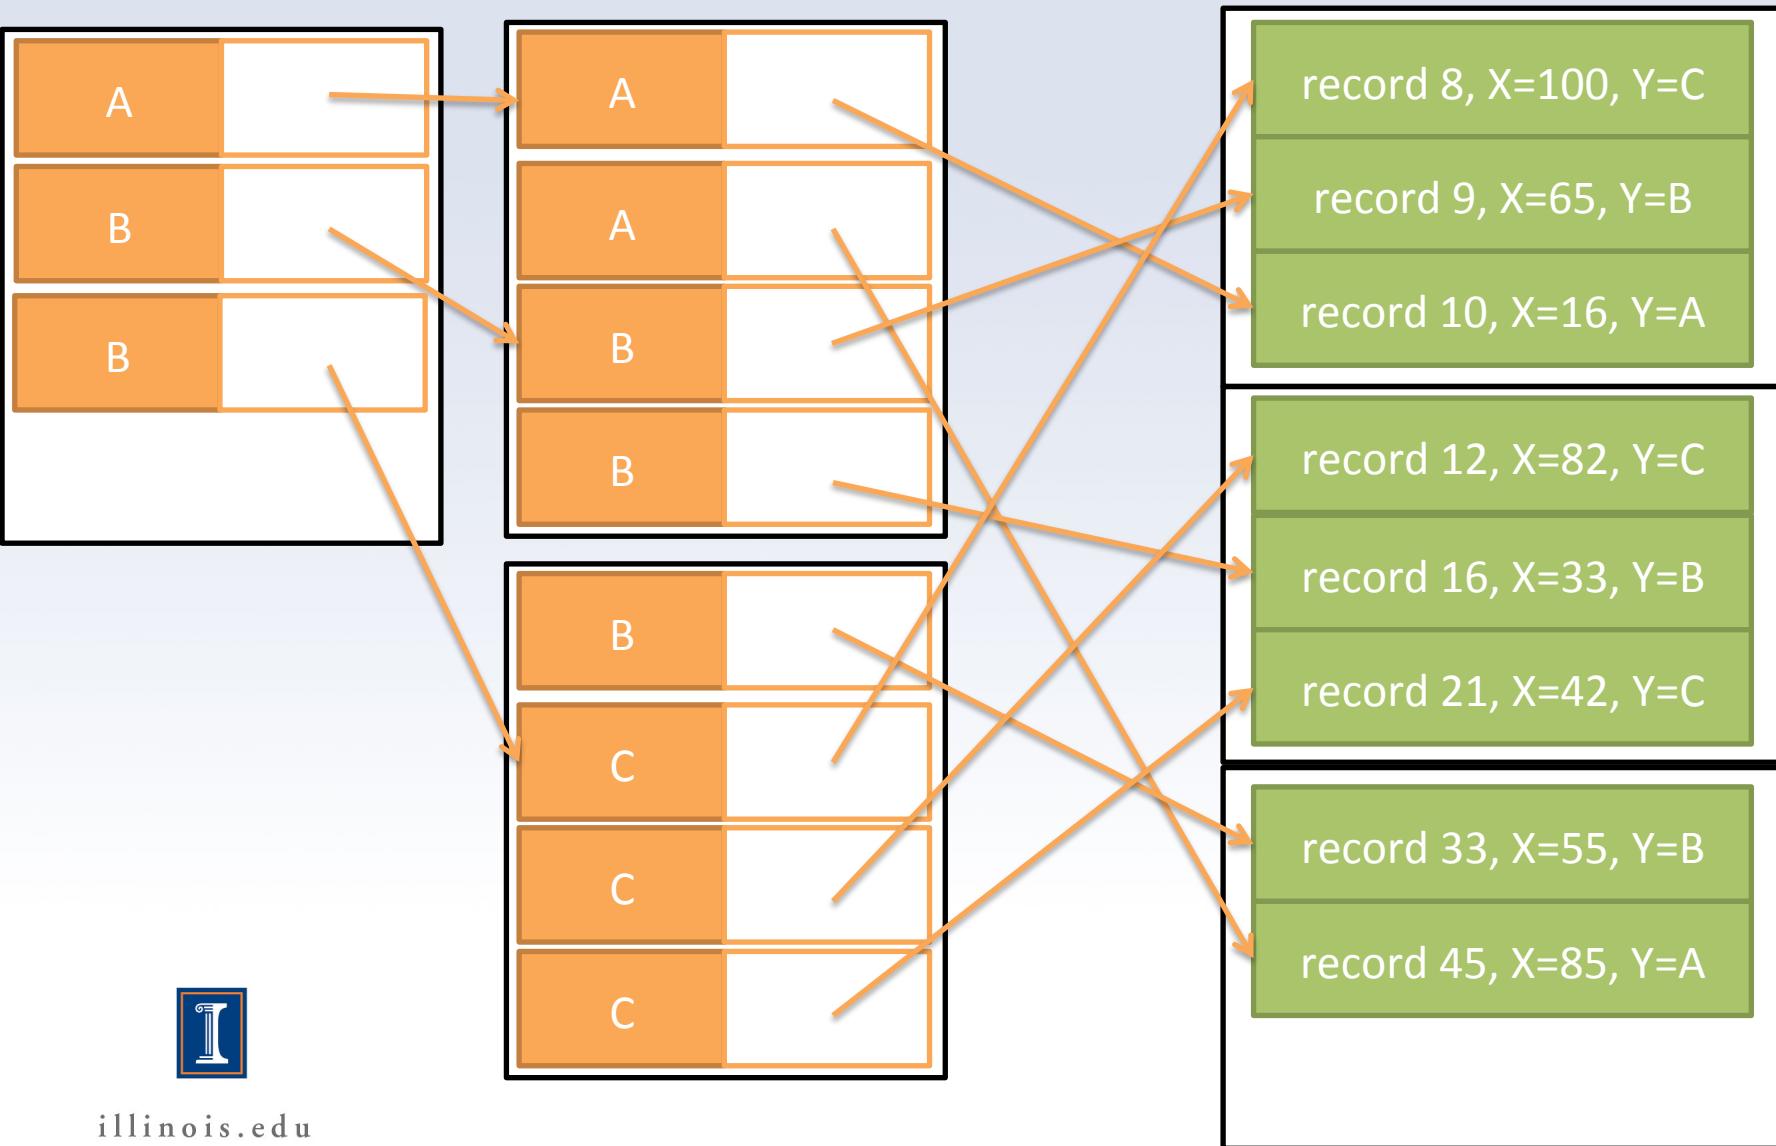
Indirection with Buckets



Indirection with Buckets



Indirection with Buckets



Bucket Intersection

```
SELECT *  
FROM R  
WHERE X=85 AND Y=A
```

record 8, X=100, Y=C

record 9, X=65, Y=B

record 10, X=16, Y=A

record 12, X=82, Y=C

record 16, X=33, Y=B

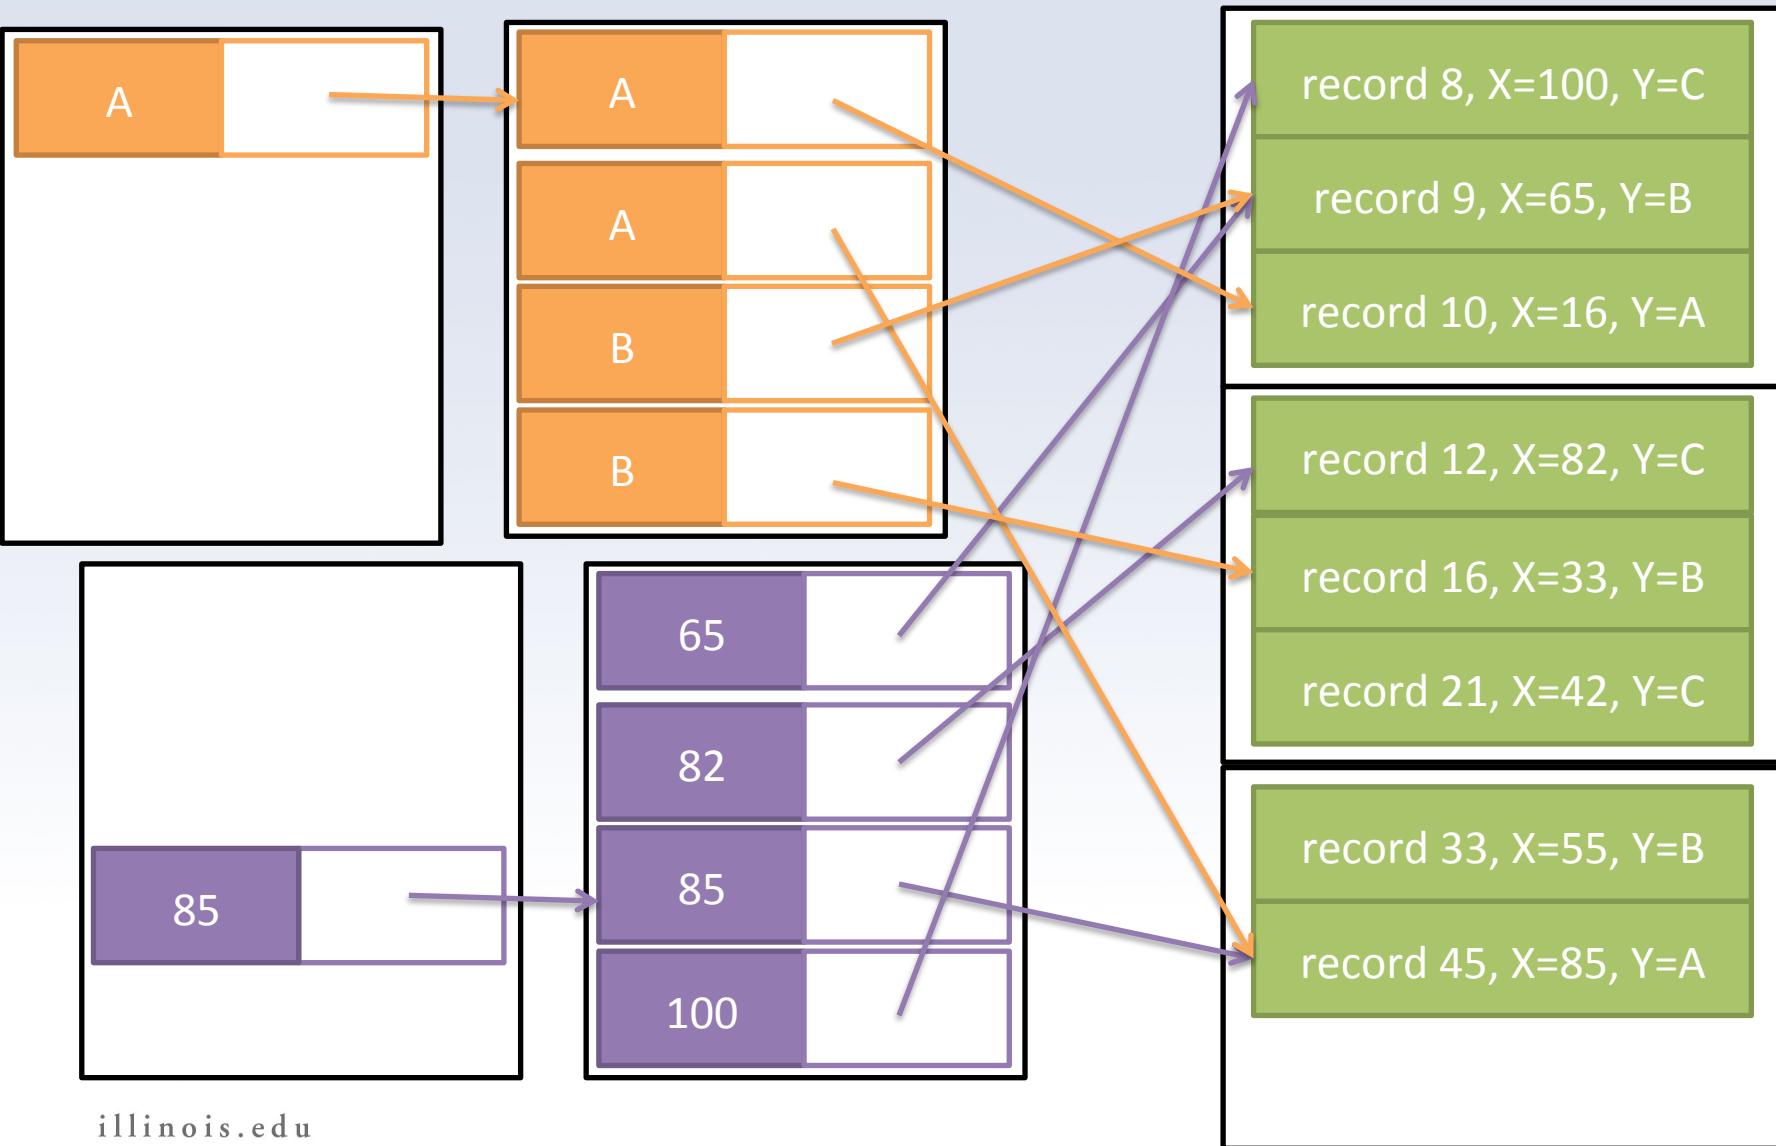
record 21, X=42, Y=C

record 33, X=55, Y=B

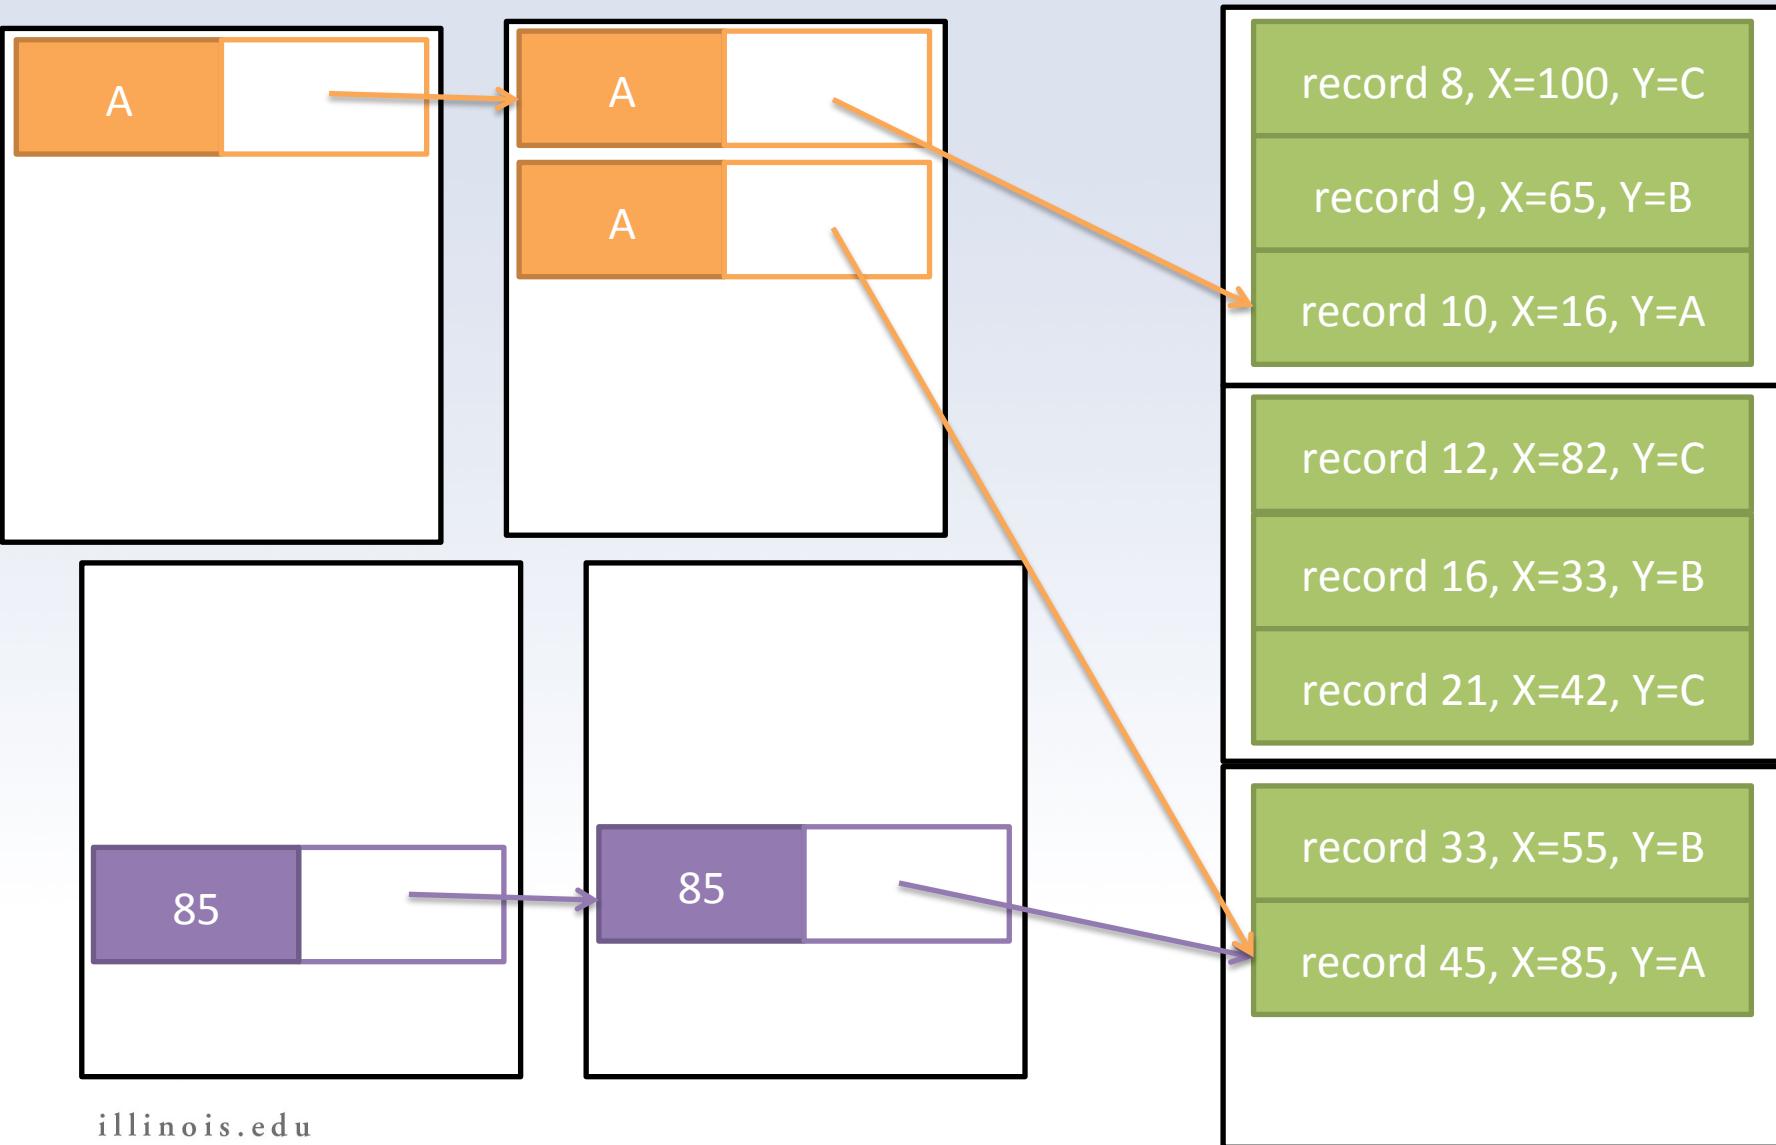
record 45, X=85, Y=A



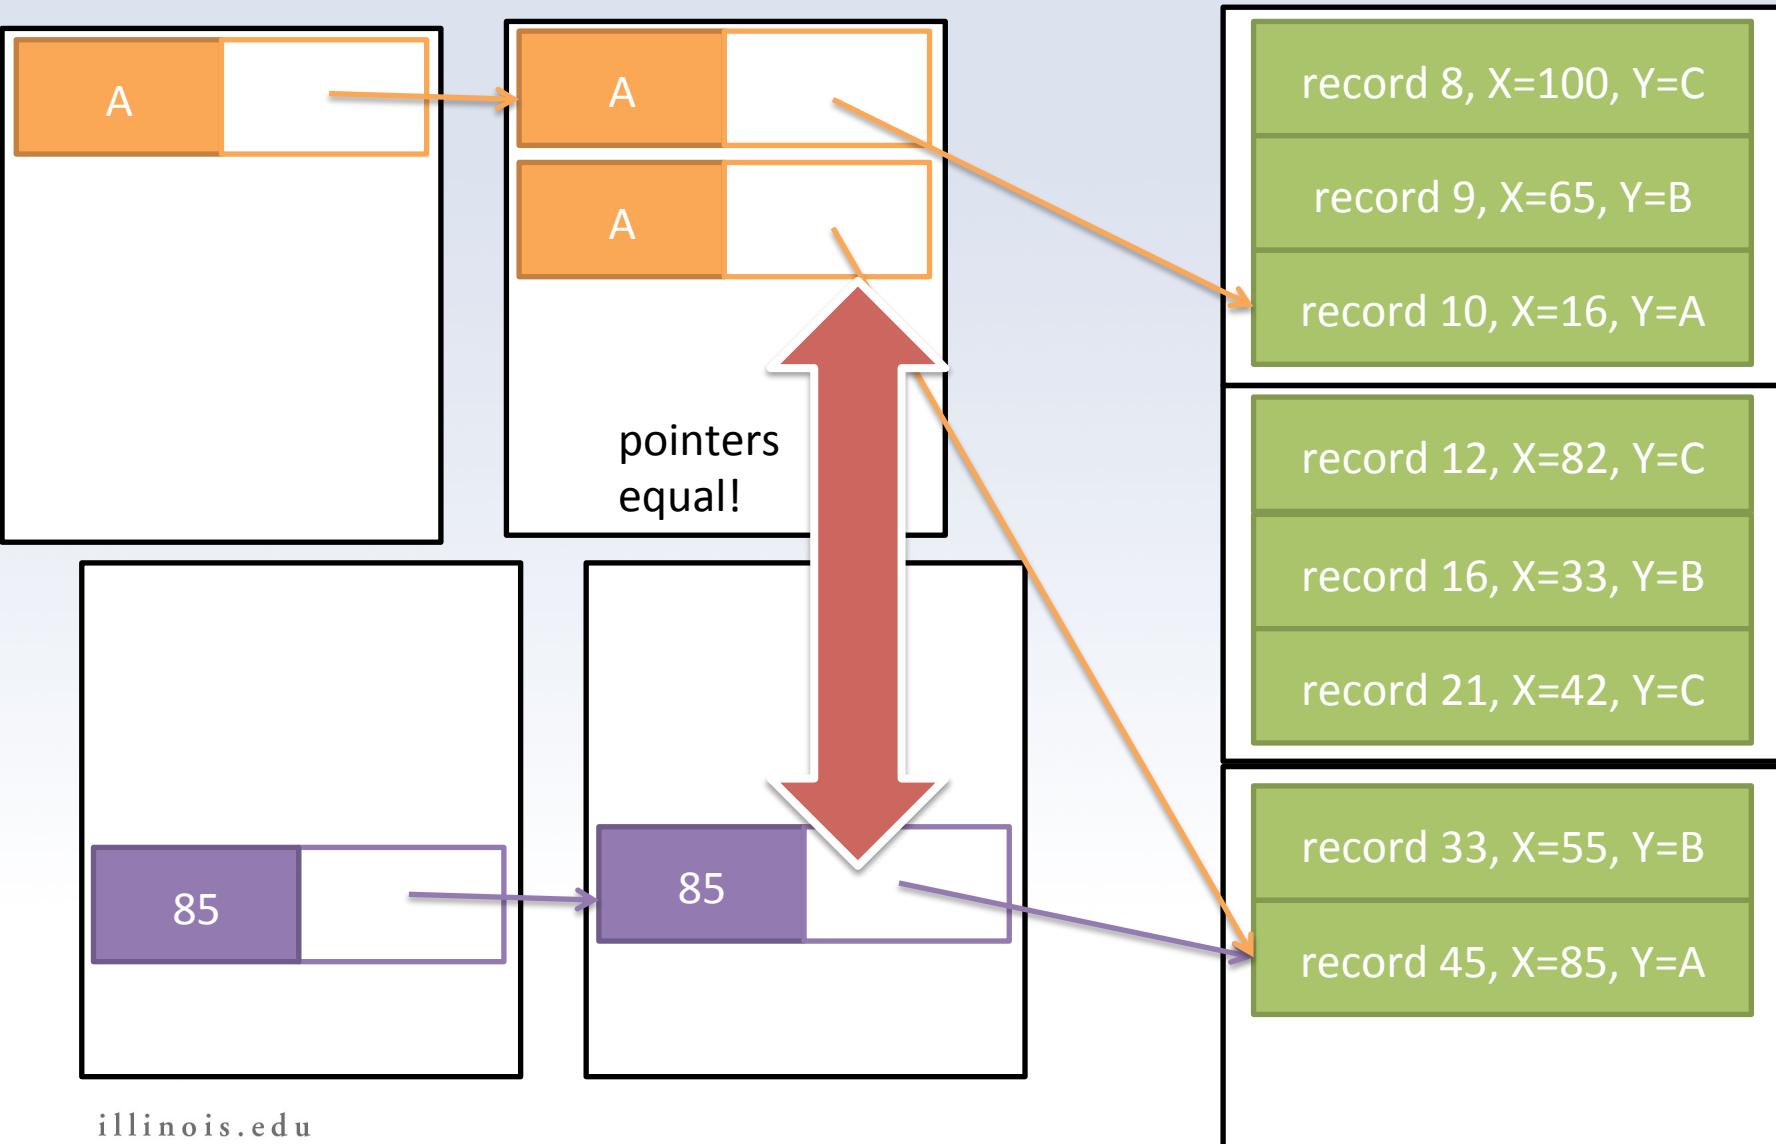
Bucket Intersection



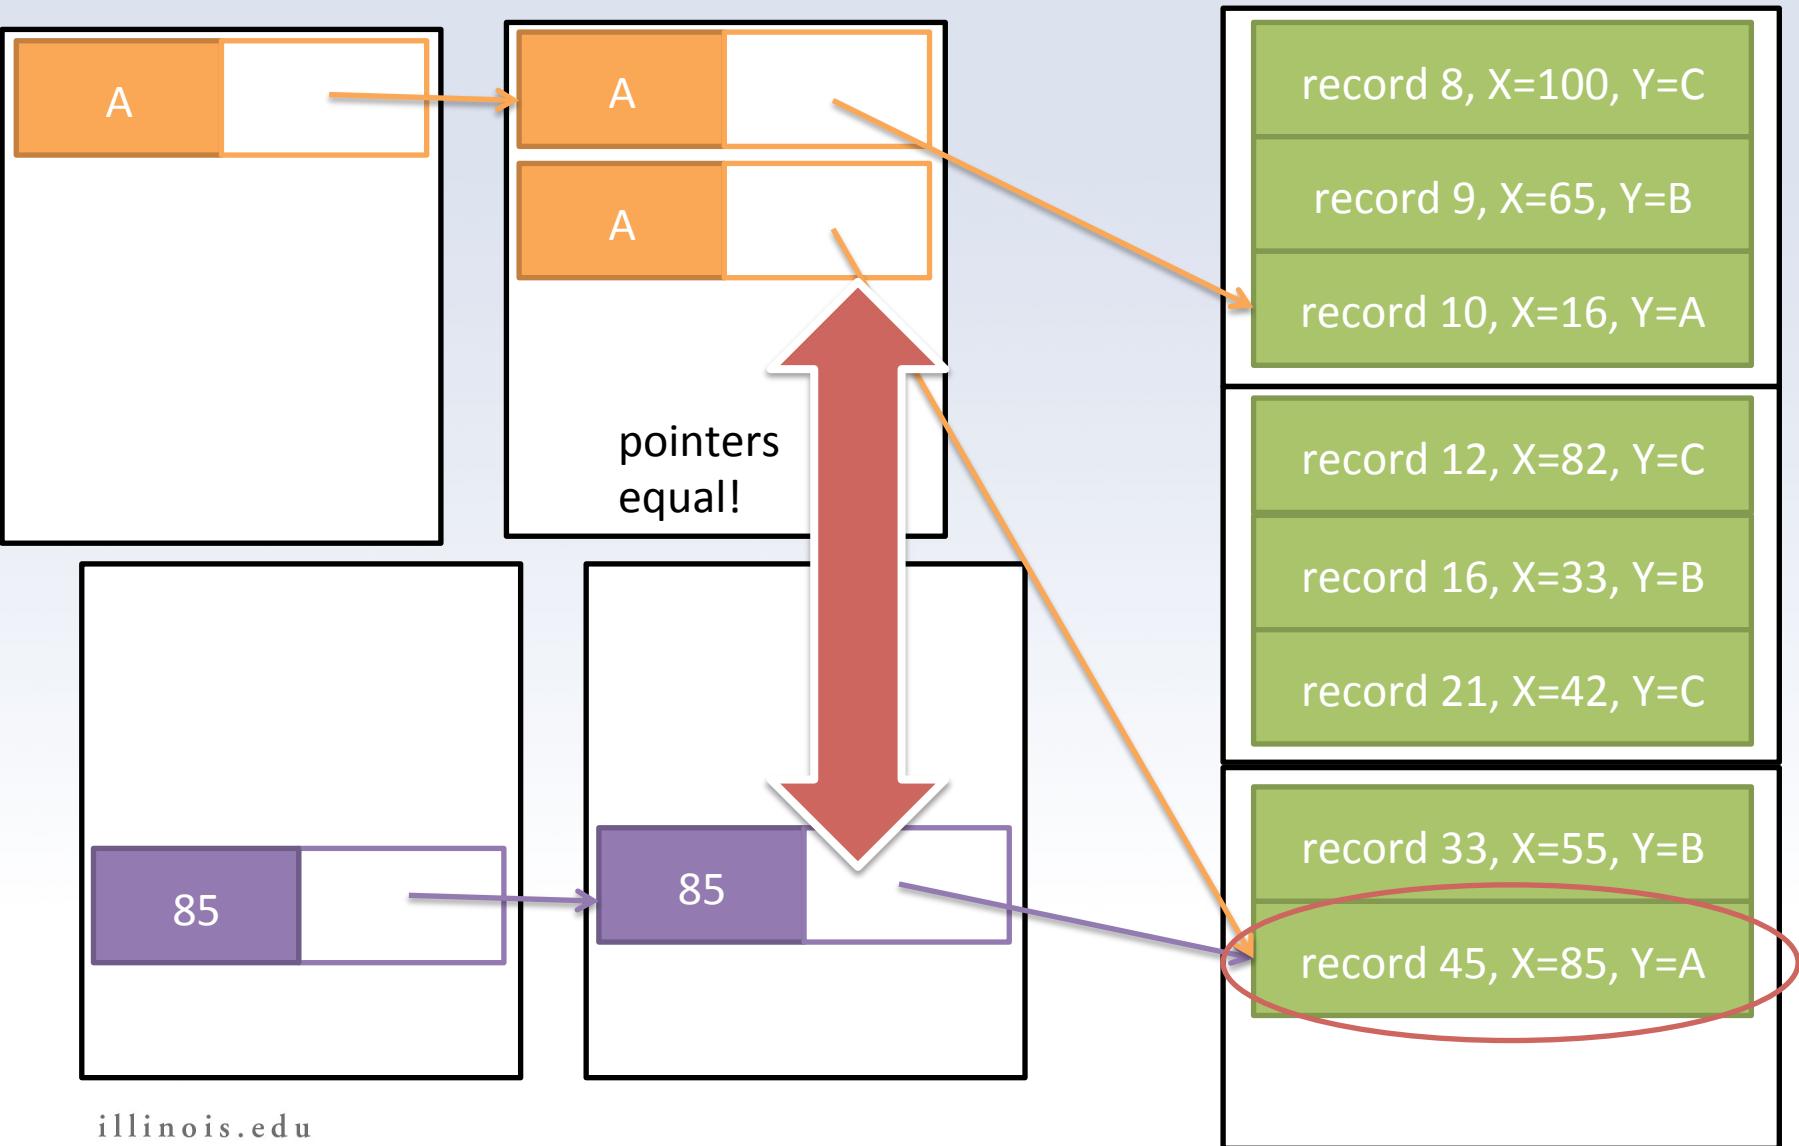
Bucket Intersection



Bucket Intersection



Bucket Intersection



B-trees

- General purpose multi-level indexing structure
 - automatically maintains appropriate number of levels
 - automatically maintains space usage in blocks so they are at least half full
 - automatically *balanced*; all paths from root to leaves are equal



B-trees

- There are many variants of the B-tree
 - What we will call a “B-tree” is actually a “B+ tree”
 - This is the only variation we will discuss, so we’ll just refer to it as a “B-tree” from now on



B-trees

- Has a single parameter n
 - Each node in the tree has n search key values
 - Each node also has $n+1$ pointers
 - We pick n to fill up a block as best we can



Example

- The search key is a 4 byte integer
- Pointers are 8 bytes
- Pages are 4096 bytes
- Find largest n such that $4n+8(n+1) \leq 4096$
- $12n+8 \leq 4096$
- $n=340$



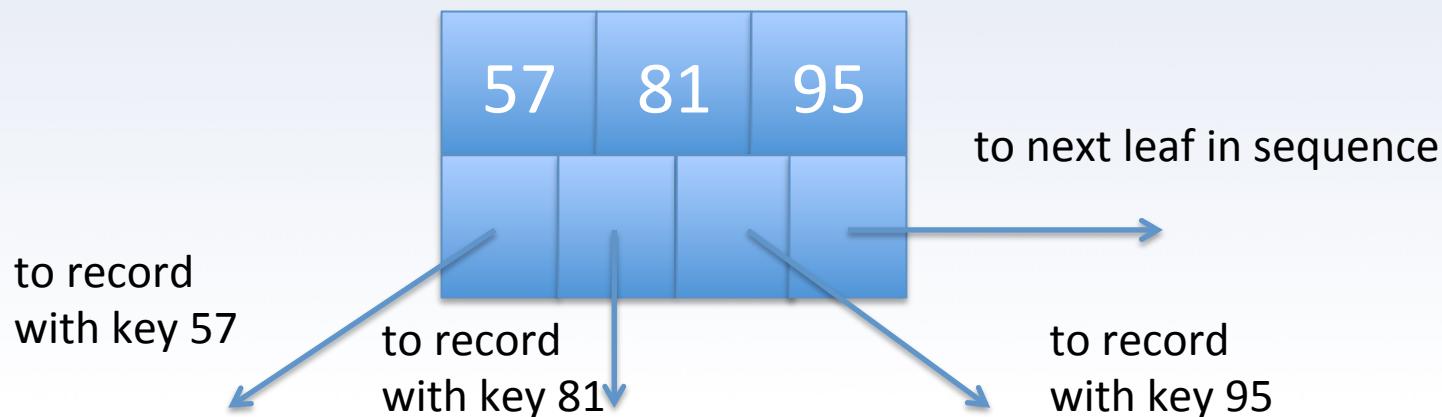
B-trees

- In our running example
 - Search key will be a single integer
 - we will use the incredibly small parameter setting $n=3$



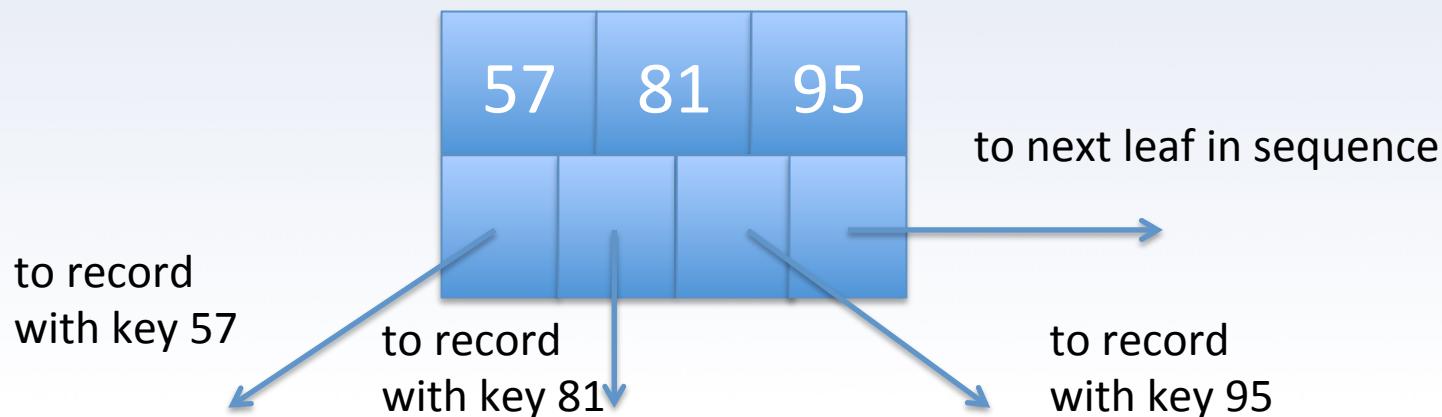
Leaf nodes

- Store key values and pointers to the actual record
- Also keep a pointer to the next leaf



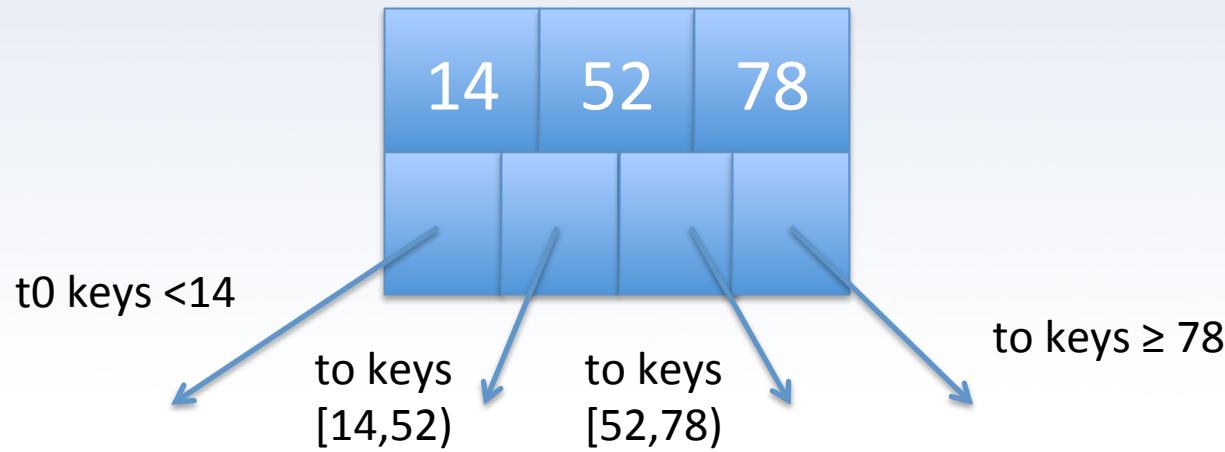
Leaf nodes

- At least $\lfloor (n + 1)/2 \rfloor$ of the pointers must be used



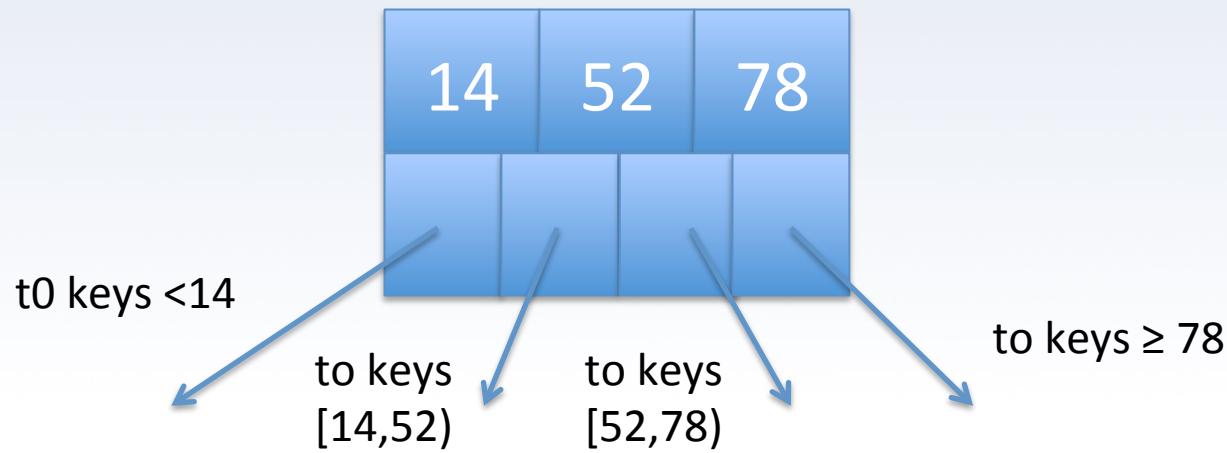
Interior nodes

- Contain pointers to blocks at lower level
- Key values indicate range of values in lower levels

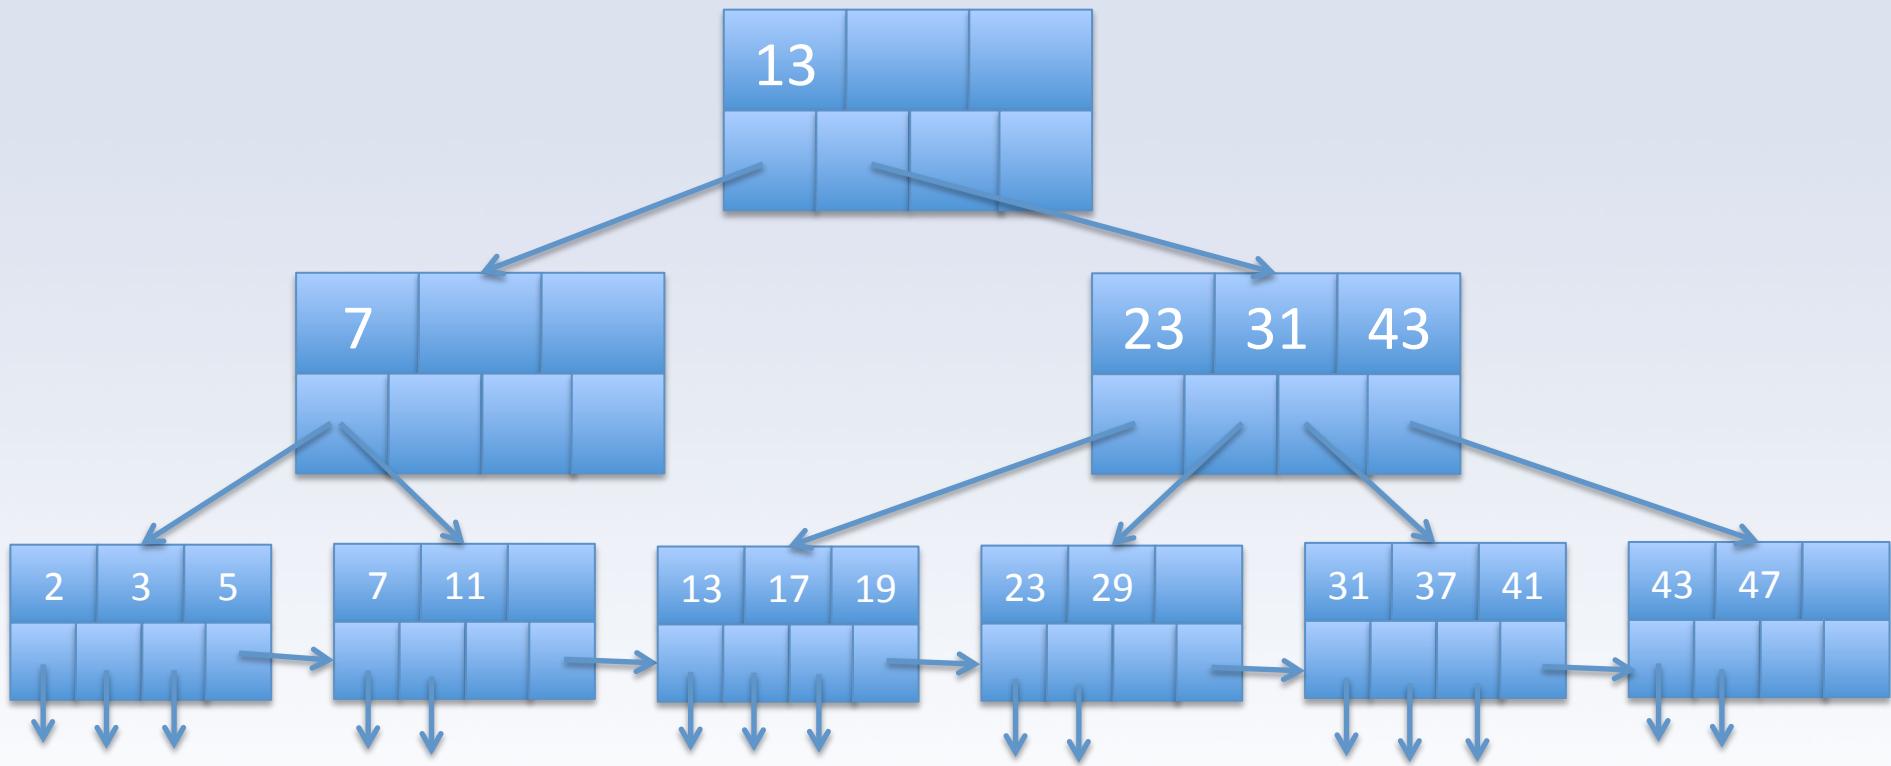


Interior nodes

- At least $\lceil (n + 1)/2 \rceil$ of the pointers must be used
- Exception: root must have at least 2 children



Example B-tree

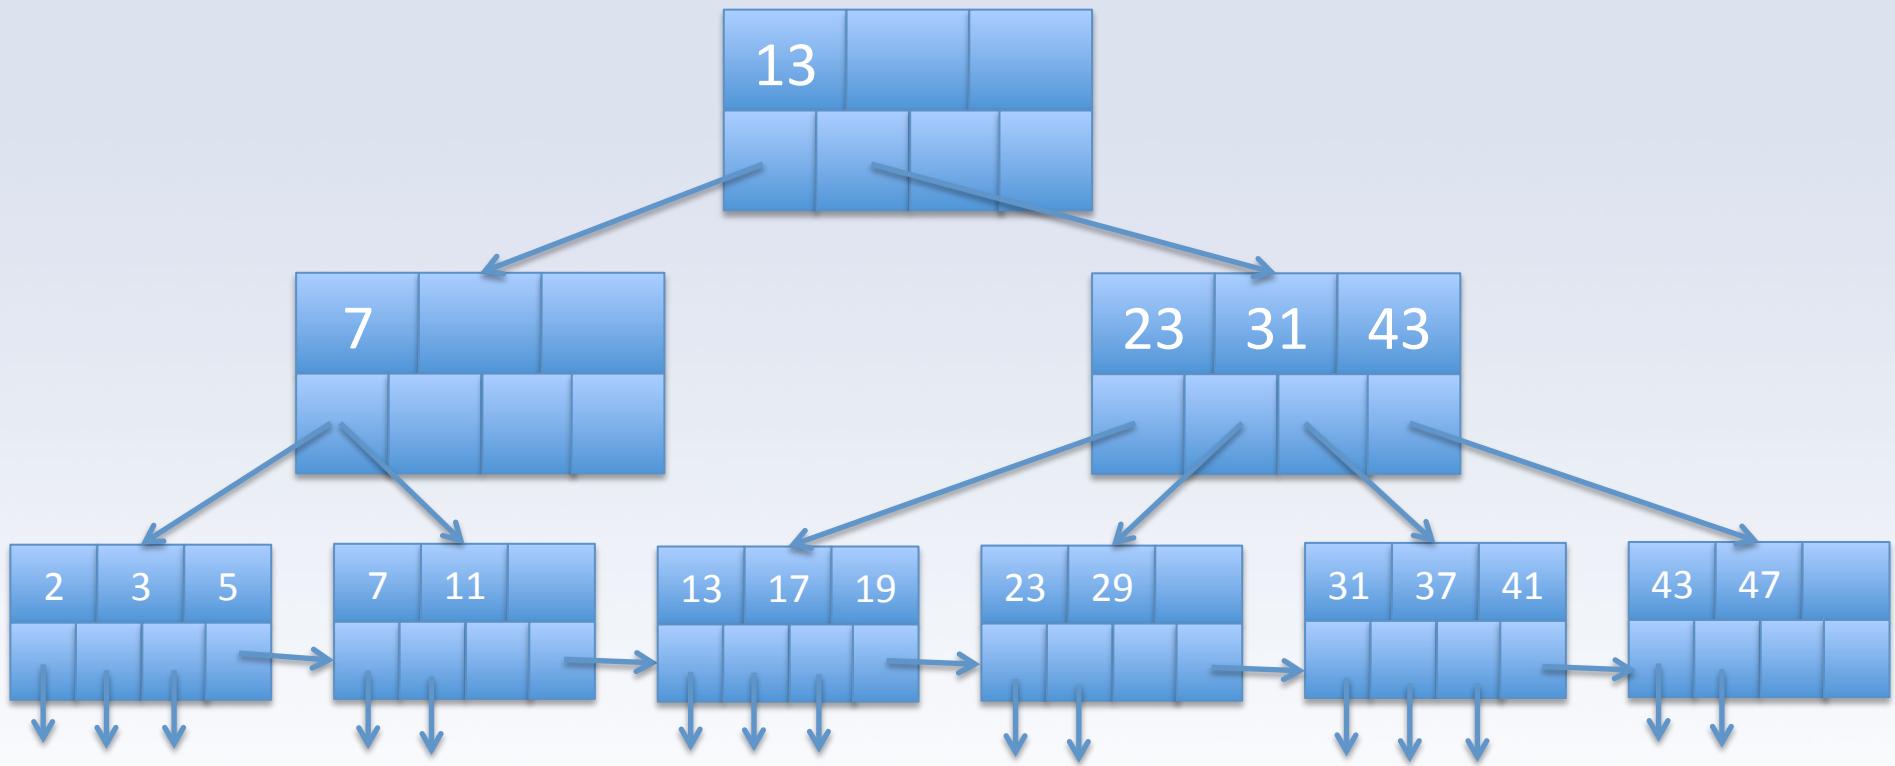


Lookup

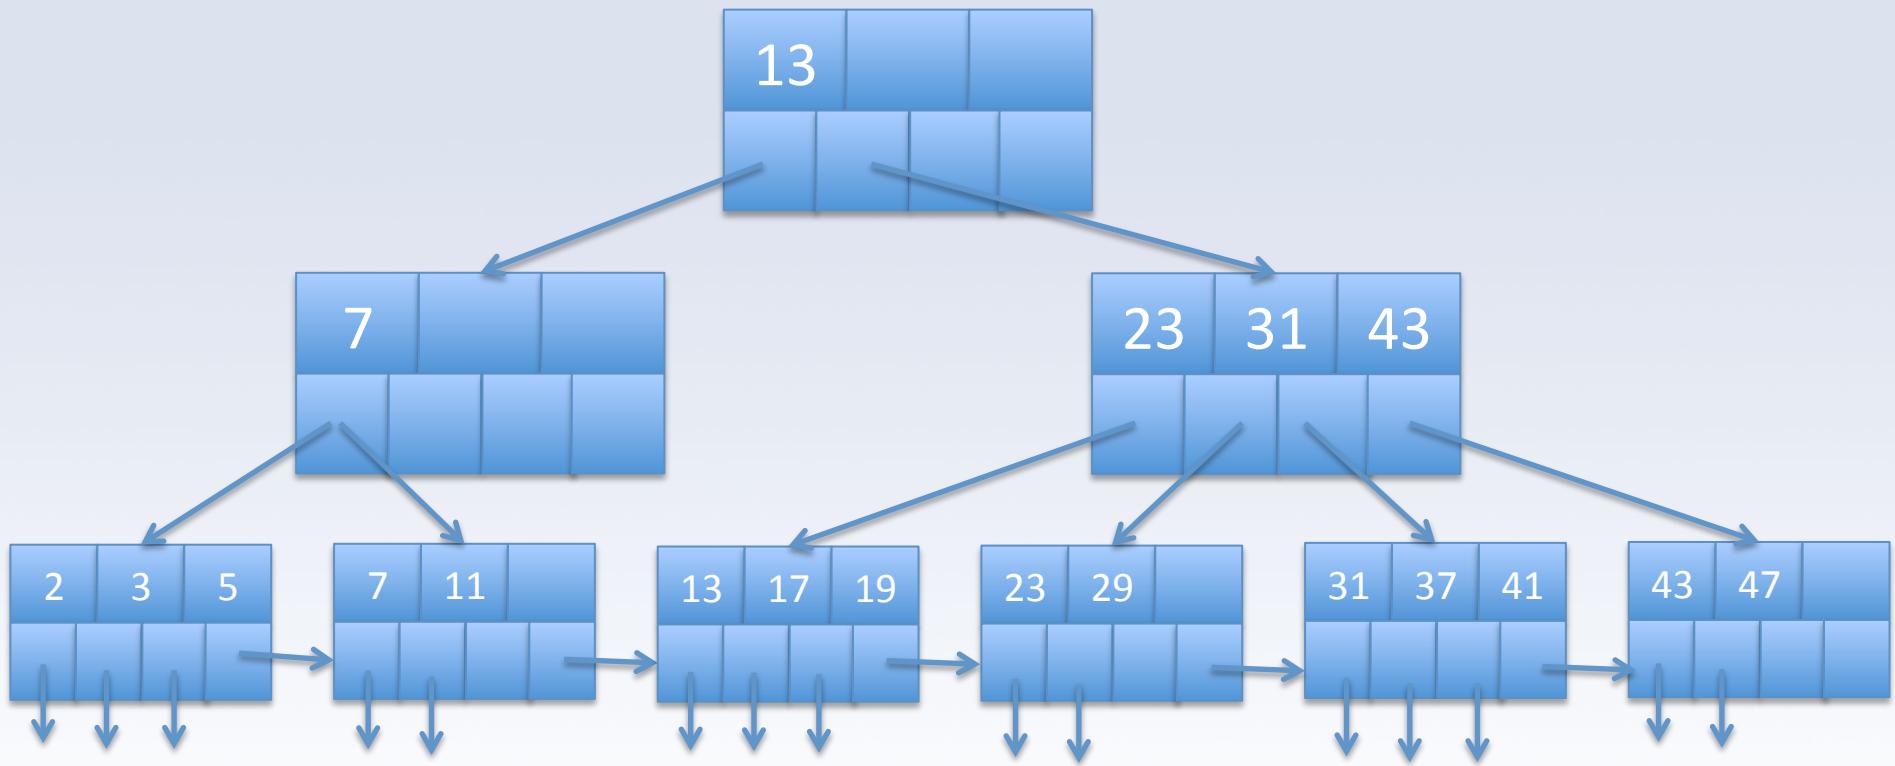
- To find a key K in the tree:
 - Start at root
 - Internal nodes: search keys to find range K belongs in and follow that pointer
 - Leaf nodes: search keys to find K and follow pointer



Lookup 40



Lookup 17



Range lookup

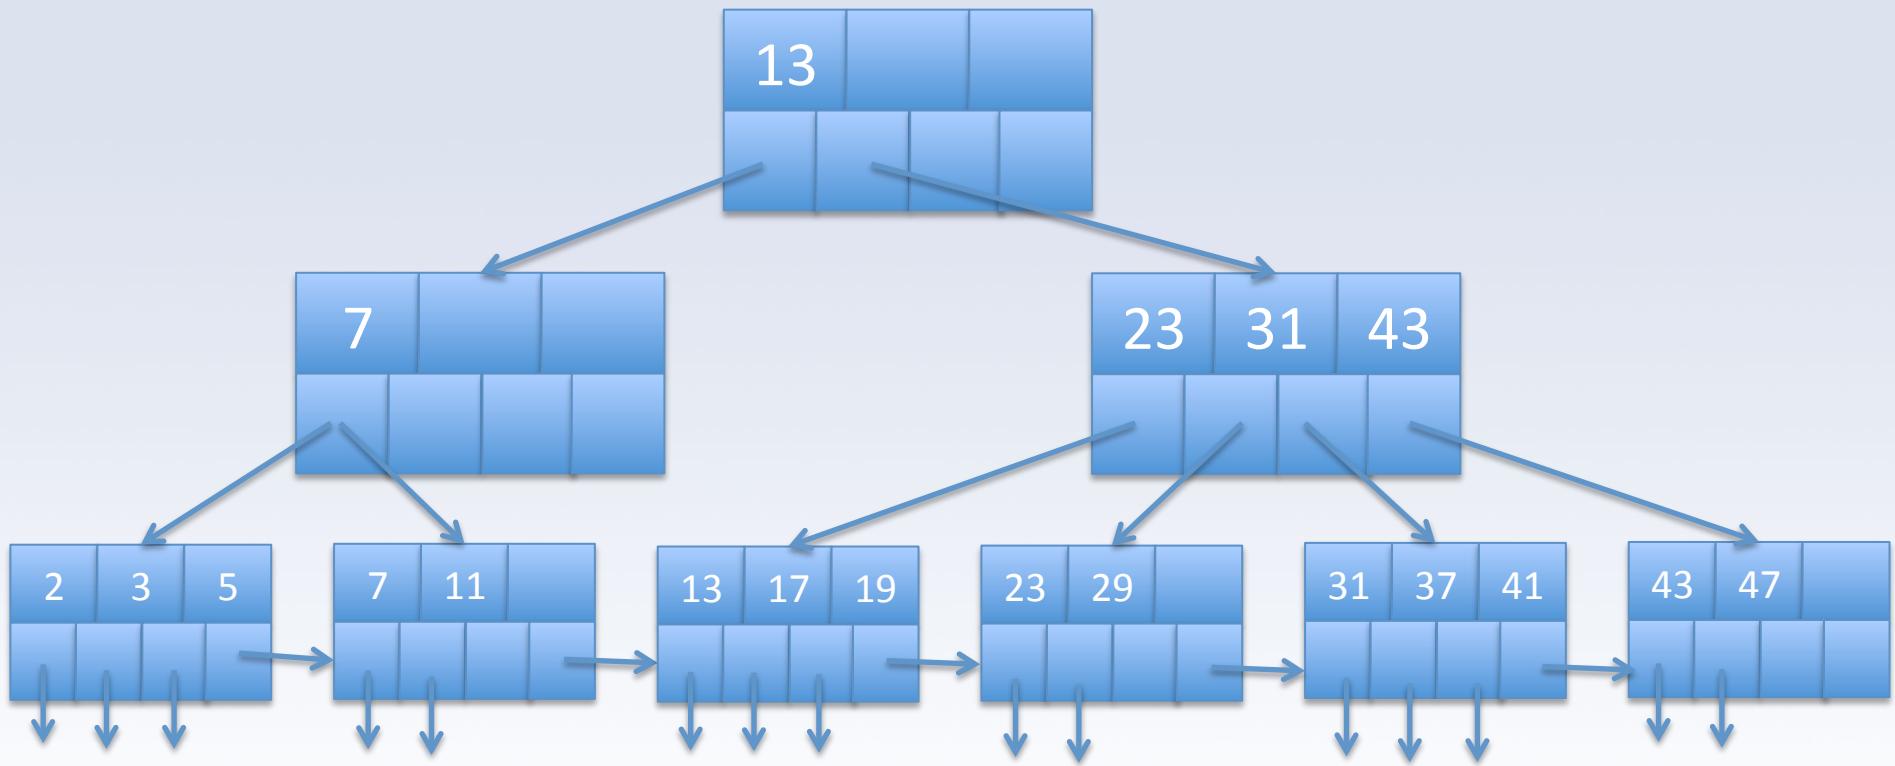
- What about queries like this

```
SELECT * FROM R
```

```
WHERE key > 10 AND k < 25;
```



Lookup (10,25)

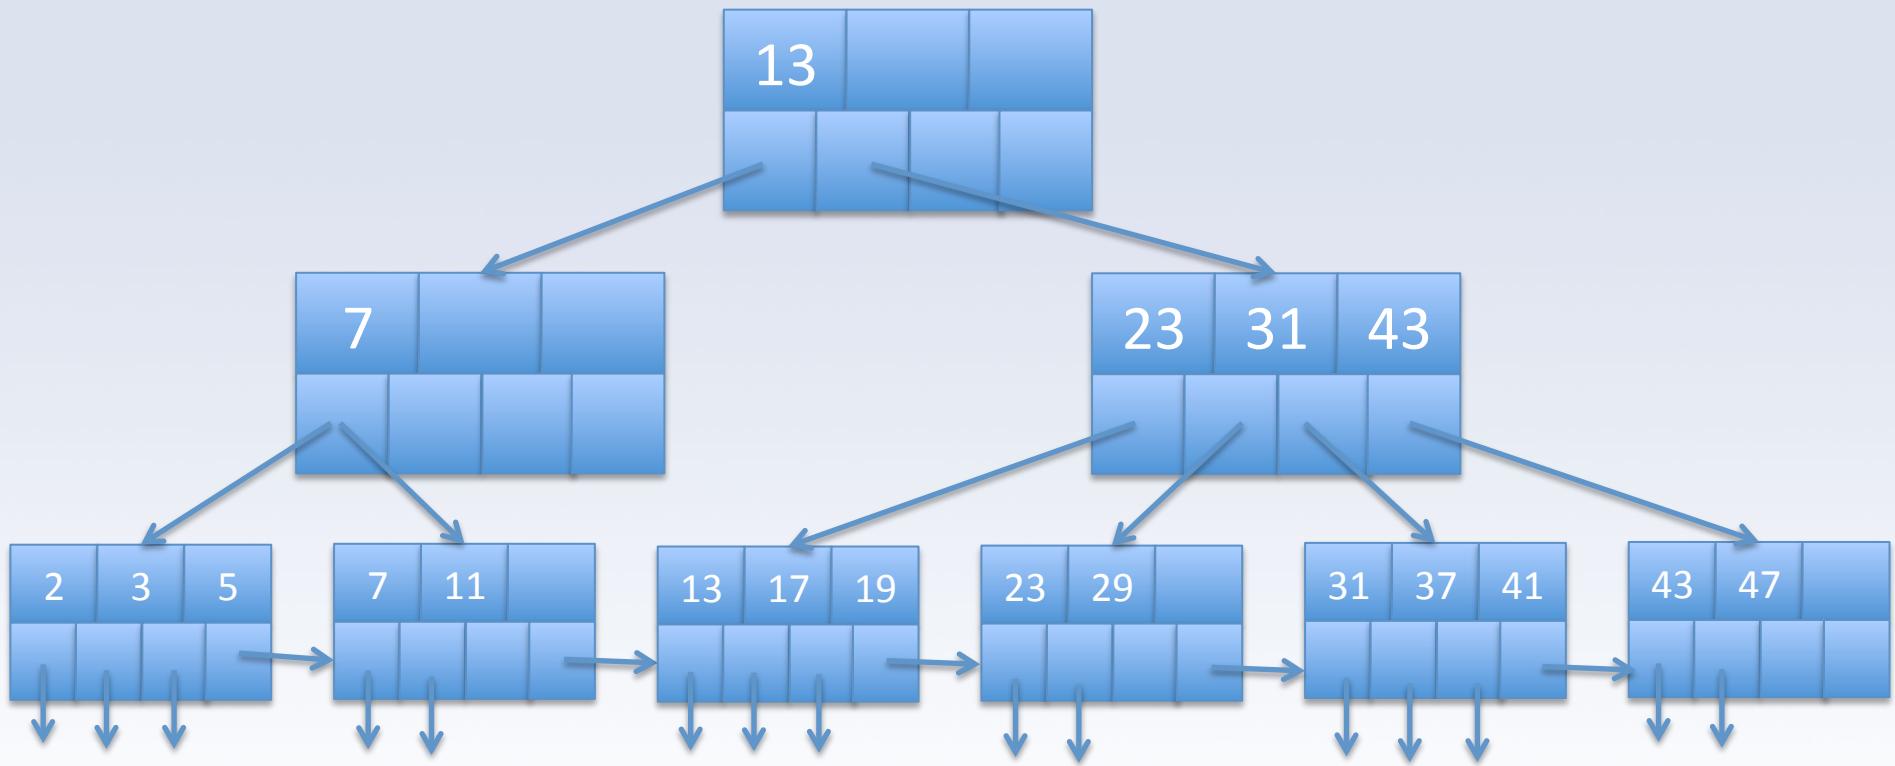


Insertion

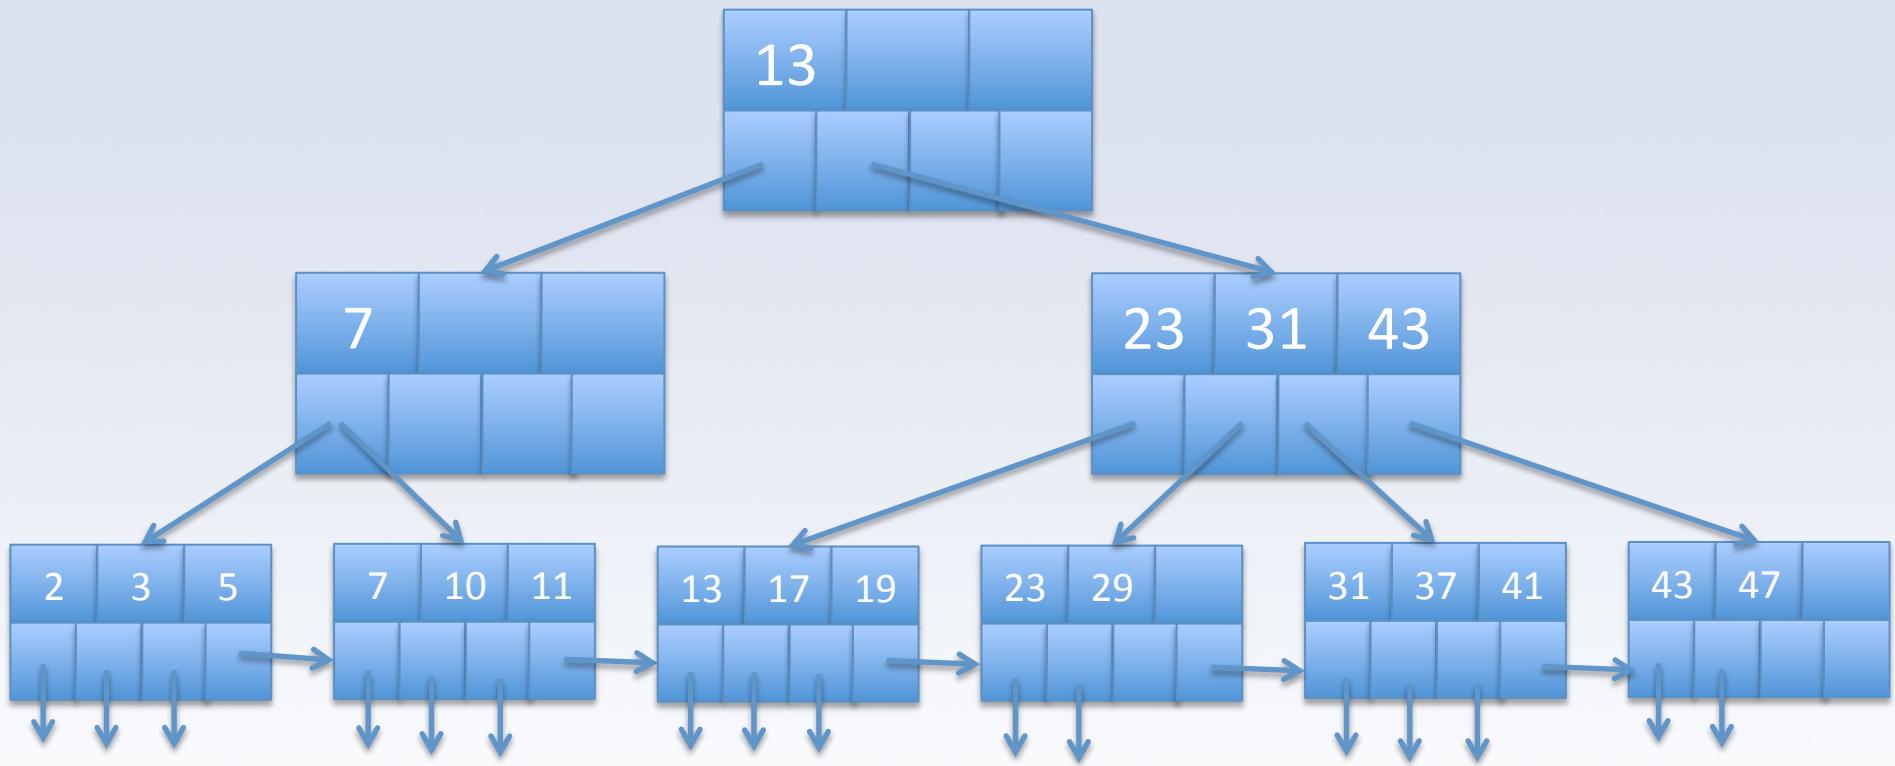
- To insert K into a B tree
 - Lookup appropriate leaf node L for K
 - Insert K into L with parent P .
 - If overflow, create new leaf N . Move half to L and half to N . Add N 's smallest key to P .
 - If P overflows, split it. Add middle key to P 's parent and recurse.
 - If P is root, split it and create a new root.



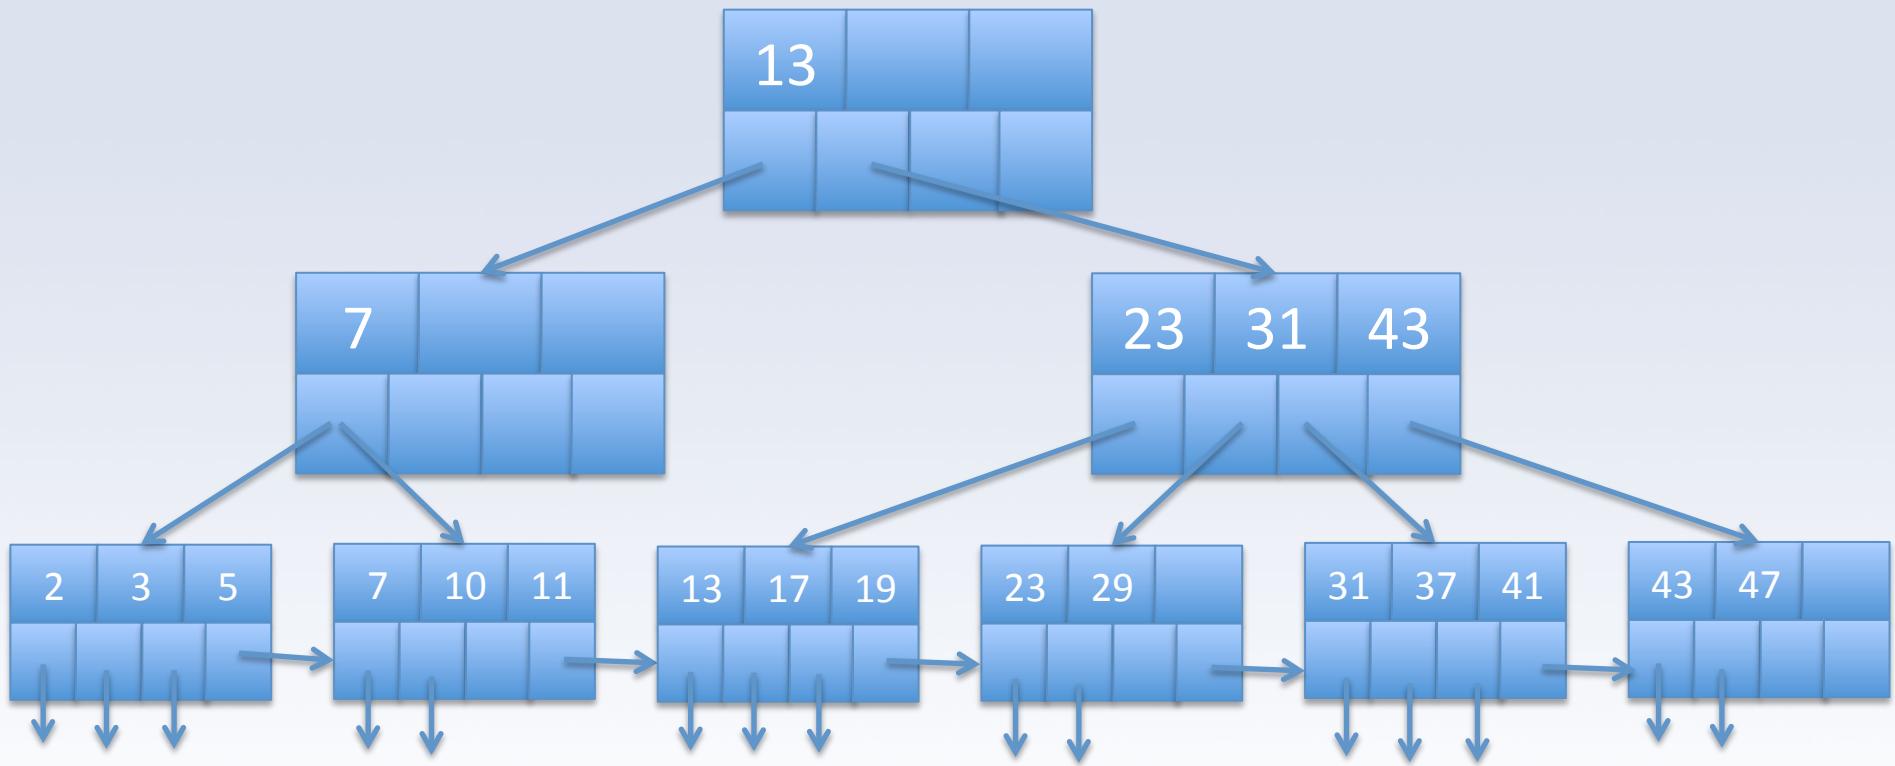
Insert 10



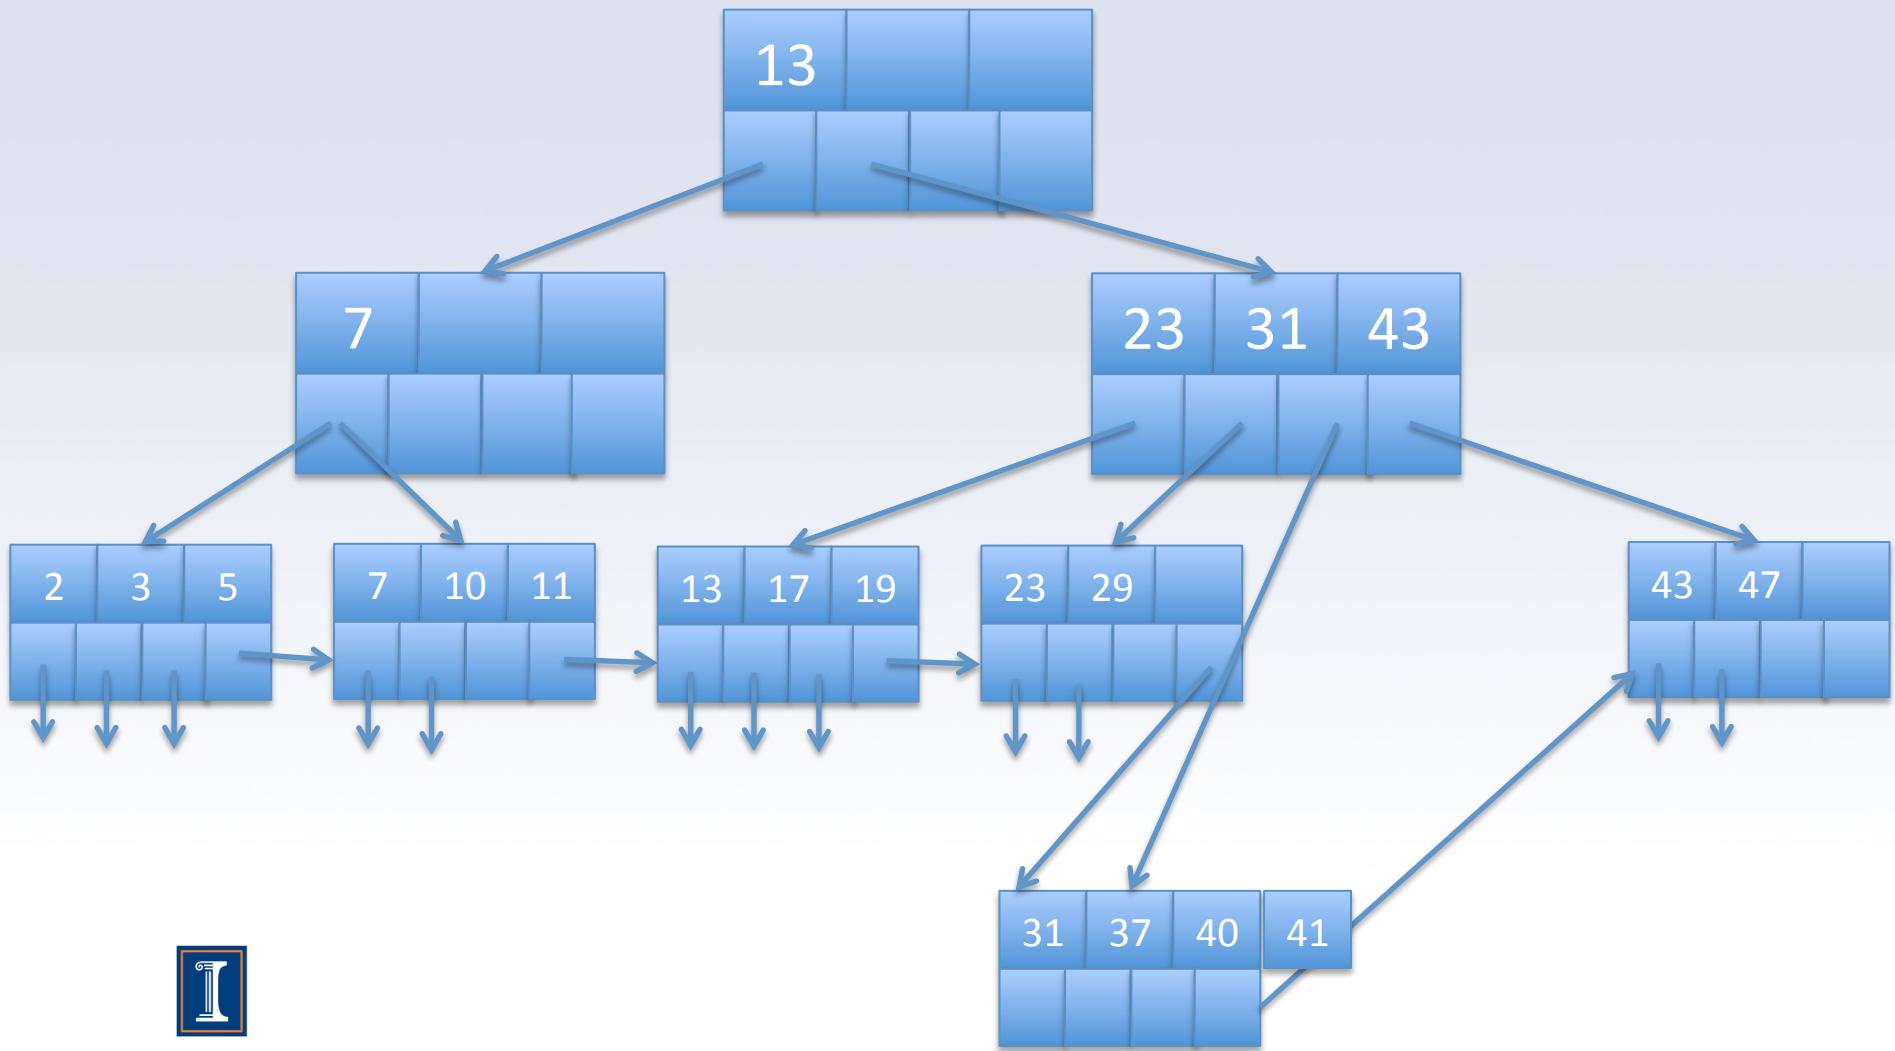
Insert 10



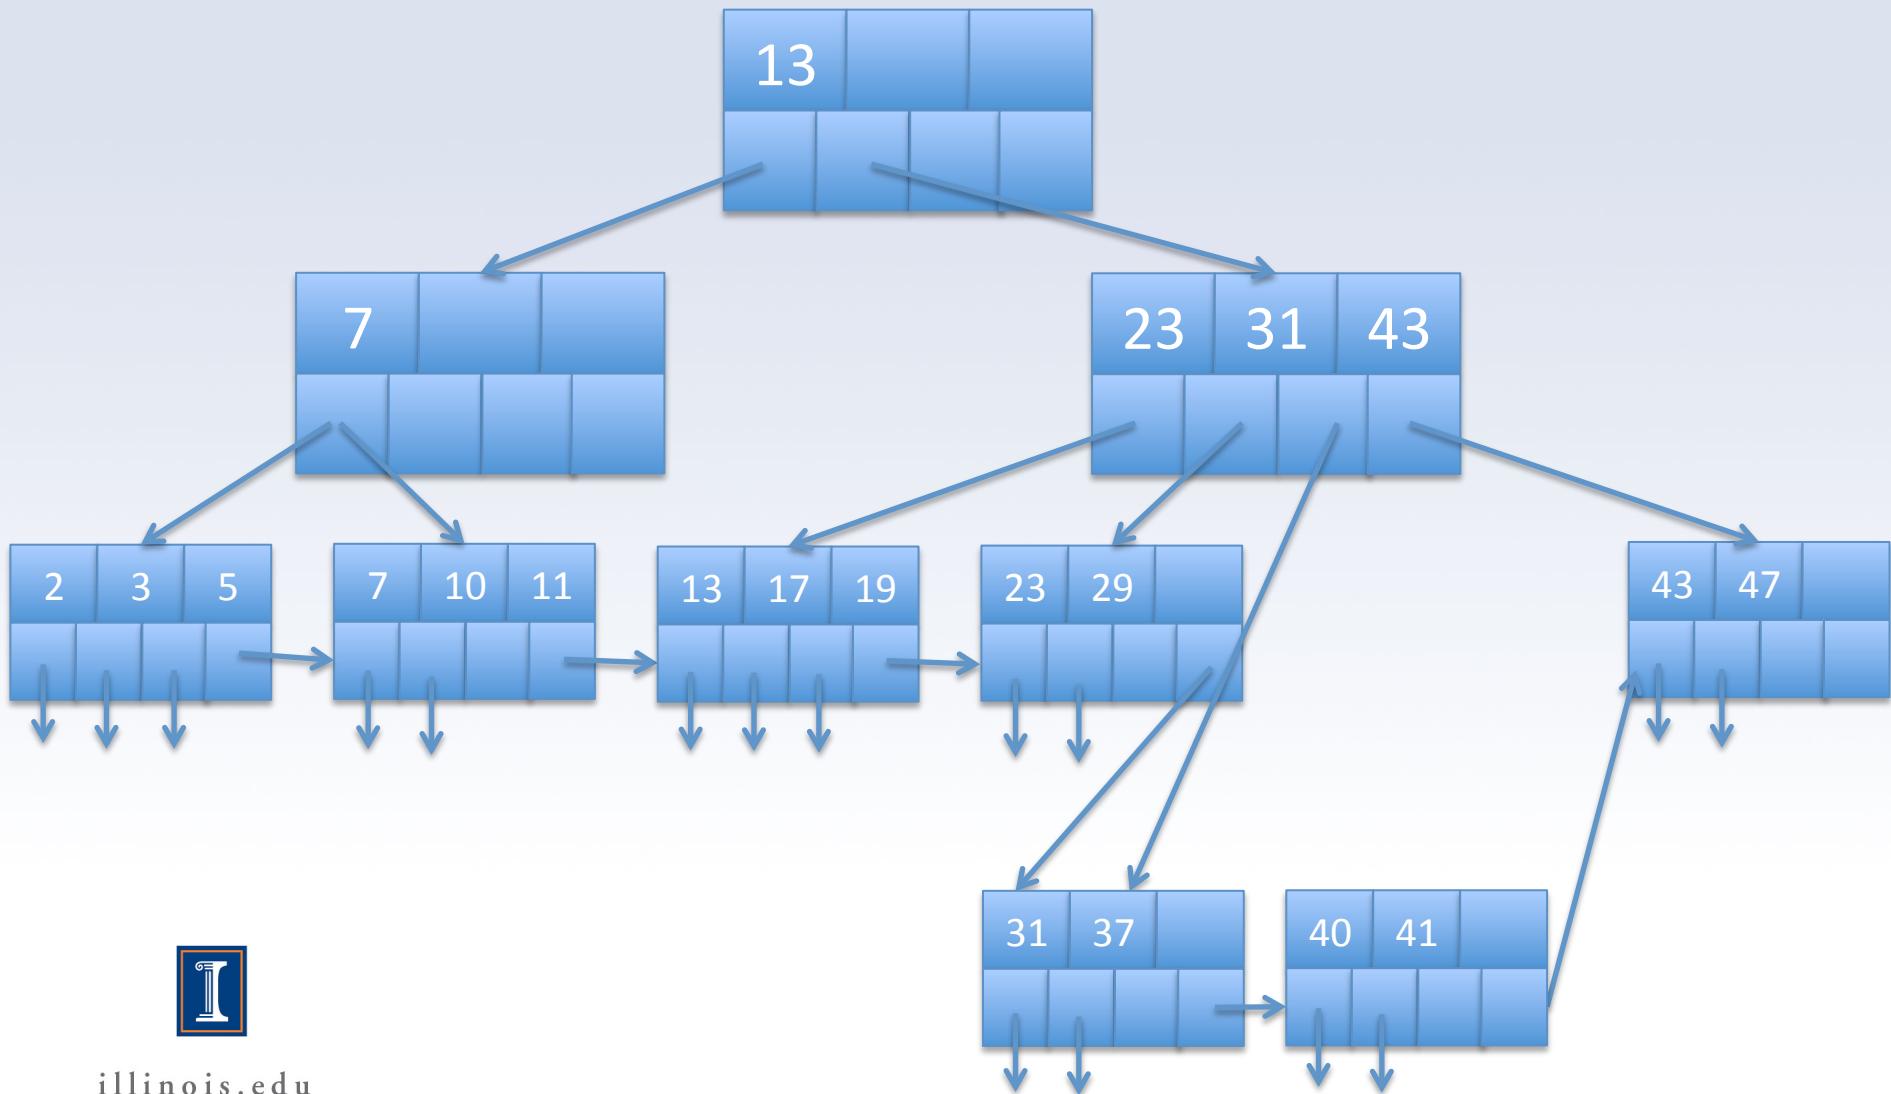
Insert 40



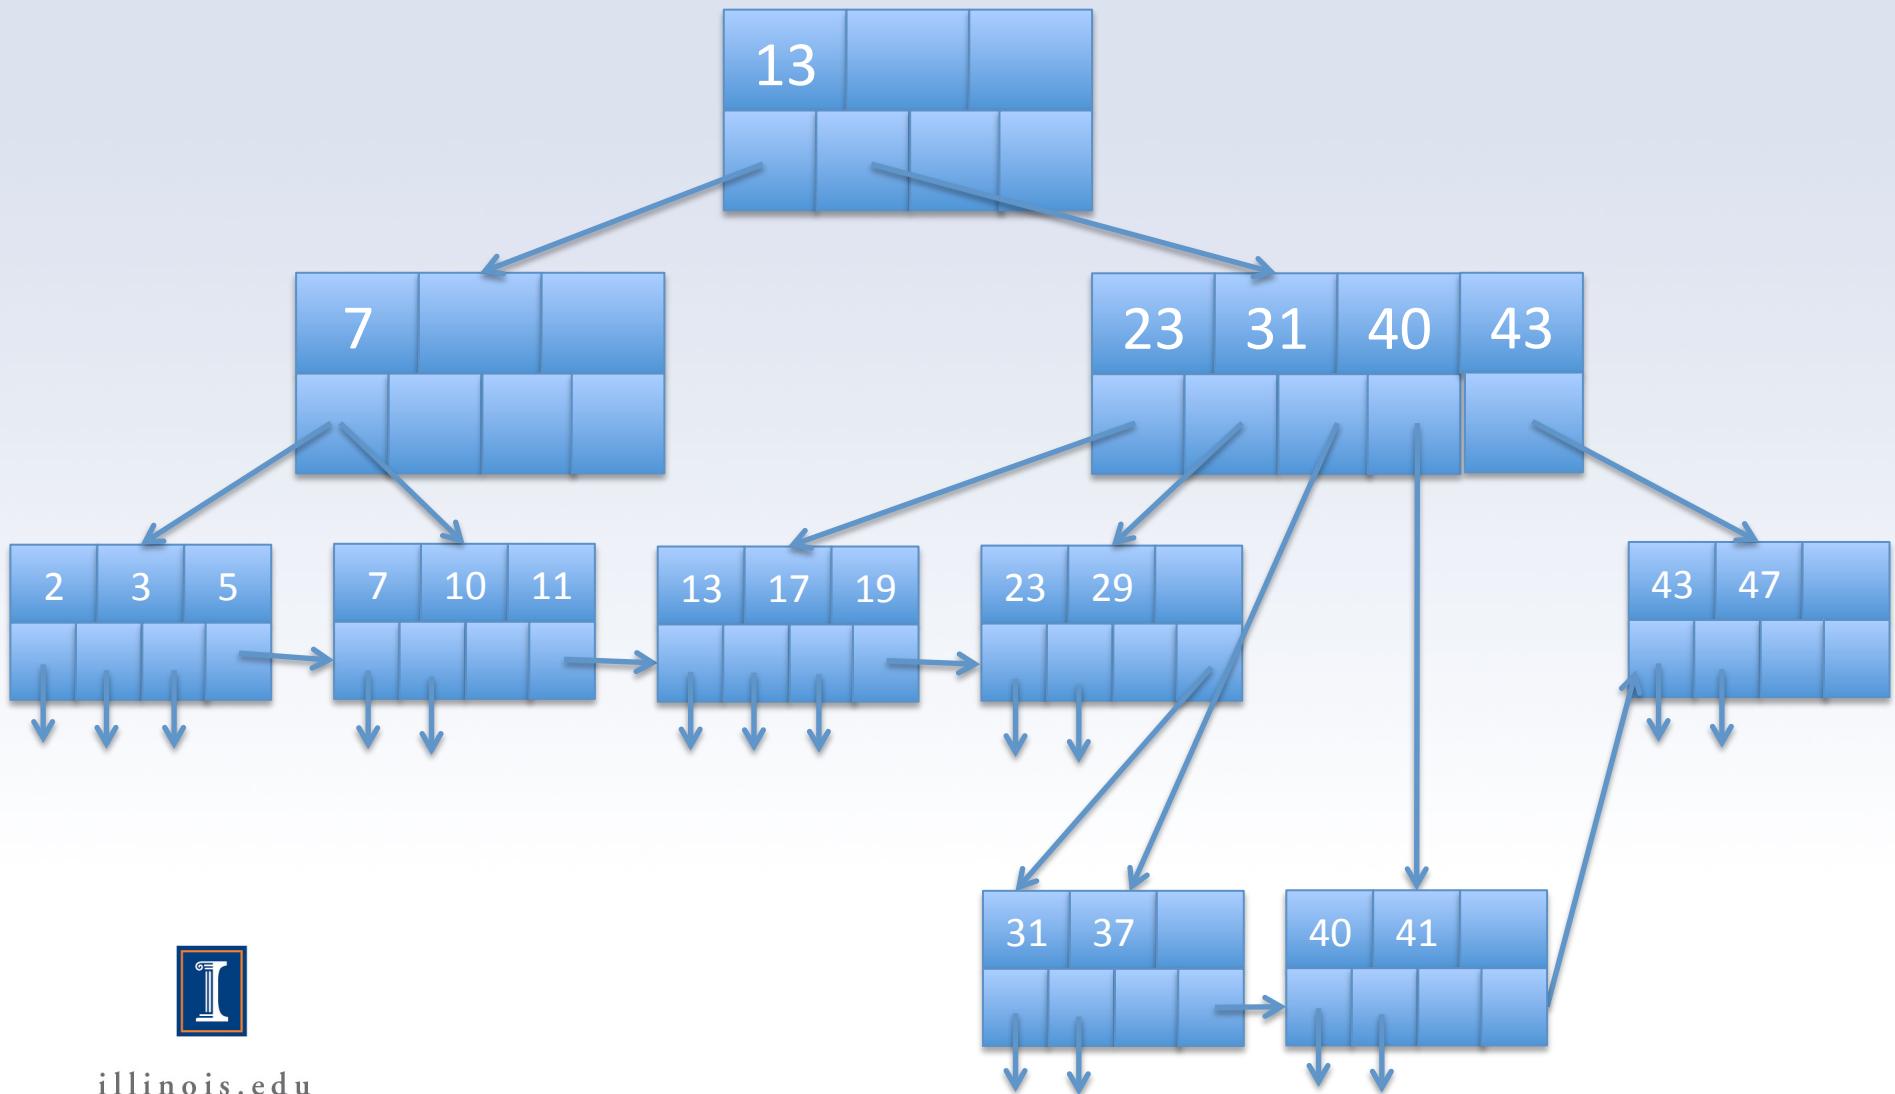
Insert 40



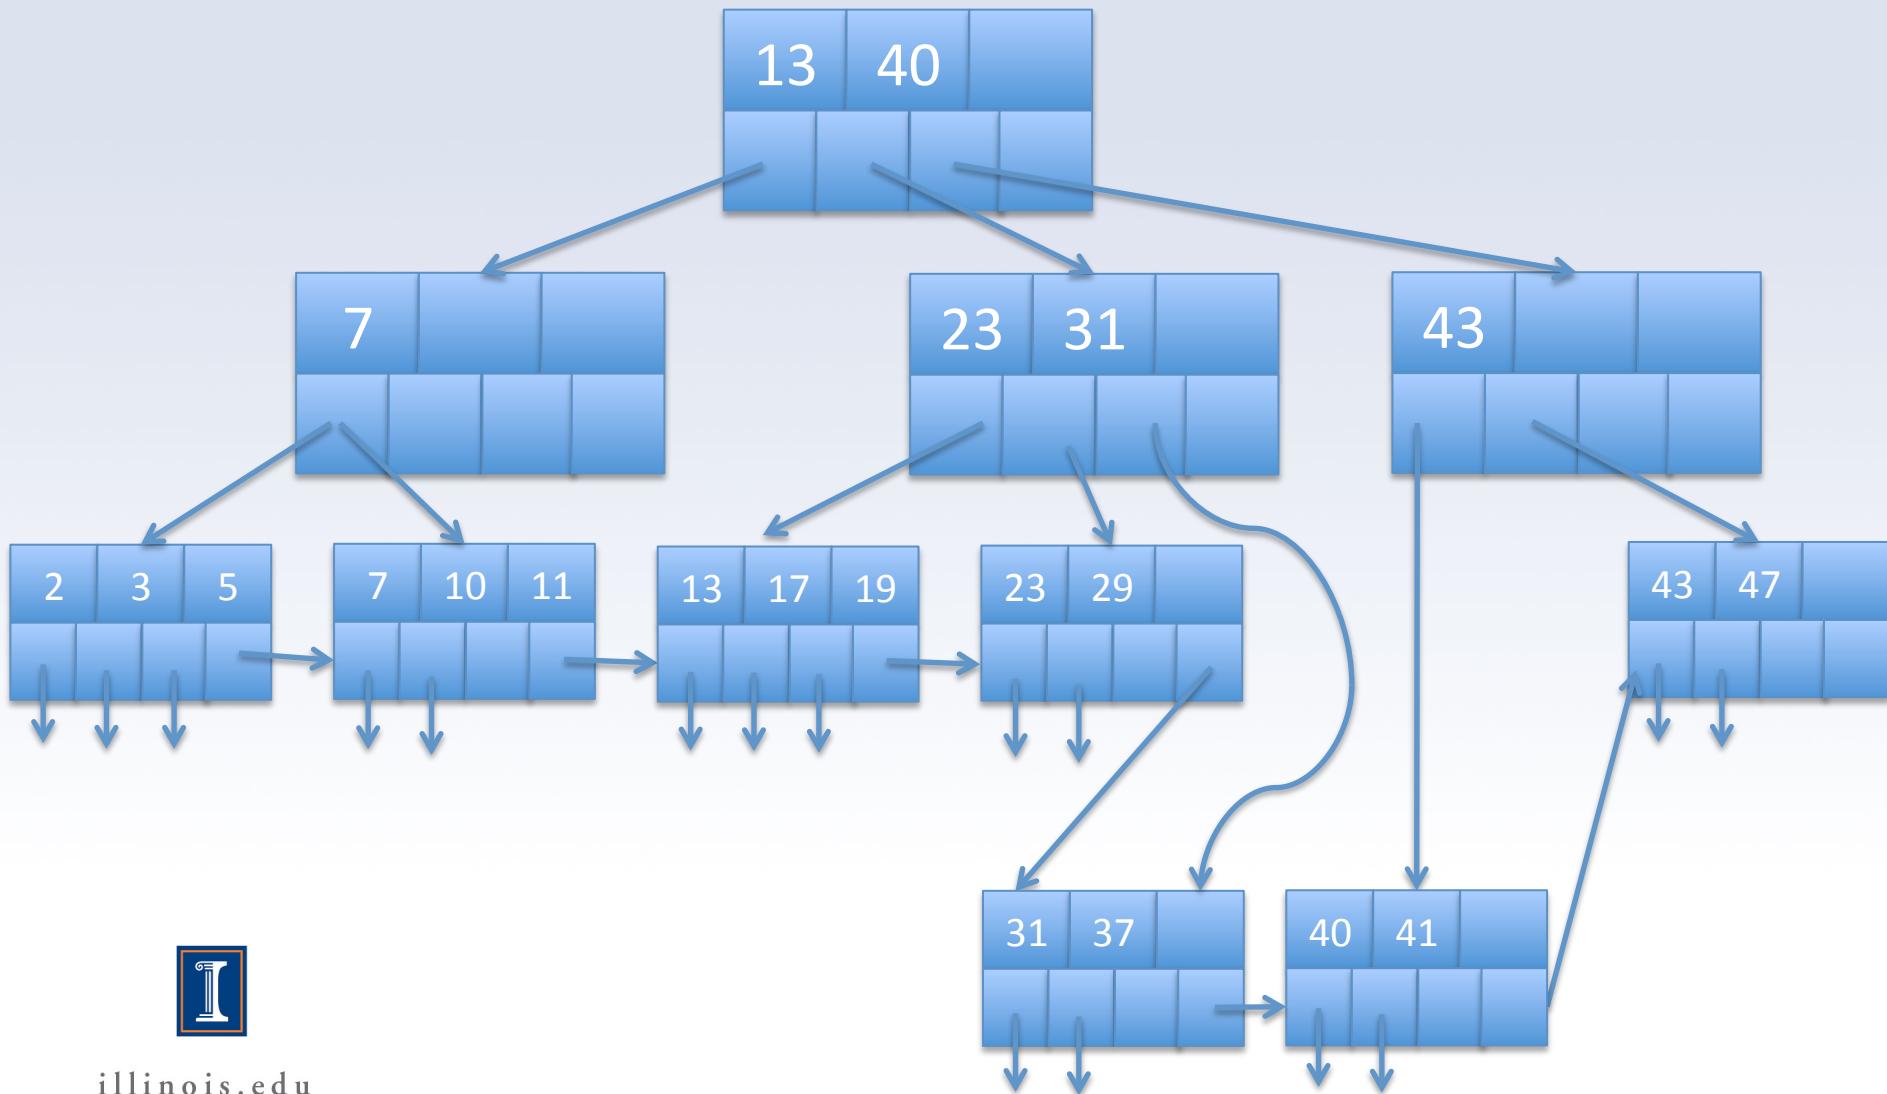
Insert 40



Insert 40



Insert 40

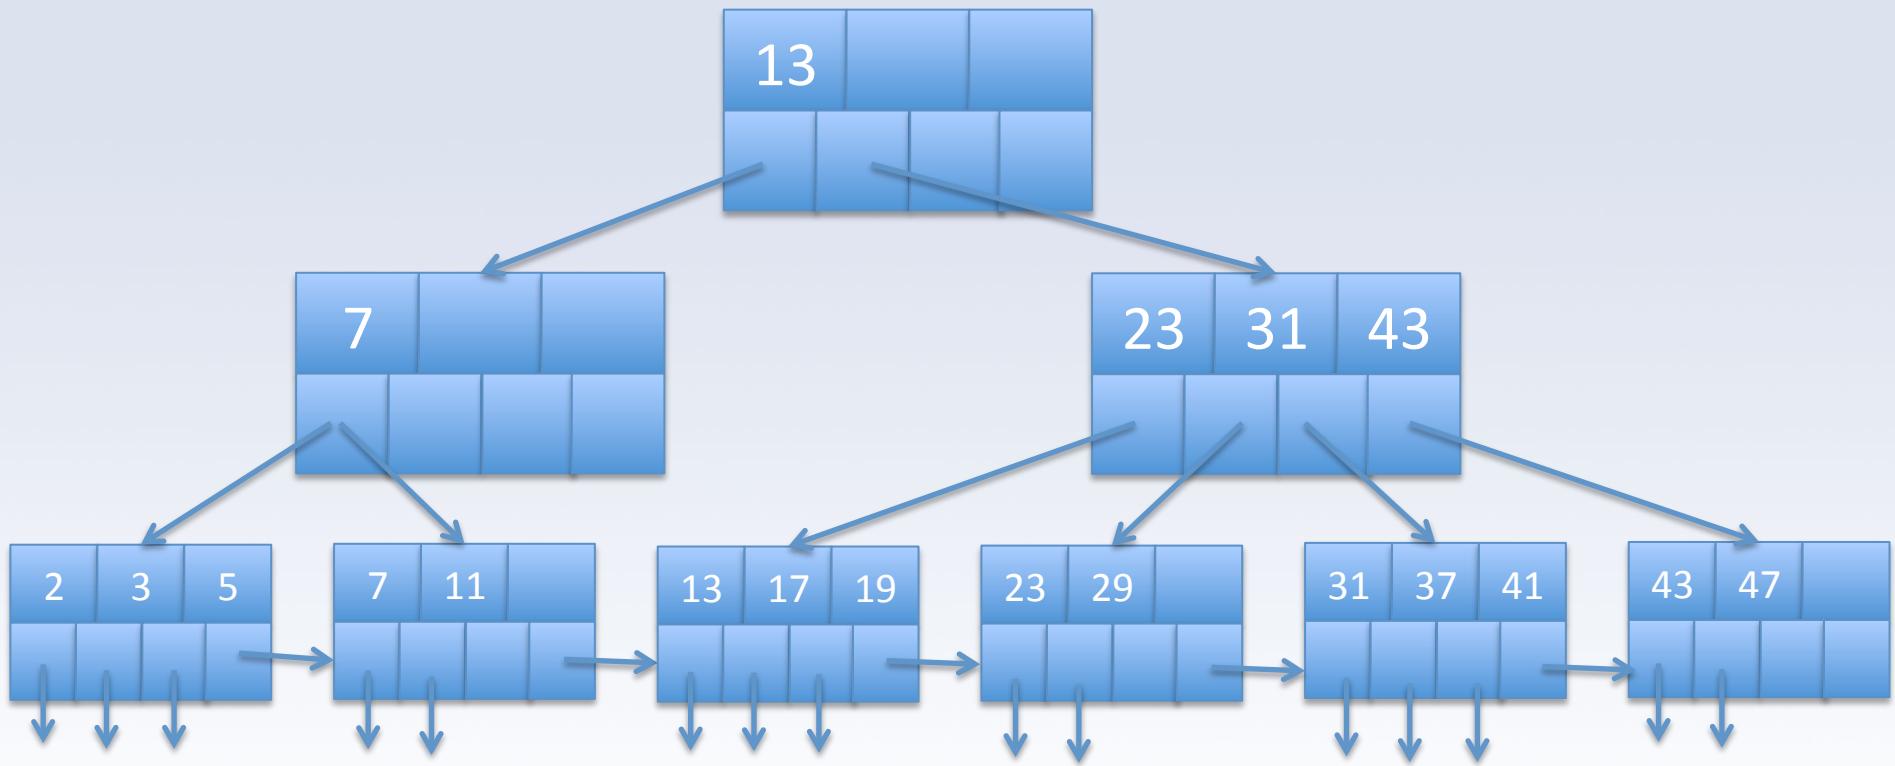


Deletion

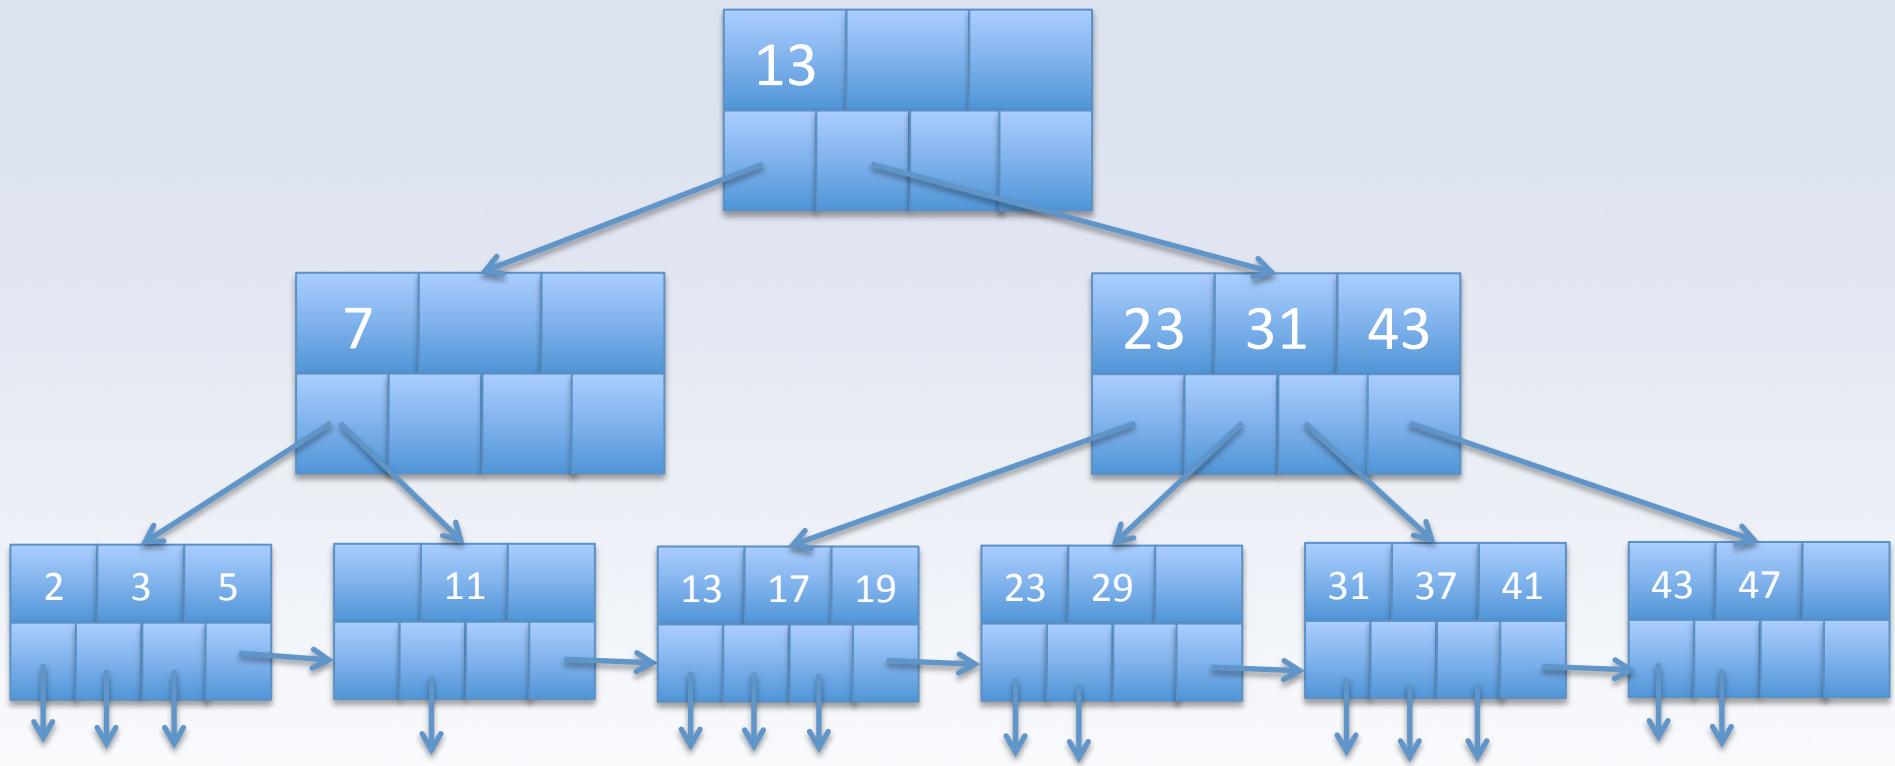
- To delete K from a B-tree
 - Lookup leaf L where K belongs
 - Remove K from L
 - If L has too few entries
 - borrow from sibling if possible
 - if cannot borrow, merge with sibling
 - If merged, delete L from parent and recurse
 - Continue until no merges or root is deleted



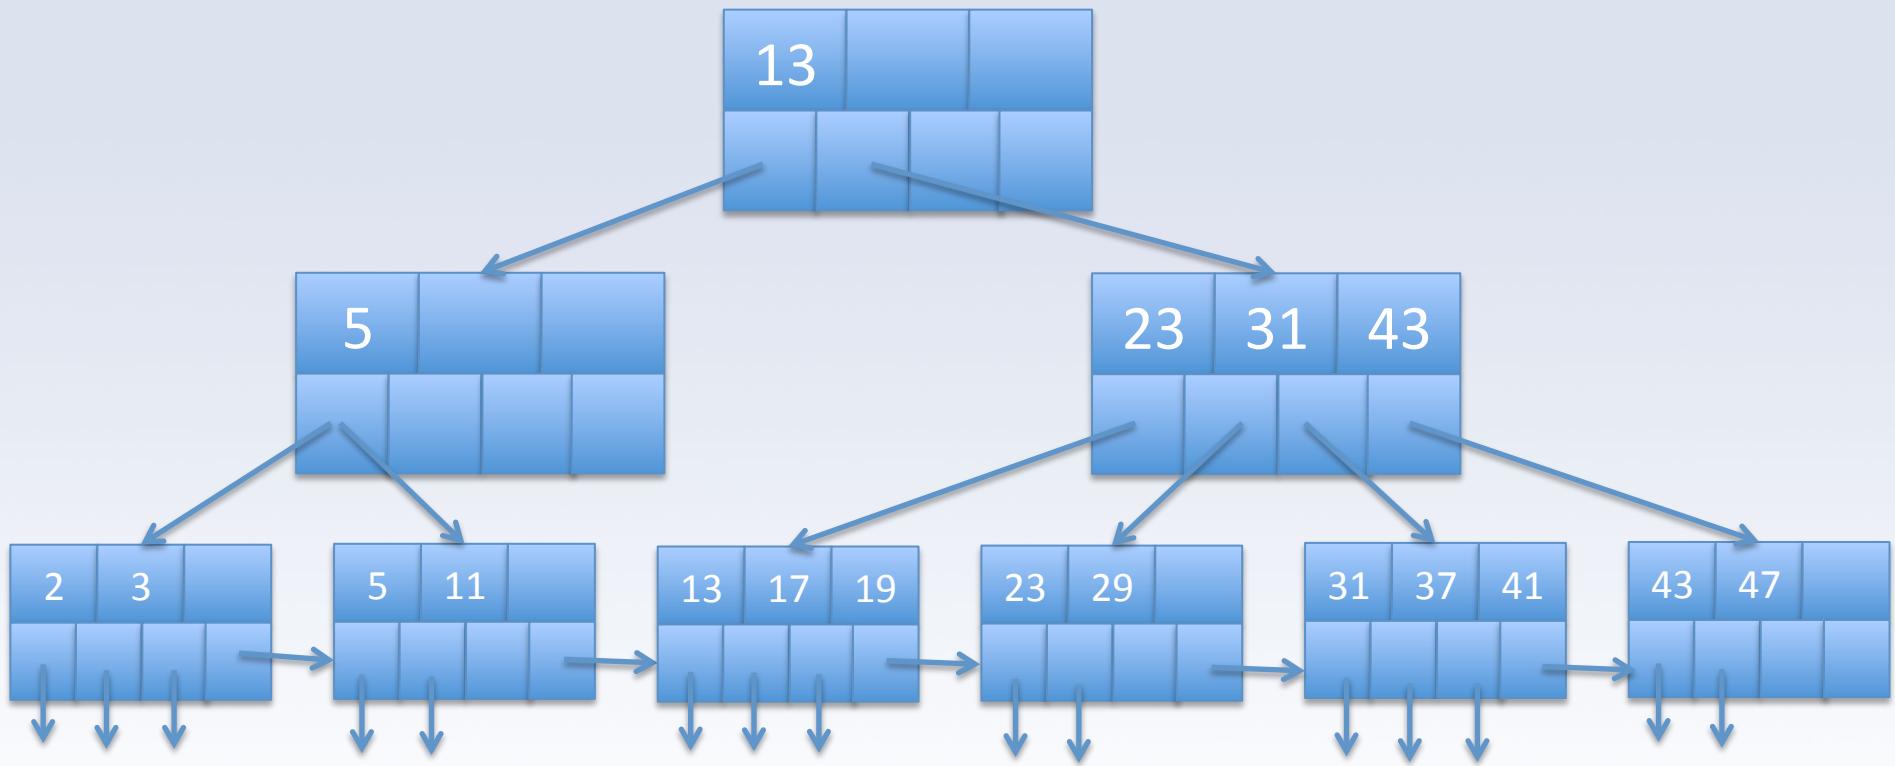
Delete 7



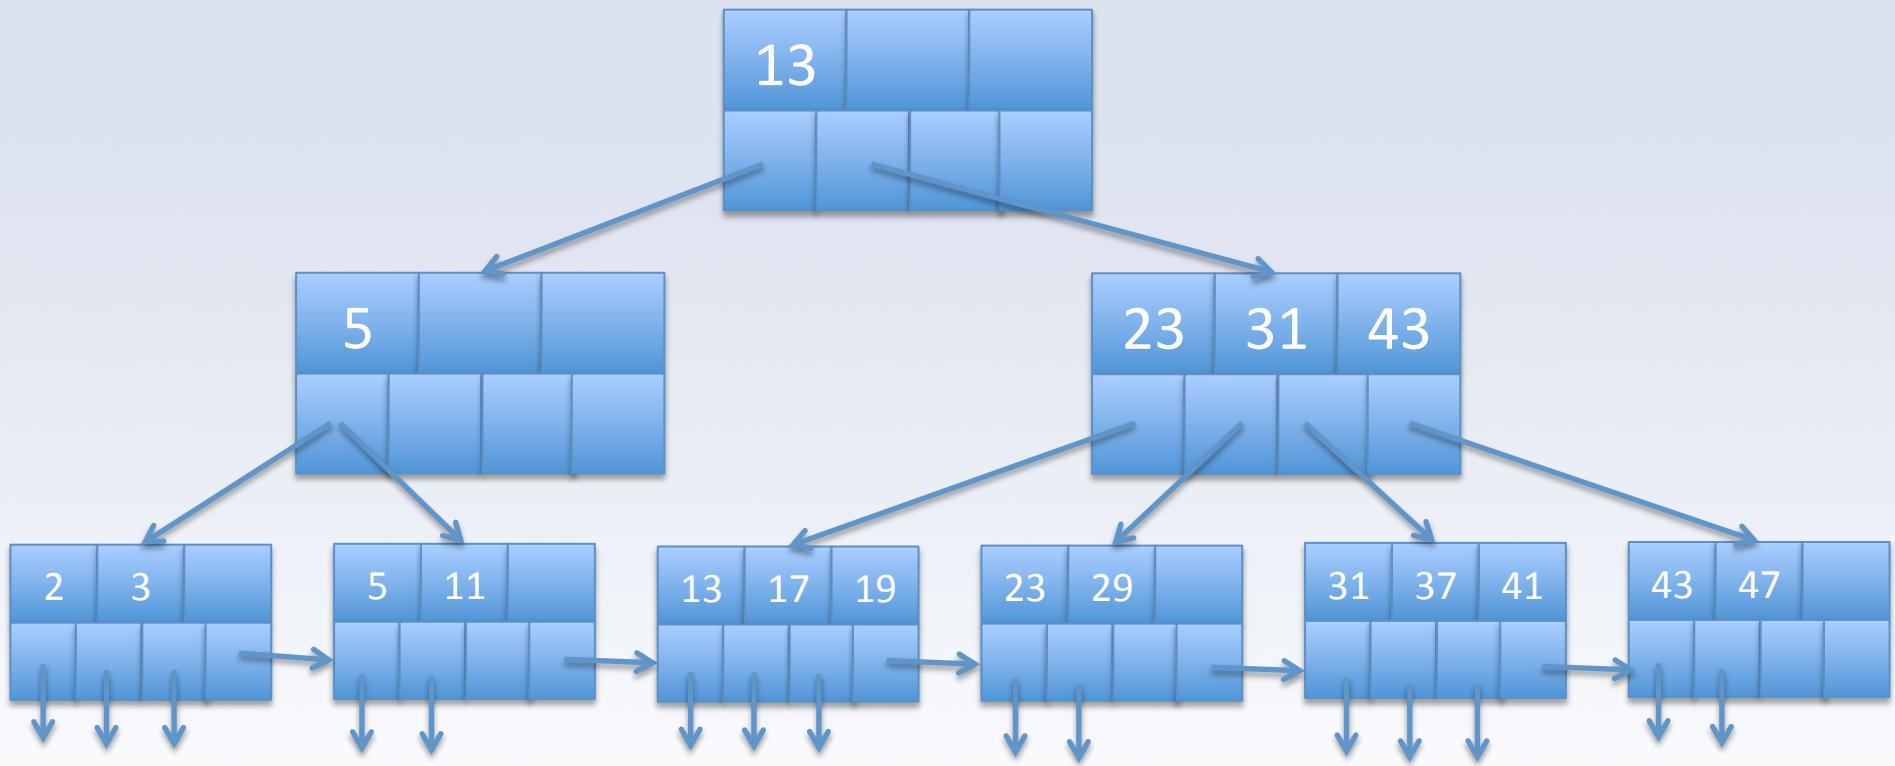
Delete 7



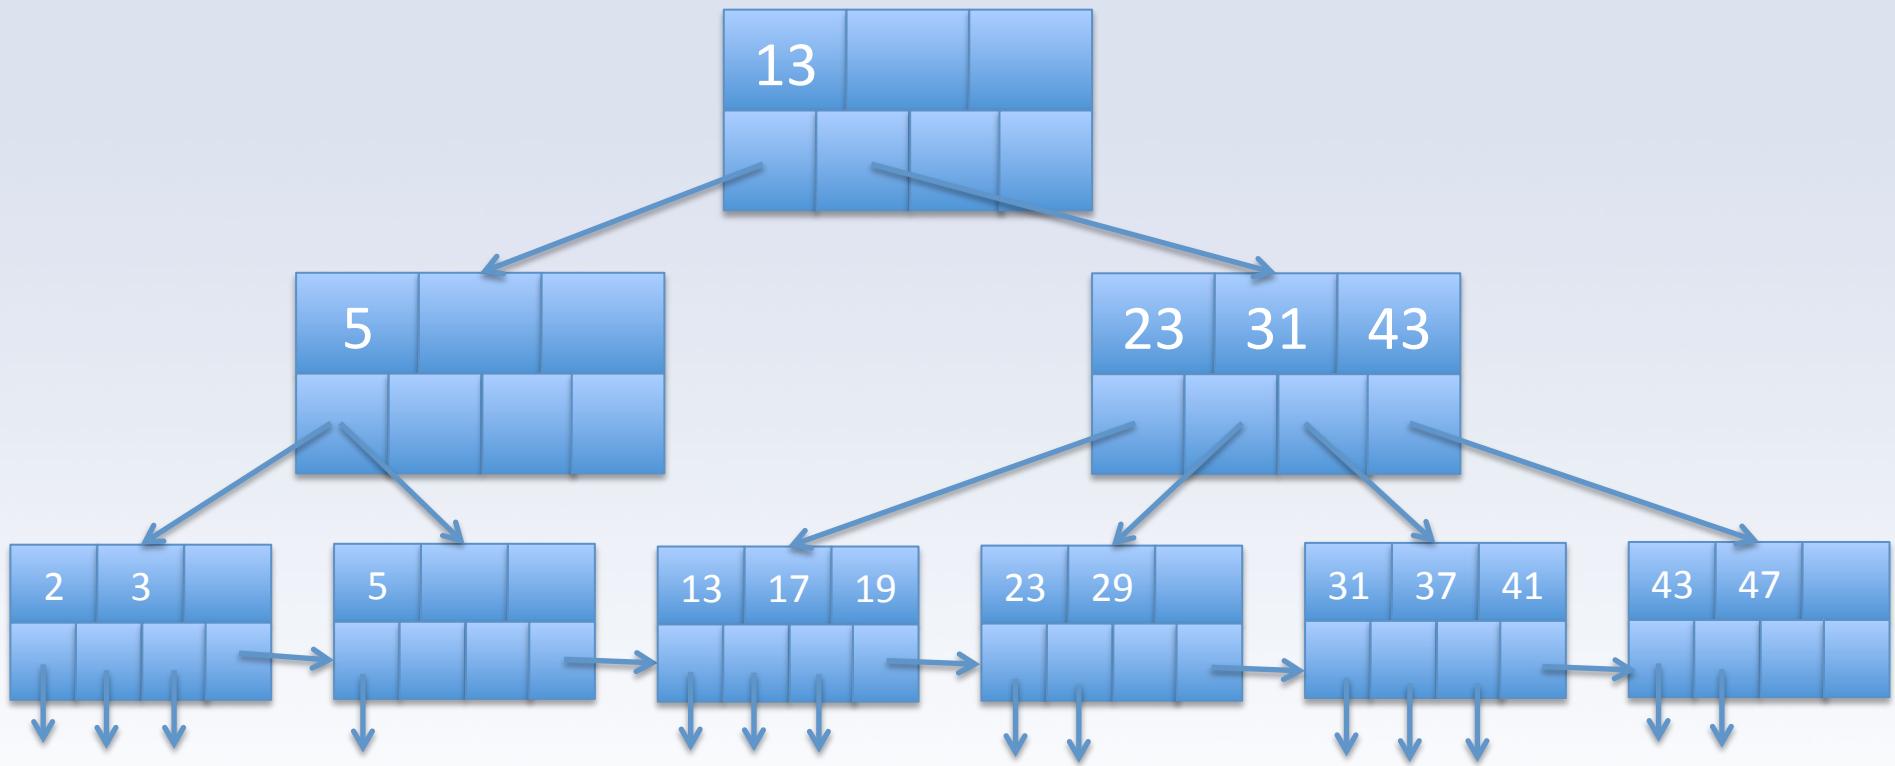
Delete 7



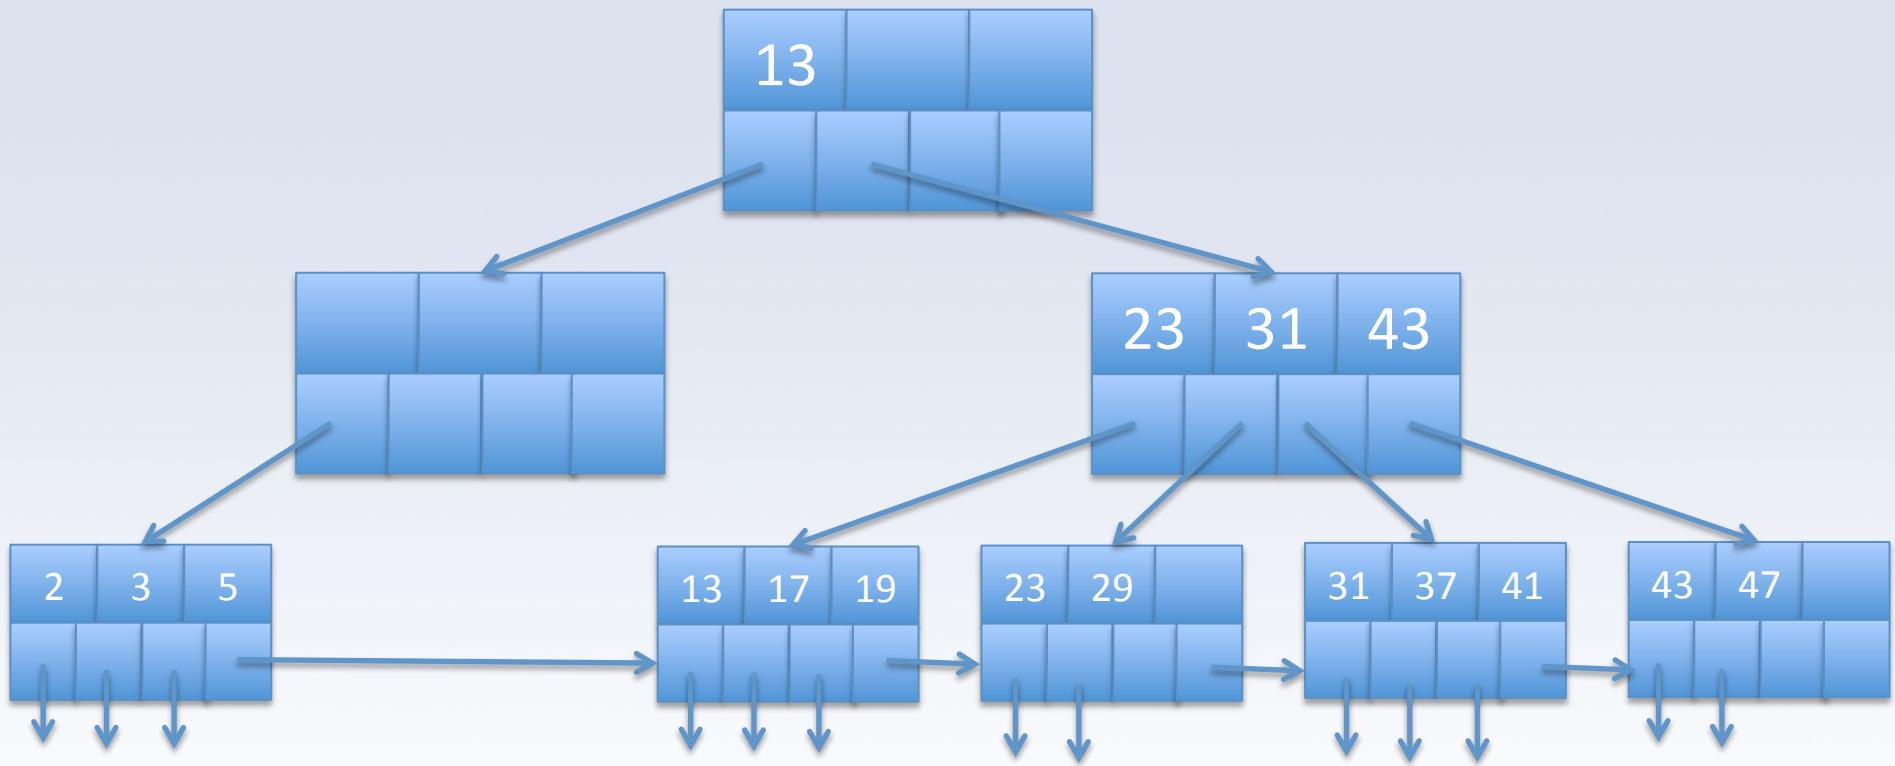
Delete 11



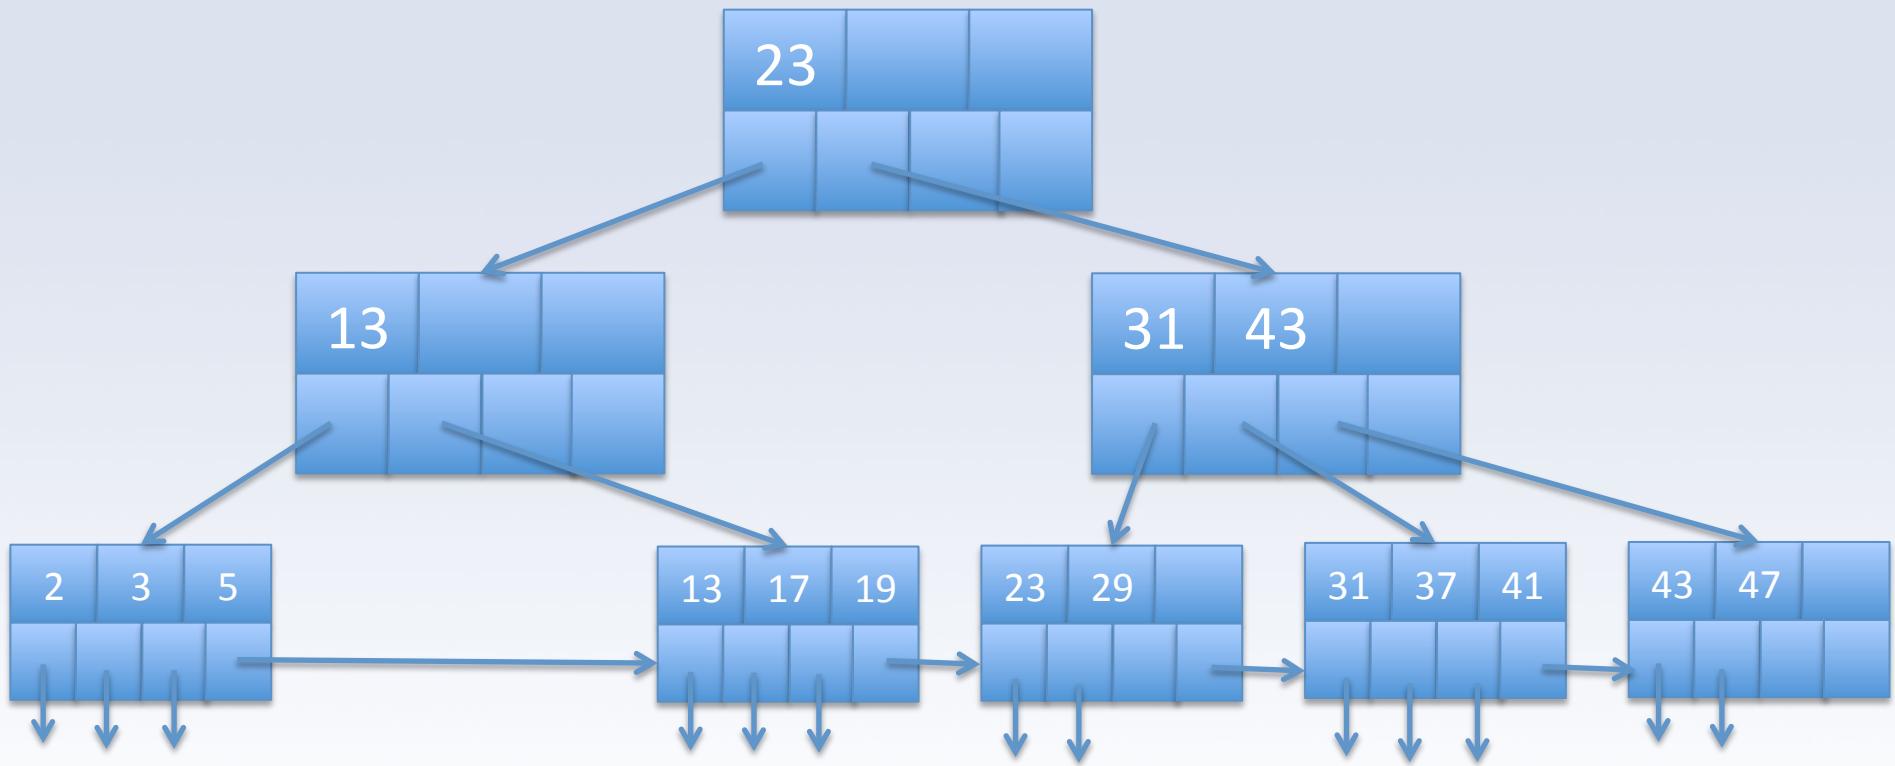
Delete 11



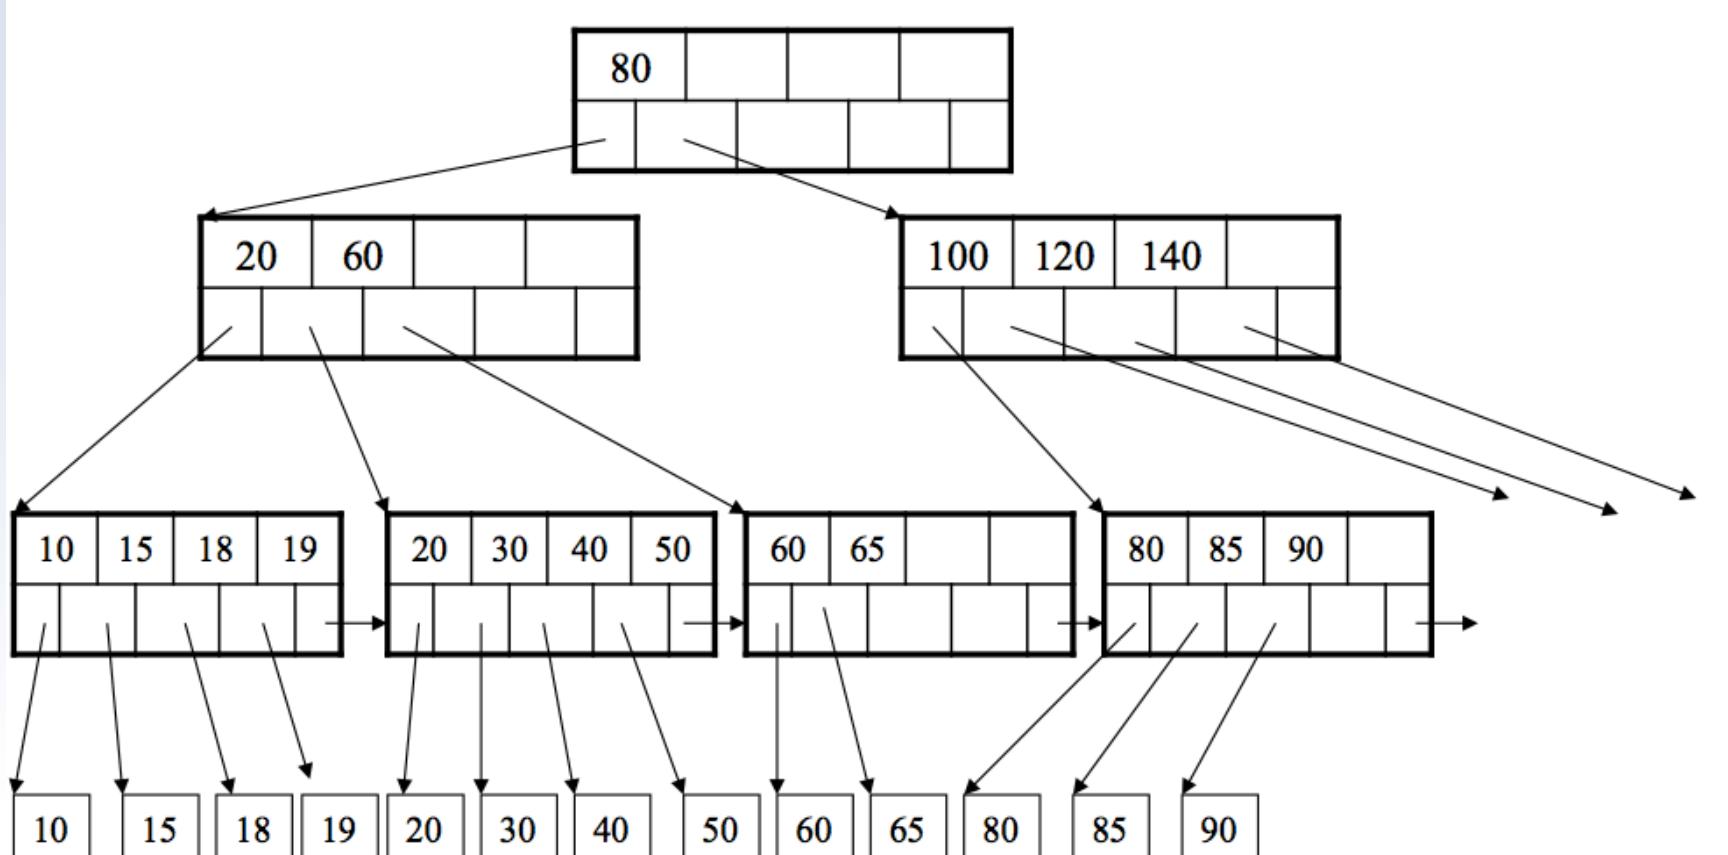
Delete 11



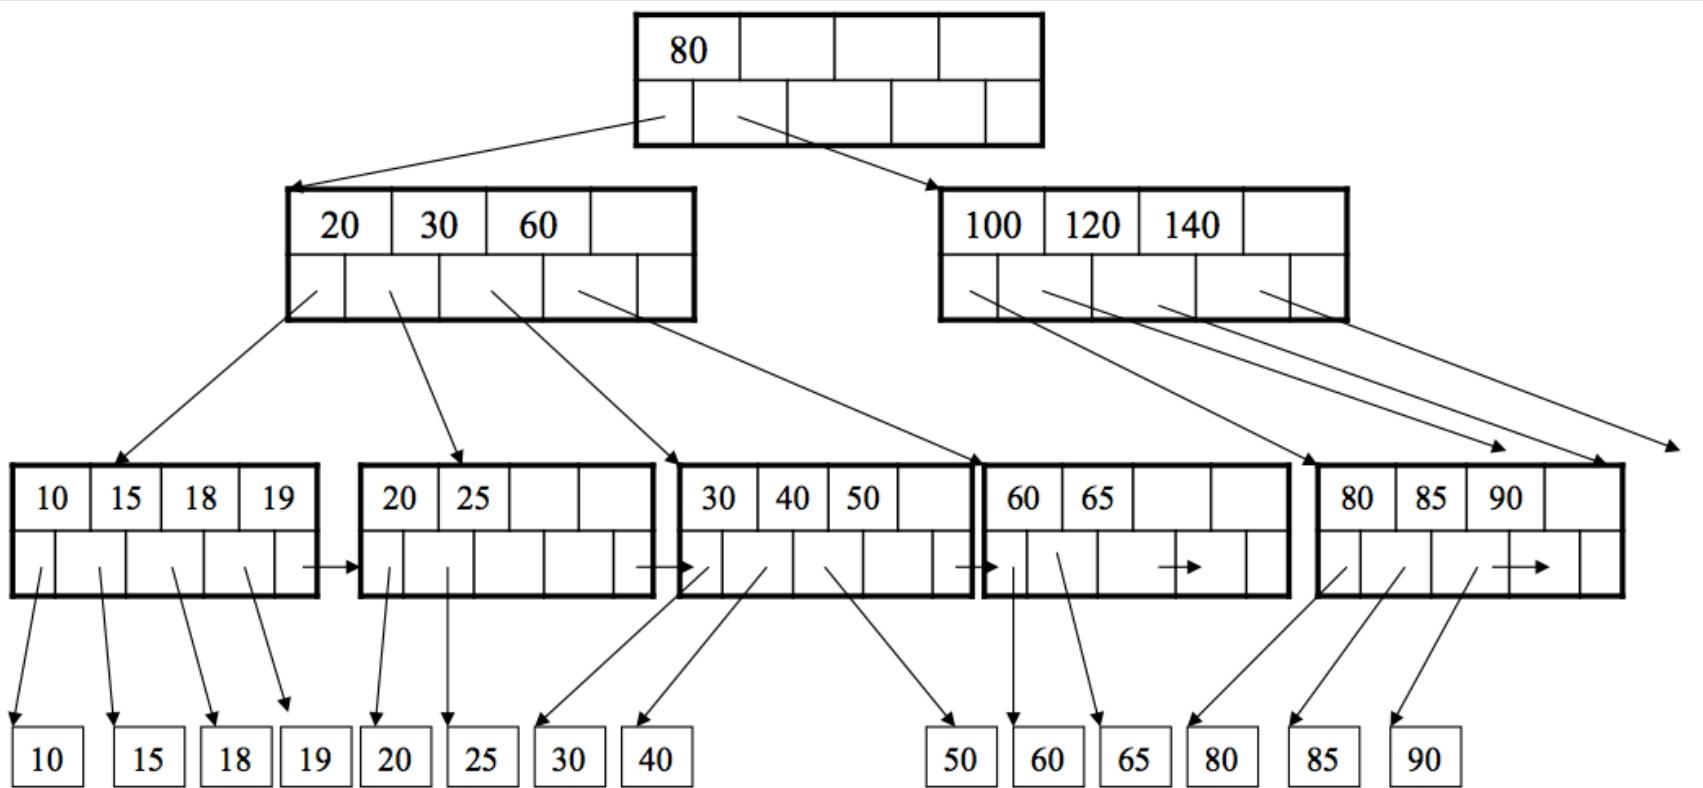
Delete 11



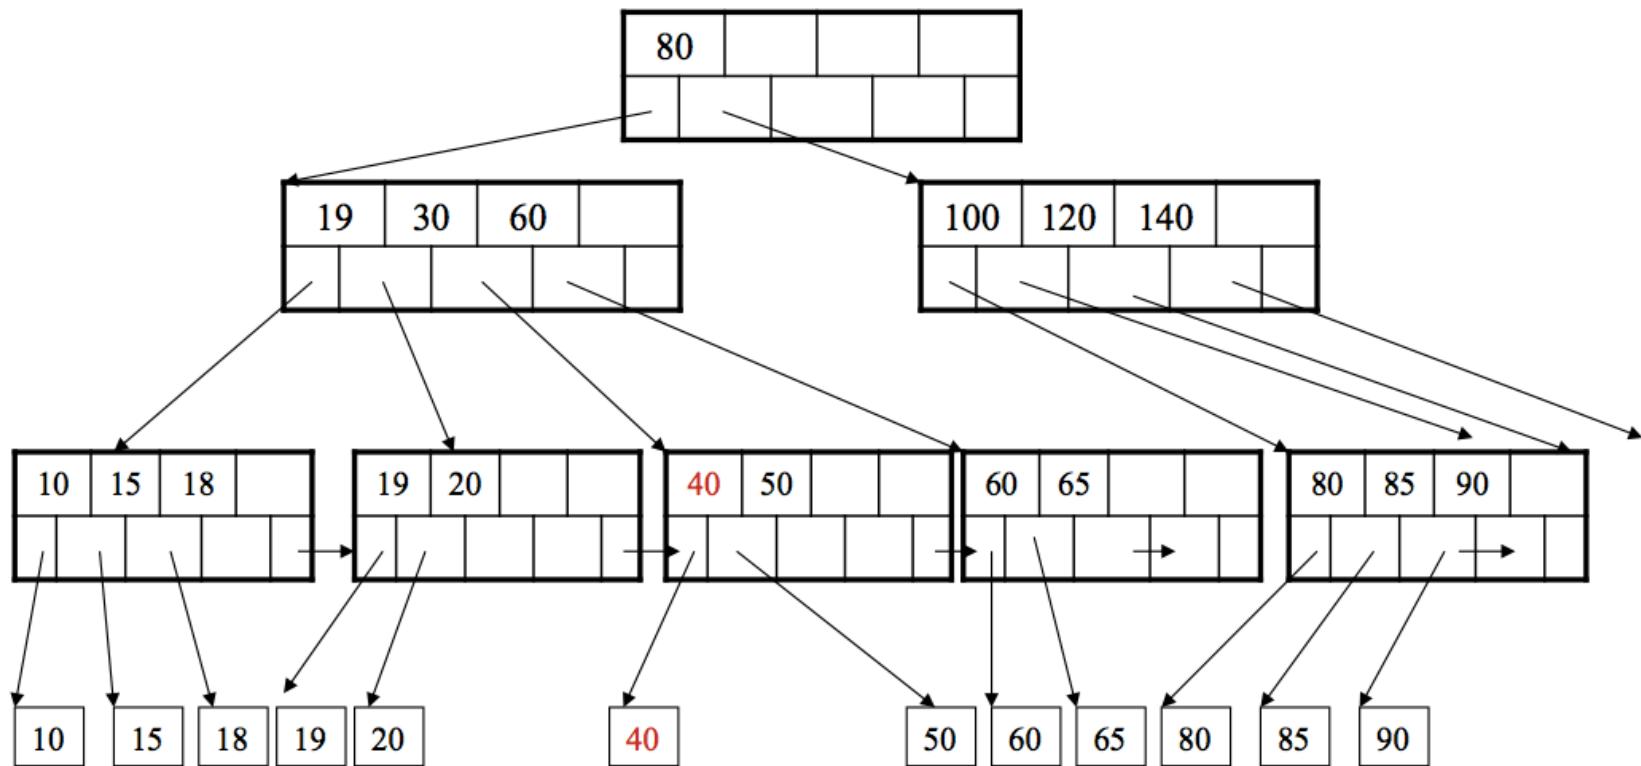
Exercise: Insert 25



Solution



Exercise: Delete 40



Solution

