

UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Indexing



illinois.edu

Announcements

- MP2 due Friday
- Track 1 initial demos Thursday, Friday
 - After this, work on your project
- HW3 assigned the break



Announcements

- Friday's lecture will be short
 - I'll cover “advanced” indexing topics
 - Multidimensional indexing
 - Large string indexing
 - Web search indexing



Review

- Why index a relation?
- What is a *search key*?
- What is a *clustered index*?
- What is a *sparse index*?



Review

- What is a *secondary index*?
- What is a *sequential file*?
- What is a *multilevel index*?
- What is *bucket intersection*?

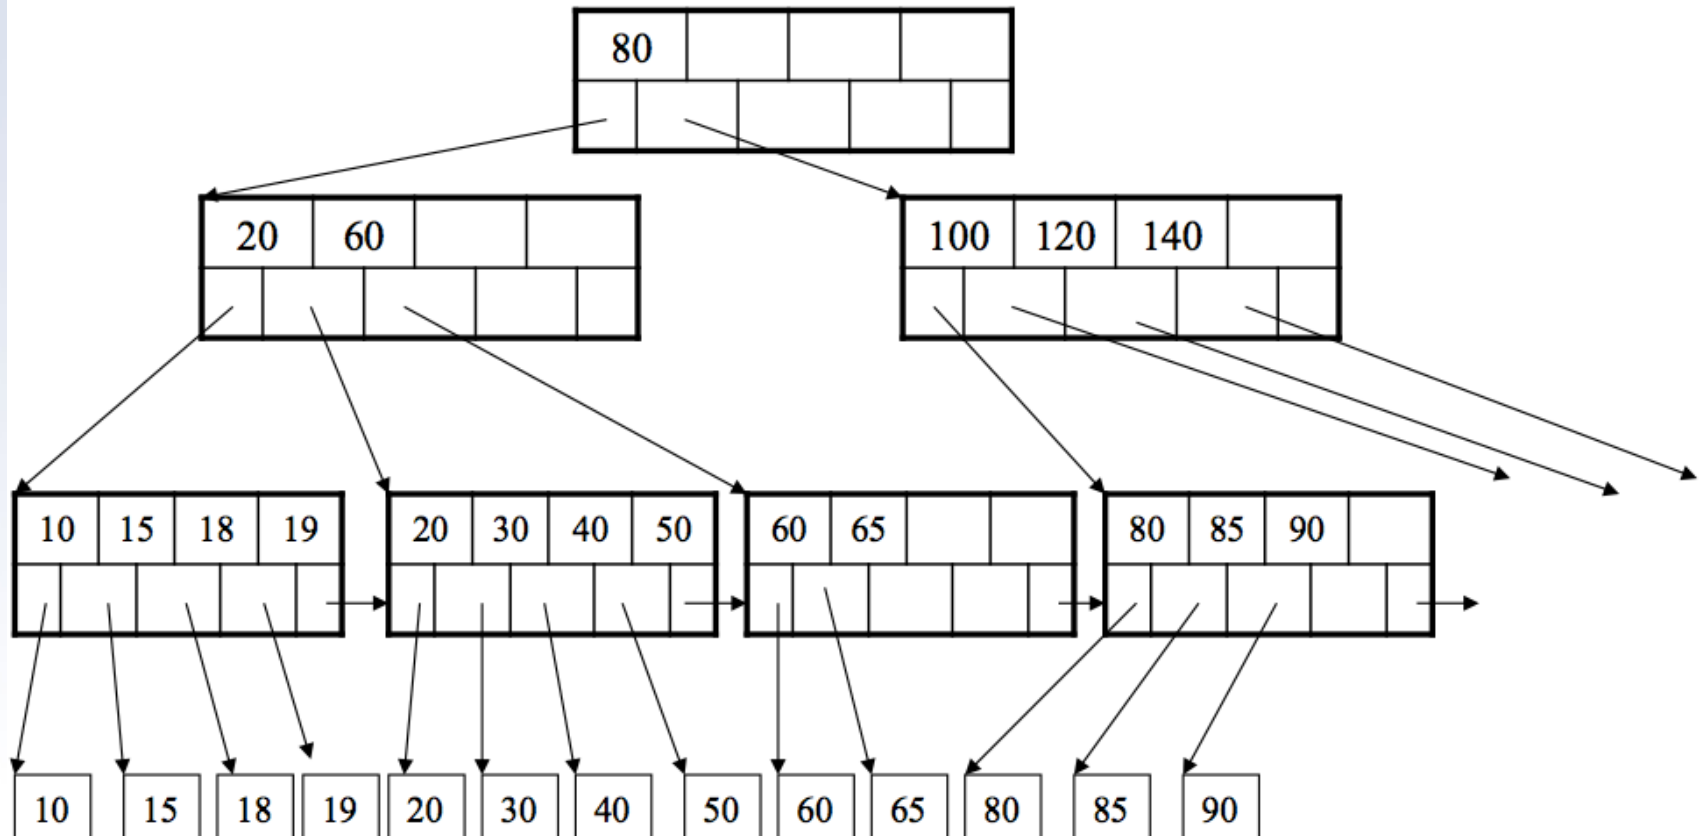


Review


- What are the advantages of a B-tree index?
- What does the parameter n of a B-tree index represent?
- Why do the leaf nodes of a B-tree have “next” pointers?



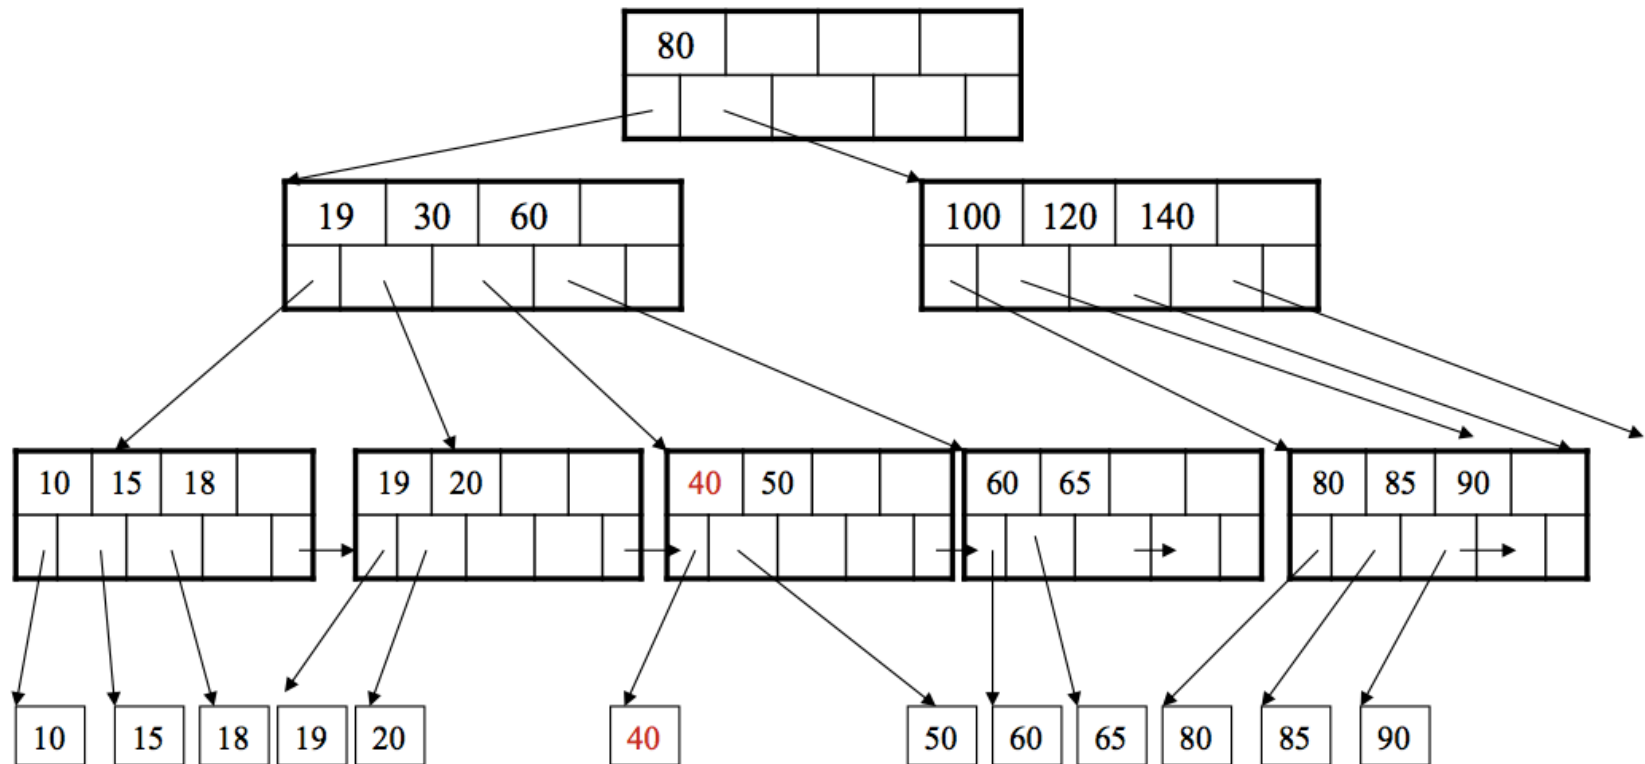
Review: Insert 25



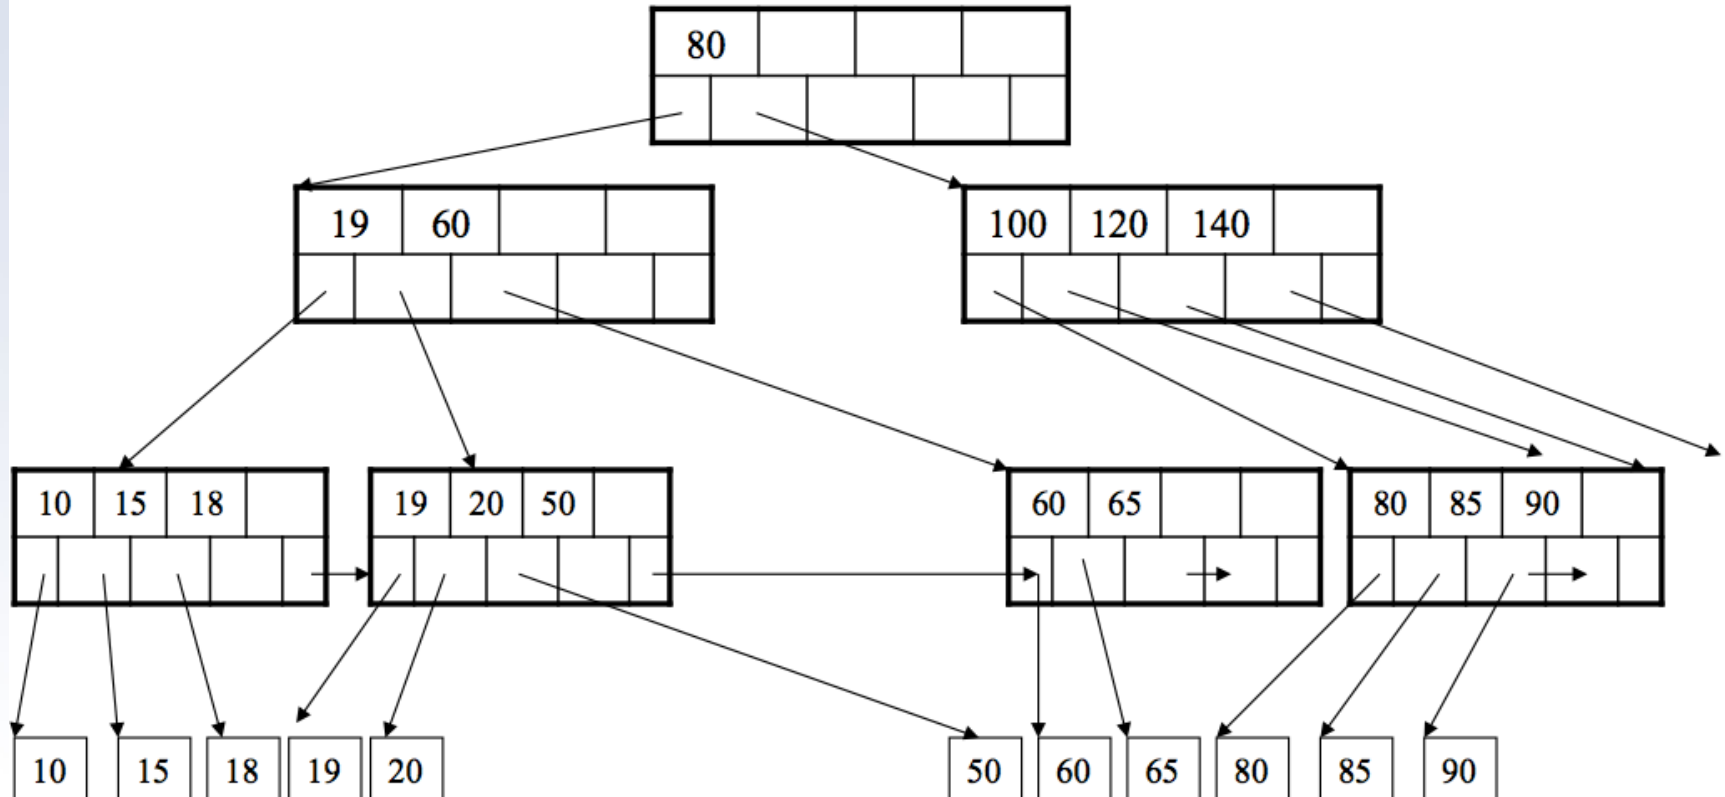
Solution



Review: Delete 40



Solution



Indexing

- Sparse indexes
- Dense indexes
- B-trees
- Hash tables



Hash Tables

- Basics:
 - n *buckets* used to store keys
 - *hash function* maps keys to buckets



Hash functions

- Maps the set of search keys to the set of buckets
 - $f: K \rightarrow \{0, 1, \dots, n - 1\}$
 - K is the set of all keys
 - n is the number of buckets



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

Bucket	Keys
0	
1	
2	
3	
4	



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

$$100\%5=0$$

Bucket	Keys
0	100
1	
2	
3	
4	



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

$$12\%5=2$$

Bucket	Keys
0	100
1	
2	12
3	
4	



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

$$59\%5=4$$

Bucket	Keys
0	100
1	
2	12
3	
4	59



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

$$3\%5=3$$

Bucket	Keys
0	100
1	
2	12
3	3
4	59



Example

- Given 5 buckets, let $f(k)=k\%5$
 - modulo (remainder after dividing by 5)
- Insert 100, 12, 59, 3, 33

$$33\%5=3$$

Bucket	Keys
0	100
1	
2	12
3	3, 33
4	59



DBMS Hash Tables

- Buckets consist of blocks
- Blocks contain records
- Overflow blocks added as needed
 - Pointers to overflow blocks go in the header



Example

block 0



record 1

block 1



$1\%3=1$

block 2



Example


block 0



record 2


$$2\%3=2$$

block 1




record 1


block 2



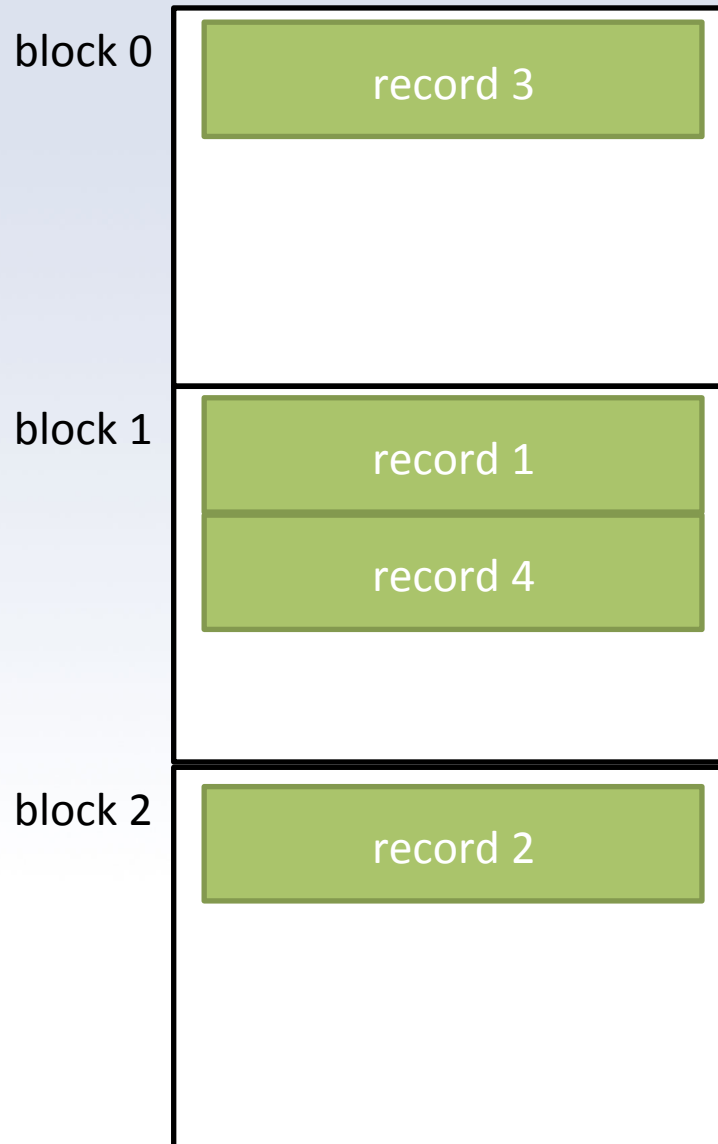
Example



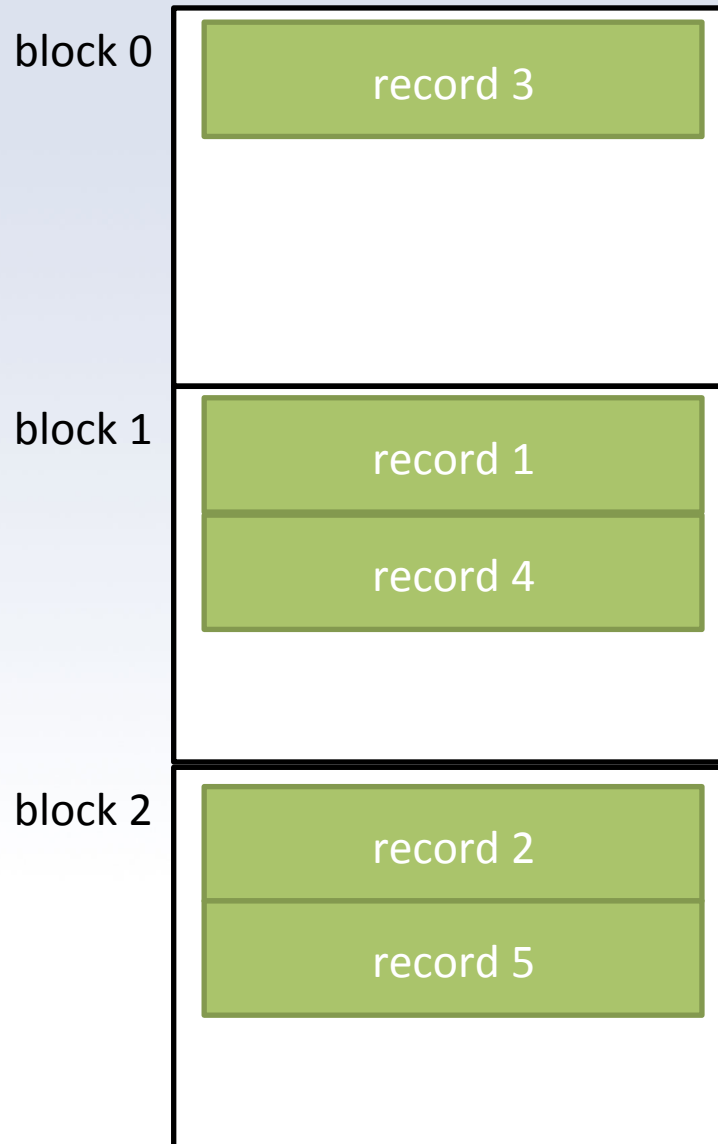
Example



Example



Example



Example

block 0

record 3

record 6

block 1

record 1

record 4

block 2

record 2

record 5



Example

block 0

record 3

record 6

block 1

record 1

record 4

record 7

block 2

record 2

record 5



Example

block 0

record 3

record 6

record 10

$$10 \% 3 = 1$$

block 1

record 1

record 4

record 7

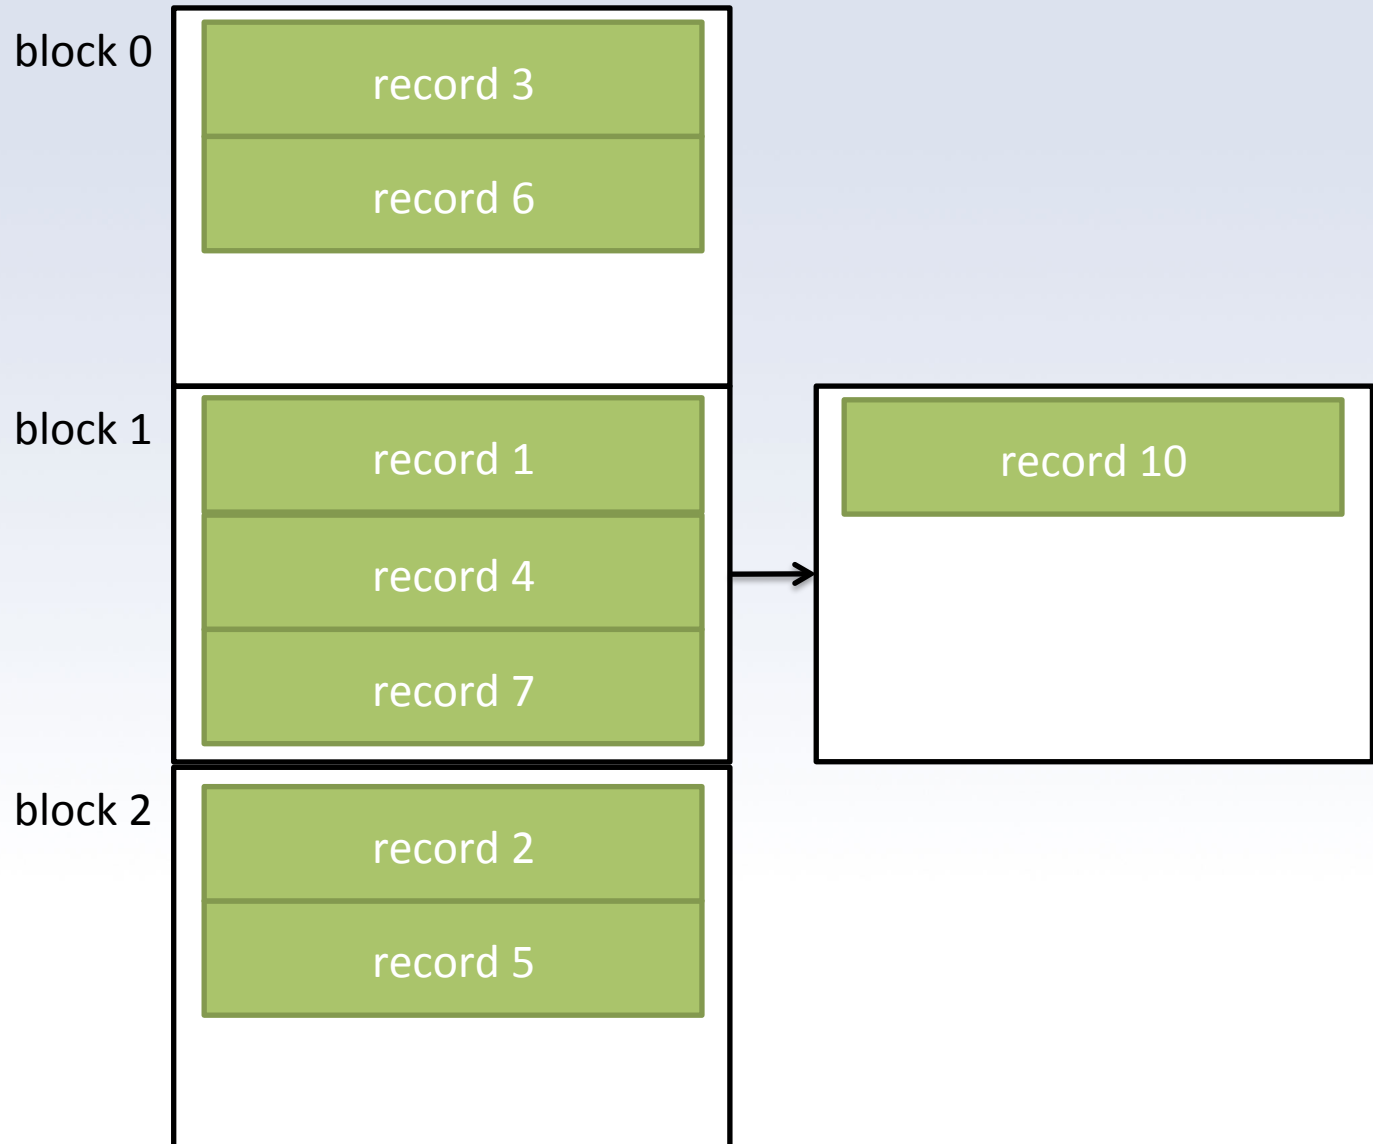
block 2

record 2

record 5



Example



Hash Table Indexes

- Efficiency depends on blocks per bucket
 - One disk I/O for lookup if one block/bucket
 - Many disk I/Os if lots of overflow buckets
- Can we improve?



Dynamic Hash Tables

- Allow number of buckets to vary
 - Try to keep number of blocks per bucket low
- Two methods
 - Extensible hashing
 - Linear hashing




Extensible Hash Tables


- Strategy: double the number of buckets when needed
 - keep counter i
 - hash size = 2^i
 - keep *hash array* = pointers to buckets
 - hash array pointers can point to same bucket




Example




Example: Look up 1001



Example: Look up **1**001



Example: Look up **1001**




Extensible Hash Tables


- Insertion
 1. Lookup new key. If room, insert.
 2. No room:
 - a. If “nub” value = i , increment i . Doubles hash array (2^{i+1} entries.)
 - b. Split block. New blocks have incremented “nub” values.
 - c. Distribute original block into new blocks.
 - d. Repeat if necessary (still not enough room.)



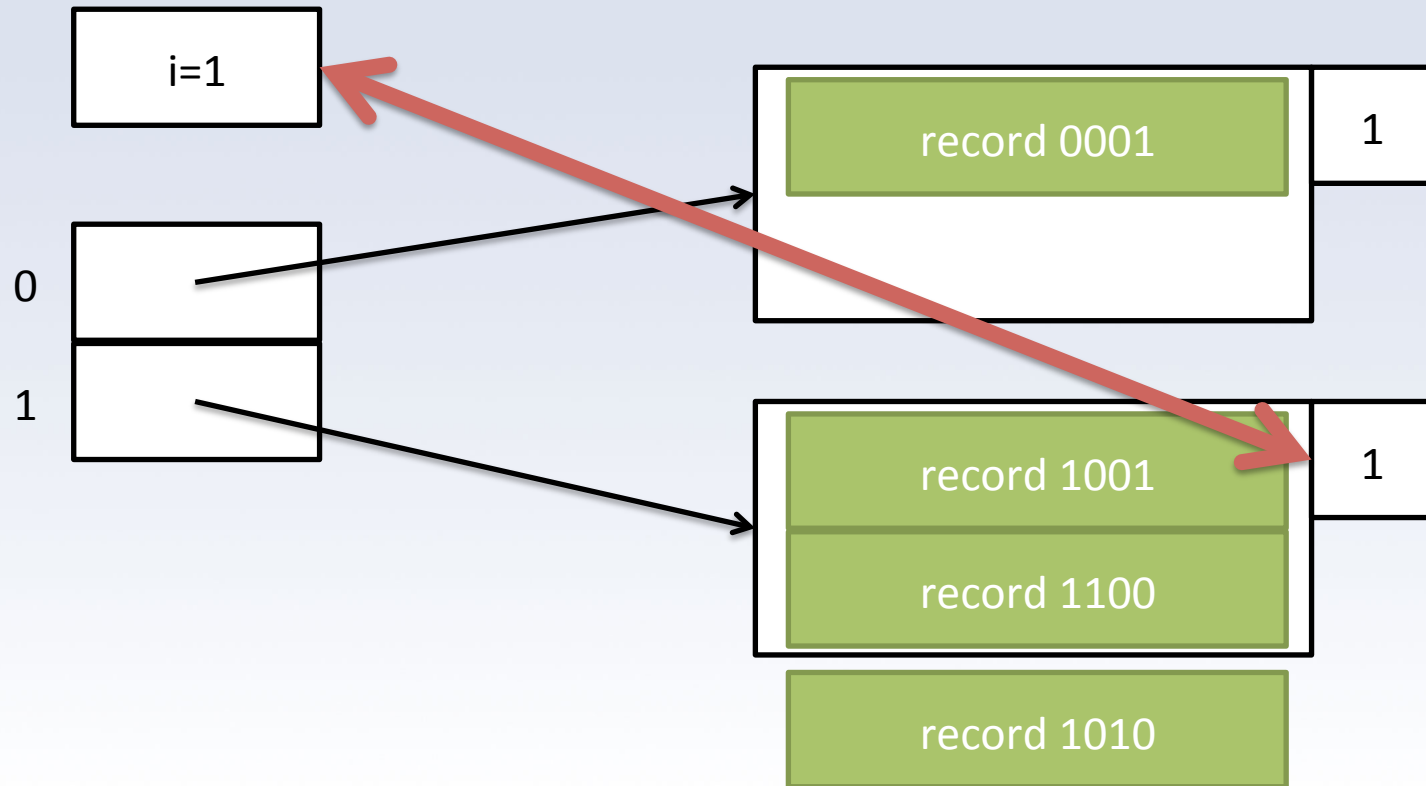
Example: Insert 1010



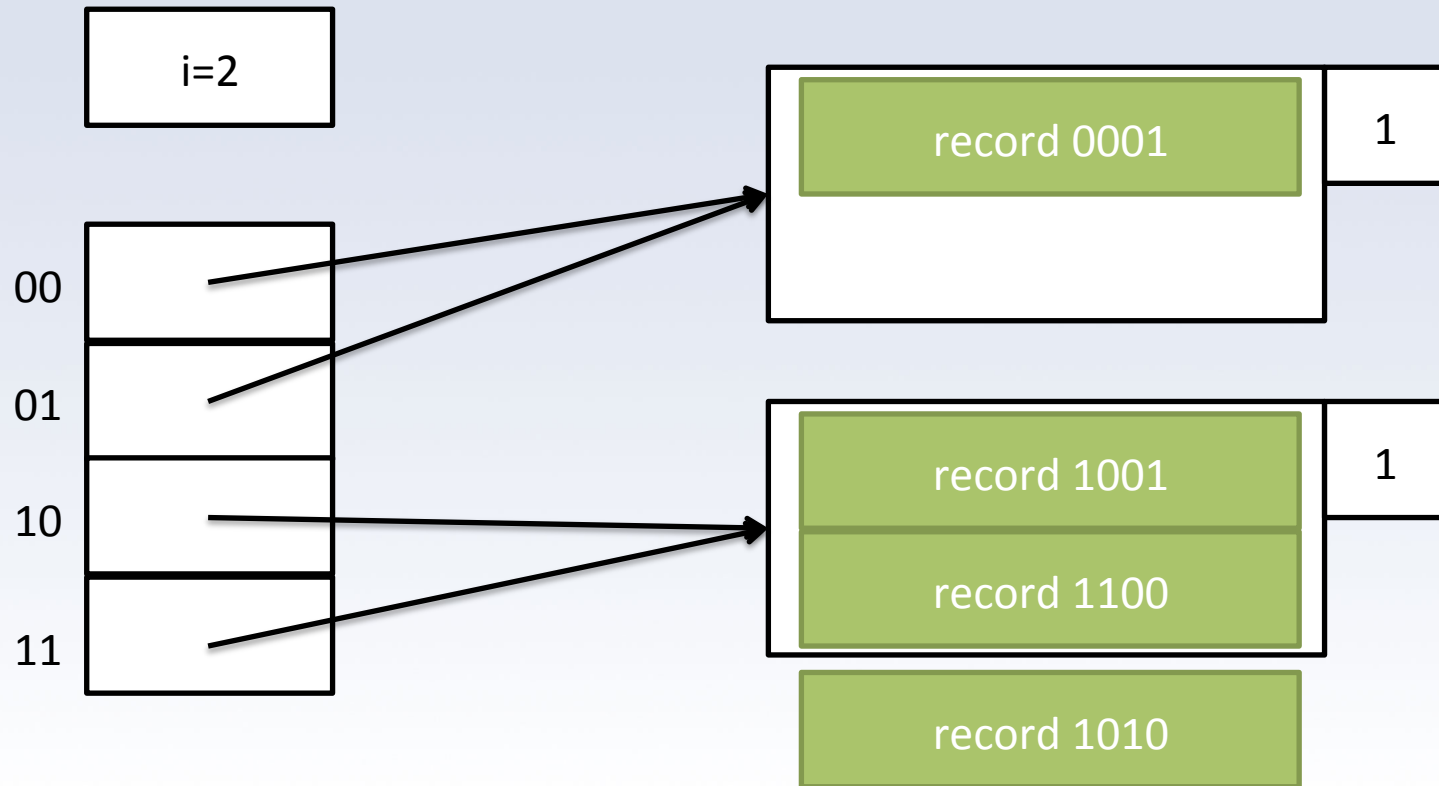
Example: Insert 1010




Example: Insert 1010



Example: Insert 1010



Example: Insert 1010



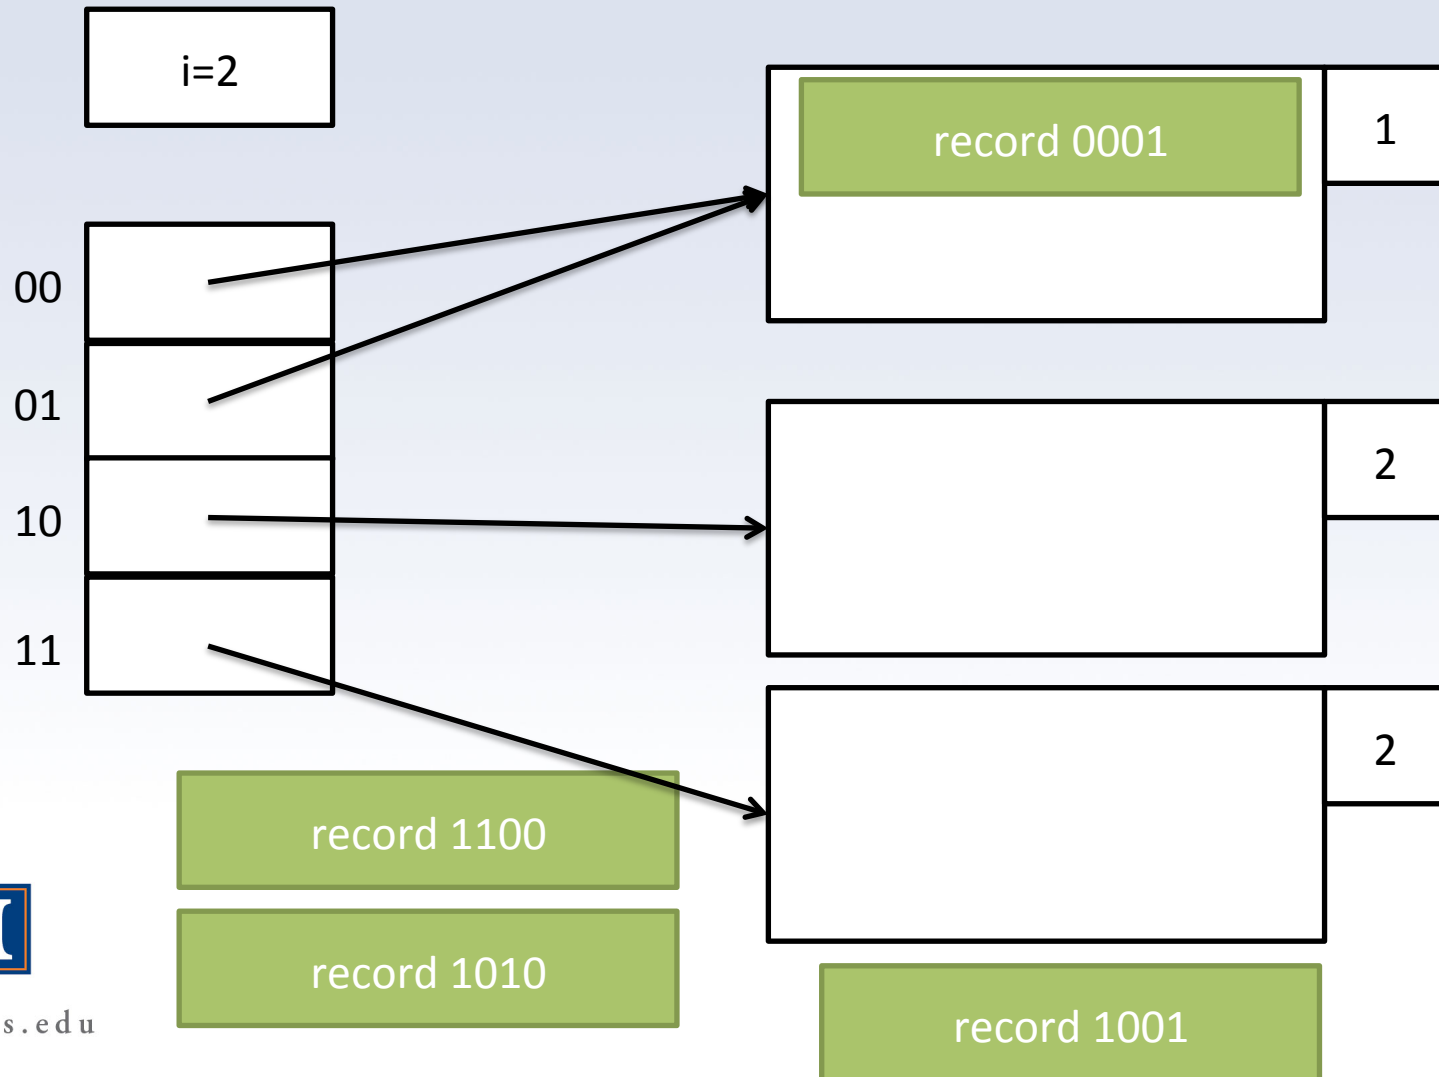
record 1100

record 1010

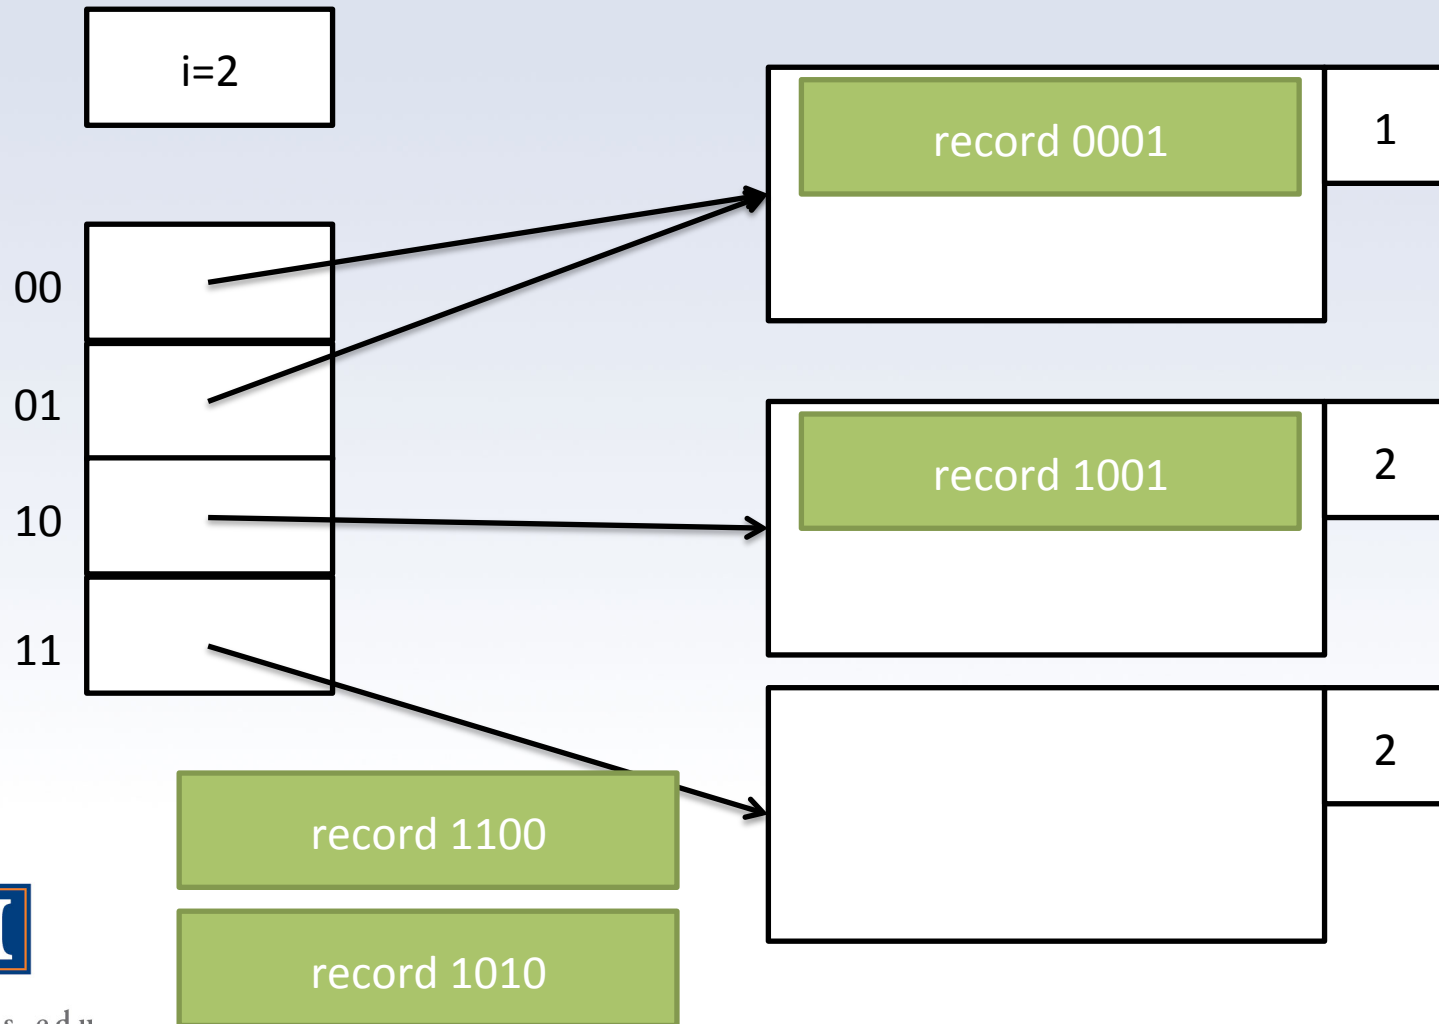
record 1001



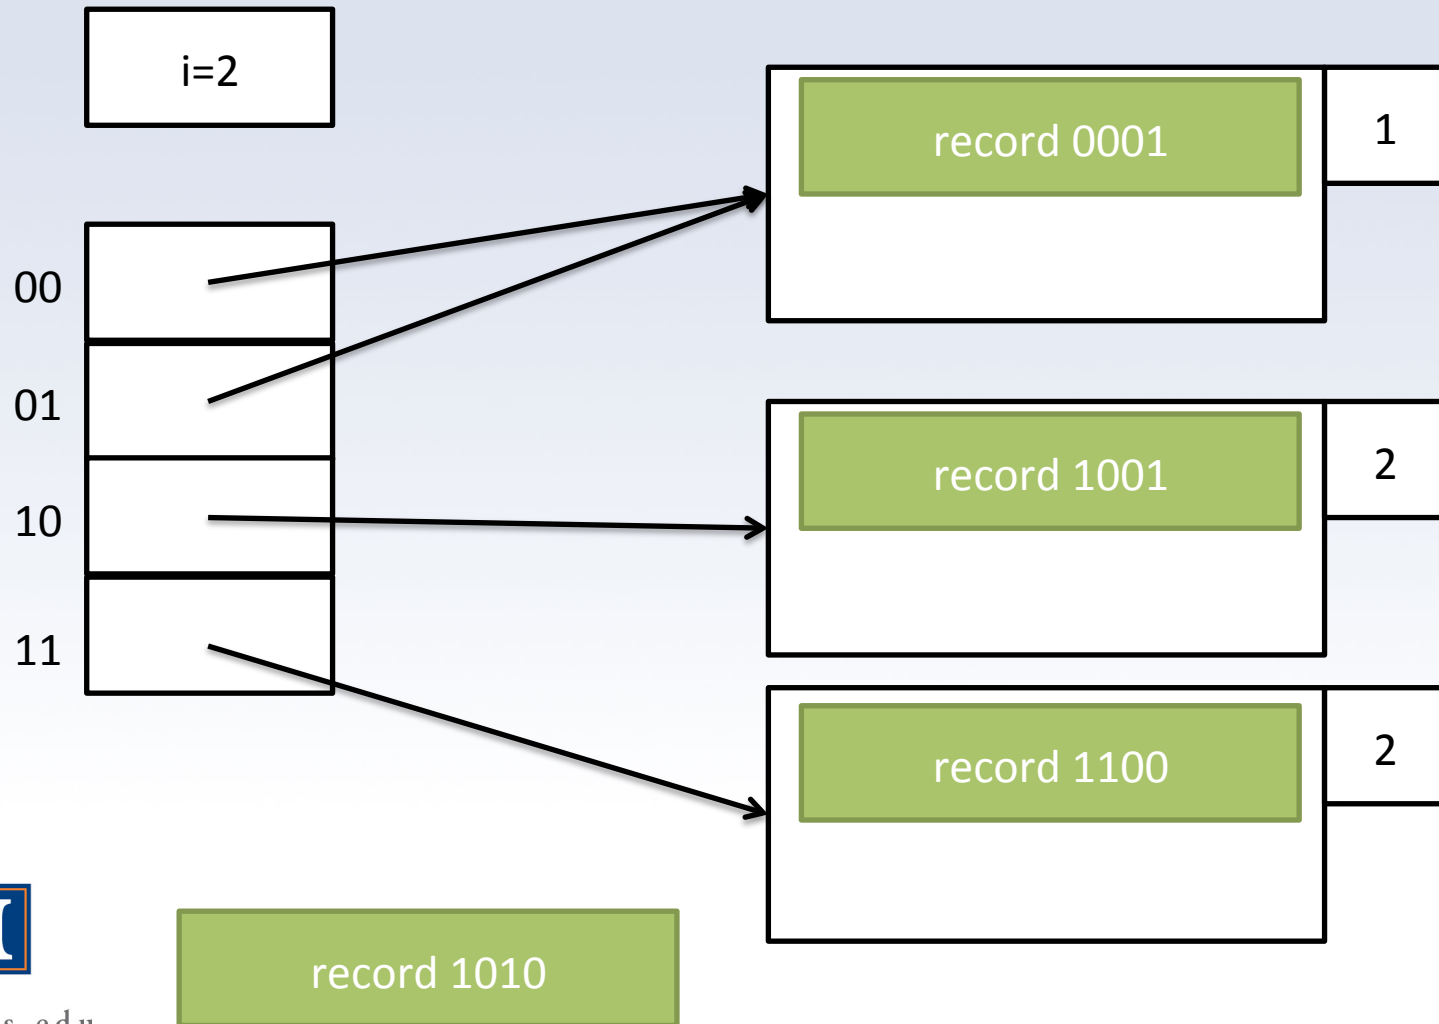
Example: Insert 1010




Example: Insert 1010




Example: Insert 1010



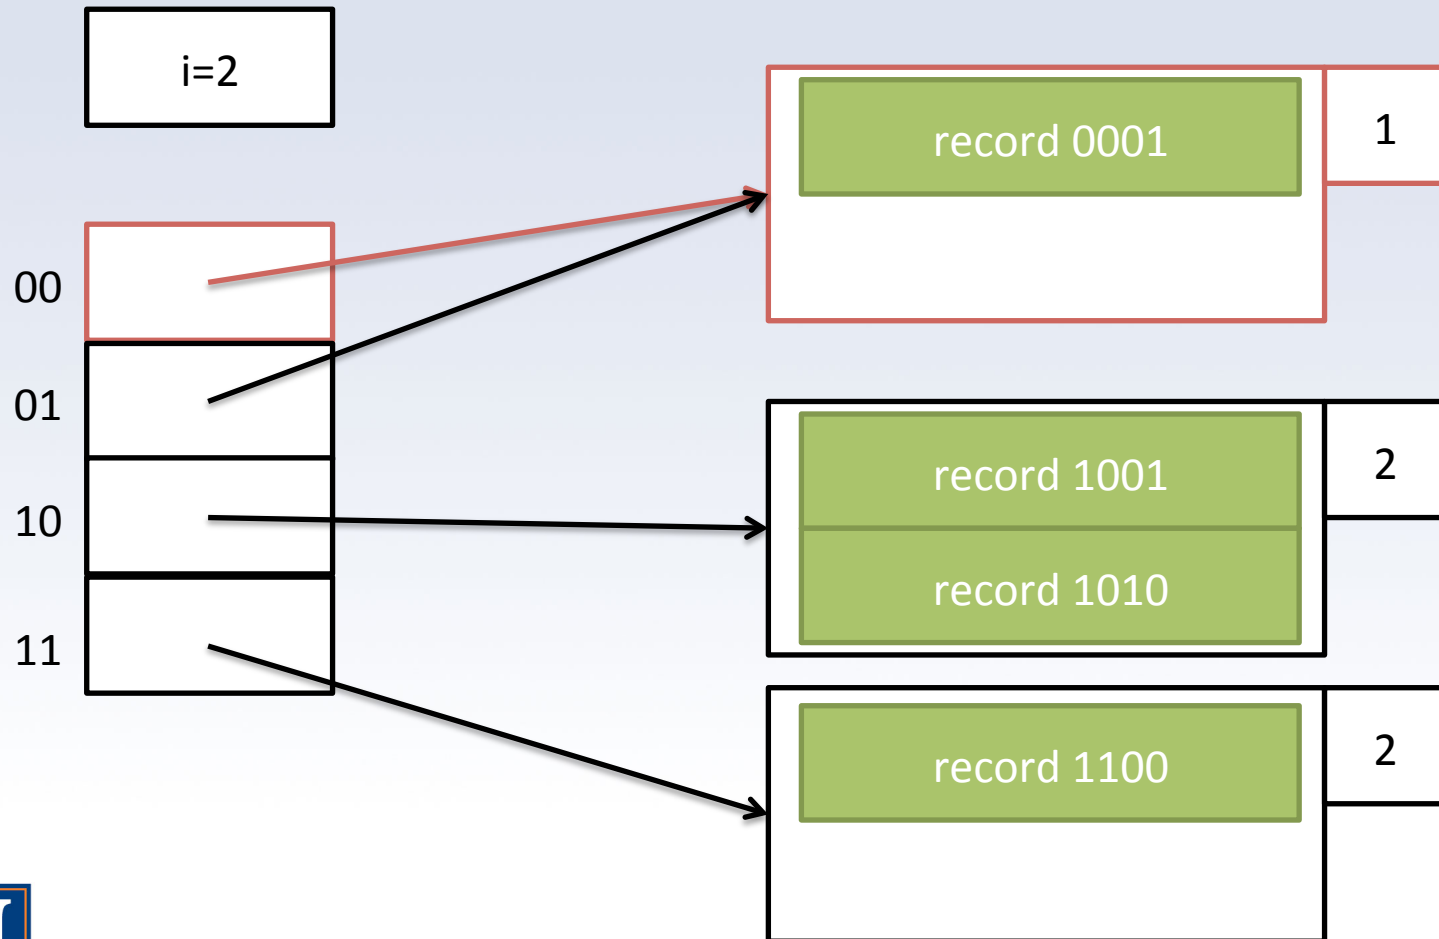
Example: Insert 1010




Example: Insert 0000



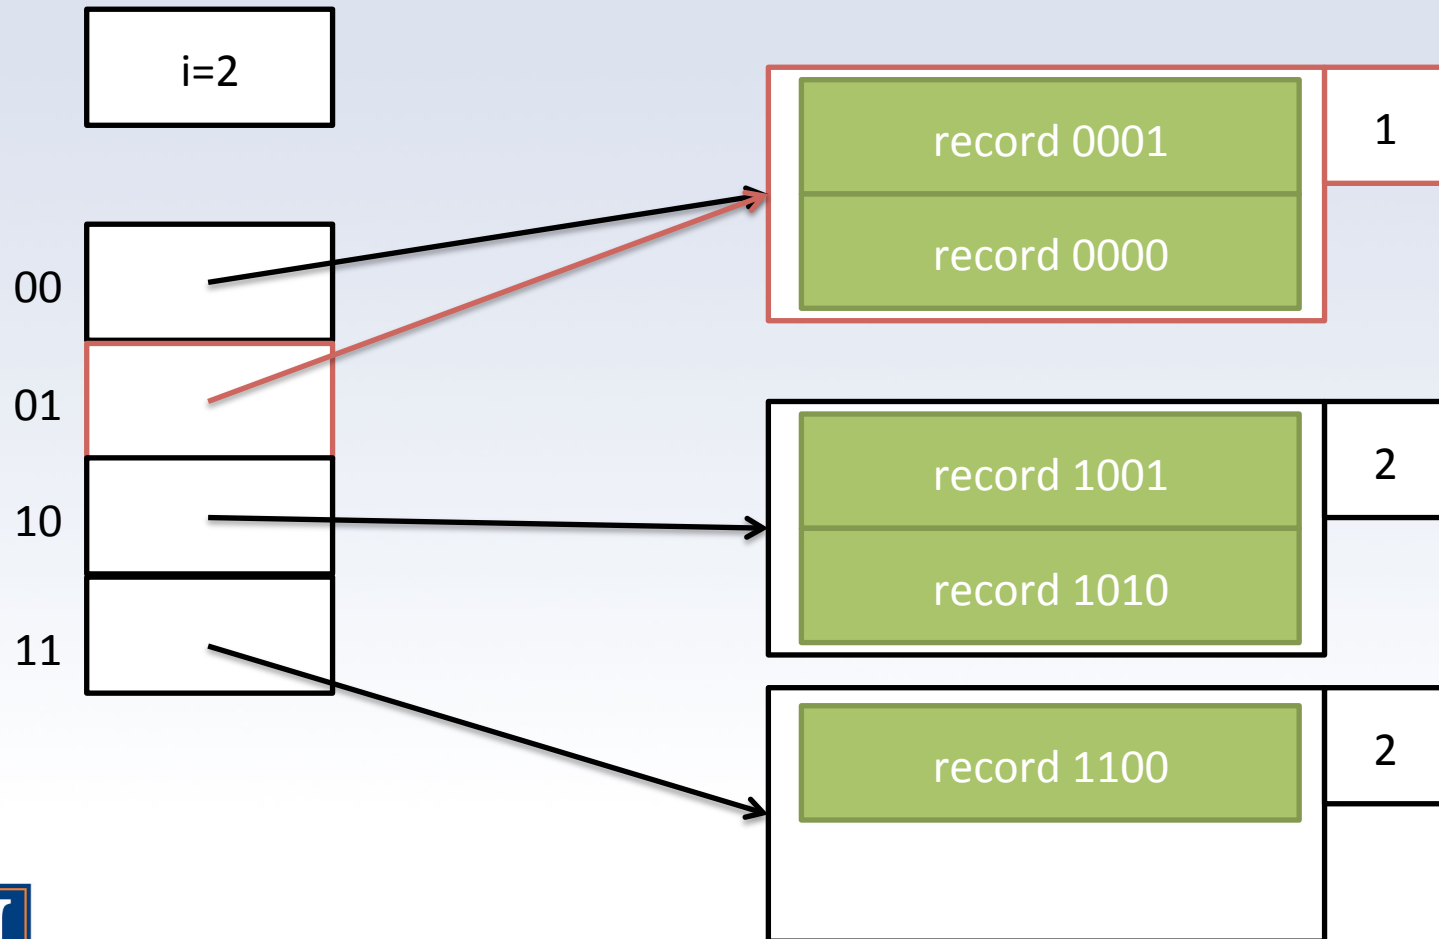
Example: Insert 0000



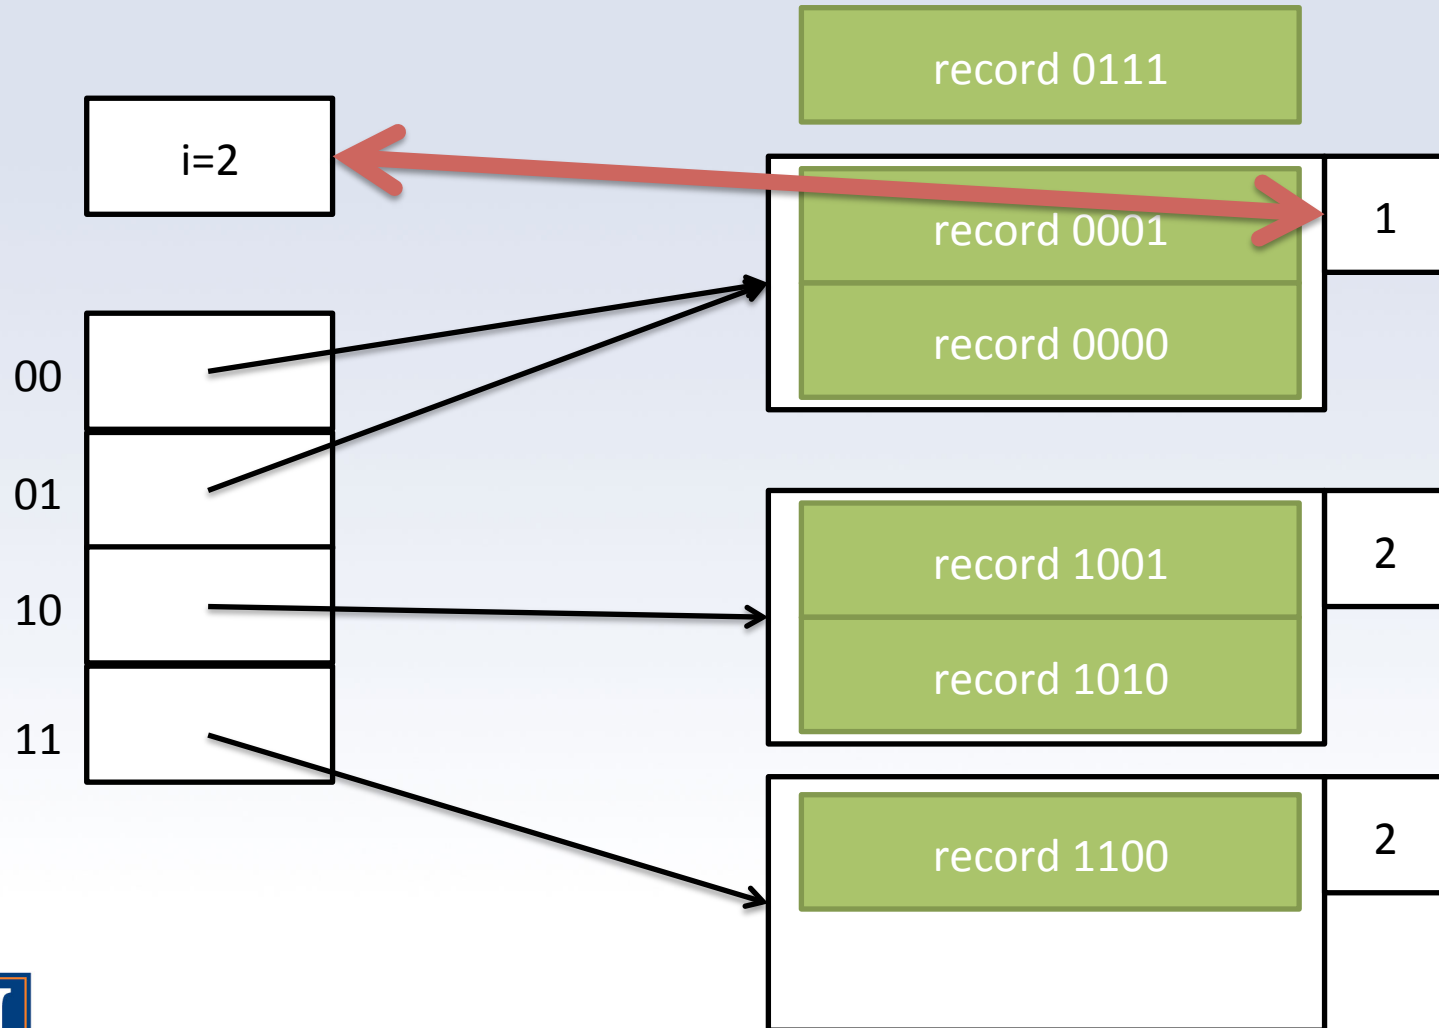
Example: Insert 0000



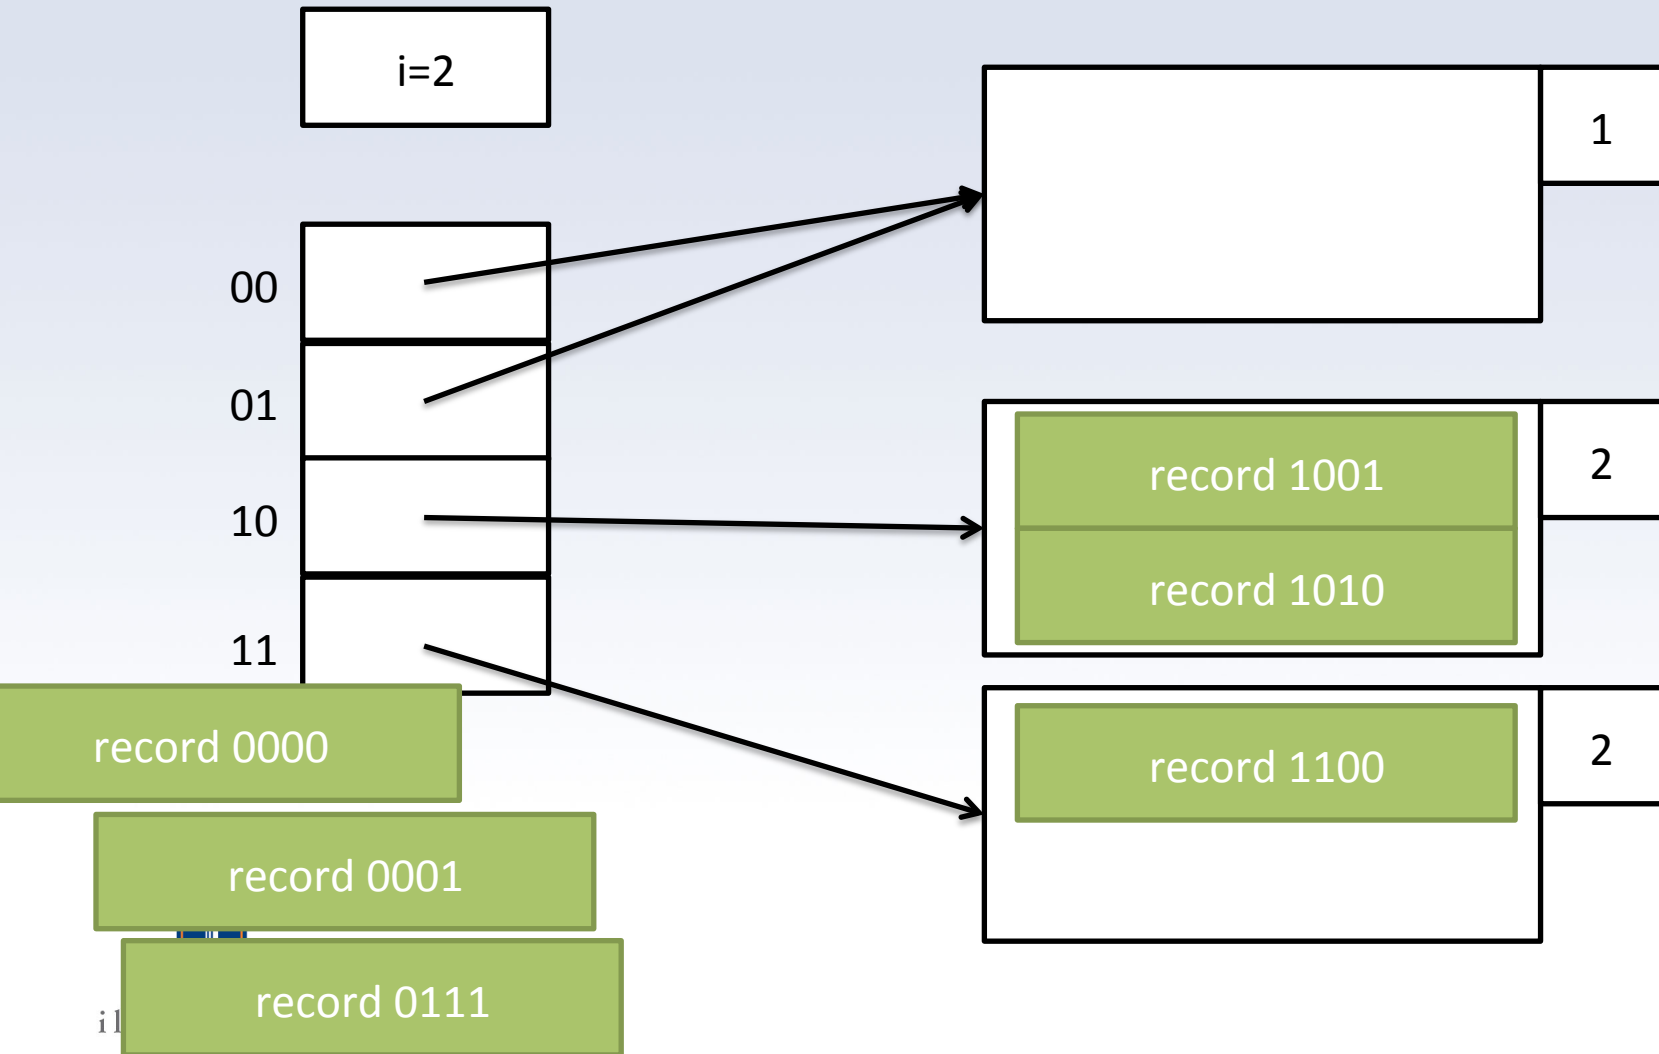
Example: Insert 0111



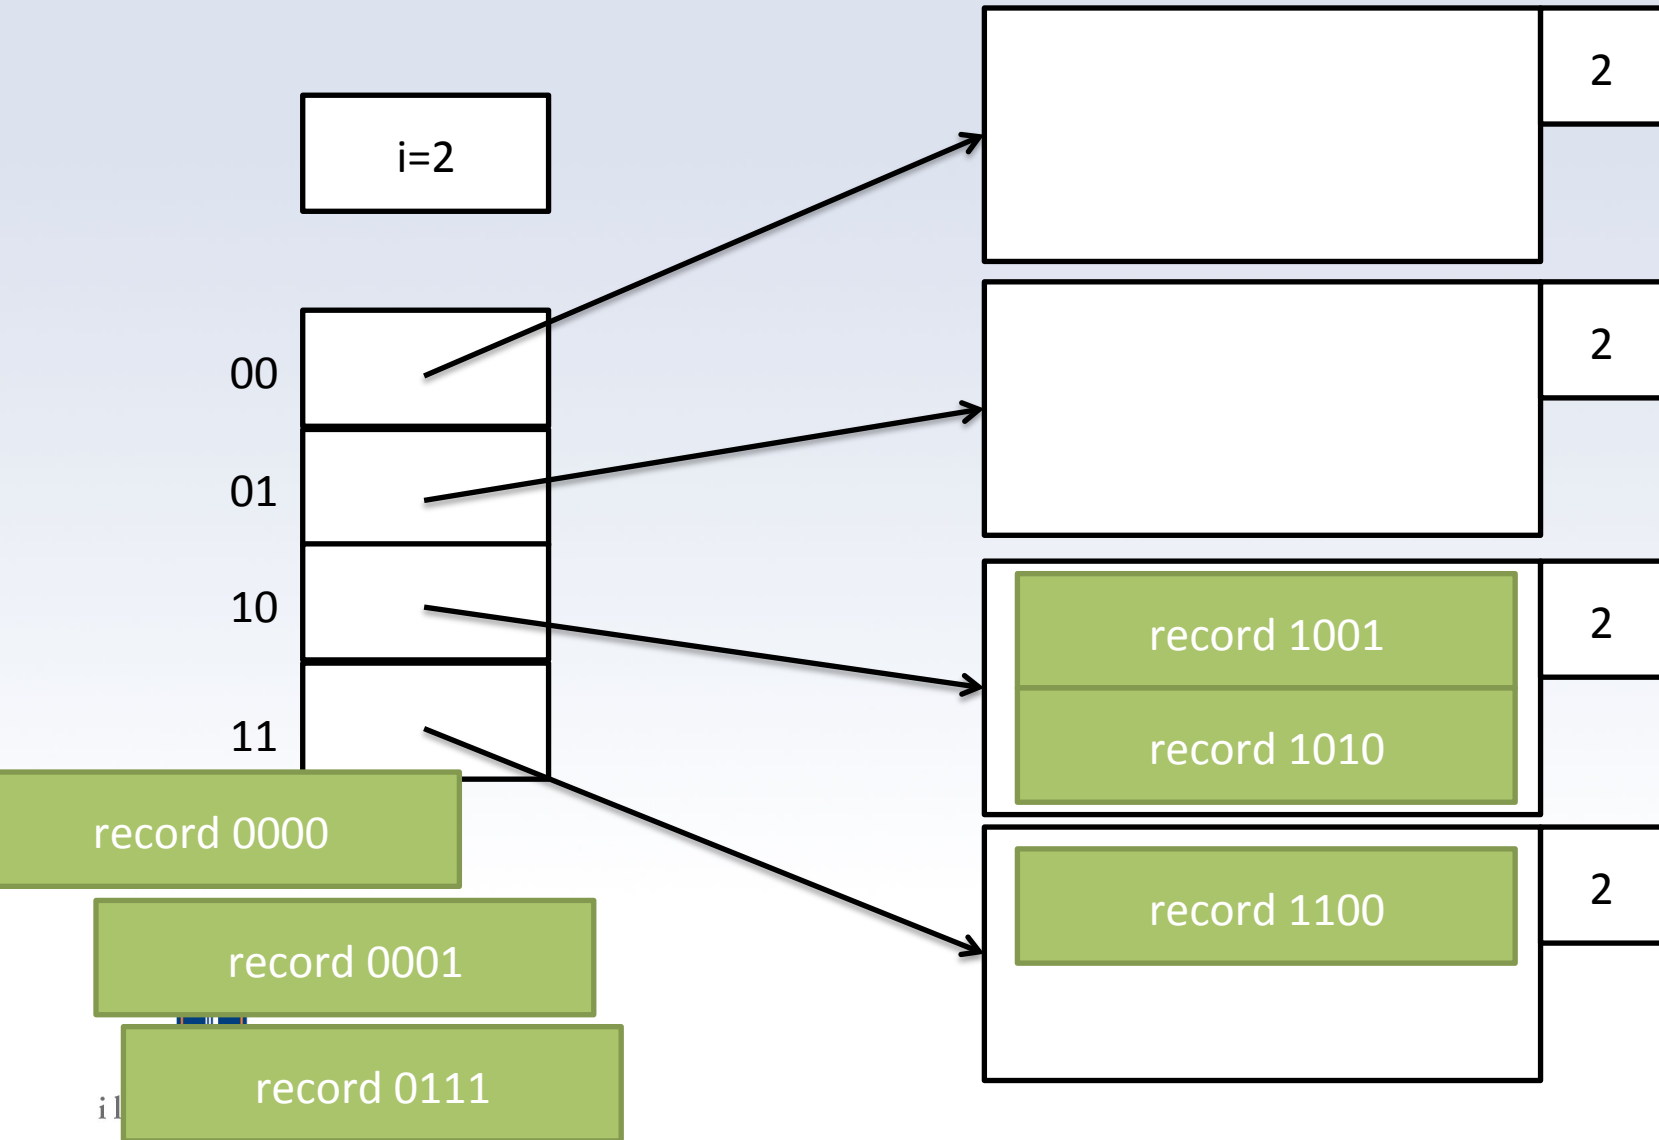
Example: Insert 0111



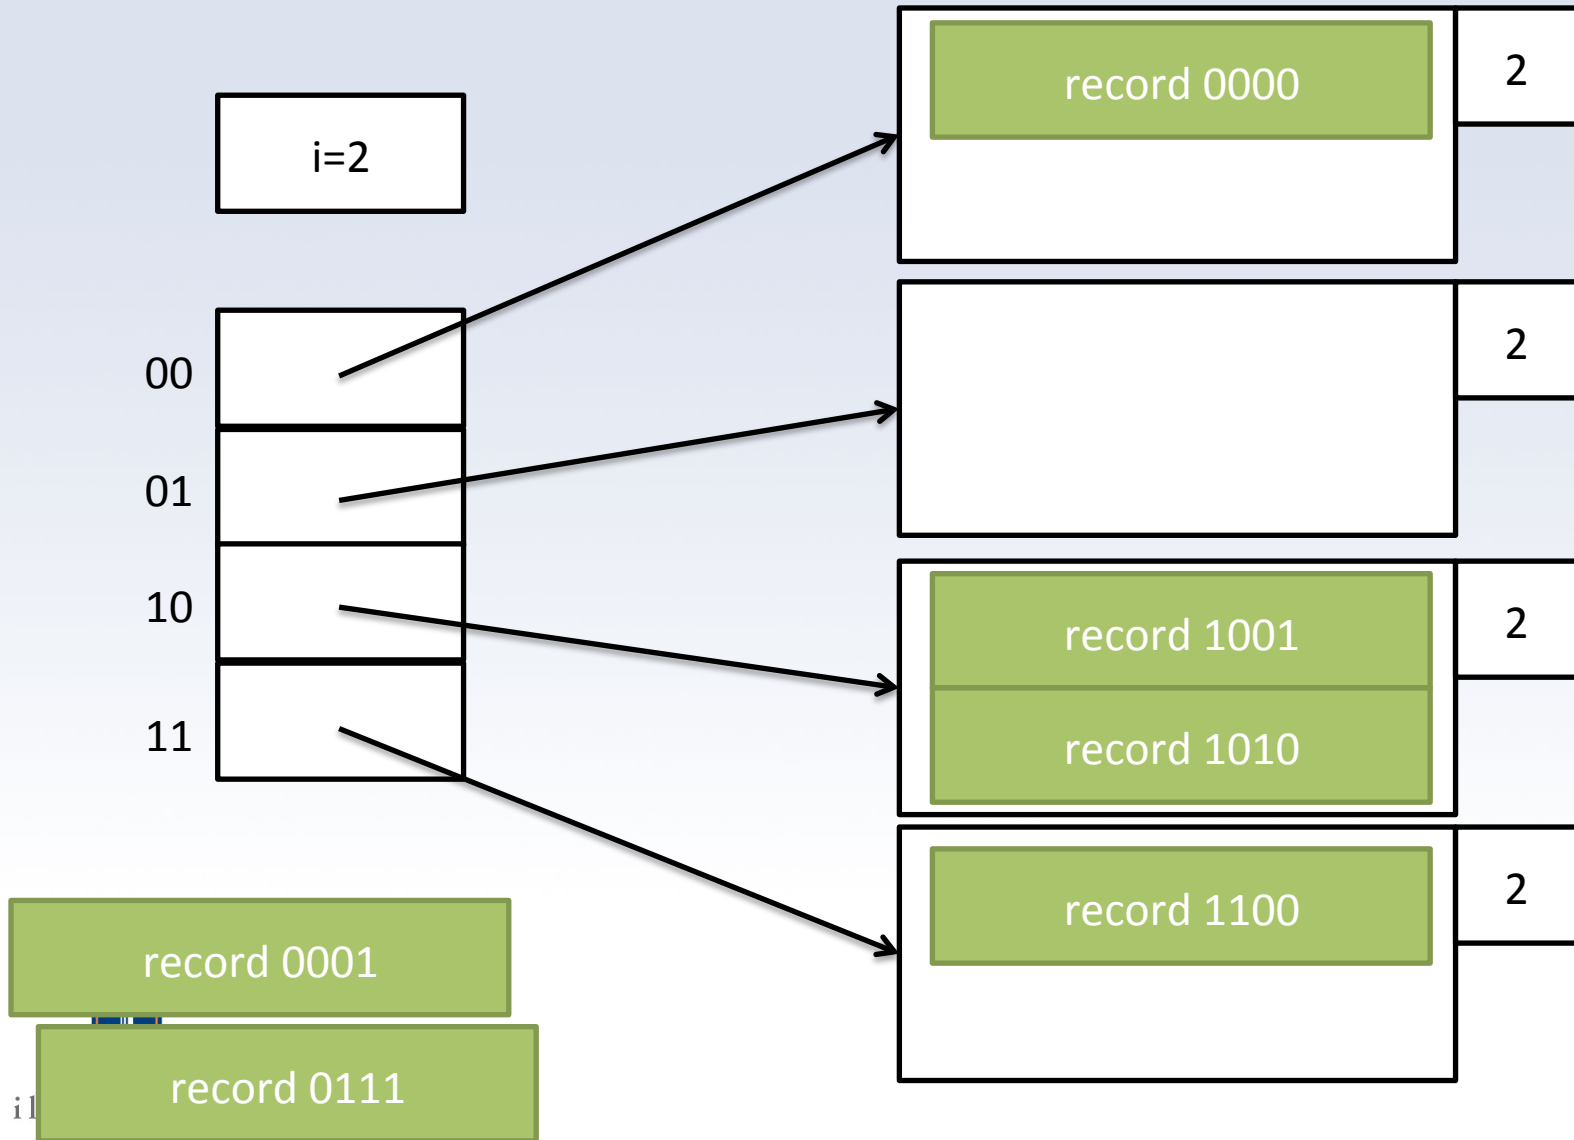
Example: Insert 0111



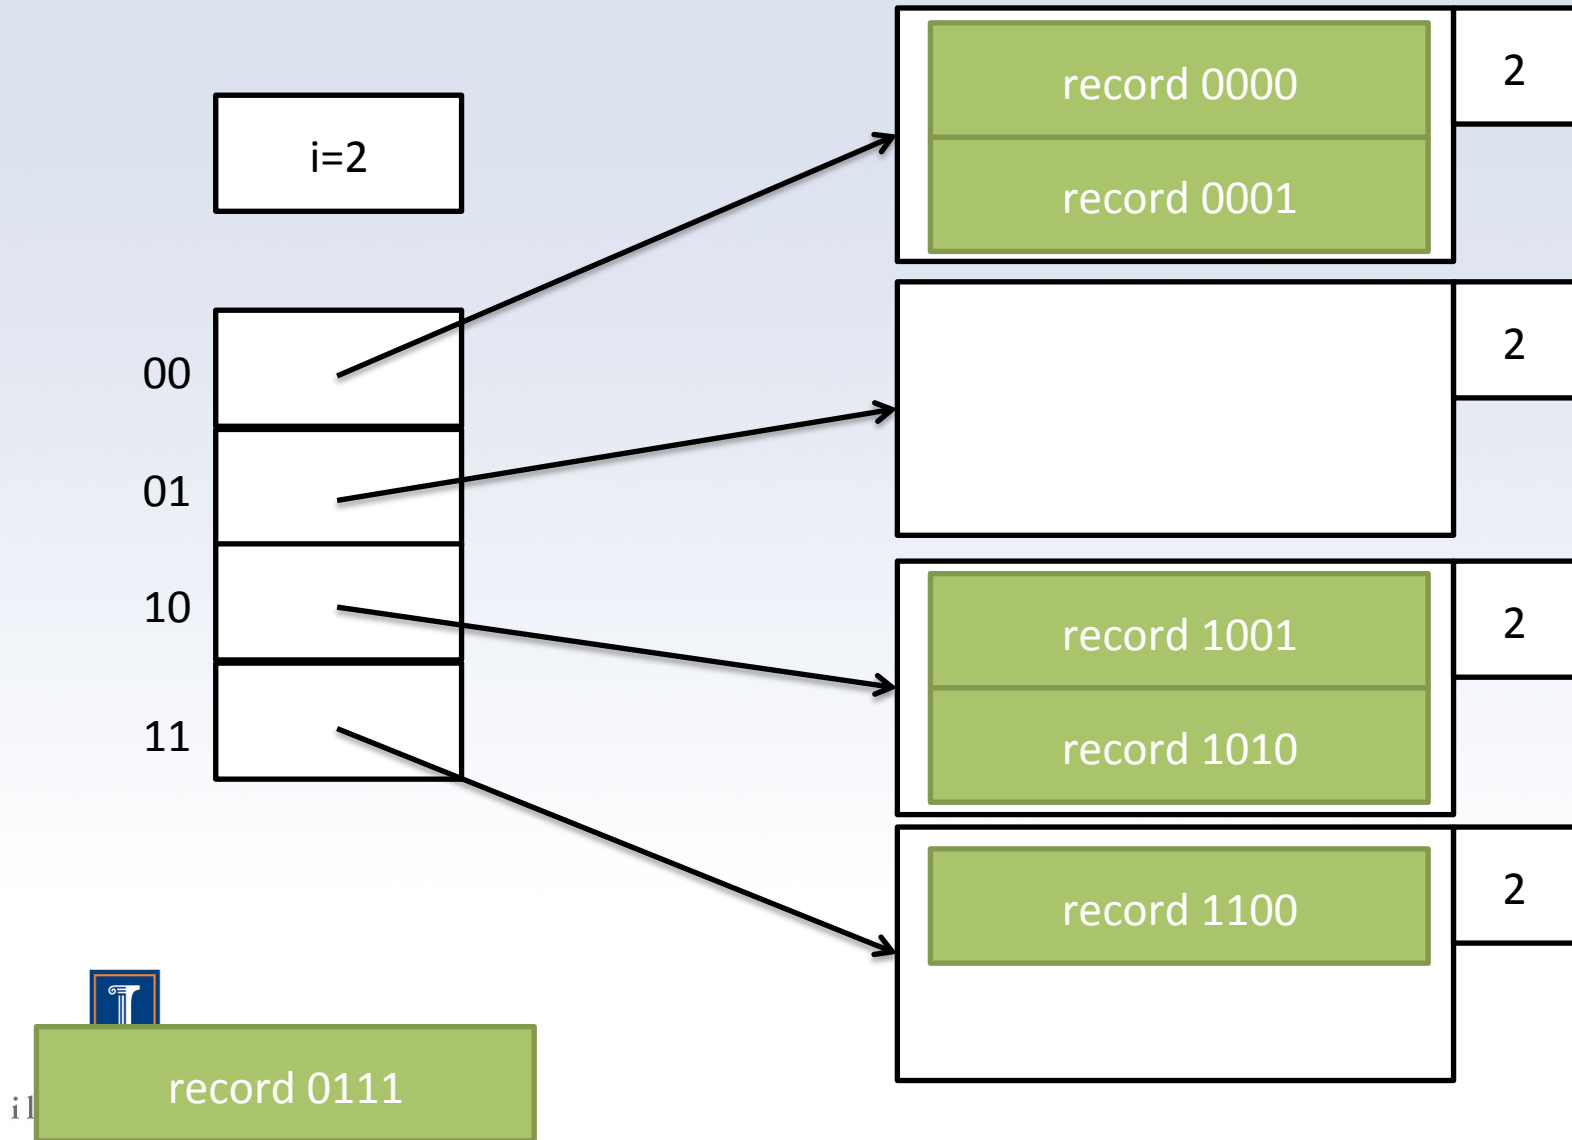
Example: Insert 0111




Example: Insert 0111




Example: Insert 0111



Example: Insert 0111



Example: Insert 1000



record 1000

i=2

00	
01	
10	
11	

record 0000	2
record 0001	

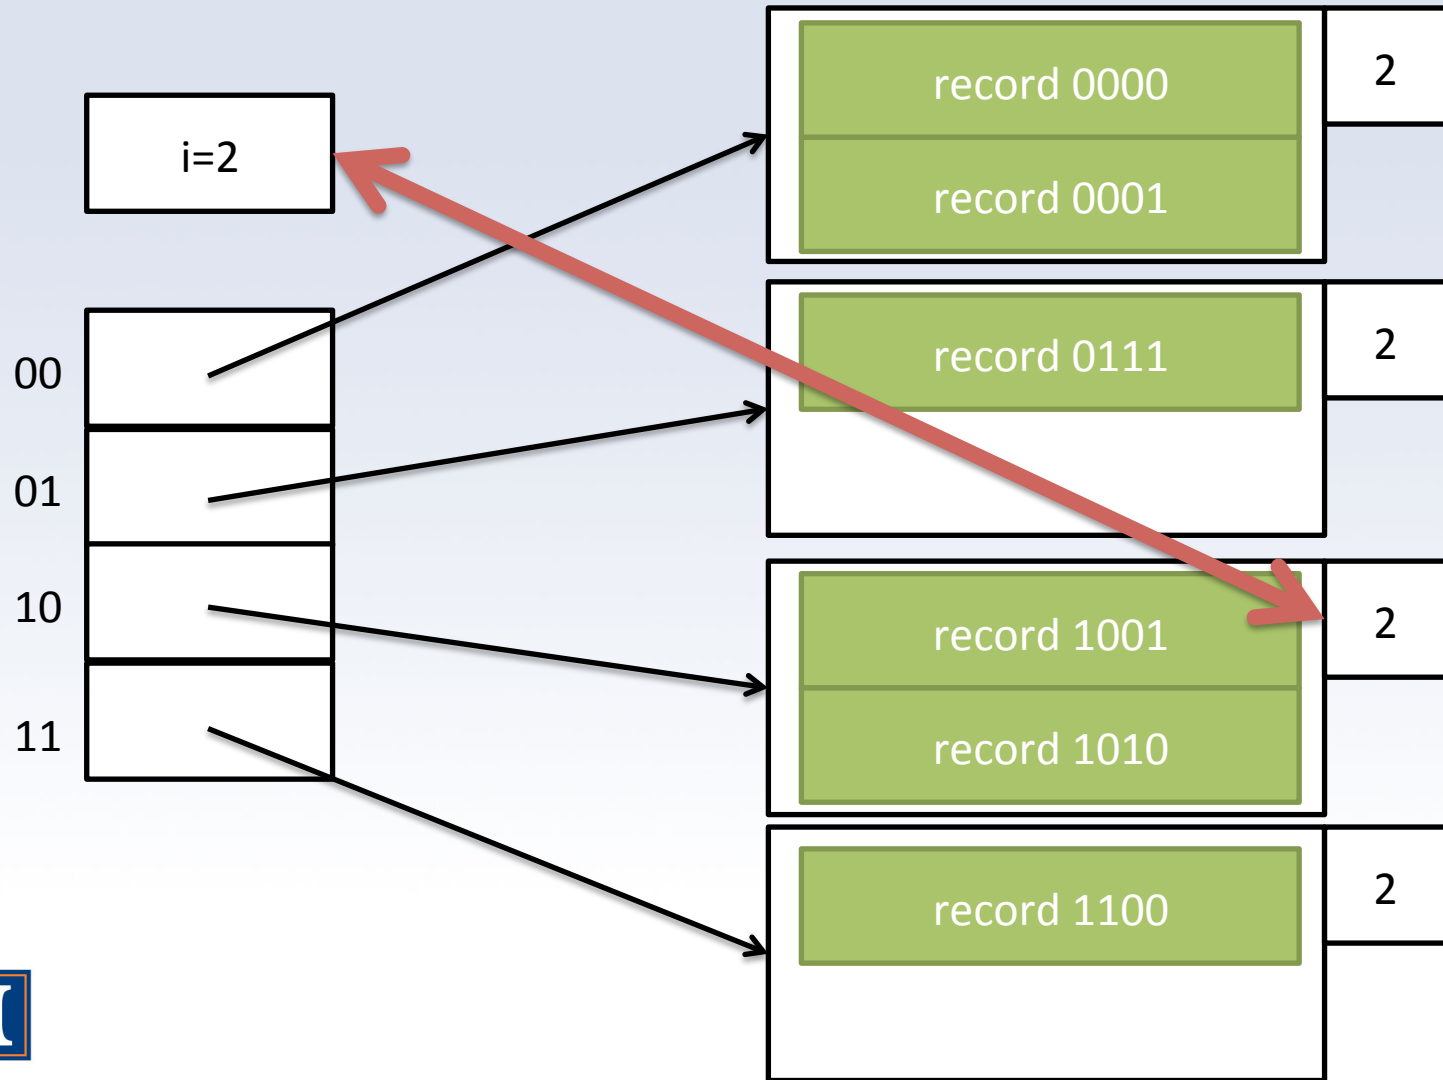
record 0111	2

record 1001	2
record 1010	

record 1100	2



record 1000



record 1000

i=3

000

001

010

011

100

101

110

111

record 0000
record 0001

2

record 0111

2

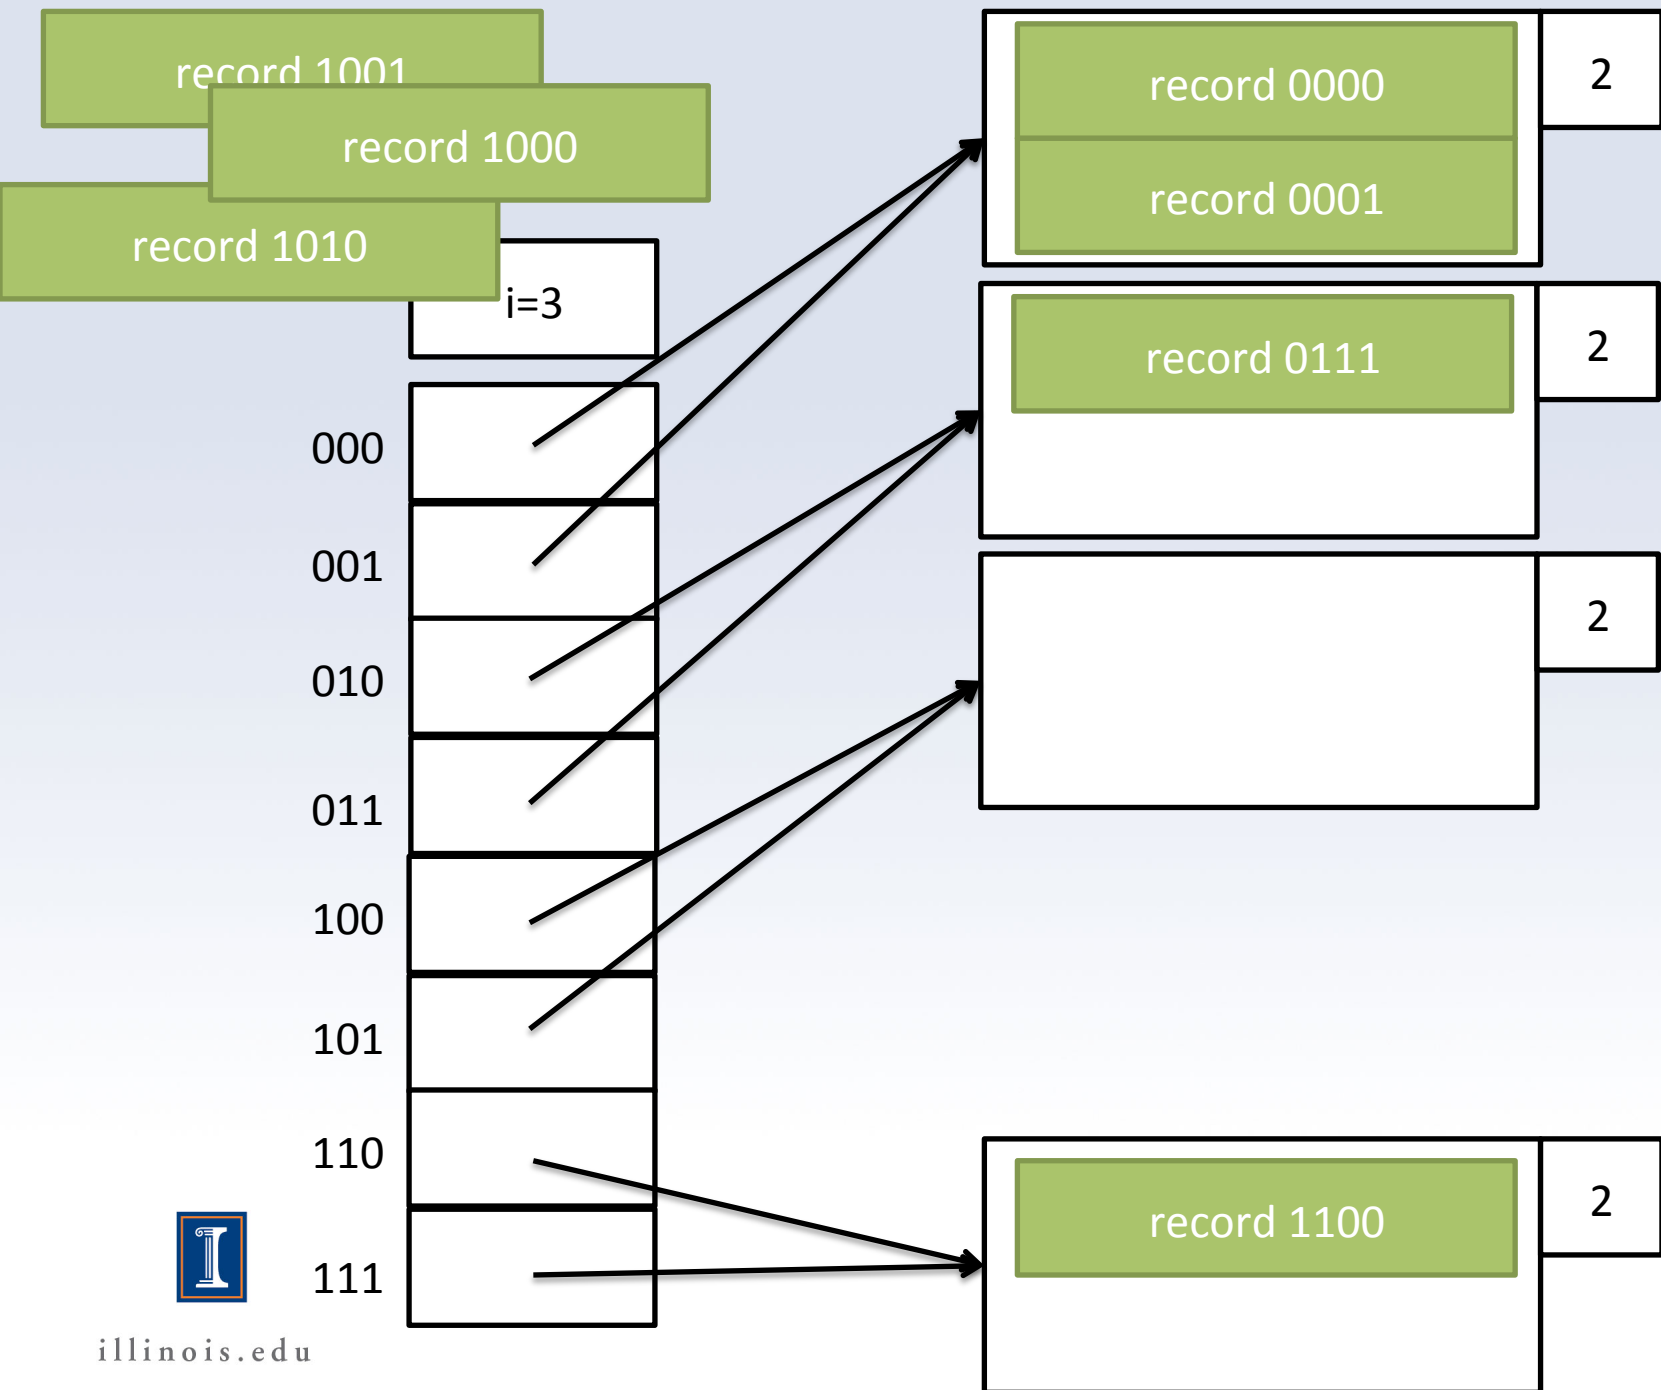
record 1001
record 1010

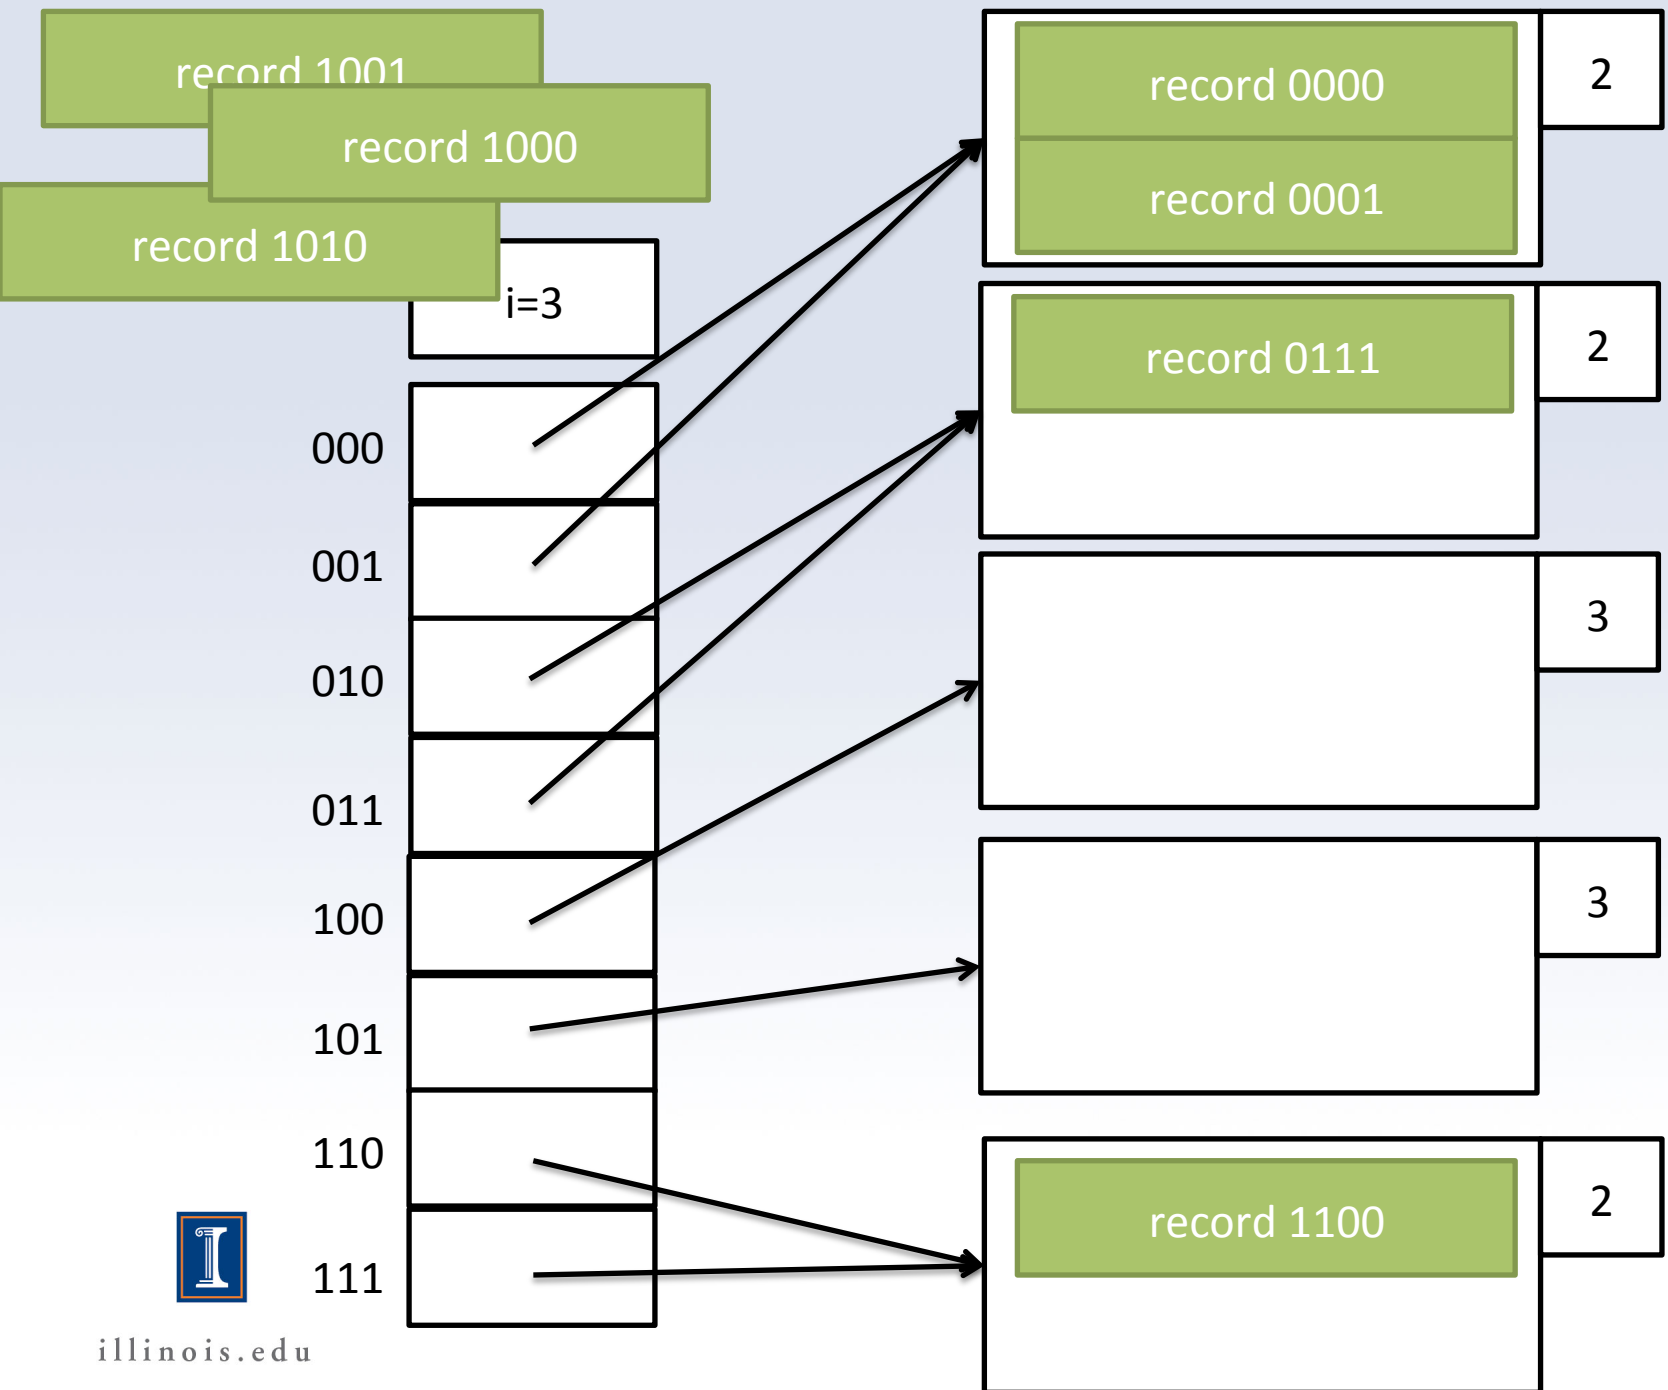
2

record 1100

2







record 1001

record 1000

i=3

000

001

010

011

100

101

110

111

record 0000

record 0001

2

record 0111

2

3

record 1010


3

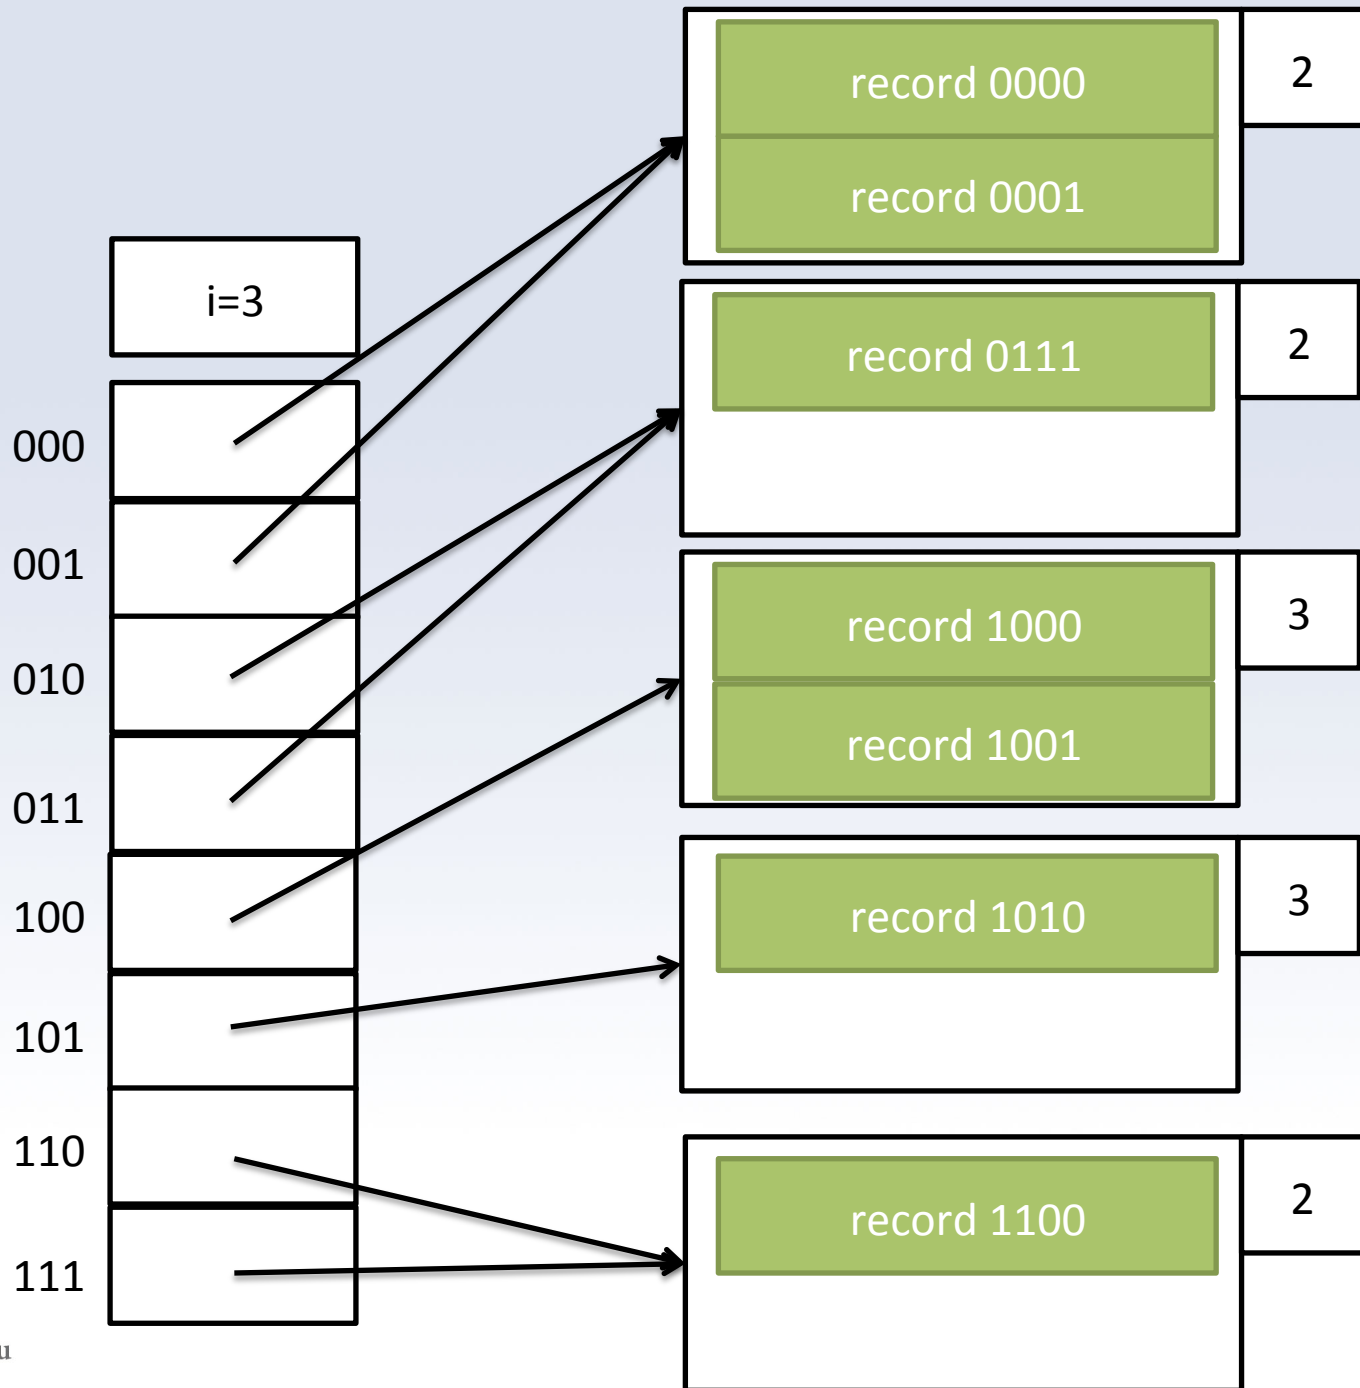
record 1100

2



record 1001





Extensible Hash

- Disadvantages:
 - If many keys have the same hash, hash array size explodes
 - Lots of work when hash array doubles
 - Doubling of hash array might make it too big for main memory



Linear Hash

- Average number of records per bucket fixed (e.g. $r/n \leq 1.7$)
- Number of buckets grows linearly
- Overflow blocks permitted



Linear Hash


- need to track 3 variables
 - i = the number of bits of the hash we use
 - r = the number of records inserted
 - n = the number of buckets in the hash
- $i = \text{ceiling}(\log_2 n)$
- increment n if r/n exceeds threshold



Example

i=1
n=2
r=3

$$r/n=1.5$$




Linear Hash Lookup

- Interpret the last i blocks of the key as an integer m
- Find block m
 - won't exist if $m \geq n$
 - switch the high order bit of m to 0
 - look in that block



Example: Lookup 1010

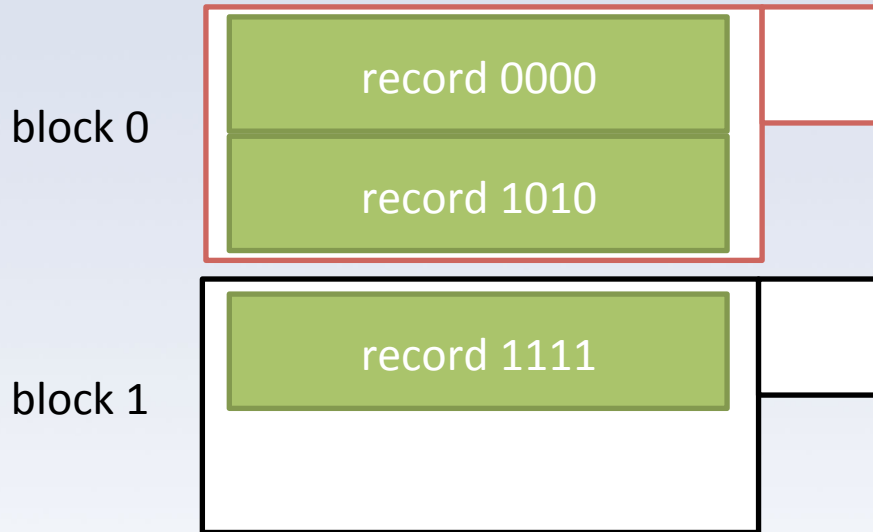
i=1
n=2
r=3



Example: Lookup 1010

i=1
n=2
r=3

$$m=0 < n$$

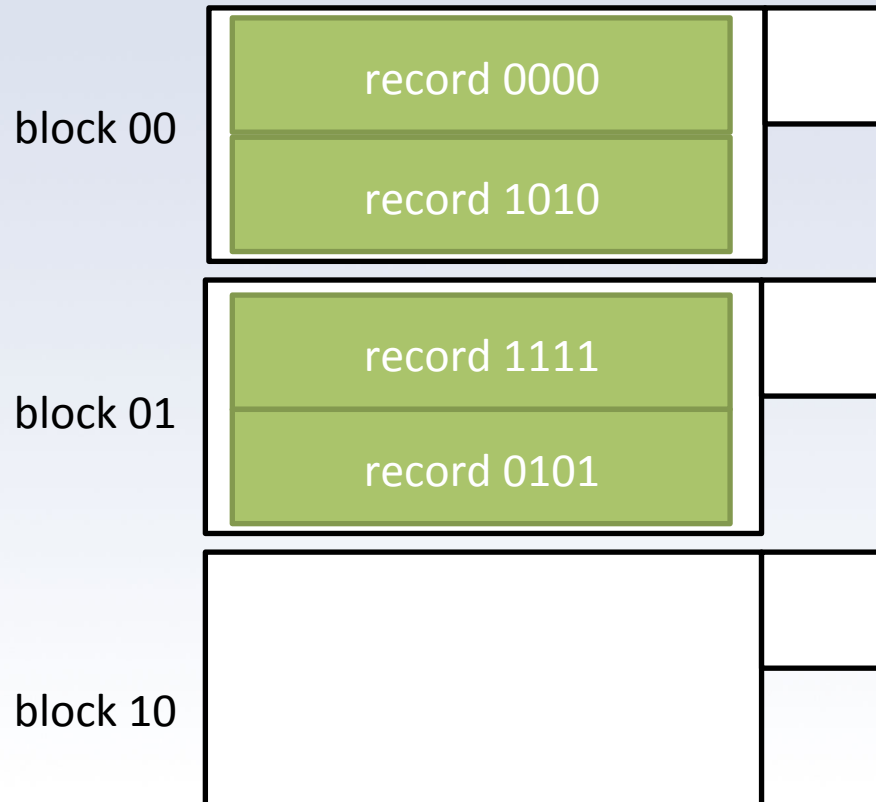


Block 0 exists



Example: Lookup 1111

i=2
n=3
r=4

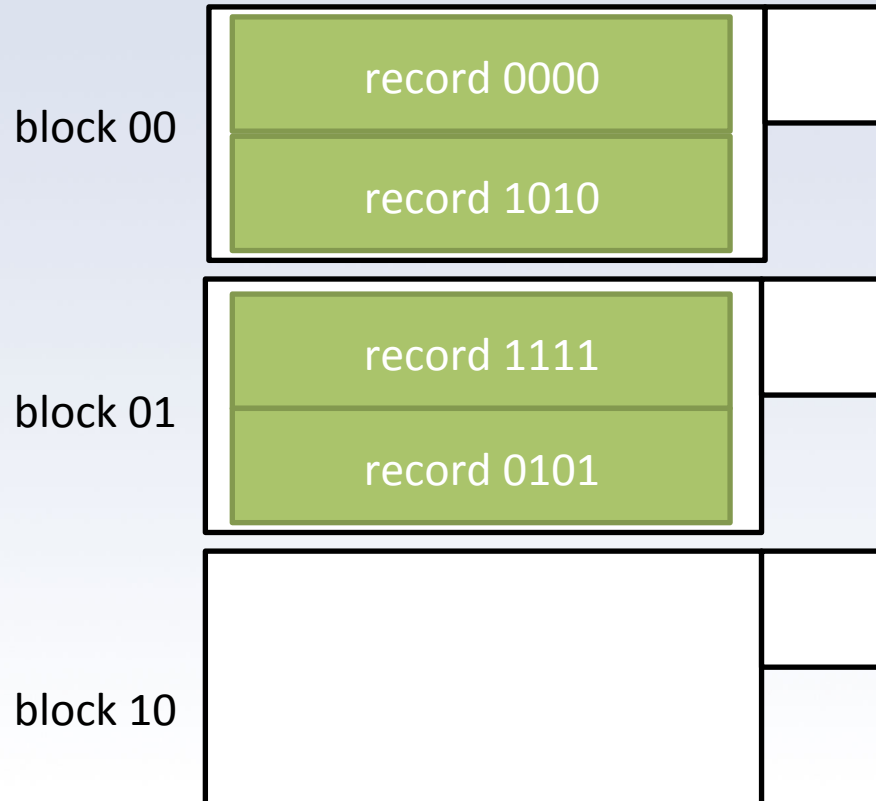


Example: Lookup 1111

i=2
n=3
r=4

$$m = 3 \geq n$$

Block 11 does
not exist

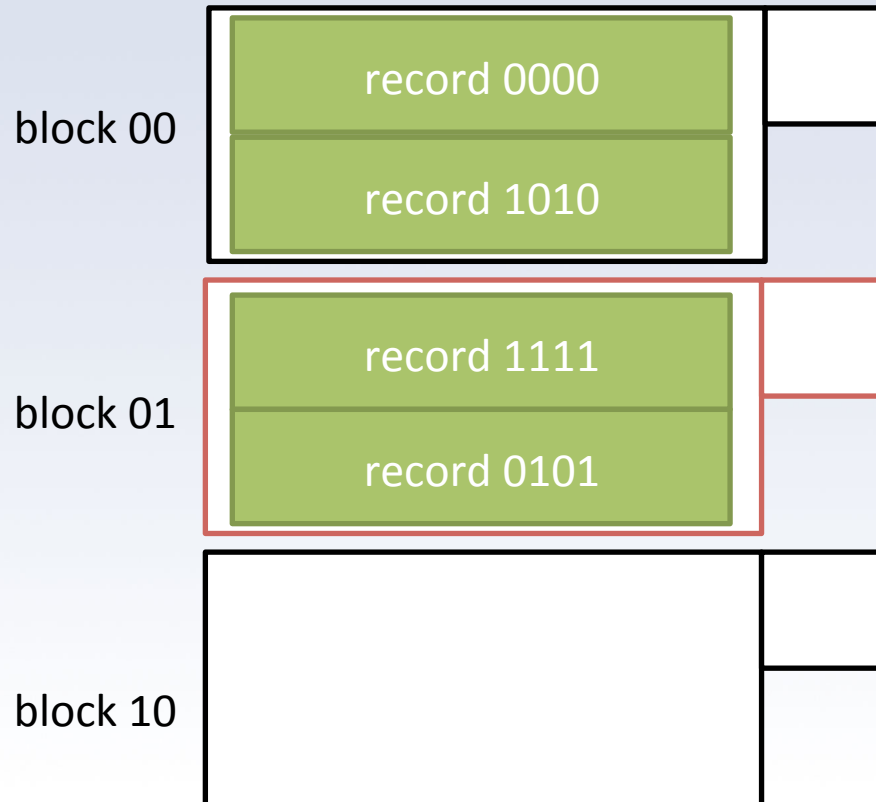


Example: Lookup 1111

i=2
n=3
r=4

$$m = 3 \geq n$$

Look in block
01 instead




Linear Hash Insertion

- Lookup correct bucket
 - If room, insert. If not, create overflow.
 - If r/n is too big, add a new bucket
 - New bucket number is $1x$. Split bucket $0x$.
 - If n is too big, increment i



Example: Insert 0101

i=1
n=2
r=3




Example: Insert 010¹



i=1
n=2
r=3

1 < 2

block 0




block 1



Example: Insert 0101


i=1
n=2
r=4

$$r/n=2 > 1.7$$



Example: Insert 0101

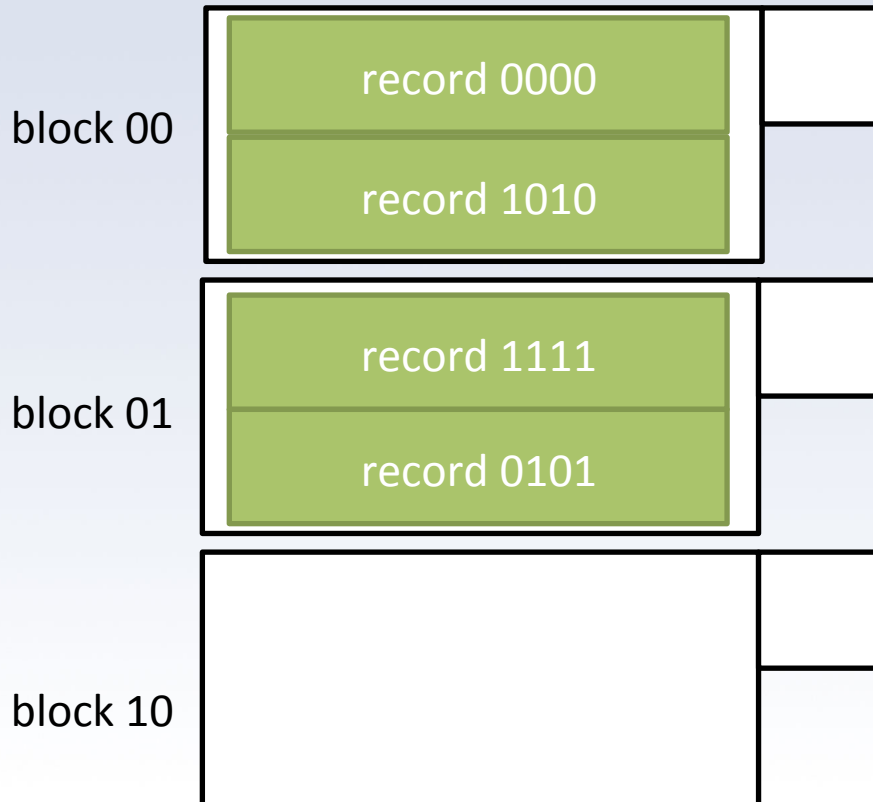
i=2
n=3
r=4



Example: Insert 0101

i=2
n=3
r=4

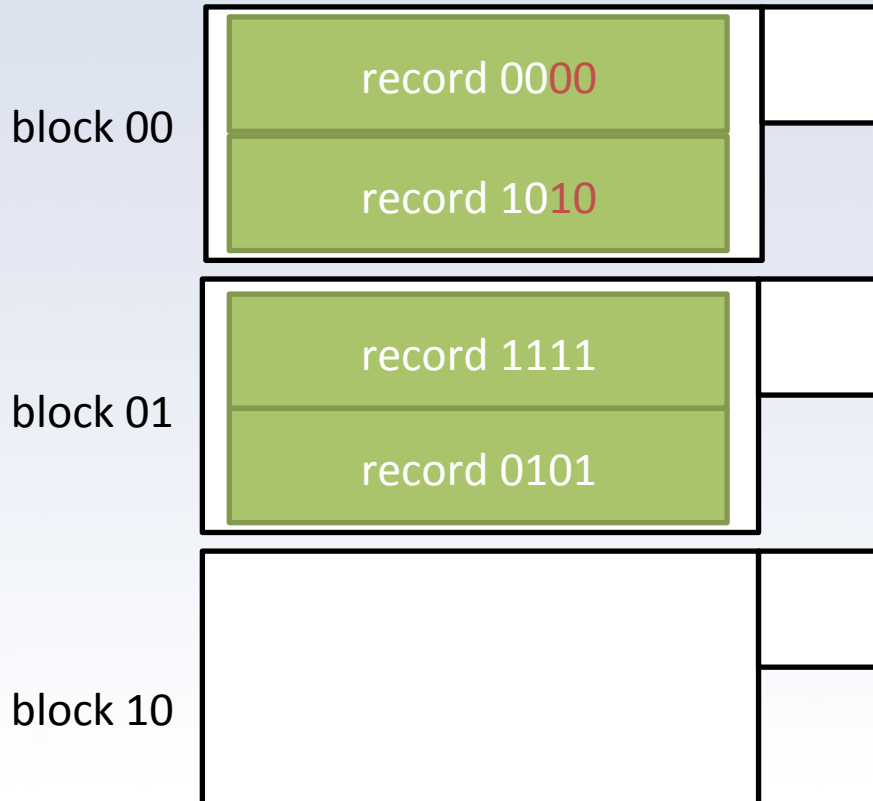
Split bucket 00



Example: Insert 0101

i=2
n=3
r=4

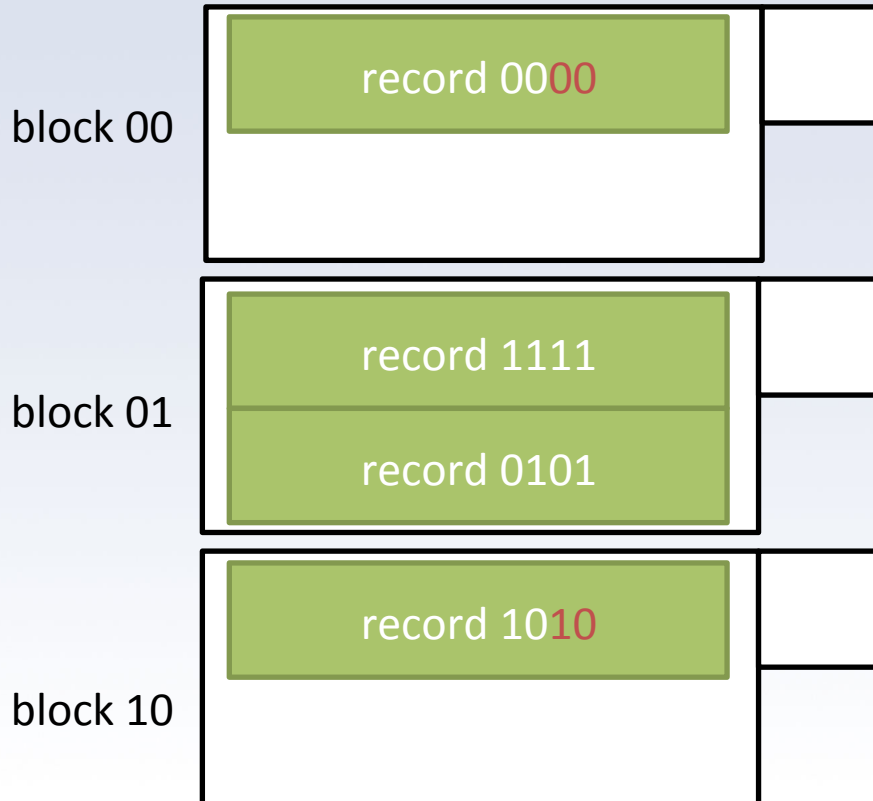
Split bucket 00



Example: Insert 0101

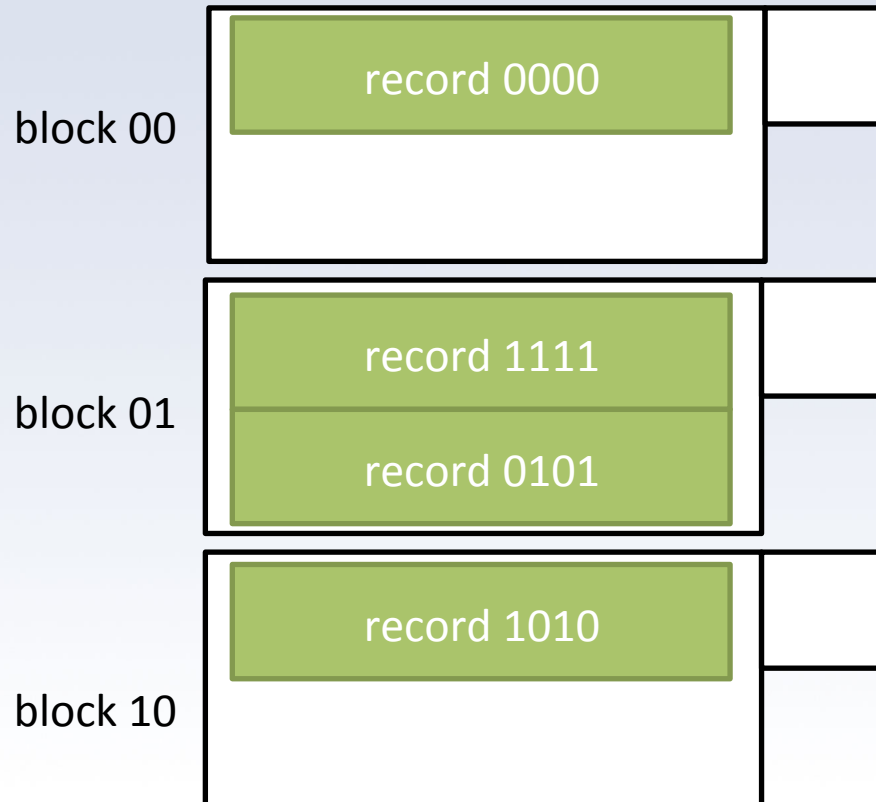
i=2
n=3
r=4

Split bucket 00



Example: Insert 0001

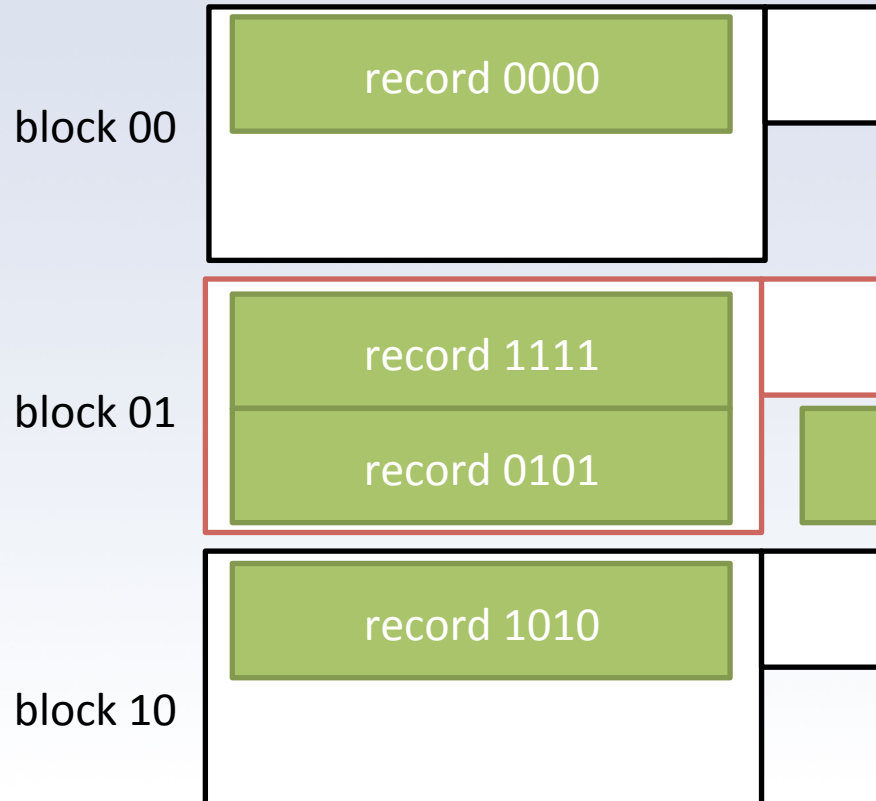
i=2
n=3
r=4



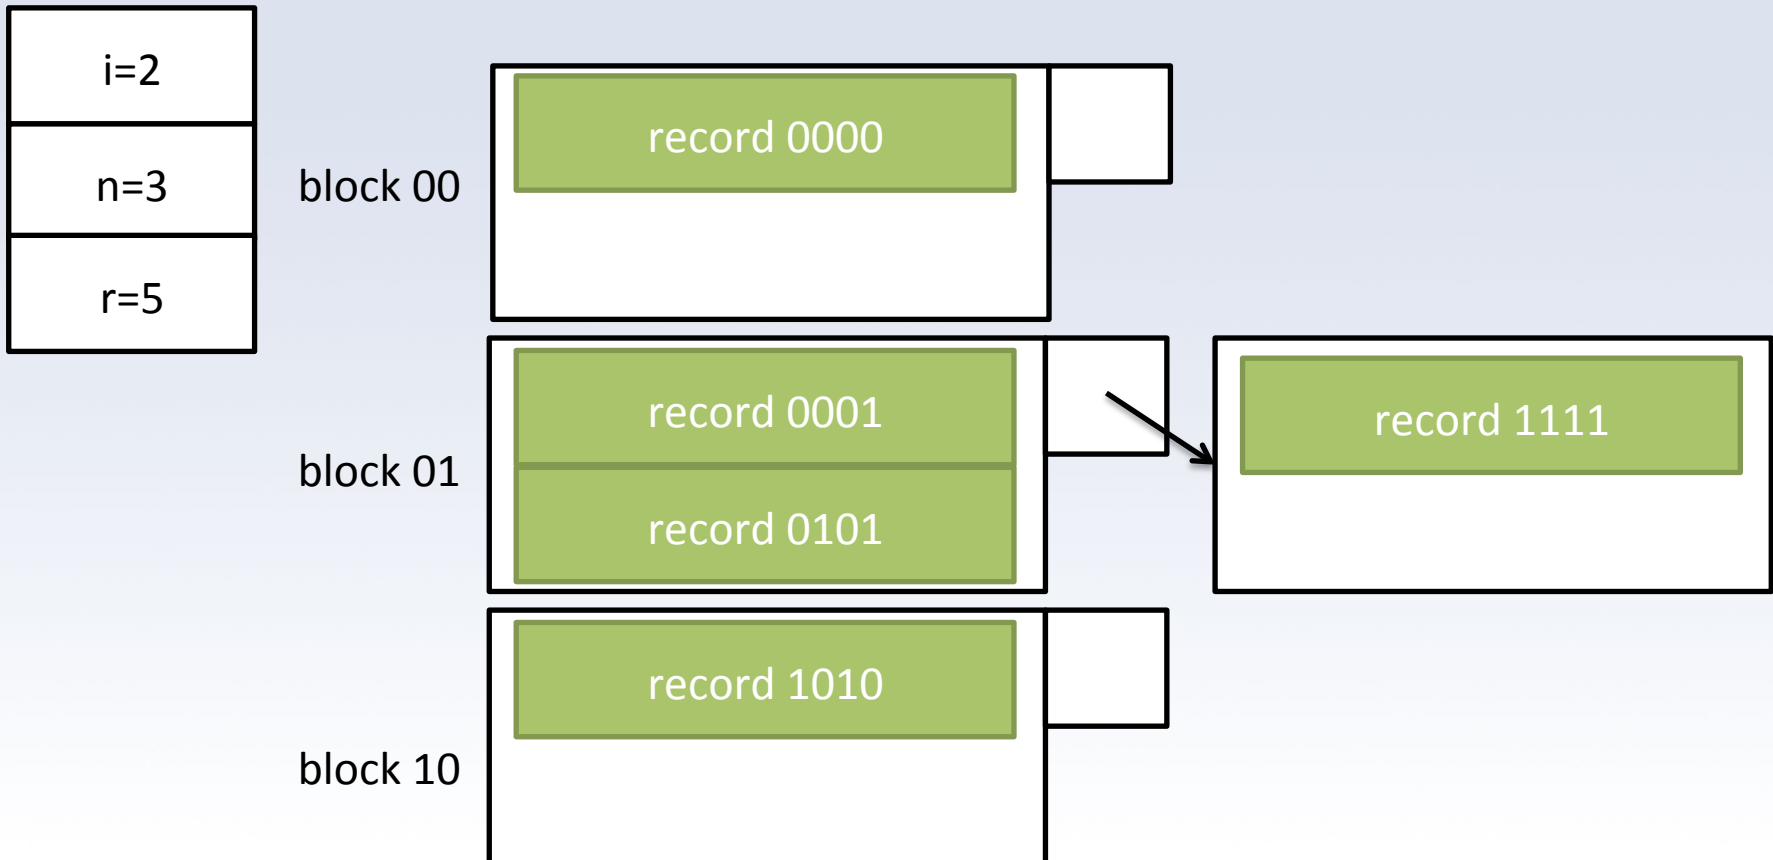
Example: Insert 0001

i=2
n=3
r=5

$$r/n = 1.67 < 1.7$$

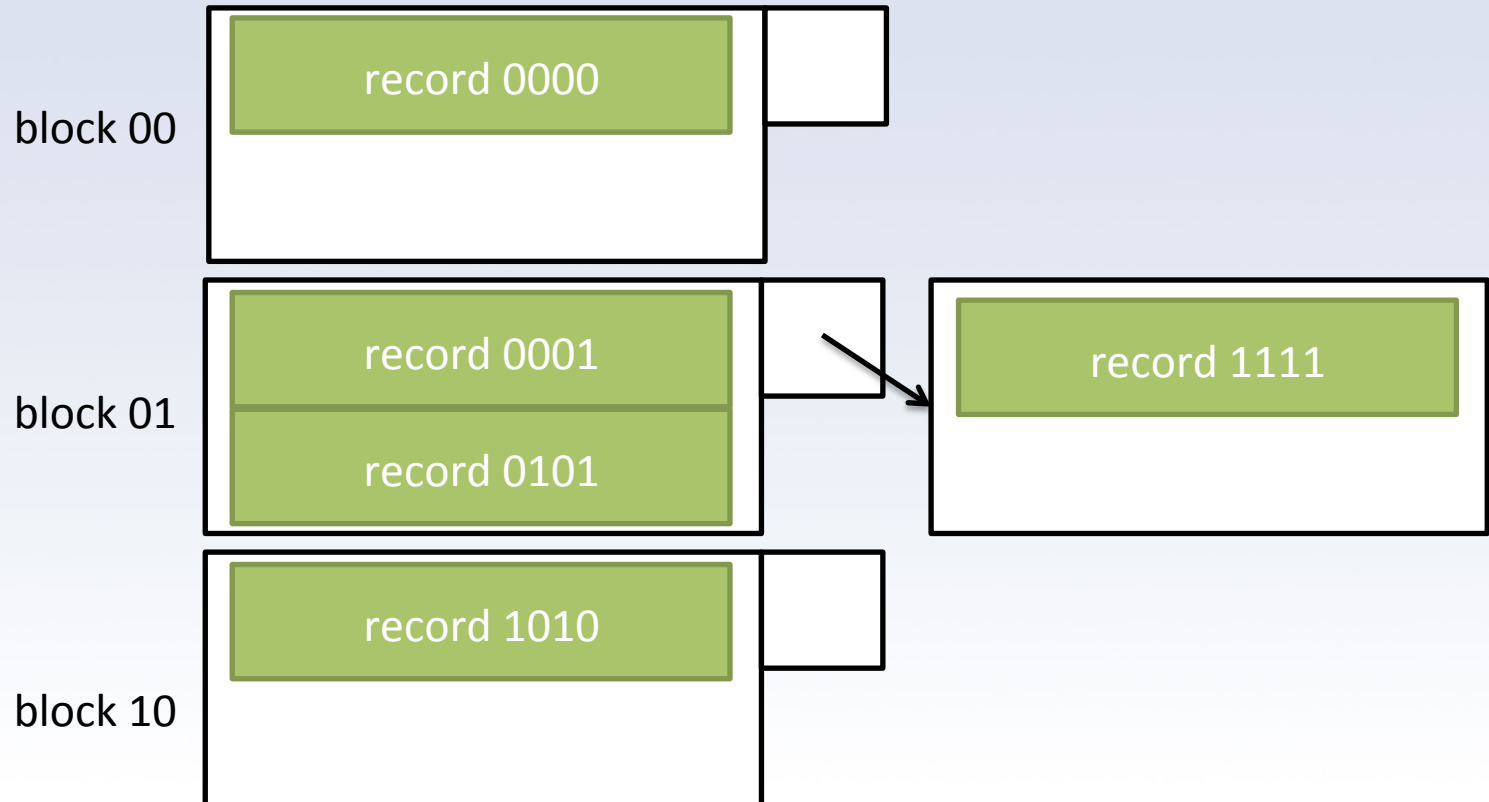


Example: Insert 0001



Example: Insert 0111

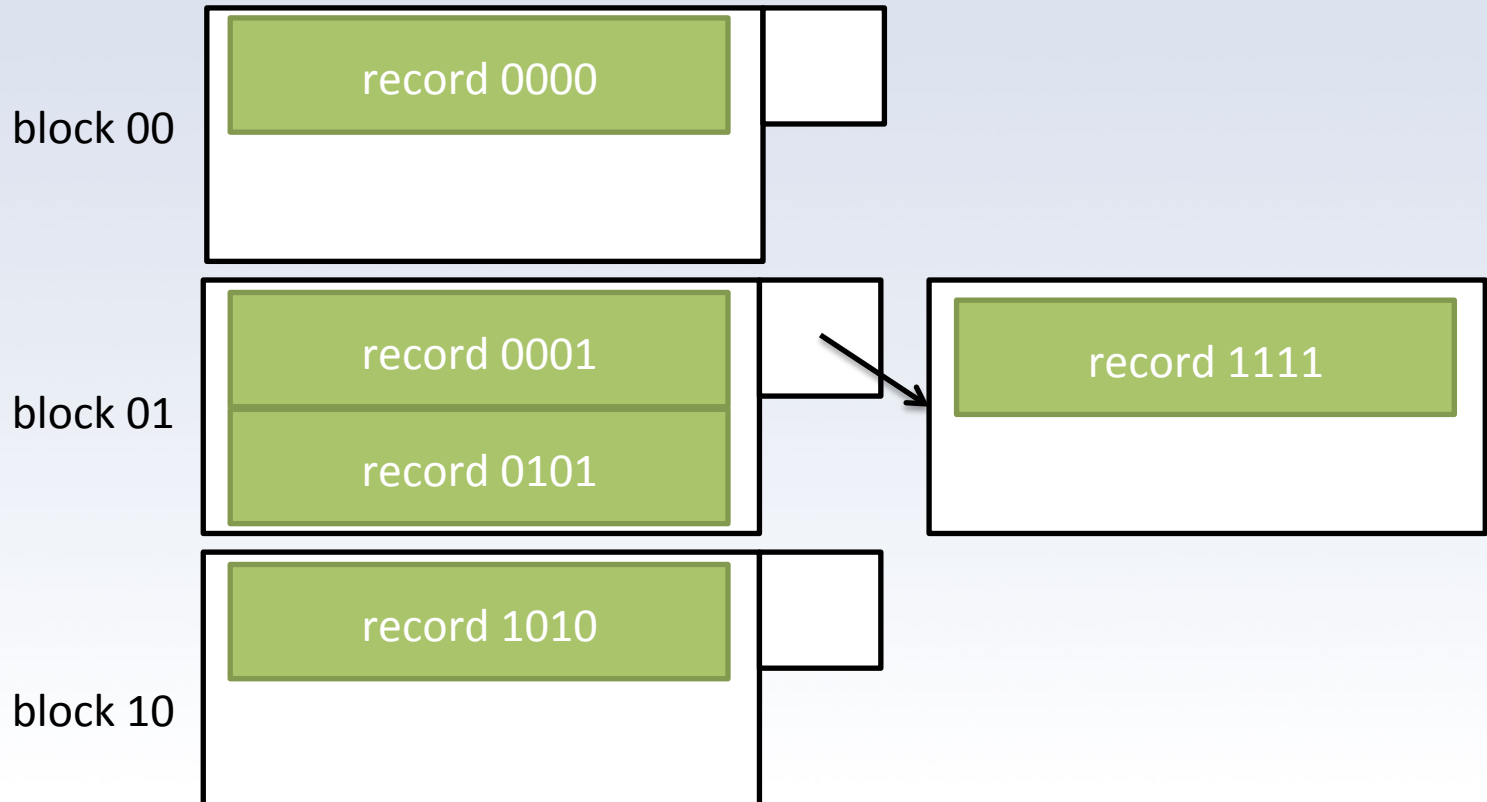
i=2
n=3
r=5



Example: Insert 0111

i=2
n=3
r=5

3 \geq n

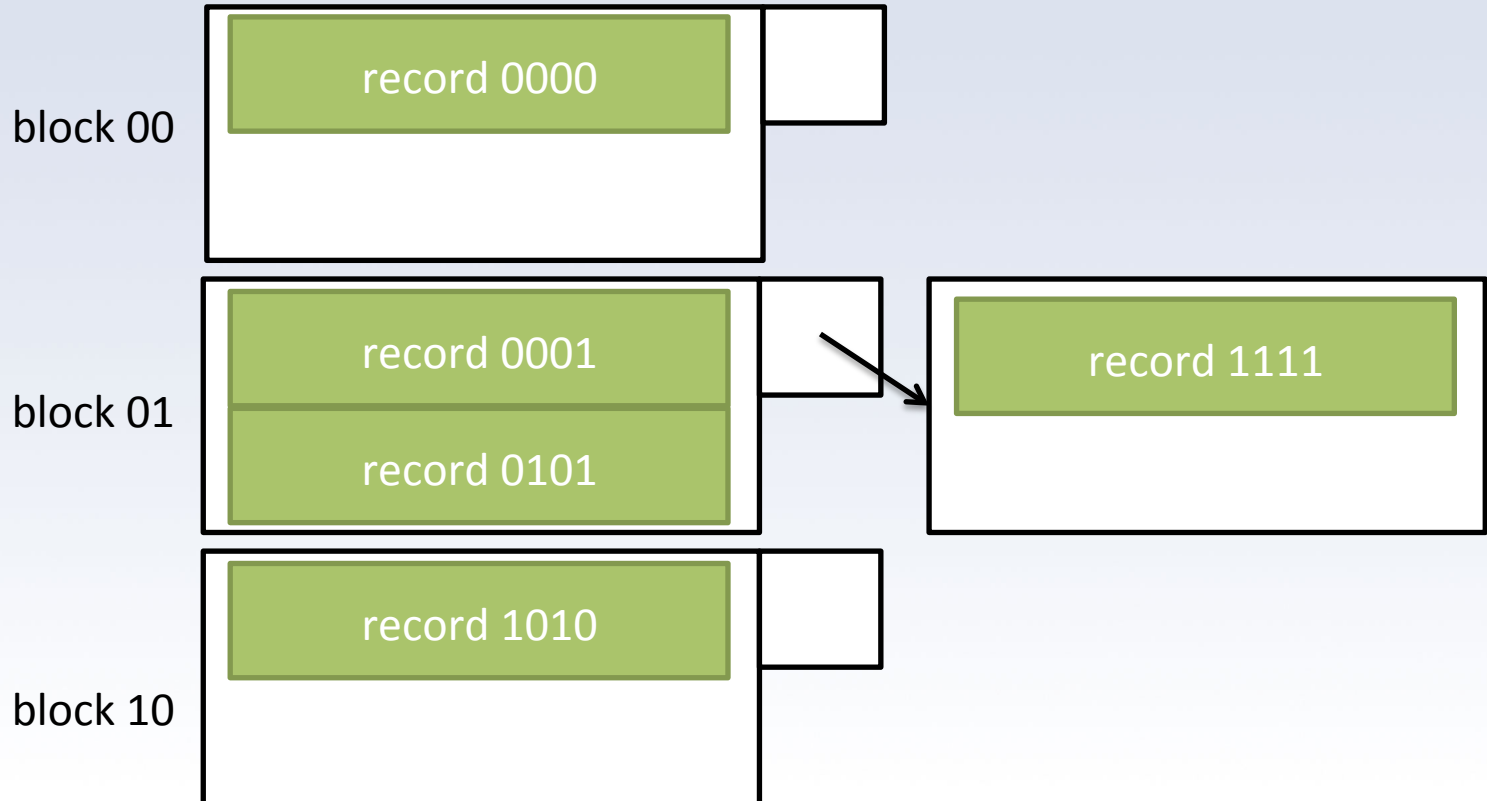


Example: Insert 0111

i=2
n=3
r=5

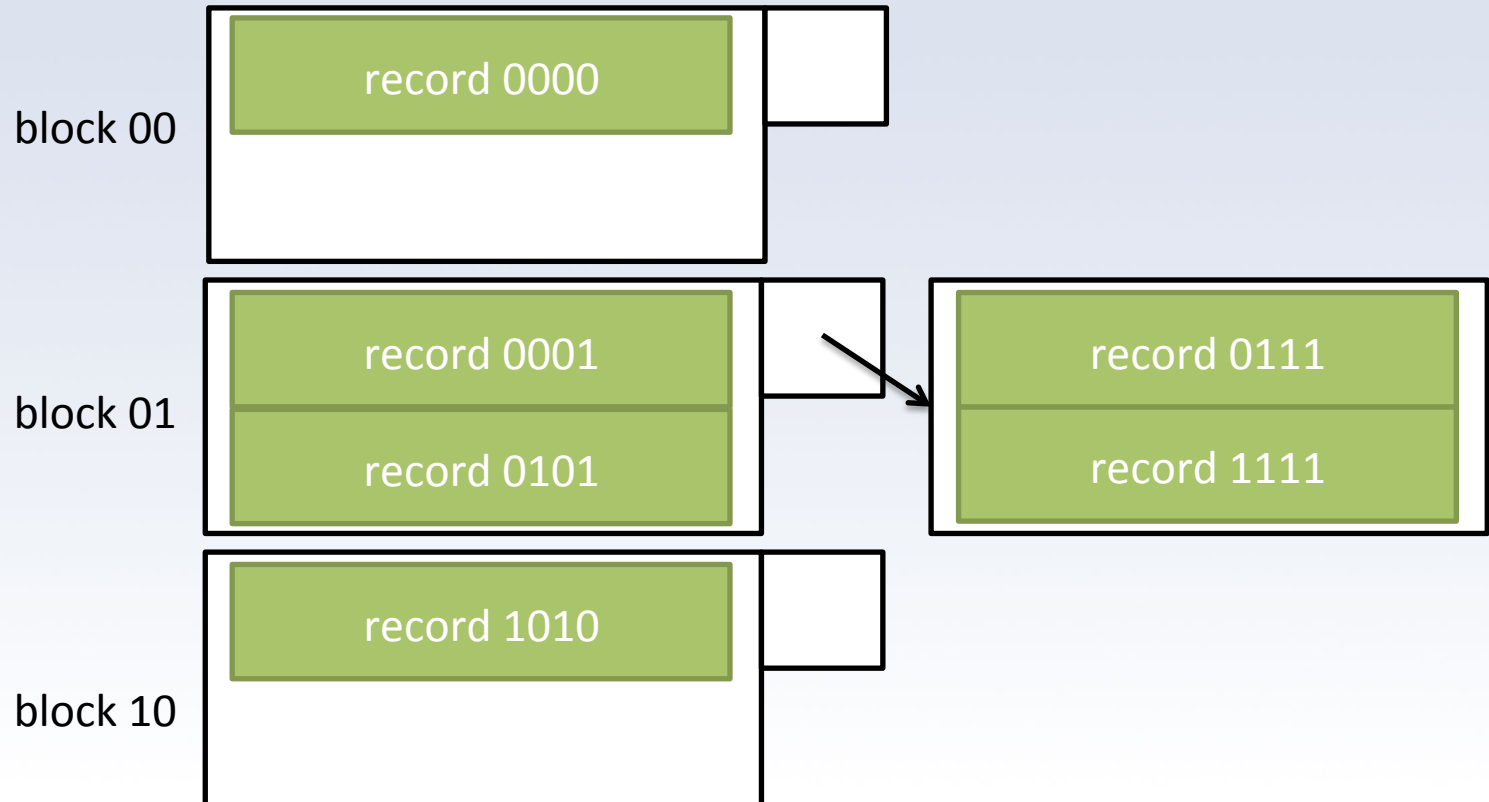
3 \geq n

Use block
block 01



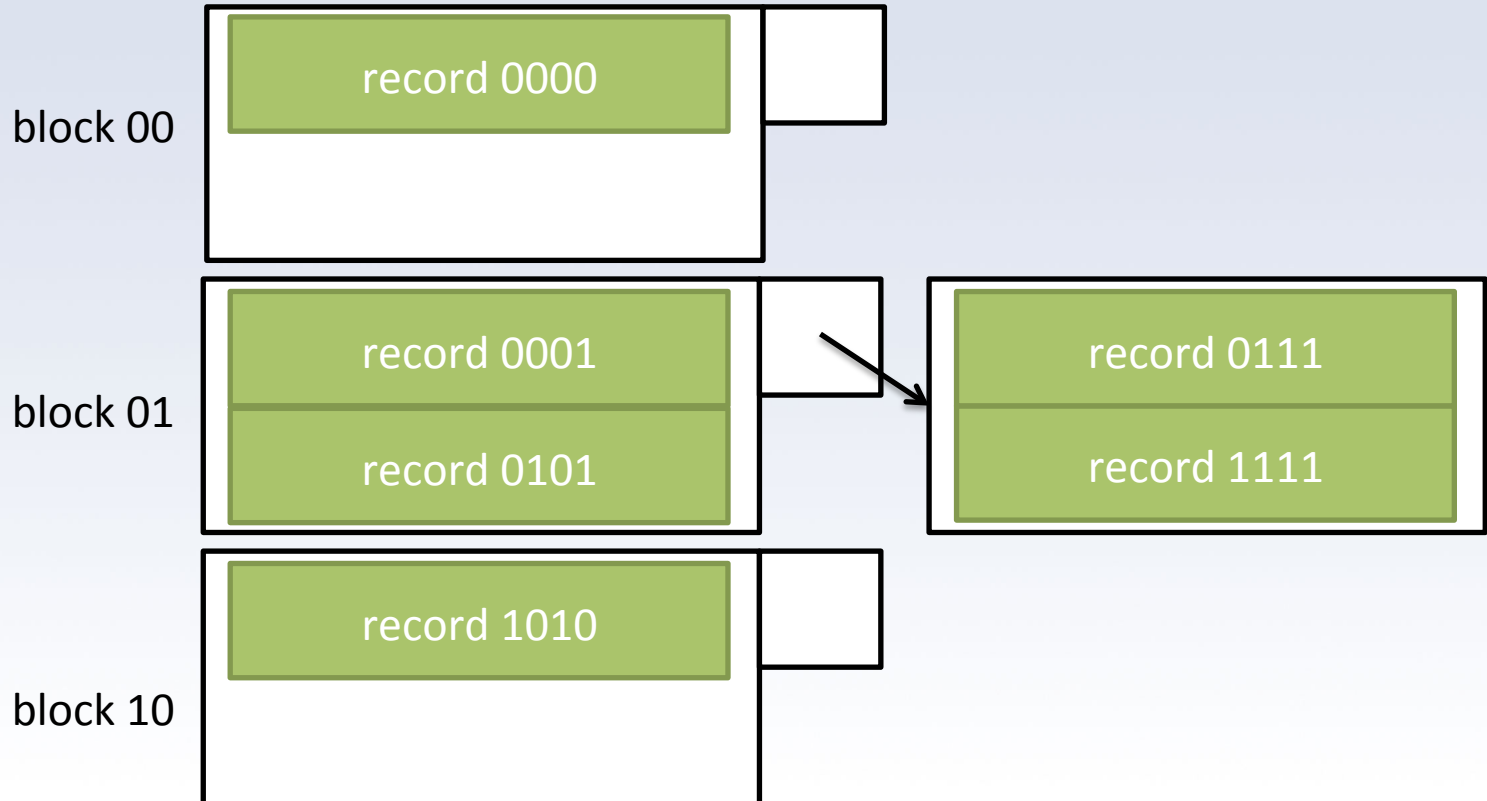
Example: Insert 0111

i=2
n=3
r=6



Example: Insert 0111

i=2
n=3
r=6

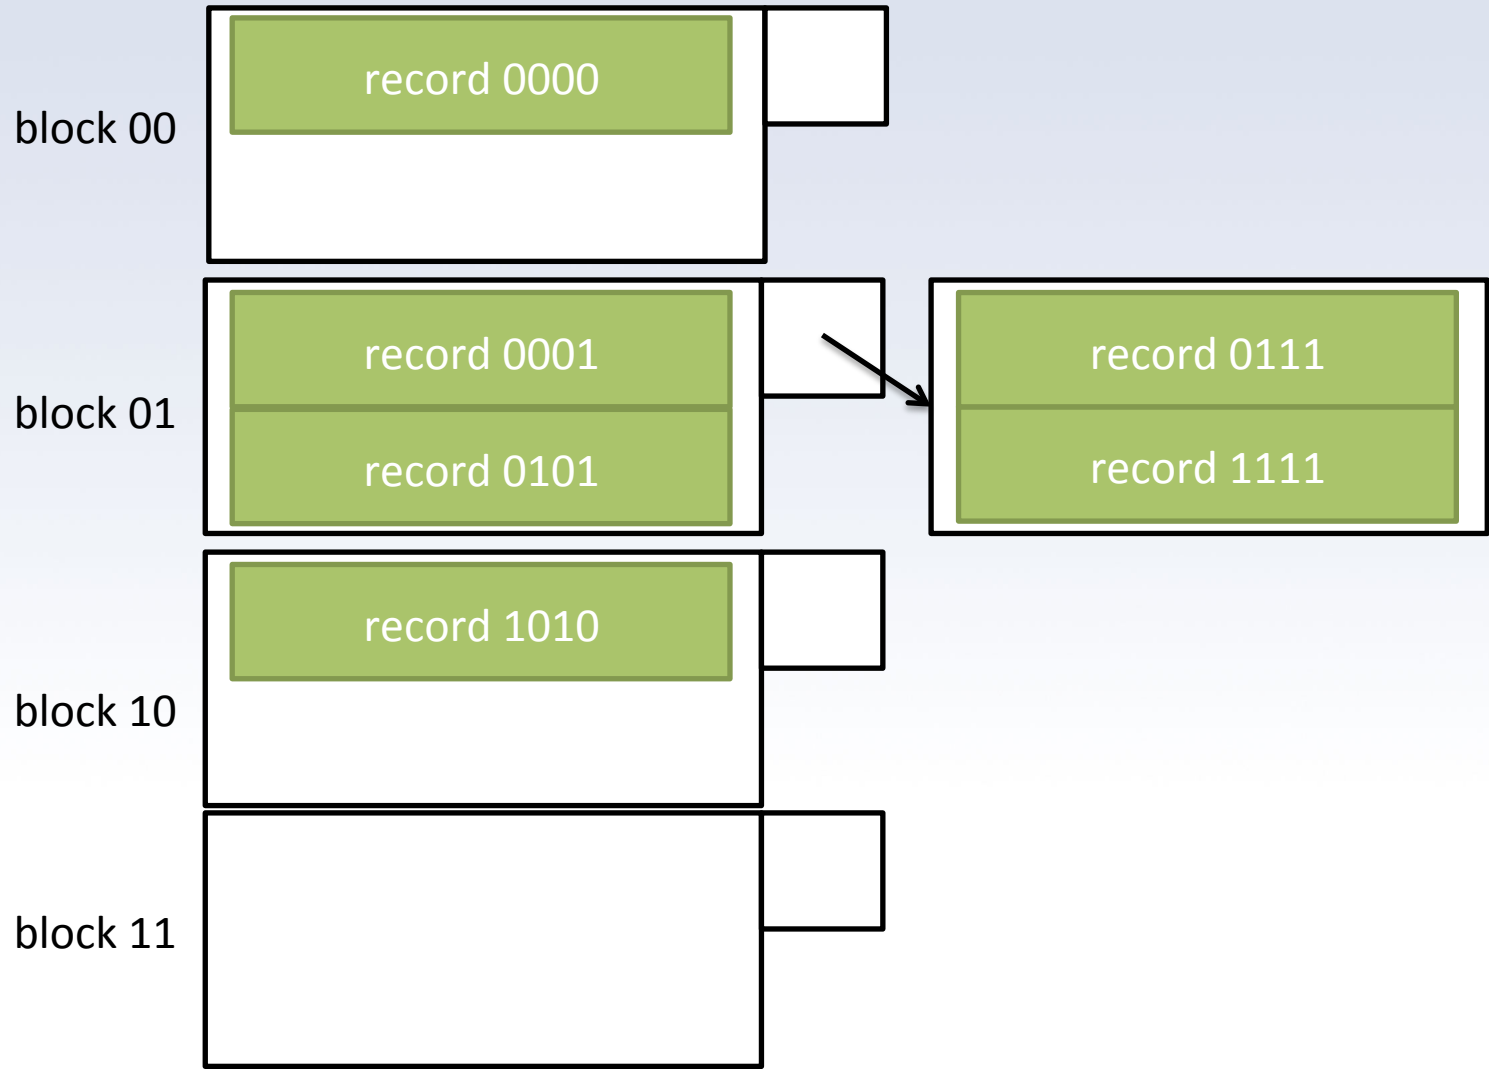


$$r/n = 2 > 1.7$$



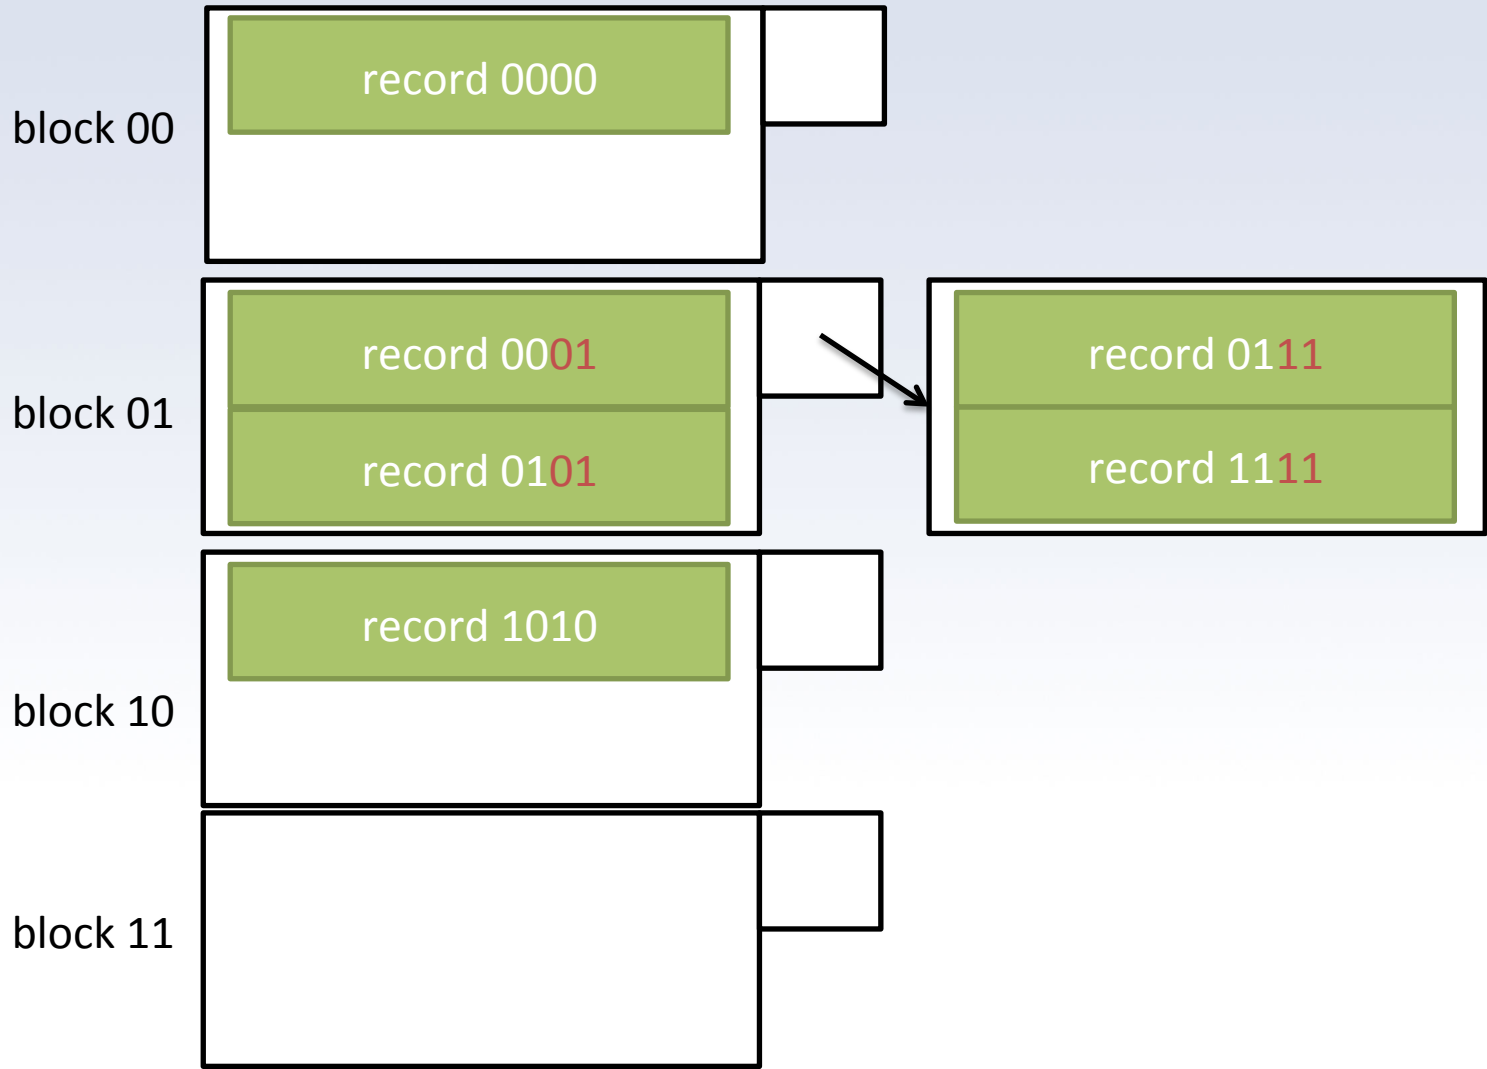
Example: Insert 0111

i=2
n=3
r=6



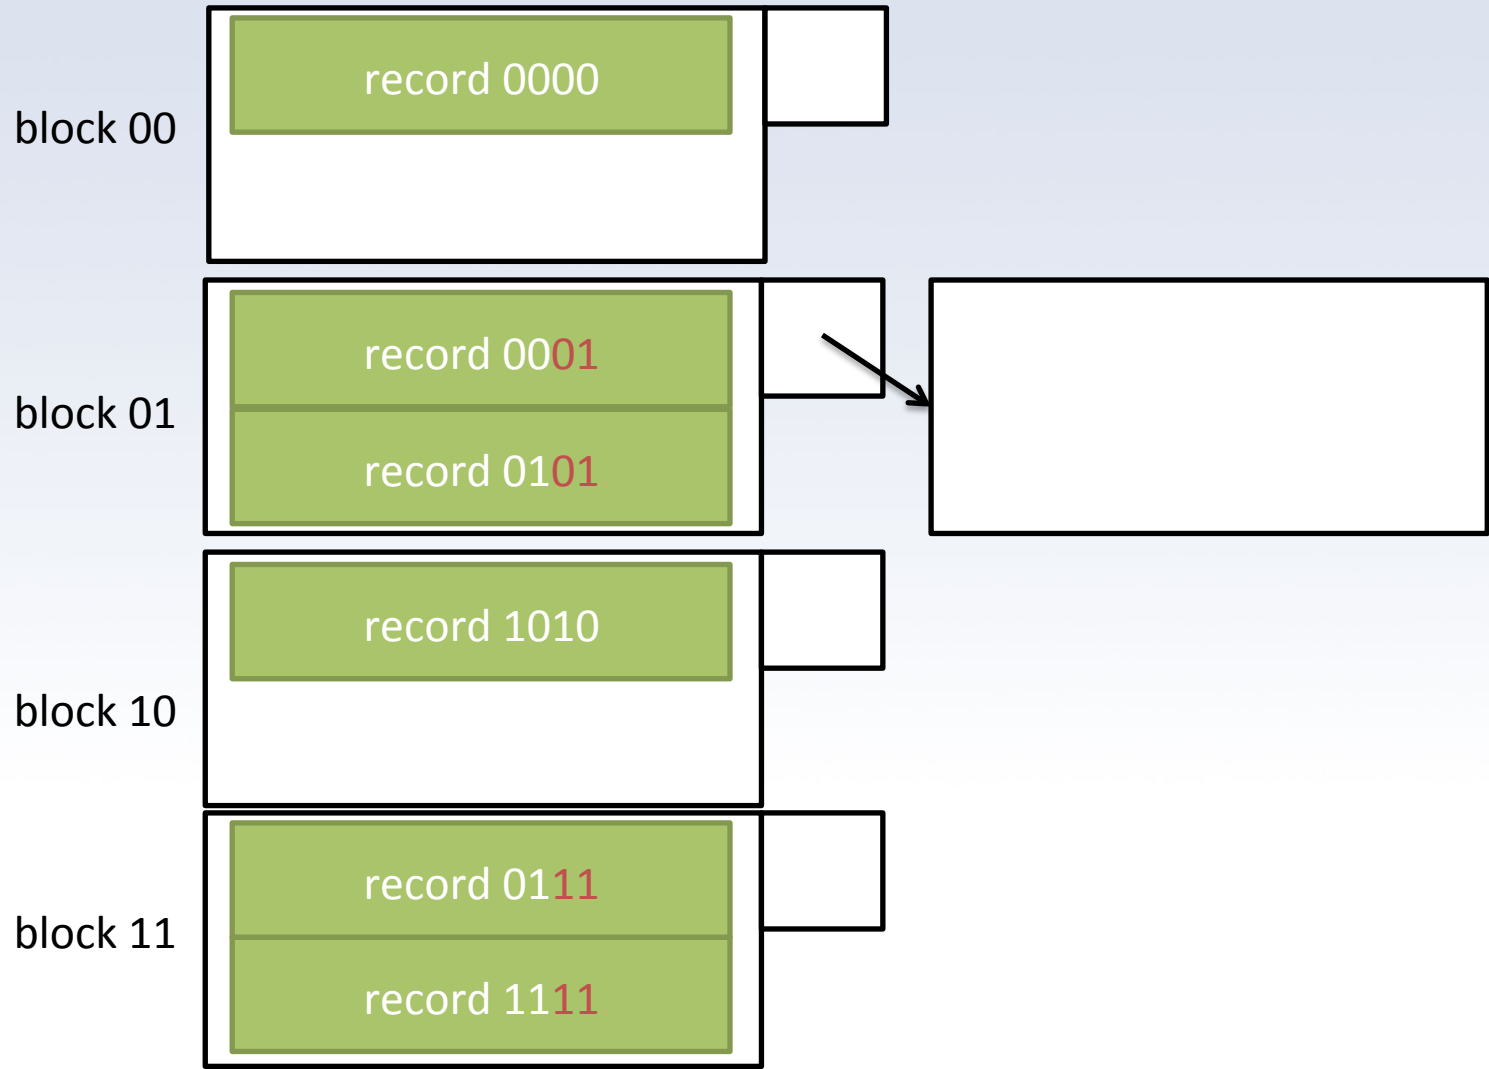
Example: Insert 0111

i=2
n=3
r=6



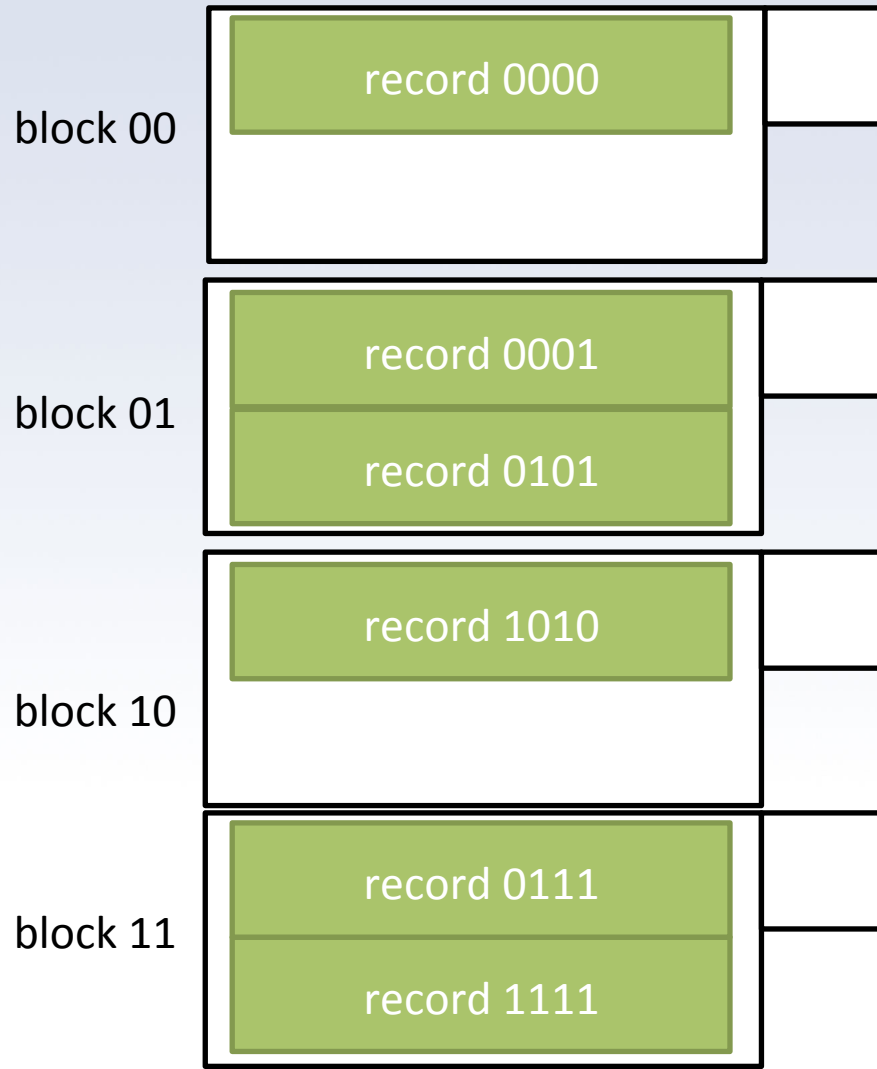
Example: Insert 0111

i=2
n=3
r=6



Example: Insert 0111

i=2
n=3
r=6



Hashing Strings

- Strings vary dramatically in length
- How can we hash them?
 - Use CRC
 - Use a rolling hash
 - Use cryptographic hash (e.g. MD5/SHA1)



Next time...

- We'll talk about
 - Multidimensional indexing (how Foursquare and other location based searching works)
 - Inverted indexing (how web search works)
 - Suffix arrays (how to index very large strings for search)

