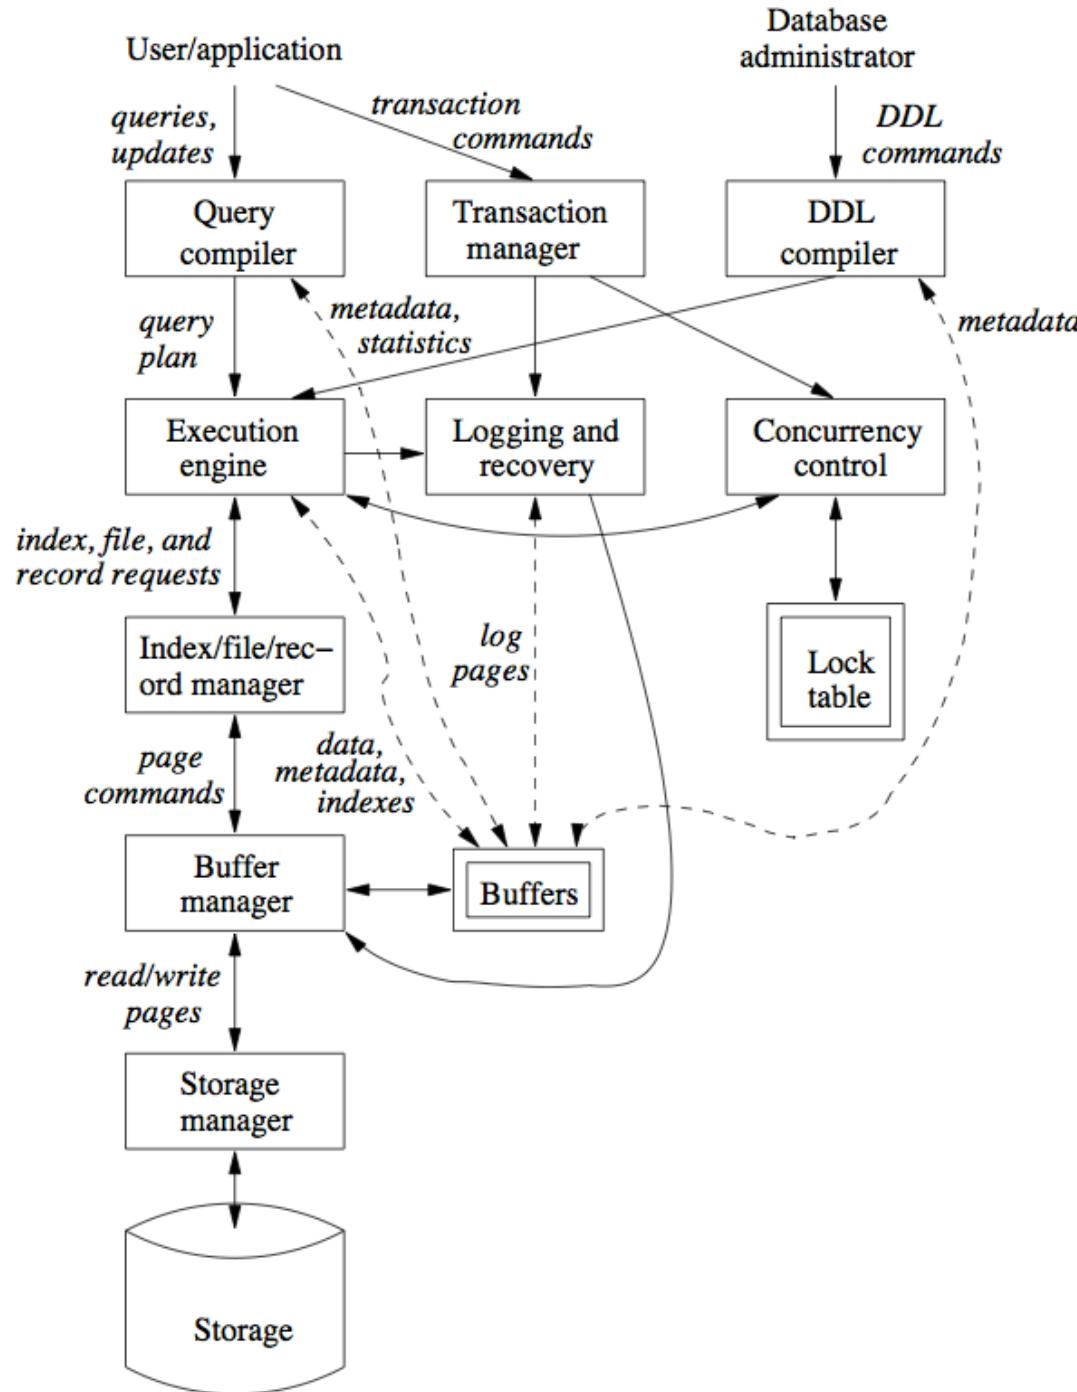


UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Query Execution



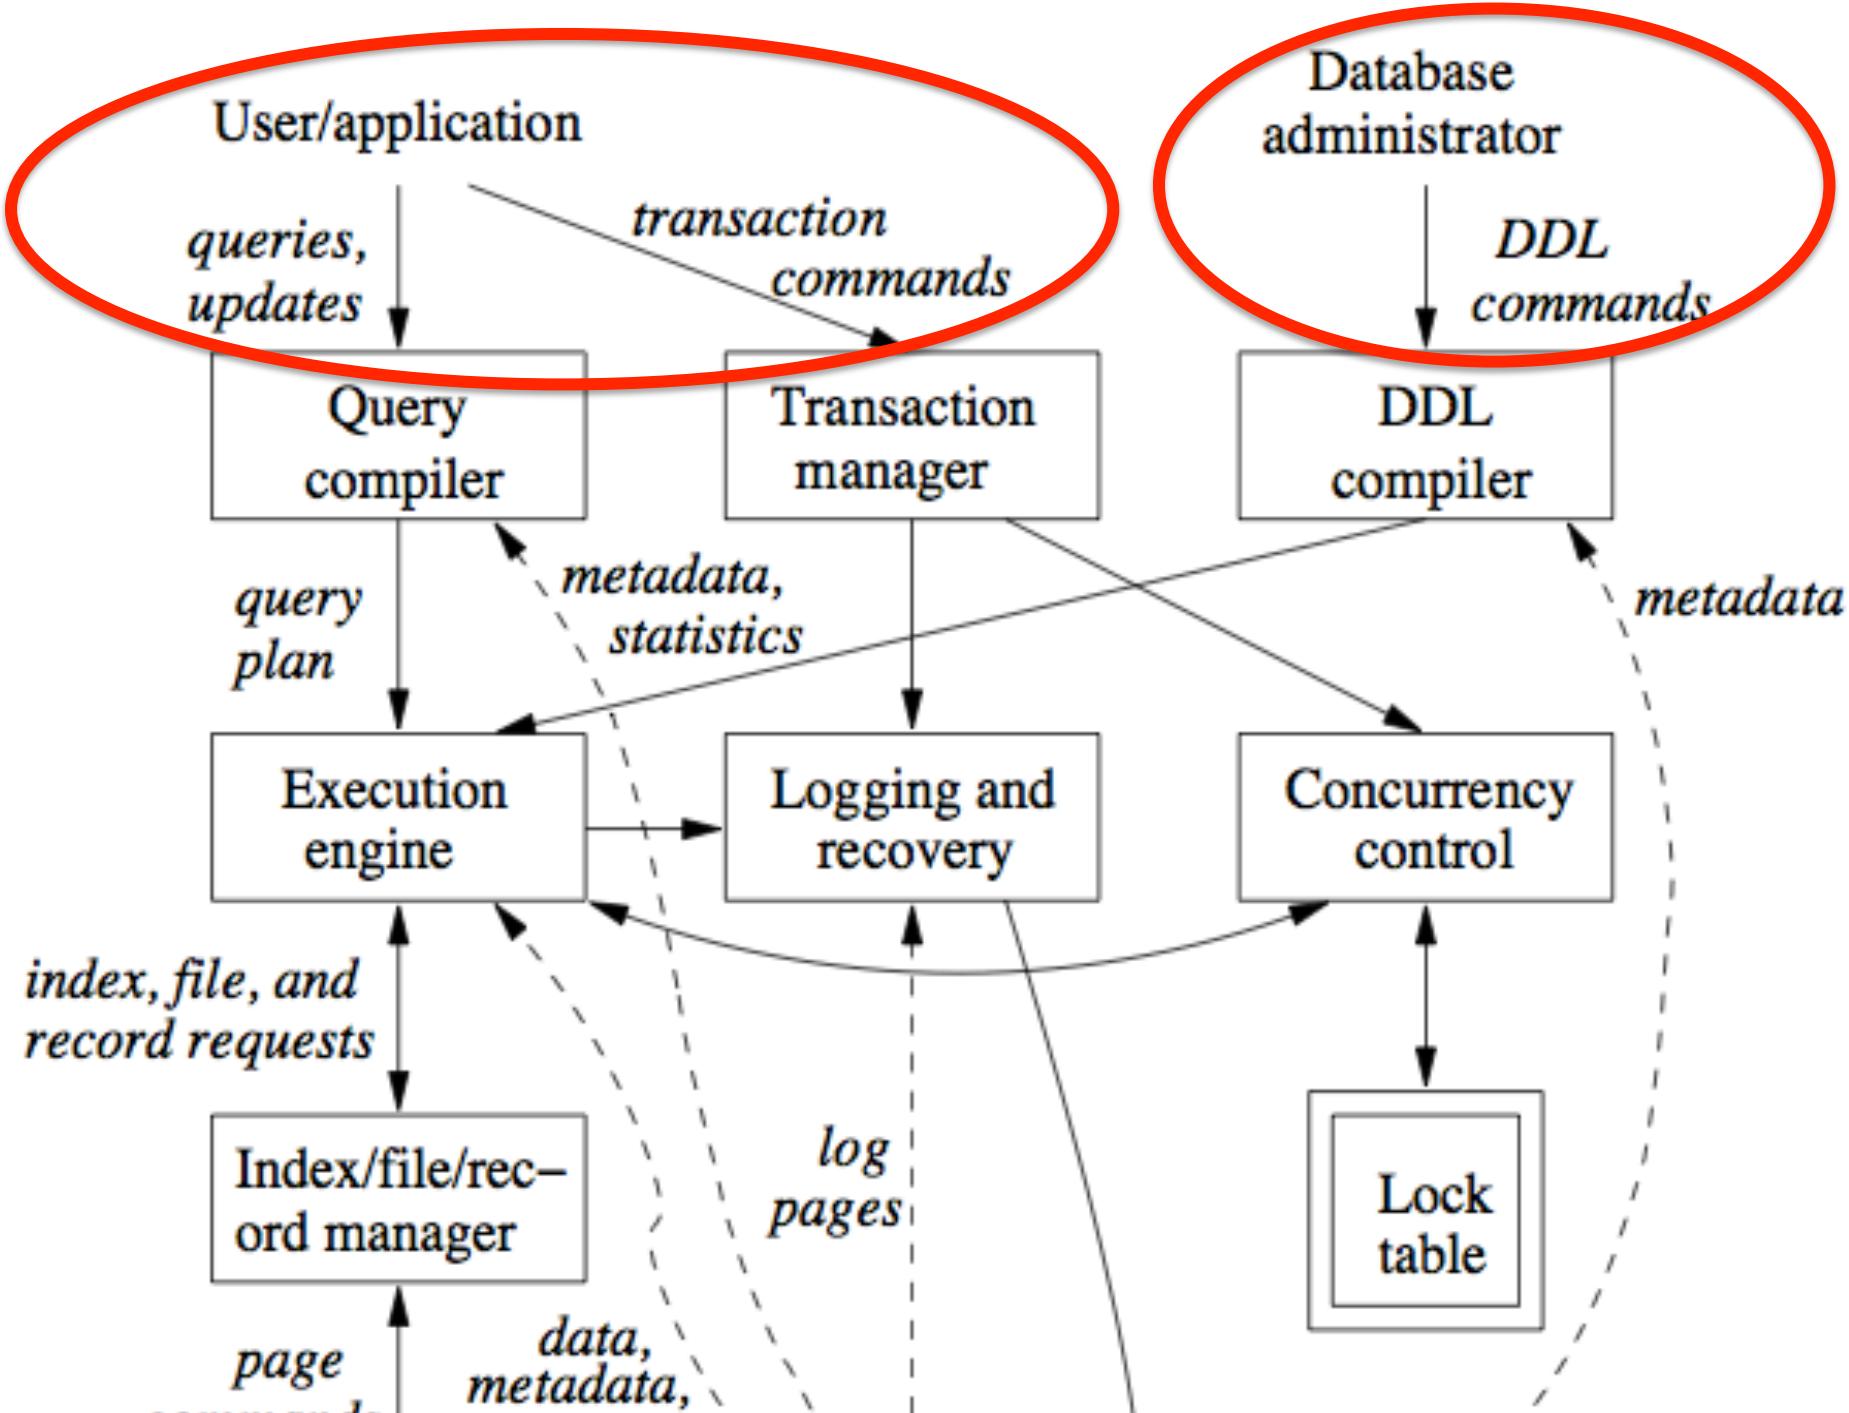
illinois.edu

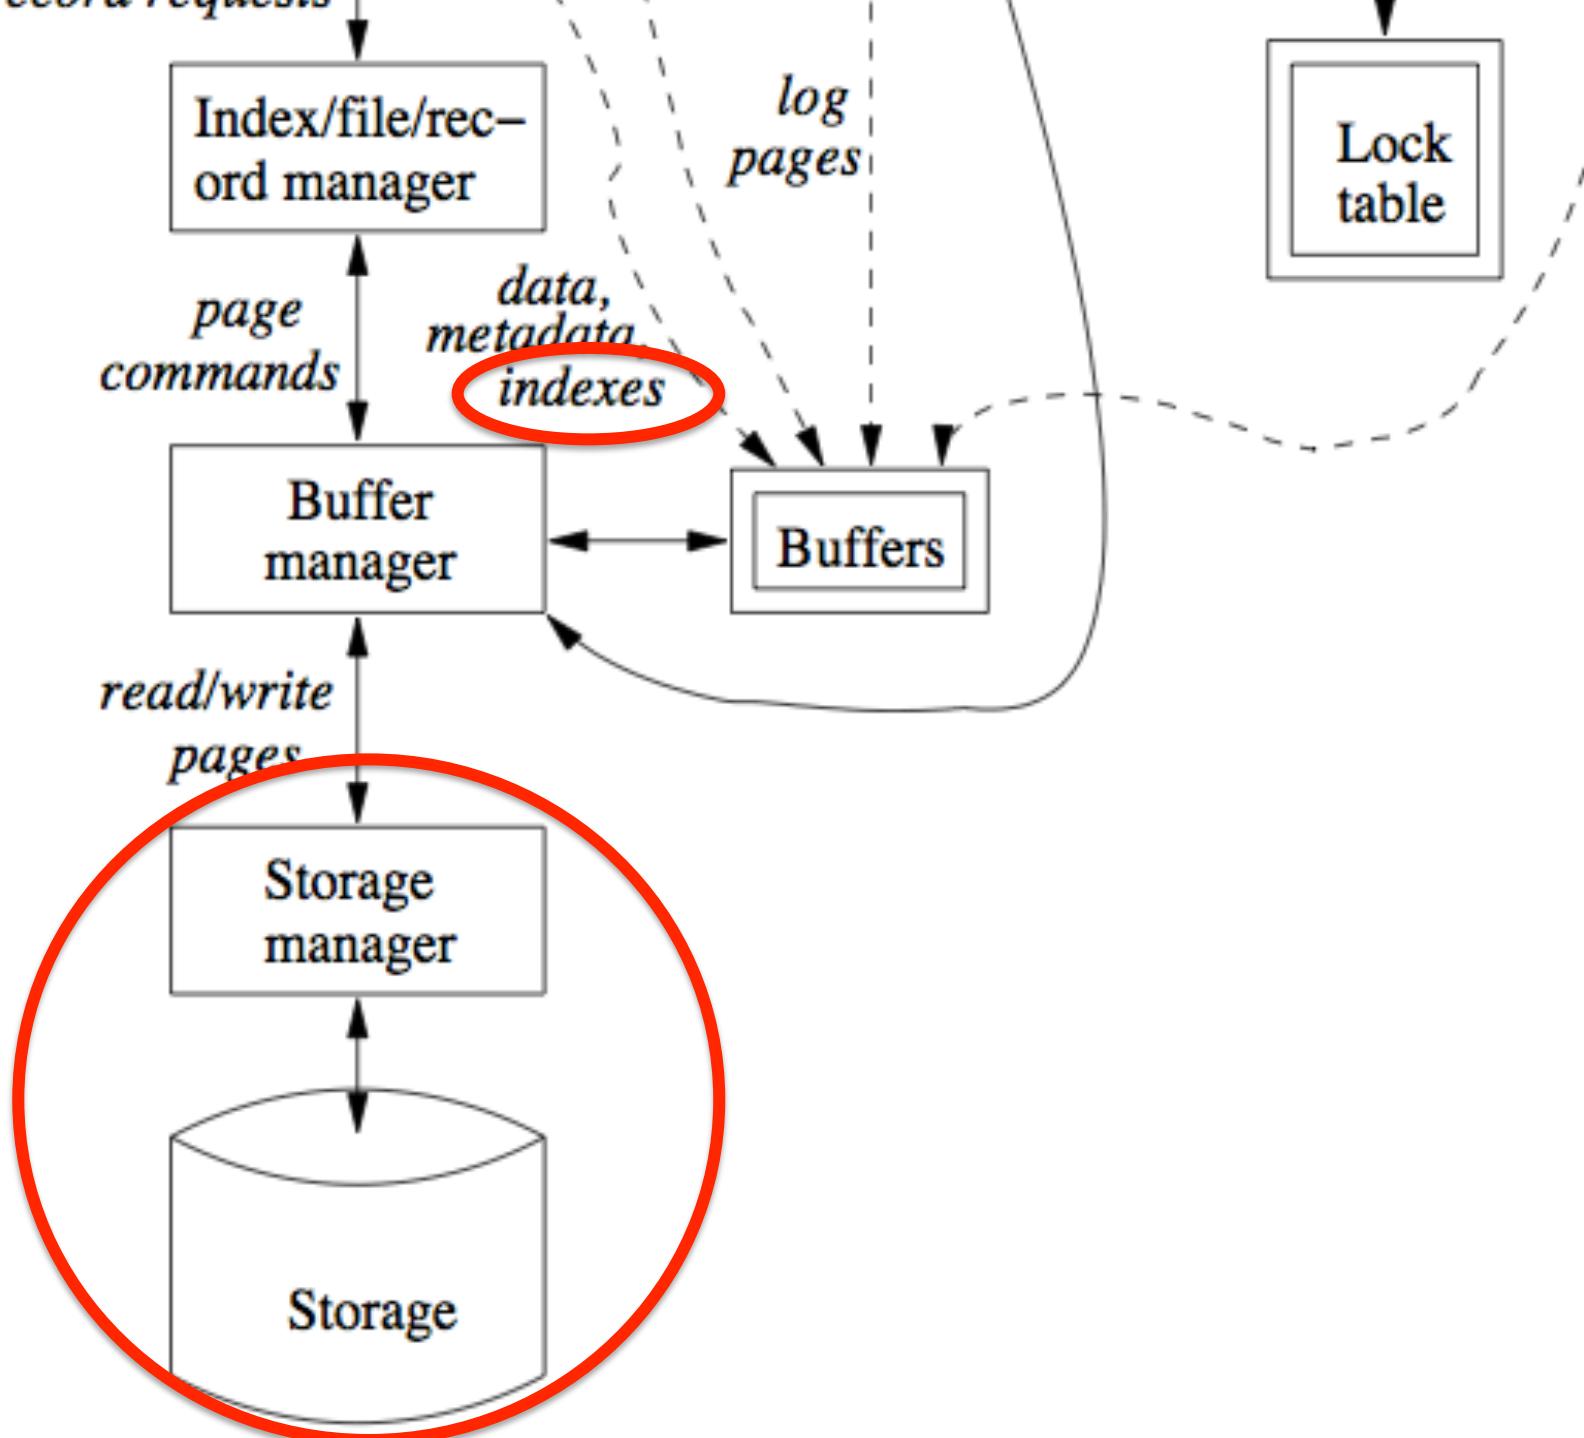


So far, we've seen

- How users/admins interact with a DBMS
- How records and relations are stored
- How to create structures to access records more quickly



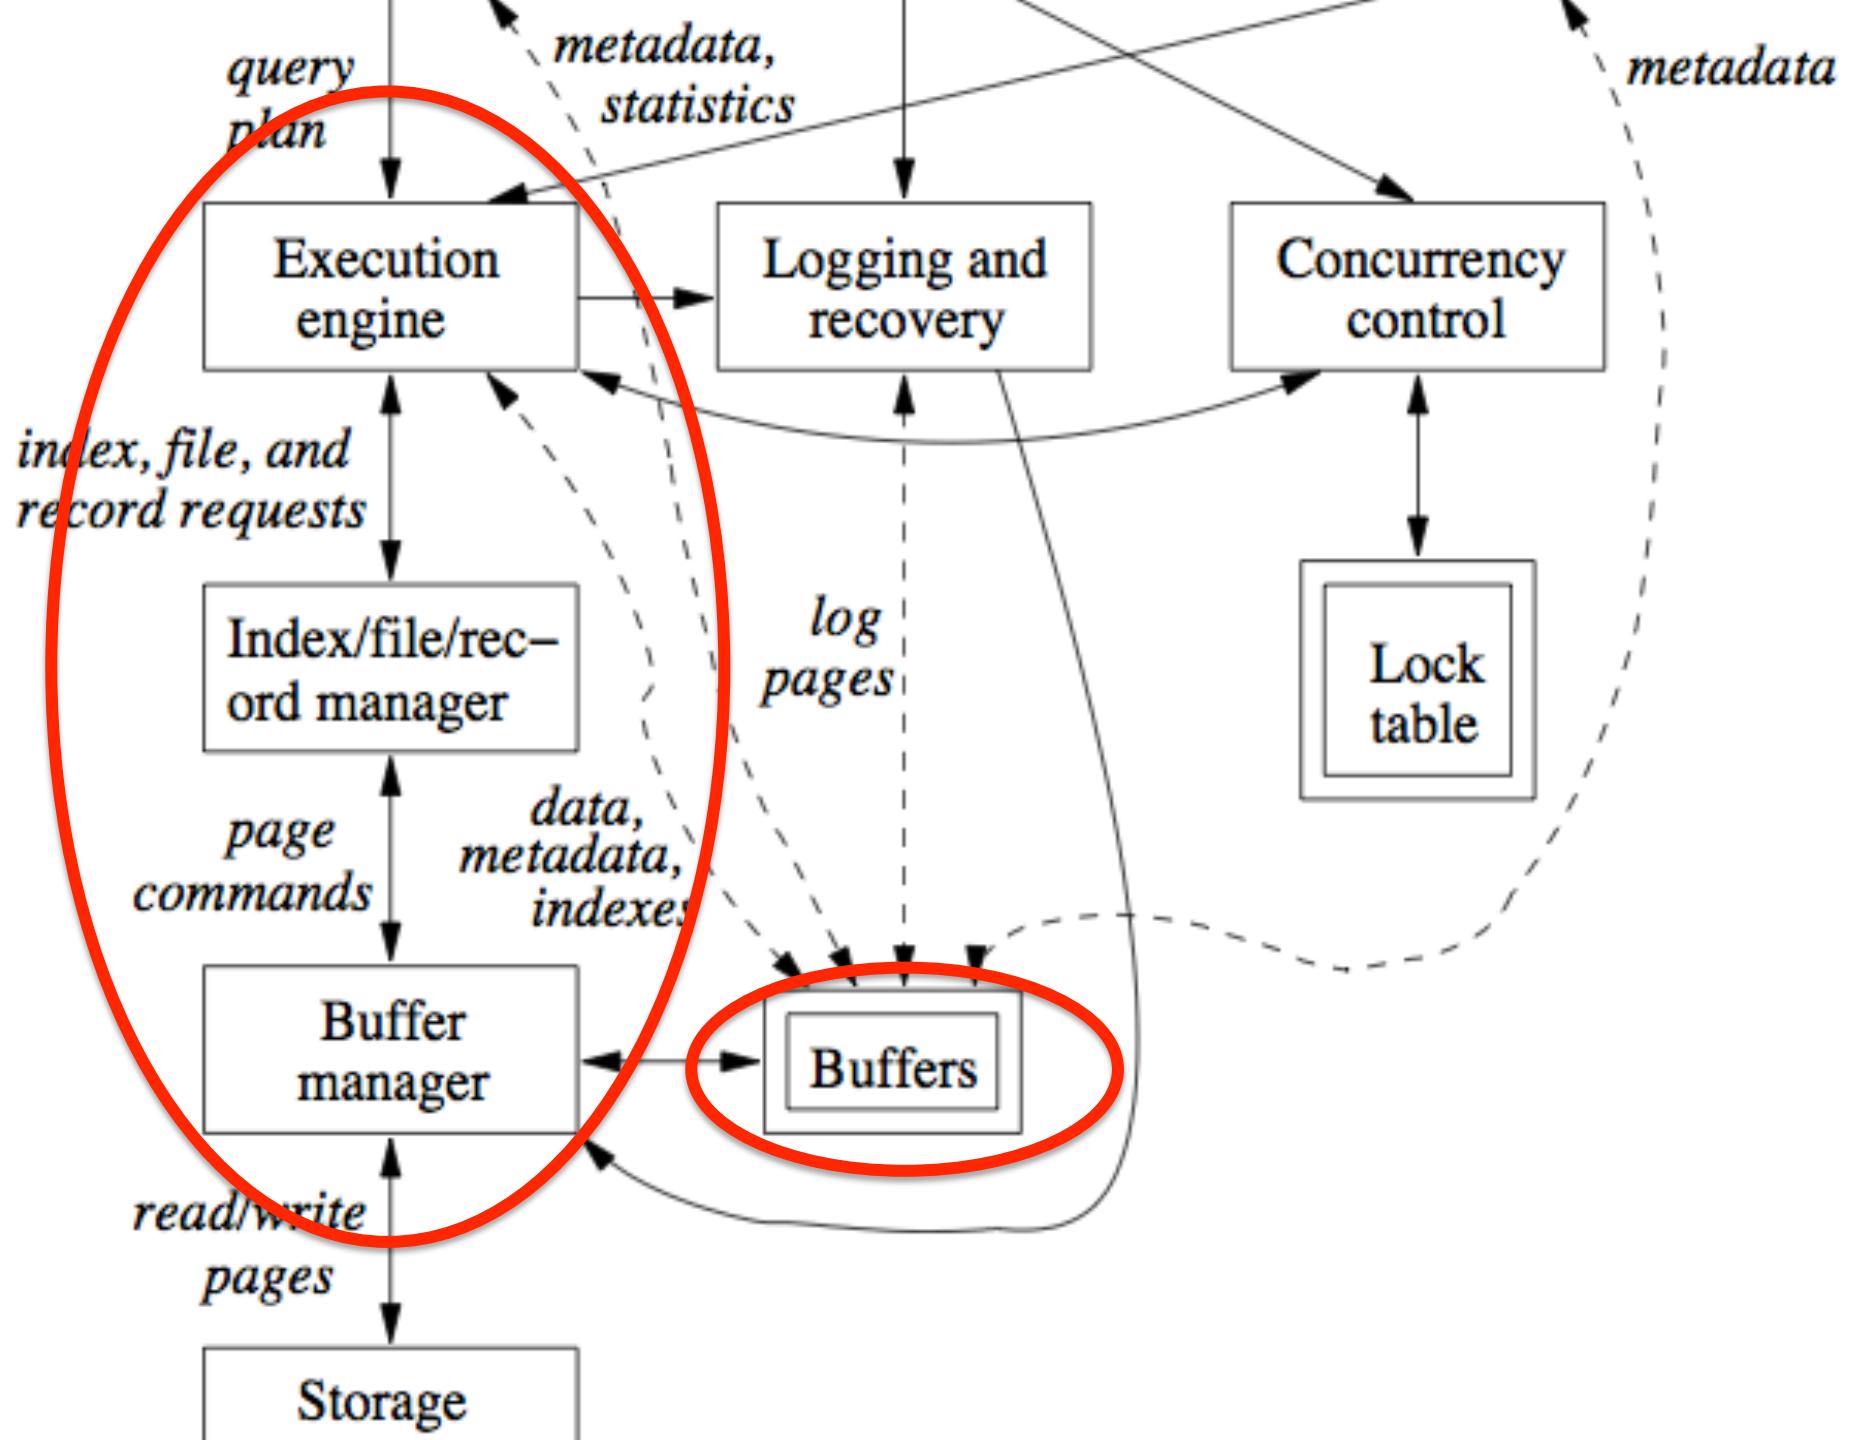




Next, we will see

- Many ways to execute query operations on our data
 - Today: without using indexes
 - Friday: taking advantage of indexes
- Next week:
 - How user queries become query plans
 - How we can efficiently execute query plan



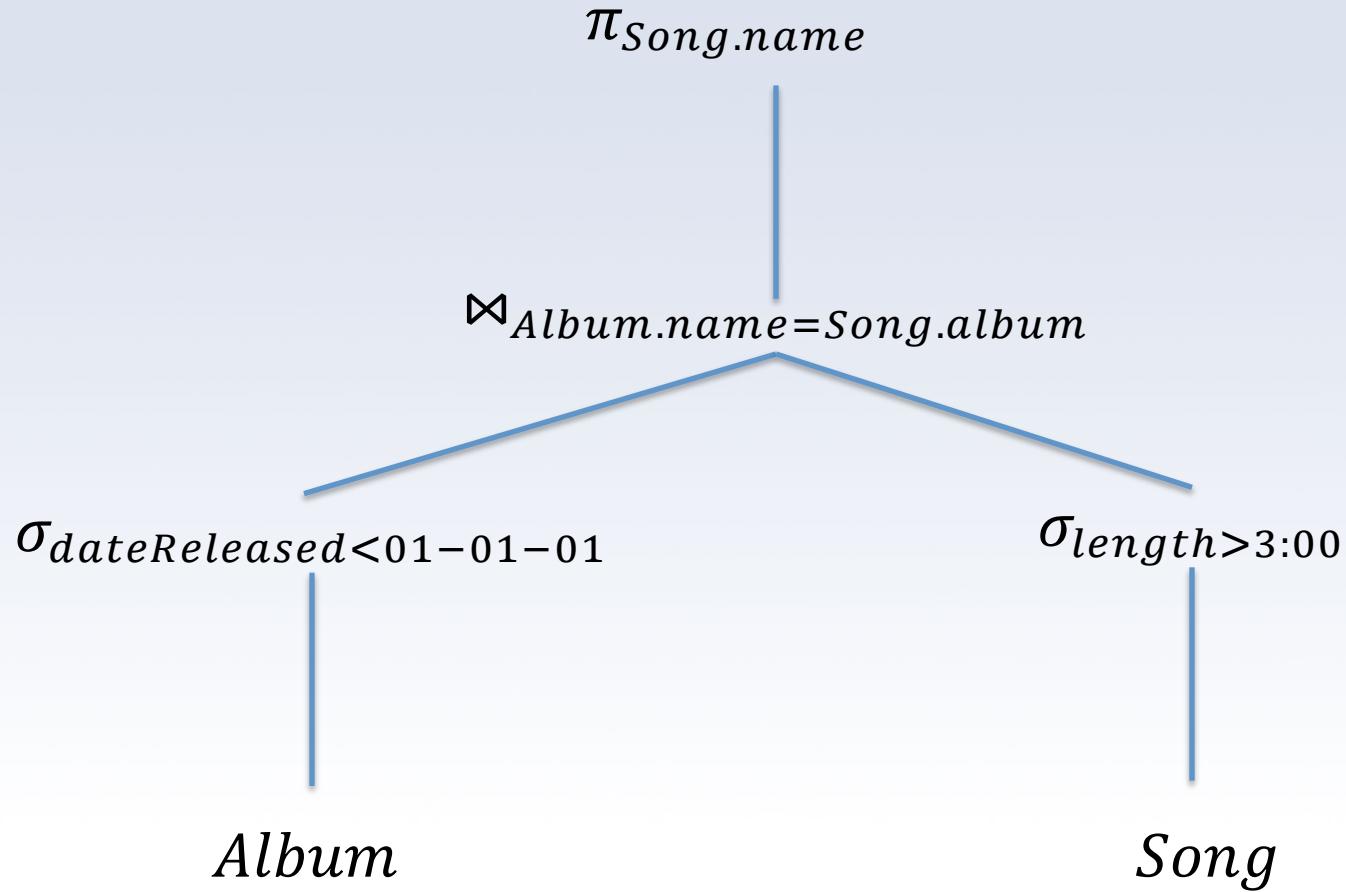


Query Compilation

1. Parsing - the query is transformed into a parse tree
2. Parse tree converted into a logical query plan
3. Logical query plan is converted into a physical query plan



Query Plan



Query Compilation

1. Parsing - the query is transformed into a parse tree
2. Parse tree converted into a logical query plan
3. Logical query plan is converted into a **physical query plan**



Physical Query Plan

- Given a tree of relational algebra operators, how can we physically execute them?
- How do we access the data efficiently to compute the desired result?



Review

- Let's take a quick step back

Disk access

- Fundamental unit of transfer is a *block*
- Blocks are typically 4-64kb
- When in main memory, blocks are sometimes called *pages*



Fixed Length Records

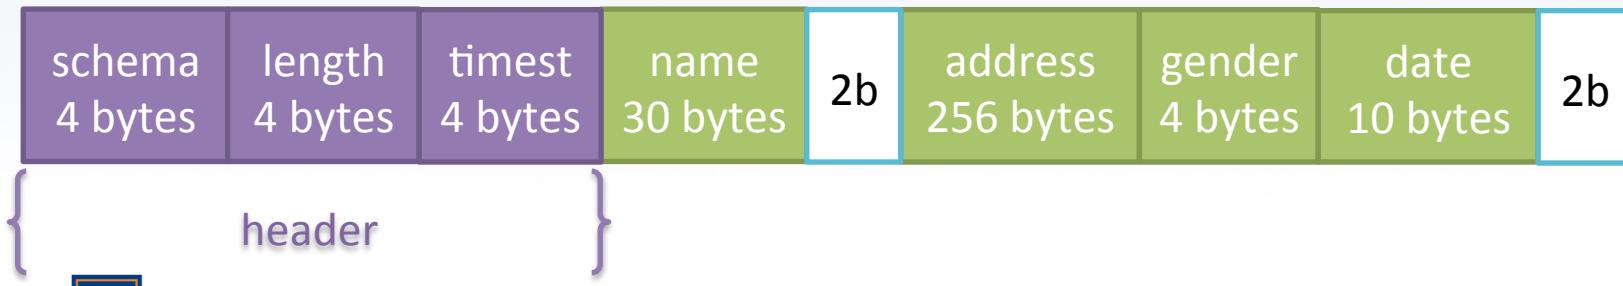
- Align fields along ***word boundaries***
 - minimum unit addressable in memory
 - usually 4 or 8 bytes
- Need to include a header
 - pointer to the schema for the record
 - length of the record
 - last modified timestamp/access metadata
 - pointers to the fields



Fixed Length Records

```
CREATE TABLE Person(  
    name CHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    gender CHAR(1),  
    birthdate DATE  
);
```

Actually takes 316 bytes



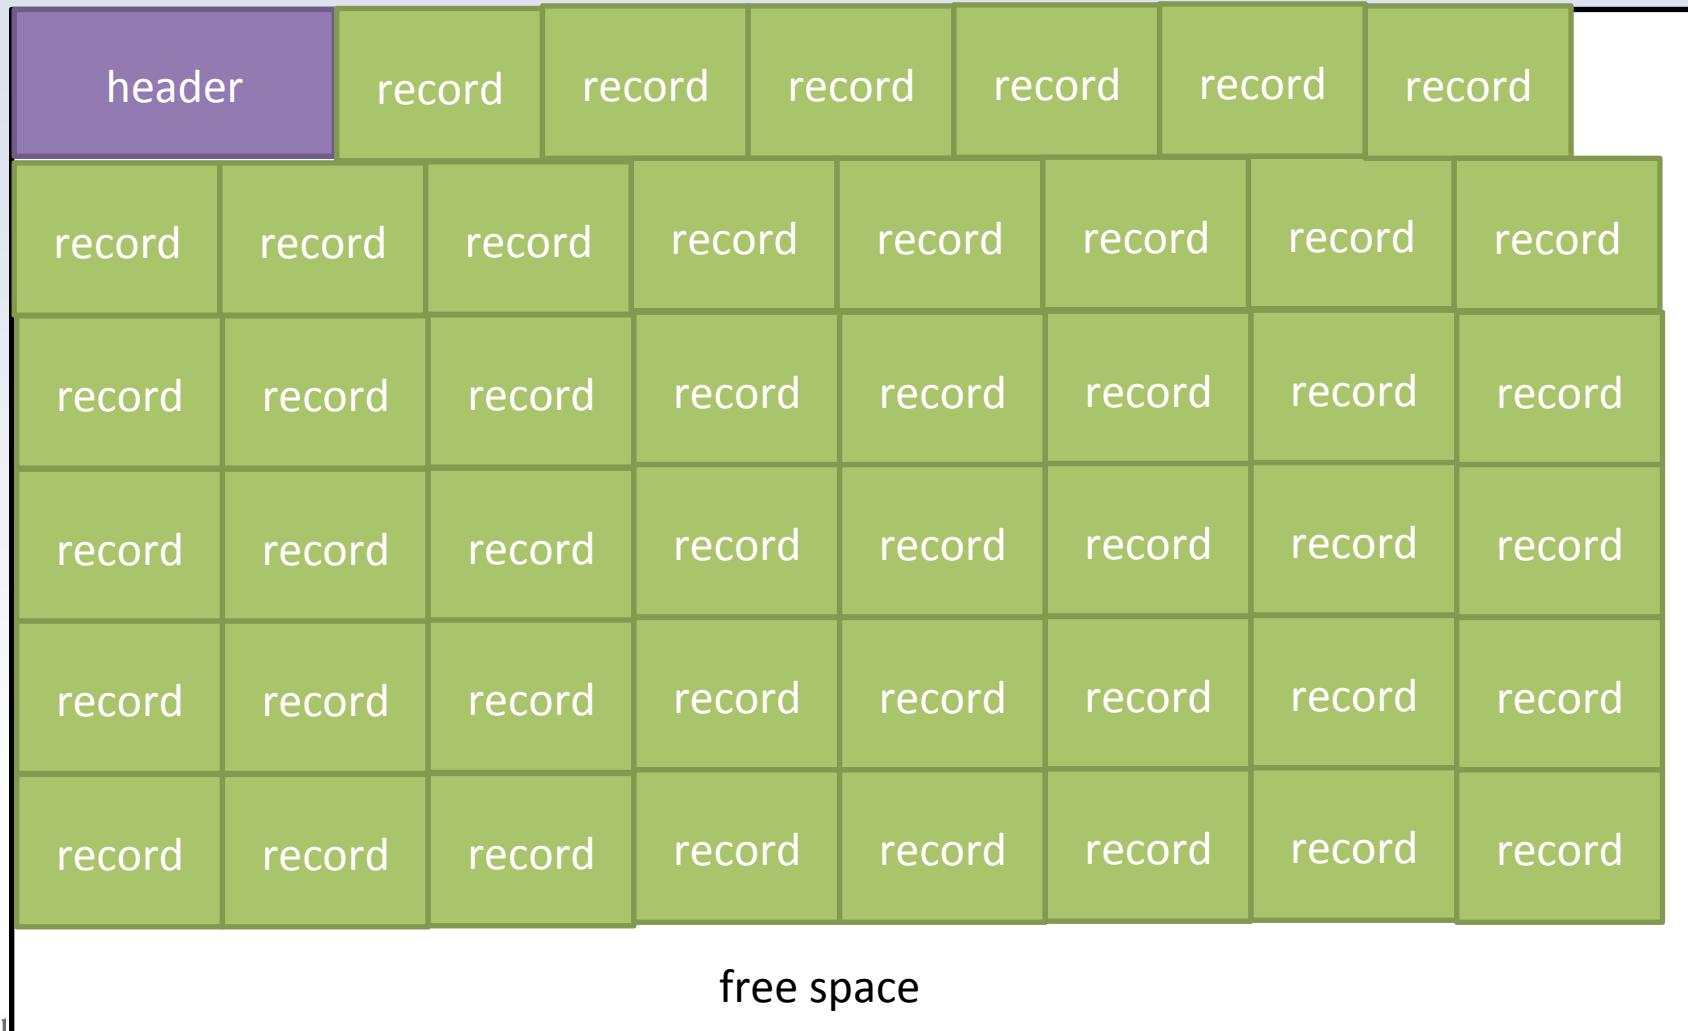
Putting records into blocks

- Doesn't necessarily fill the block
- Has a header:
 - Information about the relation
 - “Directory” of records in the block
 - Links to the “previous” and “next” blocks
 - Last modified timestamp/access metadata



Putting records into blocks

block



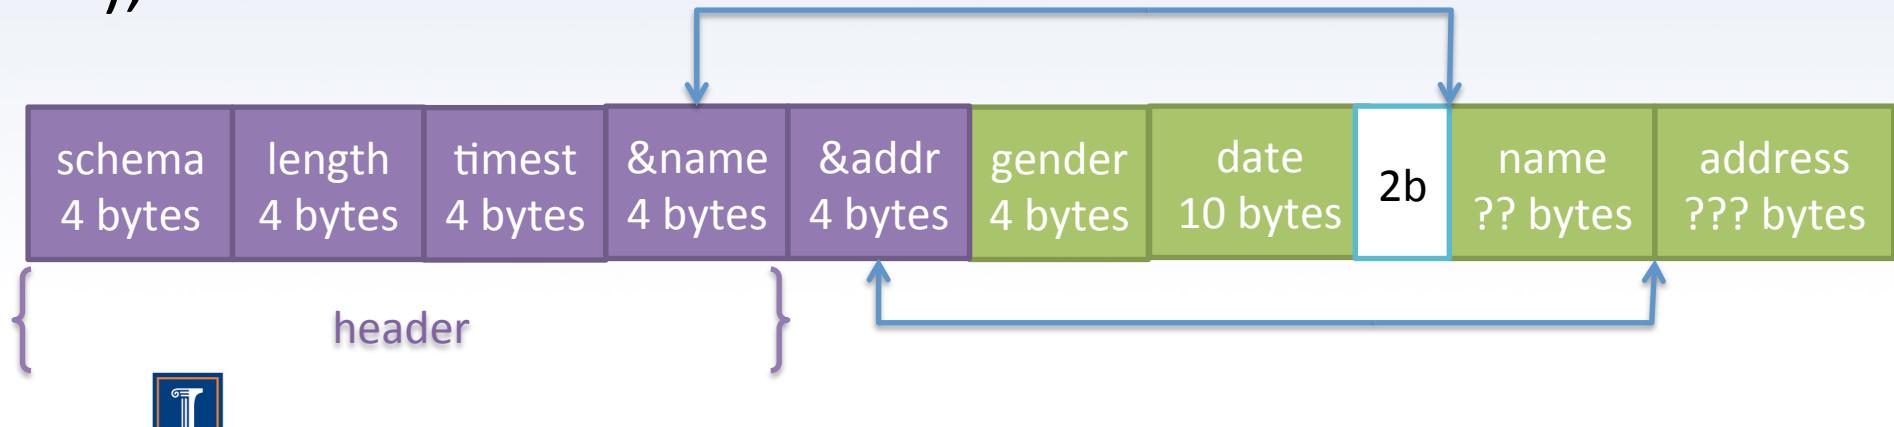
Variable length records

- Sometimes, fields have varying lengths
 - e.g. name, address
- Solution:
 - store fixed length fields first
 - store pointers to beginnings of variable length fields in header



Variable length records

```
CREATE TABLE Person(  
    name VARCHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    gender CHAR(1),  
    birthdate DATE  
);
```



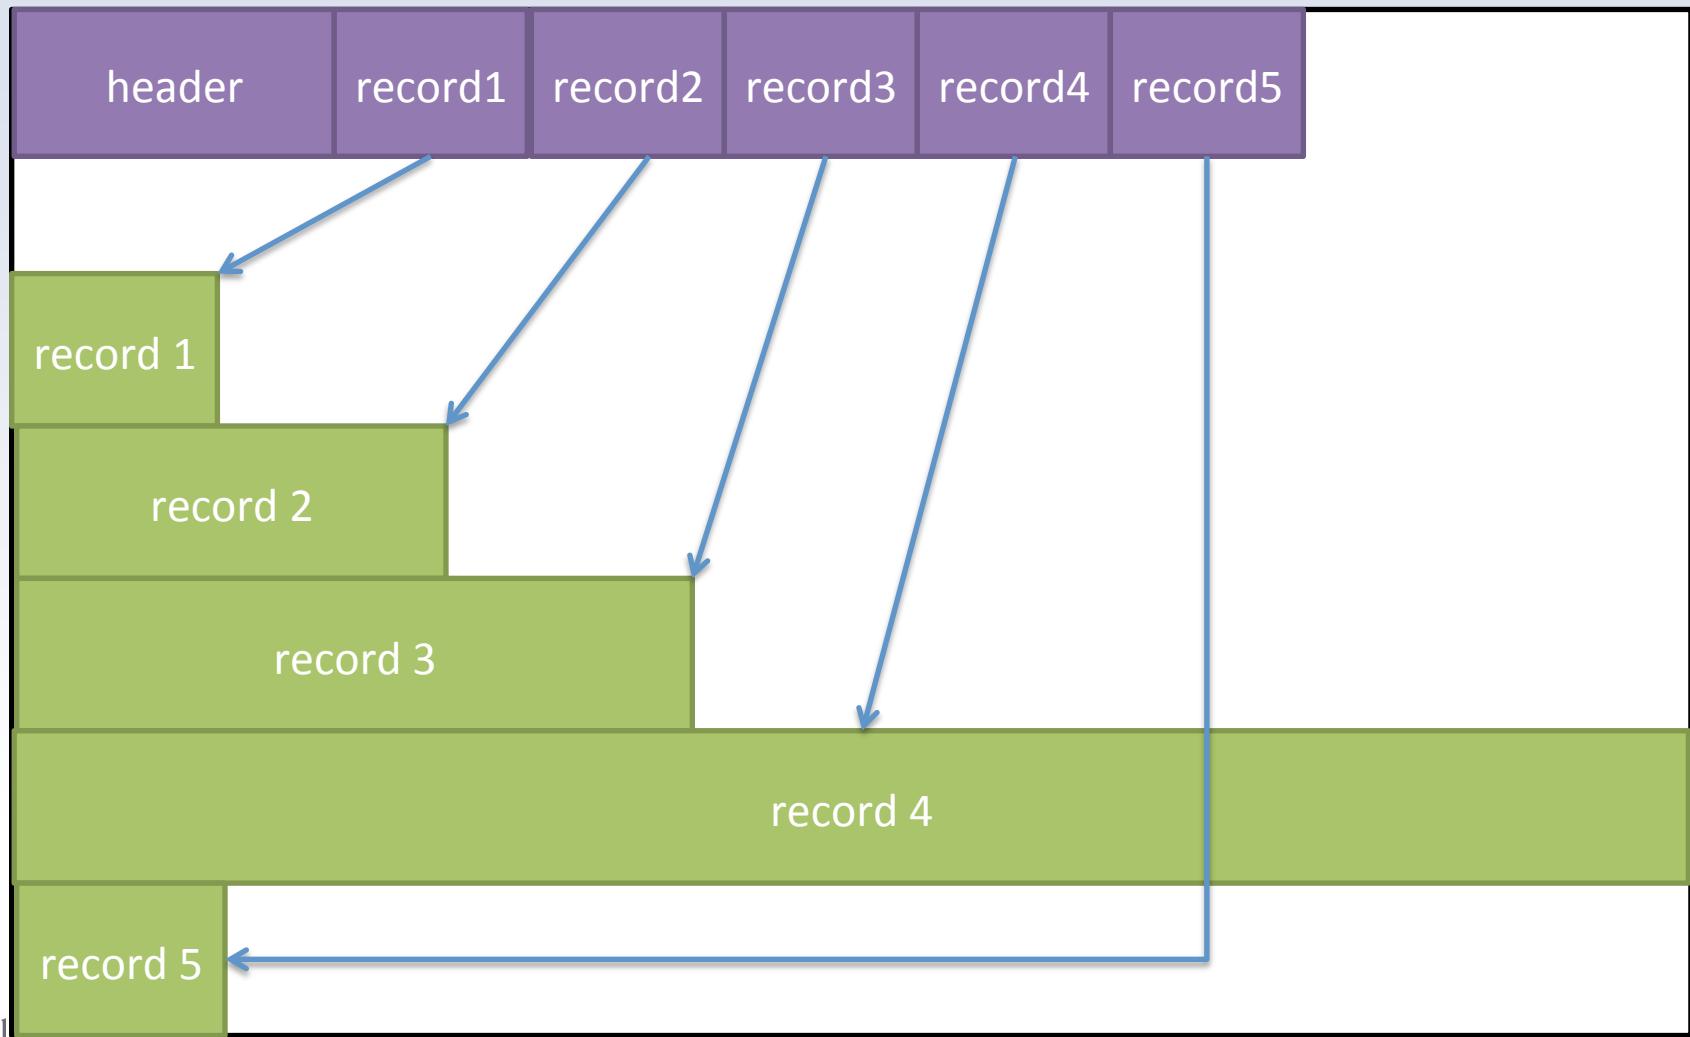
Variable length into blocks

- Added offset and length index to header for variable length records



Variable length into blocks

block



Big picture

- To access a relation, we pull the corresponding blocks into memory
- With a block in memory, we can access all of the records in that block



Query Execution

- Implement relational operators on our stored records
- Will implement in many different ways
 - Depending on the size of memory and our data, some methods better than others
 - Gives flexibility for query optimization

Goal

- We care about these operators:
 - σ selection and π projection
 - δ duplicate-elimination and γ grouping
 - \cup union, \cap intersection, and $-$ difference
 - \bowtie join and \times cross product
- Always assume relations are ***bags***



Three types of operators

- Tuple-at-a-time unary - need one tuple to perform computation (e.g. σ and π)
- Full-relational unary - need entire relation to perform computation (e.g. δ and γ)
- Full-relational binary - need two entire relations to perform computation (e.g. \cup , \cap , $-$, \bowtie , \times)

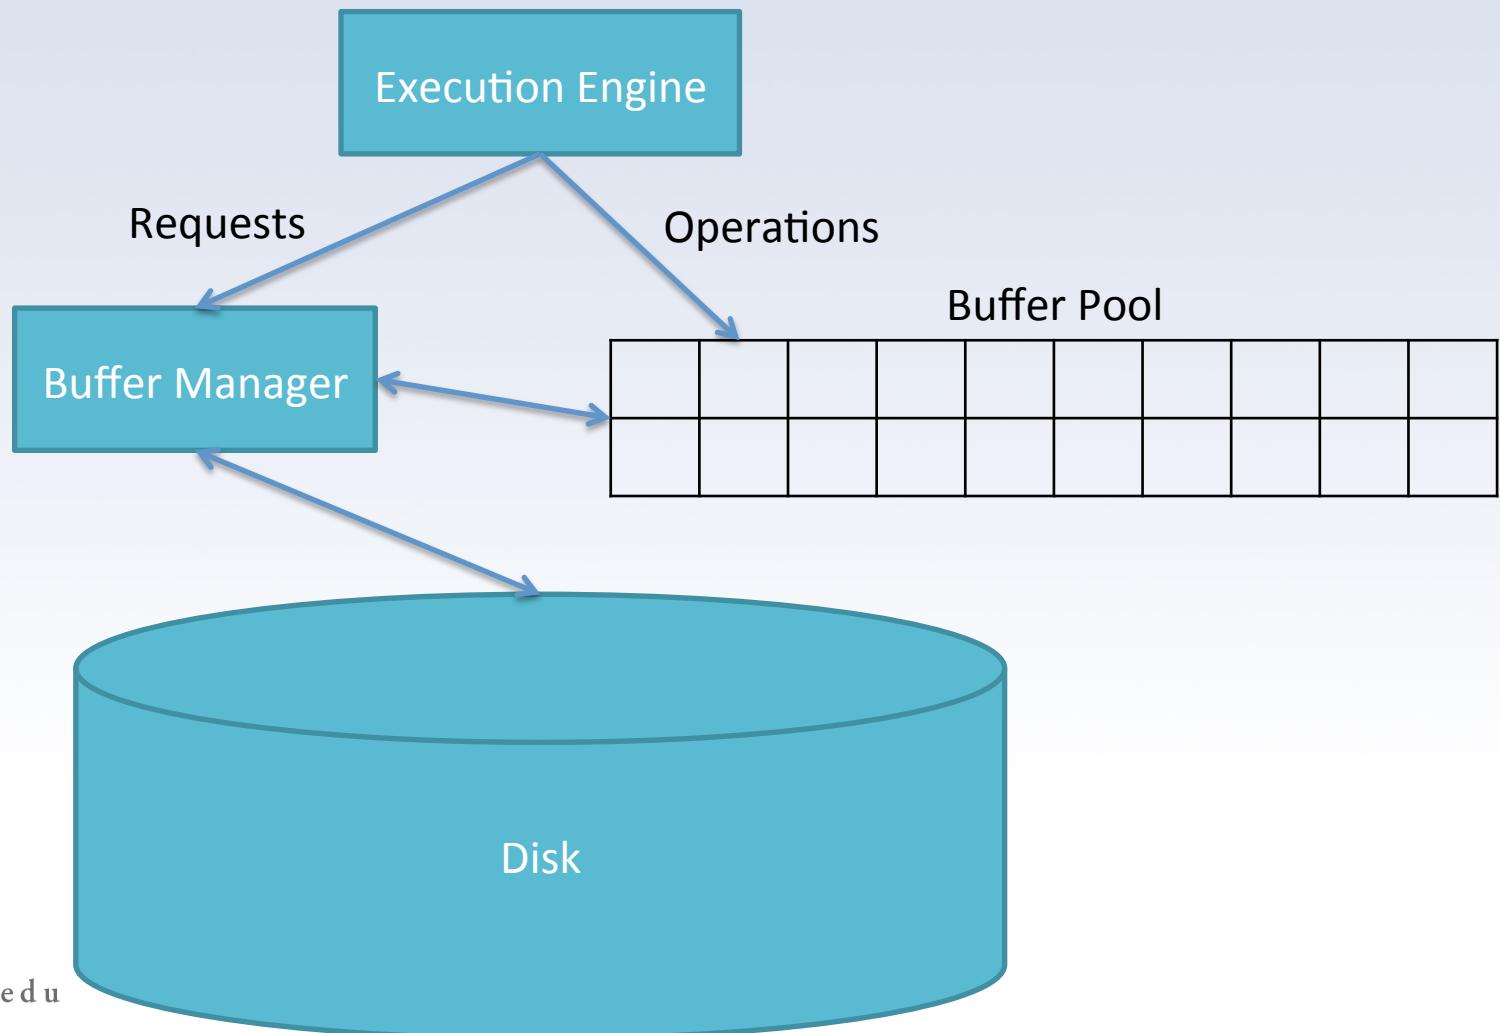


Buffer Management

- We divide memory up into block sized *buffers*
- The *buffer manager* is responsible for allocating buffers for query execution
- Query execution brings blocks into and out of the *buffer pool* for processing



Buffer management



Iterator Model

- Operators can all be implemented with three functions:
 - `Open()`: initializes the data structures
 - `GetNext()`: finds the next tuple
 - `Close()`: clean up
- Easy to combine operations
- Enables pipelining



Measuring Cost of Operation

- We only care about I/O operations
- We also ignore output (write) costs:
 - Assume we can pipeline operations
 - Result of query is independent of execution strategy



Measuring Cost of Operation

- M - the number of blocks that fit in memory (called *buffers*)
- Given a relation R
 - $B(R)$ - number of blocks to store R
 - $T(R)$ - number of tuples in R
 - $V(R,a)$ - number of distinct values for attribute a in relation R



Table Scan

- Pull records of a relation into memory one at a time
- Important operator not in relational algebra
- Example:

```
SELECT * FROM R
```



Table Scan

```
void open() {
    block* b=R.first_block;
    record* t=b.first_tuple;
}
```



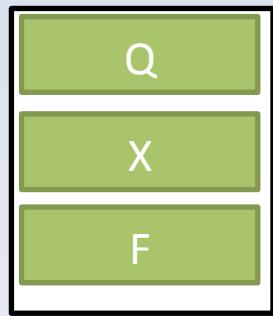
Table Scan

```
record* getNext( ){  
    if(t==NULL){  
        b=b.next_block;  
        if(b==NULL){return NULL; }  
        t=b.first_tuple;  
    }  
    record* old_t=t;  
    t+=(sizeof(record));  
    return old_t;  
}
```

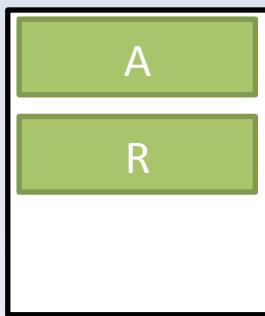


Example

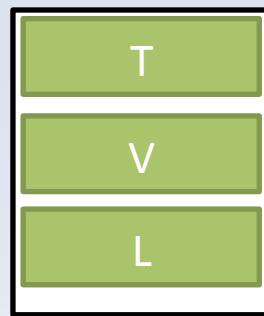
block 0



block 1



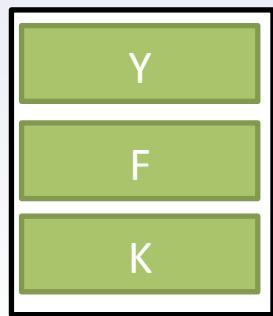
block 2



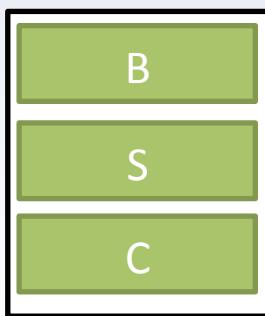
block 3



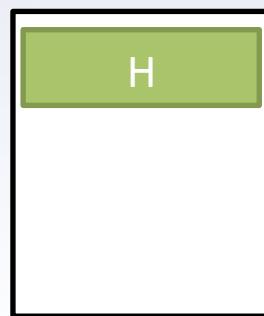
block 4



block 5



block 6



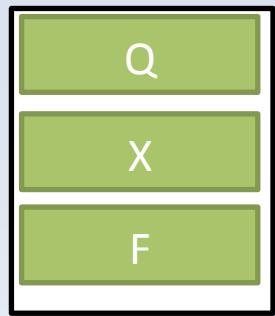
block 7



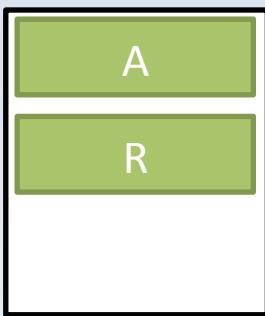
Example

main memory

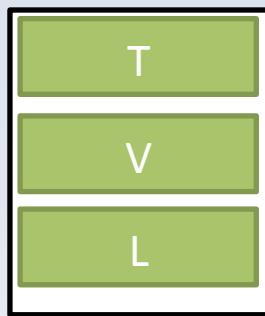
block 0



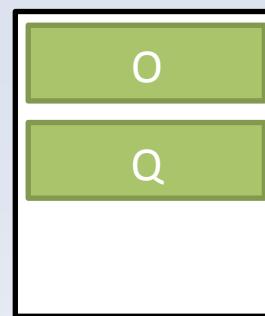
block 1



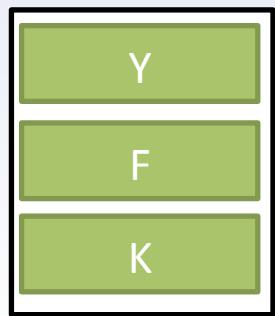
block 2



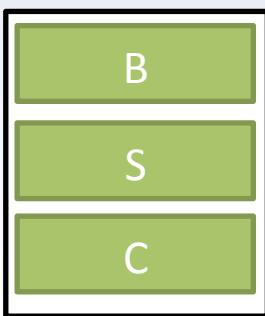
block 3



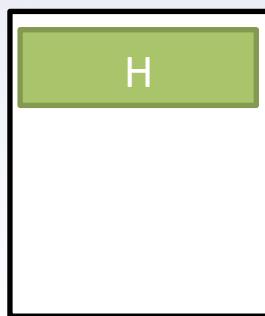
block 4



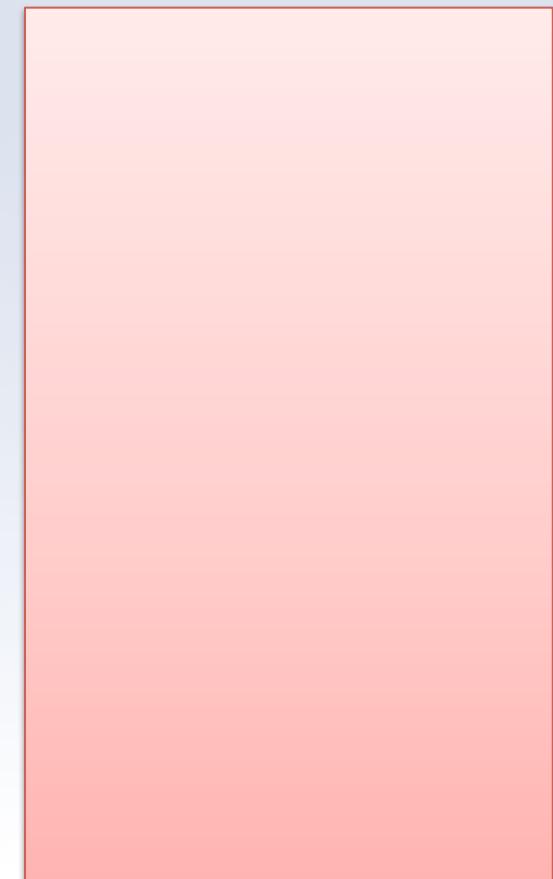
block 5



block 6



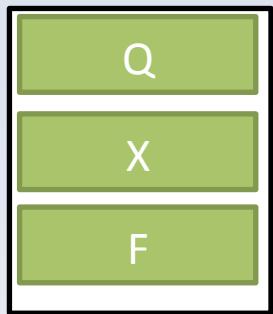
block 7



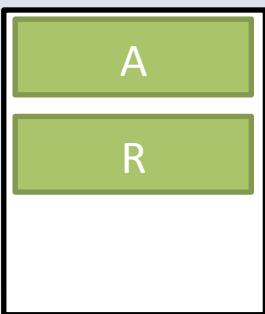
Example

main memory

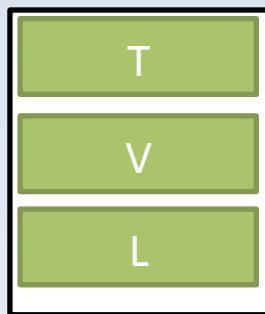
block 0



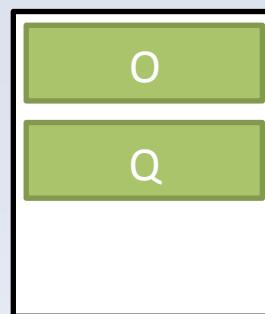
block 1



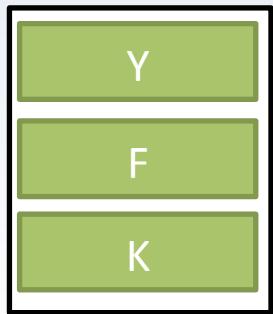
block 2



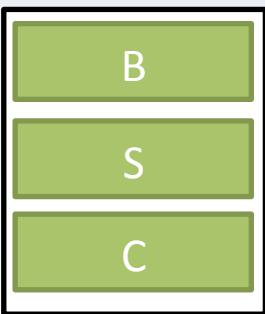
block 3



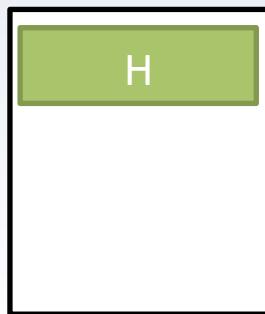
block 4



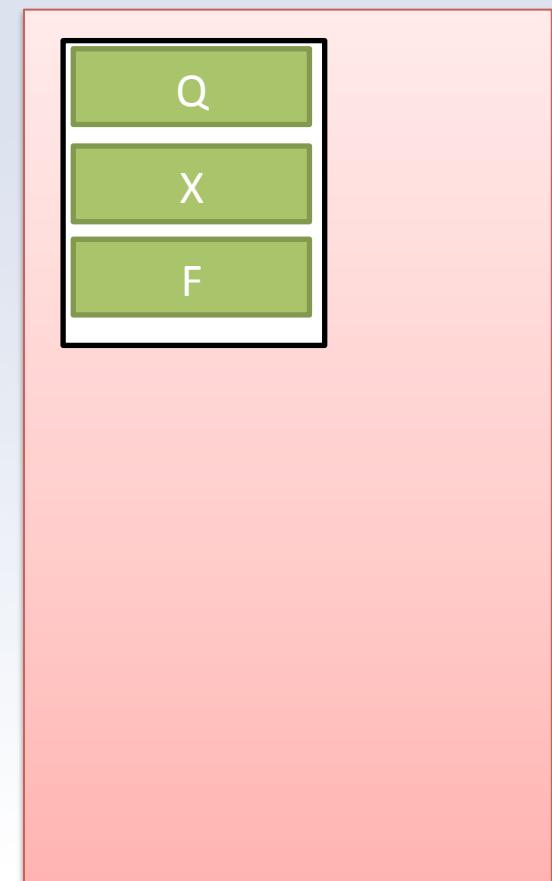
block 5



block 6



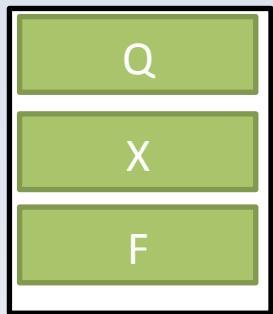
block 7



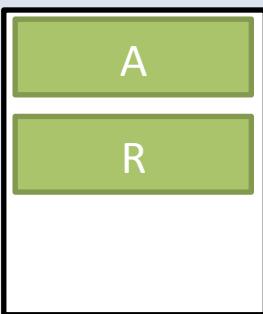
Example

main memory

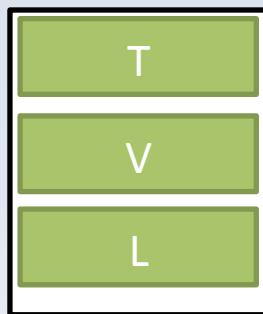
block 0



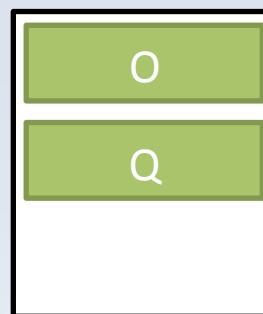
block 1



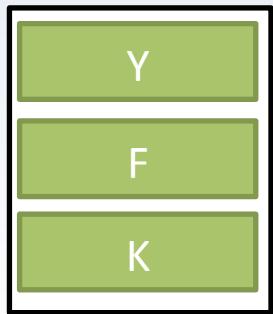
block 2



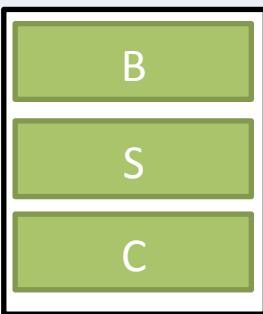
block 3



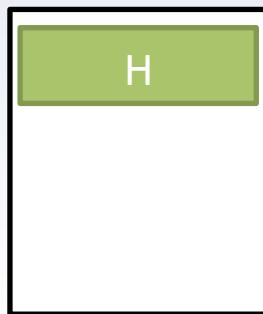
block 4



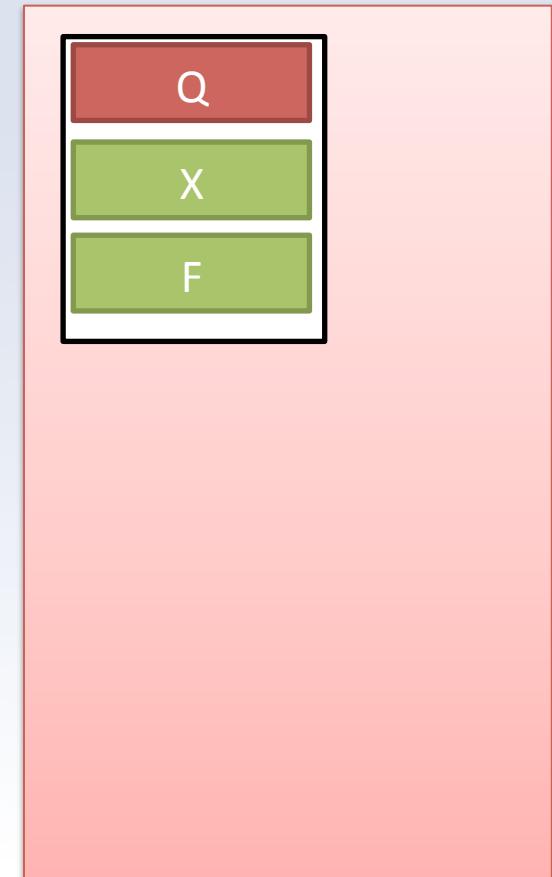
block 5



block 6



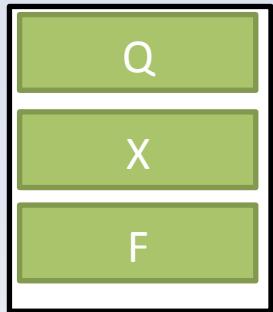
block 7



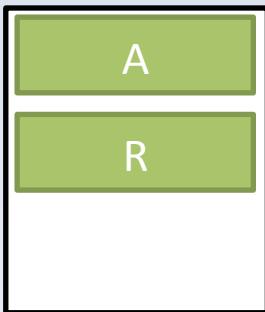
Example

main memory

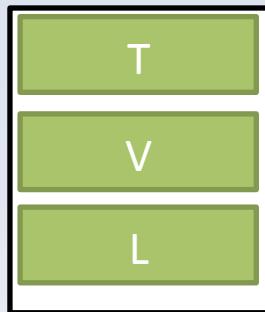
block 0



block 1



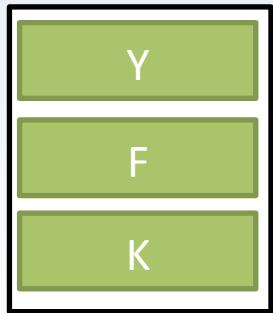
block 2



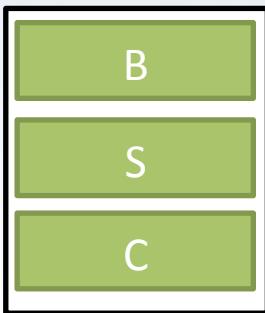
block 3



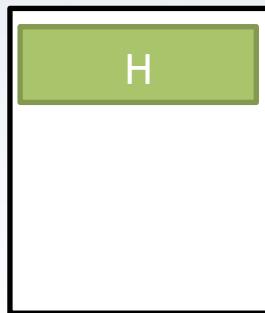
block 4



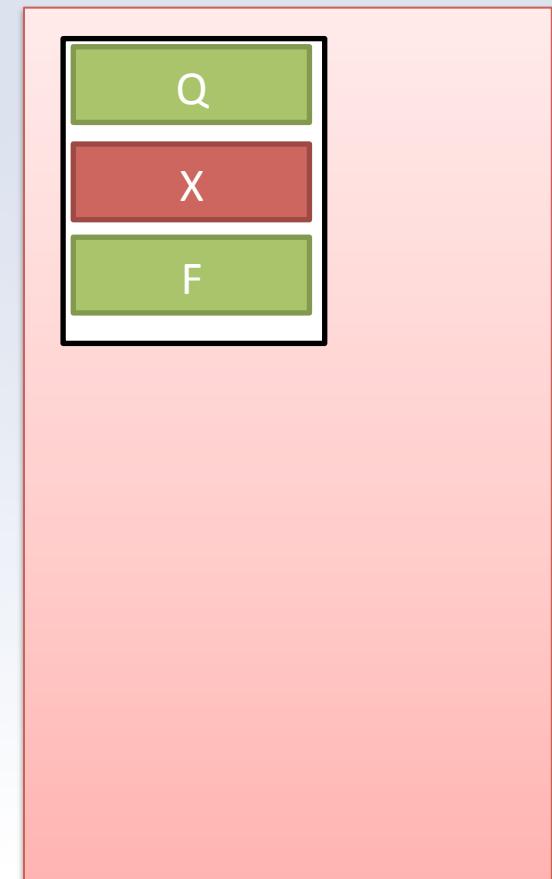
block 5



block 6



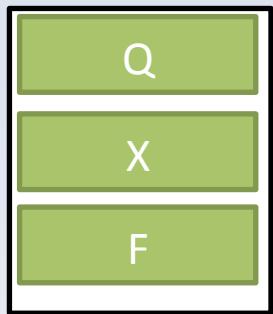
block 7



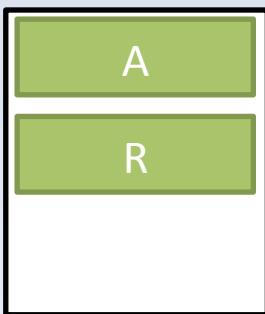
Example

main memory

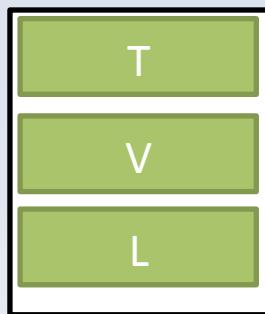
block 0



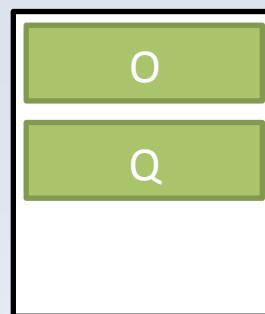
block 1



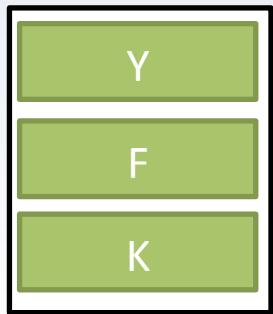
block 2



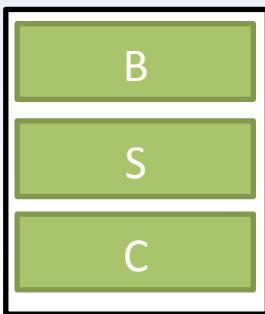
block 3



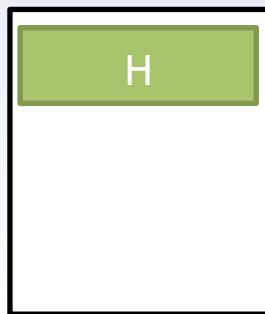
block 4



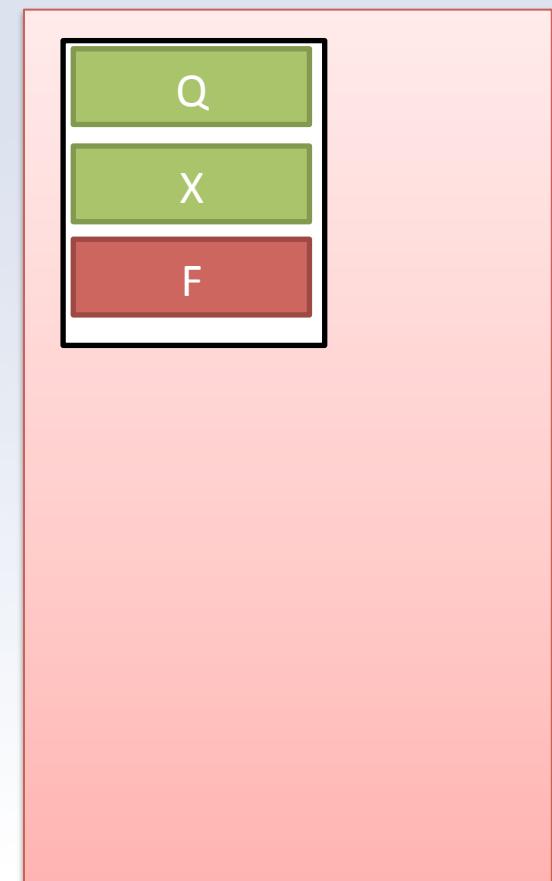
block 5



block 6



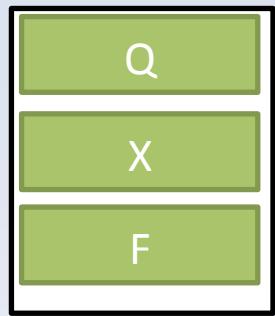
block 7



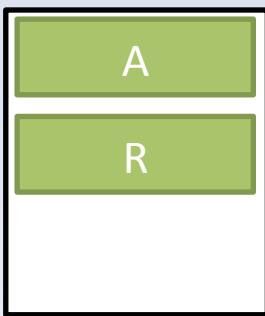
Example

main memory

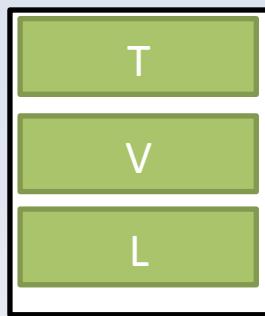
block 0



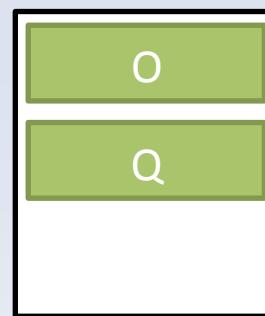
block 1



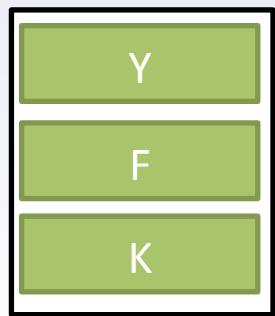
block 2



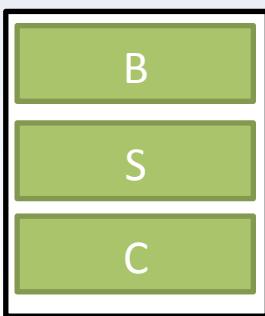
block 3



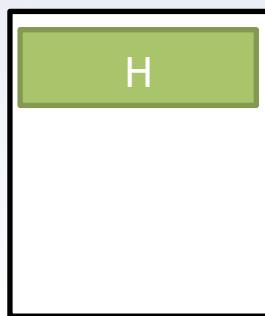
block 4



block 5



block 6



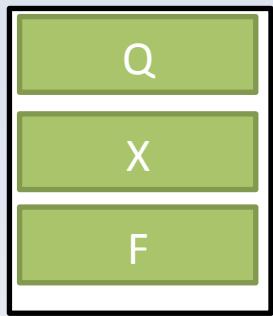
block 7



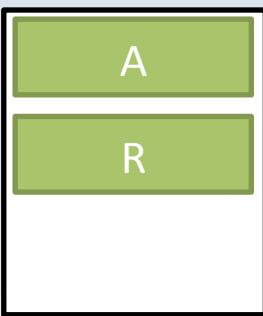
Example

main memory

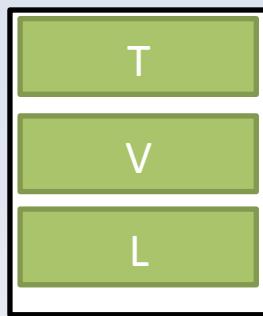
block 0



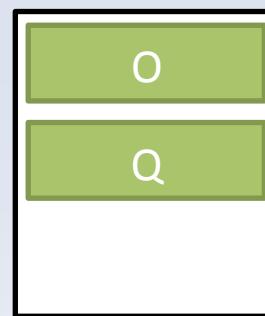
block 1



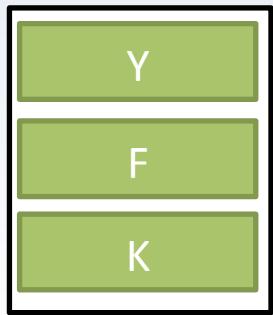
block 2



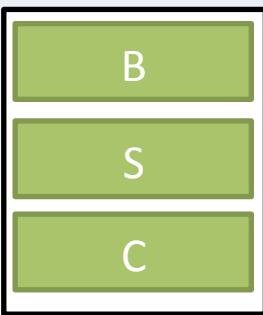
block 3



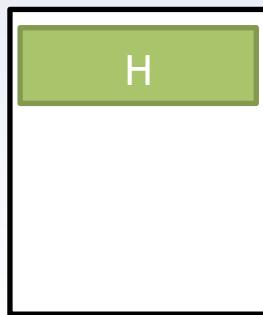
block 4



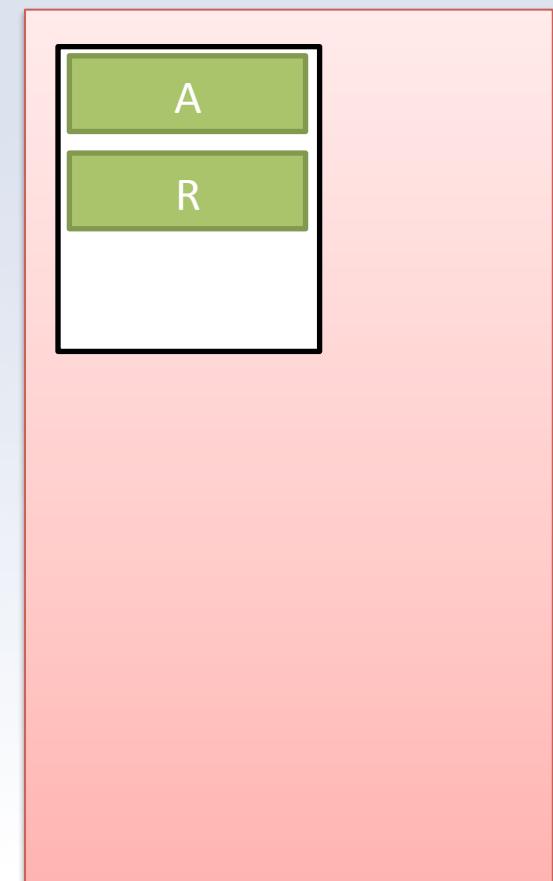
block 5



block 6



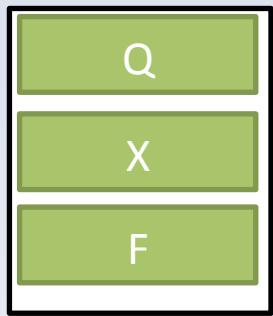
block 7



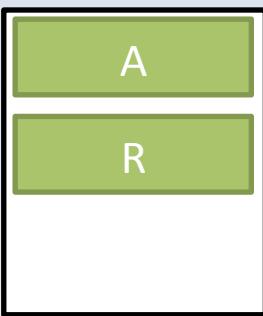
Example

main memory

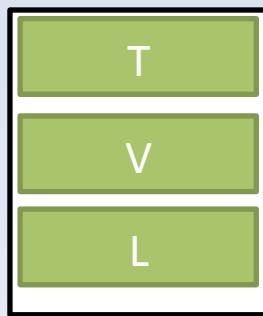
block 0



block 1



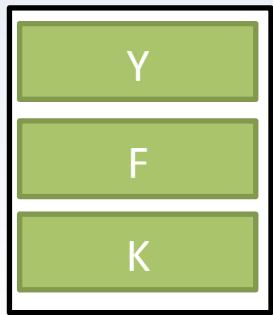
block 2



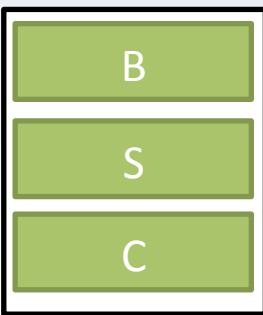
block 3



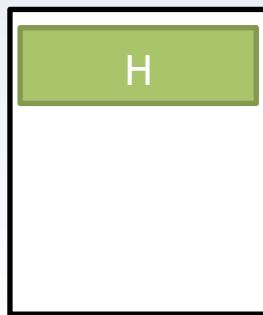
block 4



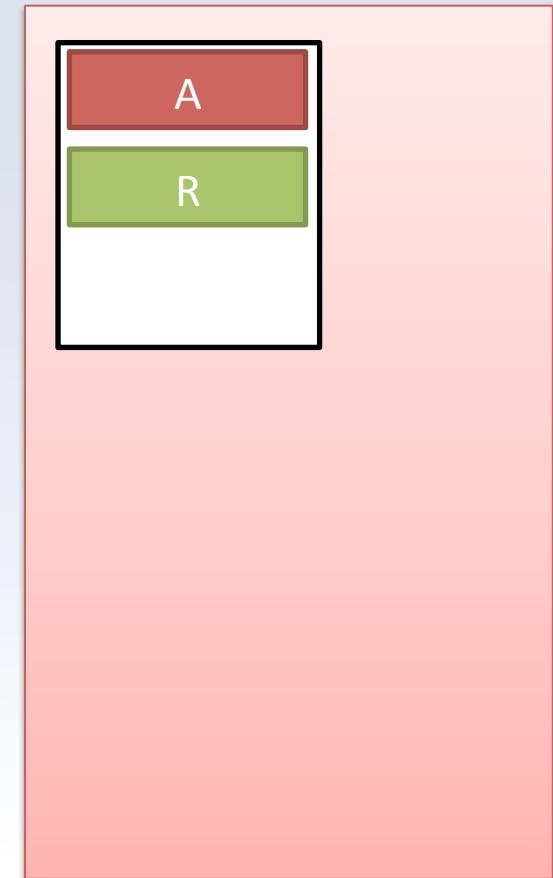
block 5



block 6



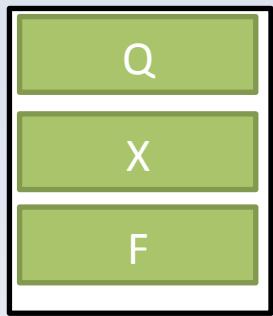
block 7



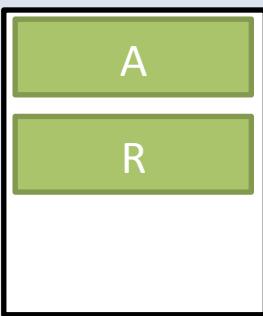
Example

main memory

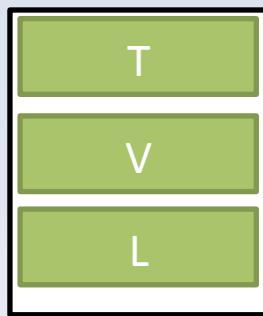
block 0



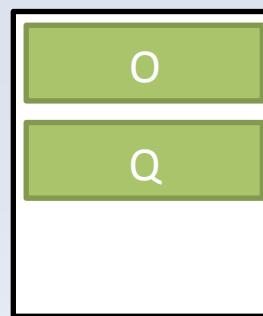
block 1



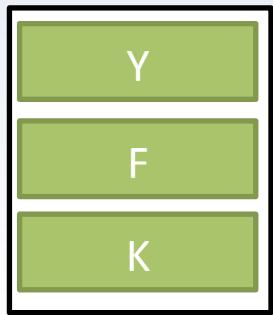
block 2



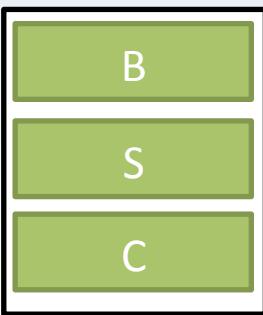
block 3



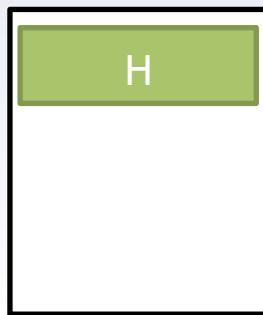
block 4



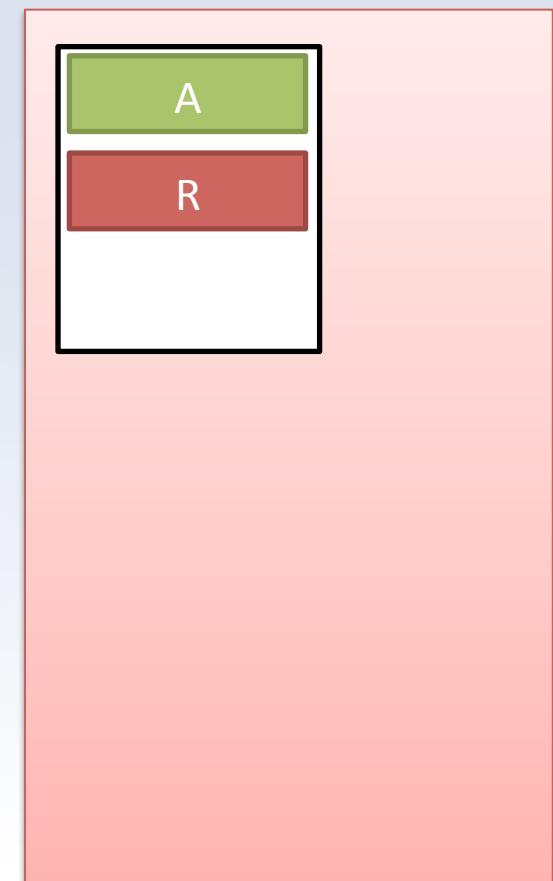
block 5



block 6



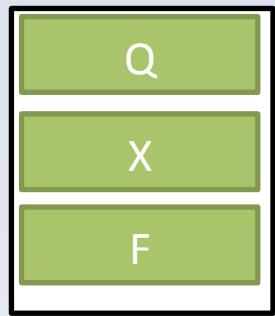
block 7



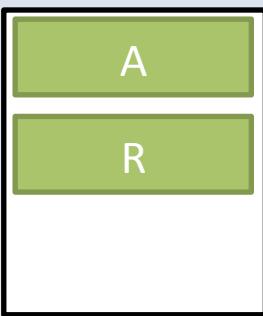
Example

main memory

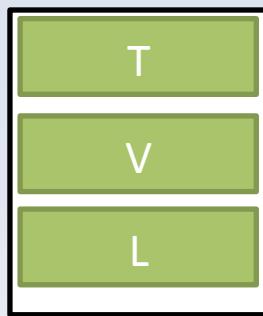
block 0



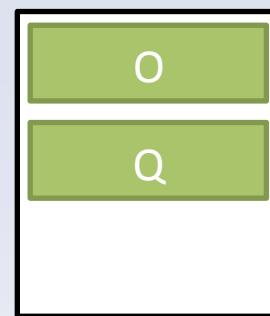
block 1



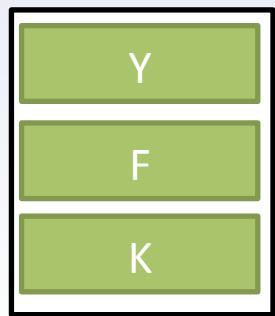
block 2



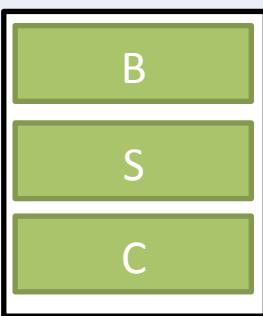
block 3



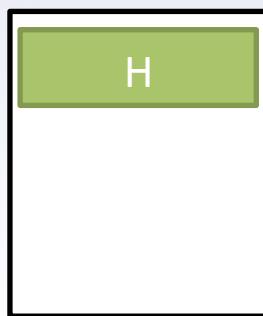
block 4



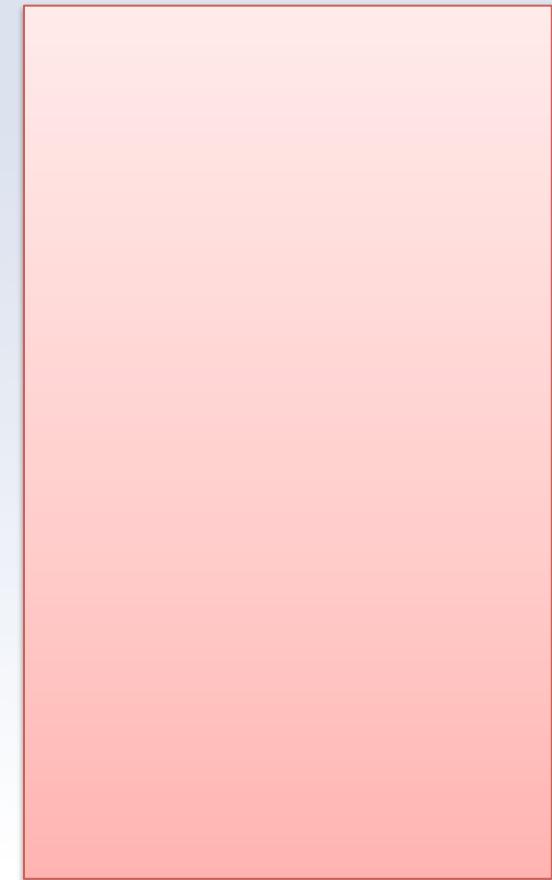
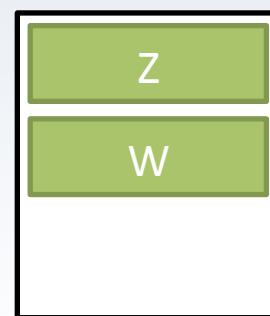
block 5



block 6



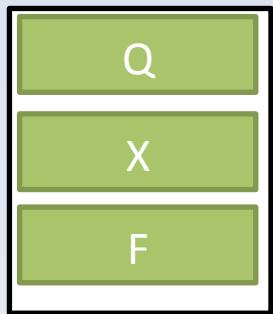
block 7



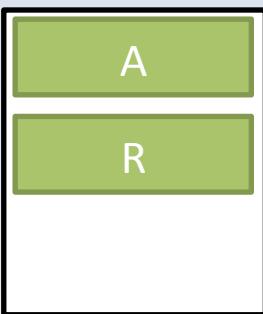
Example

main memory

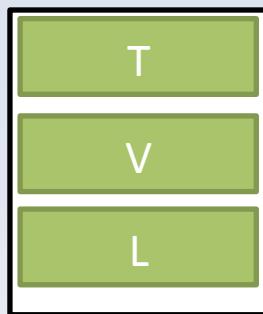
block 0



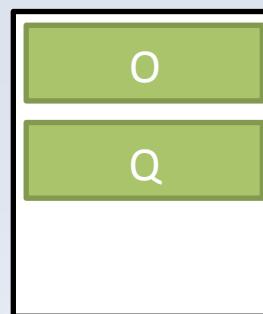
block 1



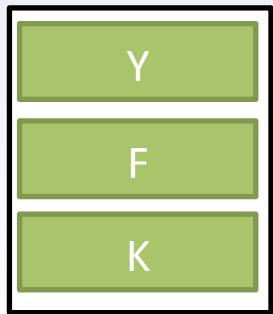
block 2



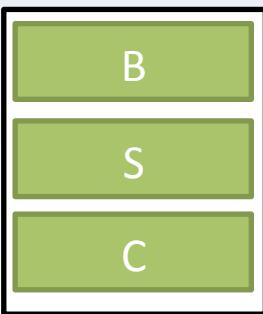
block 3



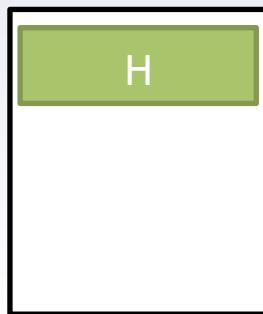
block 4



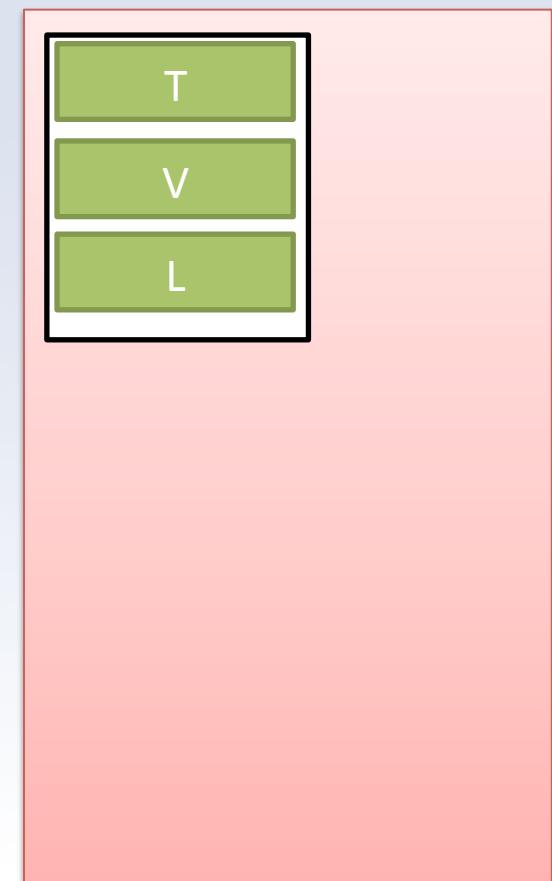
block 5



block 6



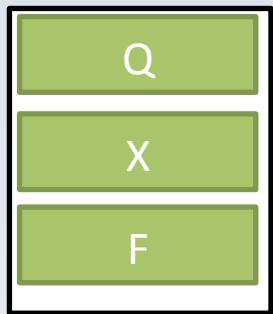
block 7



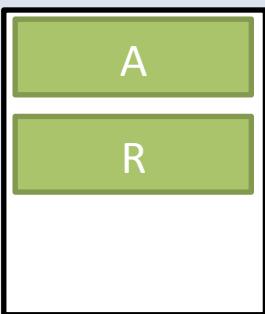
Example

main memory

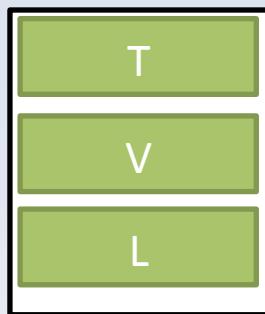
block 0



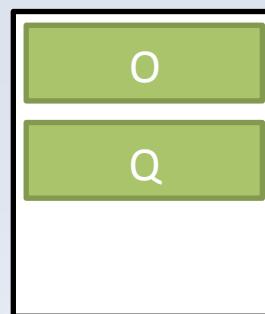
block 1



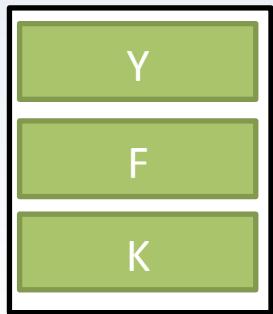
block 2



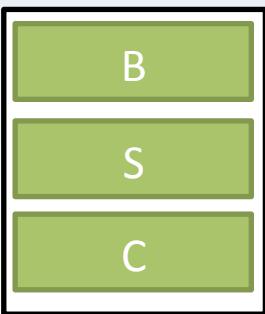
block 3



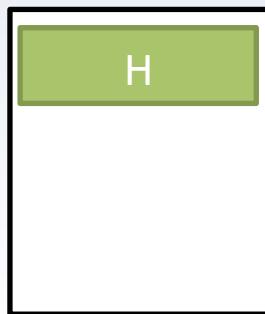
block 4



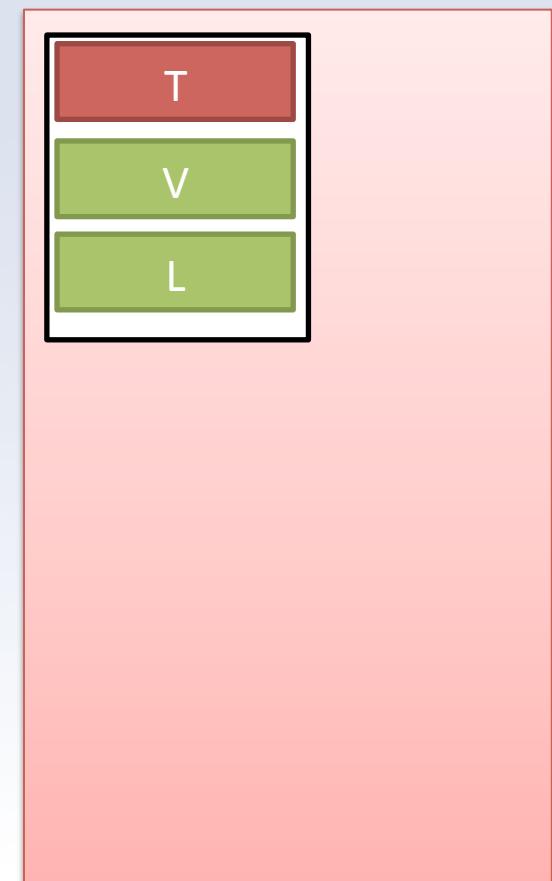
block 5



block 6



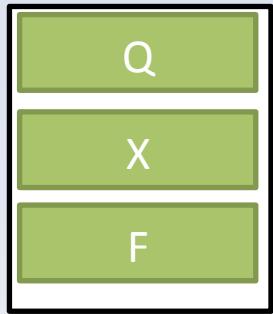
block 7



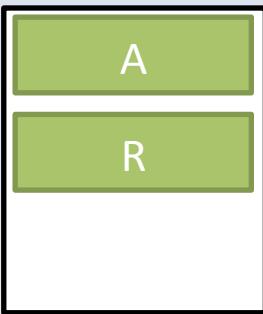
Example

main memory

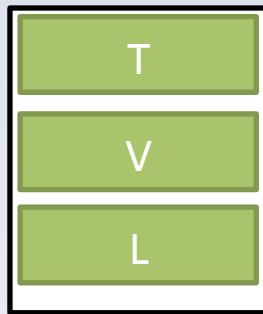
block 0



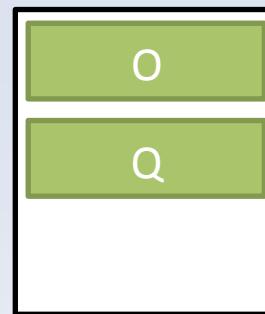
block 1



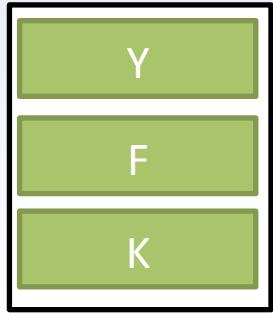
block 2



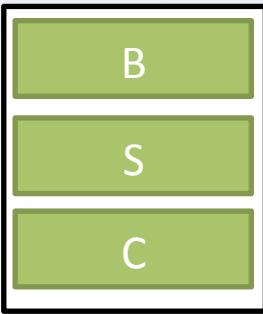
block 3



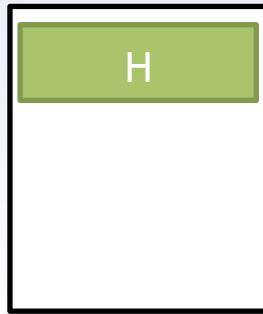
block 4



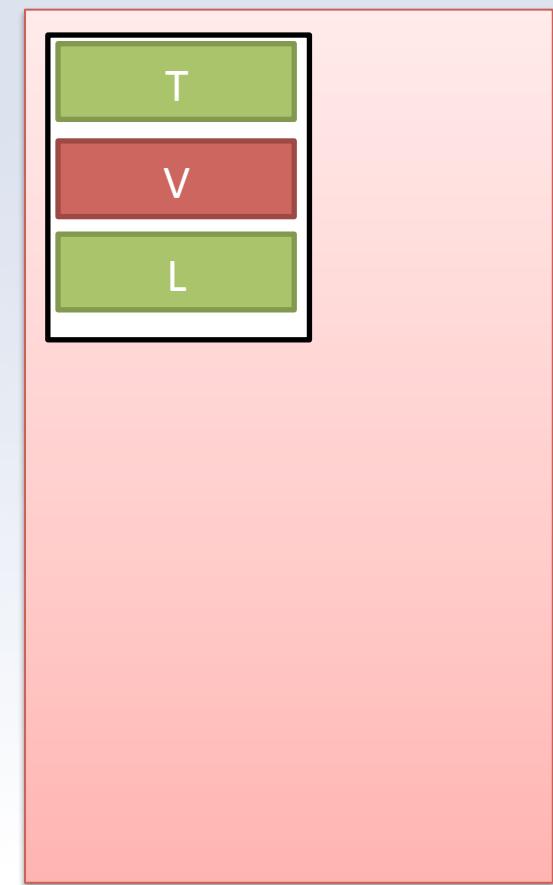
block 5



block 6



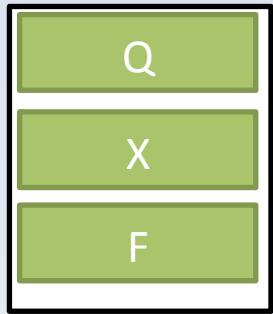
block 7



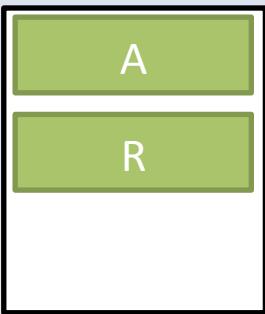
Example

main memory

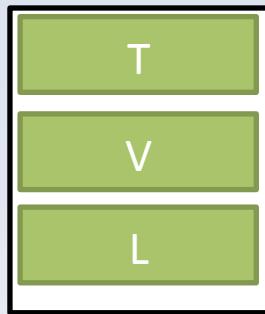
block 0



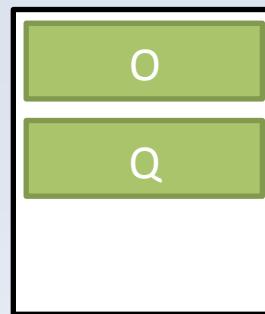
block 1



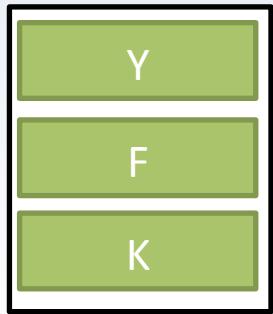
block 2



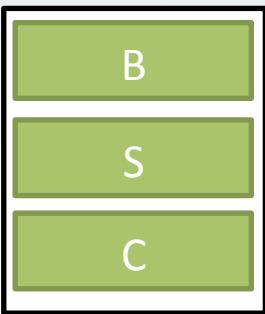
block 3



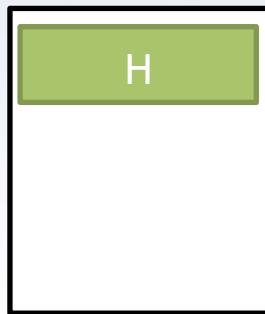
block 4



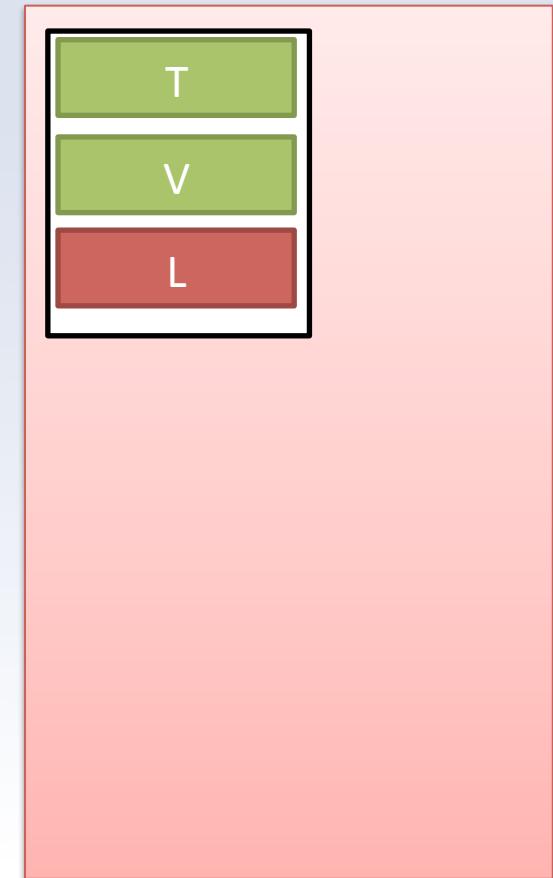
block 5



block 6



block 7



Cost of table scan

- Must read each block once
- Cost is therefore $B(R)$
 - number of blocks storing records from R



One-Pass Algorithms

- Sometimes, we can do an operation passing over the data once
 - Example: table scan
- Usually, relation has to be small
- You probably implemented some of these in MP1

Tuple-at-a-Time

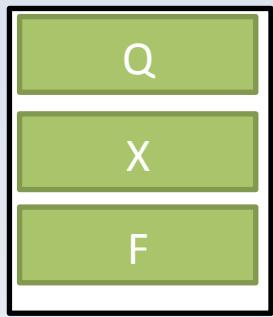
- σ and π require *at most* one pass:
 - do a table scan, but only return desired rows/ columns
- Cost: $B(R)$



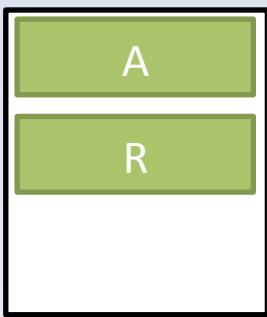
Example: <‘G’

main memory

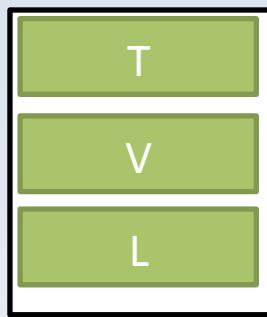
block 0



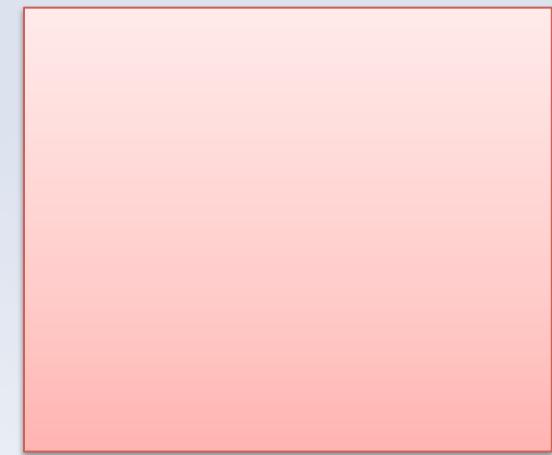
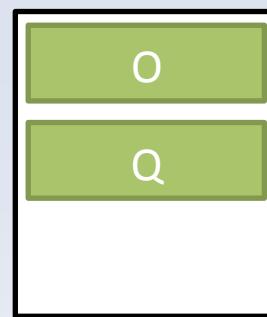
block 1



block 2



block 3



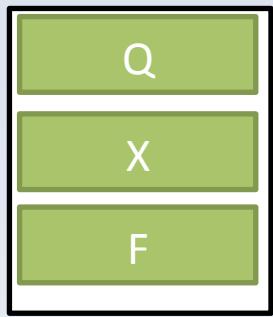
output



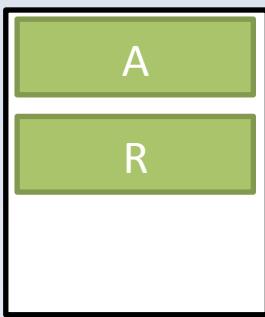
Example: <‘G’

main memory

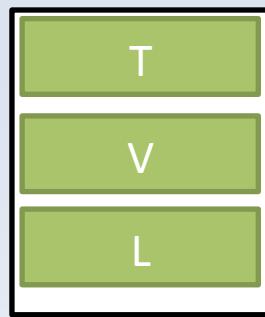
block 0



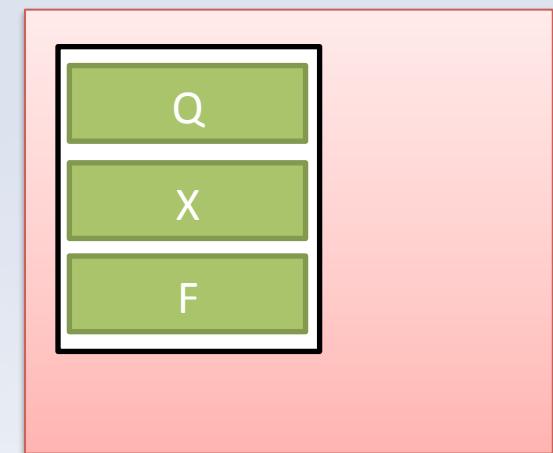
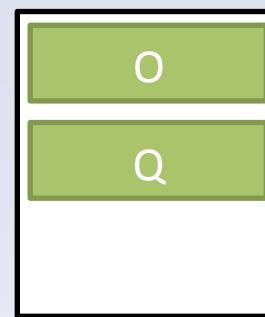
block 1



block 2



block 3



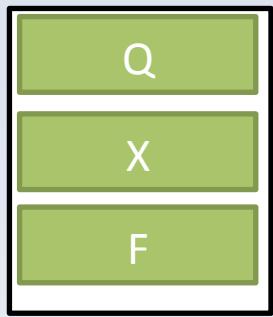
output



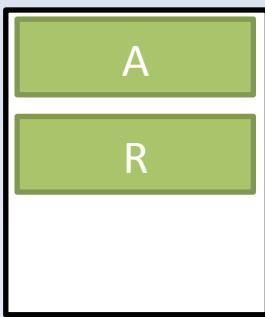
Example: <‘G’

main memory

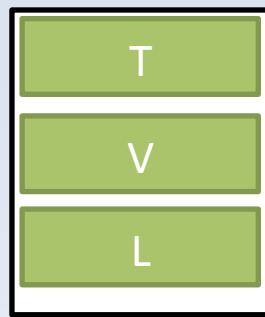
block 0



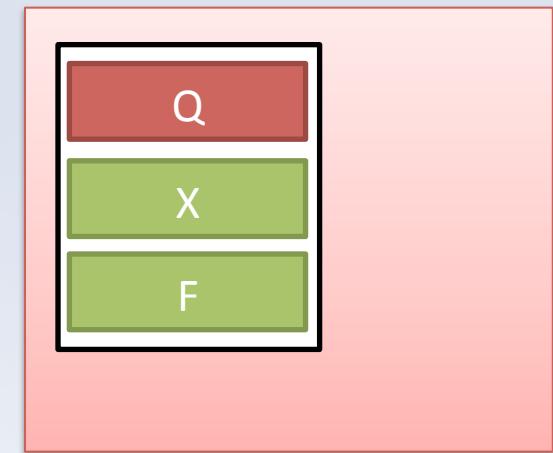
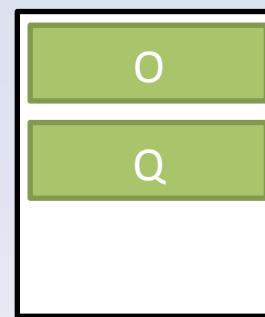
block 1



block 2



block 3



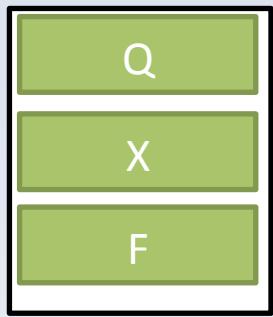
output



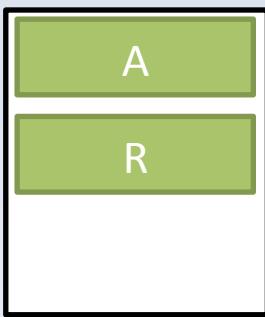
Example: <‘G’

main memory

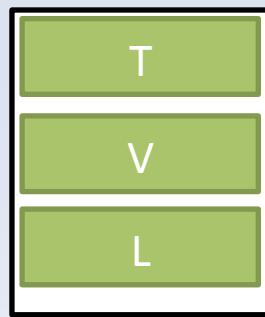
block 0



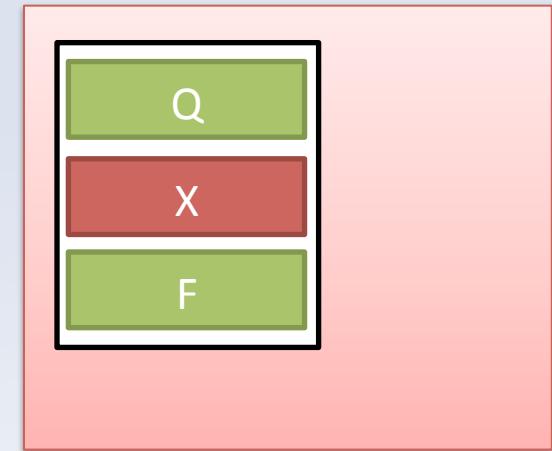
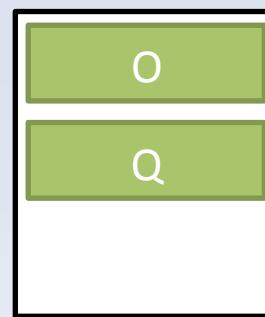
block 1



block 2



block 3



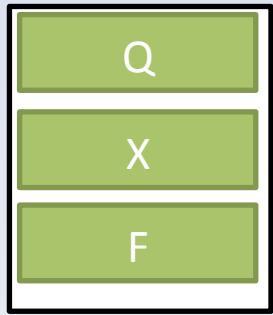
output



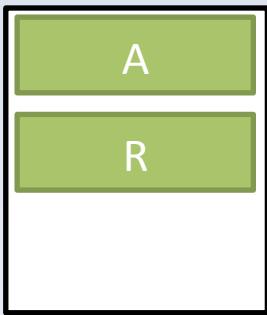
Example: <‘G’

main memory

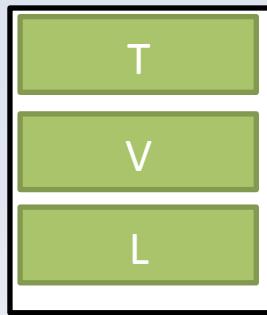
block 0



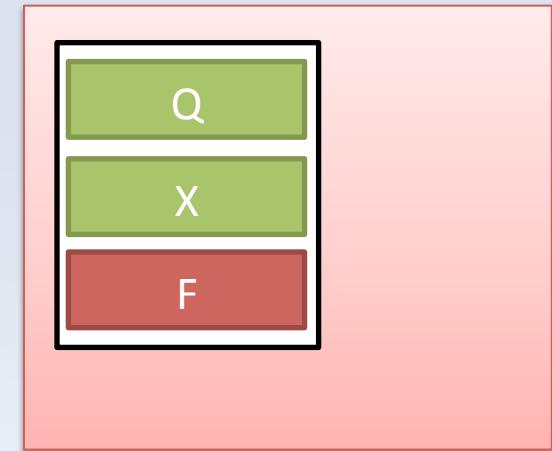
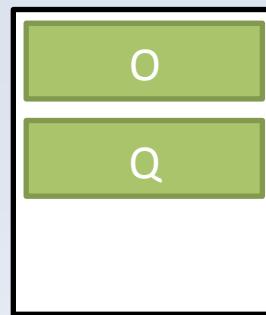
block 1



block 2



block 3



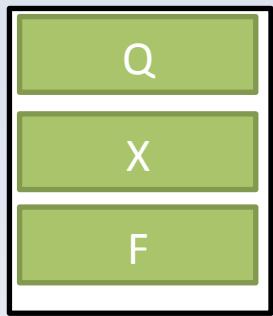
output



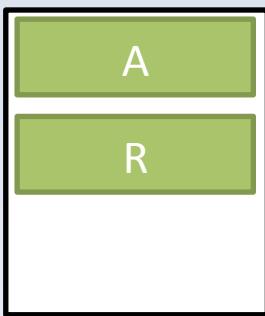
Example: <‘G’

main memory

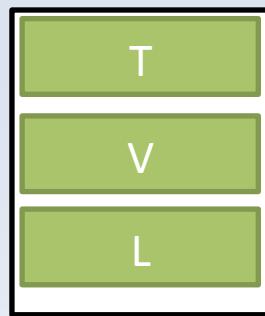
block 0



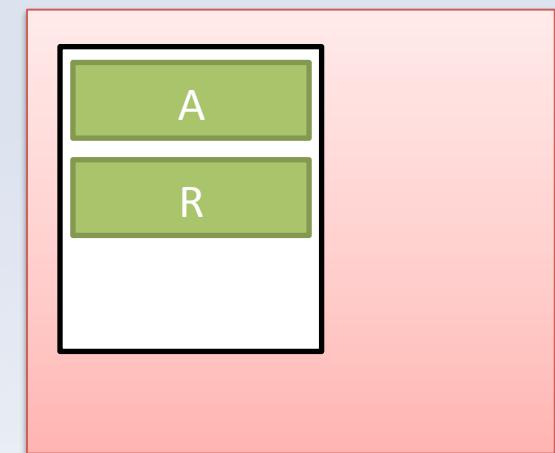
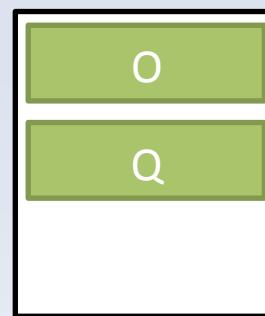
block 1



block 2



block 3



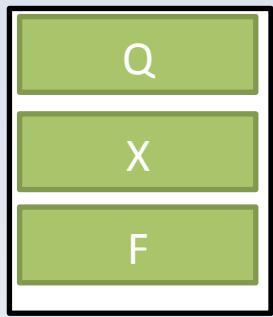
output



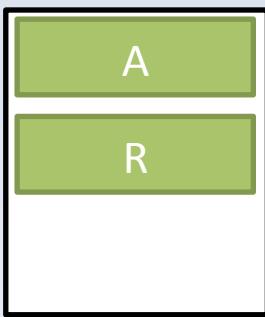
Example: <‘G’

main memory

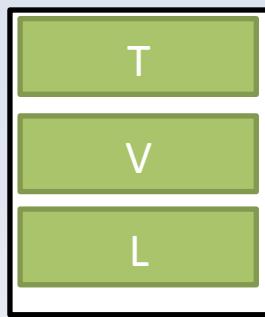
block 0



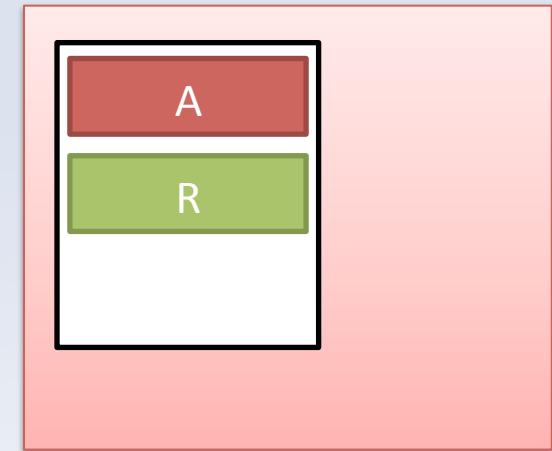
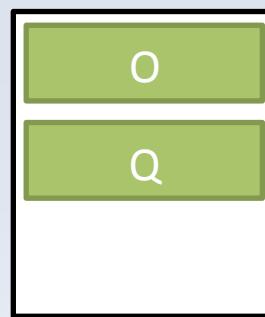
block 1



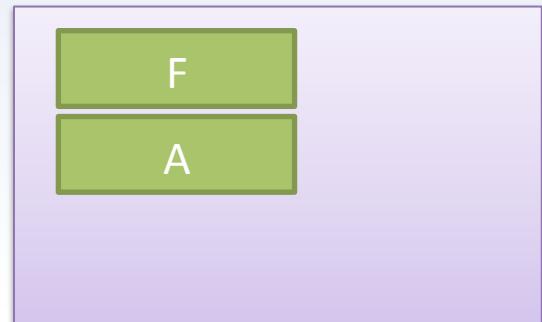
block 2



block 3



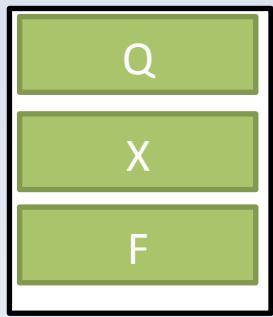
output



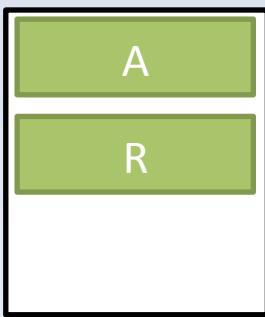
Example: <‘G’

main memory

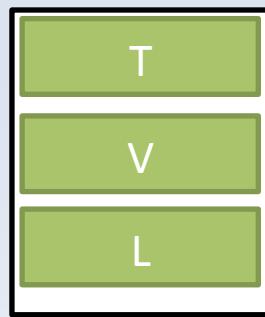
block 0



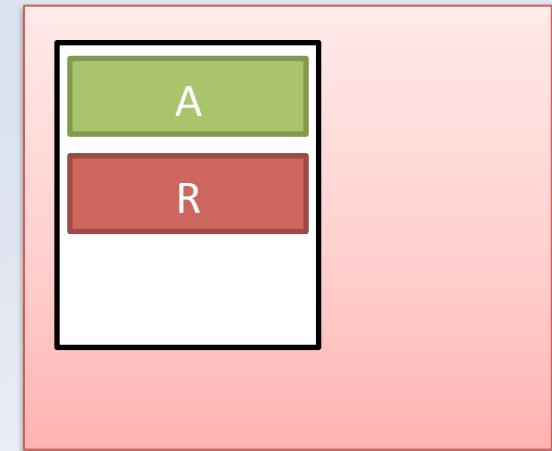
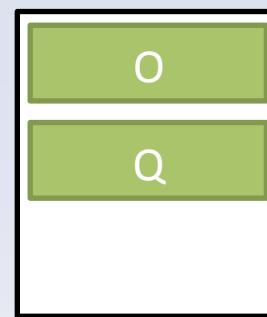
block 1



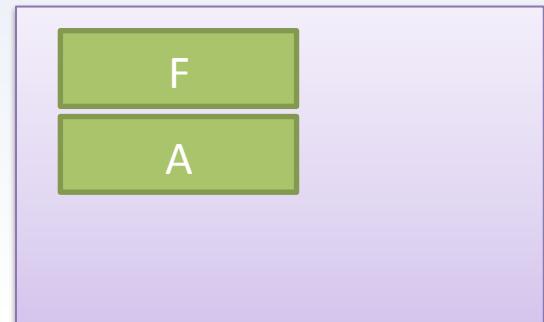
block 2



block 3

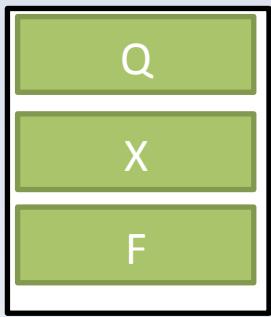


output

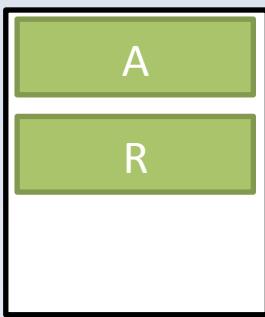


Example: <‘G’

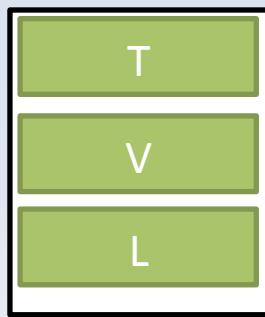
block 0



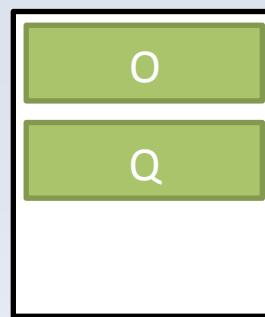
block 1



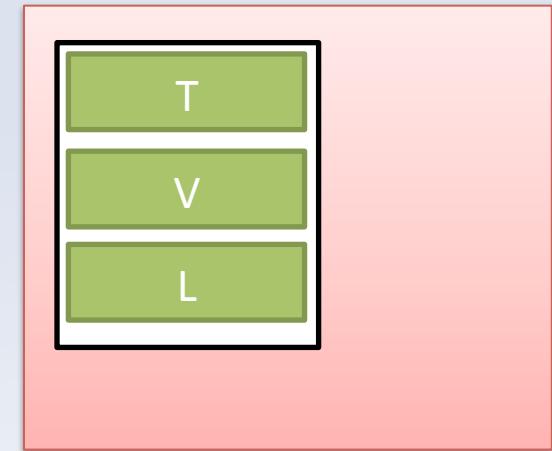
block 2



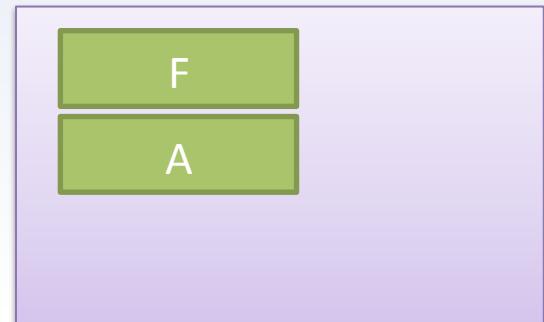
block 3



main memory

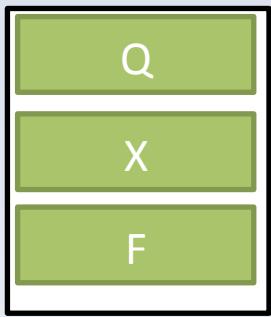


output

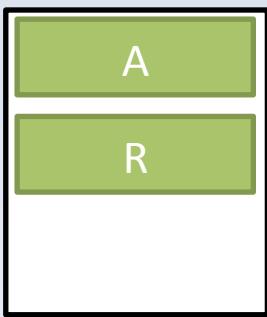


Example: <‘G’

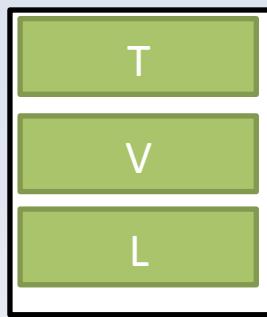
block 0



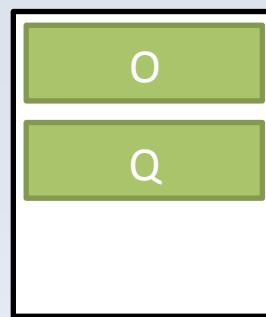
block 1



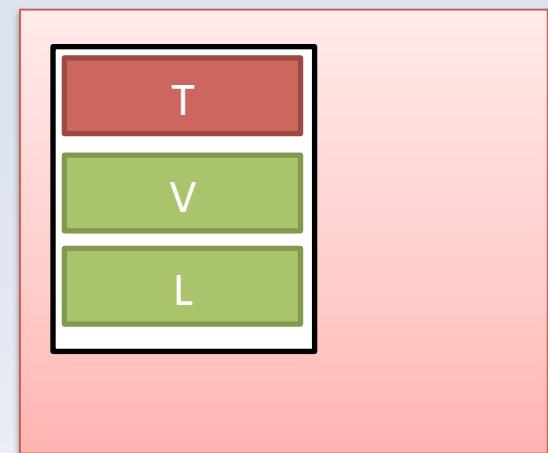
block 2



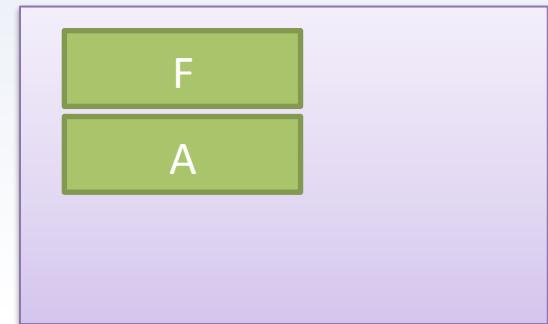
block 3



main memory



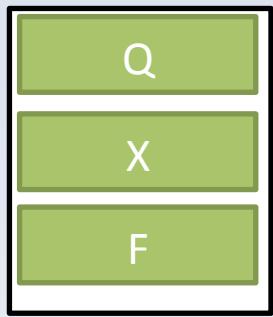
output



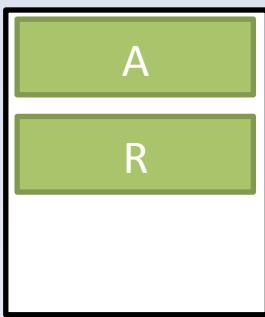
Example: <‘G’

main memory

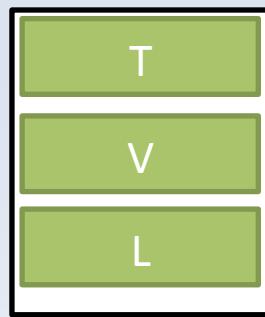
block 0



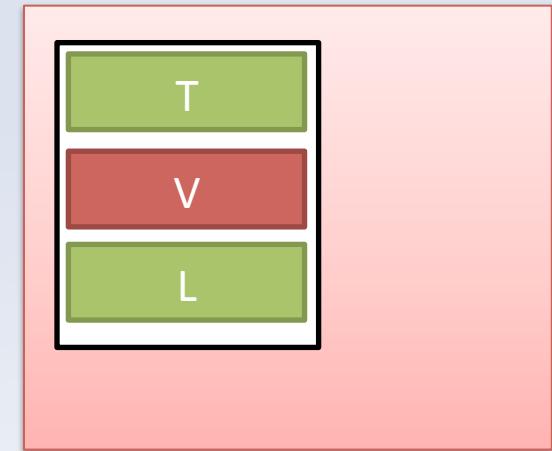
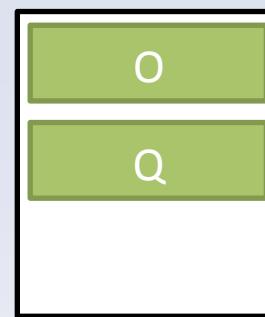
block 1



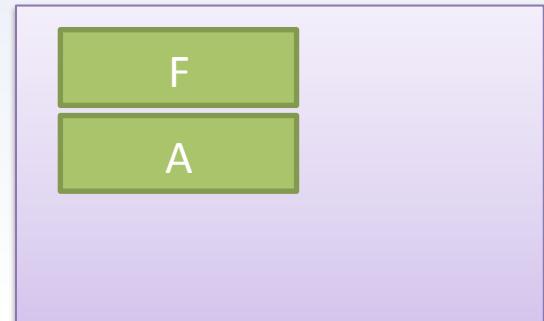
block 2



block 3

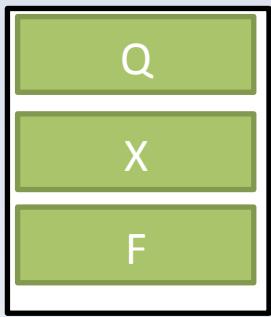


output

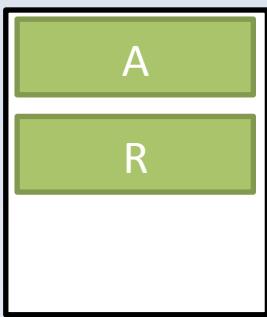


Example: <‘G’

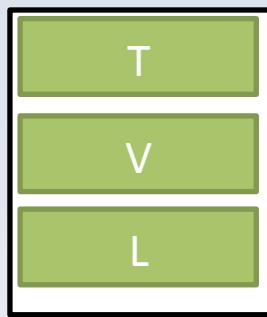
block 0



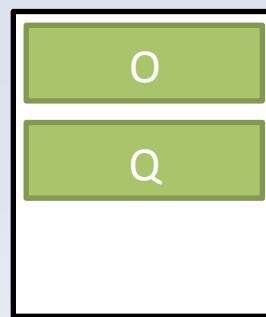
block 1



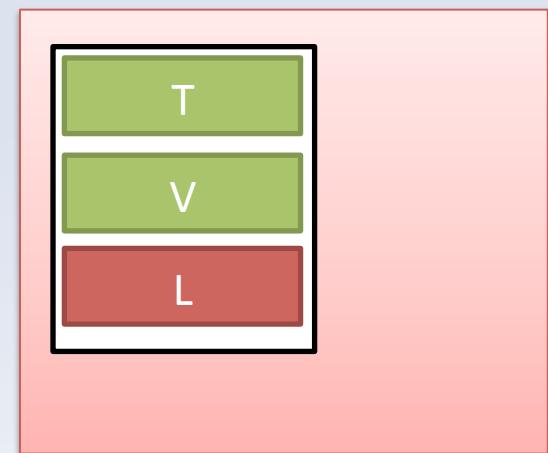
block 2



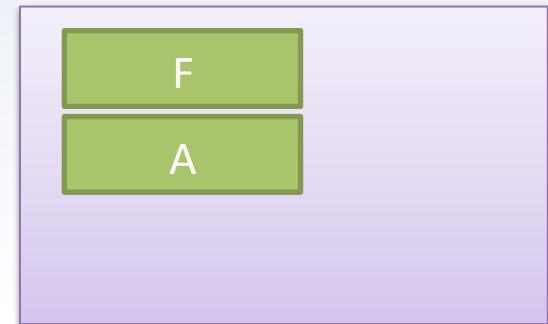
block 3



main memory

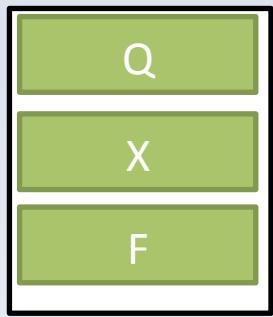


output

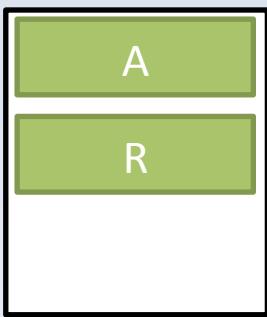


Example: <‘G’

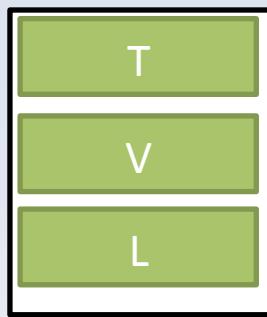
block 0



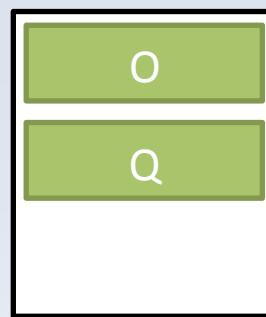
block 1



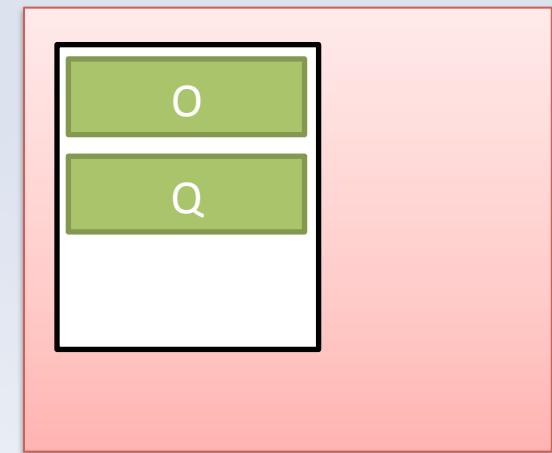
block 2



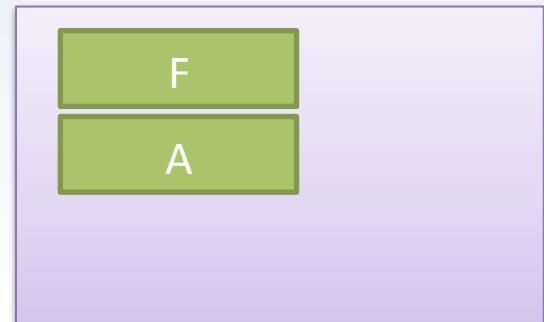
block 3



main memory

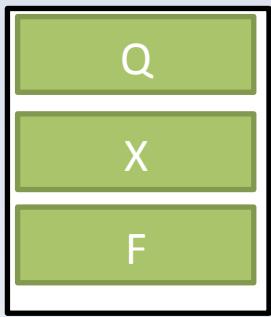


output

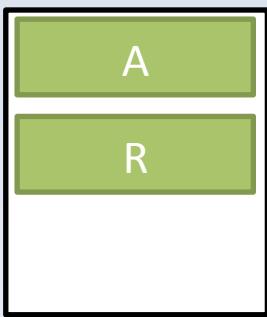


Example: <‘G’

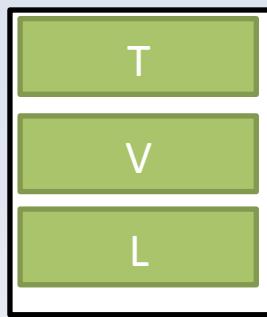
block 0



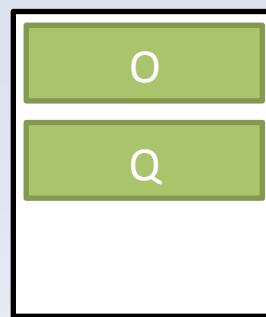
block 1



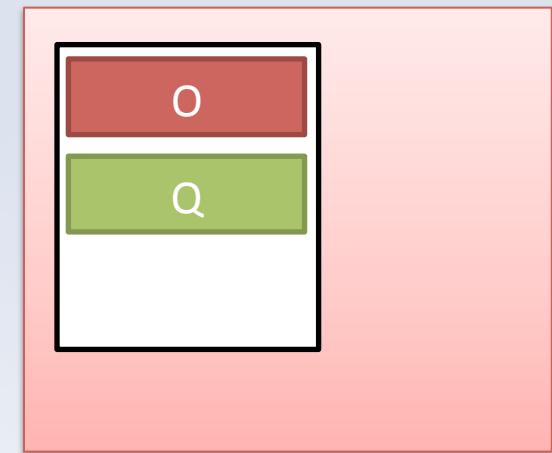
block 2



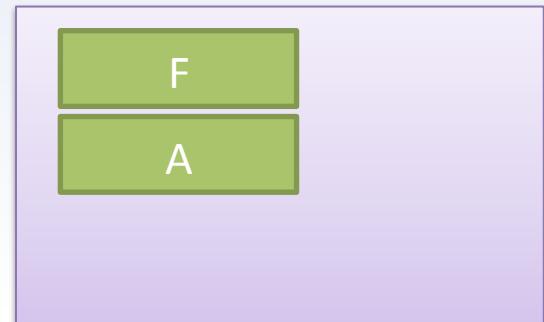
block 3



main memory

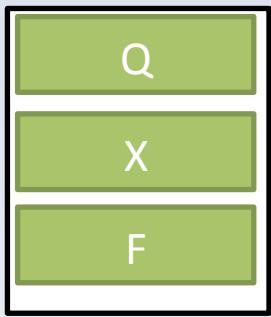


output

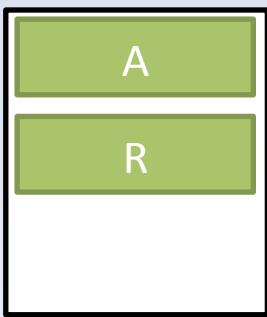


Example: <‘G’

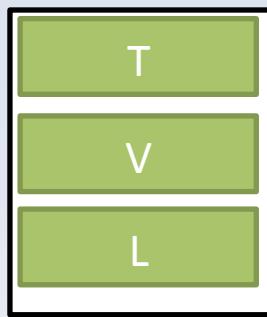
block 0



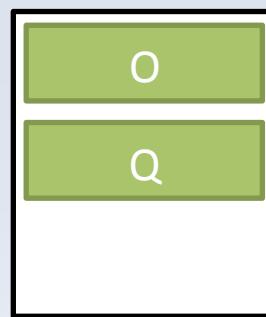
block 1



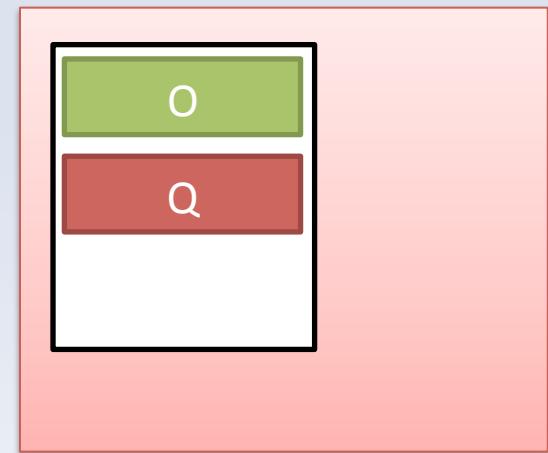
block 2



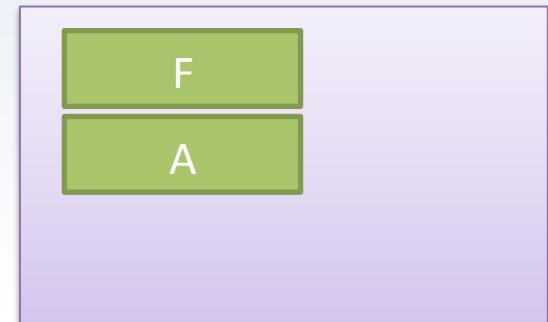
block 3



main memory



output



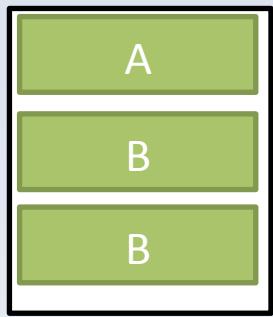
Unary Full-Relational

- For δ and γ , need to keep a *dictionary* in memory
 - Track what values we've seen so far
 - Also track aggregation results for γ
- Cost: $B(R)$
 - Assumes:
 - $B(\delta(R)) \leq M$
 - $B(\gamma(R)) \leq M$

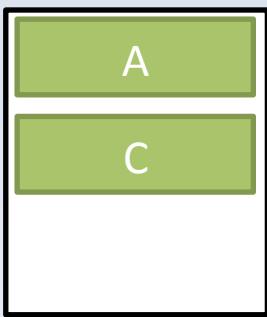


Example: δ

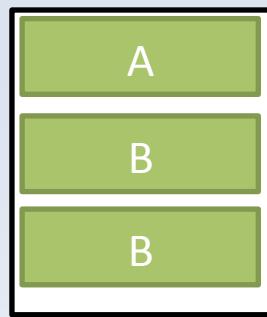
block 0



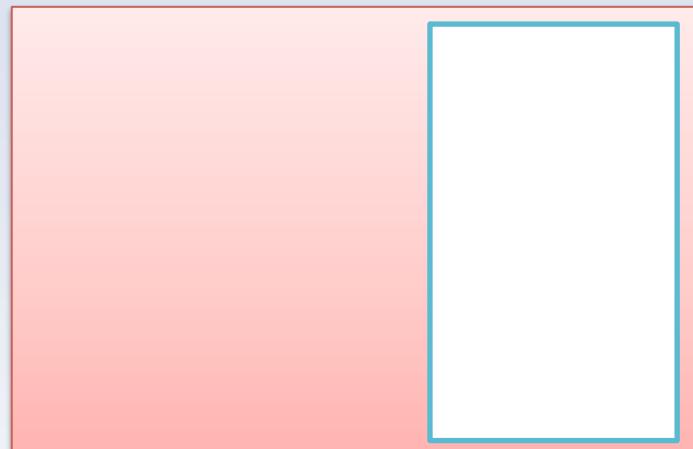
block 1



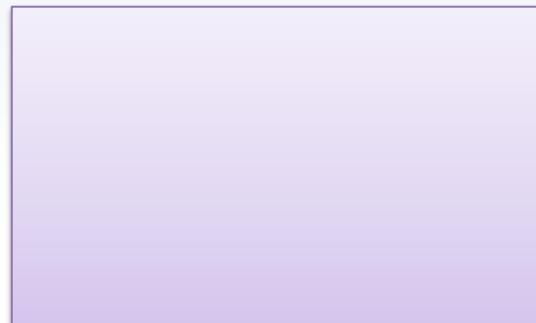
block 2



main memory

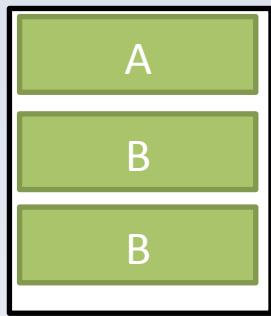


output

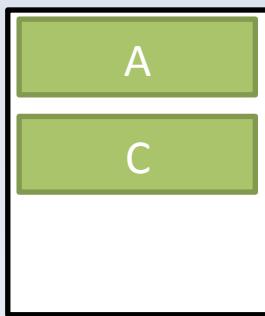


Example: δ

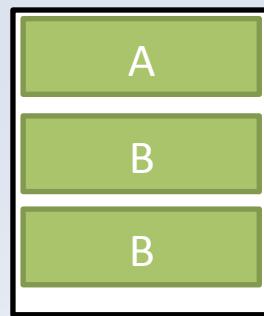
block 0



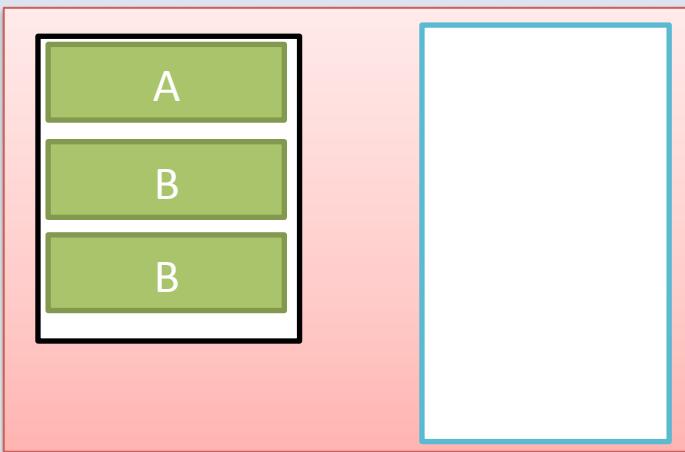
block 1



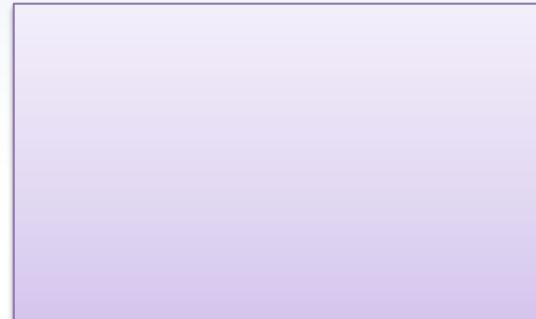
block 2



main memory

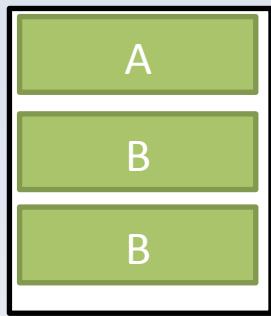


output

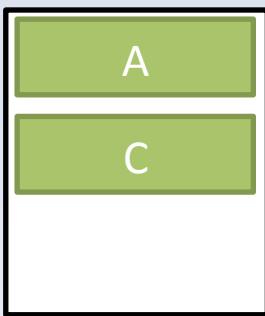


Example: δ

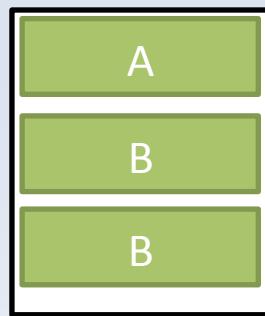
block 0



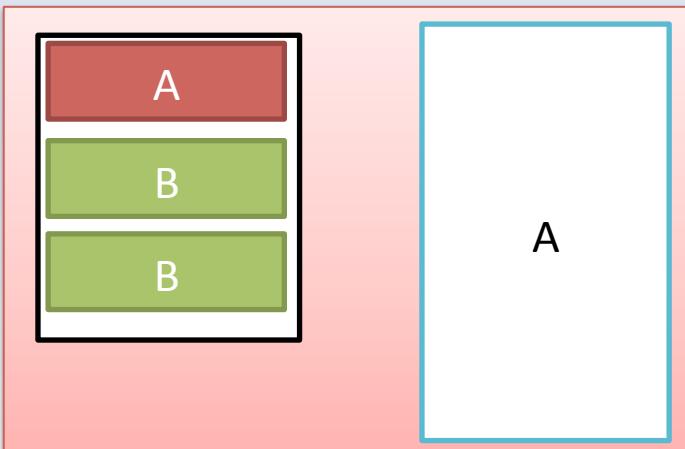
block 1



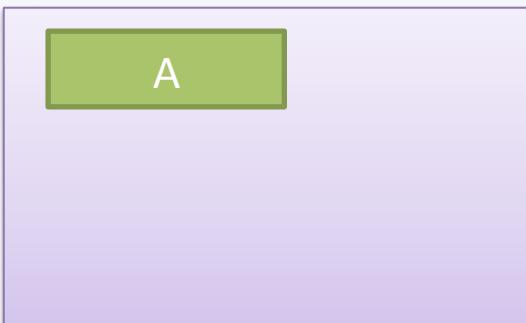
block 2



main memory

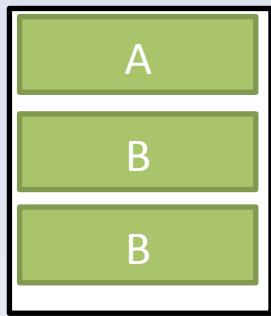


output

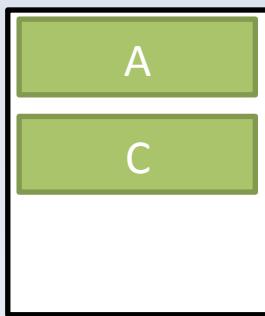


Example: δ

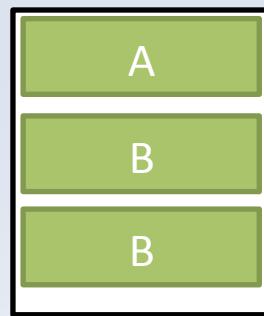
block 0



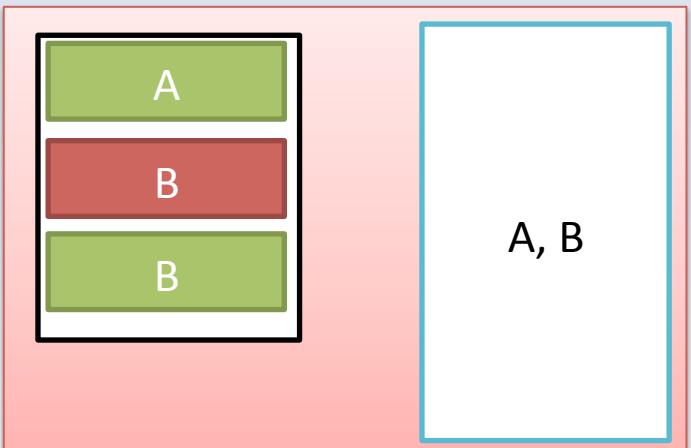
block 1



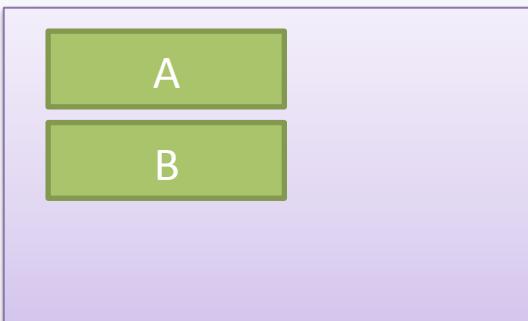
block 2



main memory

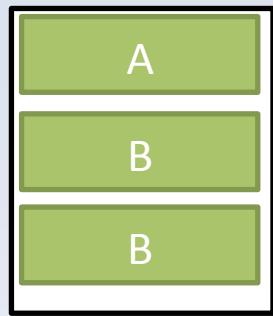


output

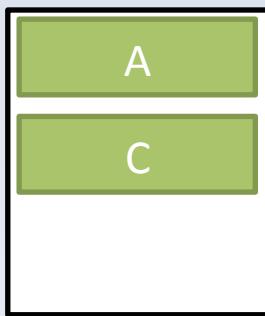


Example: δ

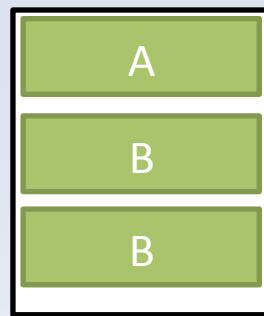
block 0



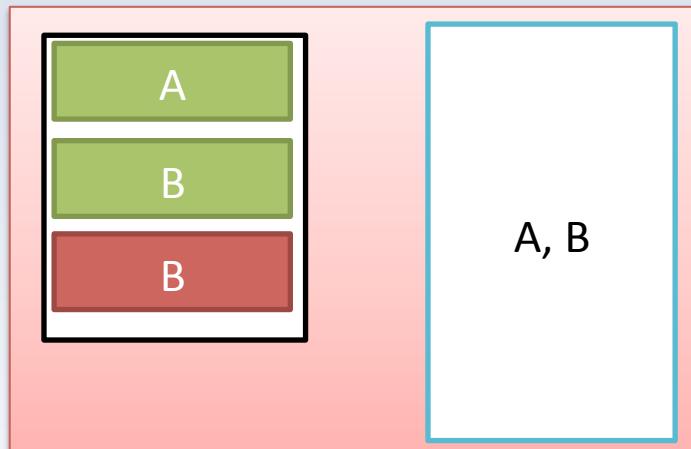
block 1



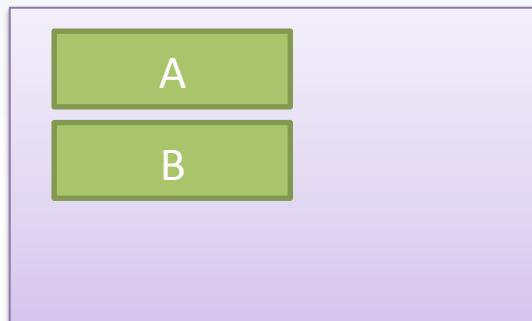
block 2



main memory

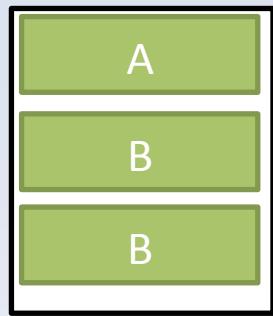


output

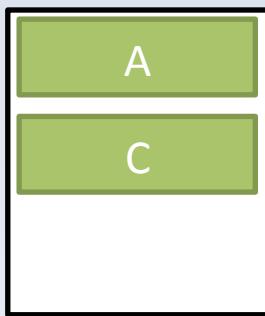


Example: δ

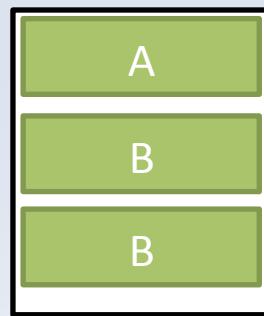
block 0



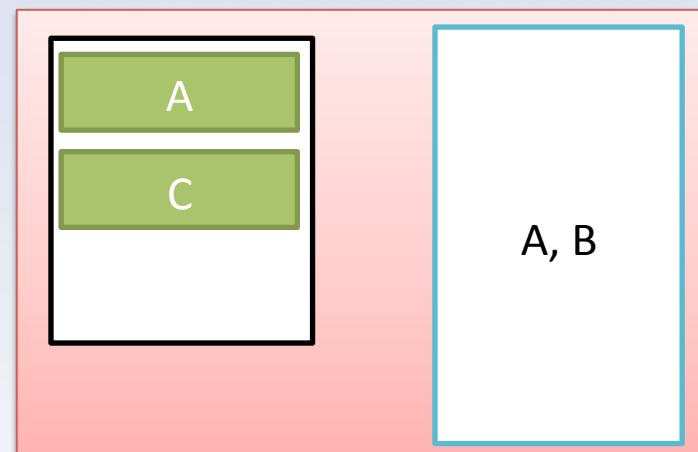
block 1



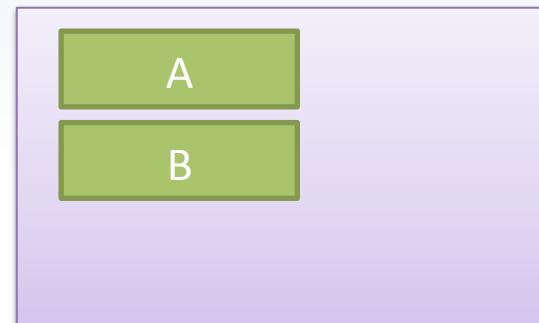
block 2



main memory

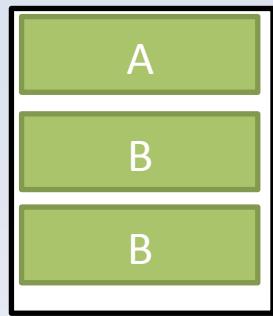


output

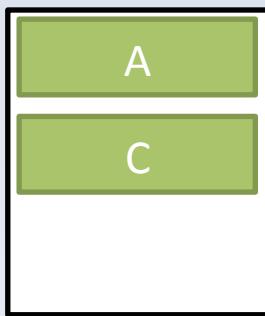


Example: δ

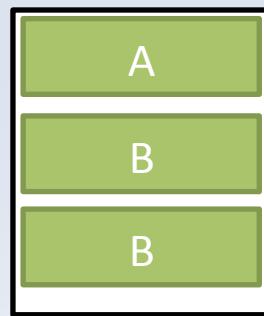
block 0



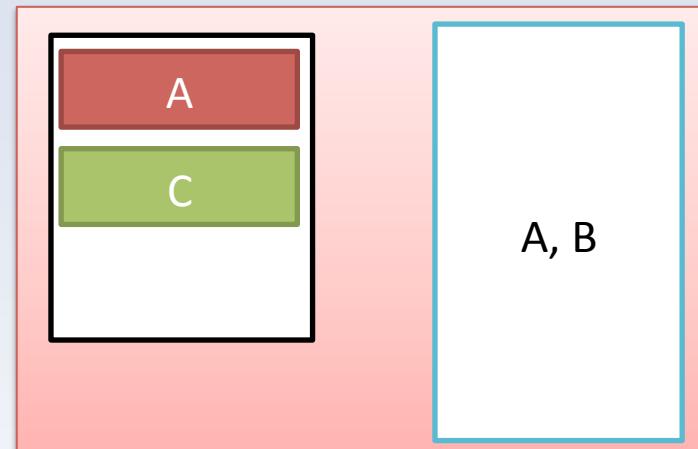
block 1



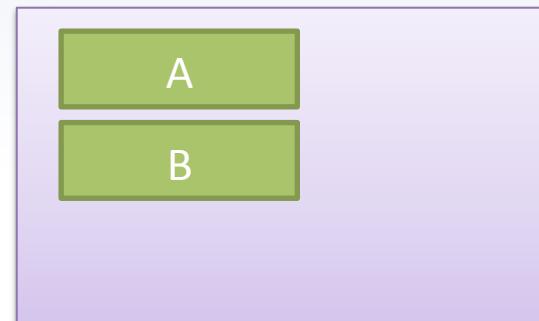
block 2



main memory

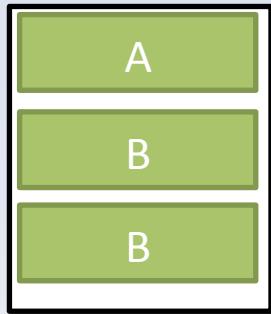


output

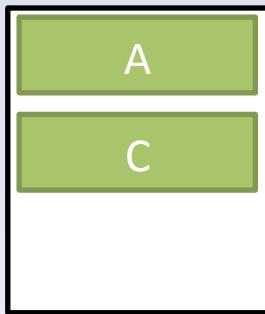


Example: δ

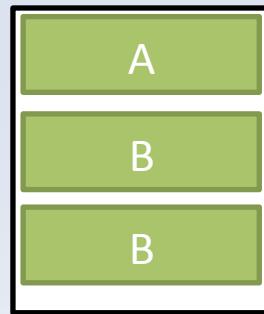
block 0



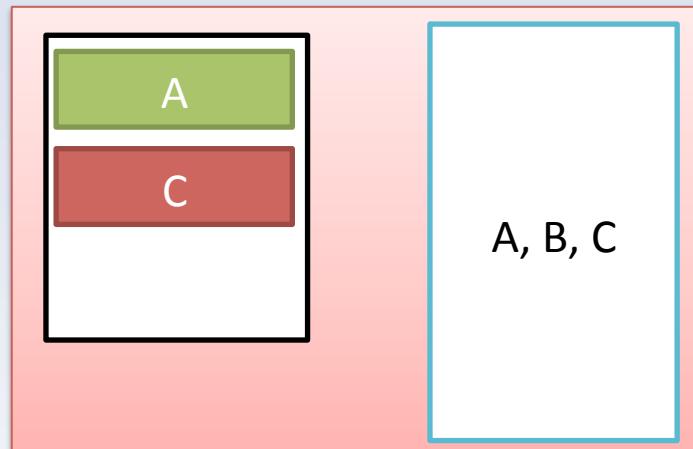
block 1



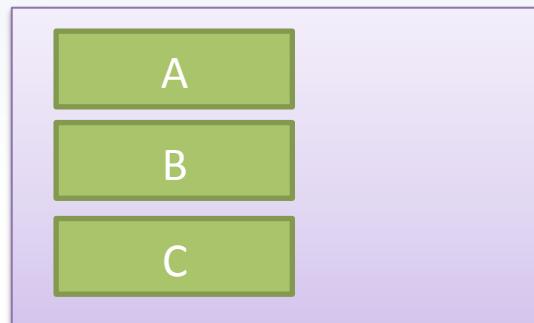
block 2



main memory

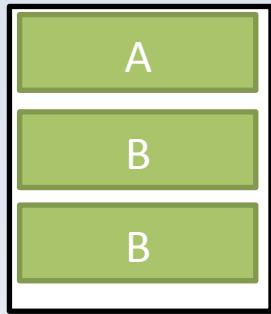


output

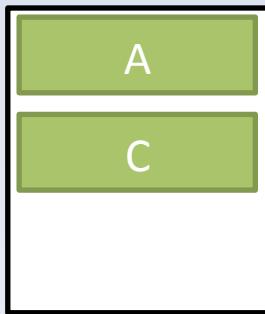


Example: δ

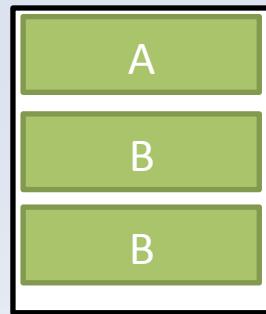
block 0



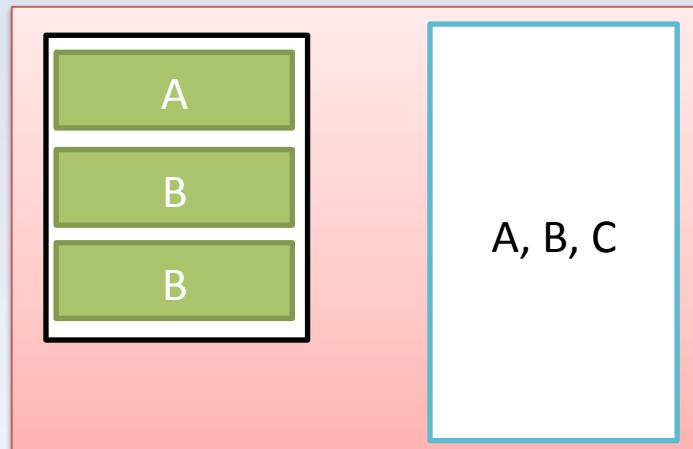
block 1



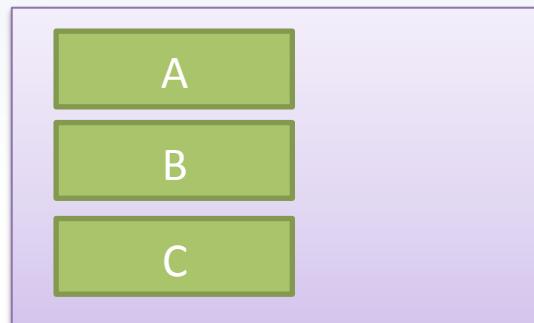
block 2



main memory

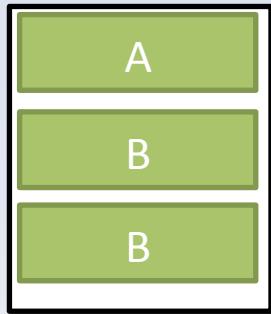


output

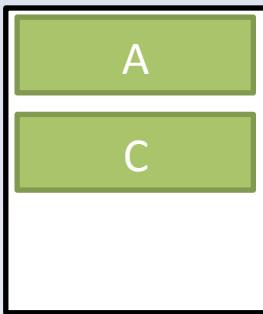


Example: δ

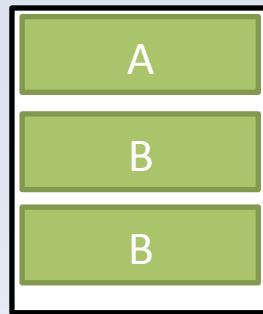
block 0



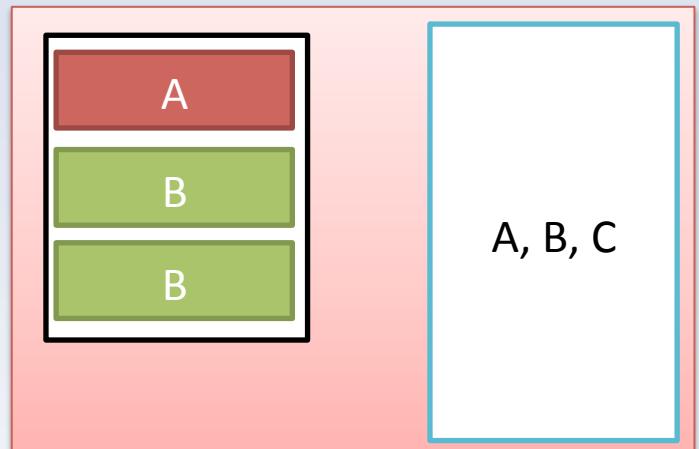
block 1



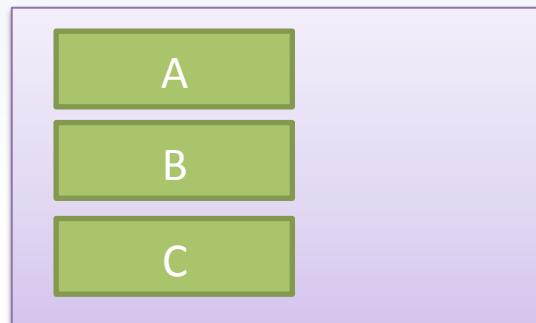
block 2



main memory

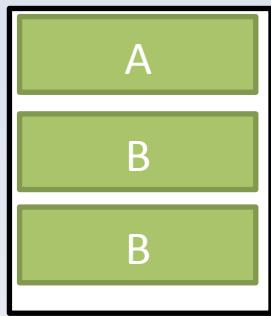


output

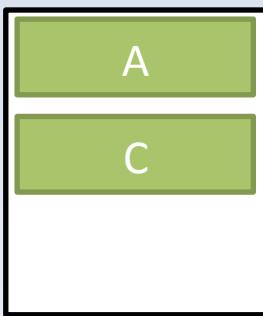


Example: δ

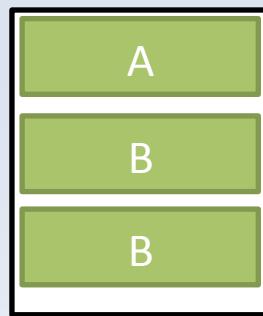
block 0



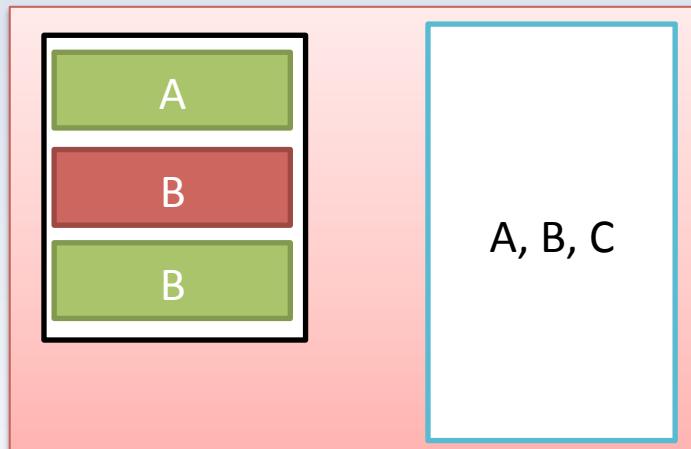
block 1



block 2

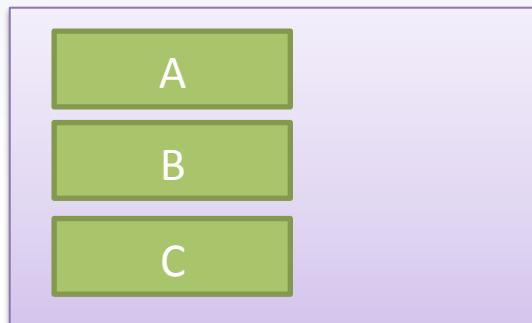


main memory



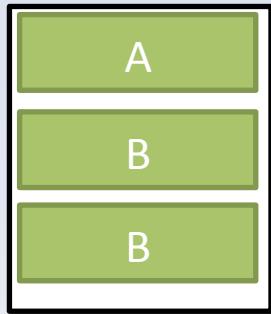
A, B, C

output

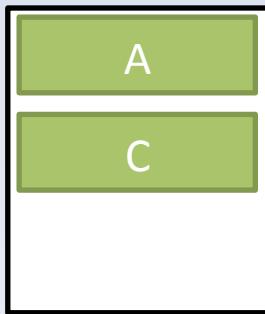


Example: δ

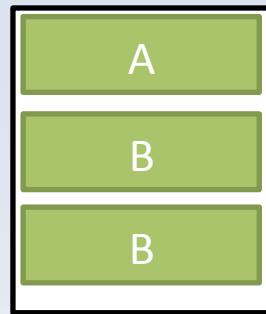
block 0



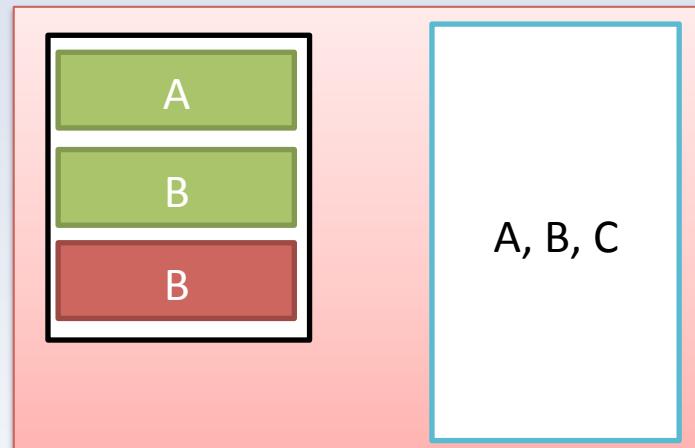
block 1



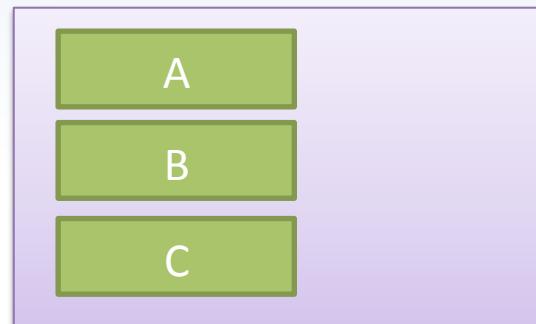
block 2



main memory



output



Binary Operations

- To do \cup , \cap , $-$, \bowtie , \times in one pass:
 - We need one of the relations to fit in memory
 - $\min(B(R), B(S)) < M$
- Assume S is smaller
- For set operations, assume $\delta(R)$ and $\delta(S)$ are already computed



Binary Operations

- \cap
 - read $\delta(S)$ into memory ($M-1$ buffers)
 - build dictionary
 - table scan $\delta(R)$ (1 buffer)
 - looking up each record in dictionary
 - output if entry in dictionary



Binary Operations

- \cup
 - read $\delta(S)$ into memory ($M-1$ buffers)
 - build dictionary and output $\delta(S)$
 - table scan $\delta(R)$ (1 buffer)
 - look up each record in dictionary
 - output if entry *not* in dictionary



Binary Operations

- R-S
 - read $\delta(S)$ into memory (M-1 buffers)
 - build dictionary
 - table scan $\delta(R)$ (1 buffer)
 - look up each record in dictionary
 - output if entry *not* in dictionary



Binary Operations

- S-R
 - read $\delta(S)$ into memory (M-1 buffers)
 - build dictionary
 - table scan $\delta(R)$ (1 buffer)
 - look up each record in dictionary
 - if found, delete entry from dictionary
 - output remaining entries



Binary Operations

- \times
 - read S into memory ($M-1$ buffers)
 - table scan R (1 buffer)
 - for each tuple in S
 - pair it with the current tuple in R
 - output the pairing



Binary Operations

- \bowtie
 - read S into memory ($M-1$ buffers)
 - build dictionary for join attributes
 - table scan R (1 buffer)
 - look up join attributes in dictionary
 - if found, output combined result



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

main memory



block 0

A,1
B,2
D,3

output



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

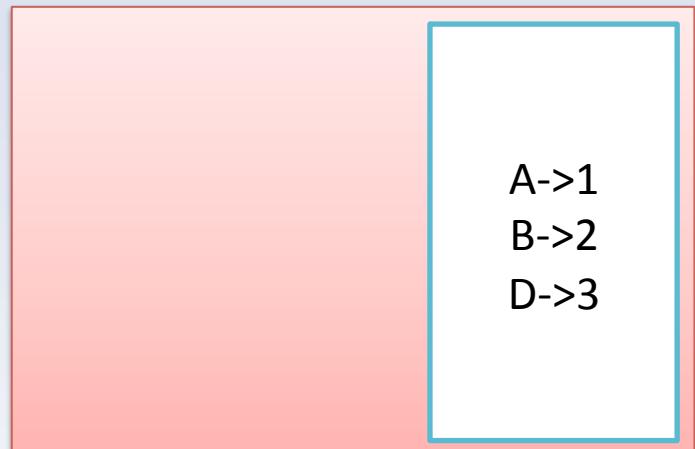
block 1

A,g
C,g

block 2

A,a
B,b
B,c

main memory



block 0

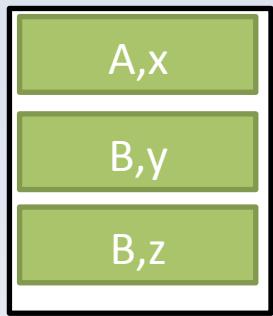
A,1
B,2
D,3

output

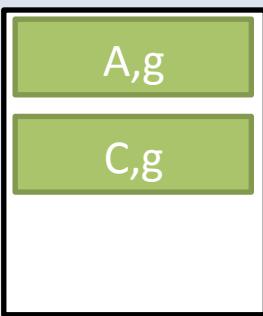


Example: $R \bowtie S$

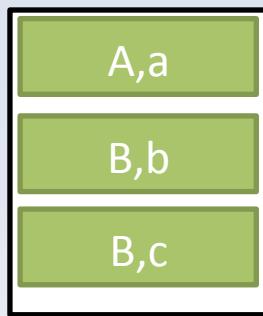
block 0



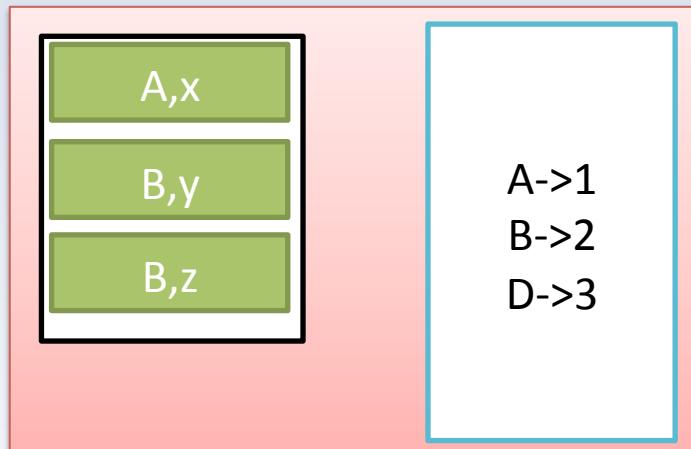
block 1



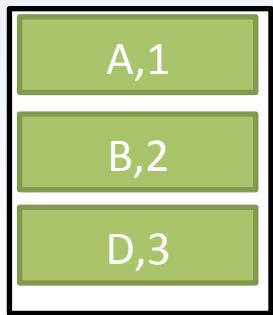
block 2



main memory



block 0

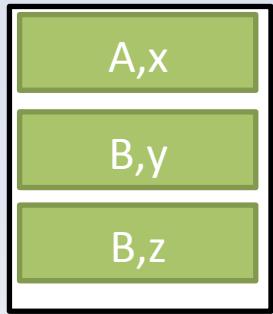


output

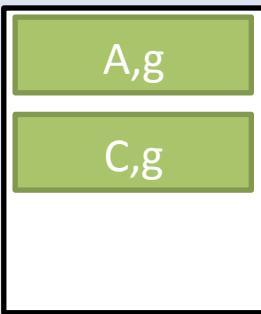


Example: $R \bowtie S$

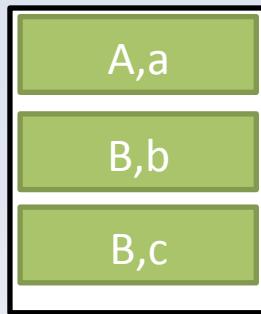
block 0



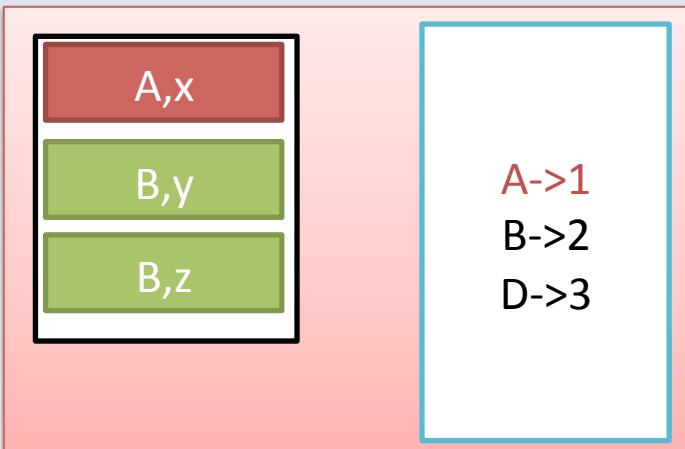
block 1



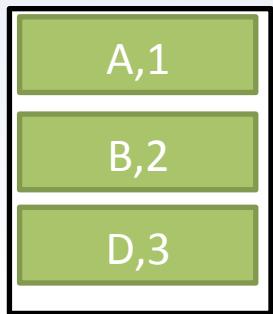
block 2



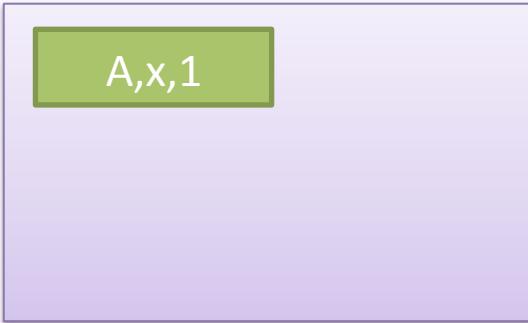
main memory



block 0

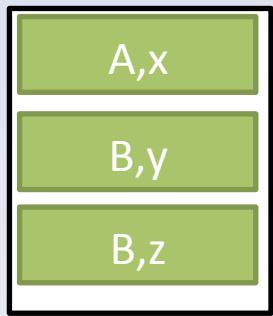


output

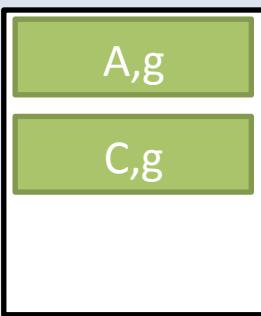


Example: $R \bowtie S$

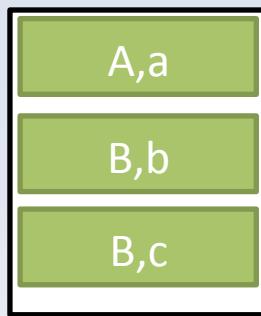
block 0



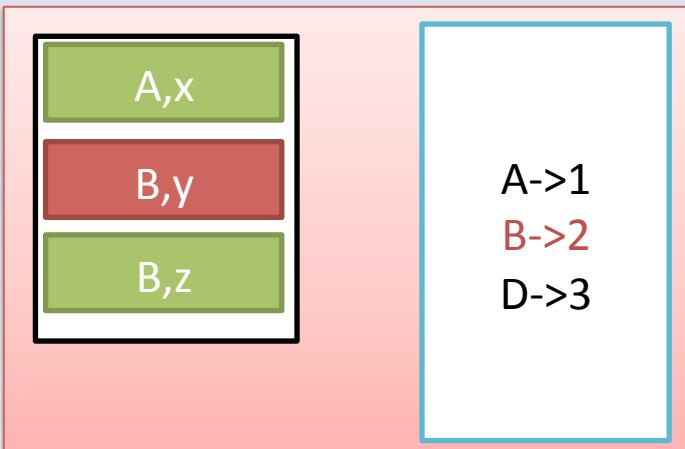
block 1



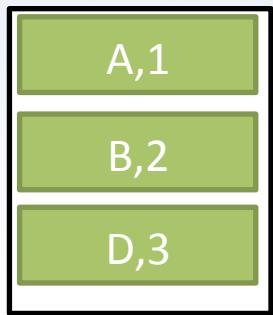
block 2



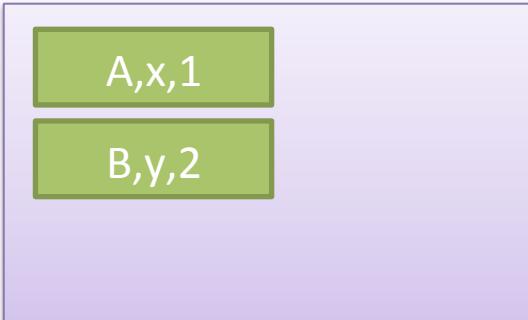
main memory



block 0

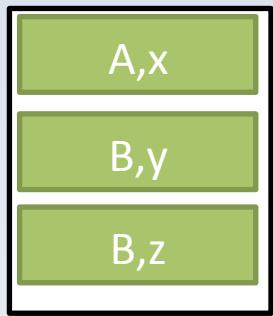


output

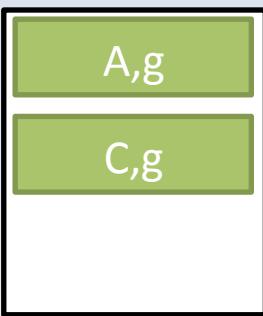


Example: $R \bowtie S$

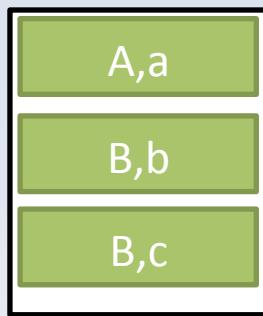
block 0



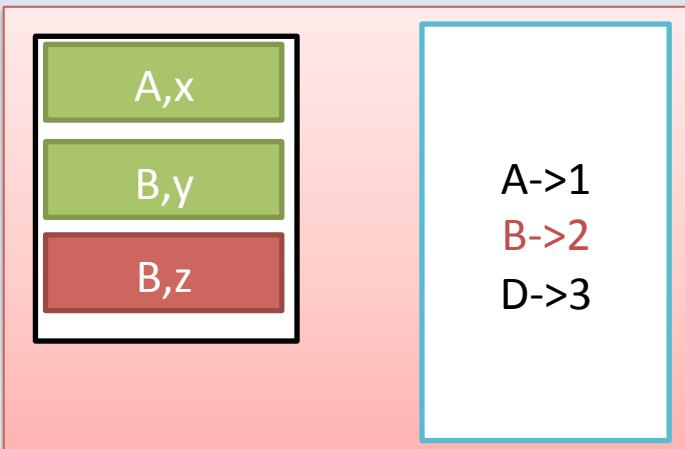
block 1



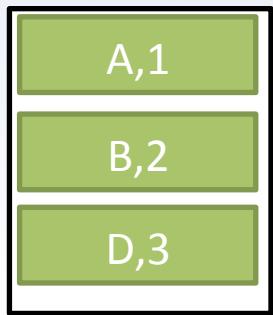
block 2



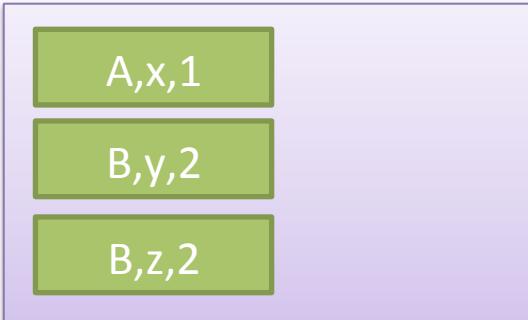
main memory



block 0

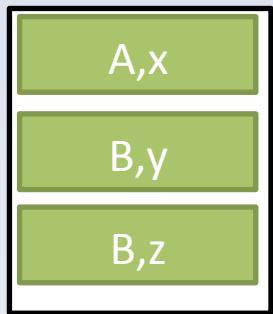


output

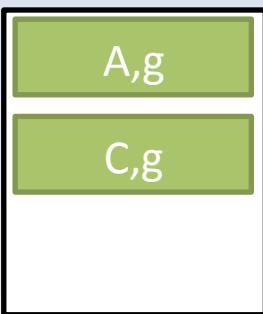


Example: $R \bowtie S$

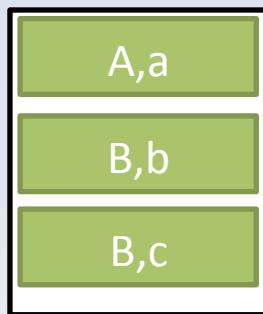
block 0



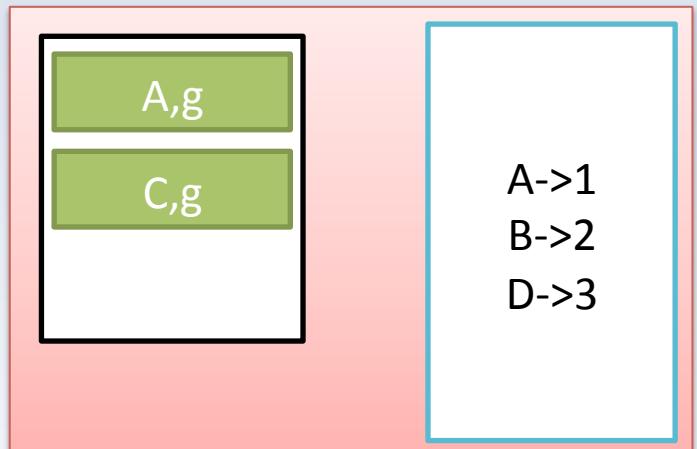
block 1



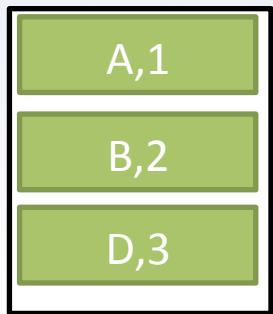
block 2



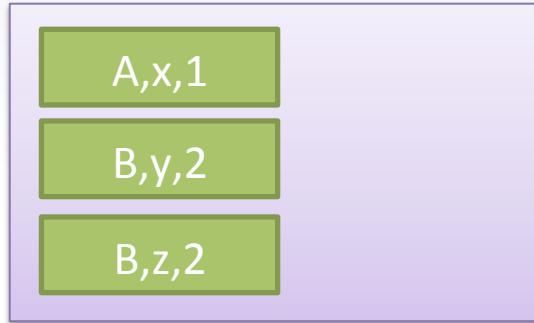
main memory



block 0

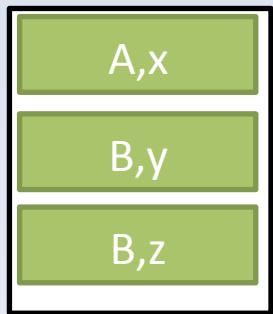


output

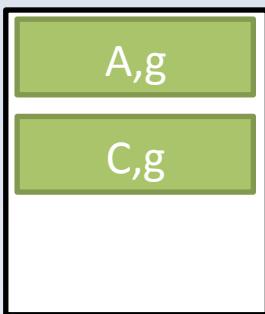


Example: $R \bowtie S$

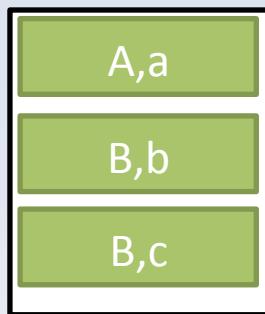
block 0



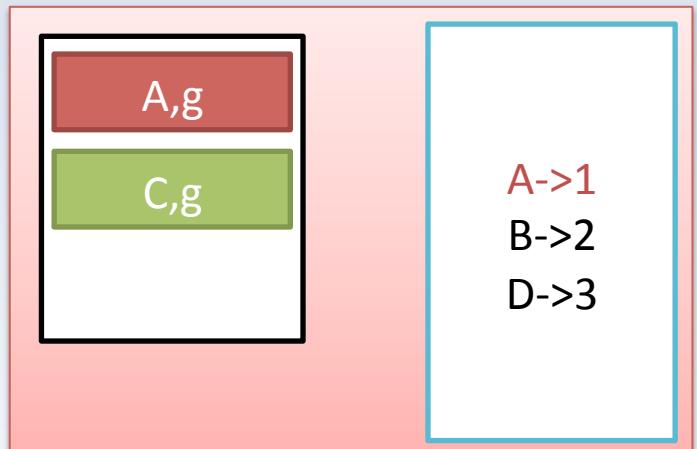
block 1



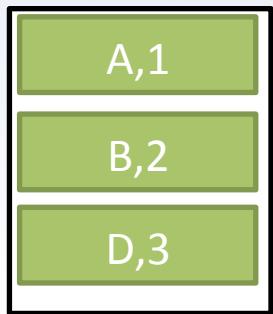
block 2



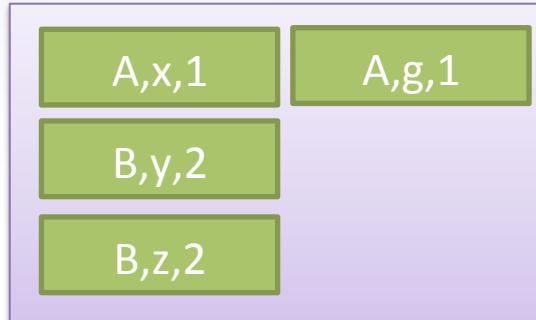
main memory



block 0

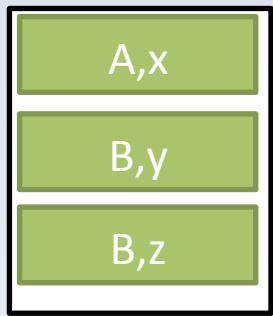


output

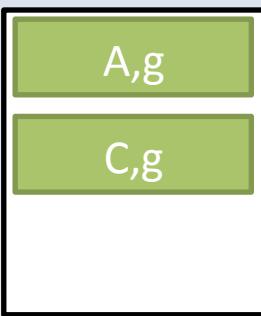


Example: $R \bowtie S$

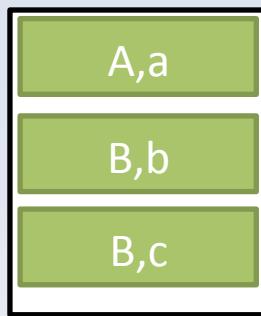
block 0



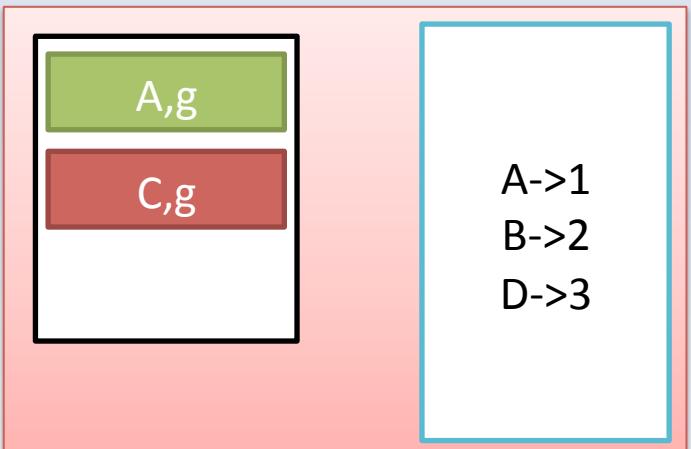
block 1



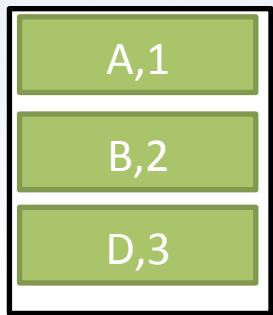
block 2



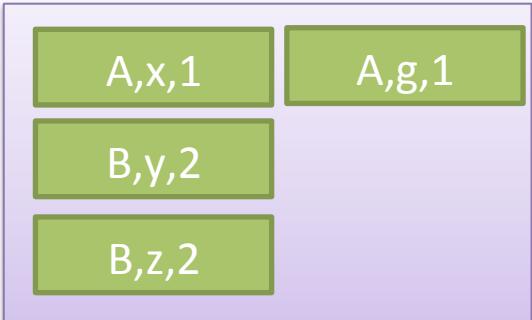
main memory



block 0

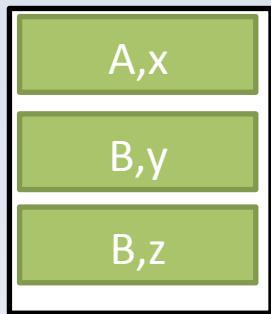


output

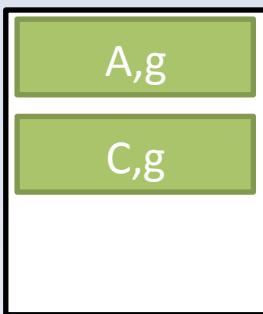


Example: $R \bowtie S$

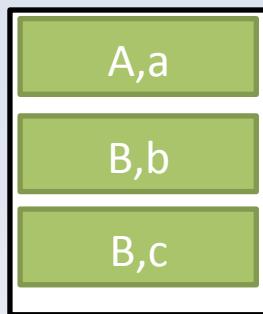
block 0



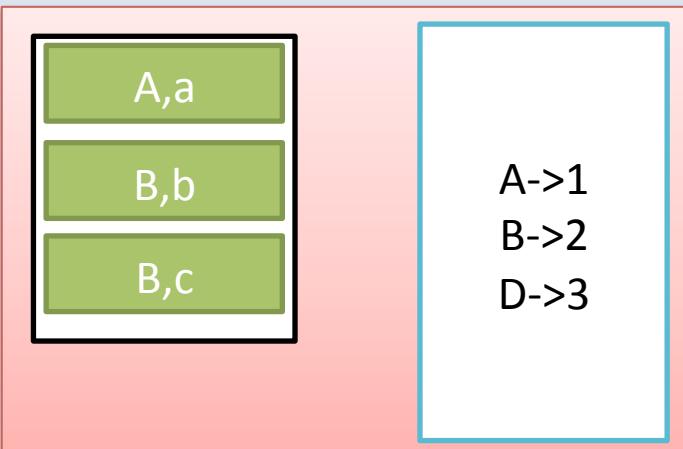
block 1



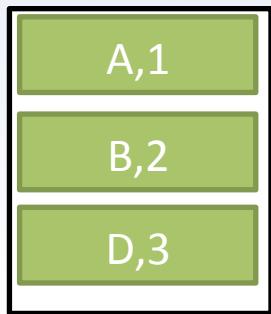
block 2



main memory



block 0

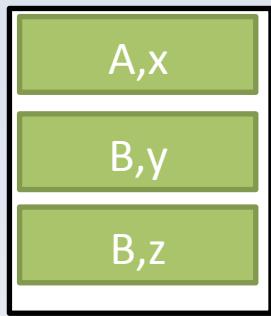


output

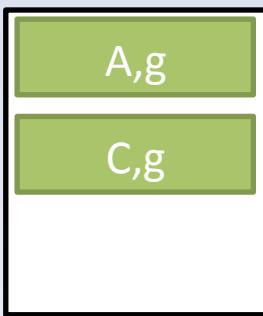


Example: $R \bowtie S$

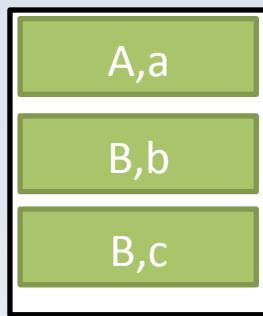
block 0



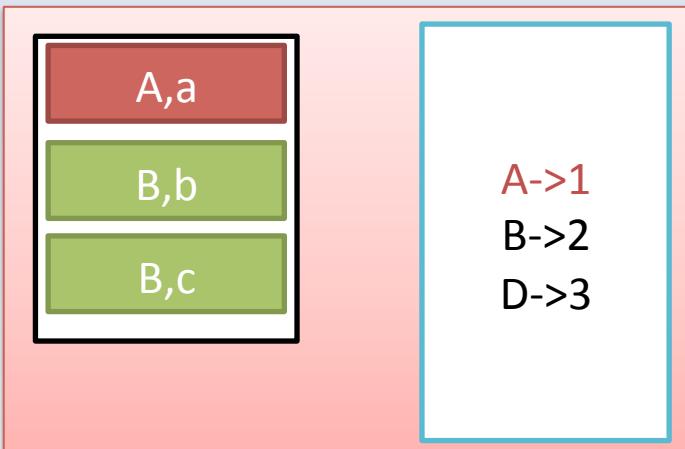
block 1



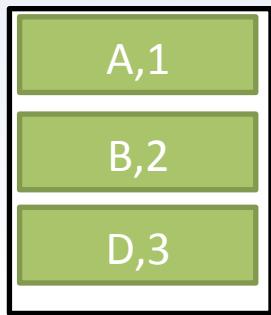
block 2



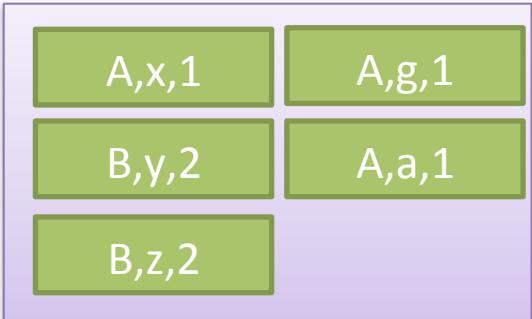
main memory



block 0



output



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

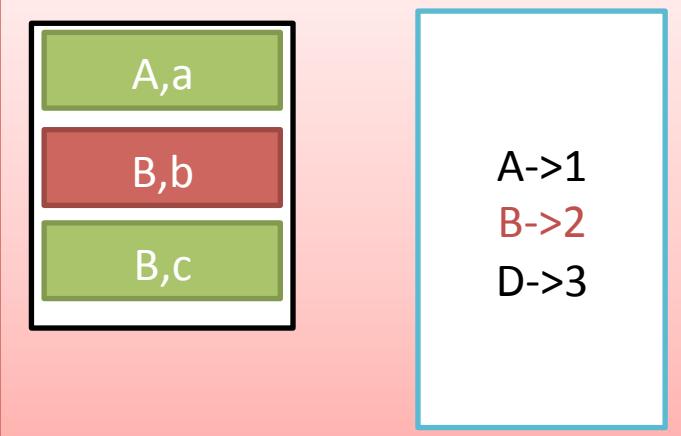
block 1

A,g
C,g

block 2

A,a
B,b
B,c

main memory



block 0

A,1
B,2
D,3

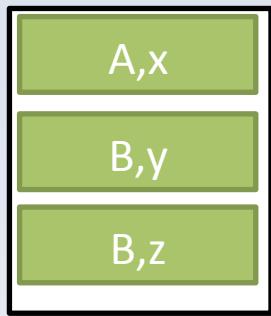
output

A,x,1	A,g,1
B,y,2	A,a,1
B,z,2	B,b,2

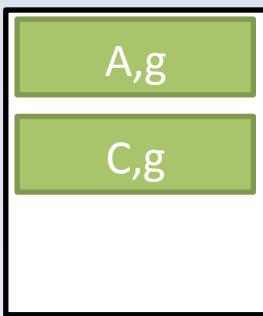


Example: $R \bowtie S$

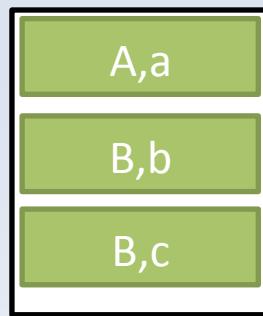
block 0



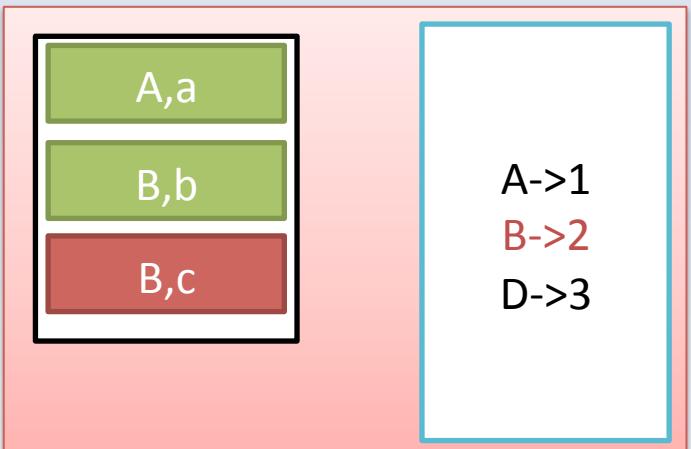
block 1



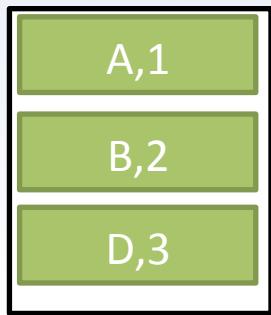
block 2



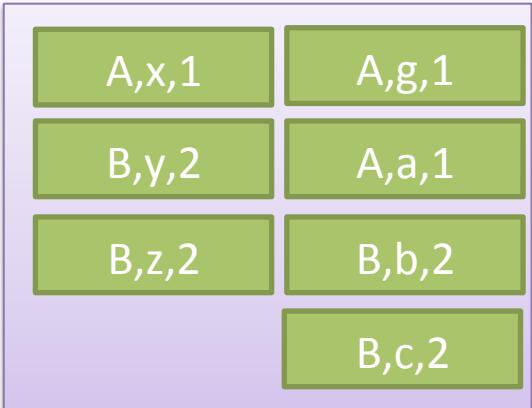
main memory



block 0



output



Binary Single Pass Cost

- Required memory:
 - Need all blocks of R or S to fit into our buffer pool
 - $\min(B(R), B(S))$
- I/O operations
 - Read smaller relation into memory takes $B(S)$
 - Table scan larger relation requires $B(R)$
 - Overall: $B(R) + B(S)$



Nested Loop Joins

- Execute the join with two for loops
- One relation is only read once, other is read multiple times

Pseudocode: tuple based

```
for s in S:  
    for r in R:  
        t=join(r,s)  
        if(t):  
            output t
```



Nested Loop-Join

- That pseudocode ignores the buffer/block model of data
- Nested-loop join is really implemented:
 - Outer loop fills $M-1$ buffers with blocks from S
 - Inner loop fills last buffer with a block from R
- S is read once
- R is read $B(S)/(M-1)$ times



Pseudocode: Block-based

```
for (M-1) blocks bs in S:  
    for block br in R:  
        for s in bs:  
            for r in br:  
                t=join(r,s)  
                if(t):  
                    output t
```



Example: $R \bowtie S$

block 0

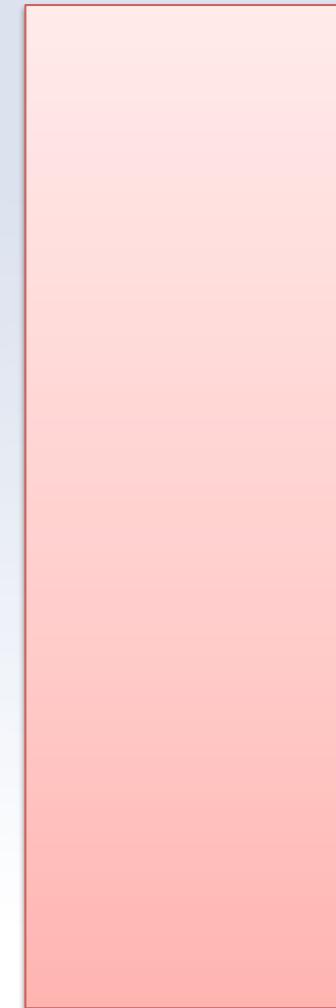
A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c



block 0

A,2
B,4
B,5

block 1

A,1
C,2

block 2

A,7
B,3
B,2

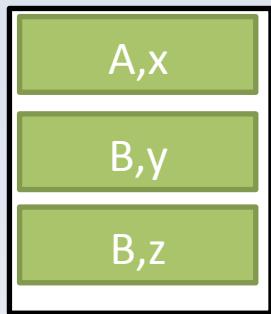
block 3

D,3
C,9
D,1

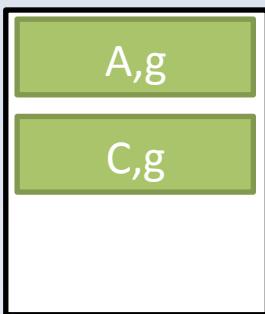


Example: $R \bowtie S$

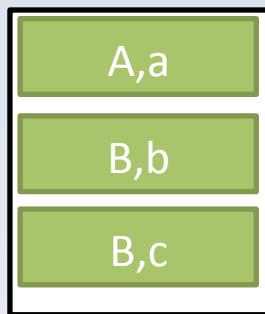
block 0



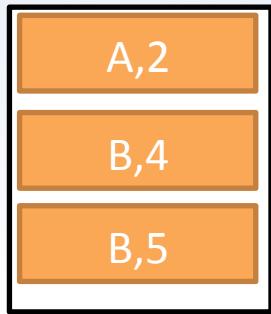
block 1



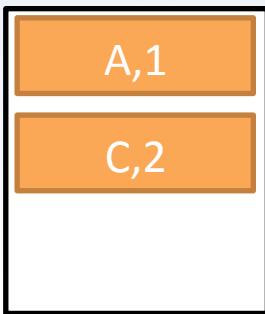
block 2



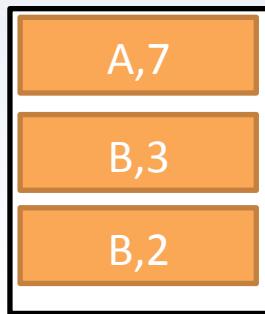
block 0



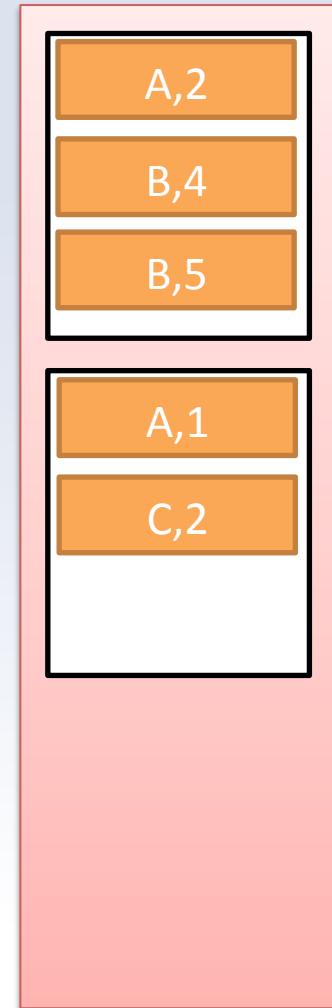
block 1



block 2

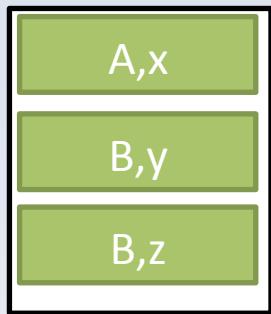


block 3

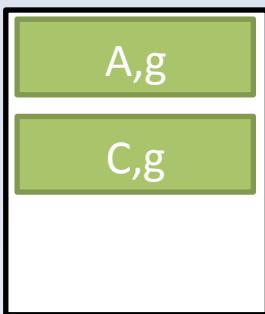


Example: $R \bowtie S$

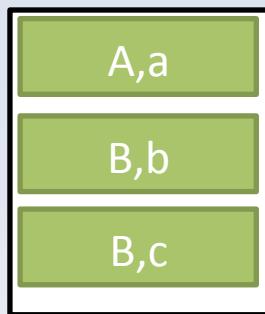
block 0



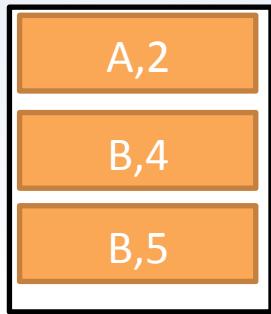
block 1



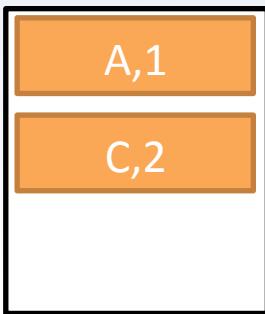
block 2



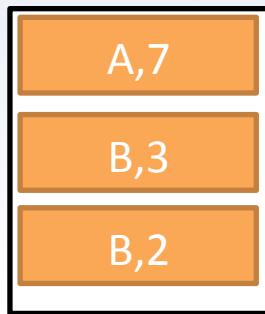
block 0



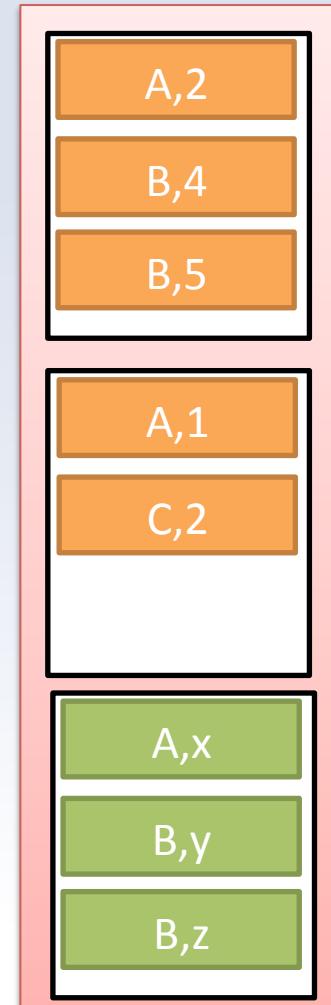
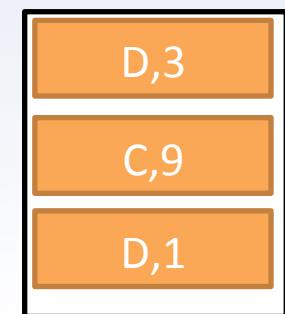
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

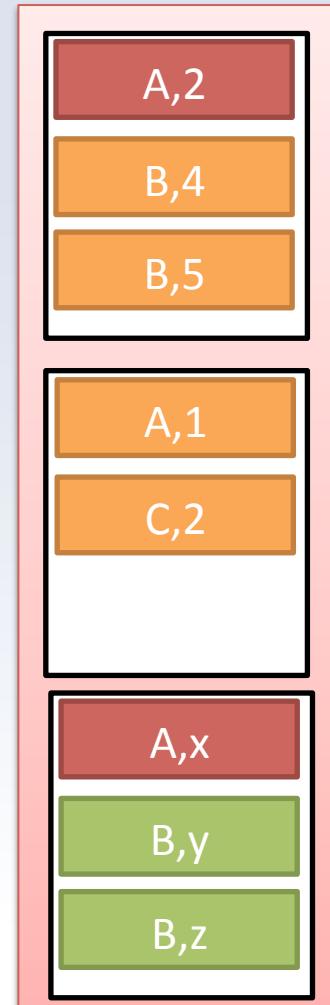
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

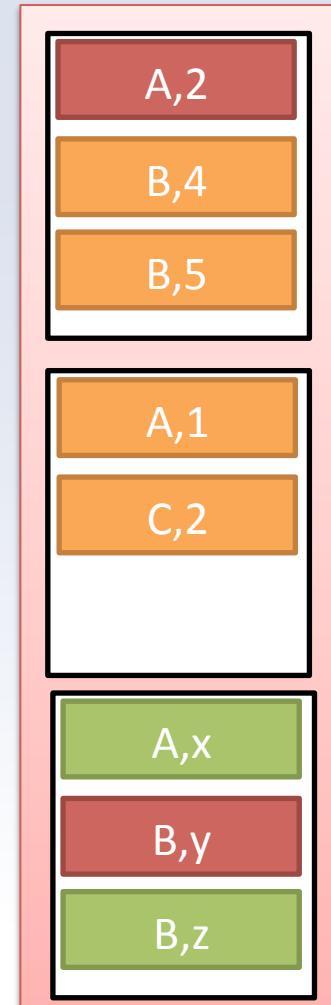
A,1
C,2

block 2

A,7
B,3
B,2

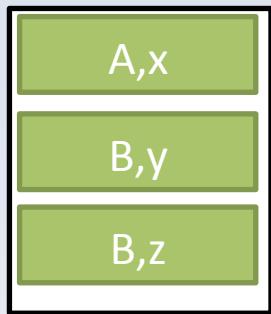
block 3

D,3
C,9
D,1

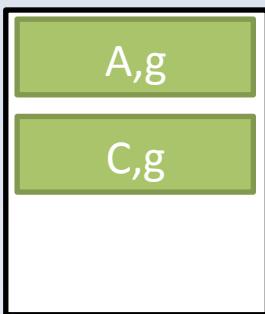


Example: $R \bowtie S$

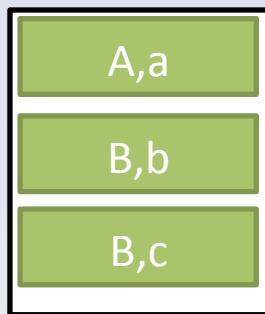
block 0



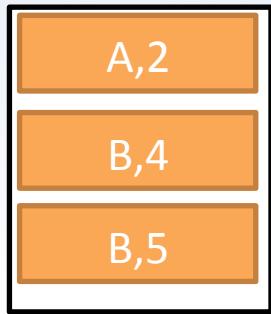
block 1



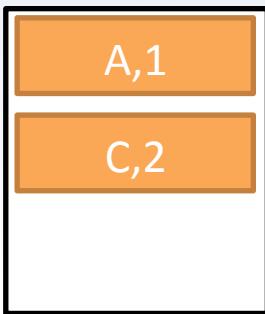
block 2



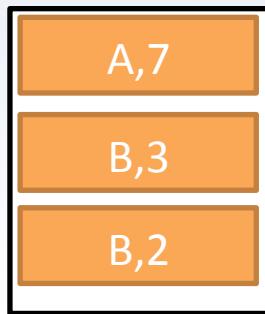
block 0



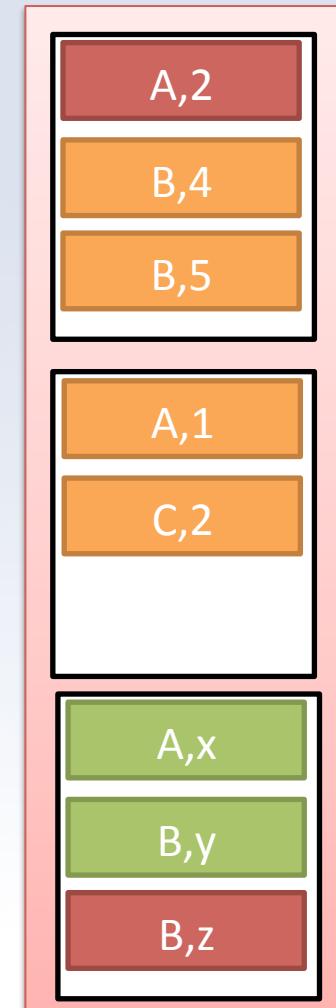
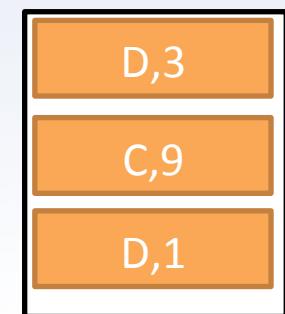
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

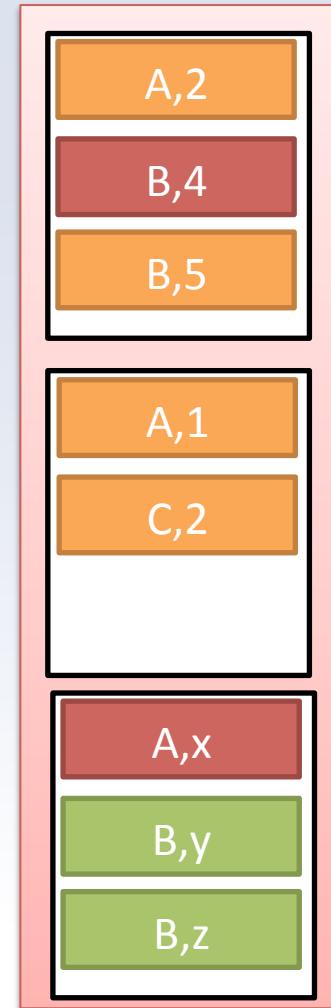
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

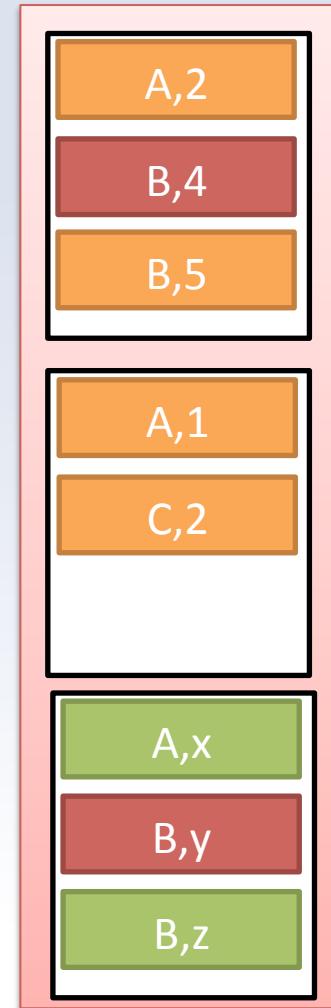
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

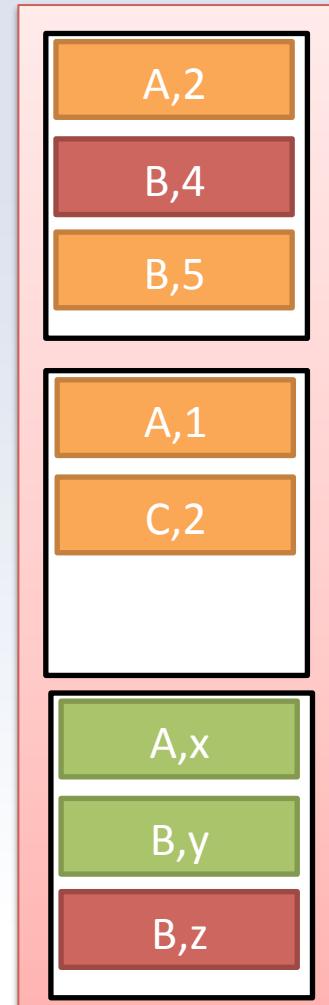
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

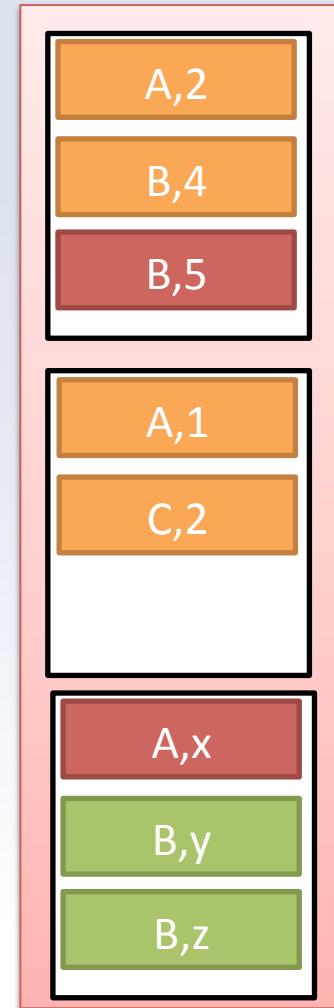
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

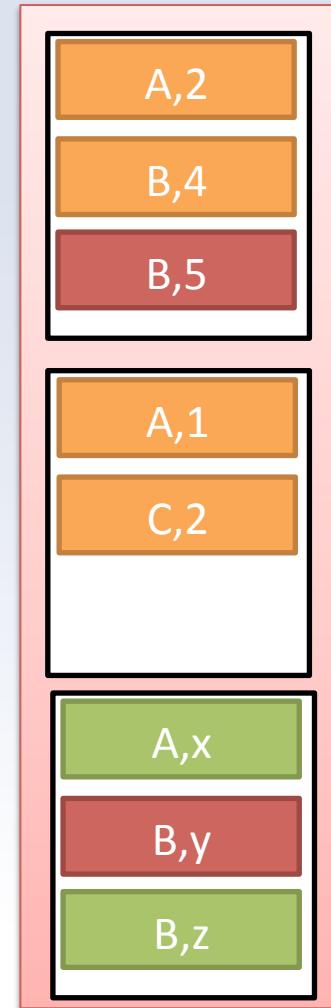
A,1
C,2

block 2

A,7
B,3
B,2

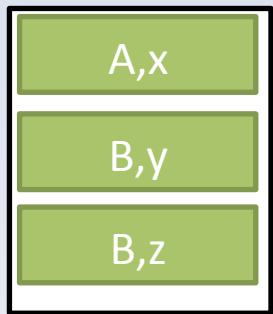
block 3

D,3
C,9
D,1

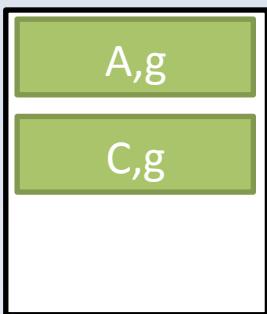


Example: $R \bowtie S$

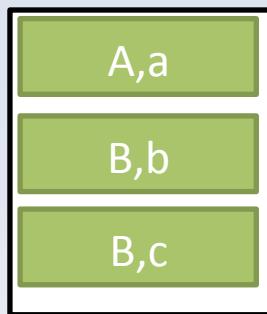
block 0



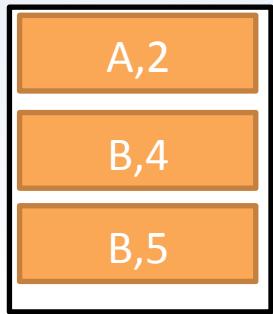
block 1



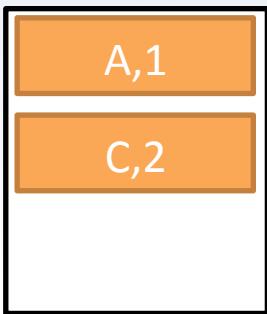
block 2



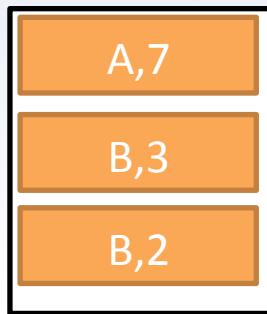
block 0



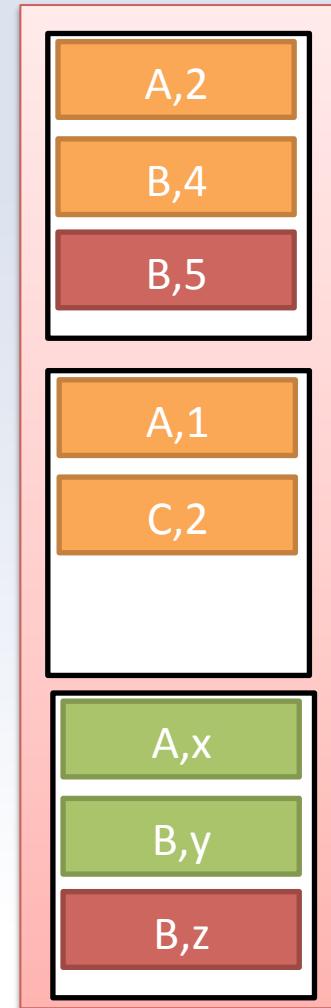
block 1



block 2

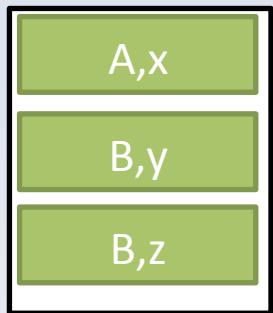


block 3

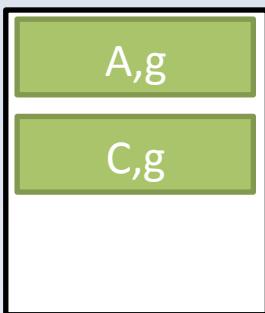


Example: $R \bowtie S$

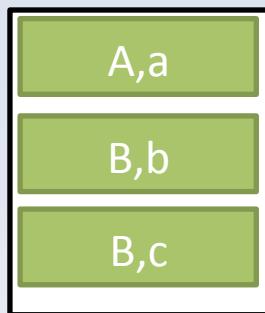
block 0



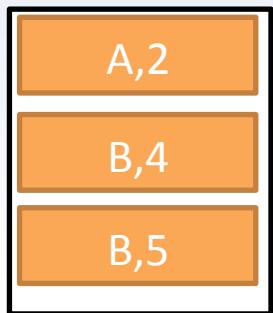
block 1



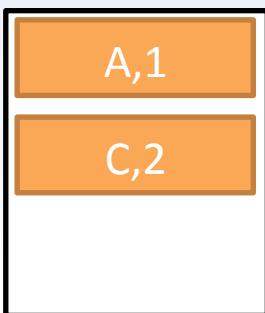
block 2



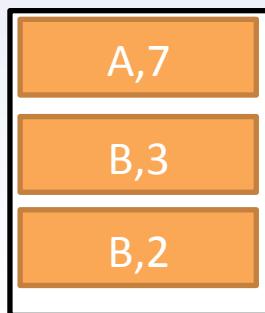
block 0



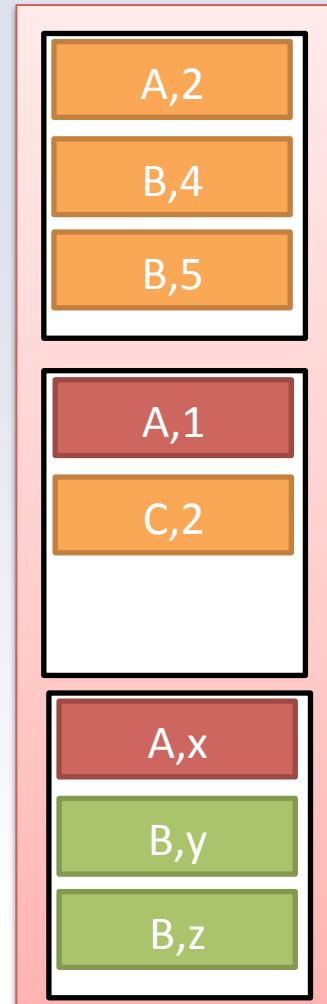
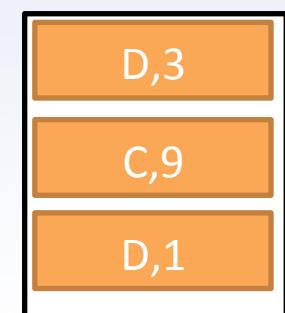
block 1



block 2

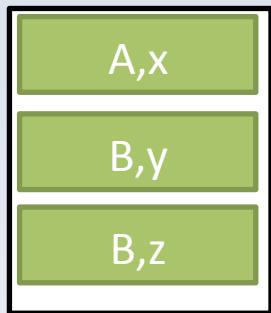


block 3

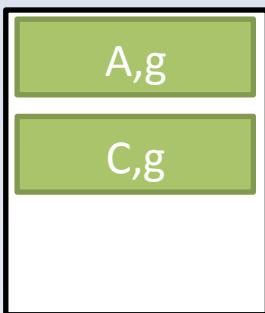


Example: $R \bowtie S$

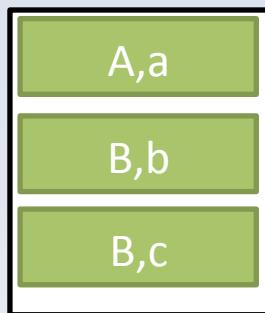
block 0



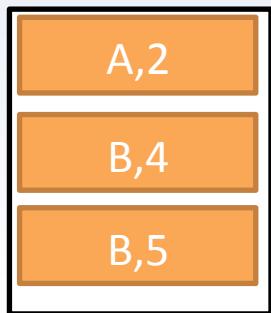
block 1



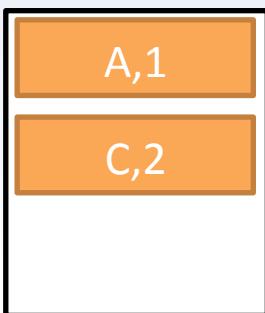
block 2



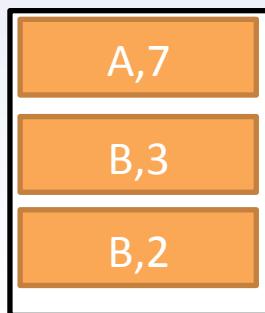
block 0



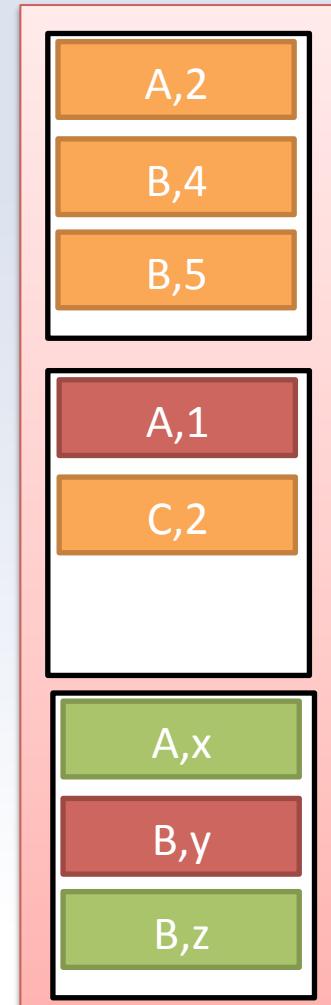
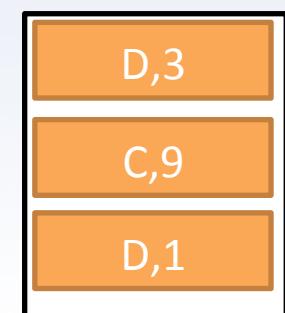
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

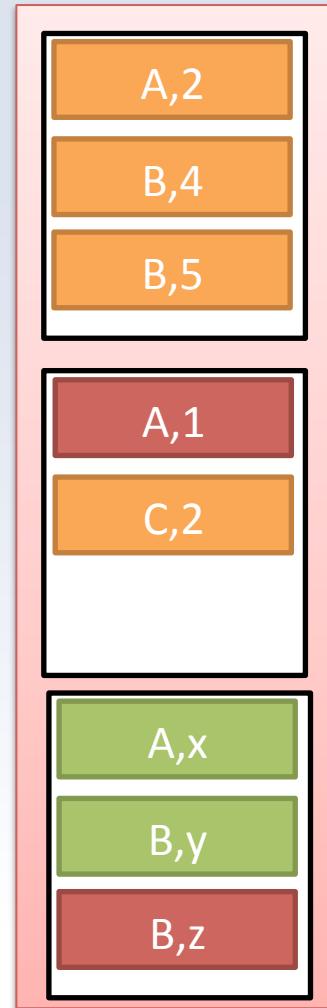
A,1
C,2

block 2

A,7
B,3
B,2

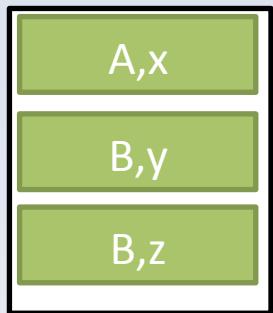
block 3

D,3
C,9
D,1

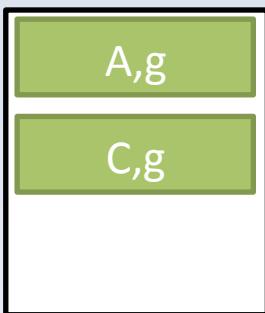


Example: $R \bowtie S$

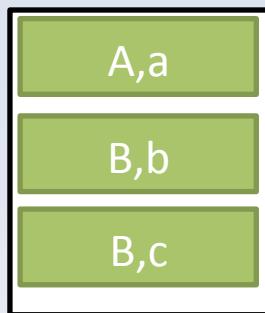
block 0



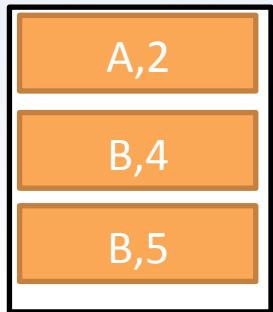
block 1



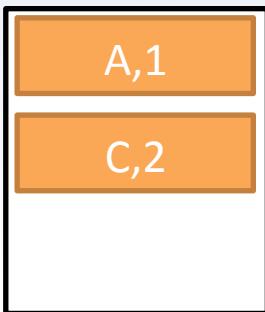
block 2



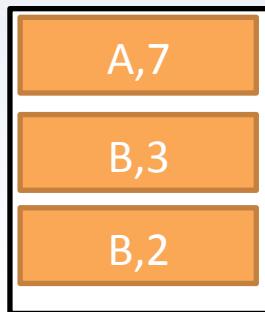
block 0



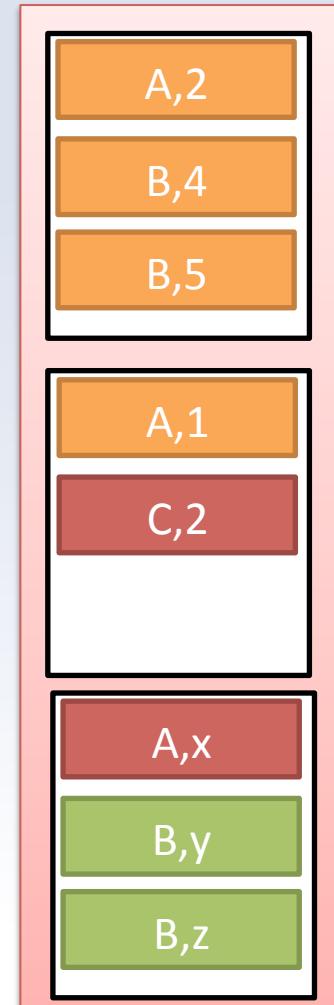
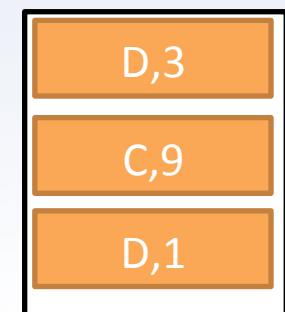
block 1



block 2

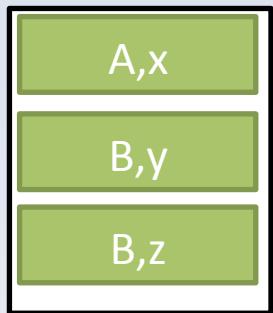


block 3

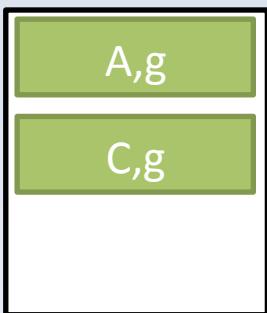


Example: $R \bowtie S$

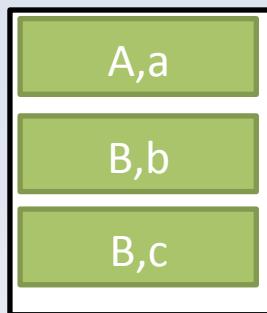
block 0



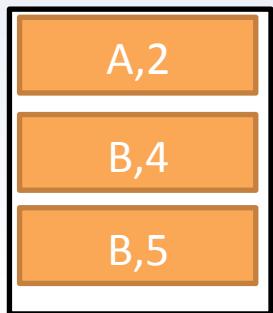
block 1



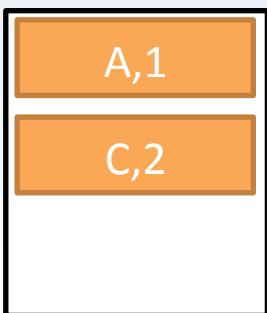
block 2



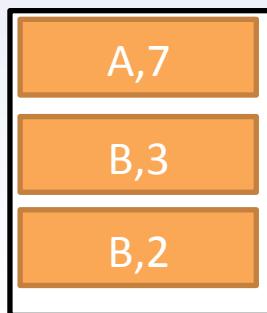
block 0



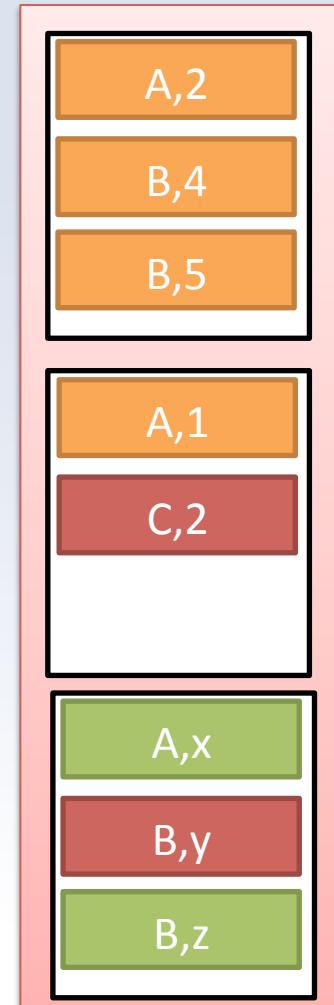
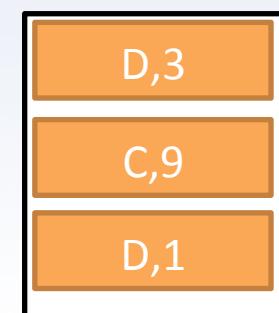
block 1



block 2

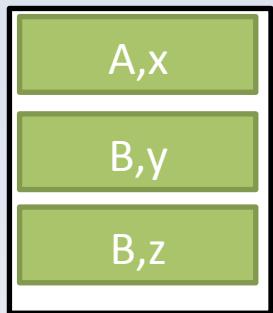


block 3

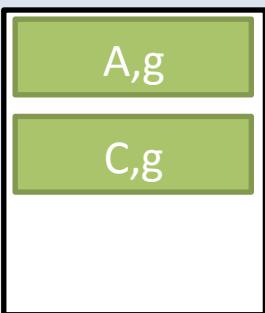


Example: $R \bowtie S$

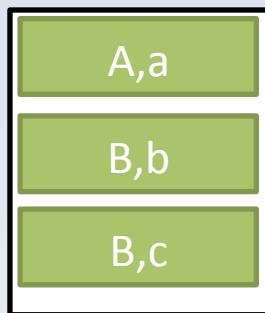
block 0



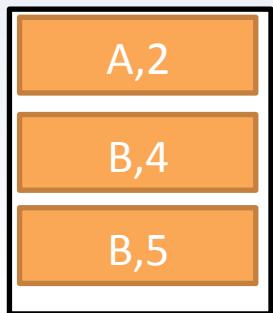
block 1



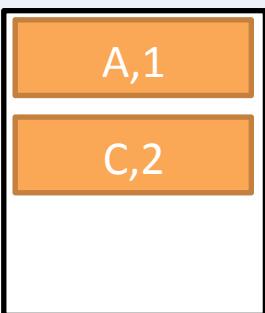
block 2



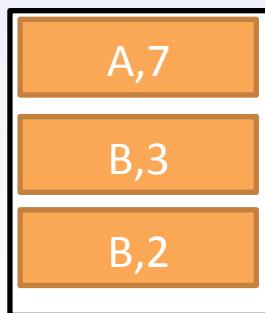
block 0



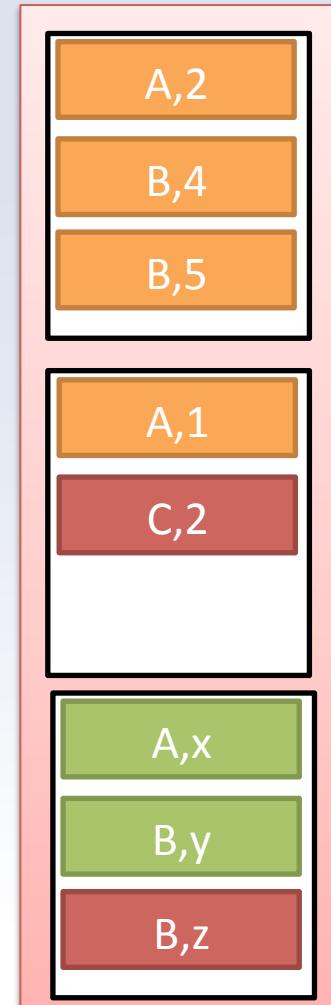
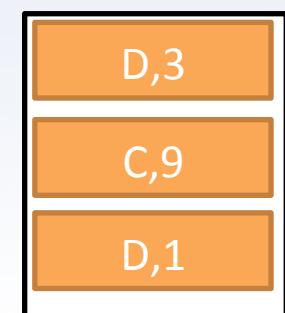
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

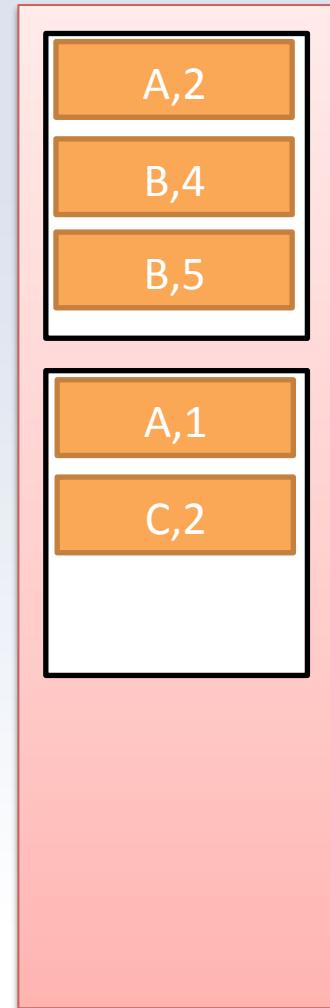
A,1
C,2

block 2

A,7
B,3
B,2

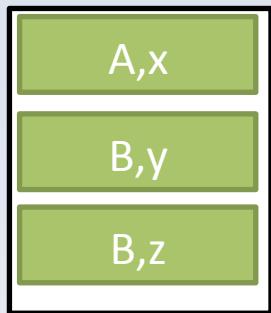
block 3

D,3
C,9
D,1

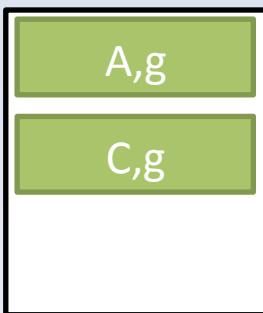


Example: $R \bowtie S$

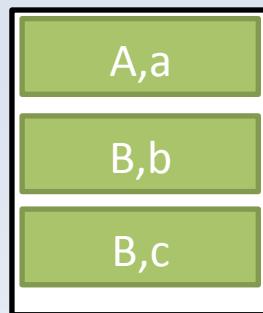
block 0



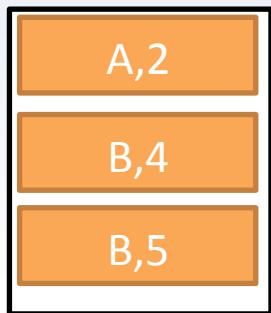
block 1



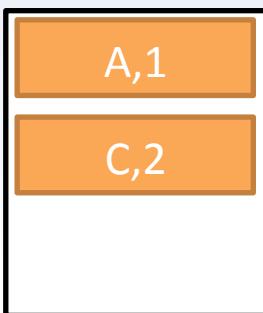
block 2



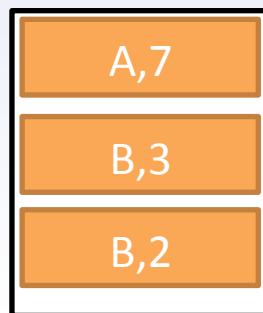
block 0



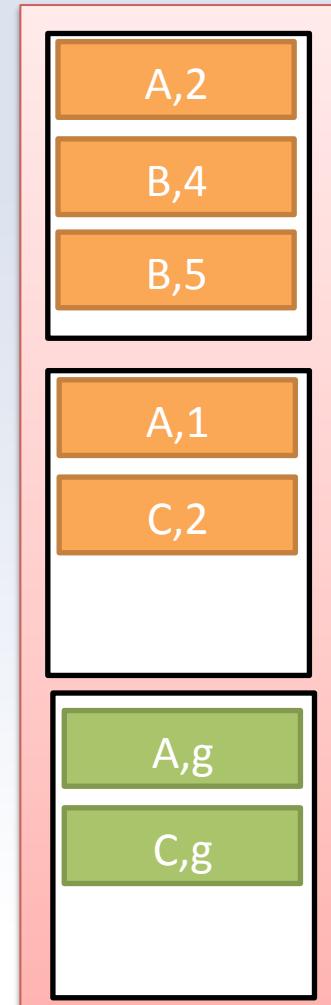
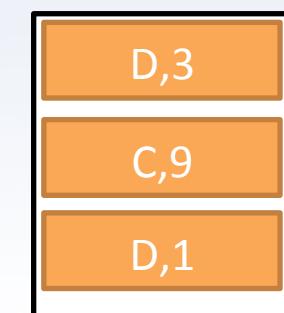
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

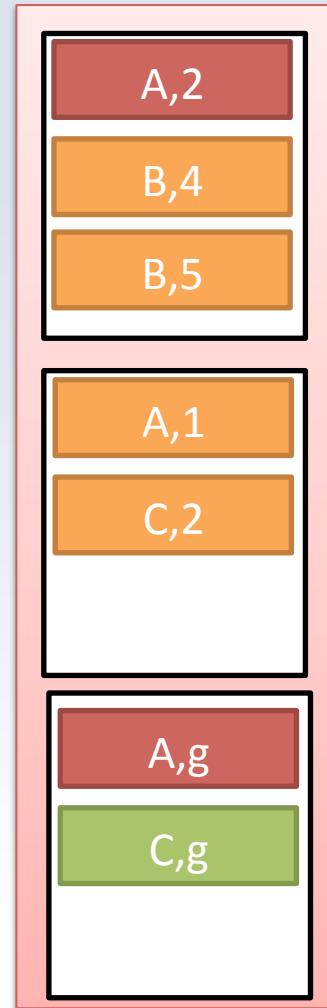
A,1
C,2

block 2

A,7
B,3
B,2

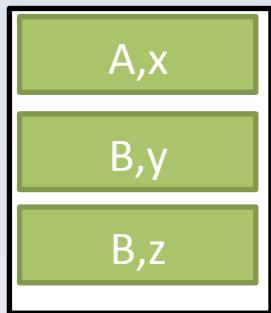
block 3

D,3
C,9
D,1

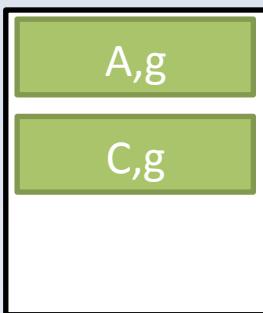


Example: $R \bowtie S$

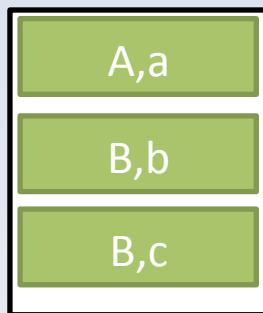
block 0



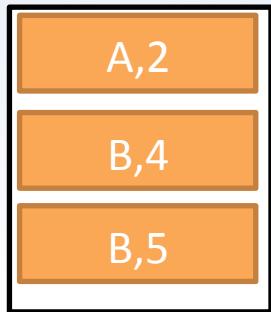
block 1



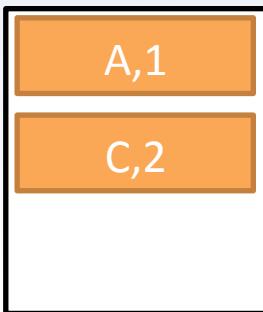
block 2



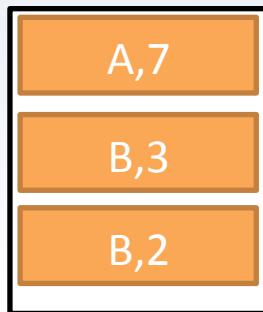
block 0



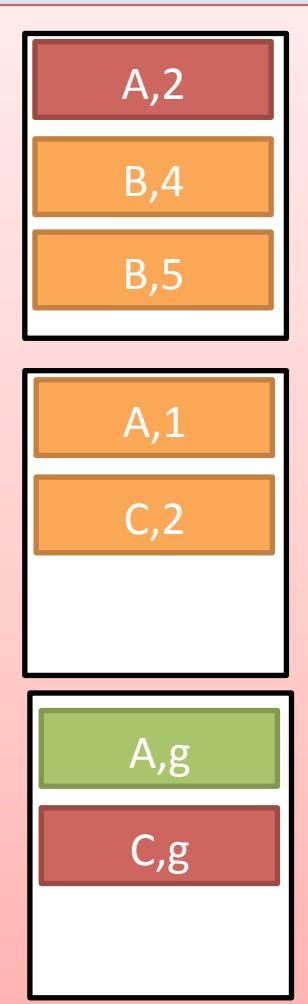
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

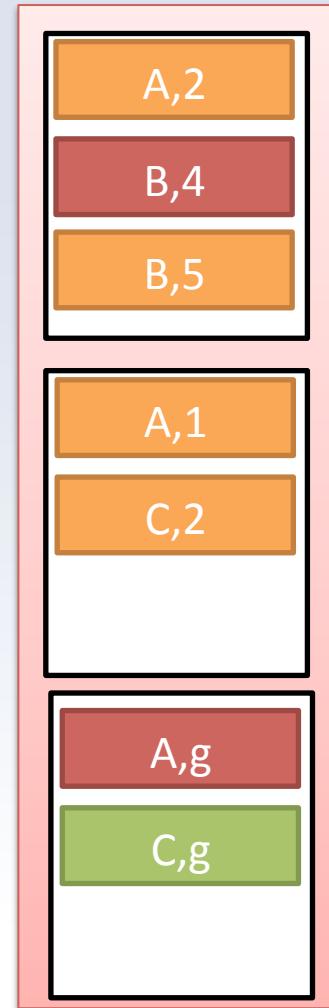
A,1
C,2

block 2

A,7
B,3
B,2

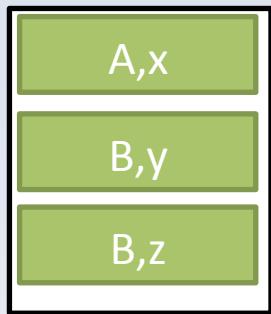
block 3

D,3
C,9
D,1

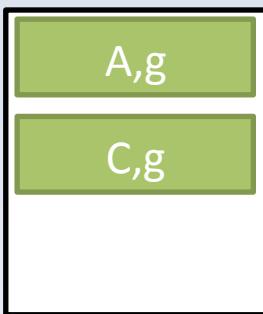


Example: $R \bowtie S$

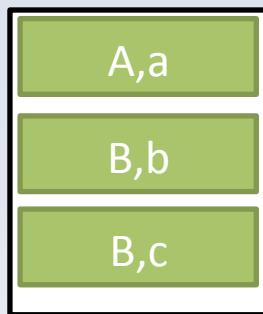
block 0



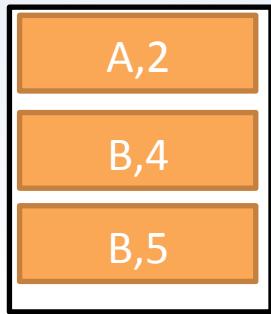
block 1



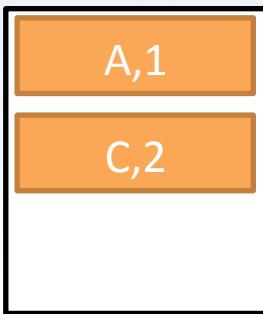
block 2



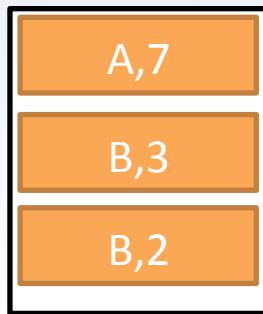
block 0



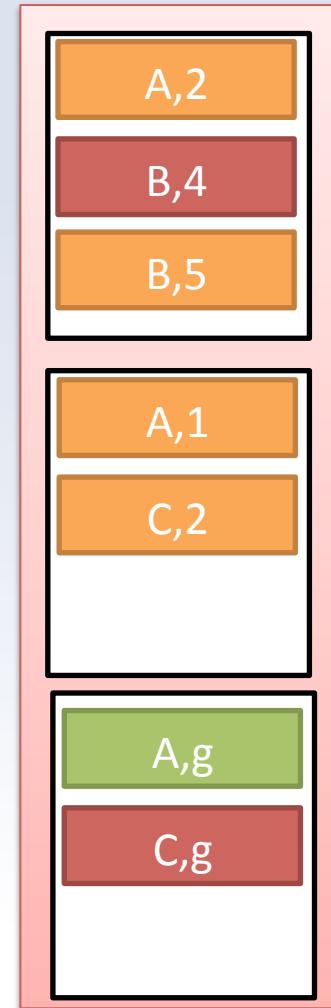
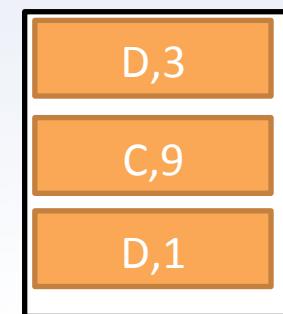
block 1



block 2

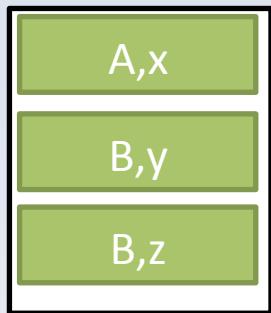


block 3

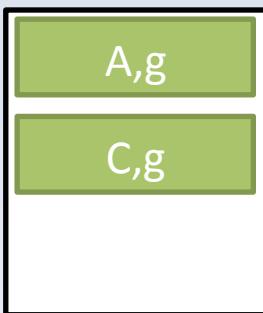


Example: $R \bowtie S$

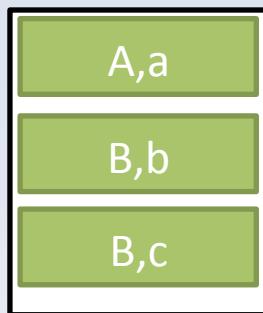
block 0



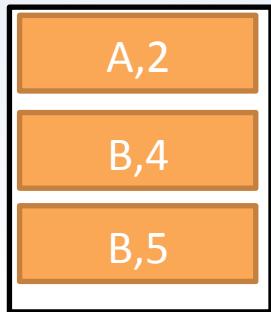
block 1



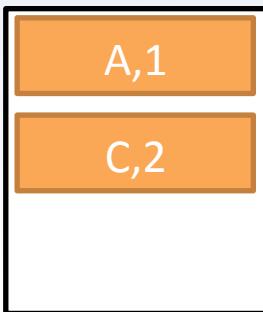
block 2



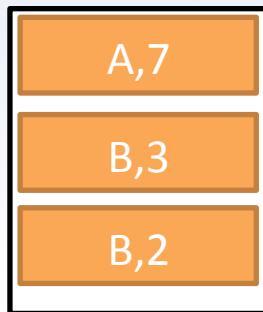
block 0



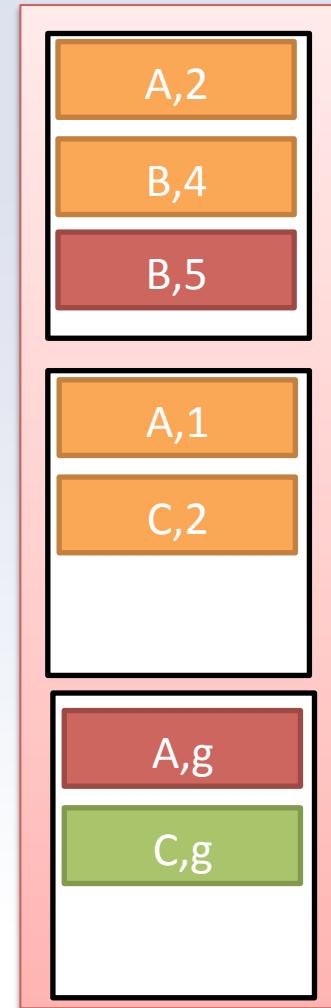
block 1



block 2

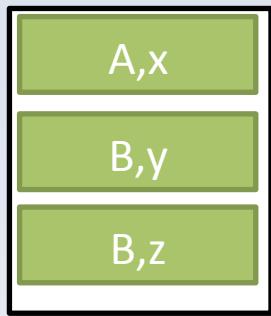


block 3

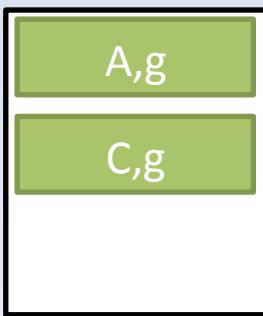


Example: $R \bowtie S$

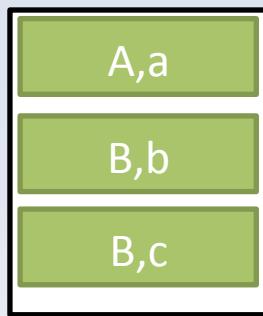
block 0



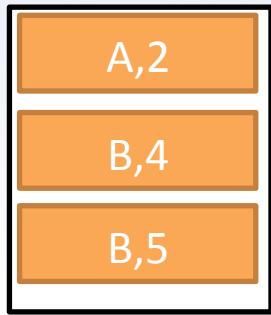
block 1



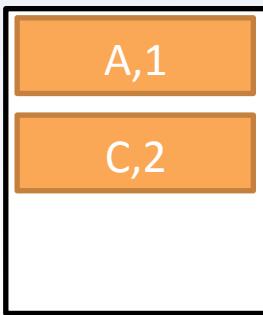
block 2



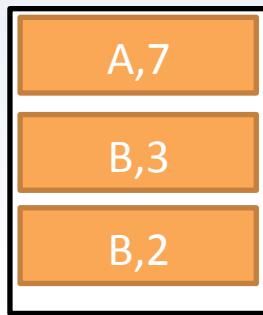
block 0



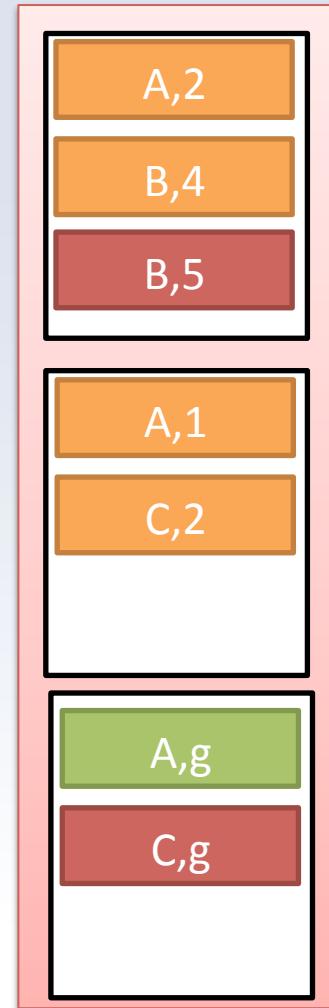
block 1



block 2

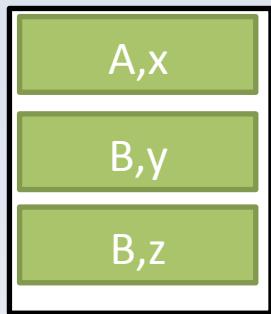


block 3

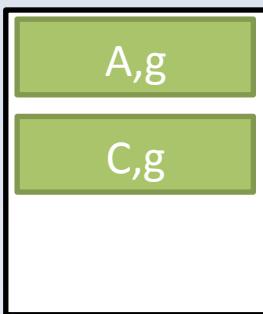


Example: $R \bowtie S$

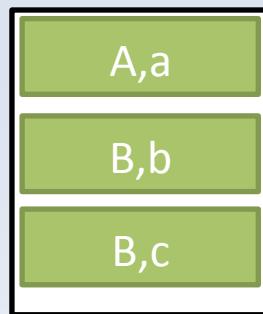
block 0



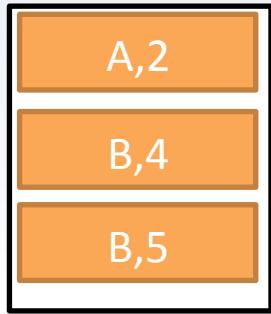
block 1



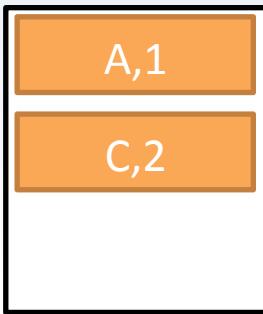
block 2



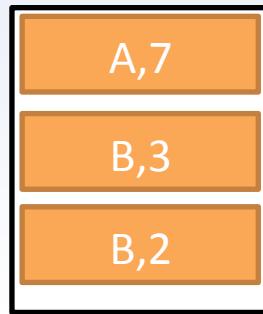
block 0



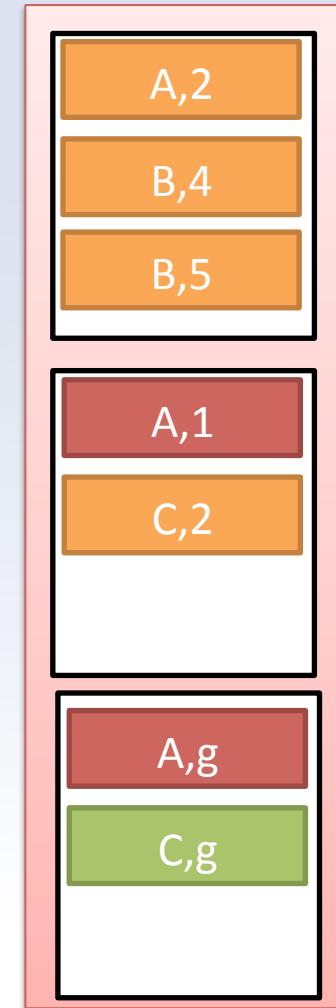
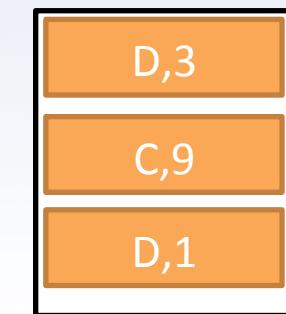
block 1



block 2

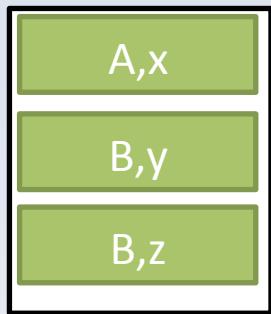


block 3

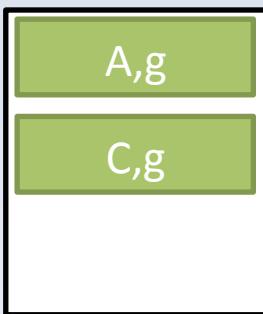


Example: $R \bowtie S$

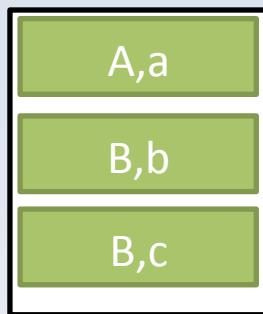
block 0



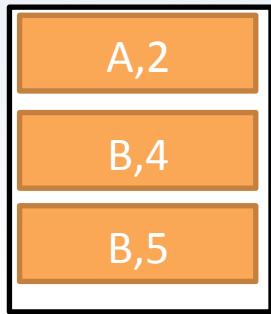
block 1



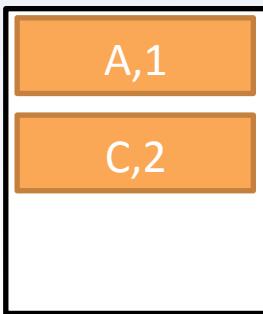
block 2



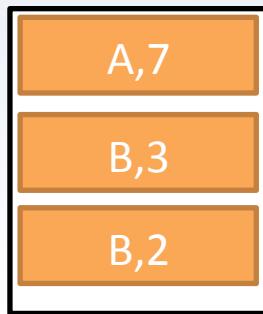
block 0



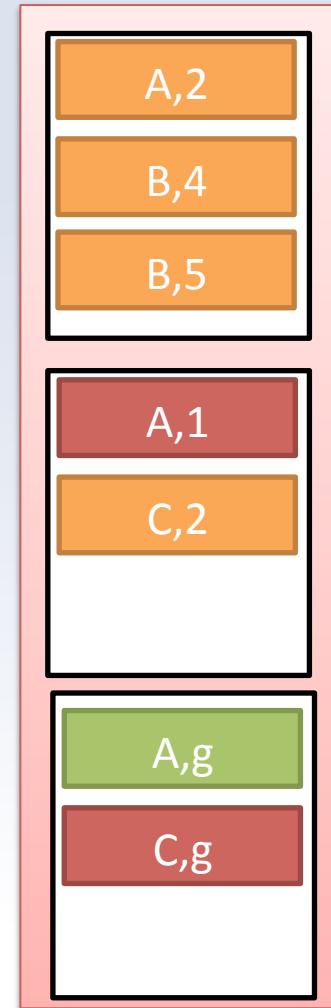
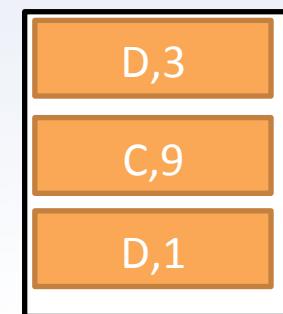
block 1



block 2

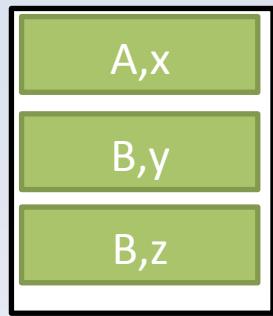


block 3

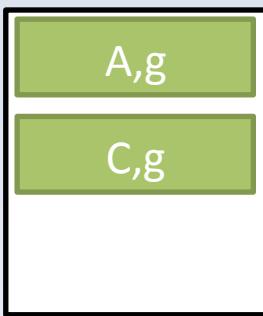


Example: $R \bowtie S$

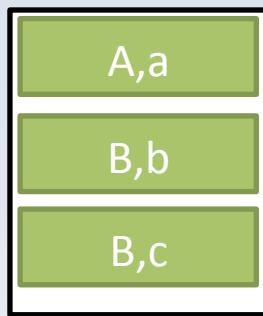
block 0



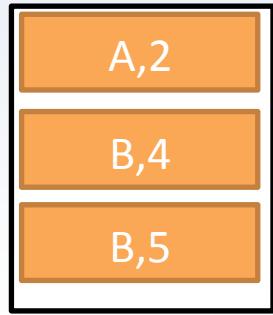
block 1



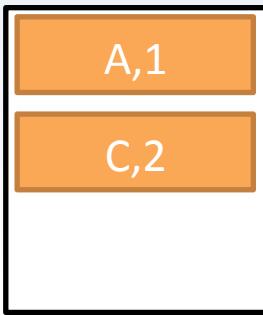
block 2



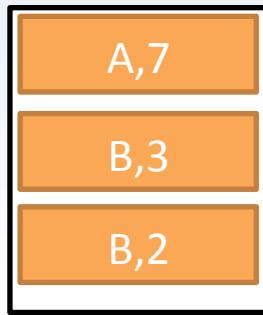
block 0



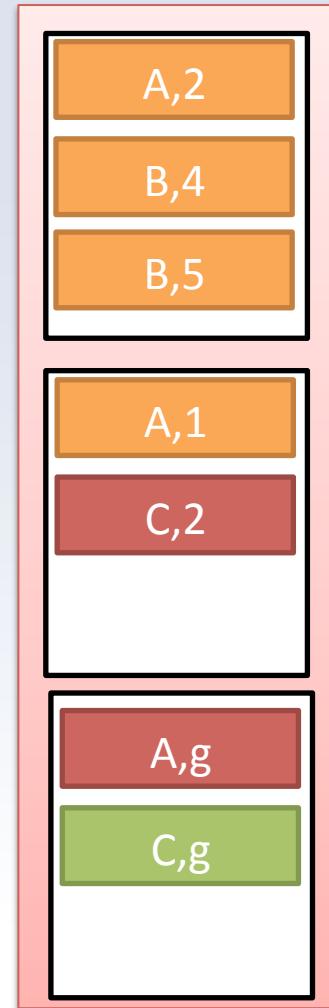
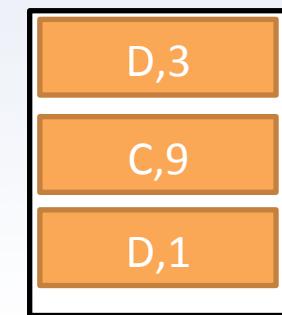
block 1



block 2

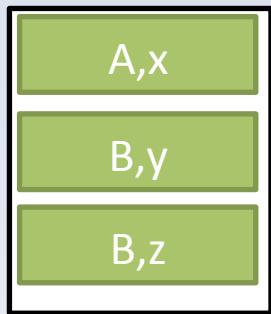


block 3

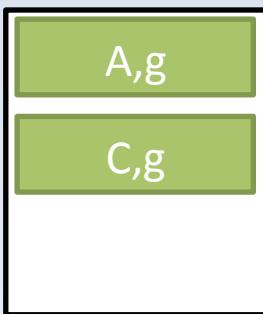


Example: $R \bowtie S$

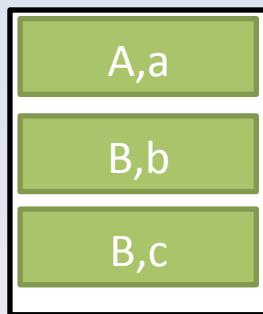
block 0



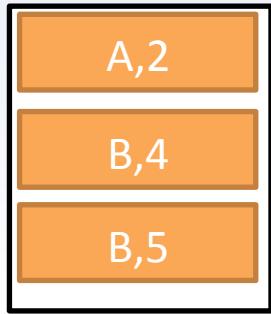
block 1



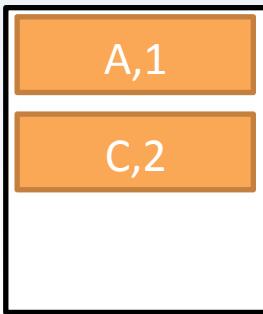
block 2



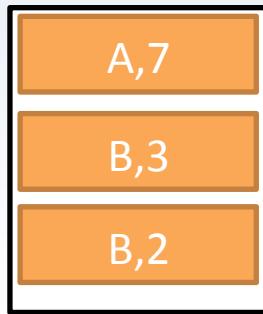
block 0



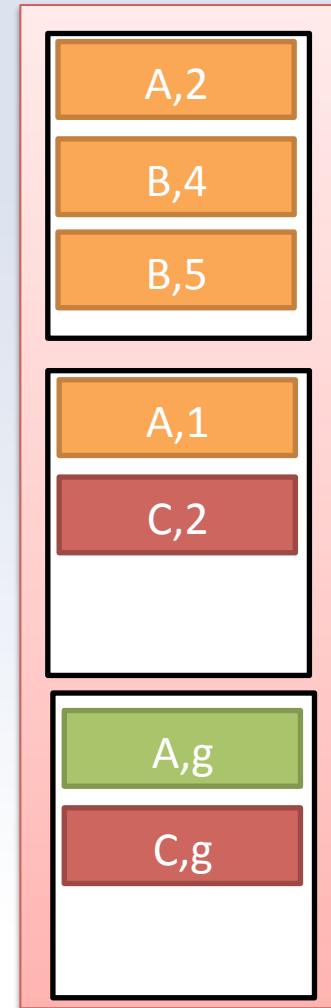
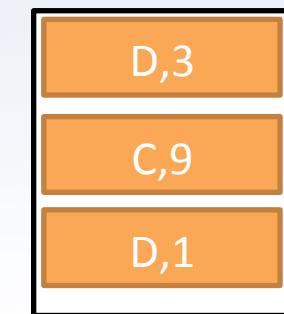
block 1



block 2

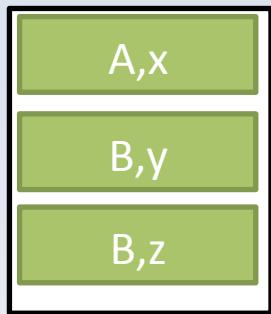


block 3

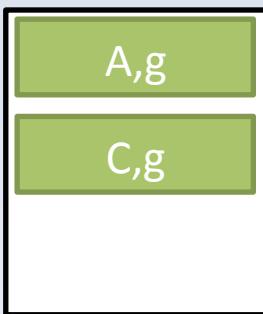


Example: $R \bowtie S$

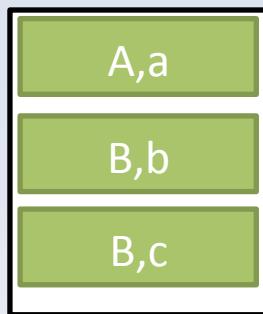
block 0



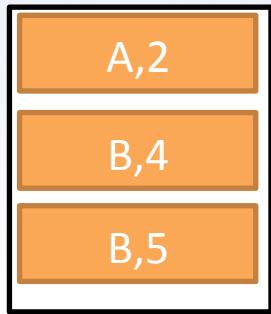
block 1



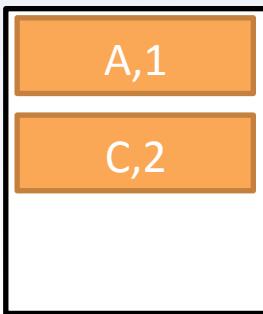
block 2



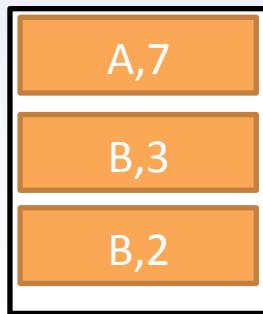
block 0



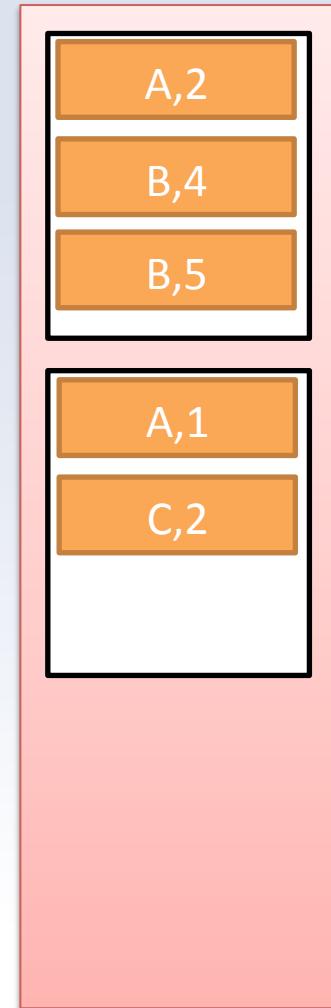
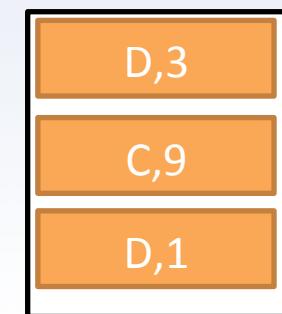
block 1



block 2

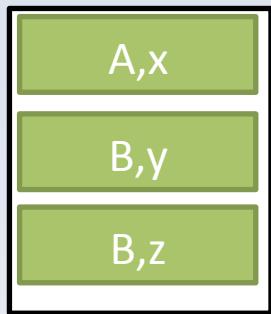


block 3

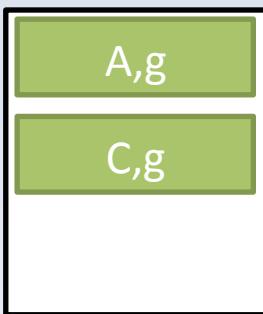


Example: $R \bowtie S$

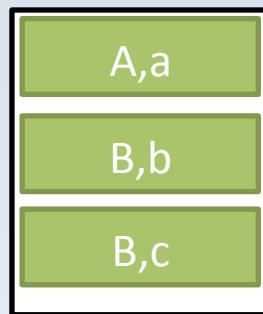
block 0



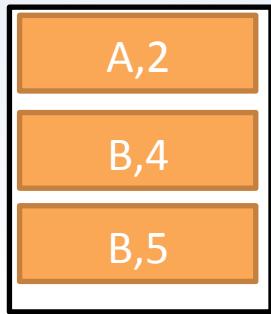
block 1



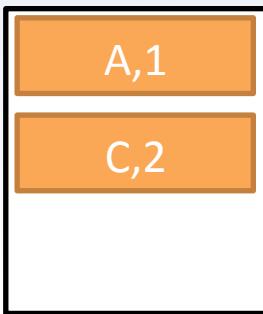
block 2



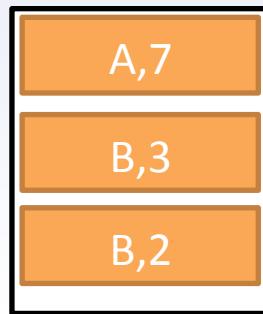
block 0



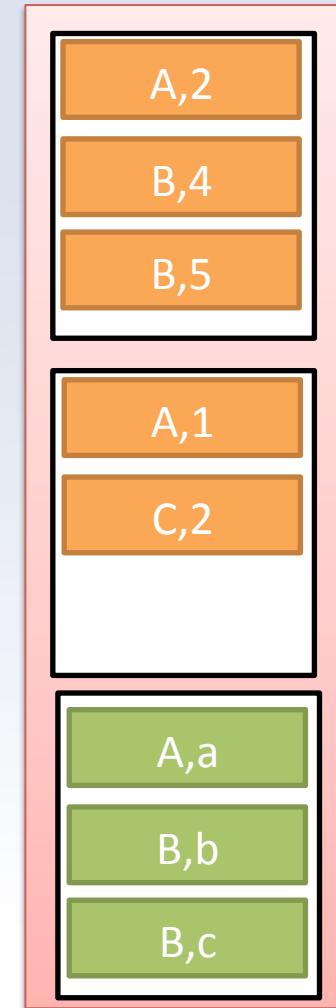
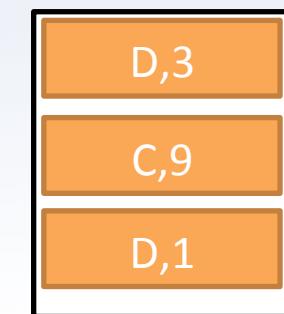
block 1



block 2

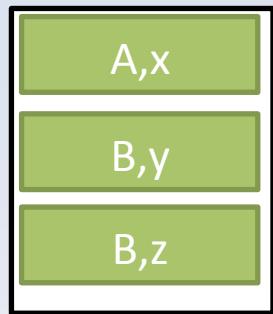


block 3

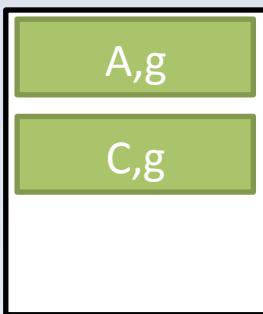


Example: $R \bowtie S$

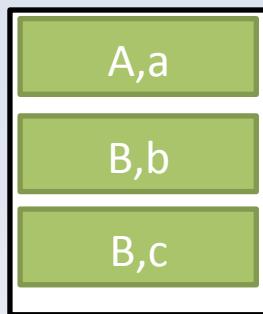
block 0



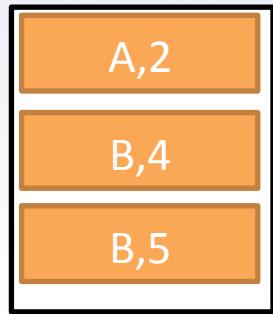
block 1



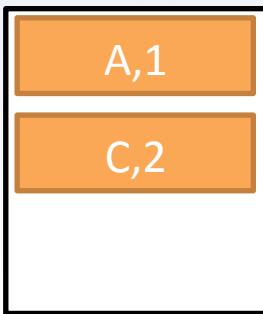
block 2



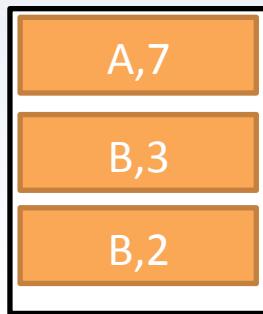
block 0



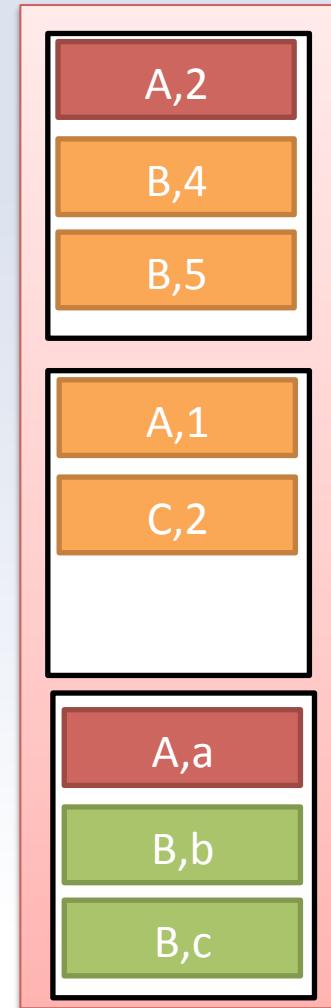
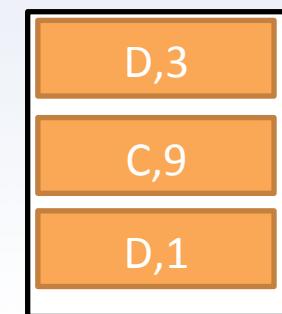
block 1



block 2

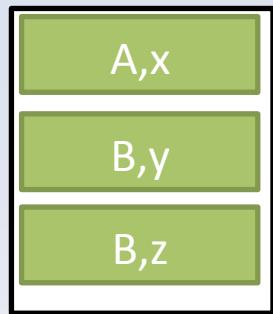


block 3

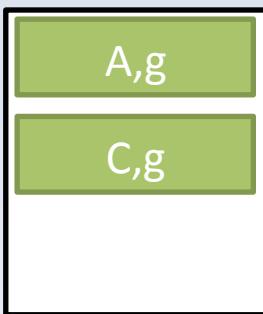


Example: $R \bowtie S$

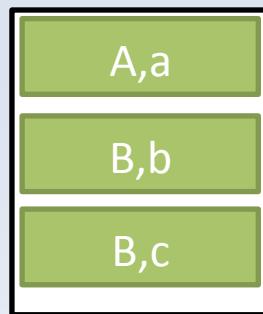
block 0



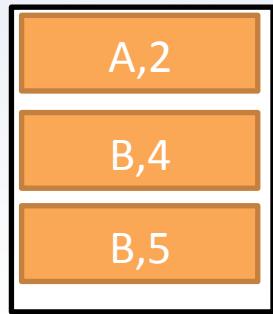
block 1



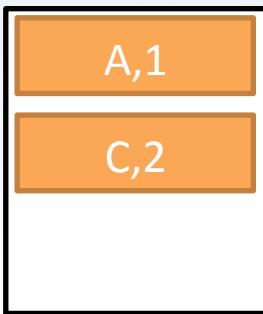
block 2



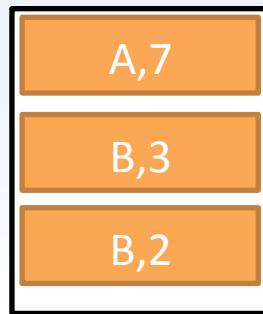
block 0



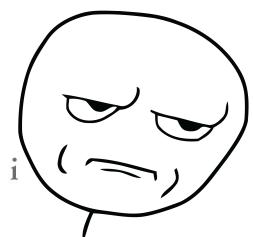
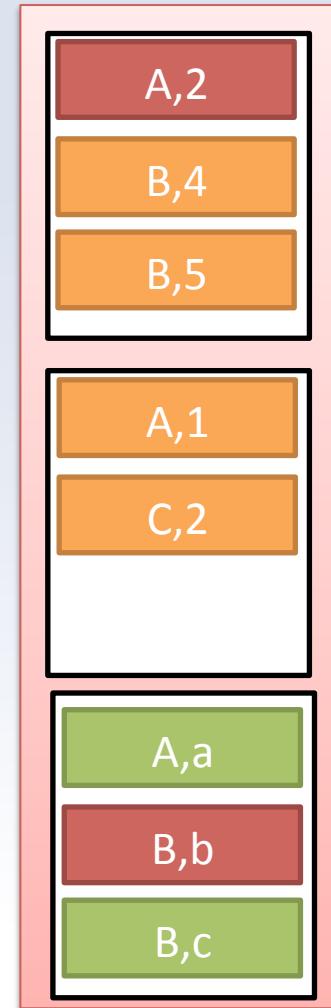
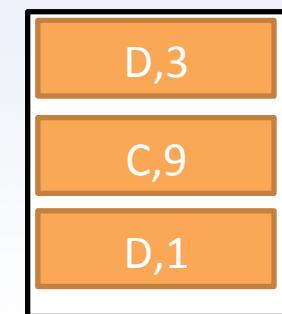
block 1



block 2

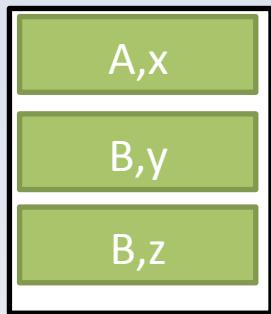


block 3

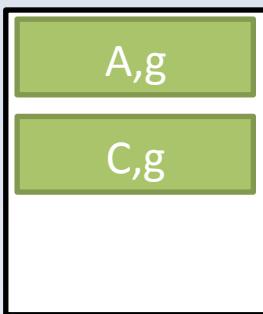


Example: $R \bowtie S$

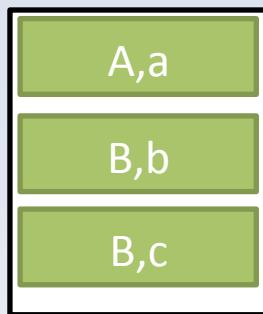
block 0



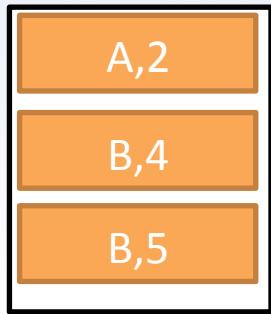
block 1



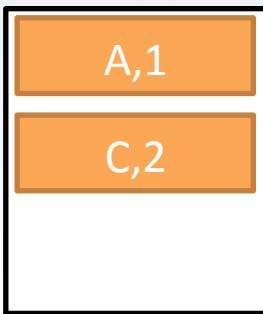
block 2



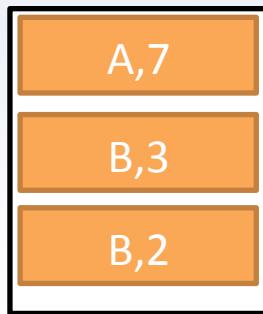
block 0



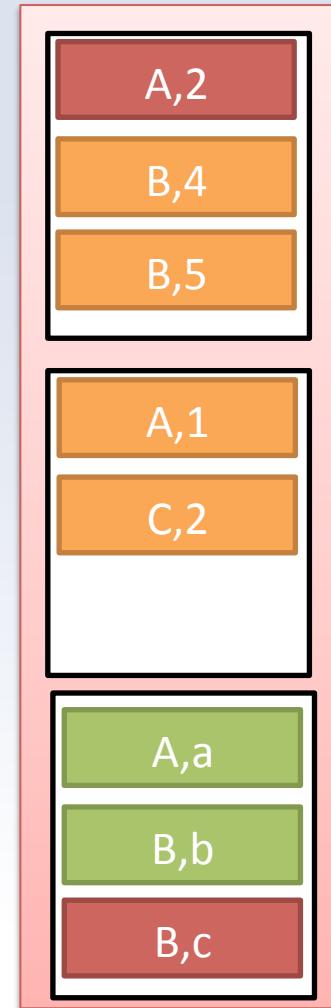
block 1



block 2

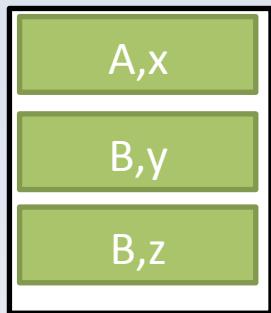


block 3

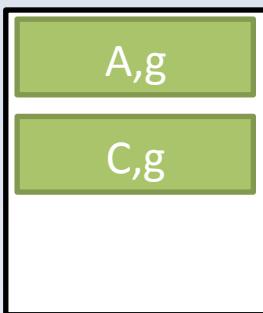


Example: $R \bowtie S$

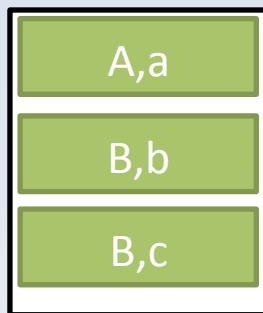
block 0



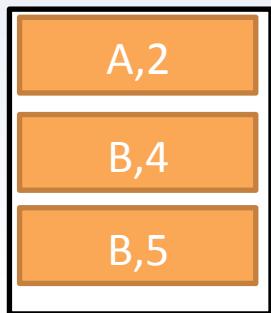
block 1



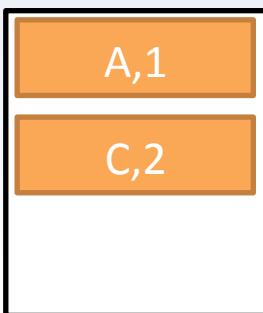
block 2



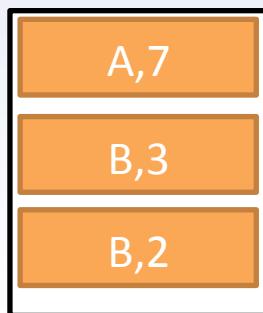
block 0



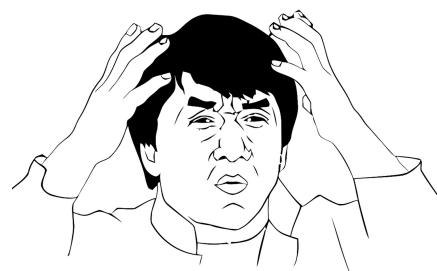
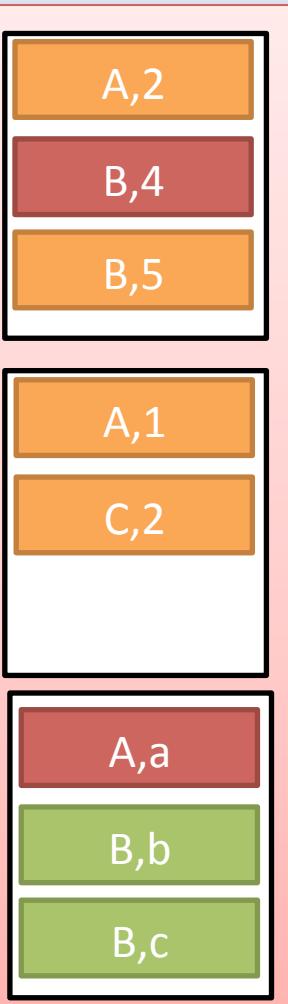
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

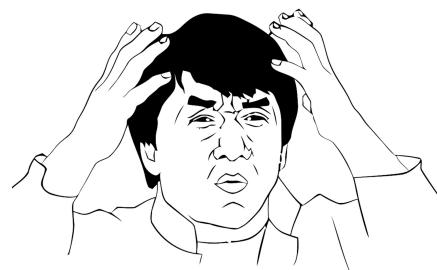
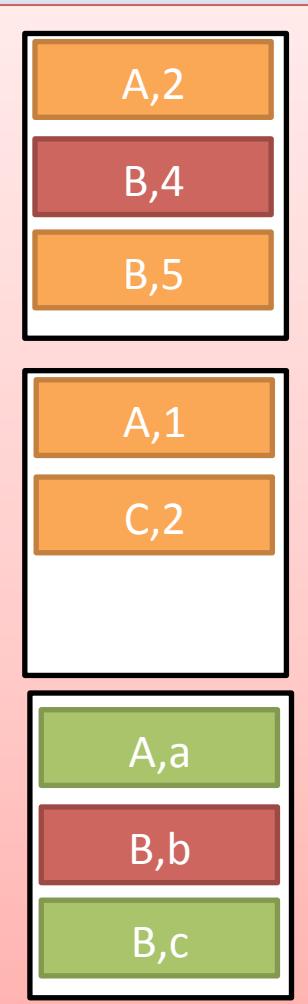
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

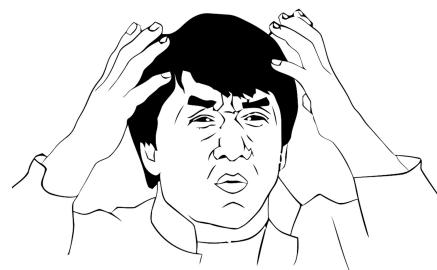
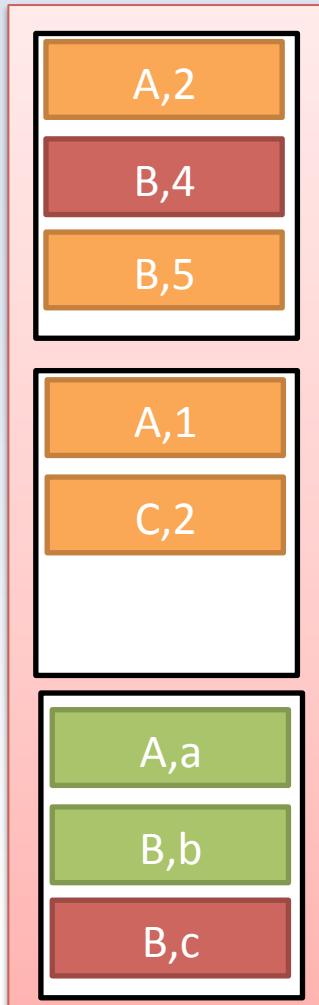
A,1
C,2

block 2

A,7
B,3
B,2

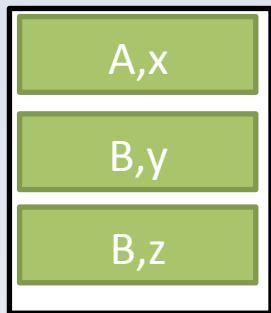
block 3

D,3
C,9
D,1

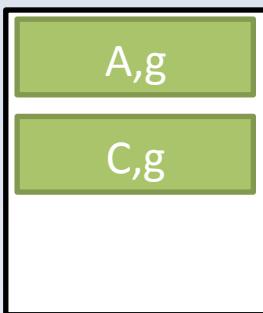


Example: $R \bowtie S$

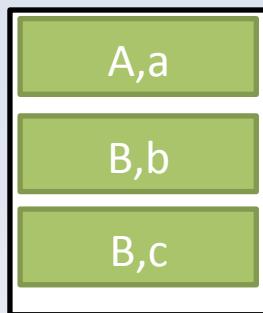
block 0



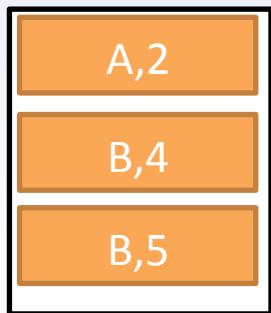
block 1



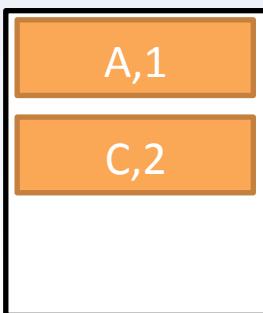
block 2



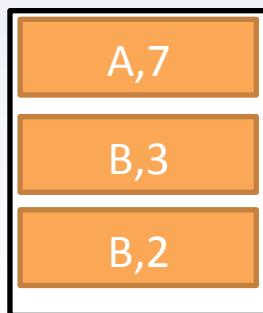
block 0



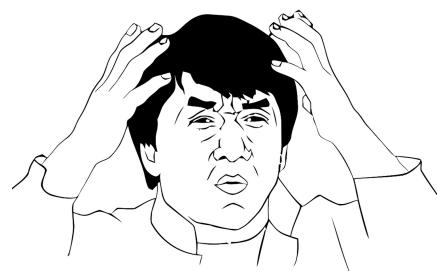
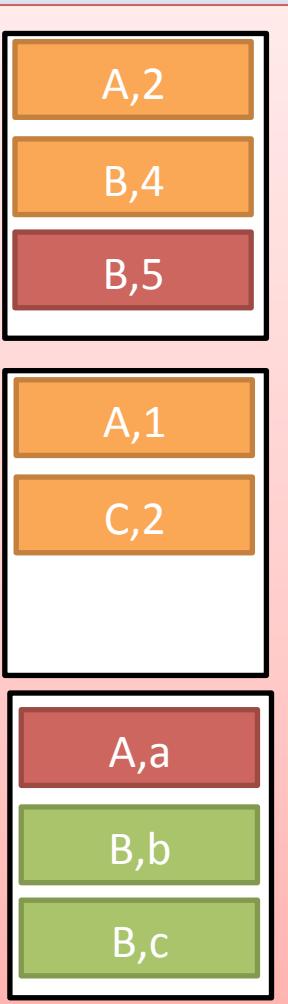
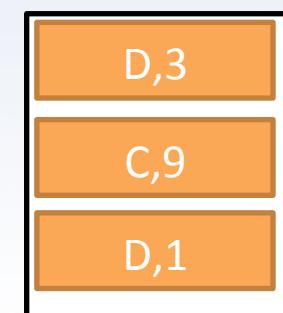
block 1



block 2

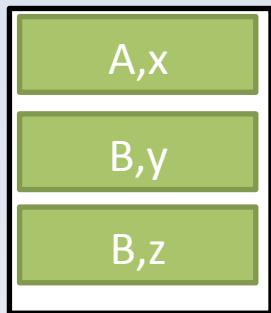


block 3

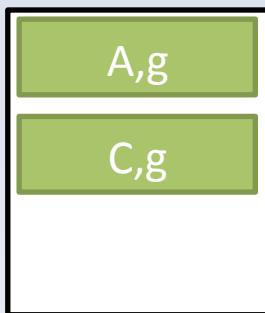


Example: $R \bowtie S$

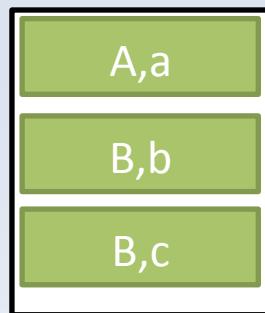
block 0



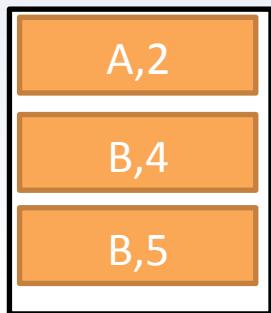
block 1



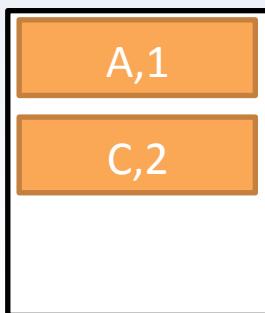
block 2



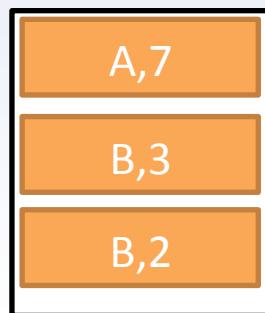
block 0



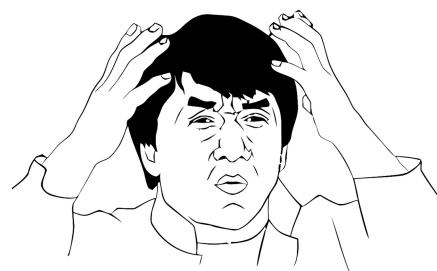
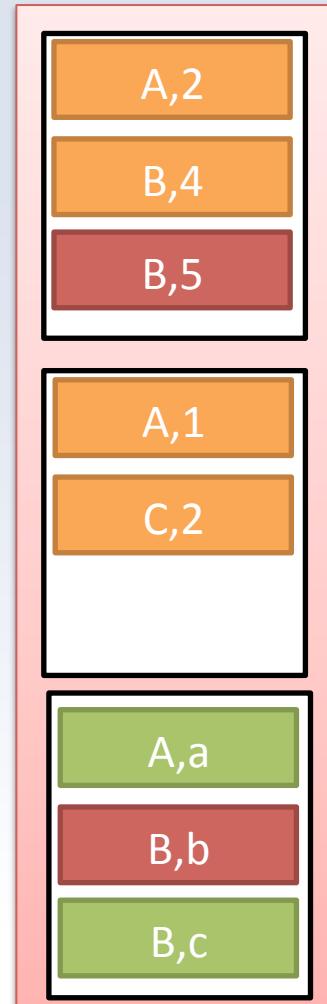
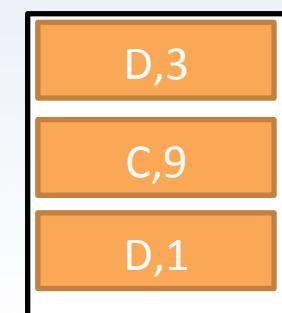
block 1



block 2

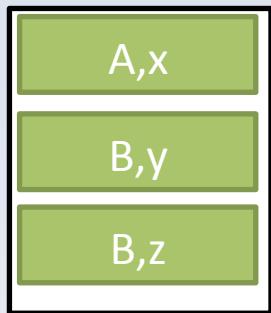


block 3

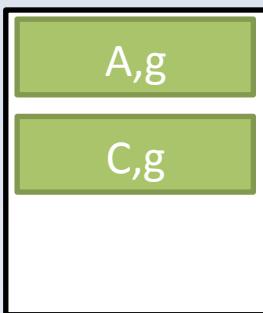


Example: $R \bowtie S$

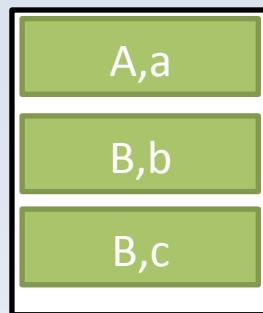
block 0



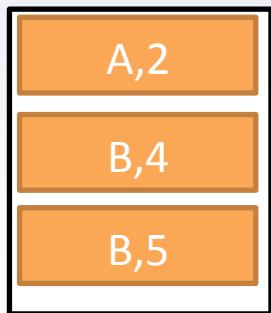
block 1



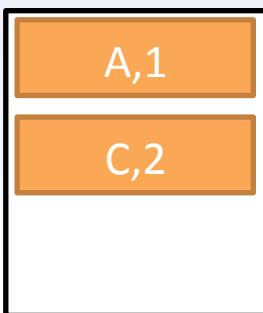
block 2



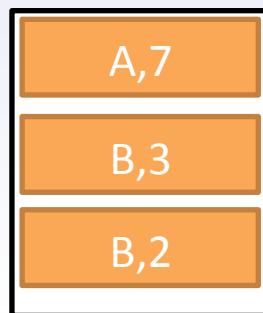
block 0



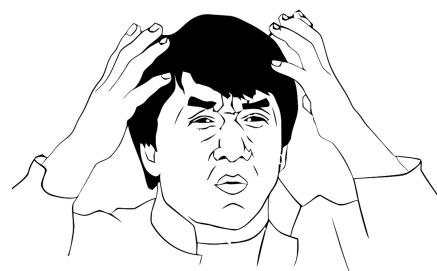
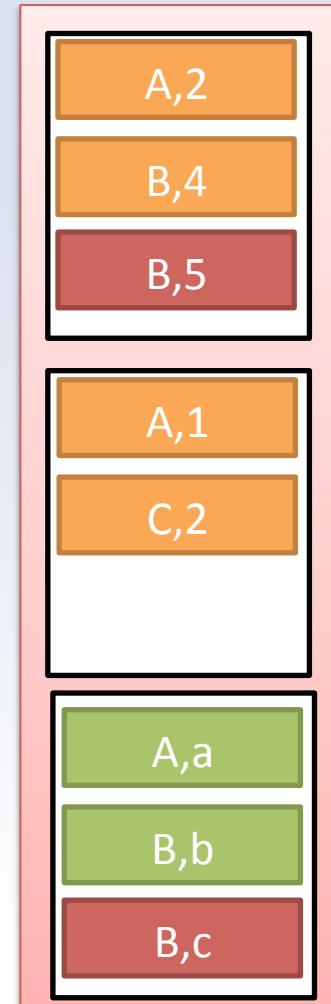
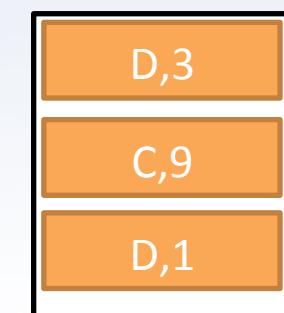
block 1



block 2

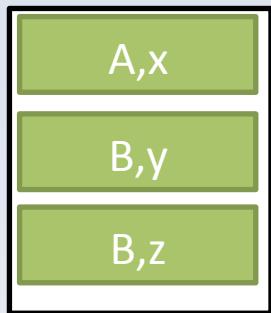


block 3

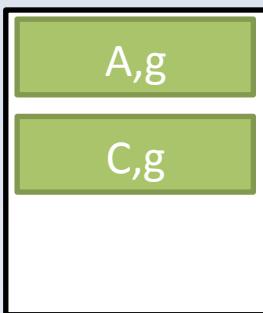


Example: $R \bowtie S$

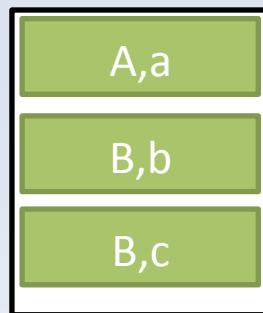
block 0



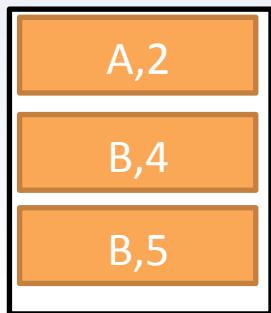
block 1



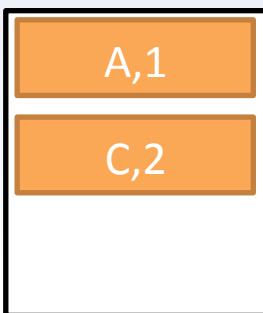
block 2



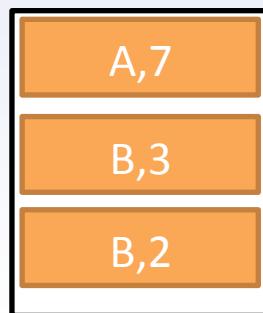
block 0



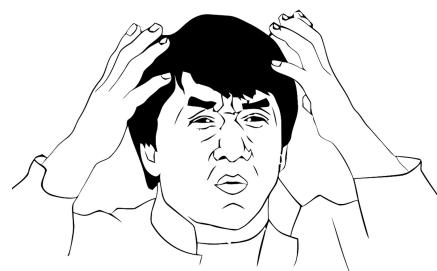
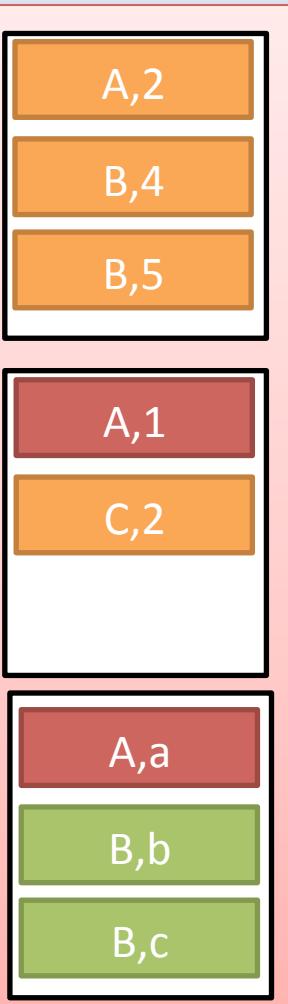
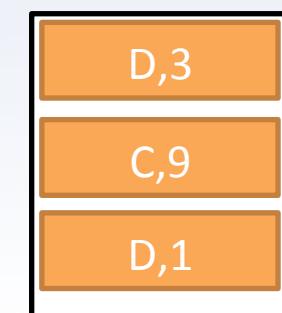
block 1



block 2

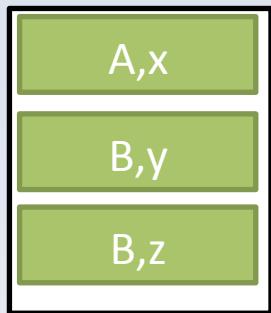


block 3

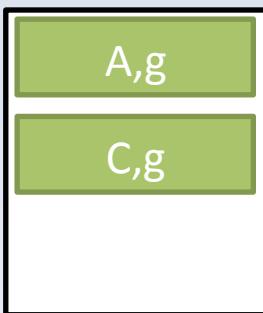


Example: $R \bowtie S$

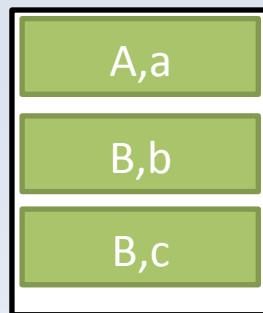
block 0



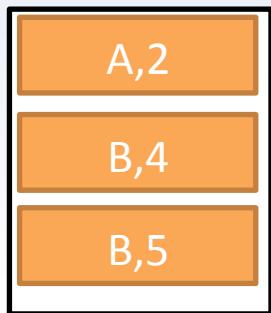
block 1



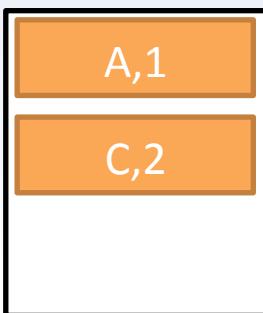
block 2



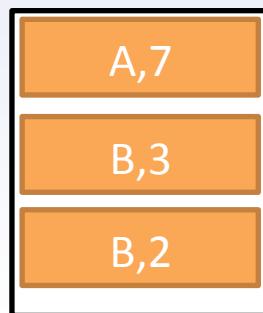
block 0



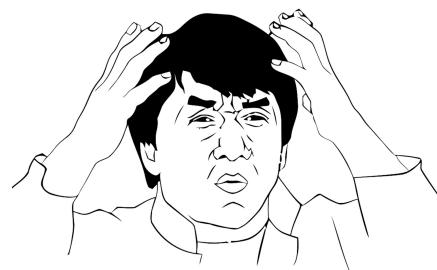
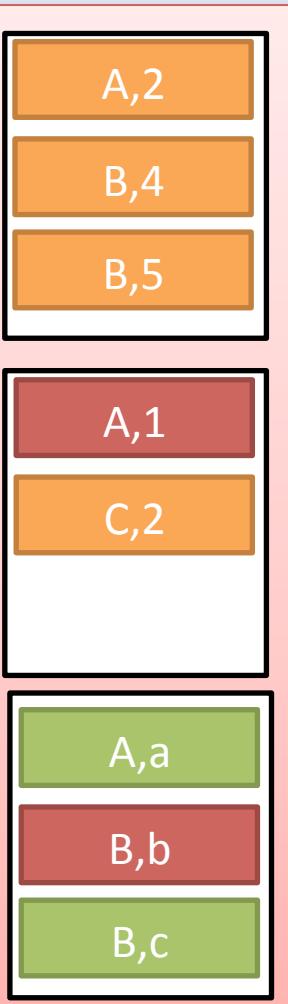
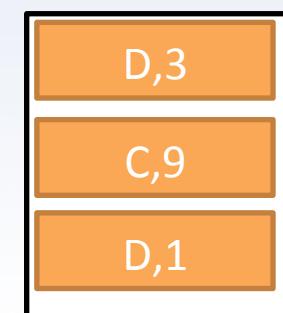
block 1



block 2

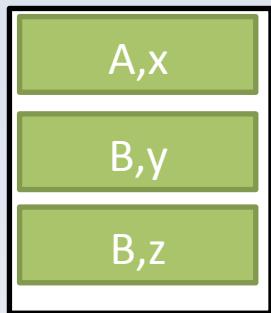


block 3

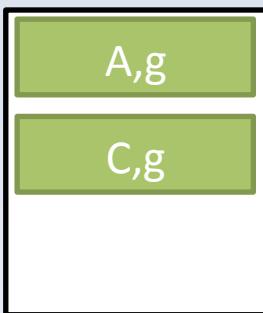


Example: $R \bowtie S$

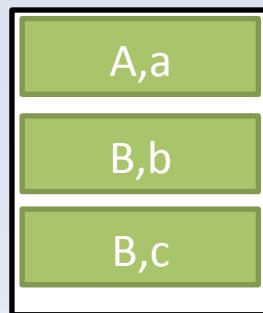
block 0



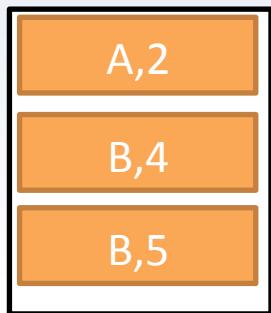
block 1



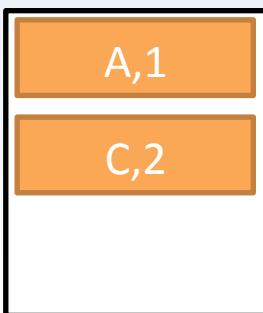
block 2



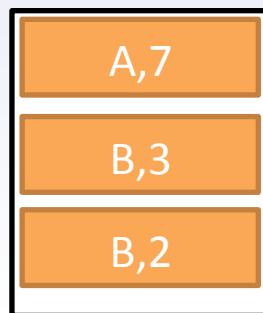
block 0



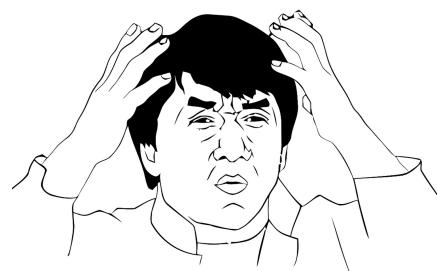
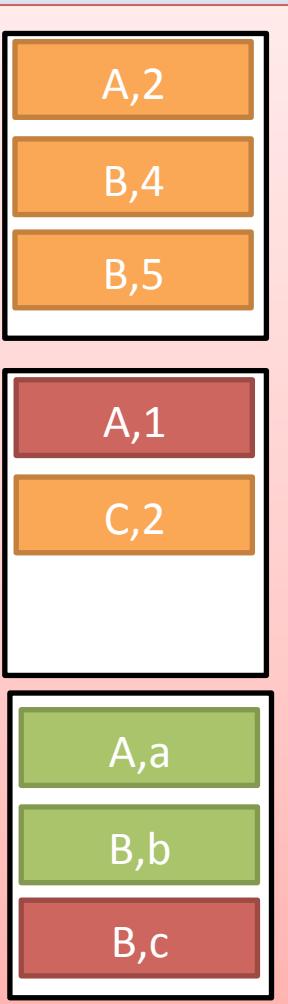
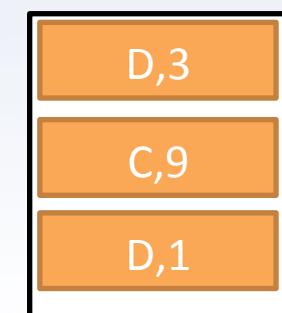
block 1



block 2

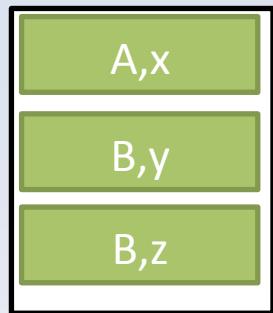


block 3

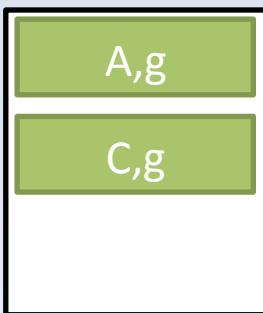


Example: $R \bowtie S$

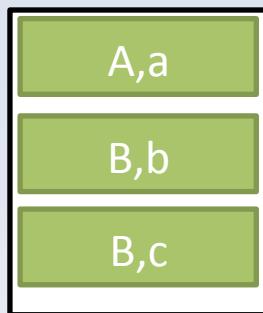
block 0



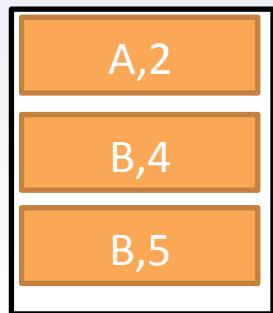
block 1



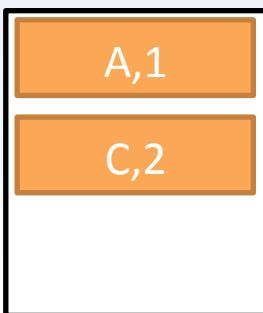
block 2



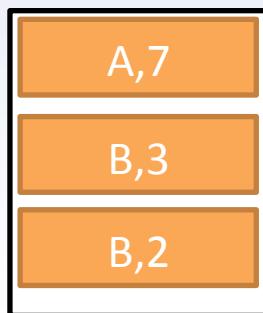
block 0



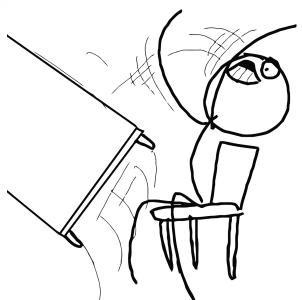
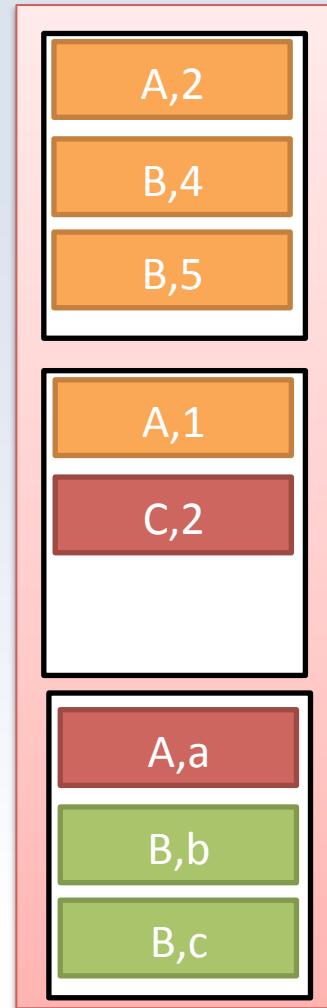
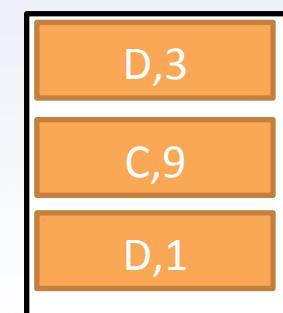
block 1



block 2

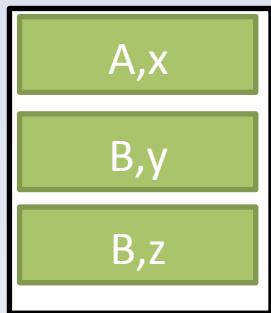


block 3

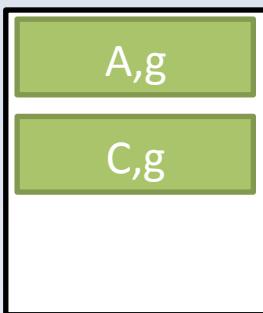


Example: $R \bowtie S$

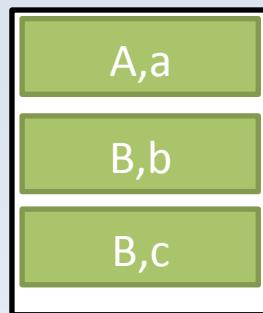
block 0



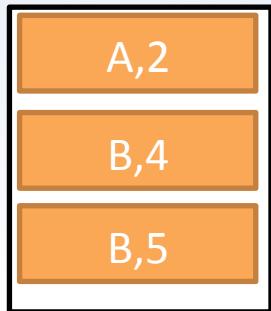
block 1



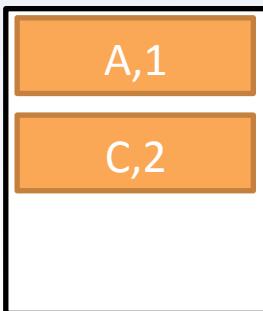
block 2



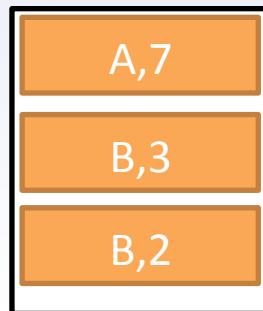
block 0



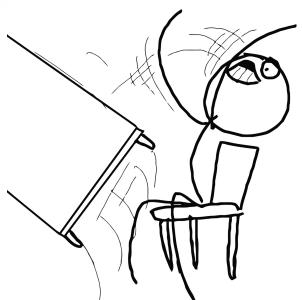
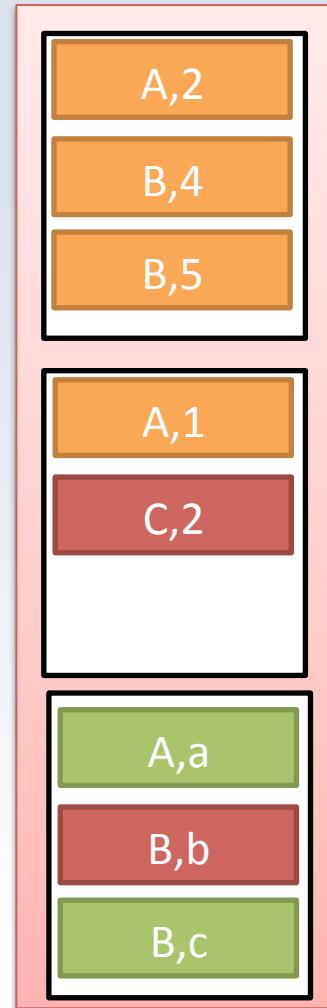
block 1



block 2

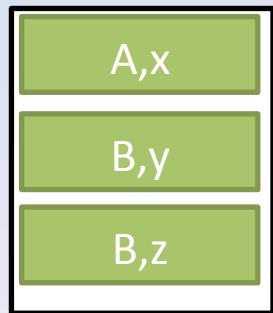


block 3

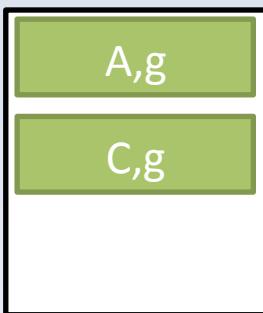


Example: $R \bowtie S$

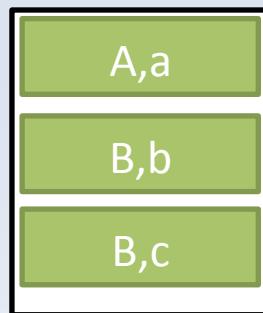
block 0



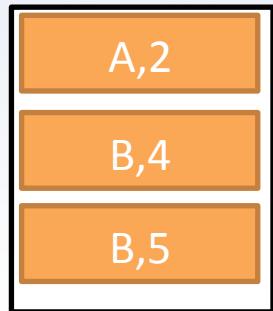
block 1



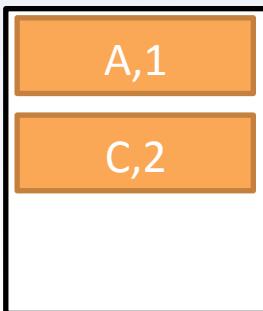
block 2



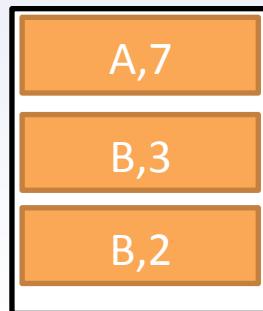
block 0



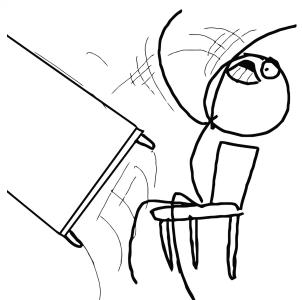
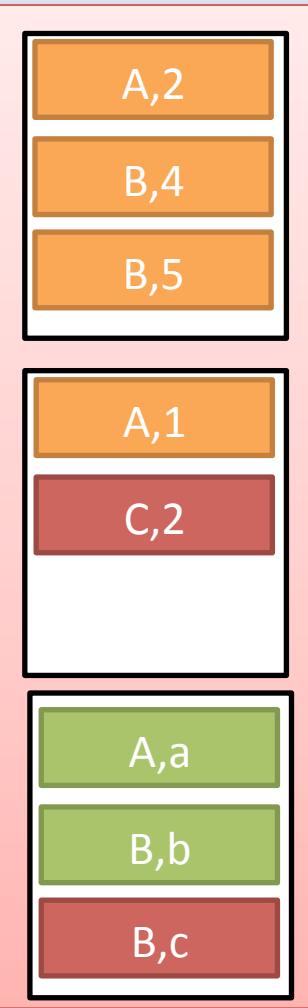
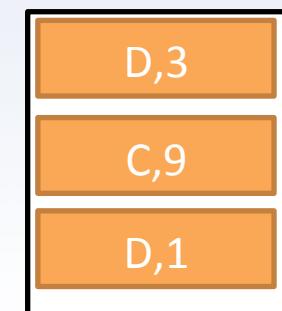
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

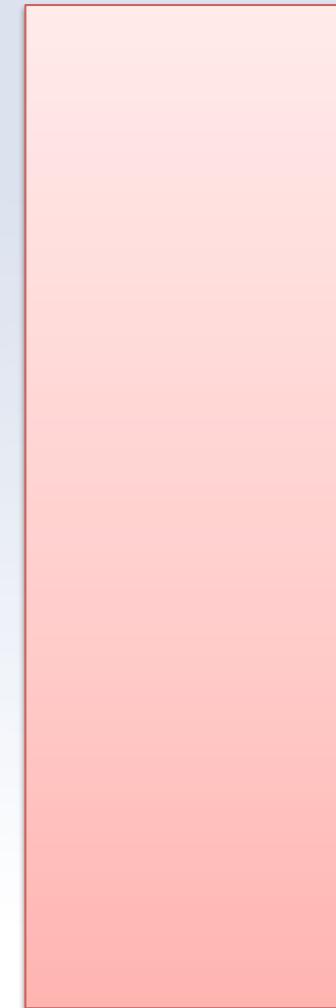
A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c



block 0

A,2
B,4
B,5

block 1

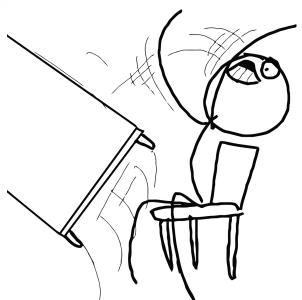
A,1
C,2

block 2

A,7
B,3
B,2

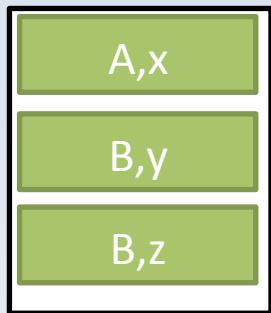
block 3

D,3
C,9
D,1

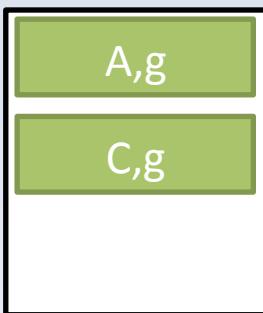


Example: $R \bowtie S$

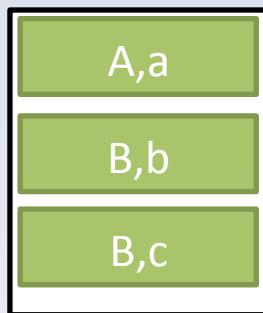
block 0



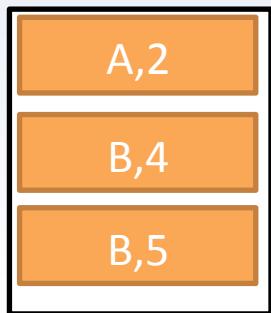
block 1



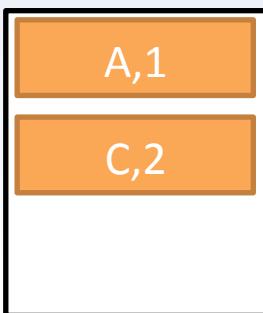
block 2



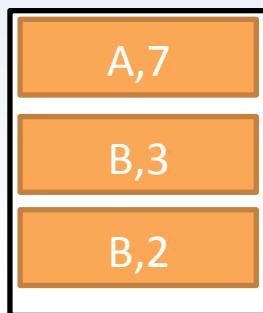
block 0



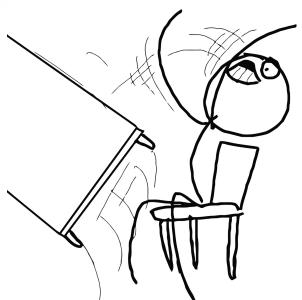
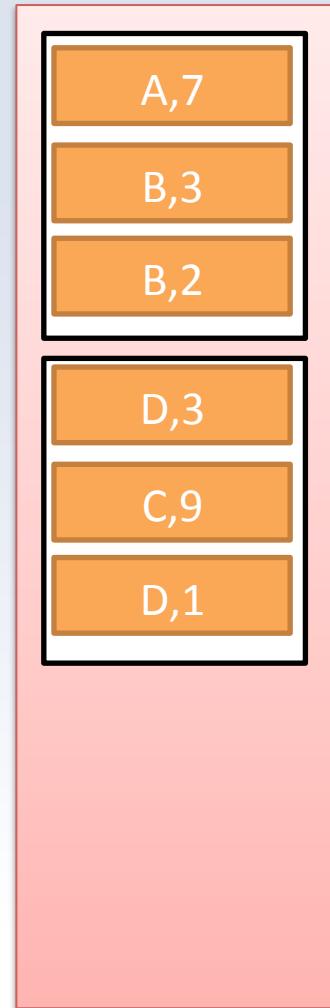
block 1



block 2

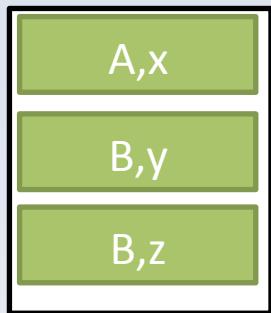


block 3

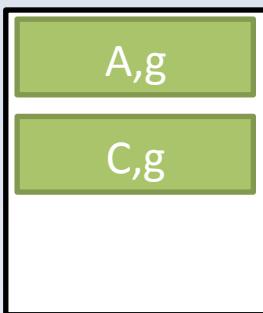


Example: $R \bowtie S$

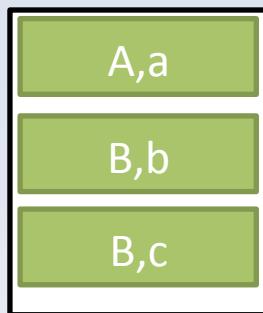
block 0



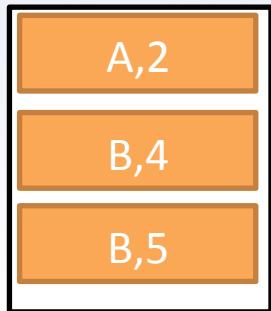
block 1



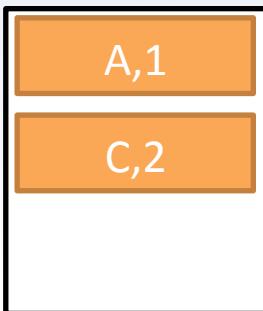
block 2



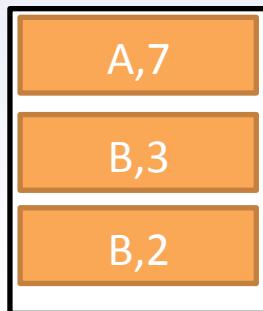
block 0



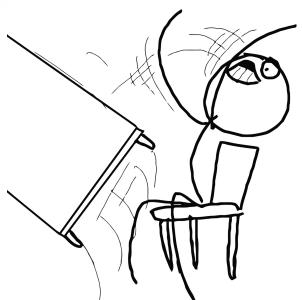
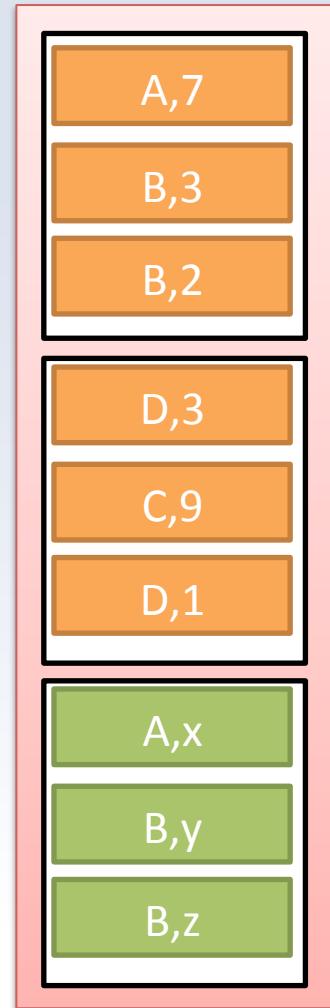
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

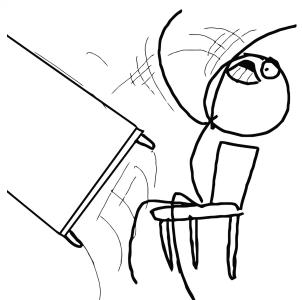
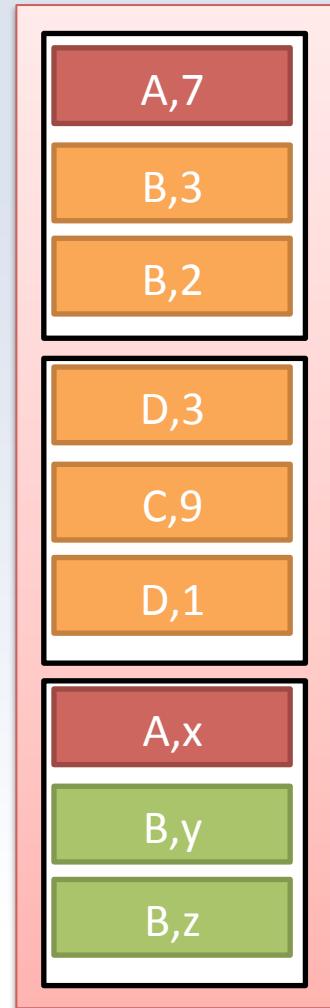
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

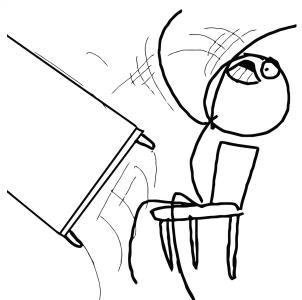
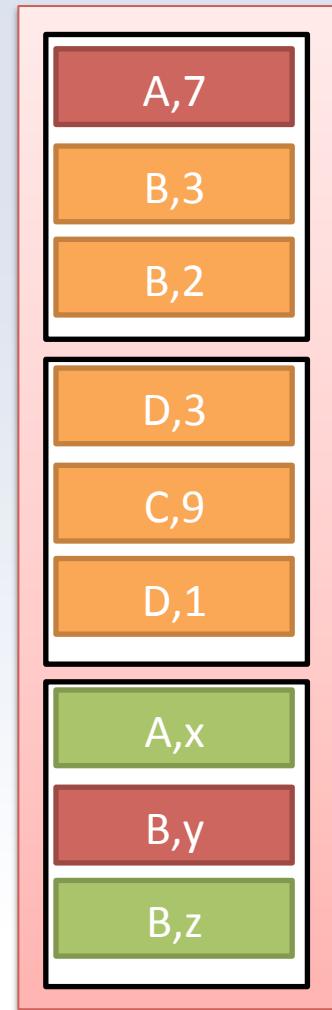
A,1
C,2

block 2

A,7
B,3
B,2

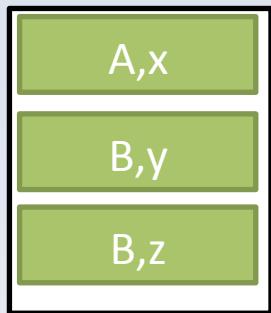
block 3

D,3
C,9
D,1

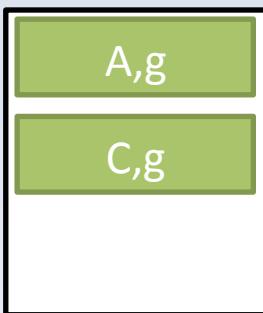


Example: $R \bowtie S$

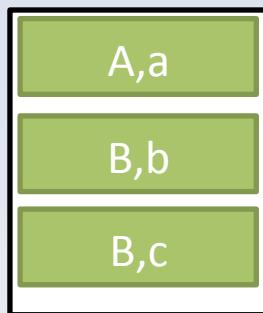
block 0



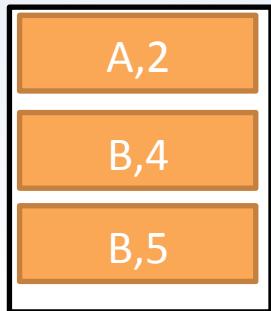
block 1



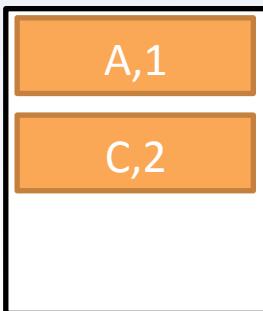
block 2



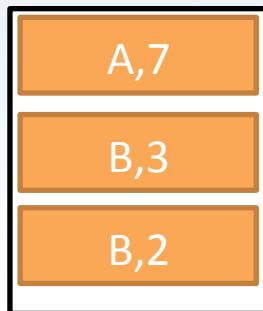
block 0



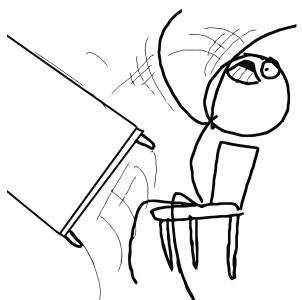
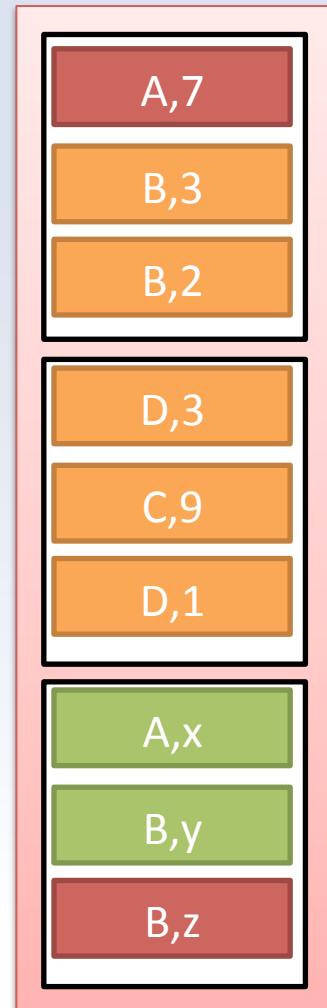
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

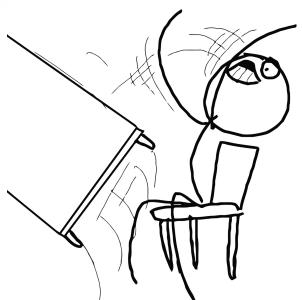
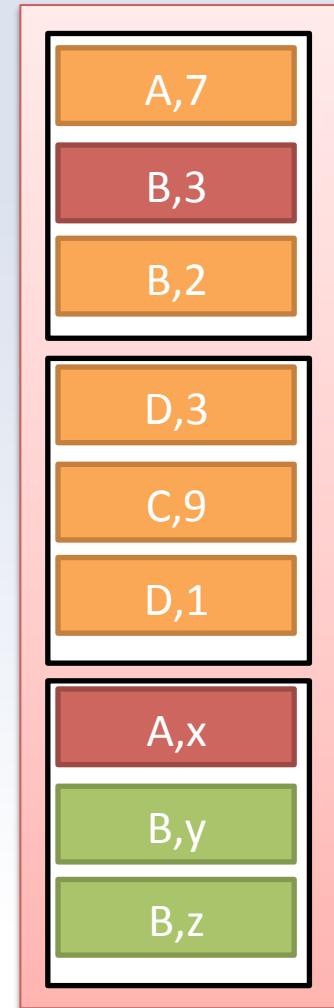
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

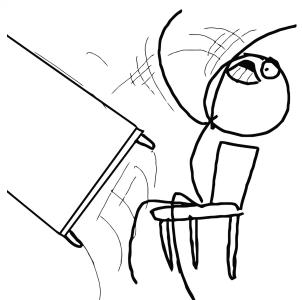
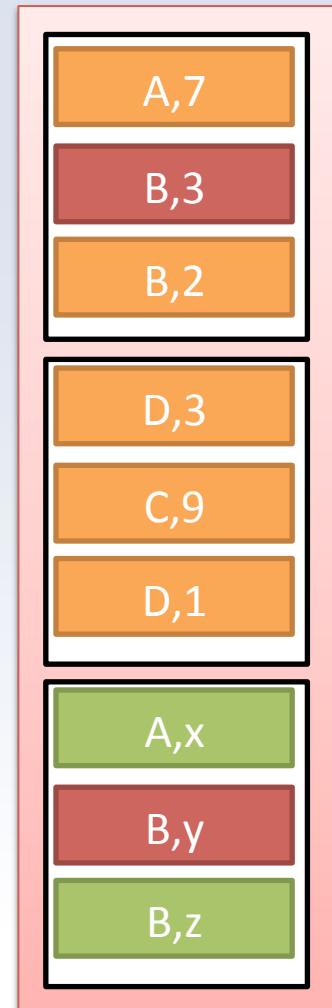
A,1
C,2

block 2

A,7
B,3
B,2

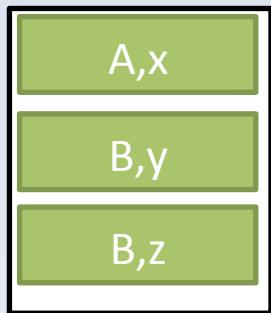
block 3

D,3
C,9
D,1

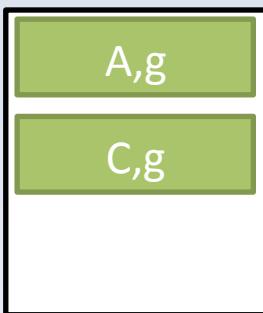


Example: $R \bowtie S$

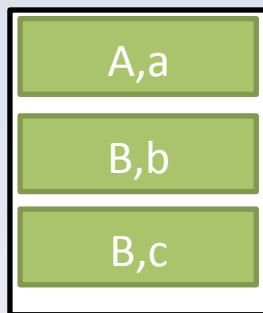
block 0



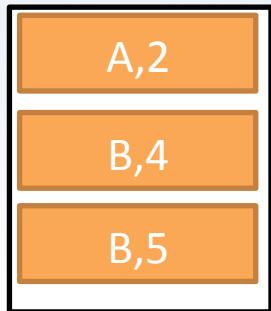
block 1



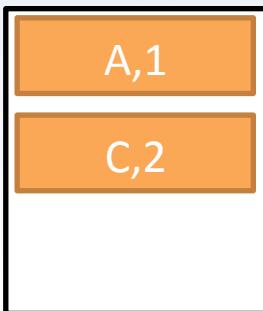
block 2



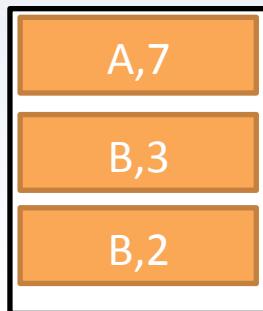
block 0



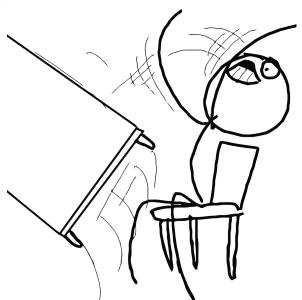
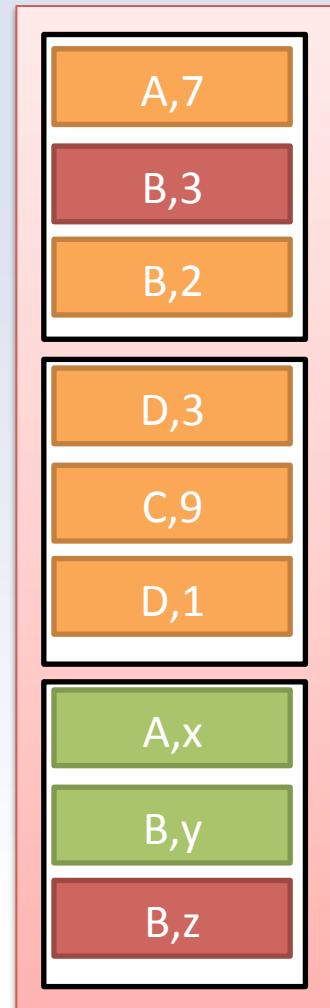
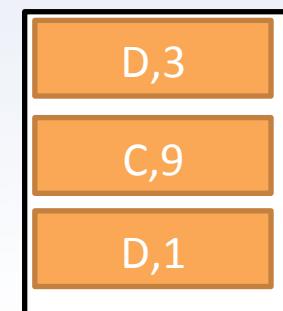
block 1



block 2



block 3



Example: $R \bowtie S$

block 0

A,x
B,y
B,z

block 1

A,g
C,g

block 2

A,a
B,b
B,c

block 0

A,2
B,4
B,5

block 1

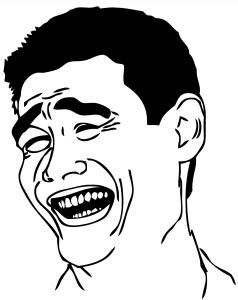
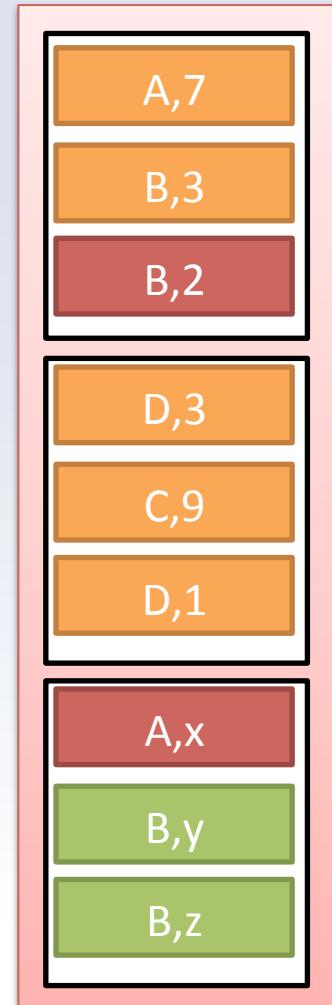
A,1
C,2

block 2

A,7
B,3
B,2

block 3

D,3
C,9
D,1



Pseudocode: Block-based

```
for (M-1) blocks bs in S:  
    for block br in R:  
        for s in bs:  
            for r in br:  
                t=join(r,s)  
                if(t):  
                    output t
```



Cost of NBLJ

- Read S once (outer loop)
 - $B(S)/(M-1)$ iterations
 - Reading *all* of R each time
- Total cost: $B(S)B(R)/(M-1)$
- Note: iterate over smaller relation first
 - keeps $B(S)/(M-1)$ small
 - fewer reads of R



One-Pass/NBLJ Summary

Operator	M required	I/O Cost
σ, π	1	B
δ, γ	B	B
$\cup, \cap, -, \bowtie, \times$	$\min(B(R), B(S))$	$B(R) + B(S)$
\bowtie	$M \geq 2$	$B(R)B(S)/M$



Sort-Based Algorithms

- Can base implementation of operations on merge sort
 - Special implementation for large data
 - Called *Two-Phase Multiway Merge-Sort*
 - TPMMS

TPMMS

- Has two phases
 - Phase 1: sort memory sized sublists
 - Phase 2: merge sorted sublists
- Requires $\sqrt{B(R)} \leq M$



TPMMS

- Phase 1: Repeat until R is exhausted:
 - Fill up *all* M buffers with blocks from R
 - Sort using quicksort
 - Write resulting sublist to disk
- Phase 2: Repeat until sublists exhausted:
 - Load buffers with smallest block for each sublist
 - Identify smallest elements, move to output block
 - When output block full, write to disk
 - If list buffer is empty, get next one from disk



Example

block 0

Q
X
F

block 1

A
R

block 2

T
V
L

block 3

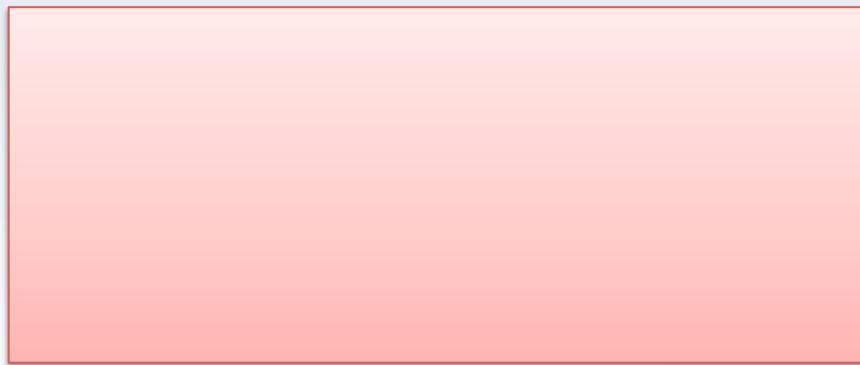
O
Q

block 4

Y
F
K

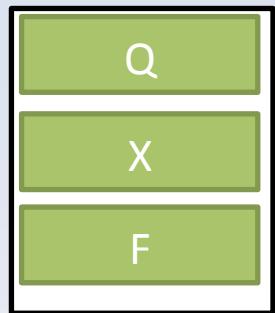
block 5

B
S
C

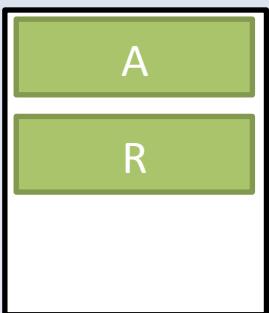


Phase 1

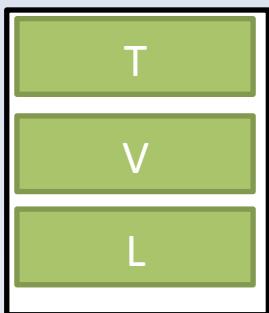
block 0



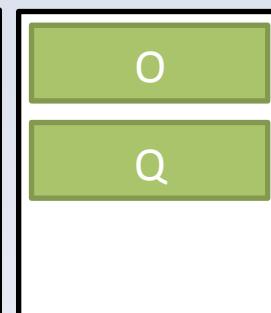
block 1



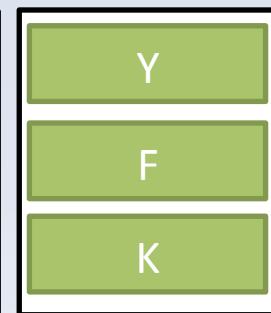
block 2



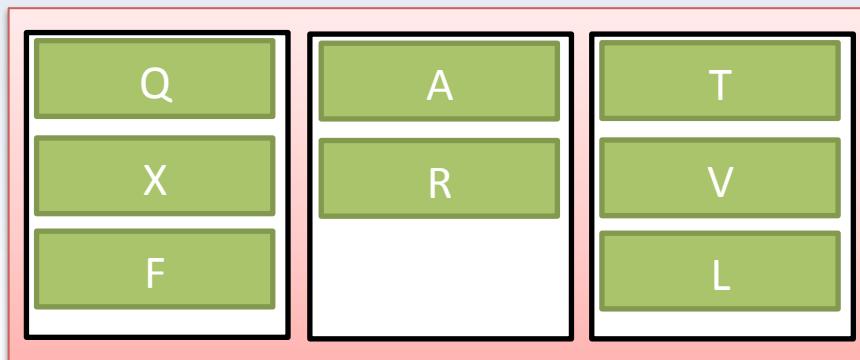
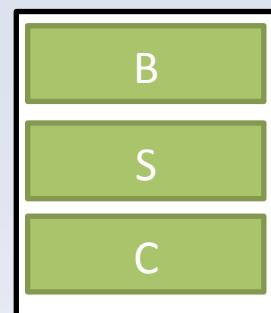
block 3



block 4

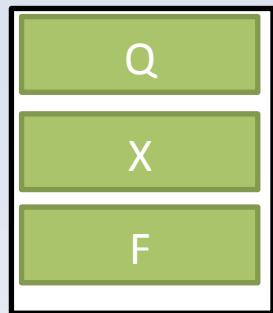


block 5

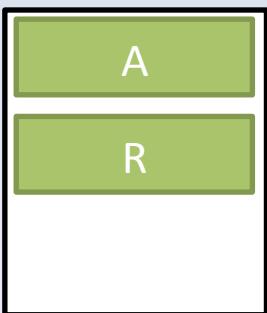


Phase 1

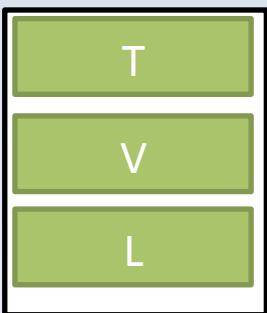
block 0



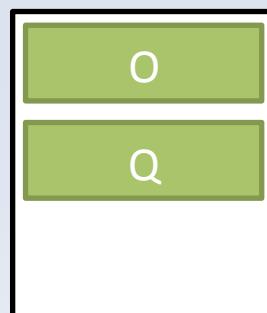
block 1



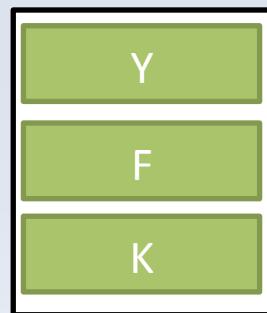
block 2



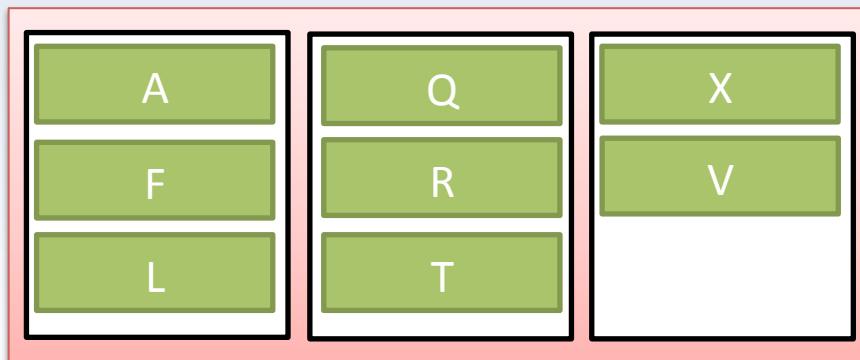
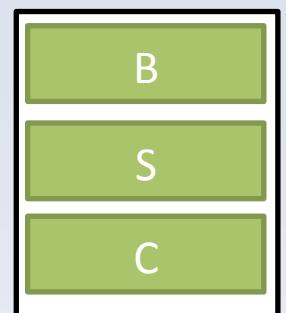
block 3



block 4

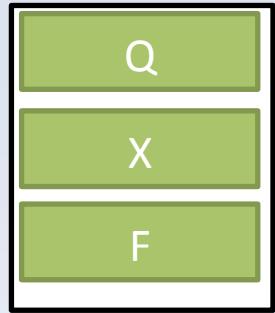


block 5

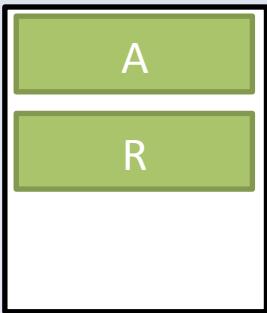


Phase 1

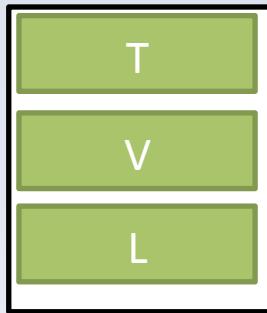
block 0



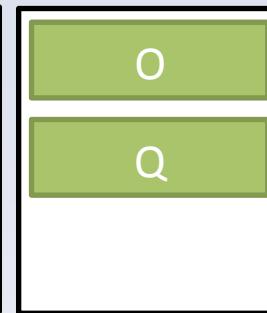
block 1



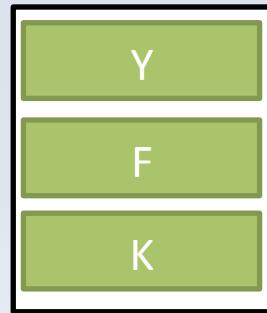
block 2



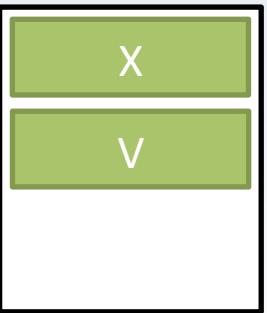
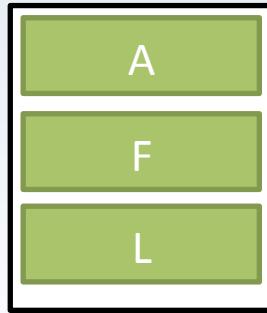
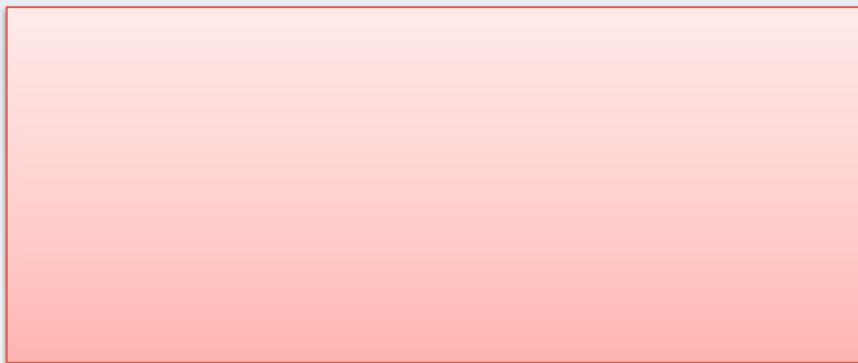
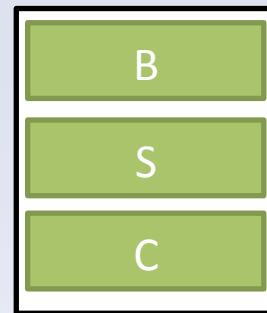
block 3



block 4

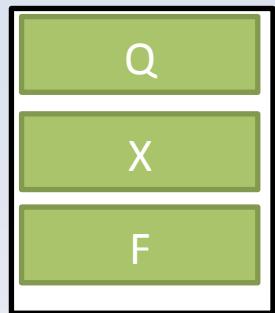


block 5

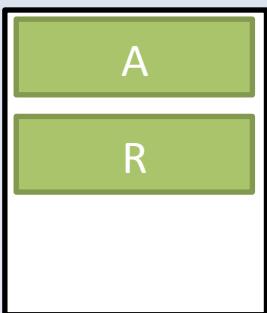


Phase 1

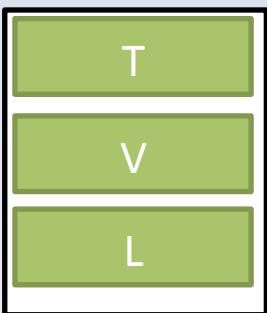
block 0



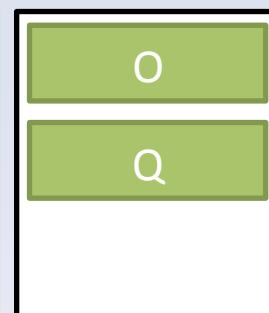
block 1



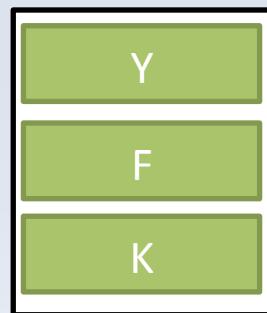
block 2



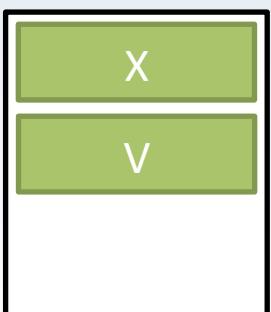
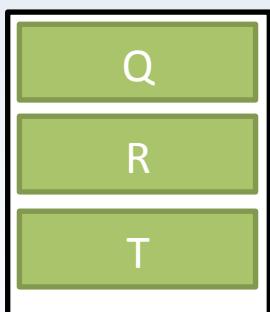
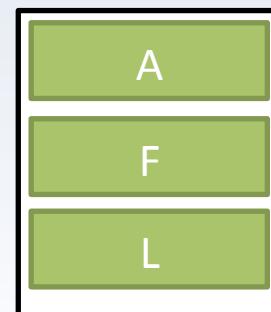
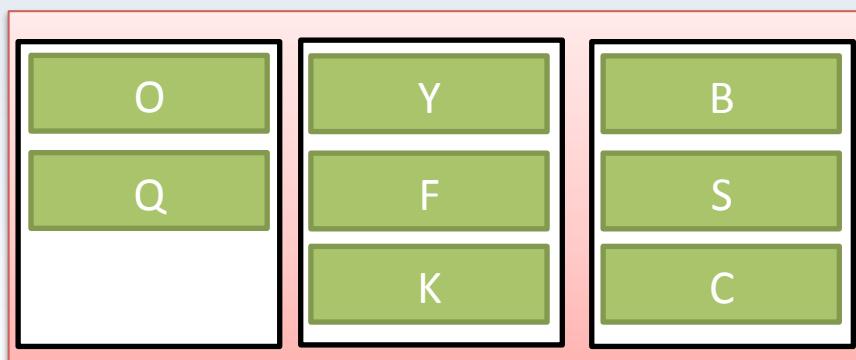
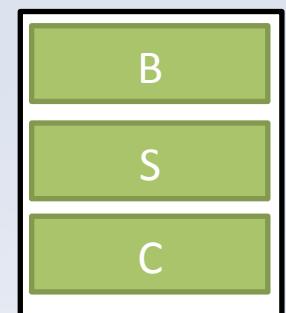
block 3



block 4

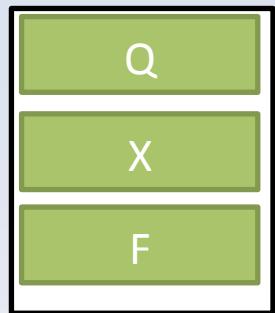


block 5

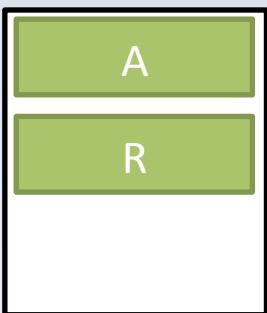


Phase 1

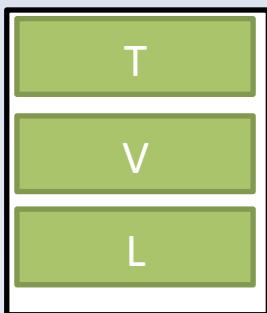
block 0



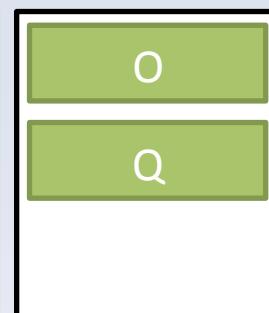
block 1



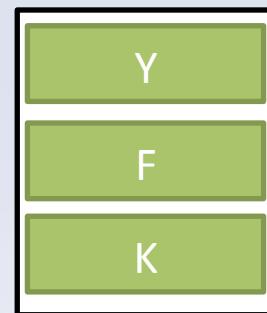
block 2



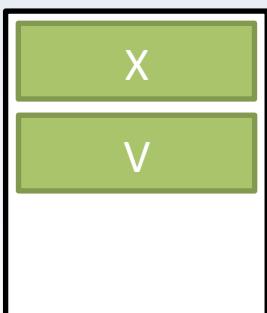
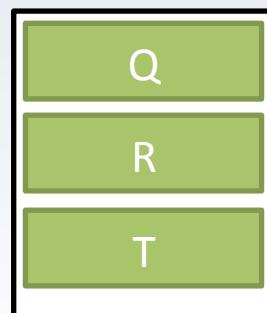
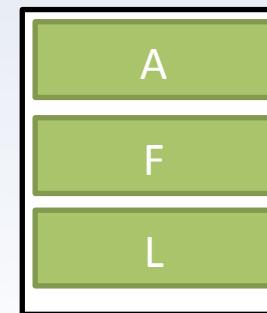
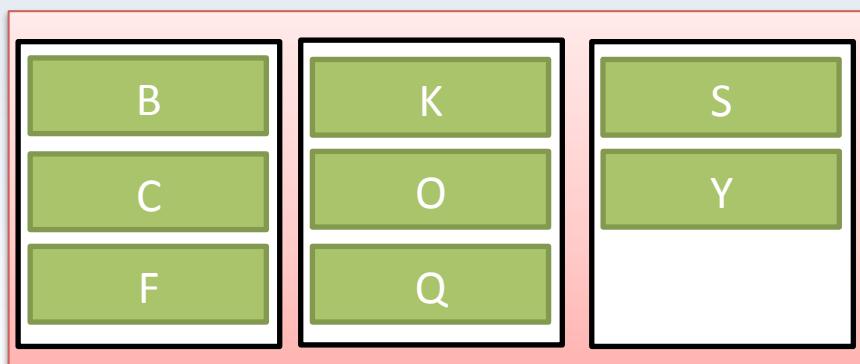
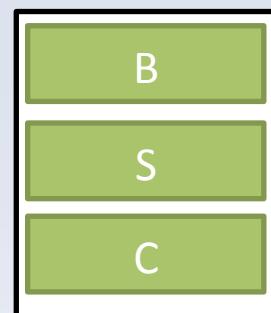
block 3



block 4

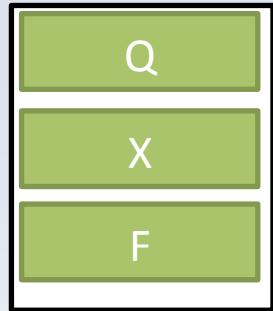


block 5

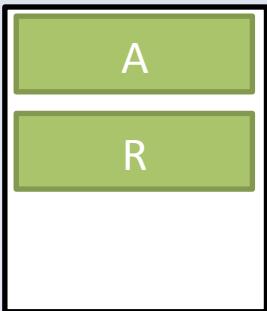


Phase 1

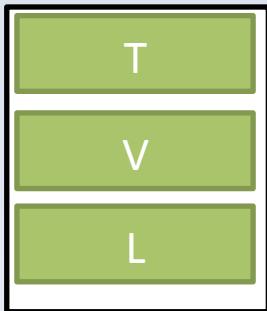
block 0



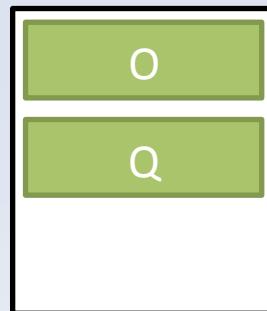
block 1



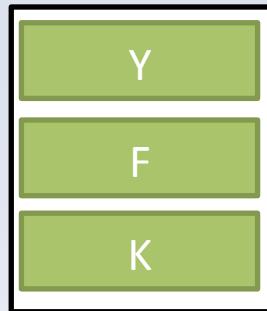
block 2



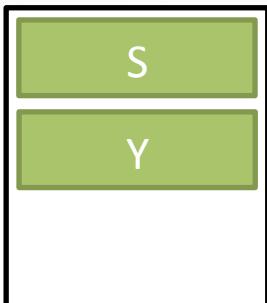
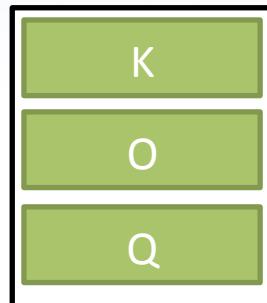
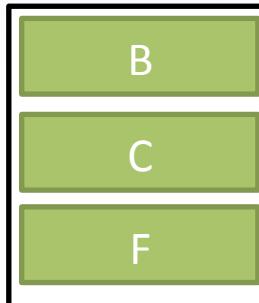
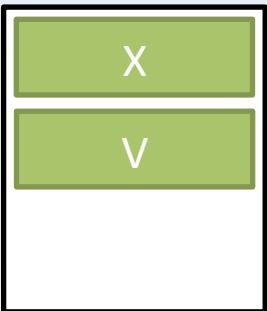
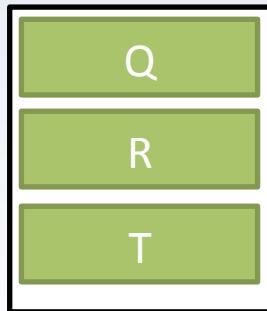
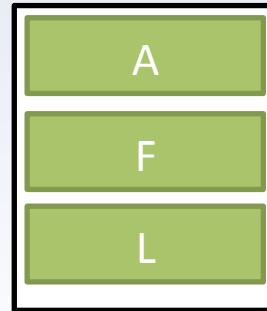
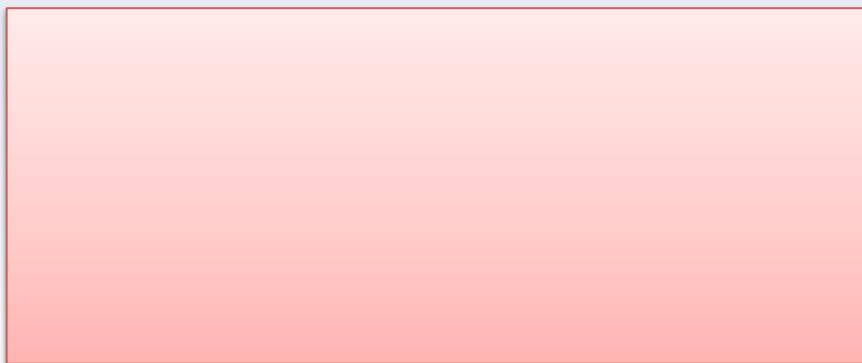
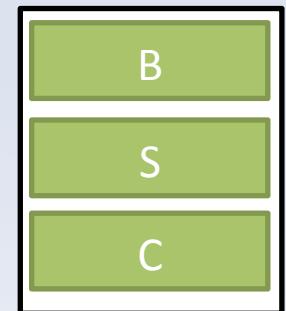
block 3



block 4

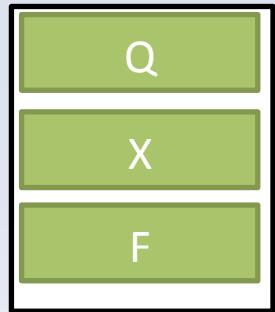


block 5

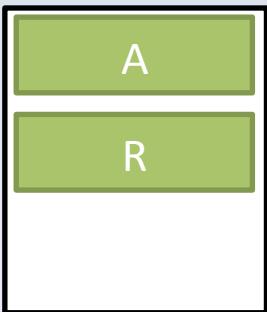


Phase 2

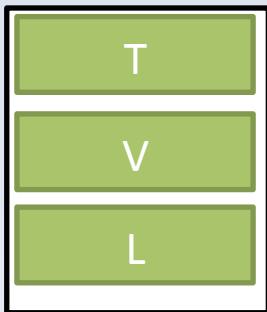
block 0



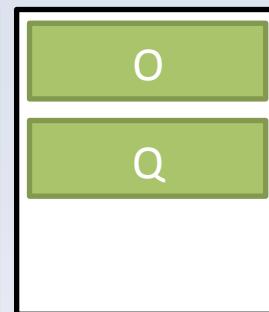
block 1



block 2



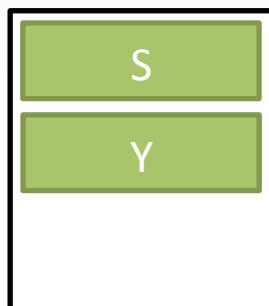
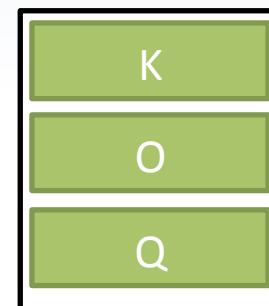
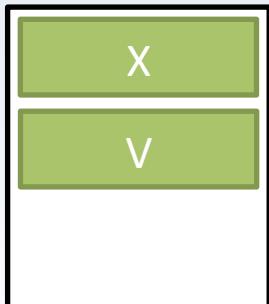
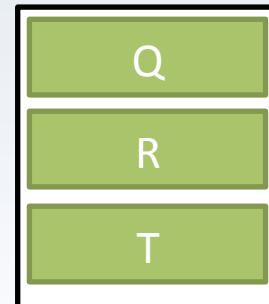
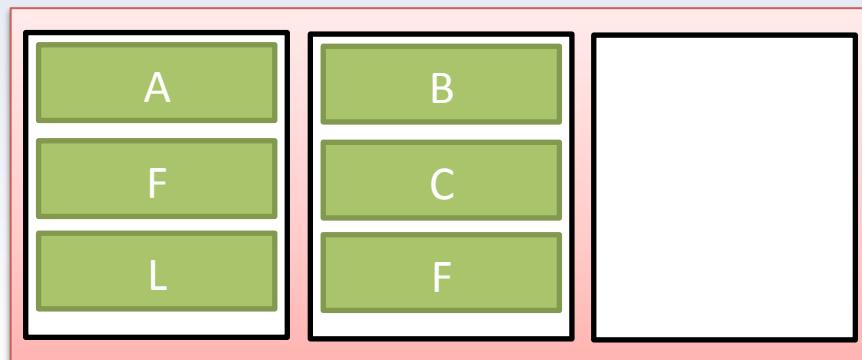
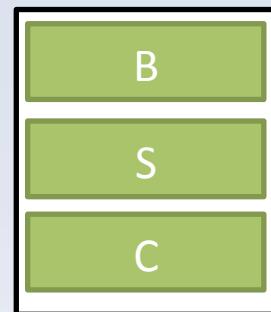
block 3



block 4

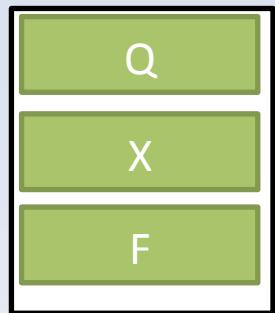


block 5

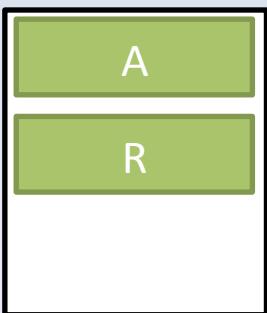


Phase 2

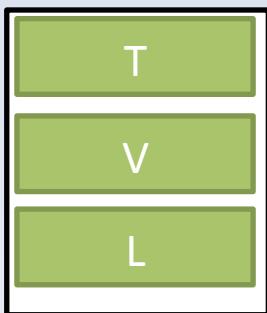
block 0



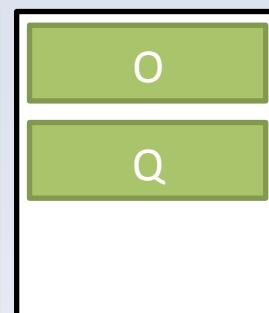
block 1



block 2



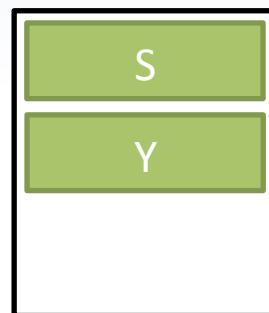
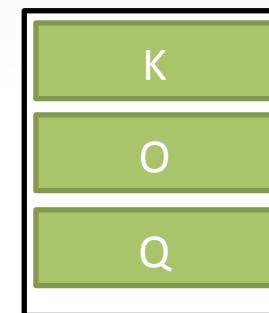
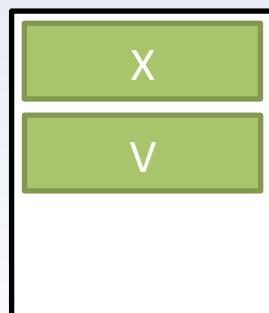
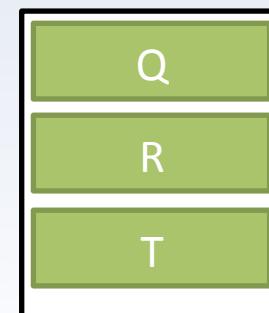
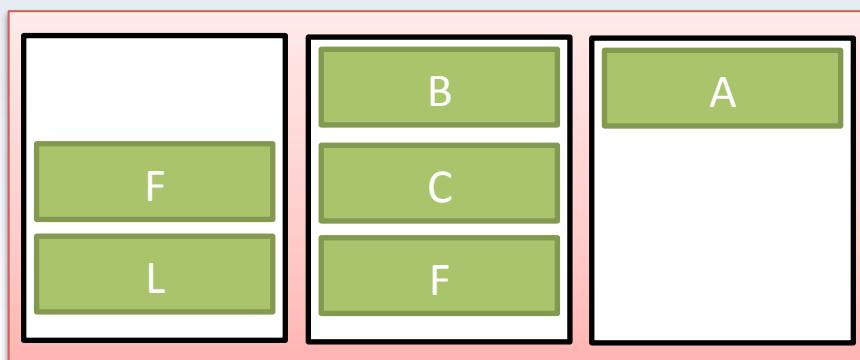
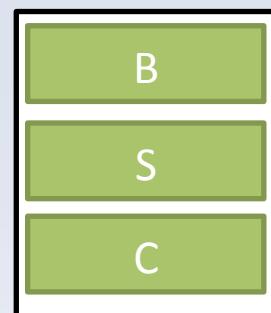
block 3



block 4

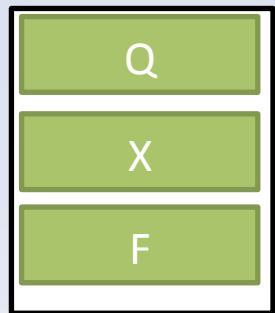


block 5

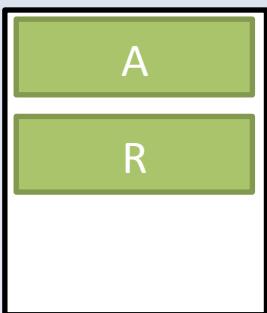


Phase 2

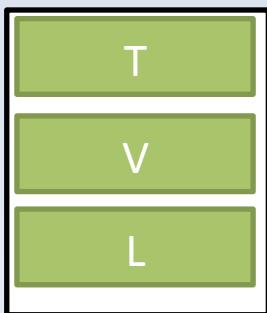
block 0



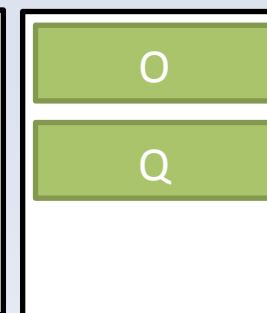
block 1



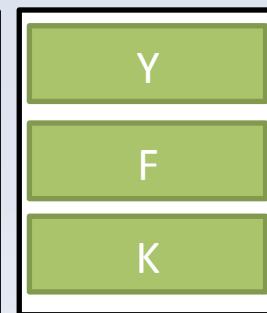
block 2



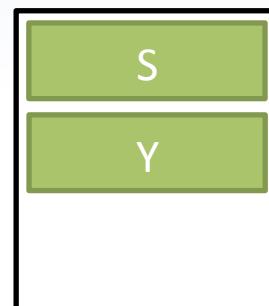
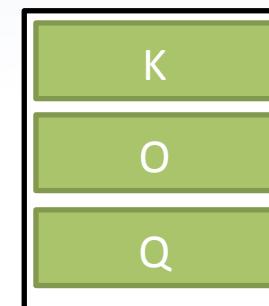
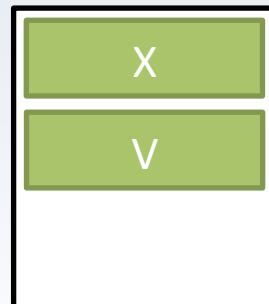
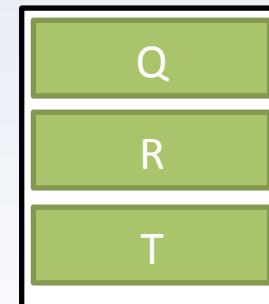
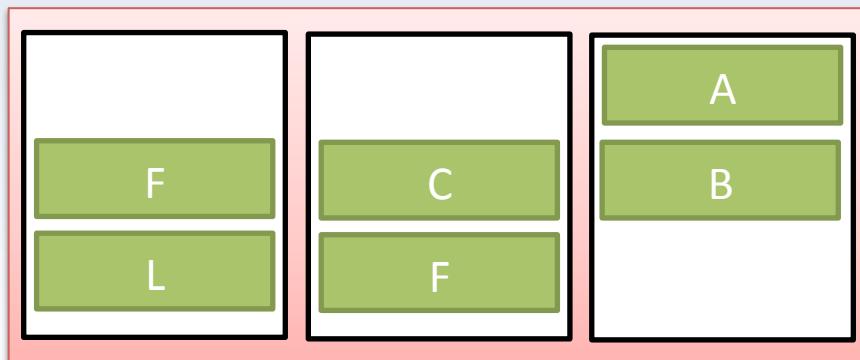
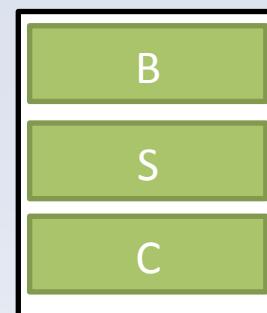
block 3



block 4

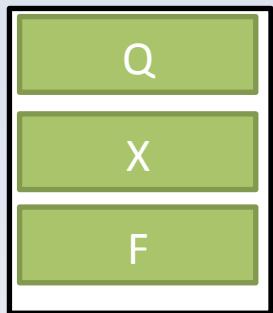


block 5

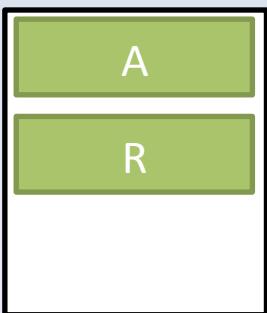


Phase 2

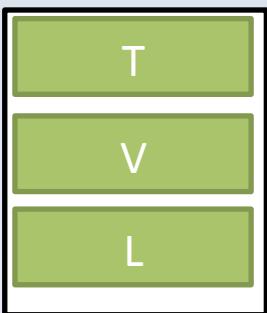
block 0



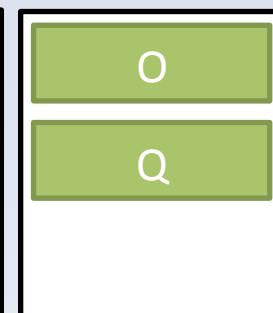
block 1



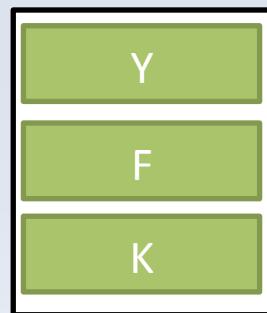
block 2



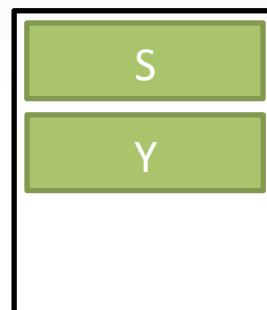
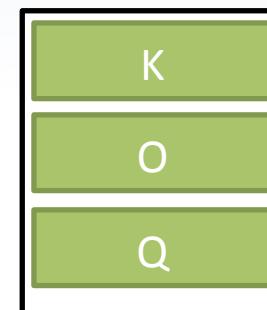
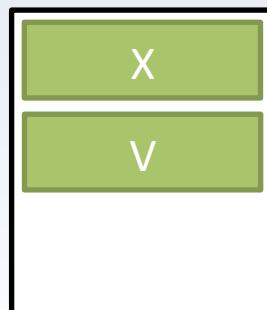
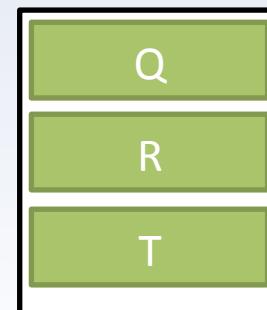
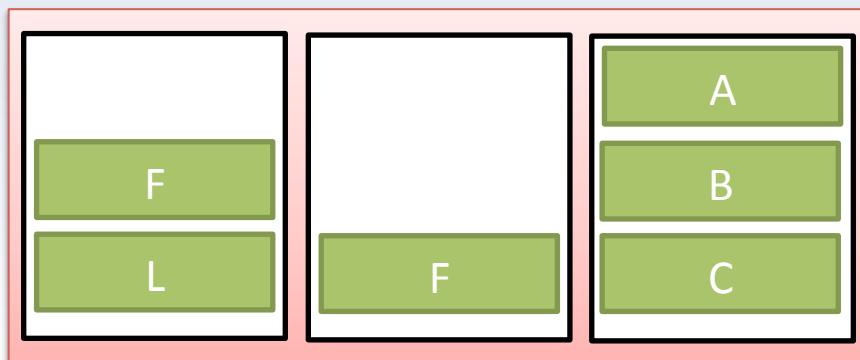
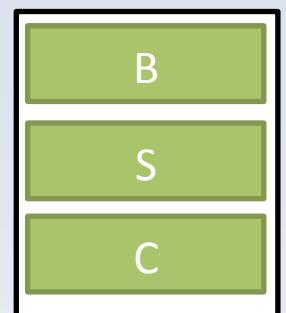
block 3



block 4

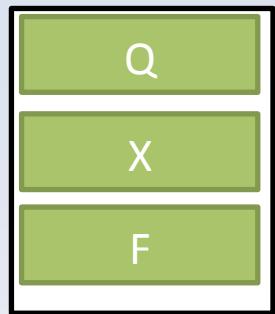


block 5

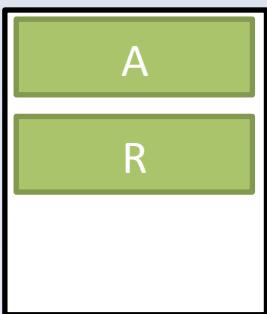


Phase 2

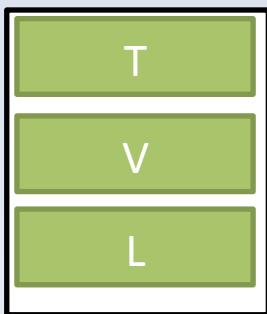
block 0



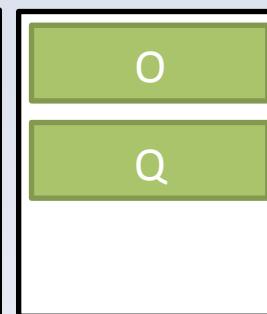
block 1



block 2



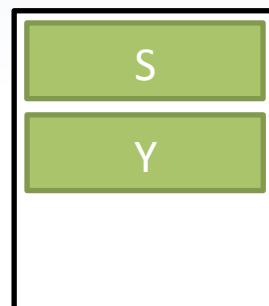
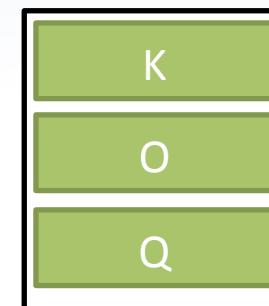
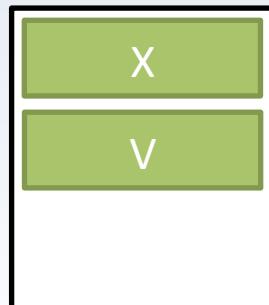
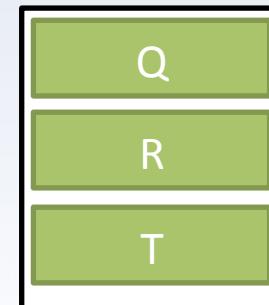
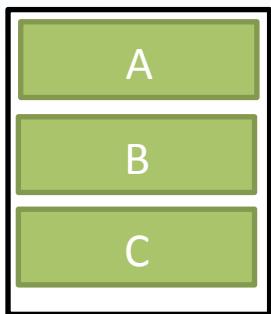
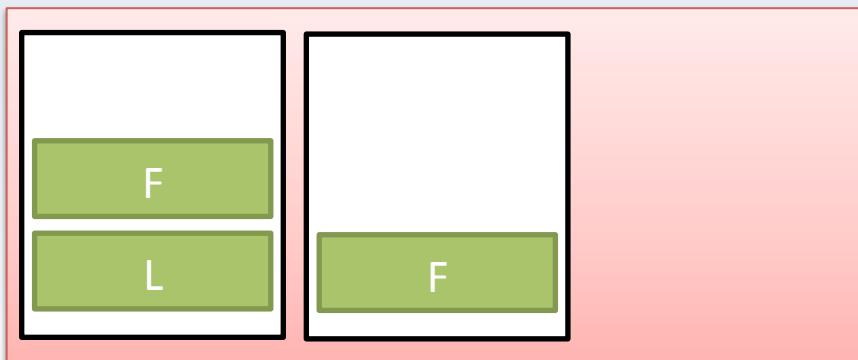
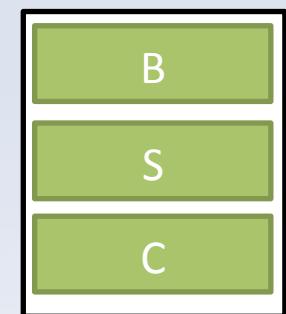
block 3



block 4

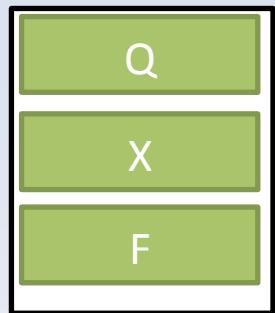


block 5

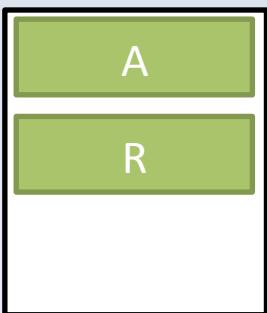


Phase 2

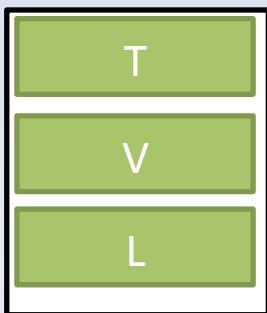
block 0



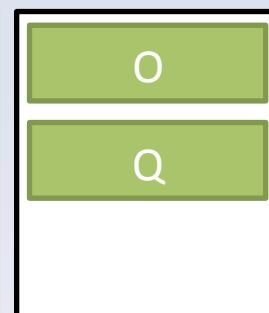
block 1



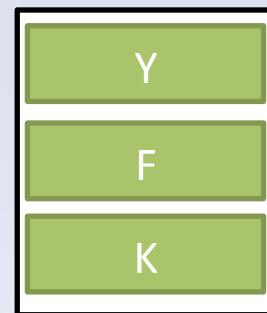
block 2



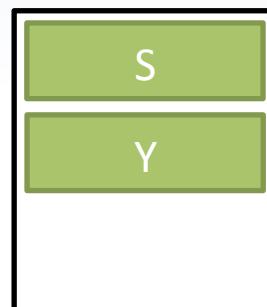
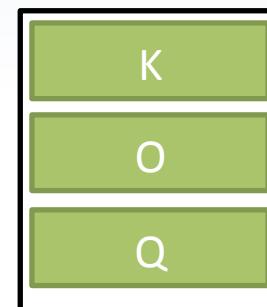
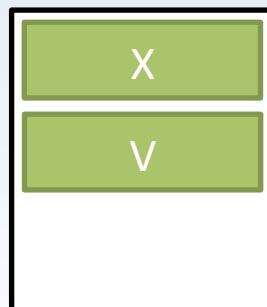
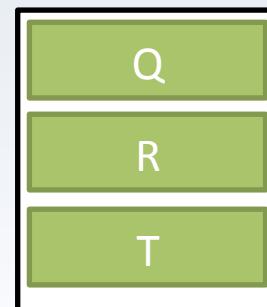
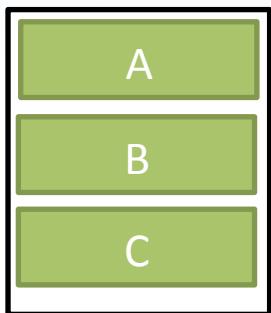
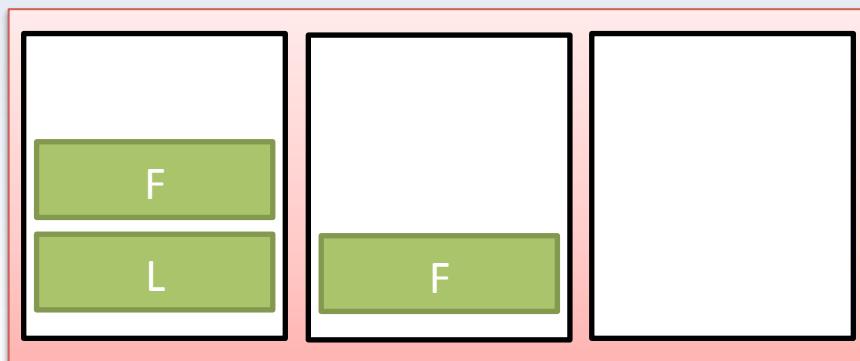
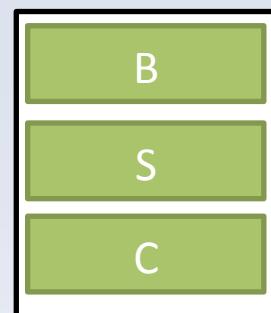
block 3



block 4

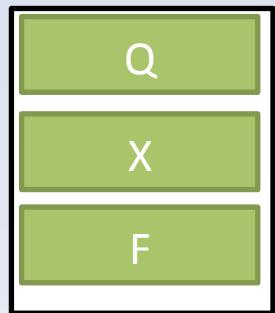


block 5

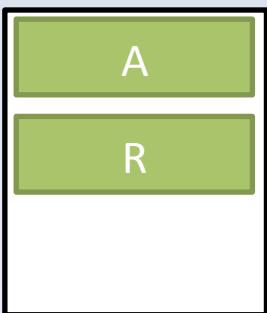


Phase 2

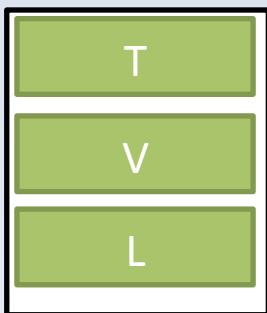
block 0



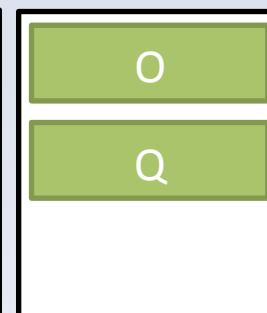
block 1



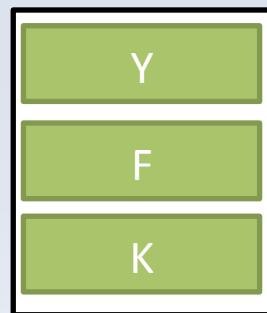
block 2



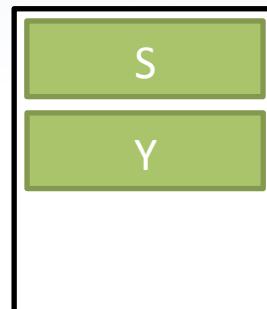
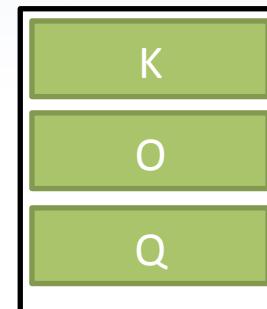
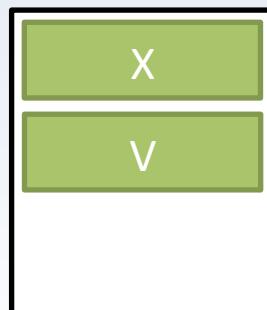
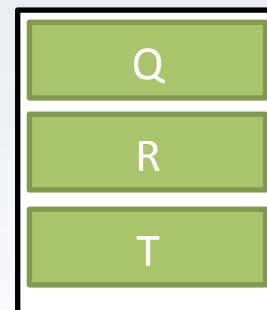
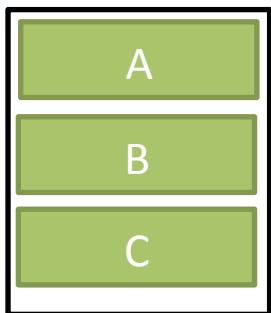
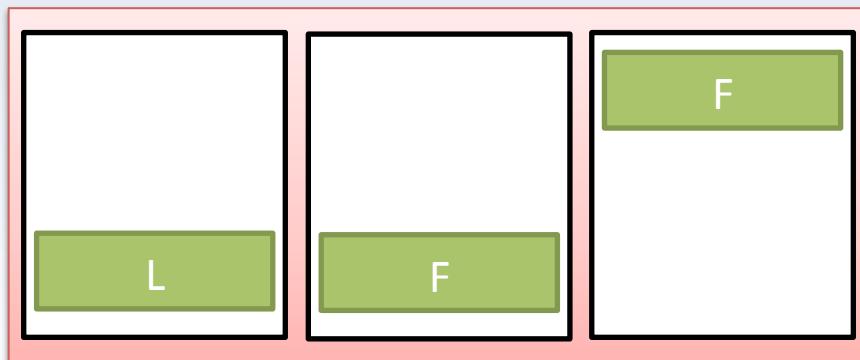
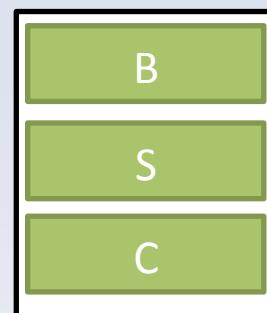
block 3



block 4

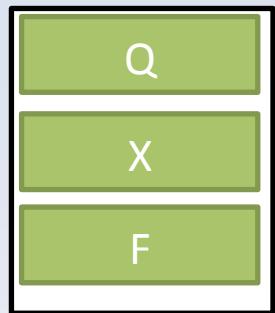


block 5

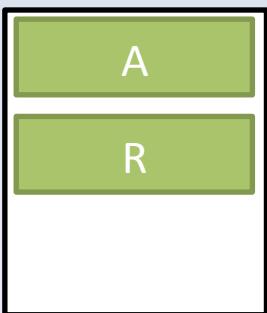


Phase 2

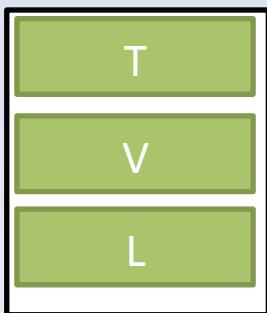
block 0



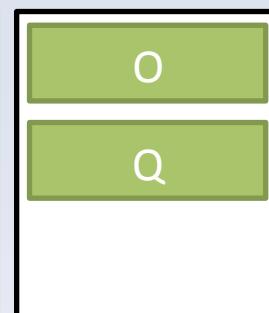
block 1



block 2



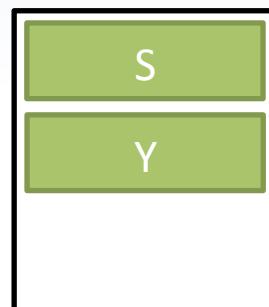
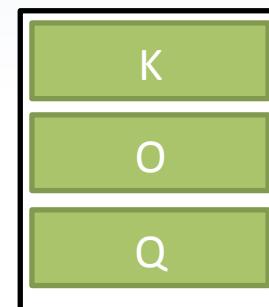
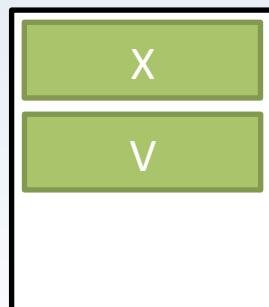
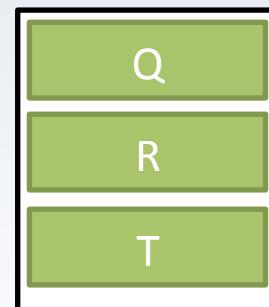
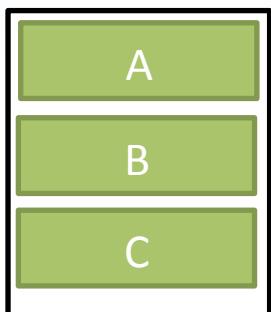
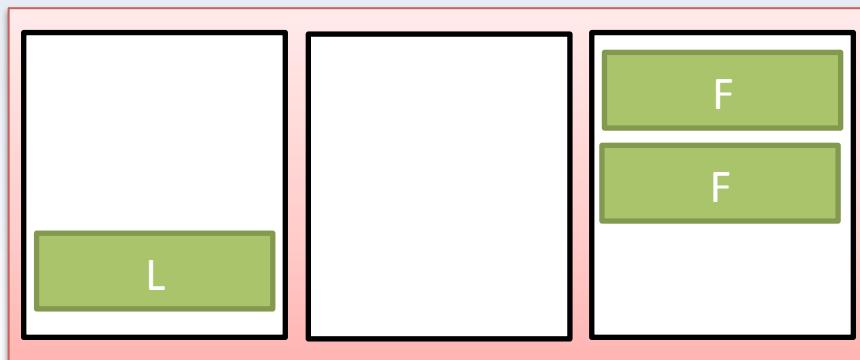
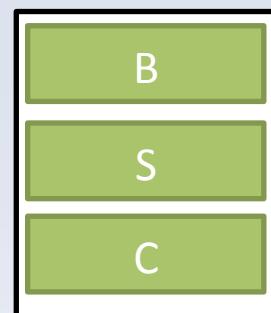
block 3



block 4

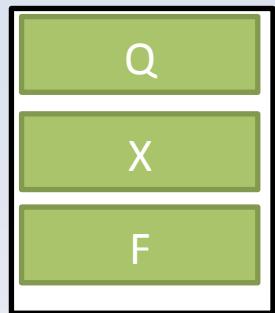


block 5

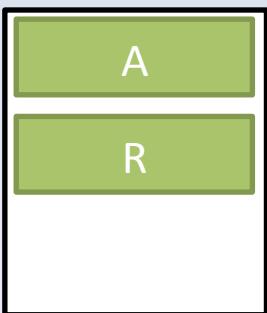


Phase 2

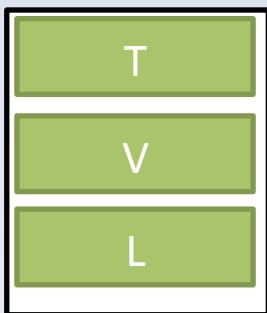
block 0



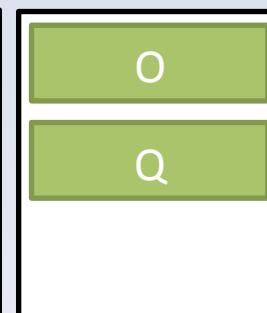
block 1



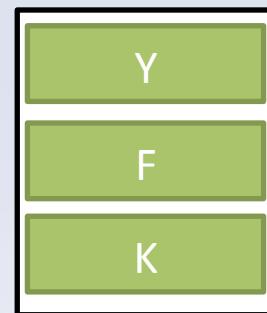
block 2



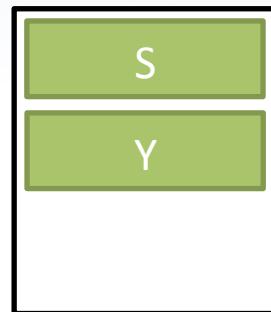
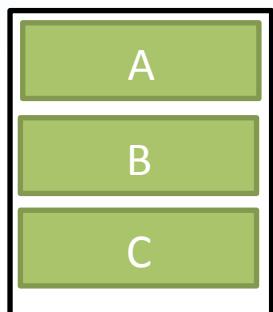
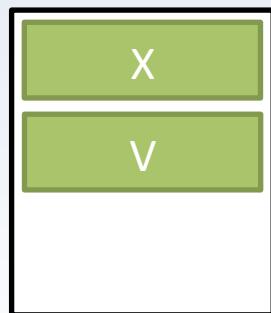
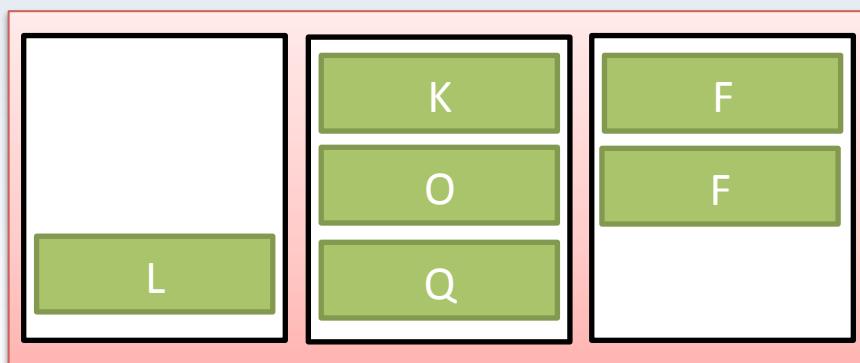
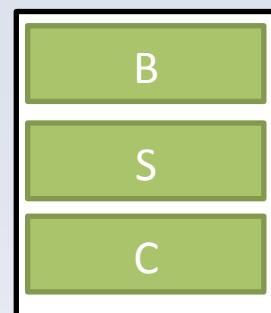
block 3



block 4

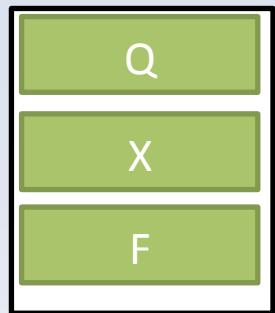


block 5

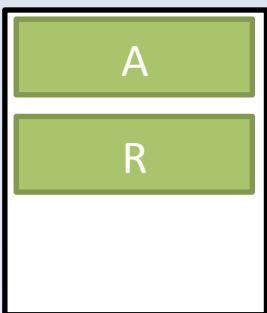


Phase 2

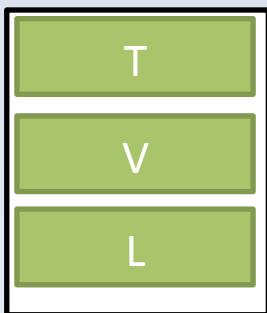
block 0



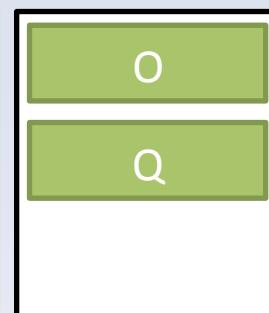
block 1



block 2



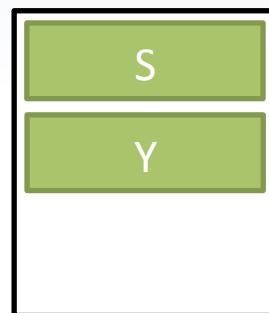
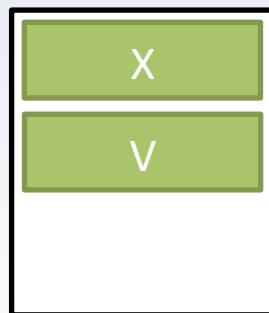
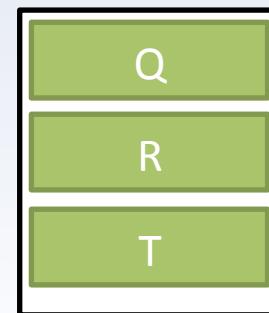
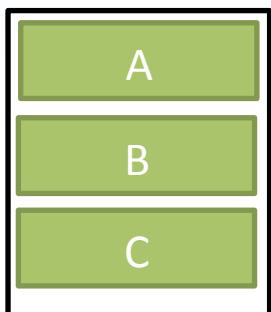
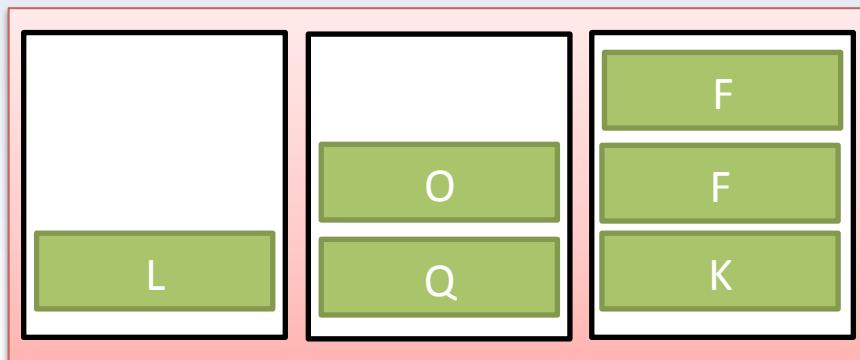
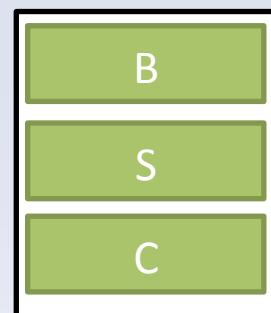
block 3



block 4

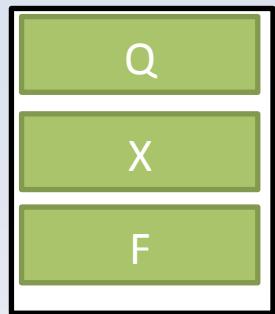


block 5

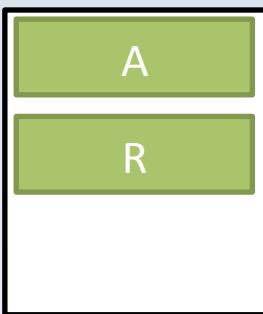


Phase 2

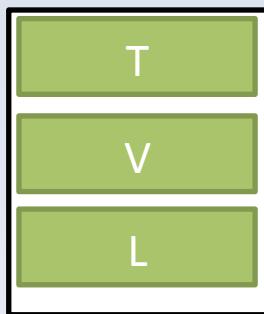
block 0



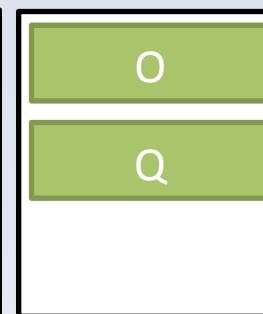
block 1



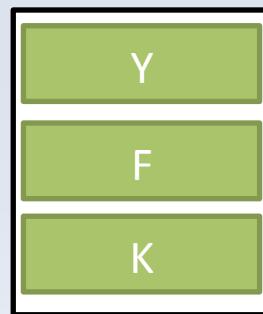
block 2



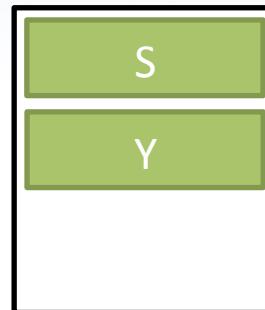
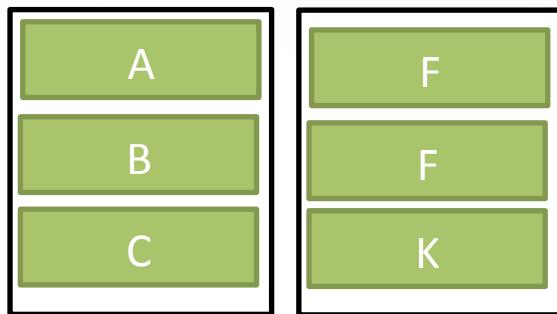
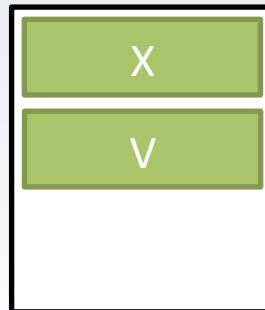
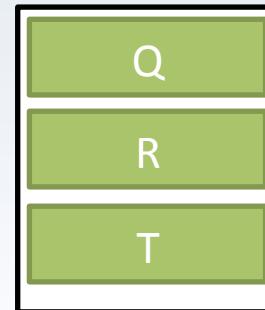
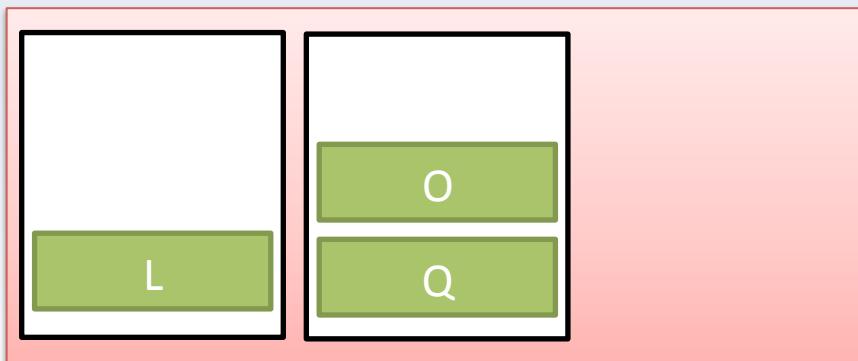
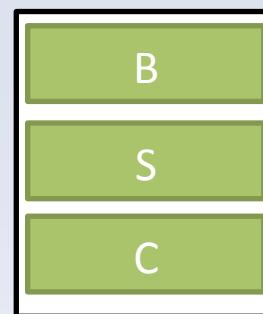
block 3



block 4

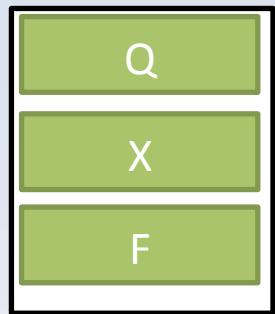


block 5

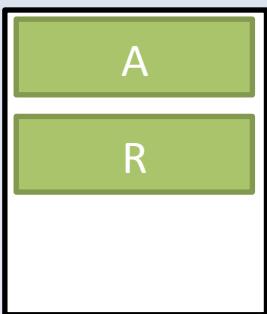


Phase 2

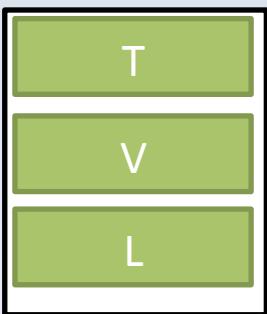
block 0



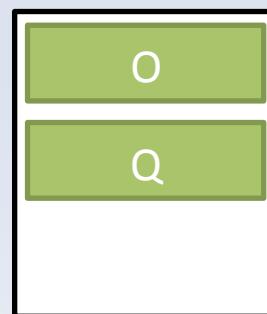
block 1



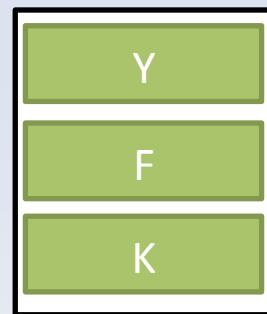
block 2



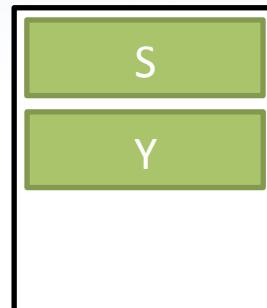
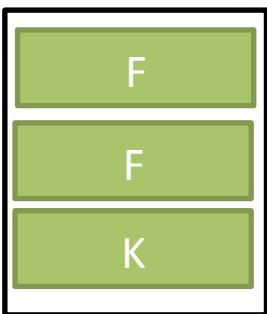
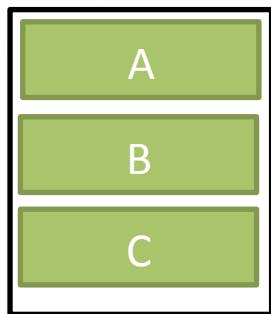
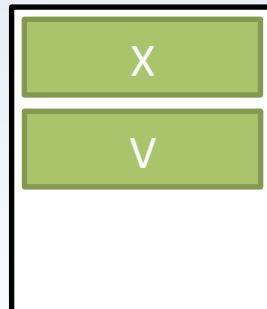
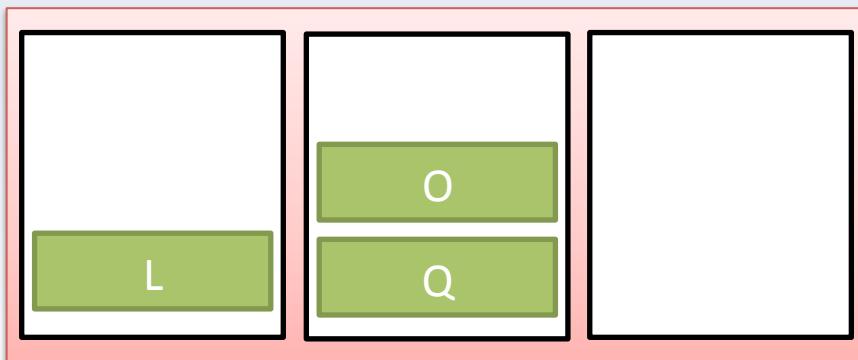
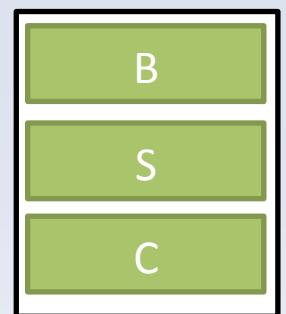
block 3



block 4

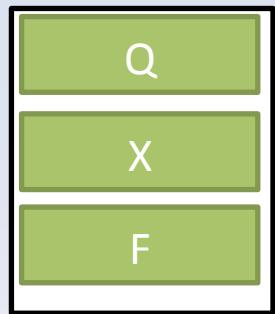


block 5

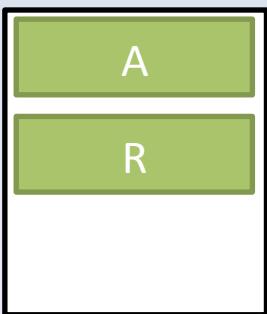


Phase 2

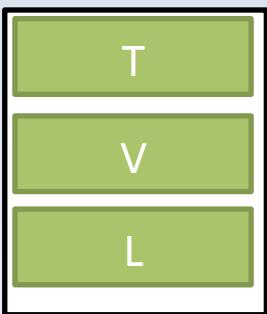
block 0



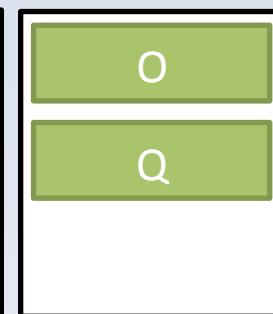
block 1



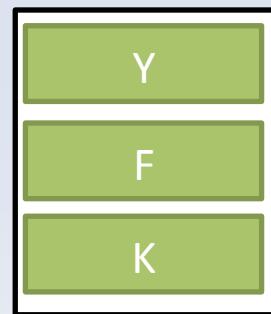
block 2



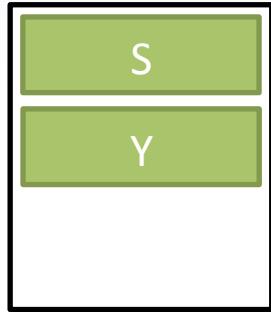
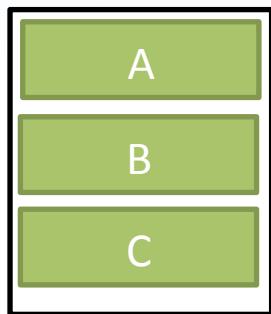
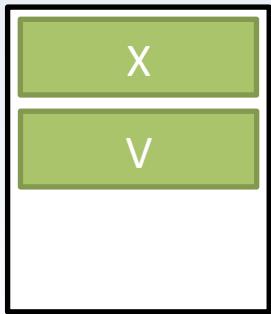
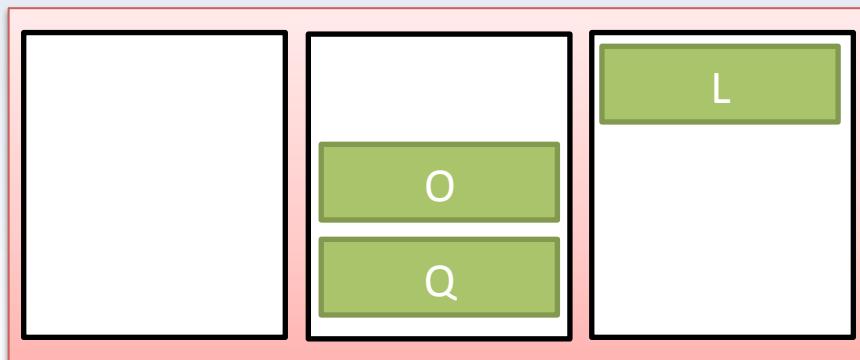
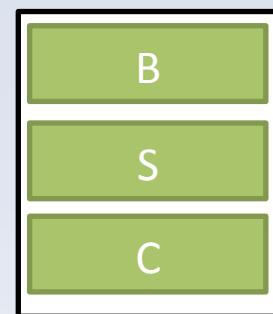
block 3



block 4

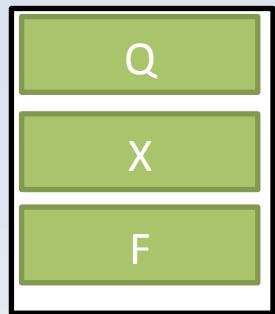


block 5

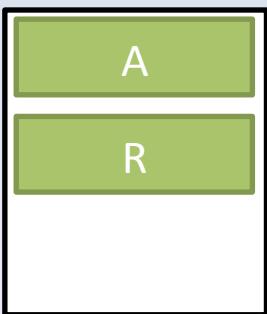


Phase 2

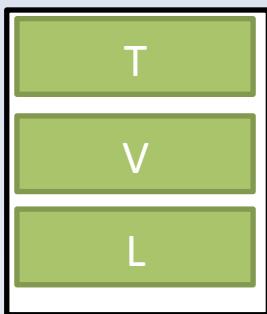
block 0



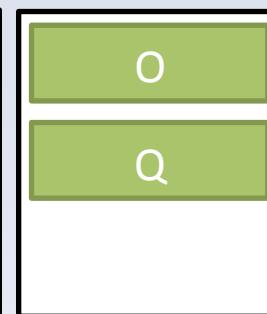
block 1



block 2



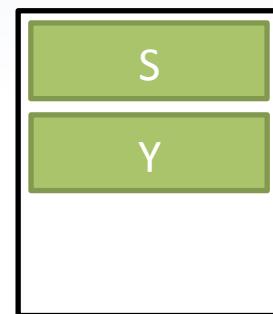
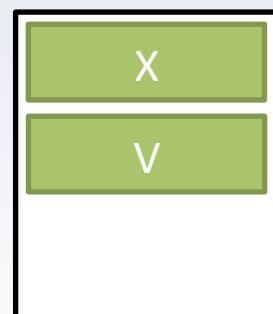
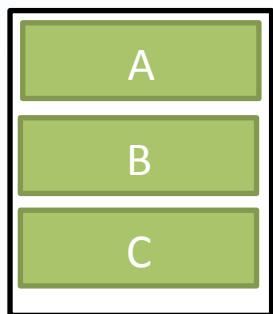
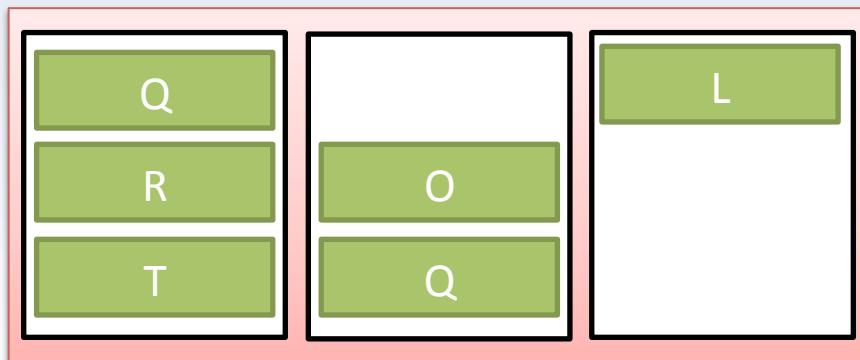
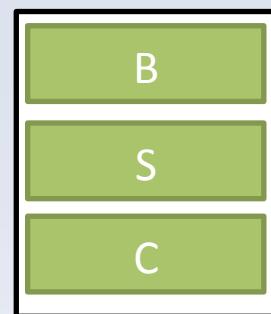
block 3



block 4

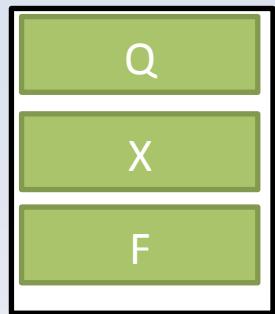


block 5

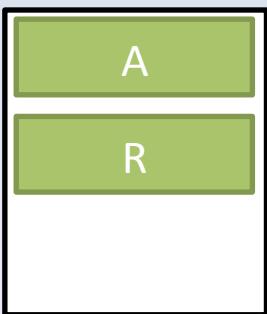


Phase 2

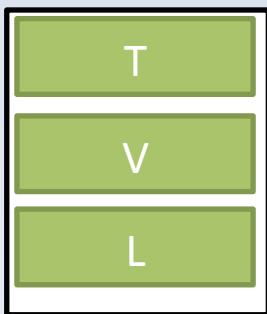
block 0



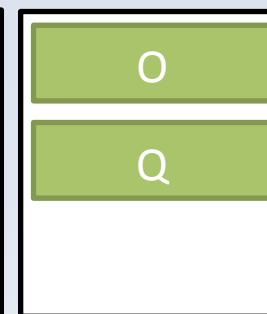
block 1



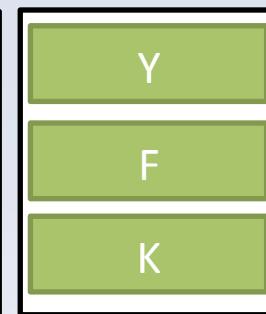
block 2



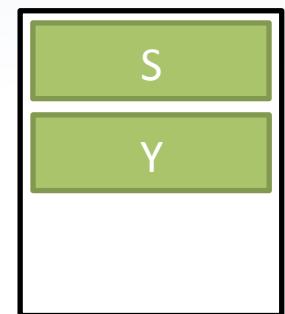
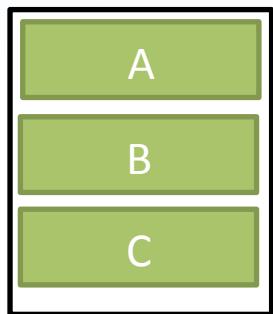
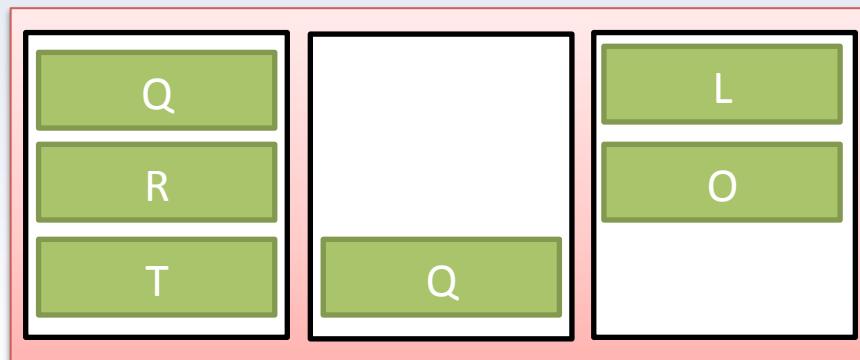
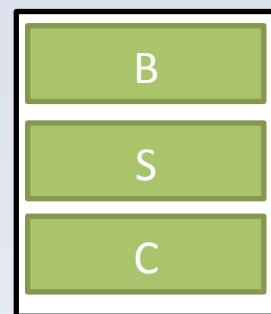
block 3



block 4

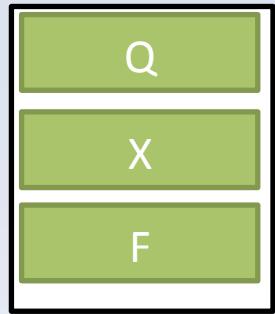


block 5

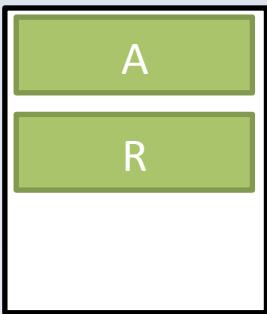


Phase 2

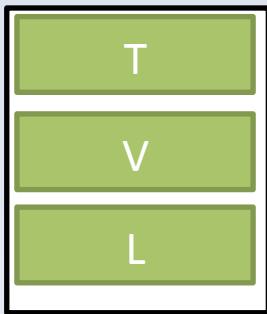
block 0



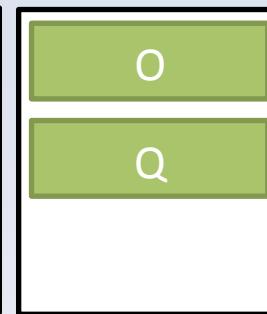
block 1



block 2



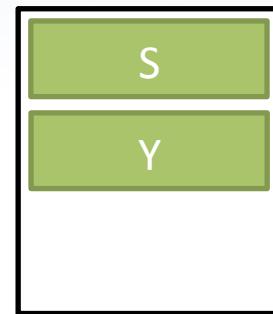
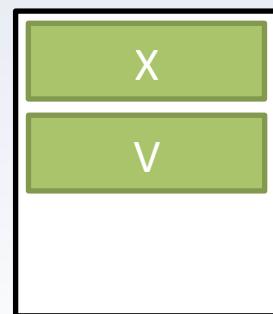
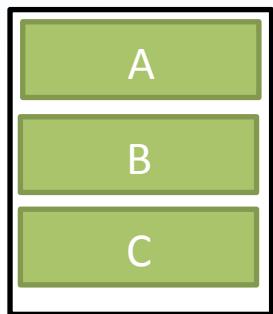
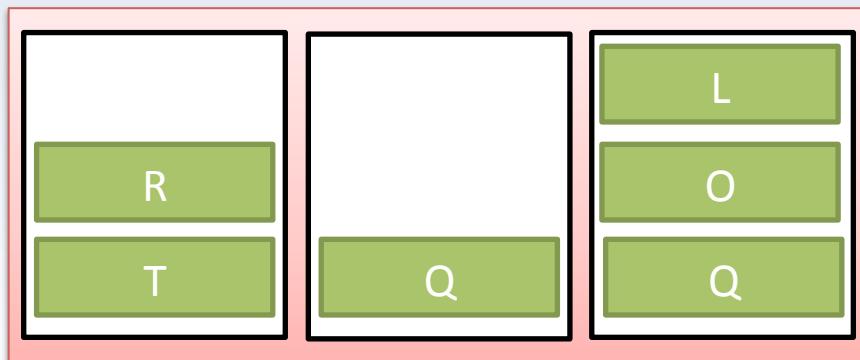
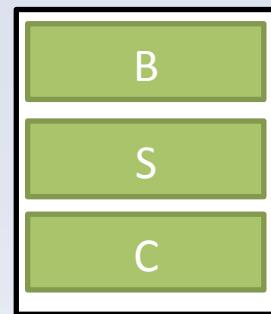
block 3



block 4

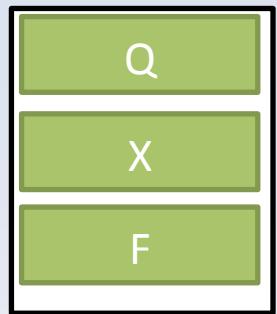


block 5

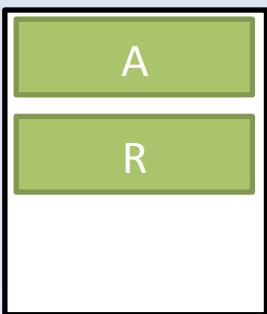


Phase 2

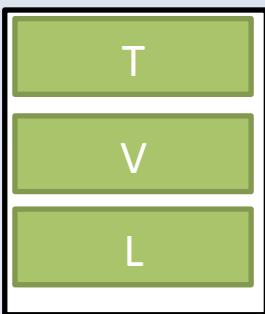
block 0



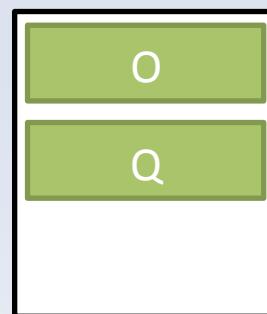
block 1



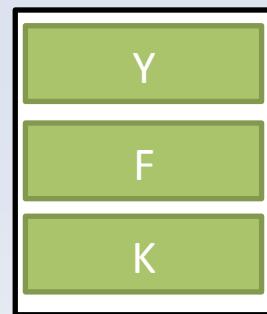
block 2



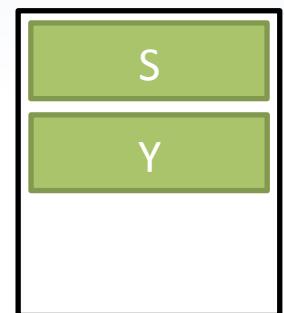
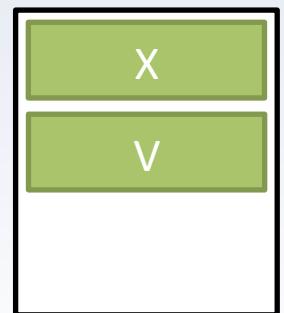
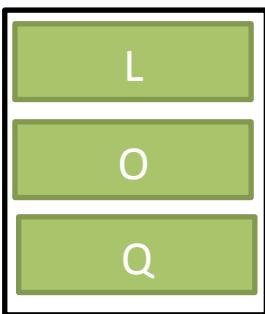
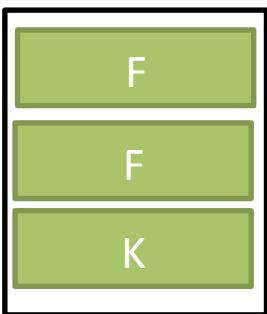
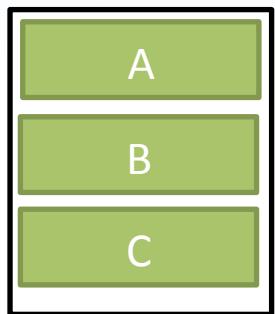
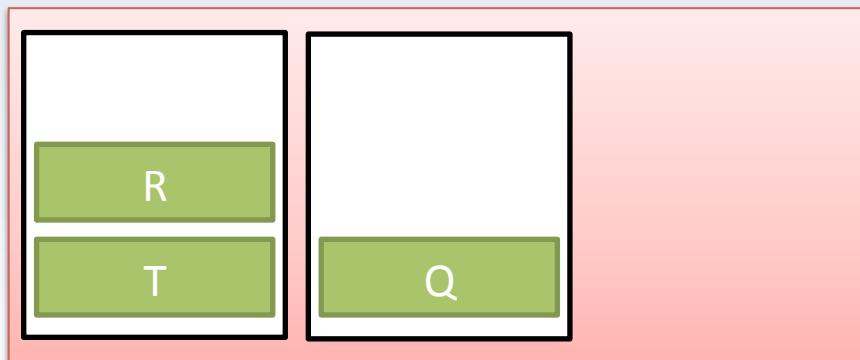
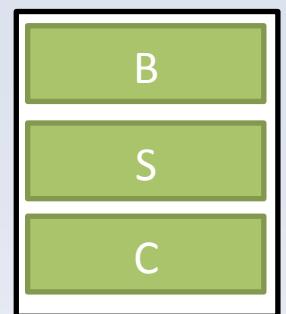
block 3



block 4

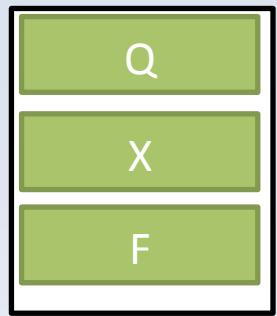


block 5

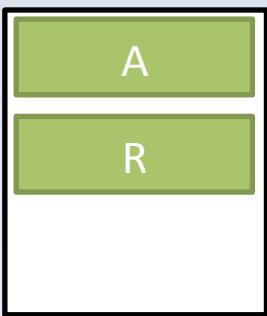


Phase 2

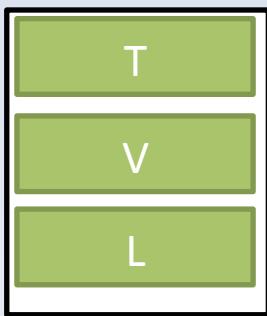
block 0



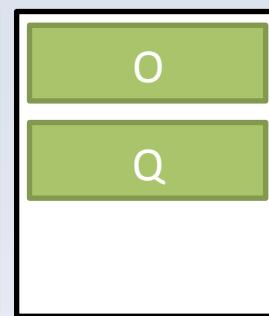
block 1



block 2



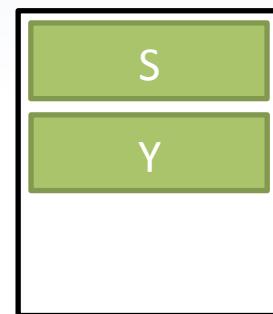
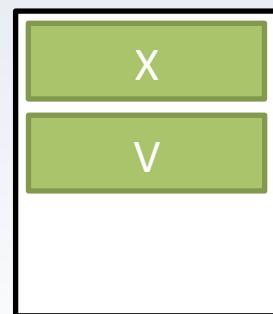
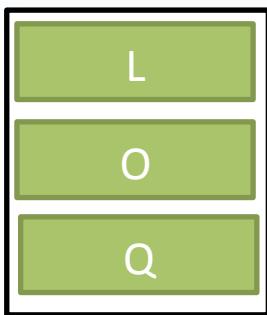
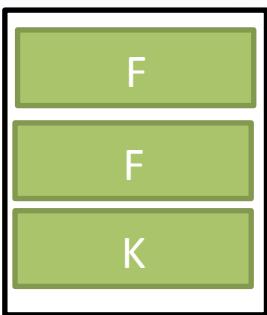
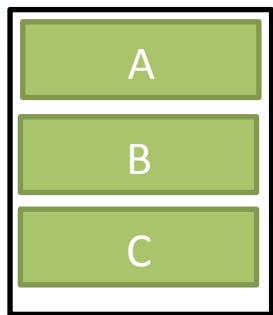
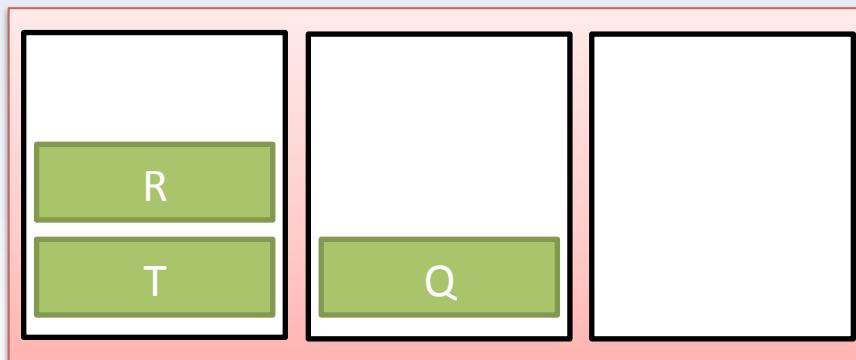
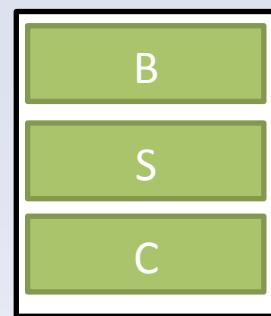
block 3



block 4

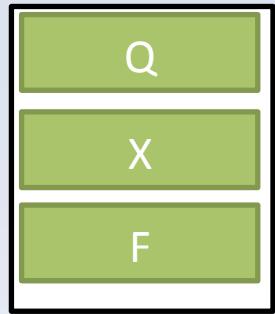


block 5

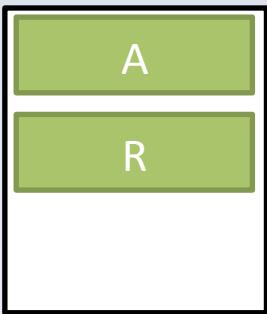


Phase 2

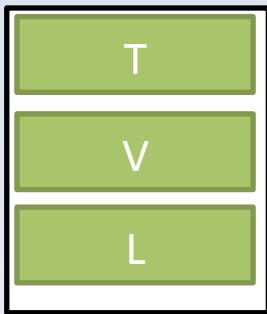
block 0



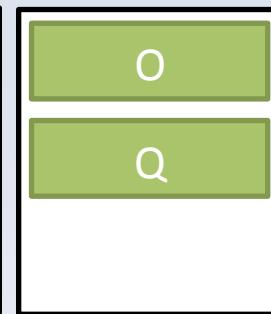
block 1



block 2



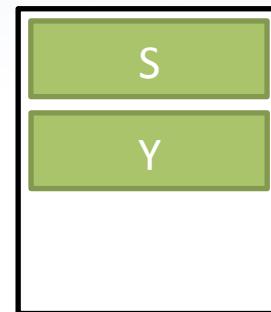
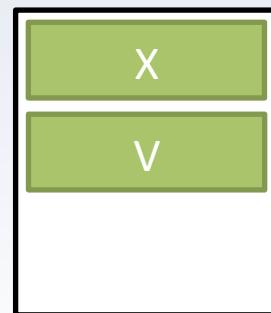
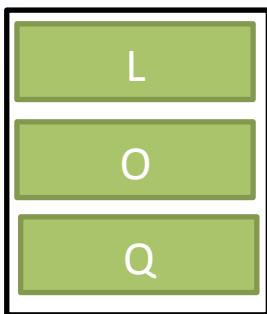
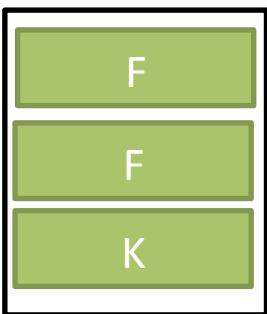
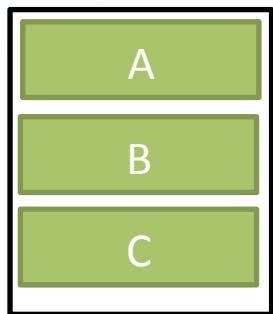
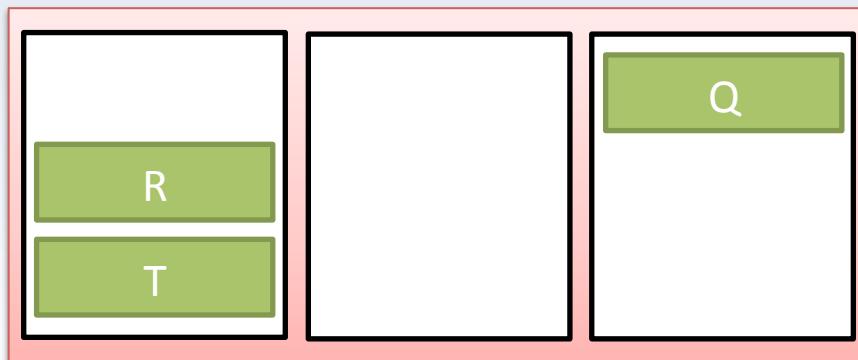
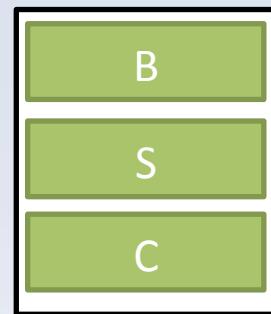
block 3



block 4

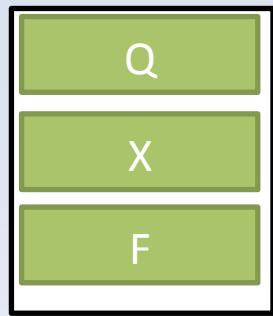


block 5

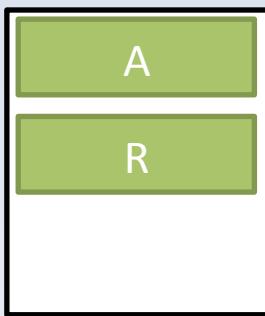


Phase 2

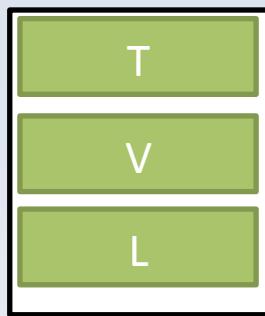
block 0



block 1



block 2



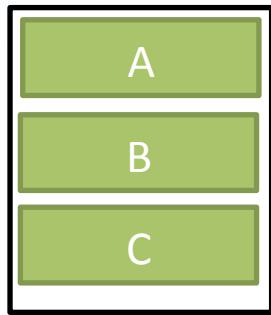
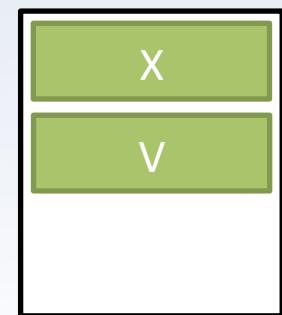
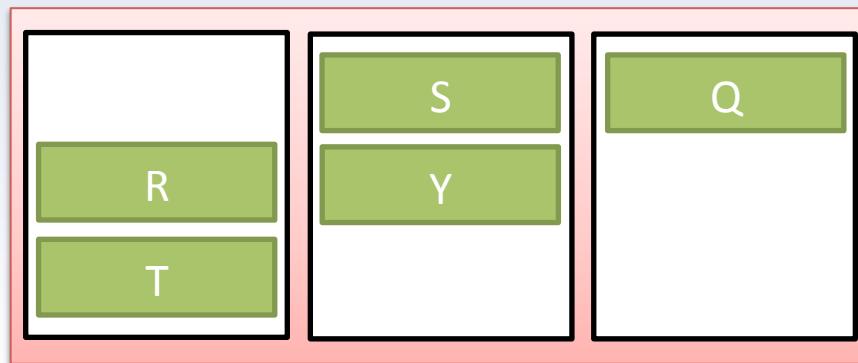
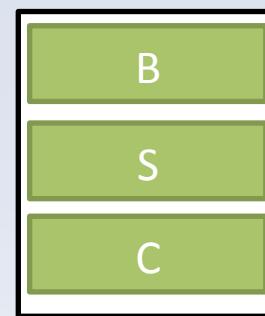
block 3



block 4

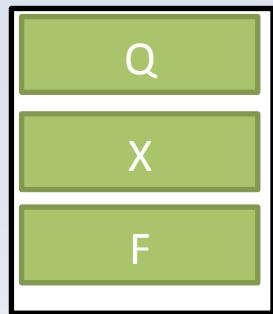


block 5

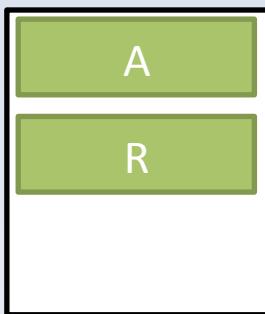


Phase 2

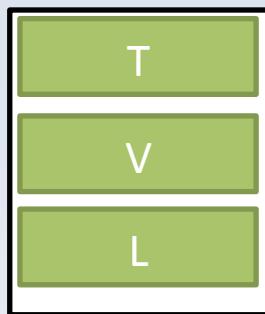
block 0



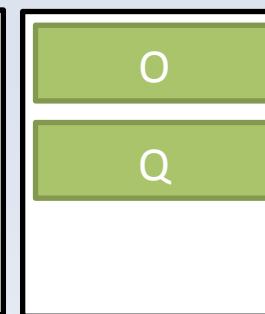
block 1



block 2



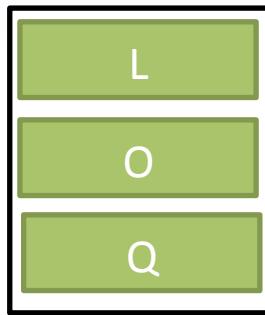
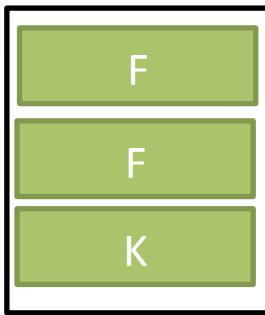
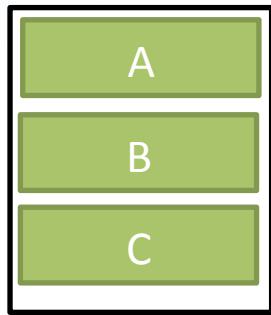
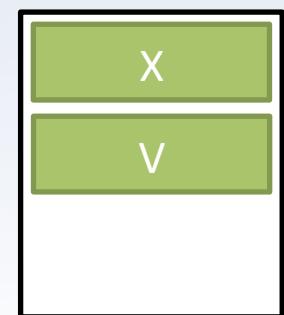
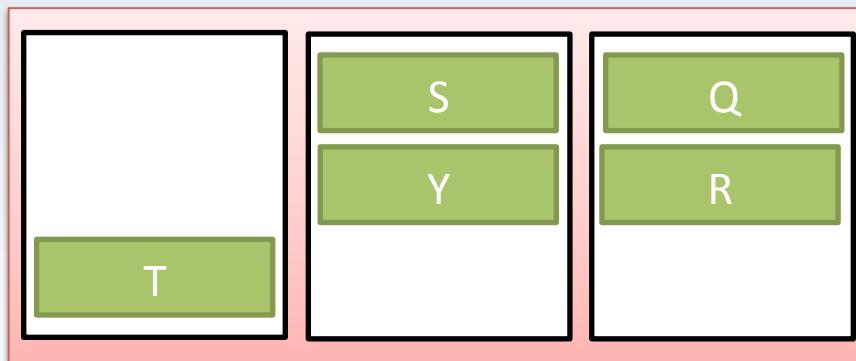
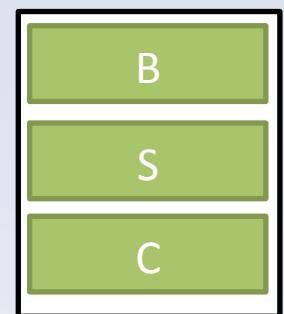
block 3



block 4

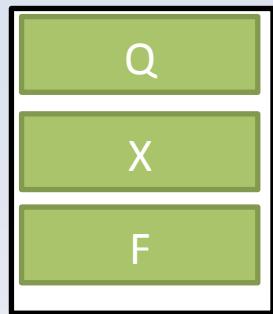


block 5

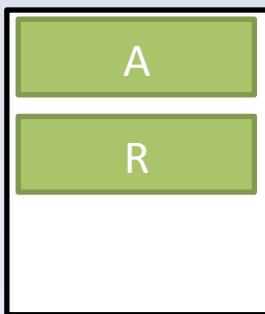


Phase 2

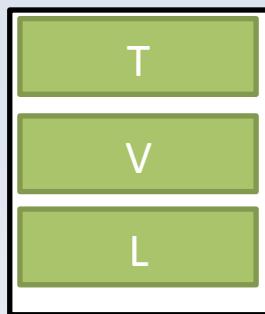
block 0



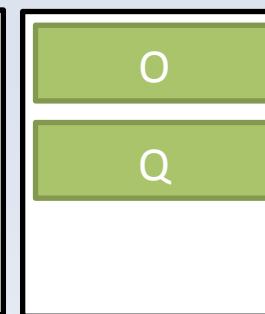
block 1



block 2



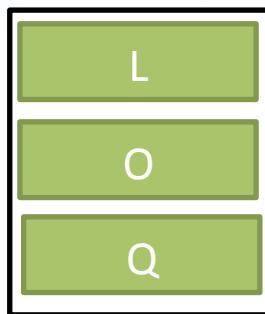
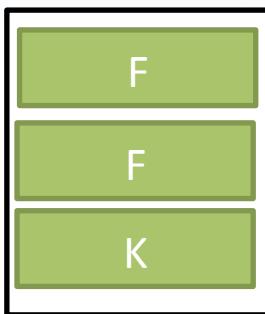
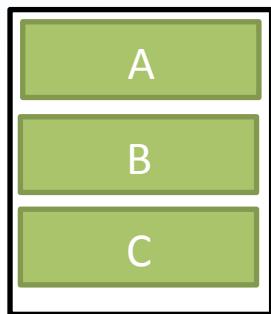
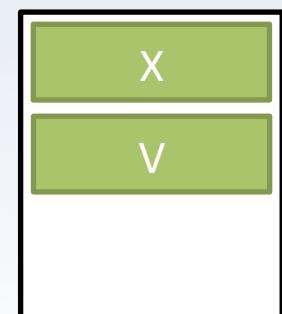
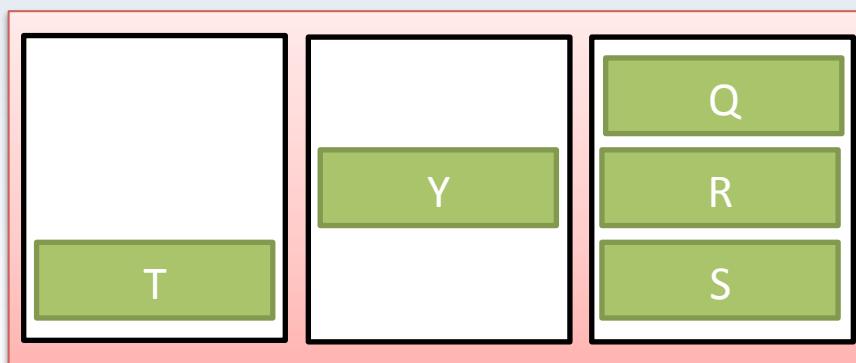
block 3



block 4

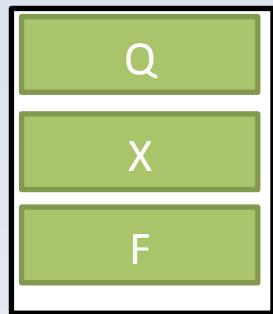


block 5

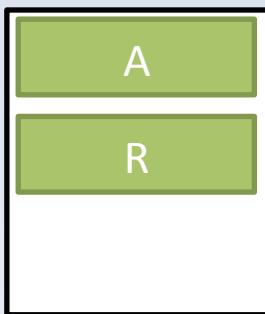


Phase 2

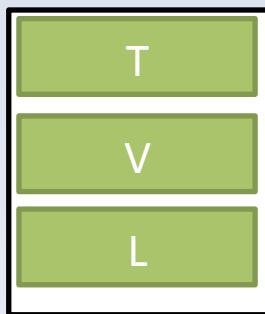
block 0



block 1



block 2



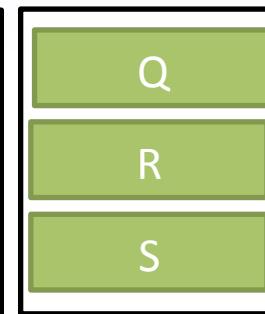
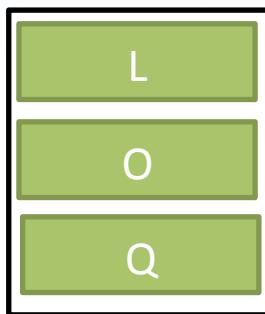
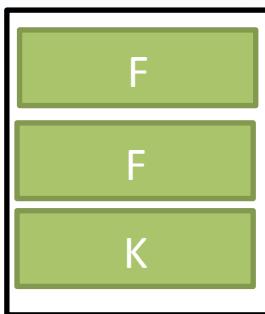
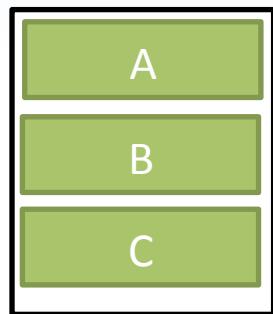
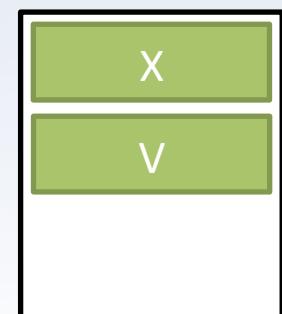
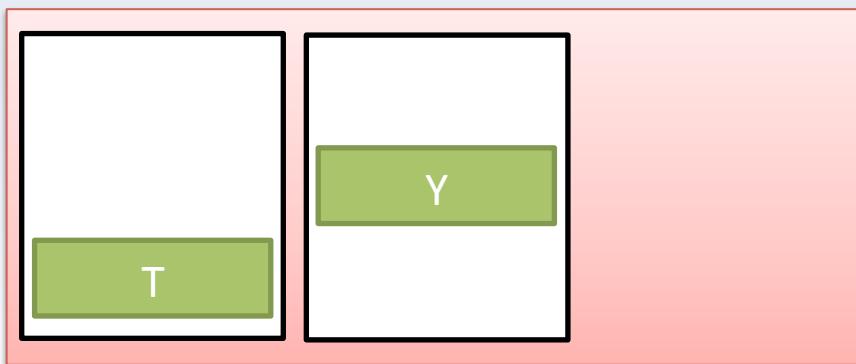
block 3



block 4

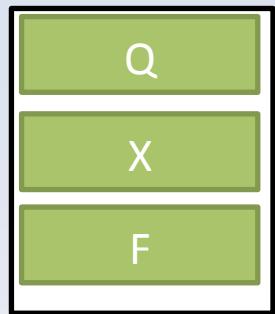


block 5

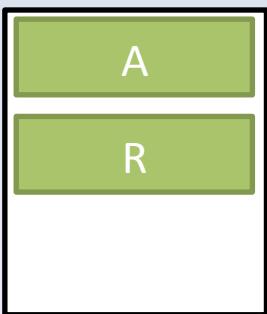


Phase 2

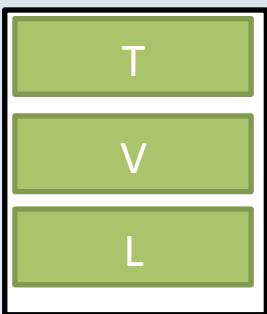
block 0



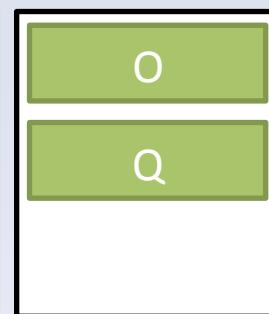
block 1



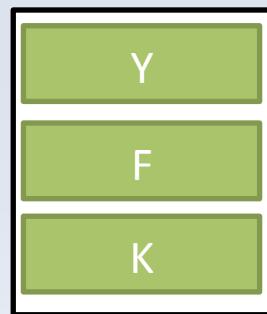
block 2



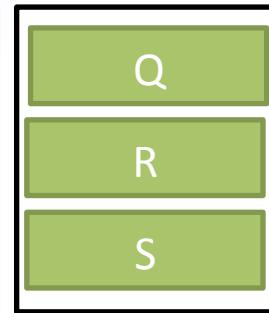
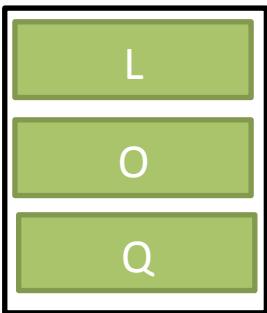
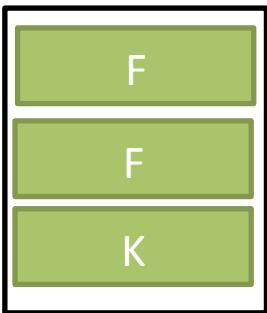
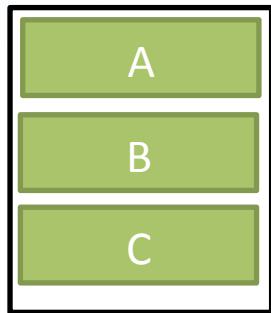
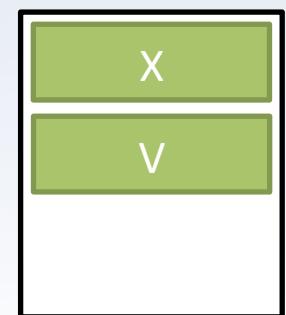
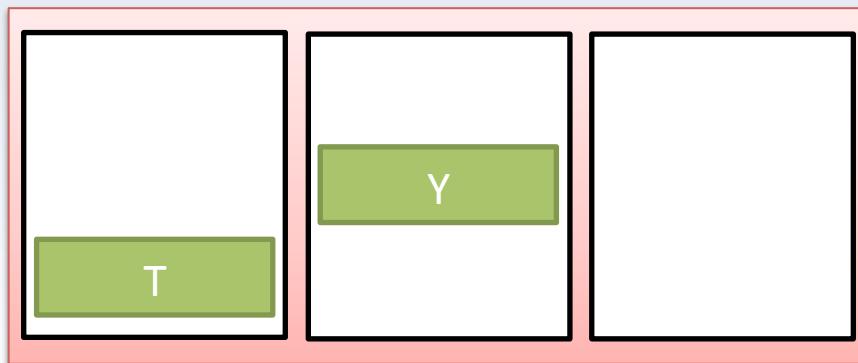
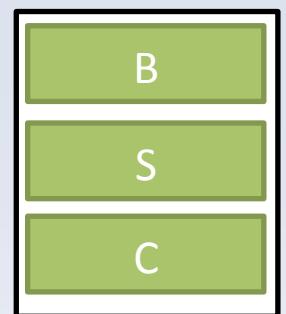
block 3



block 4

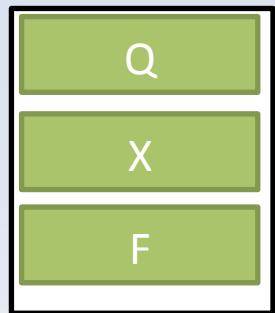


block 5

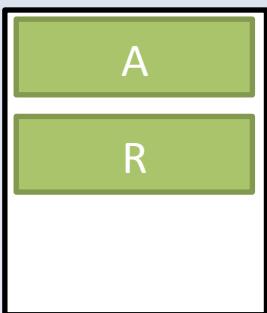


Phase 2

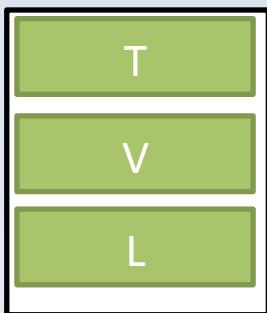
block 0



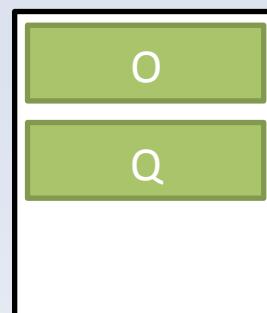
block 1



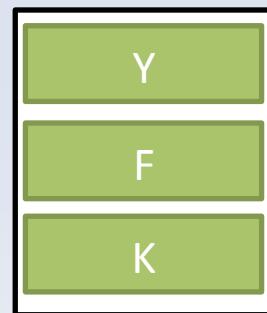
block 2



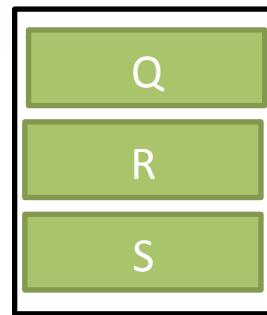
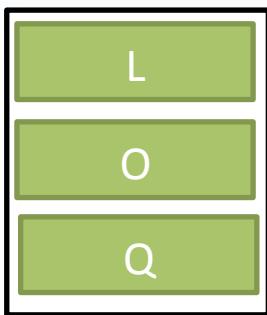
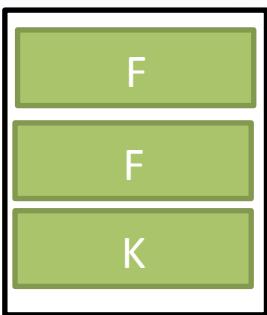
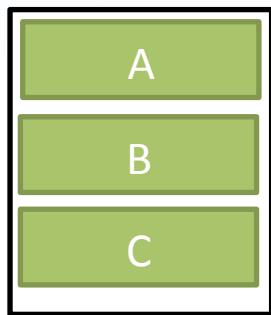
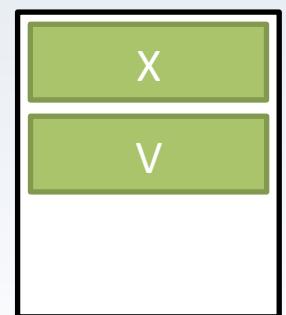
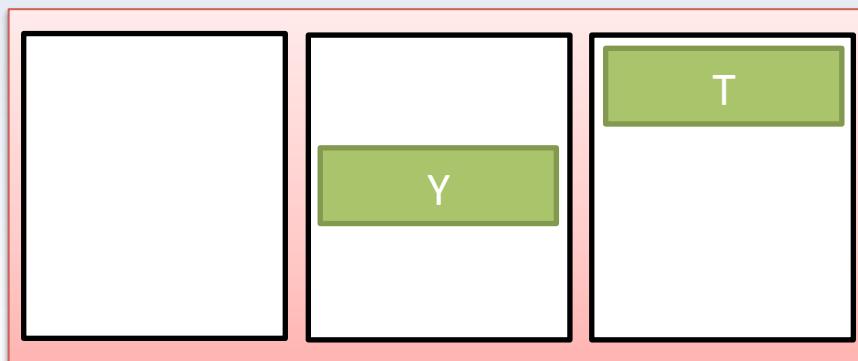
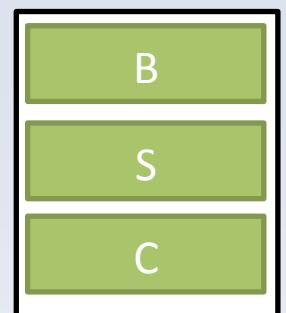
block 3



block 4

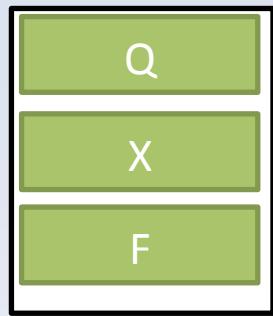


block 5

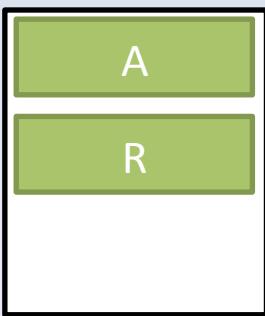


Phase 2

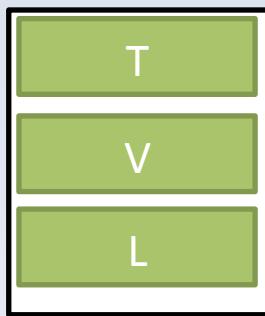
block 0



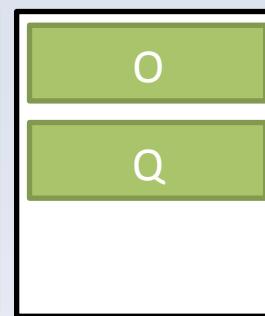
block 1



block 2



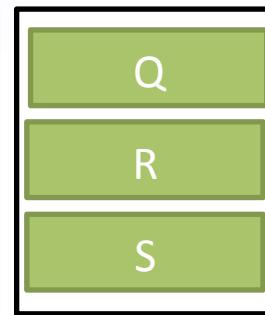
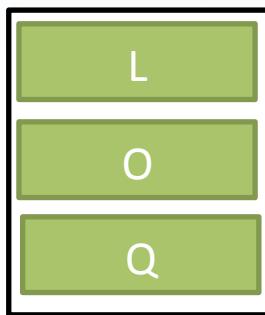
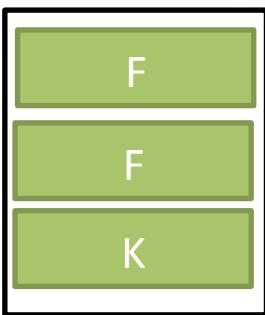
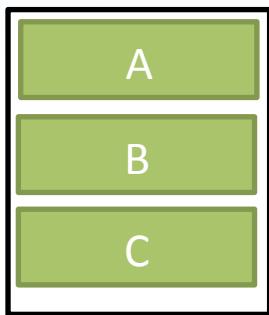
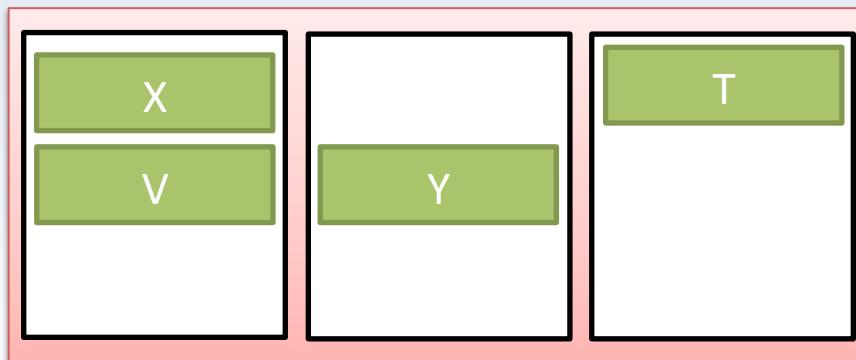
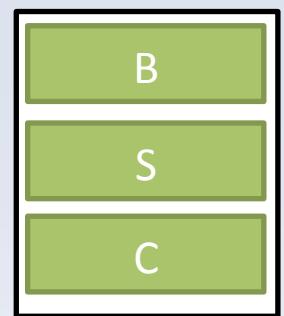
block 3



block 4

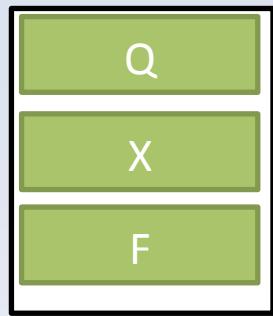


block 5

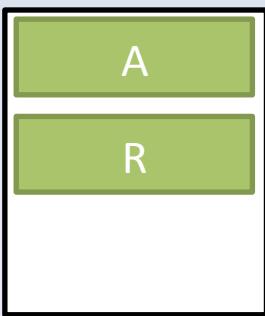


Phase 2

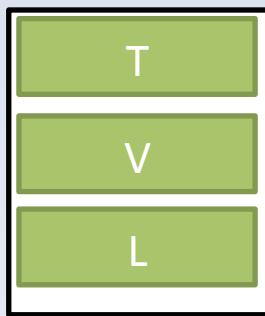
block 0



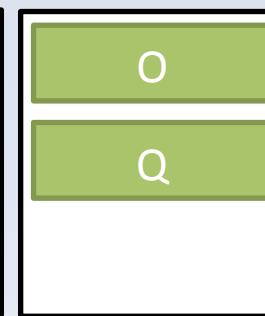
block 1



block 2



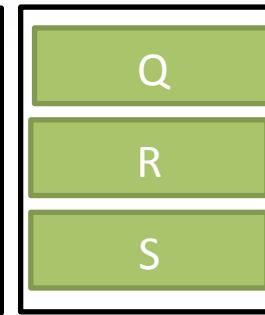
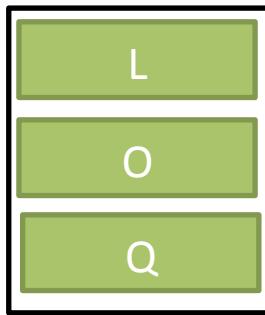
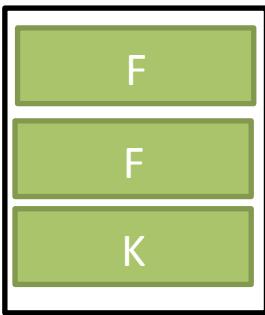
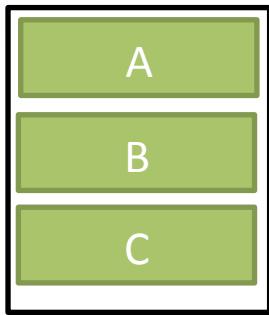
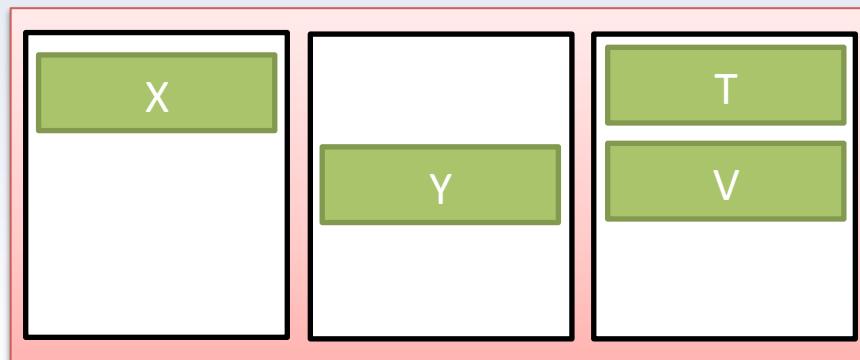
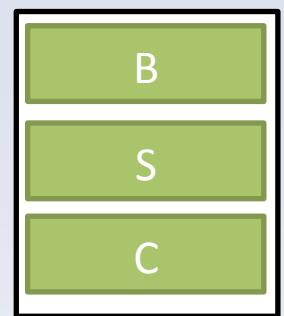
block 3



block 4

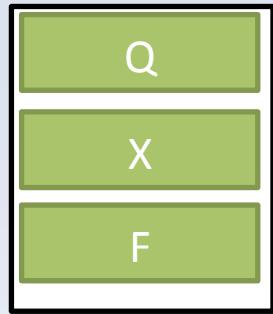


block 5

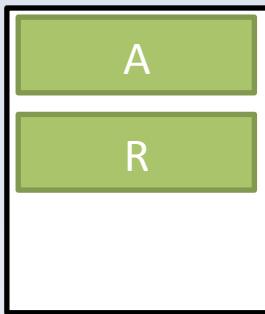


Phase 2

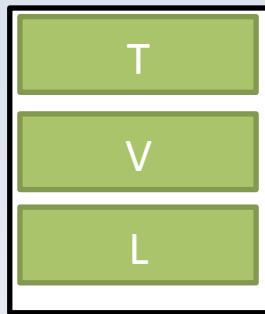
block 0



block 1



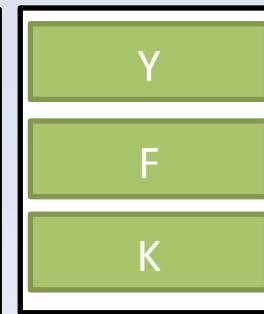
block 2



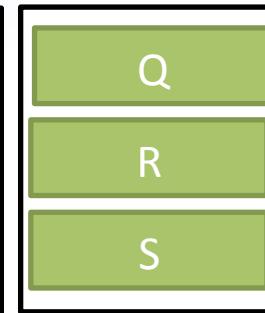
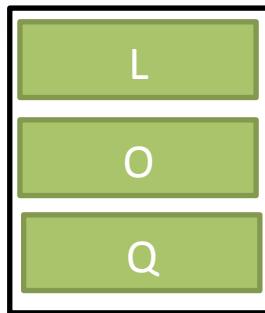
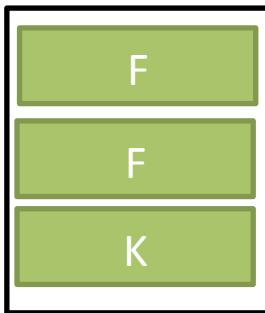
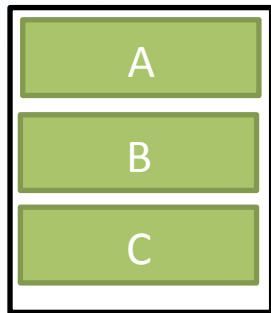
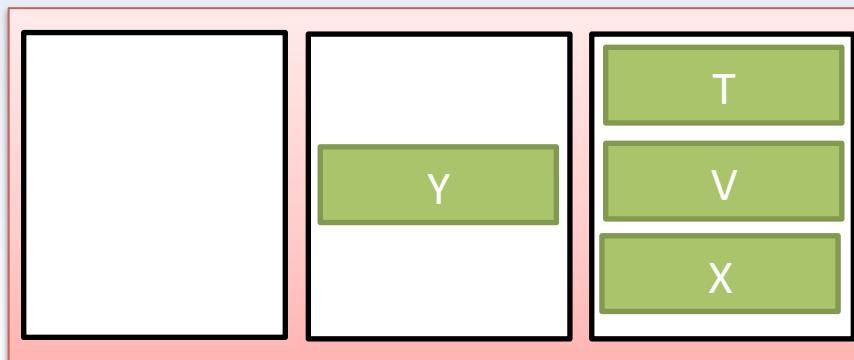
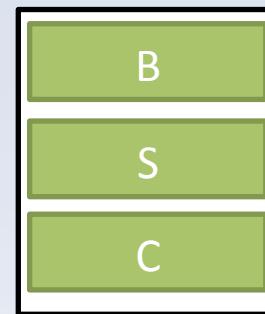
block 3



block 4

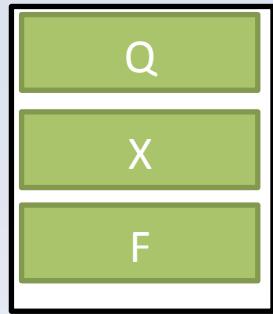


block 5

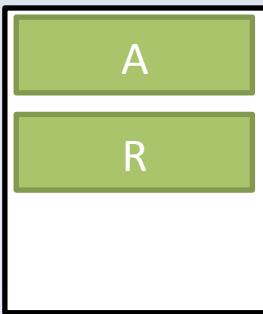


Phase 2

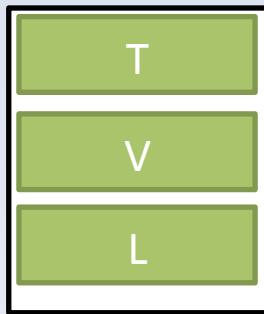
block 0



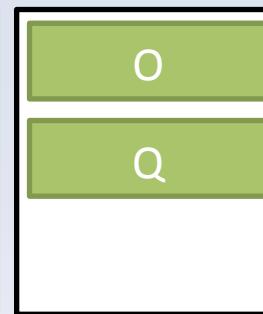
block 1



block 2



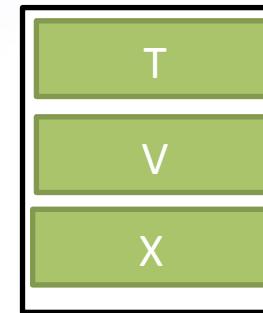
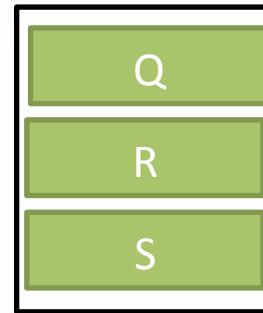
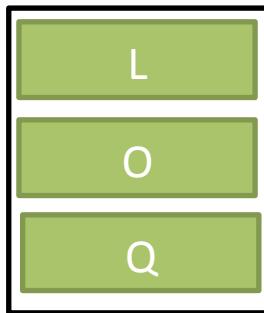
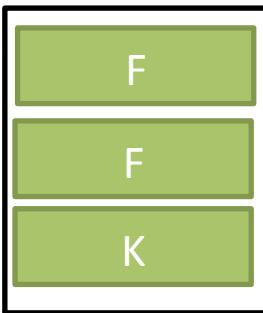
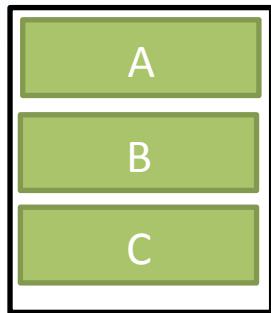
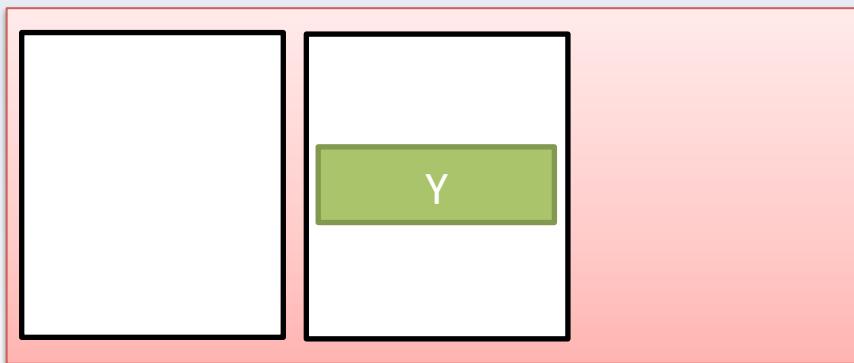
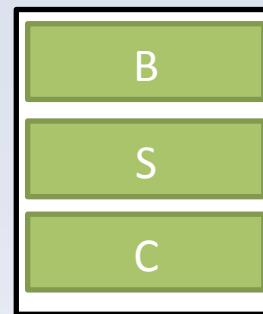
block 3



block 4

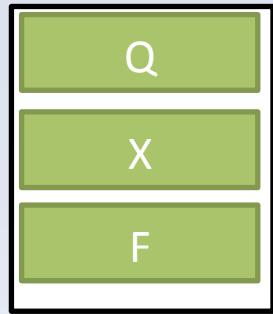


block 5

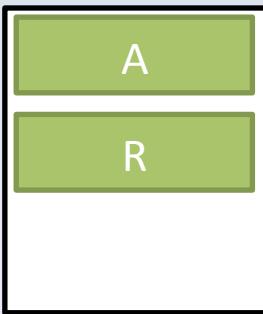


Phase 2

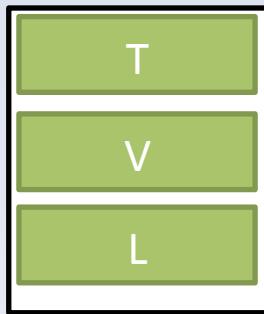
block 0



block 1



block 2



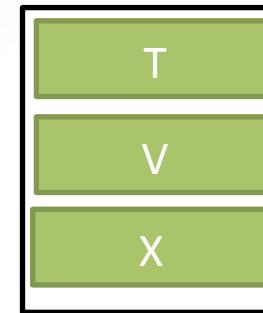
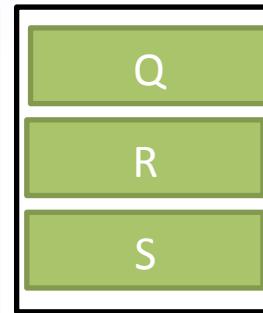
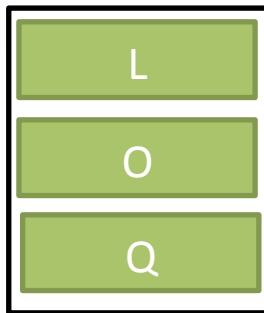
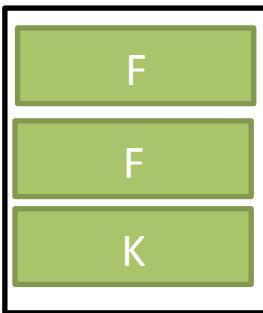
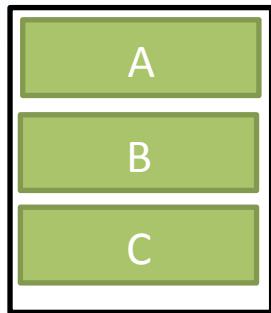
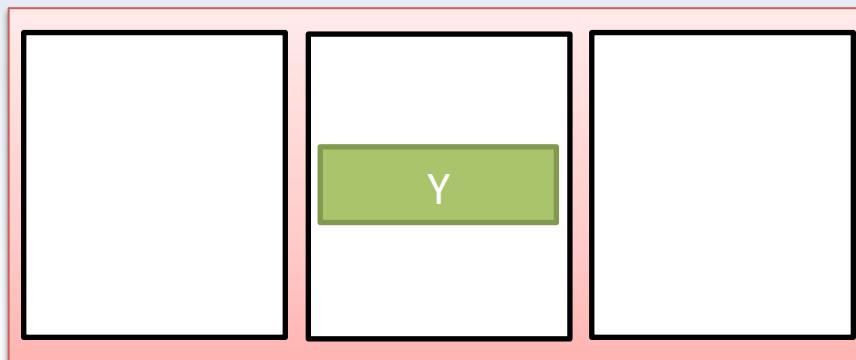
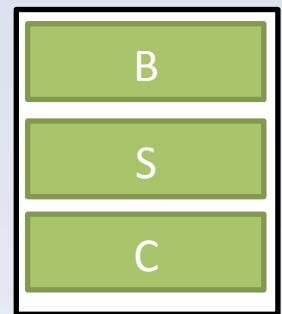
block 3



block 4

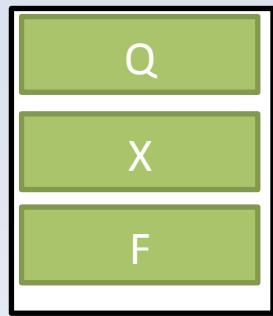


block 5

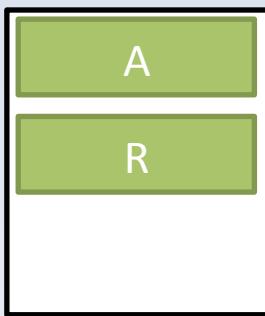


Phase 2

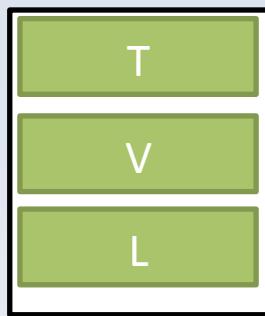
block 0



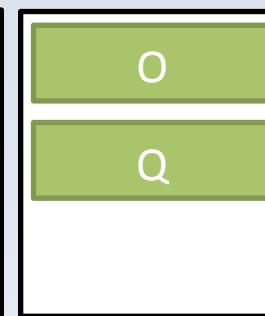
block 1



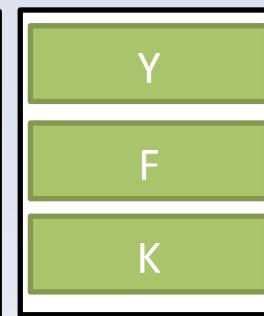
block 2



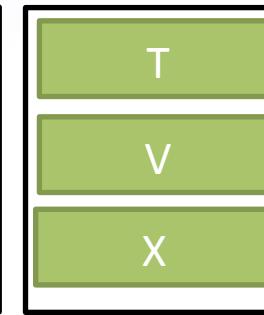
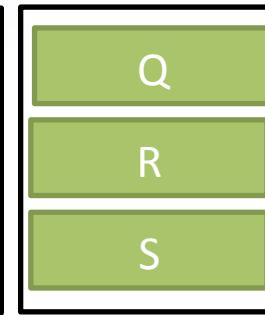
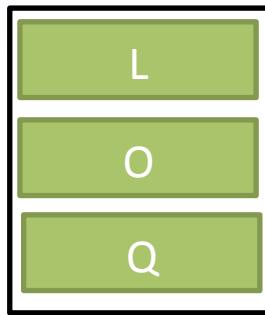
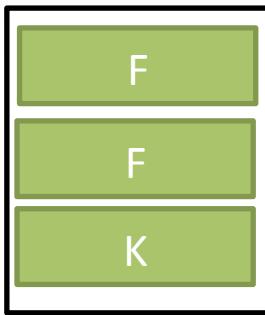
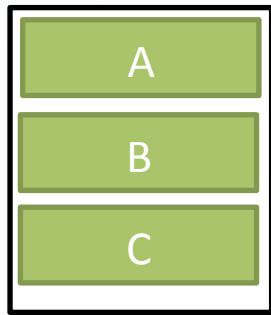
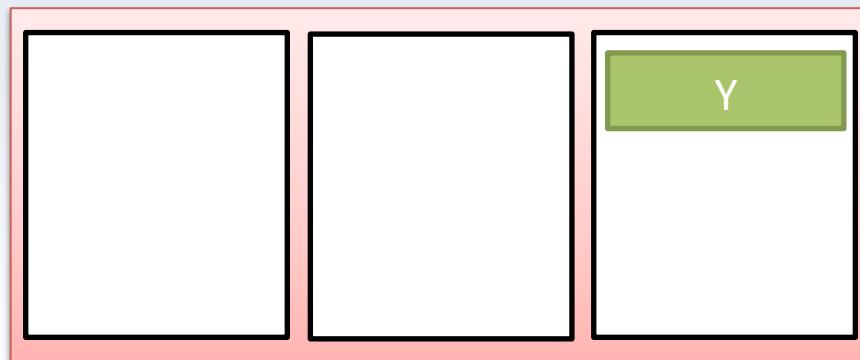
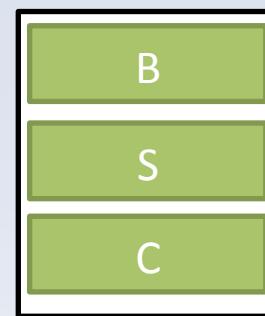
block 3



block 4

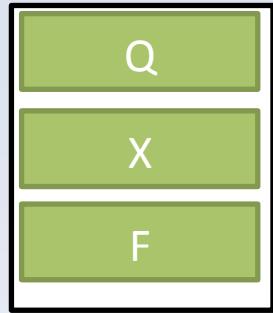


block 5

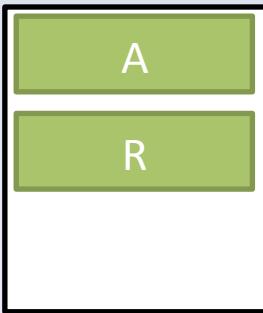


Phase 2

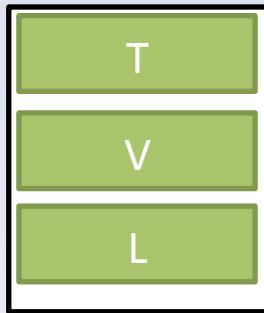
block 0



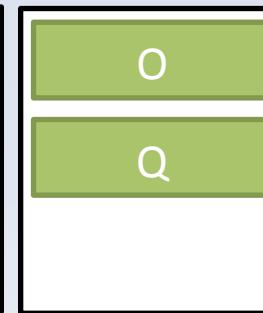
block 1



block 2



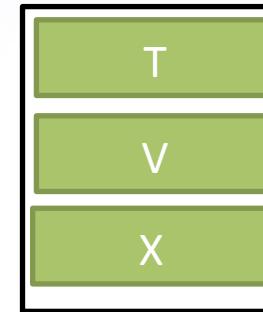
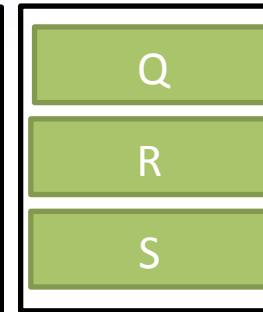
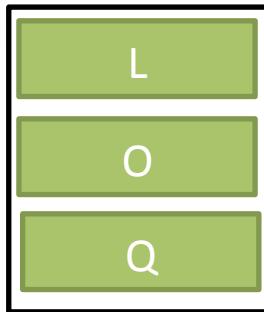
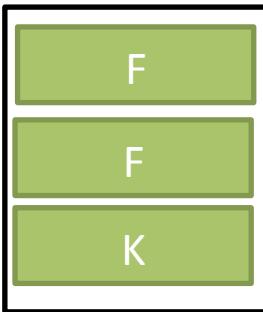
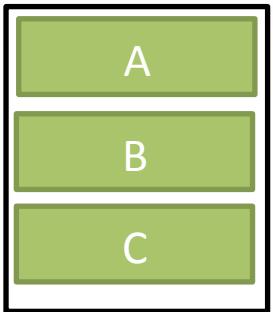
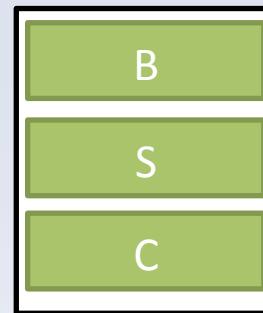
block 3



block 4



block 5



Analysis

- Number of sublists is B/M
- To fit one block for each sublist, $B/M \leq M-1$
- Rewritten, $B \leq M(M-1)$
- Requires that $\sqrt{B(R)} \leq M$



Analysis

- Read B(R) blocks in Phase 1
- Write B(R) blocks in Phase 1
- Read B(R) blocks in Phase 2
- Write B(R) blocks in Phase 2



Analysis

- Read B(R) blocks in Phase 1
- Write B(R) blocks in Phase 1
- Read B(R) blocks in Phase 2
- ~~Write B(R) blocks in Phase 2~~ Output



Analysis

- Read $B(R)$ blocks in Phase 1
- Write $B(R)$ blocks in Phase 1
- Read $B(R)$ blocks in Phase 2
- ~~Write $B(R)$ blocks in Phase 2~~ Output
- Cost is therefore $3B(R)$



Sorting-based δ

- We can implement δ operation with our TPMMS as follows:
 1. Create all of the sorted sublists as usual
 2. In phase 2, don't output duplicate entries (keep track of last record output and skip identical records)



Sorting-based γ

- We can implement γ operation with our TPMMS as follows:
 1. Create all of the sorted sublists as usual
 2. In phase 2, don't output duplicate entries. Instead, compute aggregation functions on duplicate entries



Sorting-based \cup

- We can implement \cup operation with our TPMMS as follows:
 1. Create all of the sorted sublists for both R and S
 2. In phase 2, bring sorted lists for BOTH relations into buffers. Don't output duplicate entries.
- Essentially, treat R and S as one relation



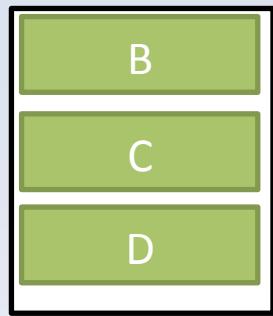
Sorting-based \cap

- We can implement \cap operation with our TPMMS as follows:
 1. Create all of the sorted sublists for both R and S
 2. In phase 2, bring sorted lists for both relations into buffers, but keep them separate
 3. Only output when both relations contain the same value

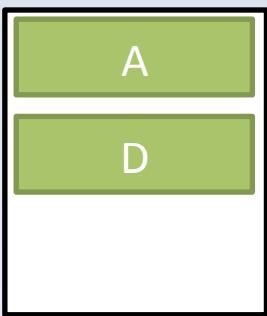


Example $R \cap S$

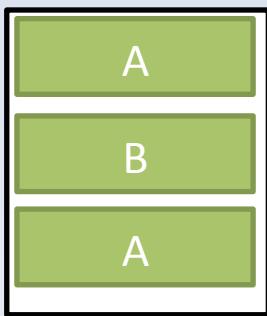
block 0



block 1



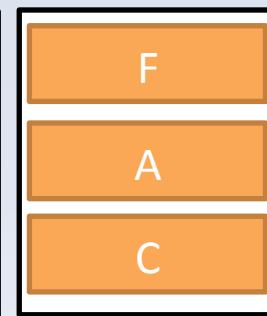
block 2



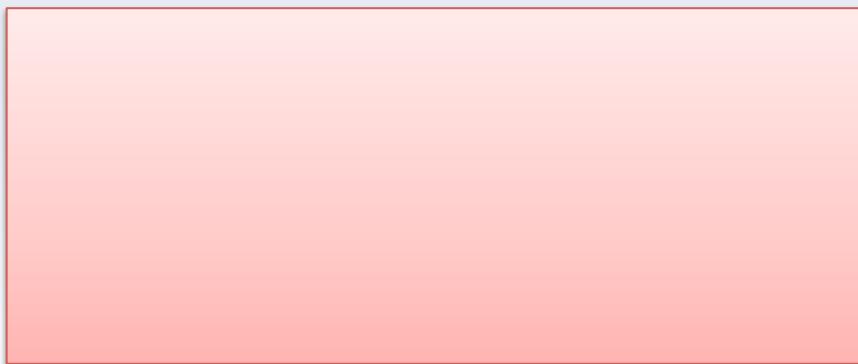
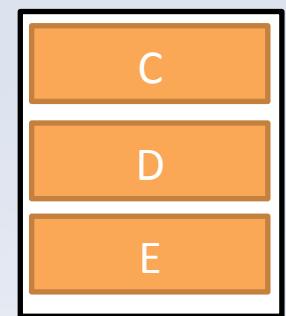
block 0



block 1

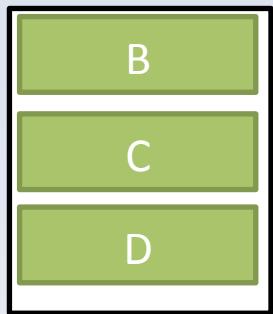


block 2

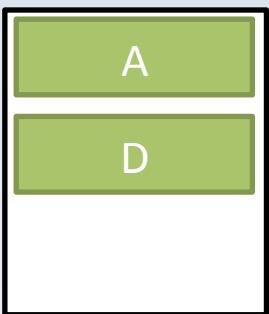


Example $R \cap S$

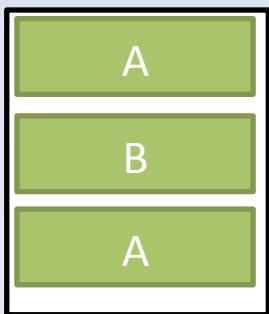
block 0



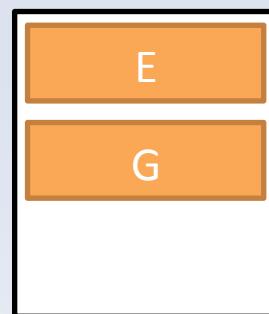
block 1



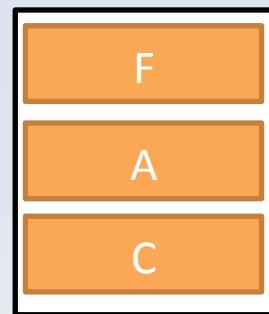
block 2



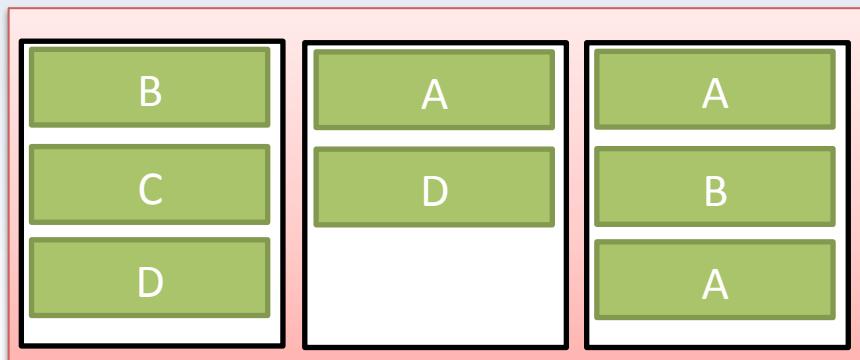
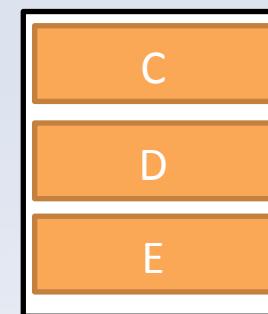
block 0



block 1

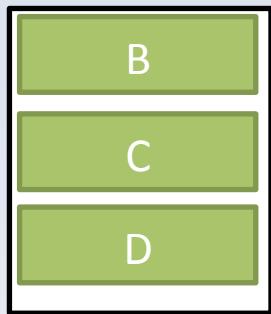


block 2

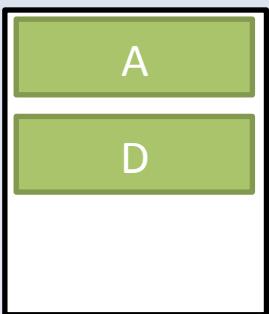


Example $R \cap S$

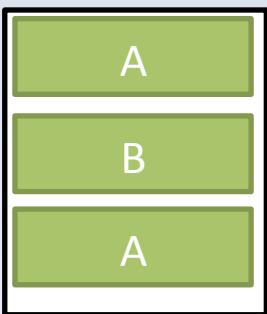
block 0



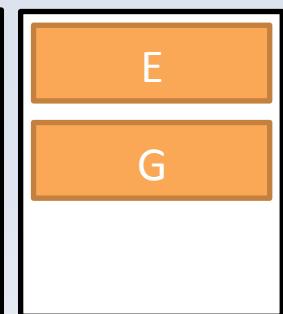
block 1



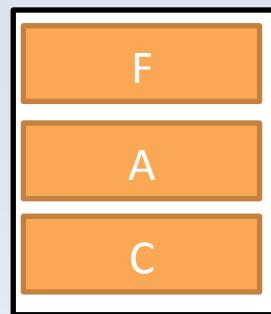
block 2



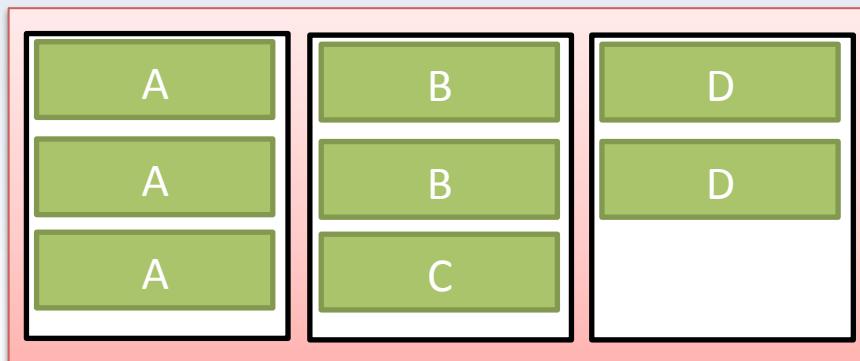
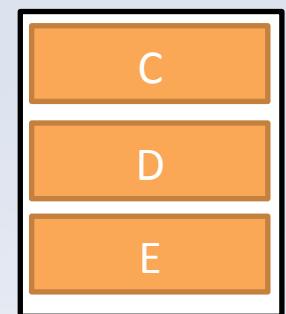
block 0



block 1

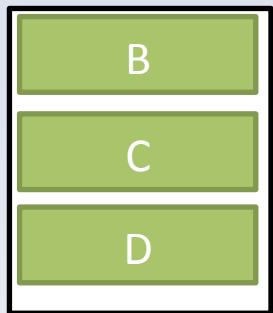


block 2

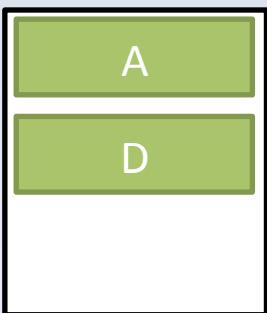


Example $R \cap S$

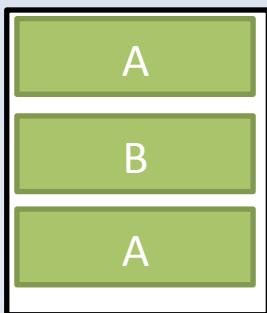
block 0



block 1



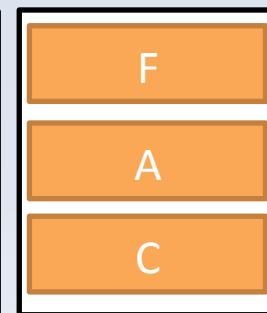
block 2



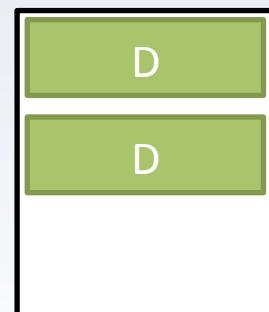
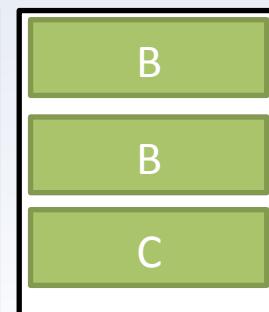
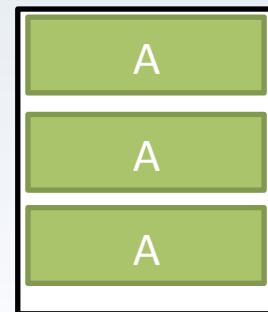
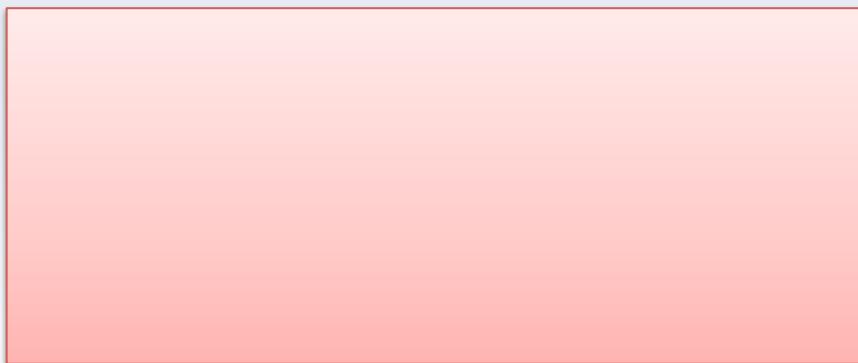
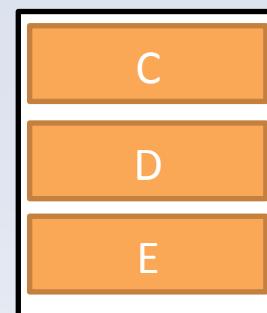
block 0



block 1

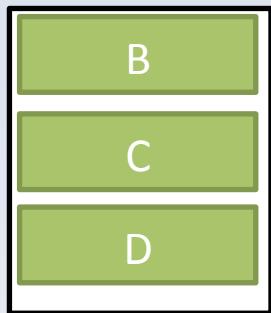


block 2

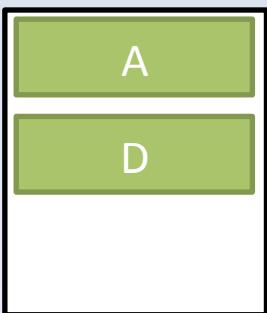


Example $R \cap S$

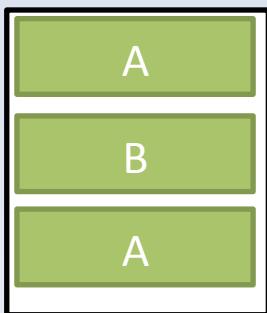
block 0



block 1



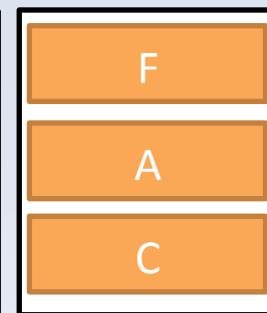
block 2



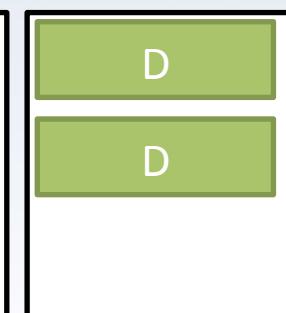
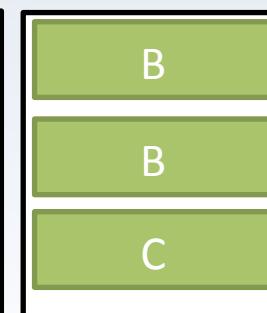
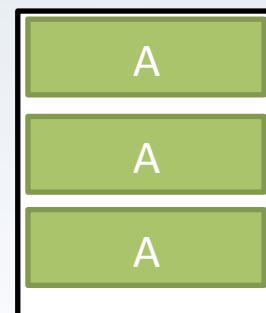
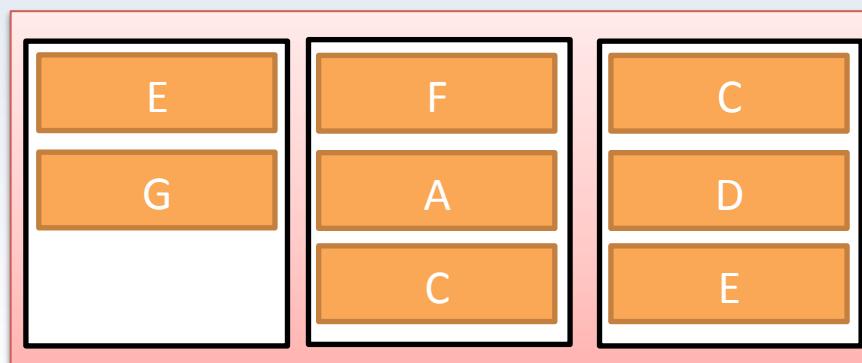
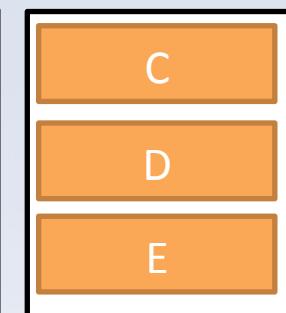
block 0



block 1

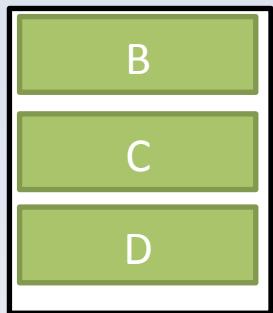


block 2

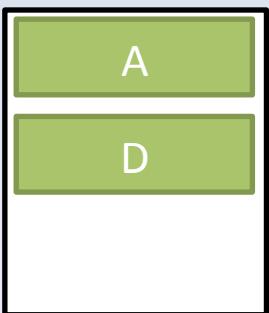


Example $R \cap S$

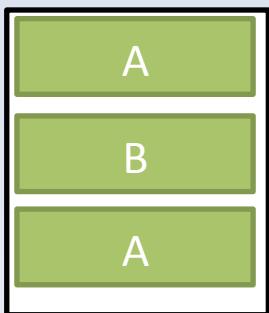
block 0



block 1



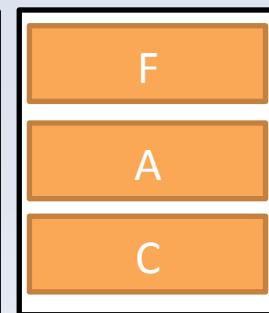
block 2



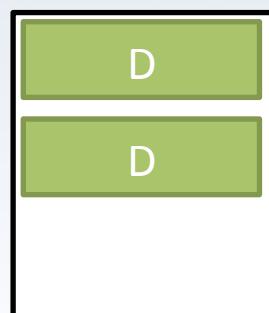
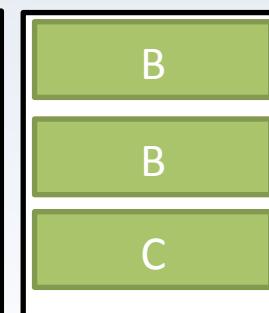
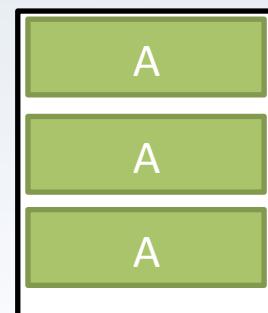
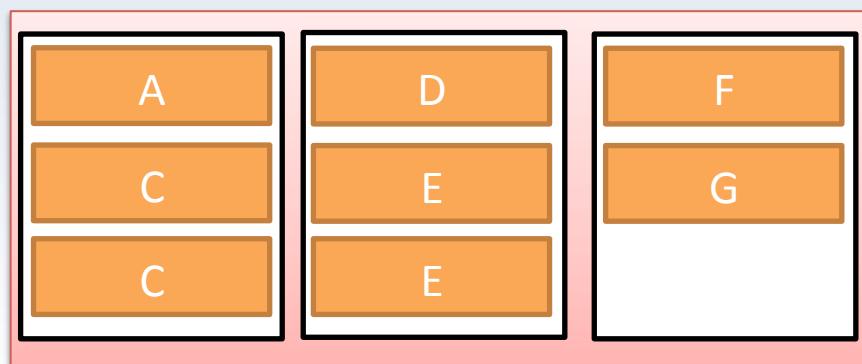
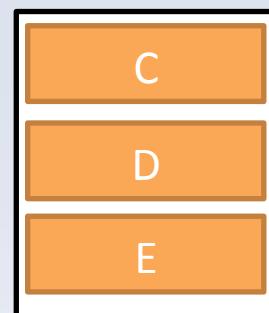
block 0



block 1

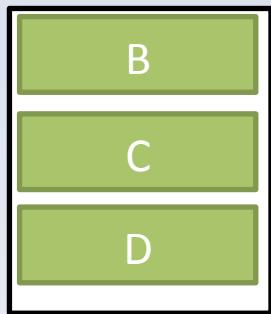


block 2

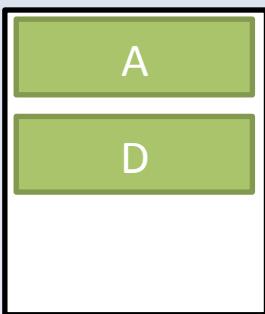


Example $R \cap S$

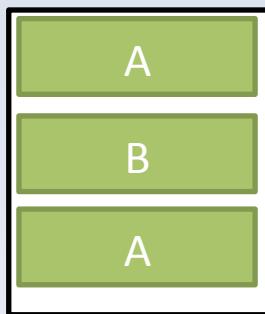
block 0



block 1



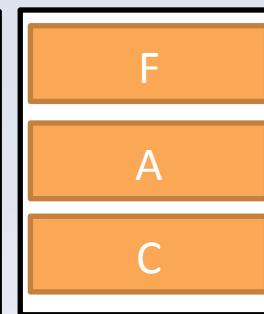
block 2



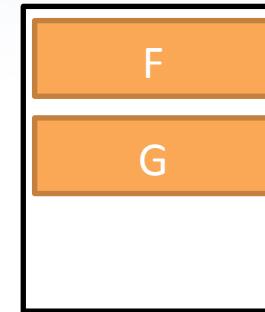
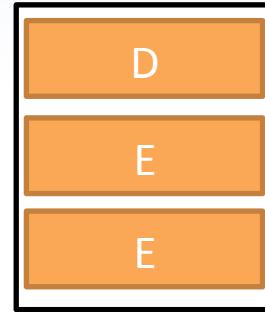
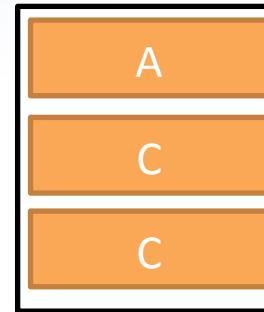
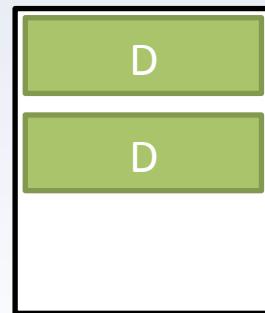
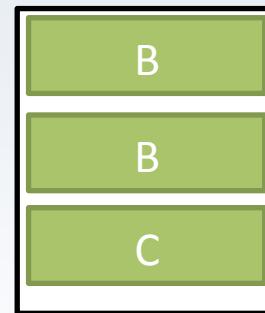
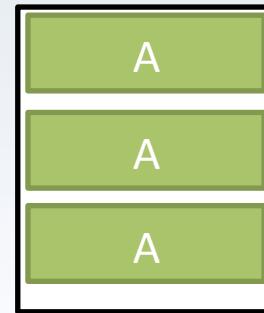
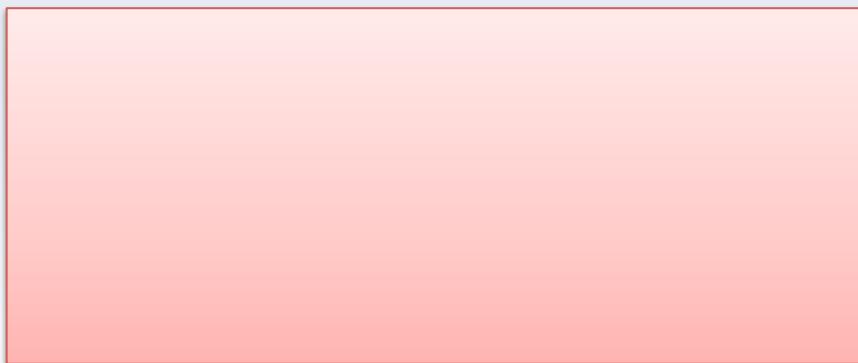
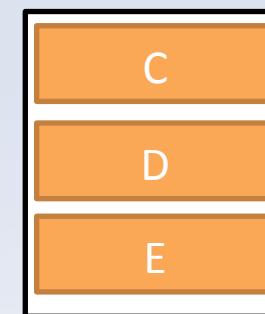
block 0



block 1

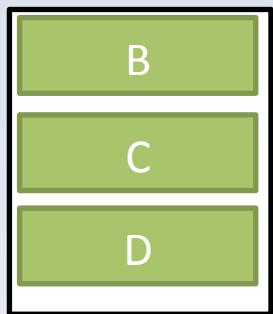


block 2

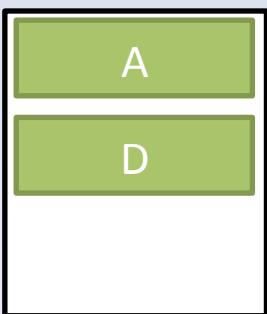


Example $R \cap S$

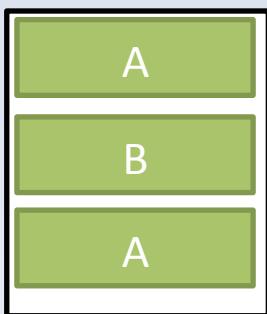
block 0



block 1



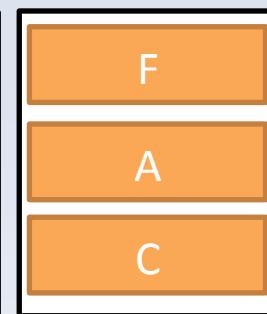
block 2



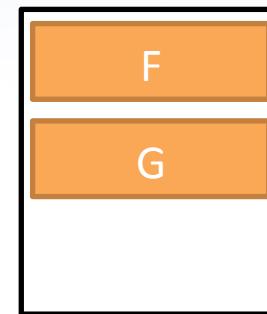
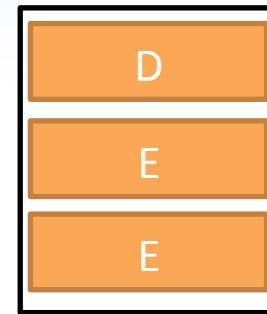
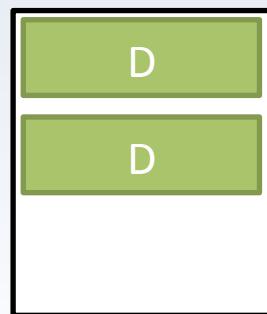
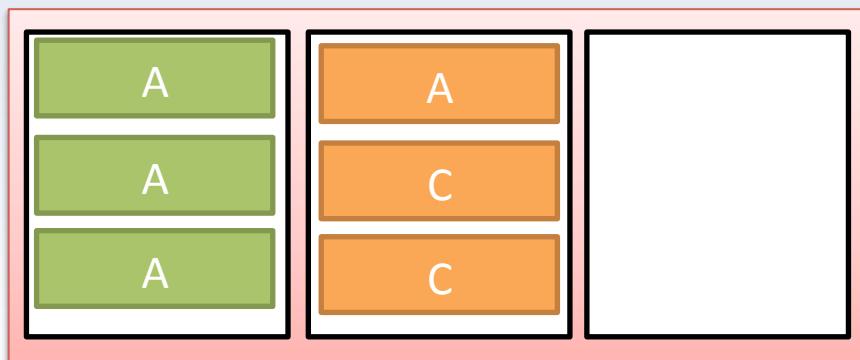
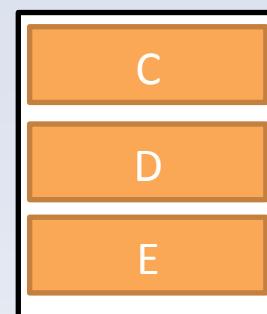
block 0



block 1

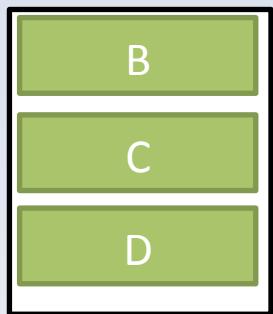


block 2

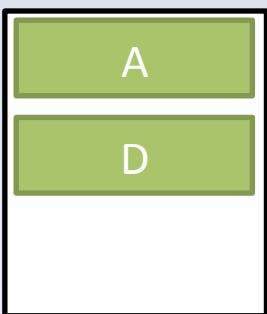


Example $R \cap S$

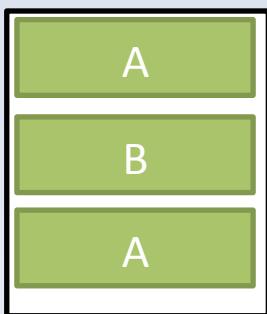
block 0



block 1



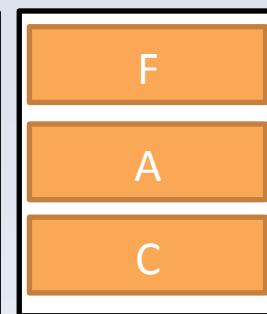
block 2



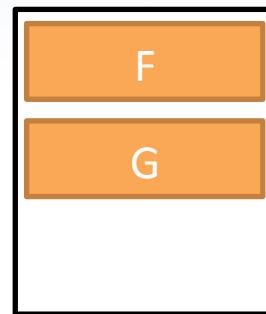
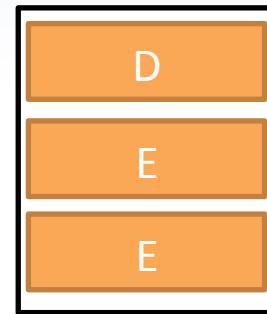
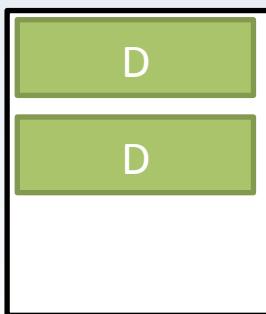
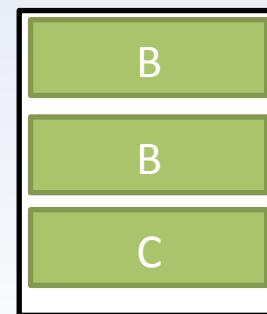
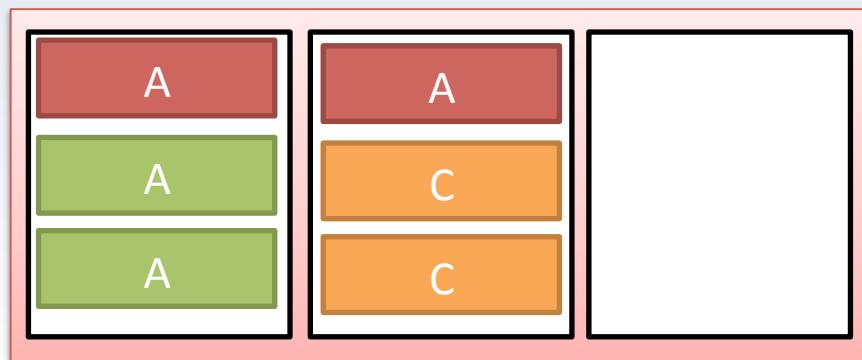
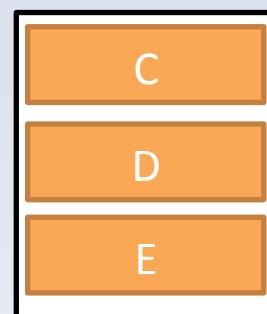
block 0



block 1

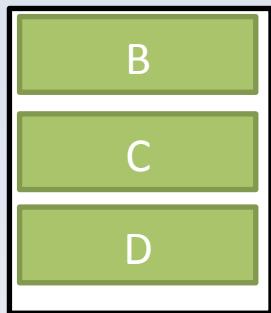


block 2

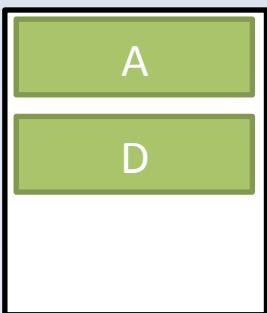


Example $R \cap S$

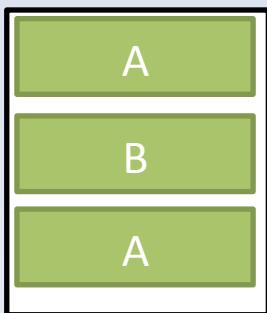
block 0



block 1



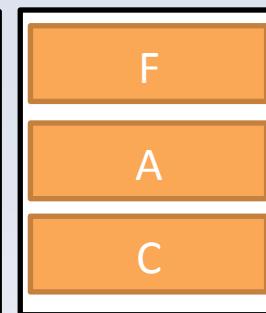
block 2



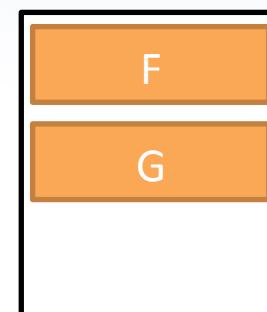
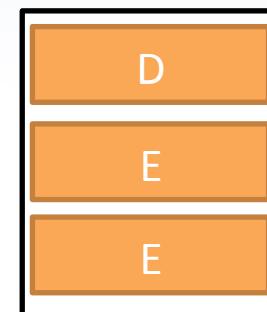
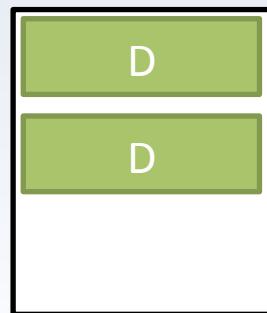
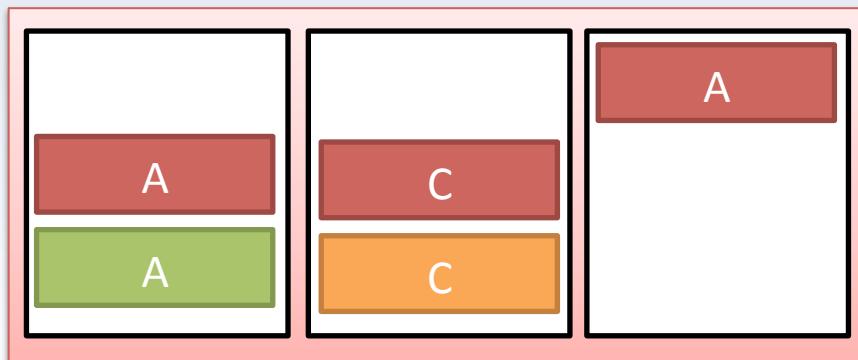
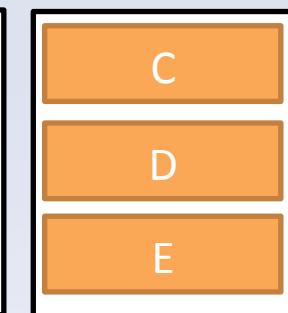
block 0



block 1

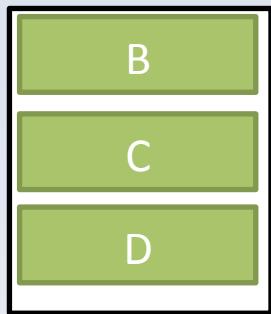


block 2

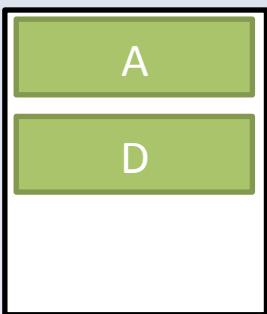


Example $R \cap S$

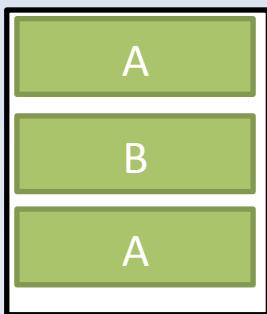
block 0



block 1



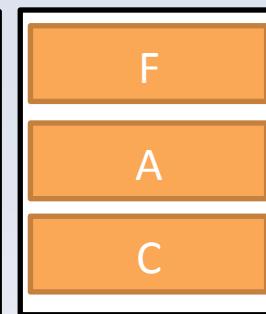
block 2



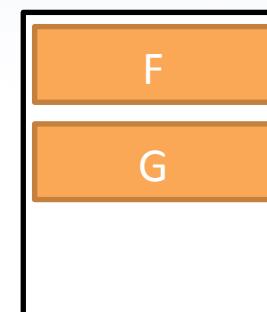
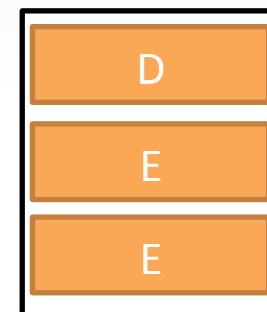
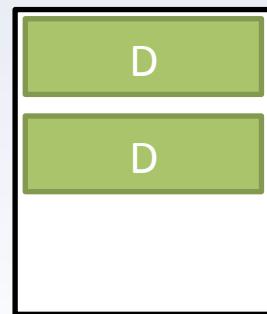
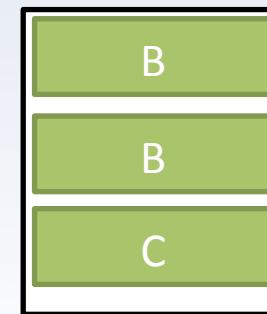
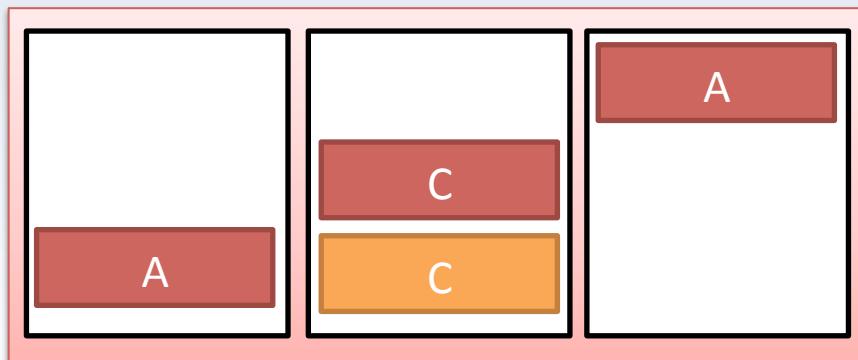
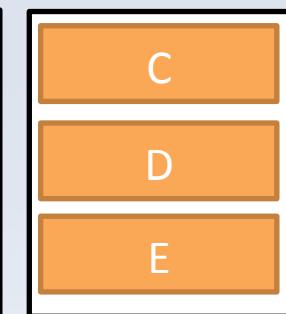
block 0



block 1

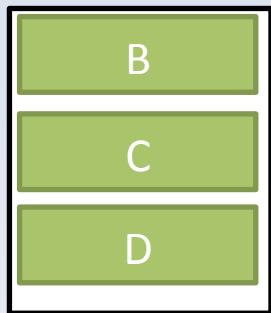


block 2

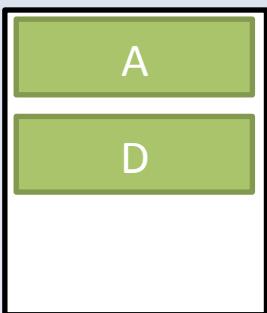


Example $R \cap S$

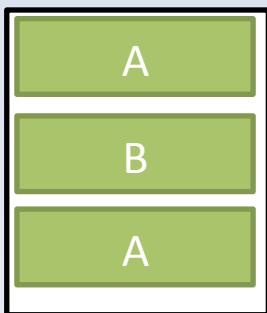
block 0



block 1



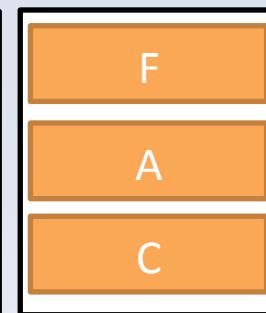
block 2



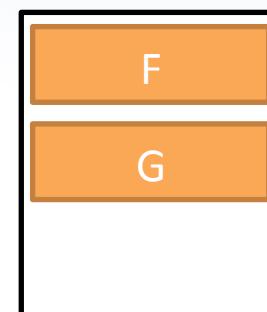
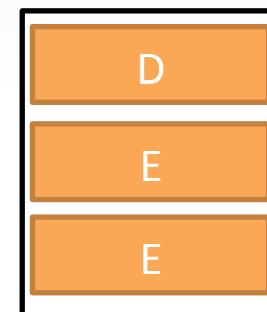
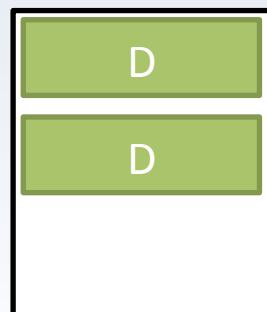
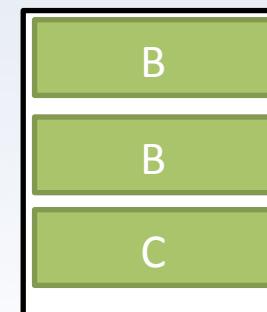
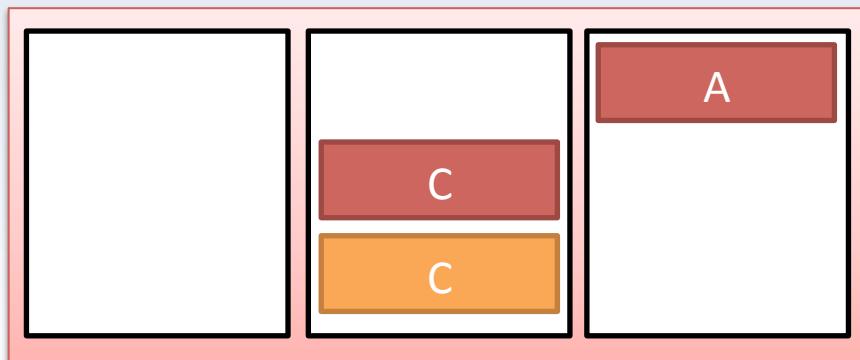
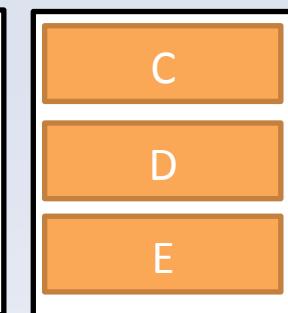
block 0



block 1

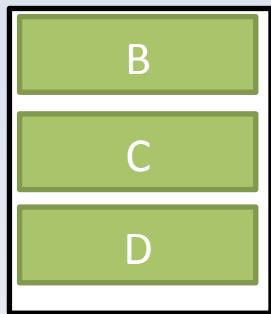


block 2

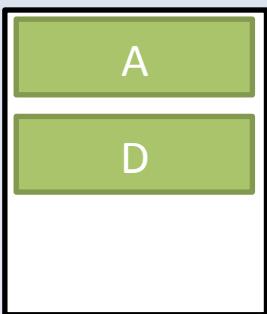


Example $R \cap S$

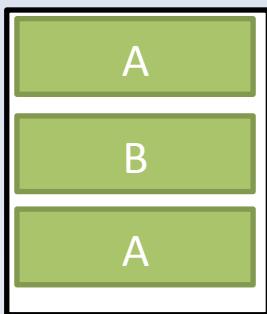
block 0



block 1



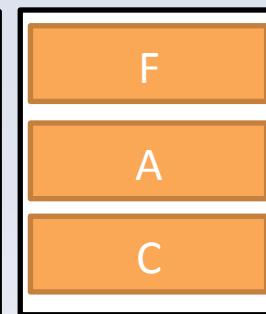
block 2



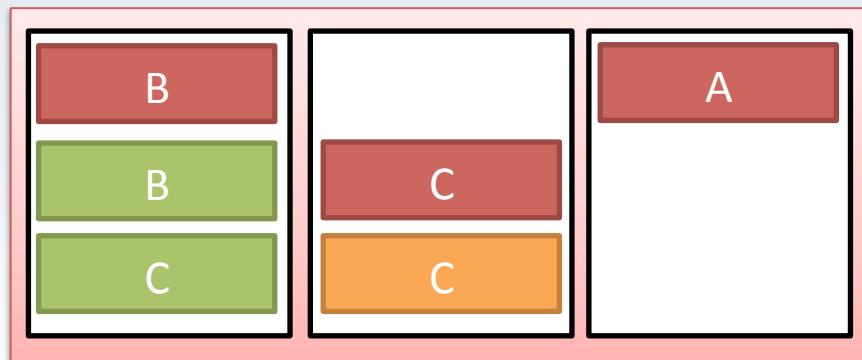
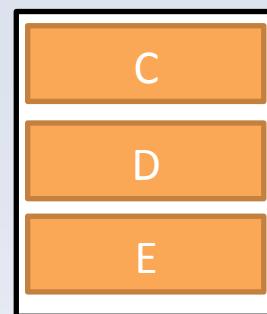
block 0



block 1

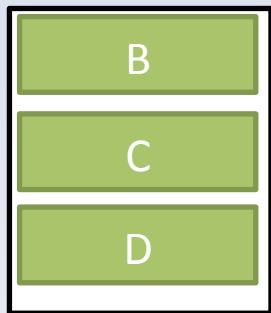


block 2

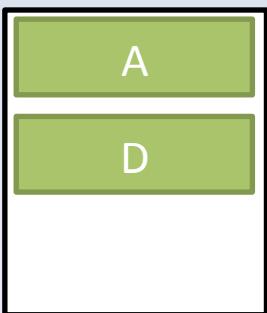


Example $R \cap S$

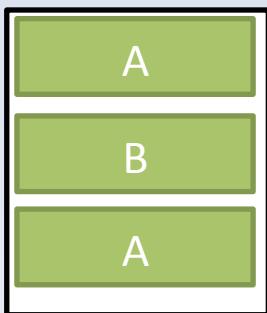
block 0



block 1



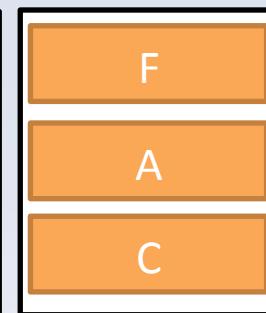
block 2



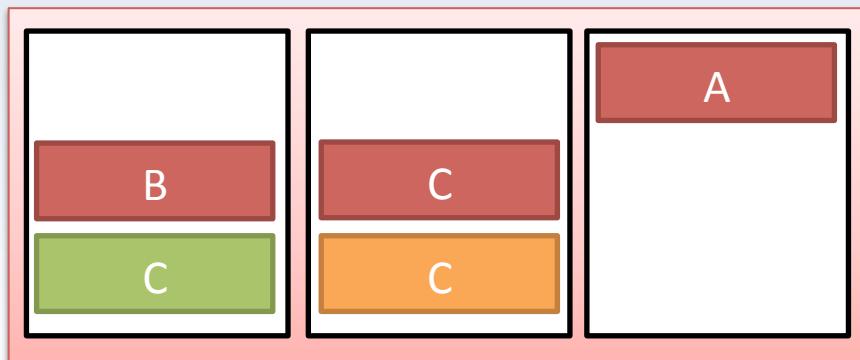
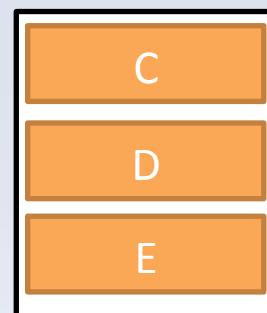
block 0



block 1

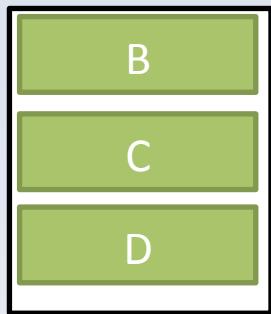


block 2

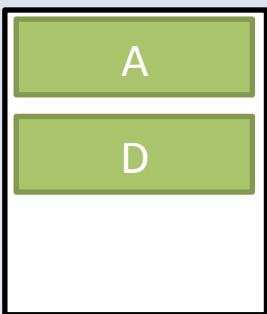


Example $R \cap S$

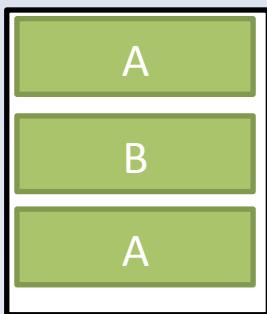
block 0



block 1



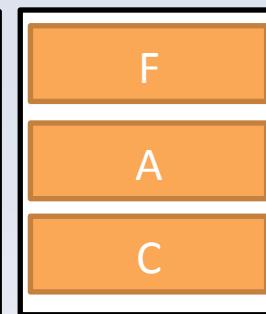
block 2



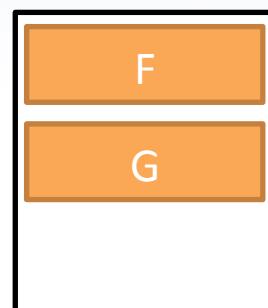
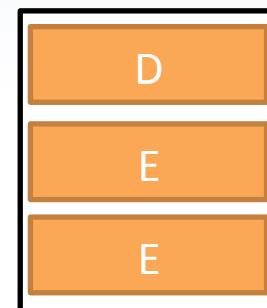
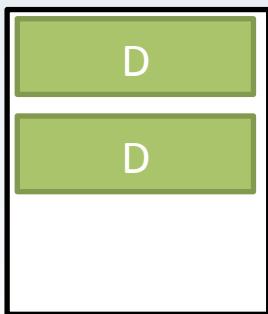
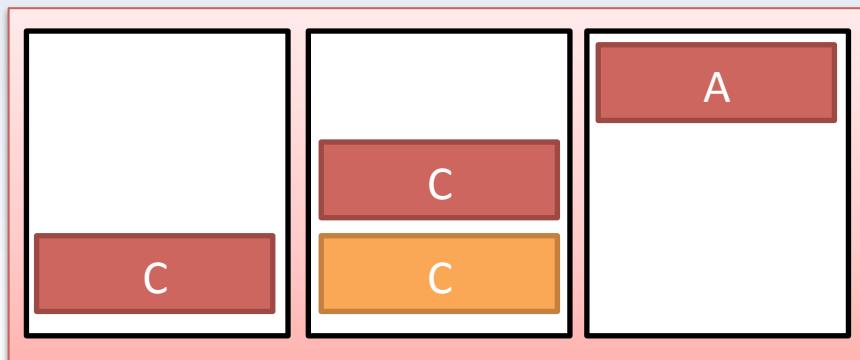
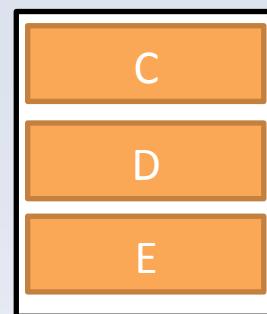
block 0



block 1

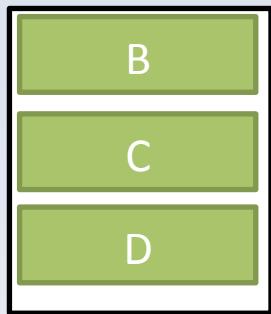


block 2

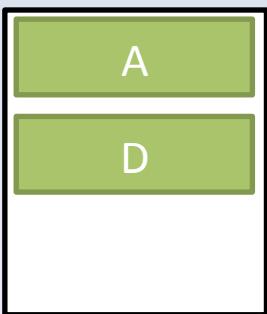


Example $R \cap S$

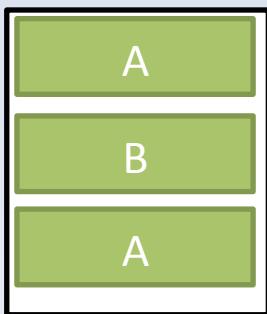
block 0



block 1



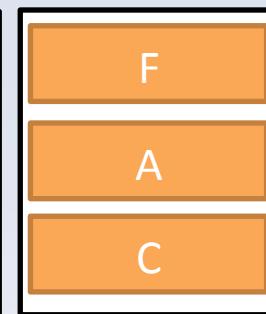
block 2



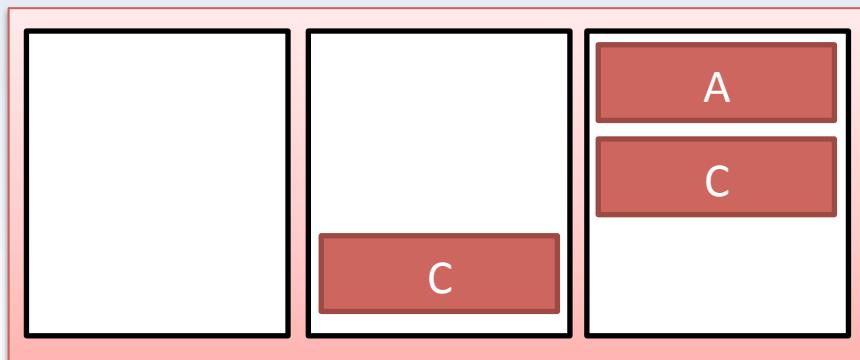
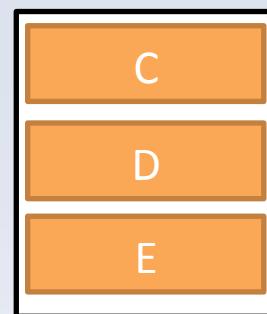
block 0



block 1

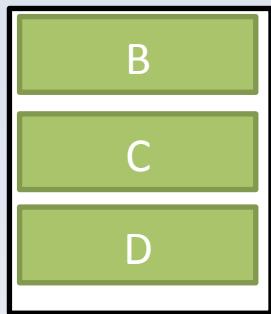


block 2

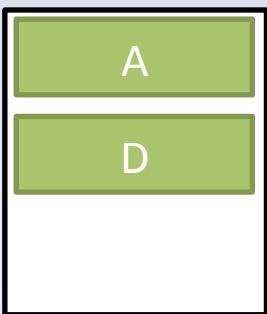


Example $R \cap S$

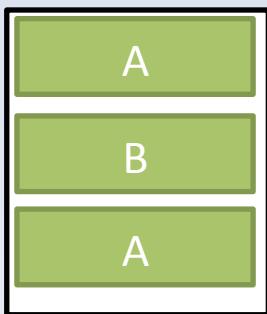
block 0



block 1



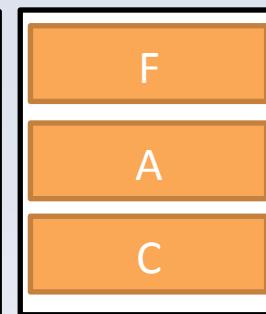
block 2



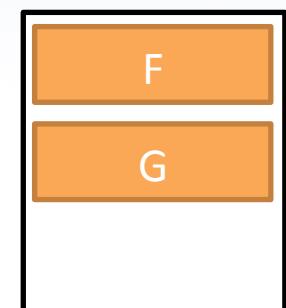
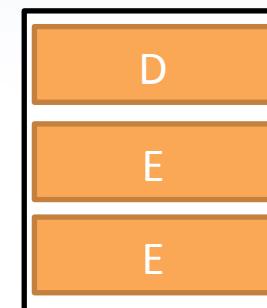
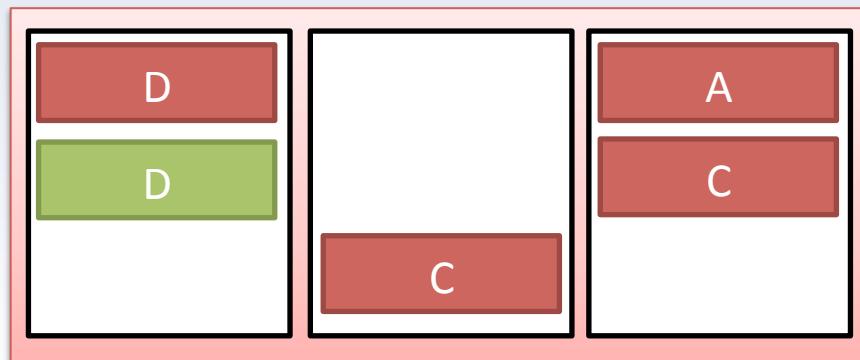
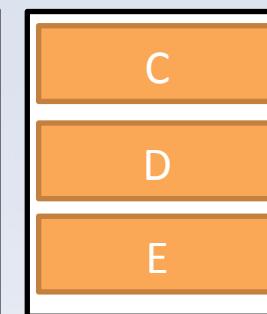
block 0



block 1

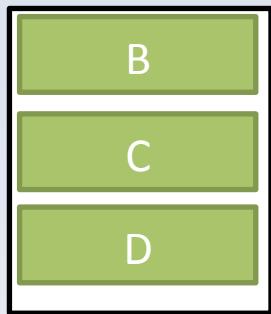


block 2

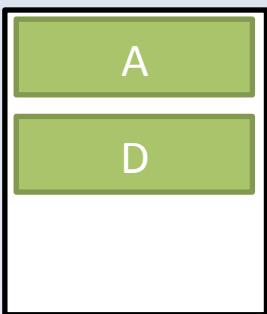


Example $R \cap S$

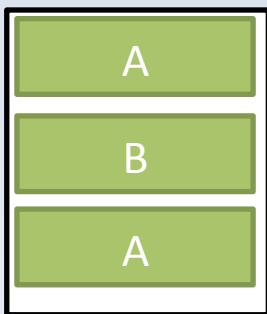
block 0



block 1



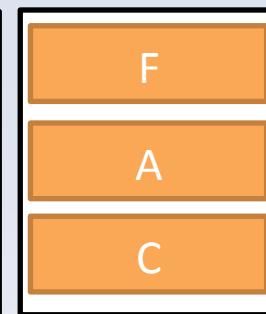
block 2



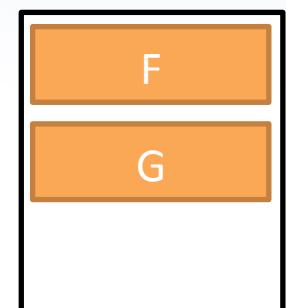
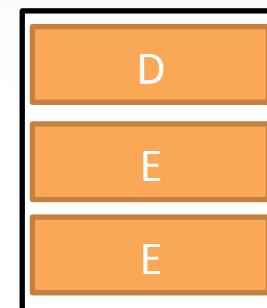
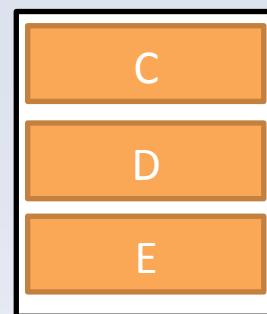
block 0



block 1

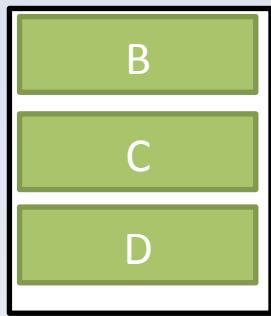


block 2

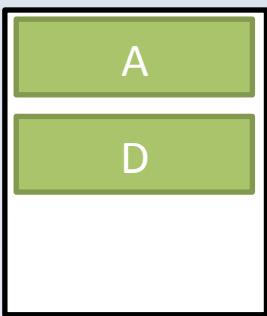


Example $R \cap S$

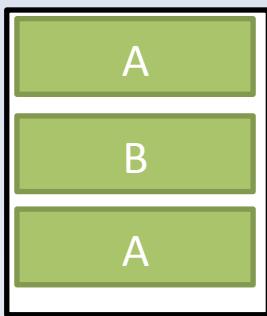
block 0



block 1



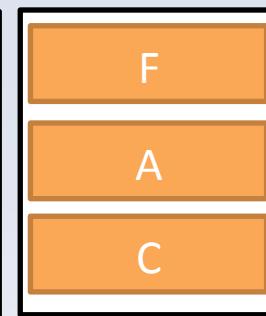
block 2



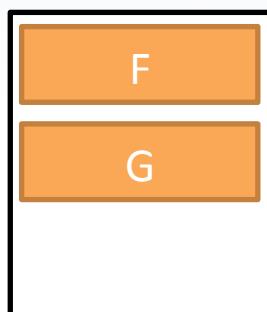
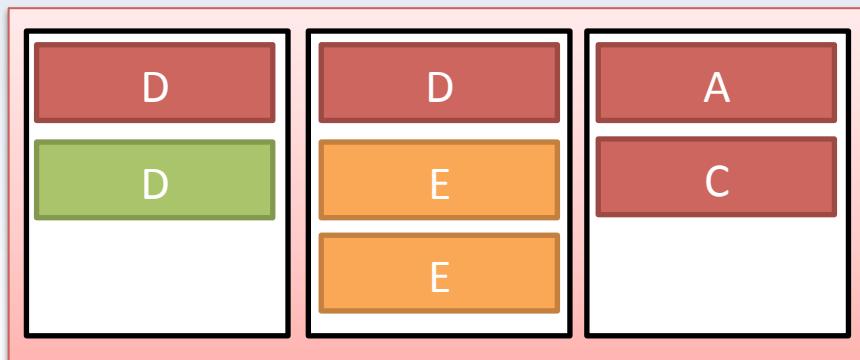
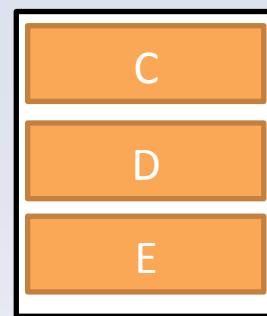
block 0



block 1

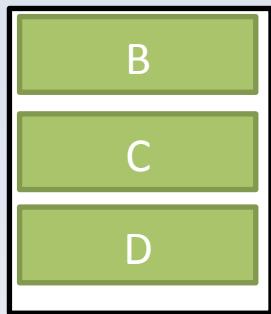


block 2

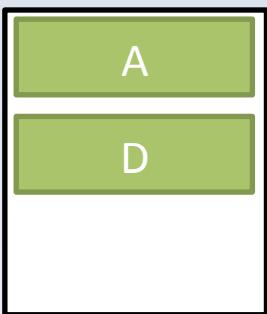


Example $R \cap S$

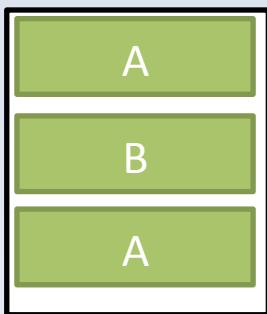
block 0



block 1



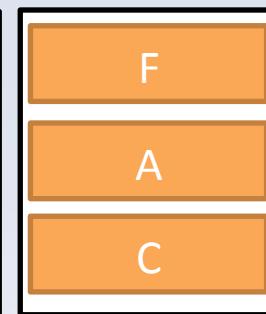
block 2



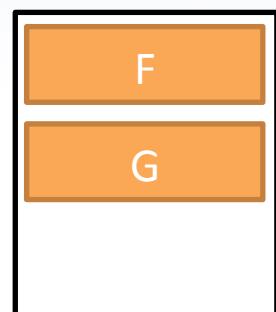
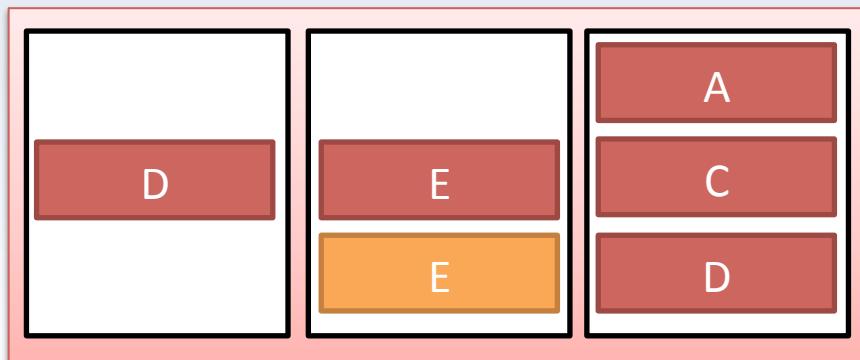
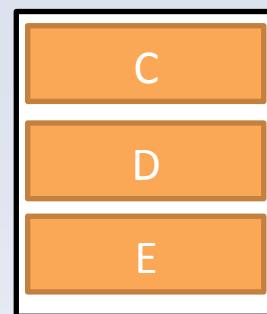
block 0



block 1

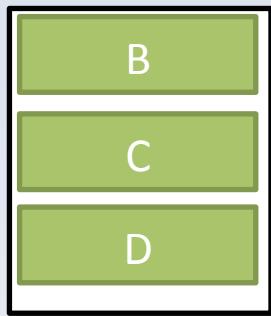


block 2

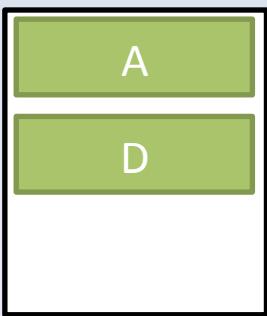


Example $R \cap S$

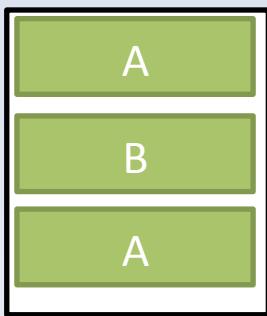
block 0



block 1



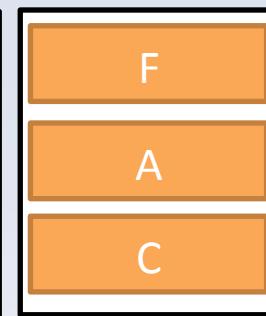
block 2



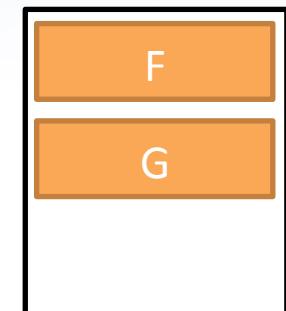
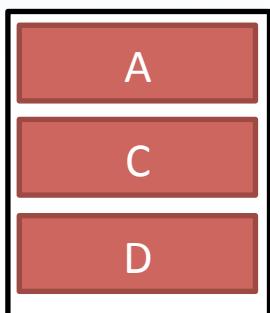
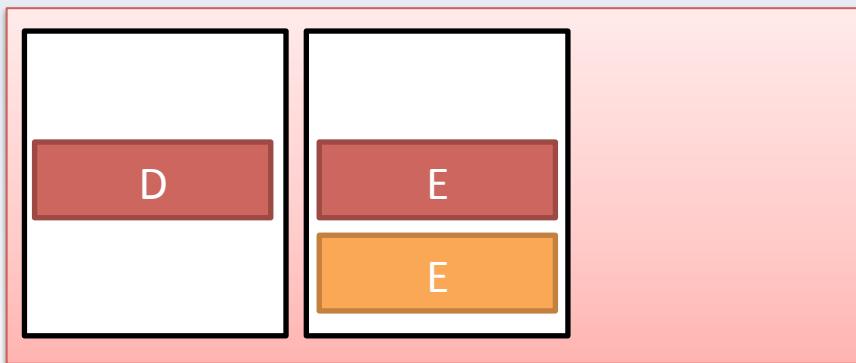
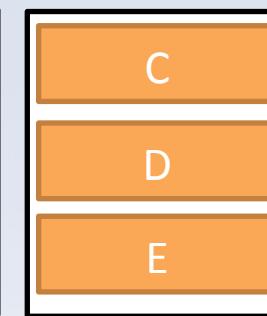
block 0



block 1

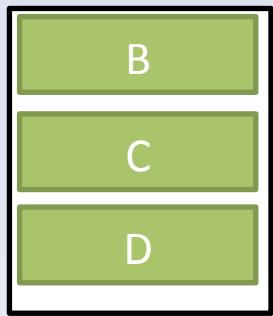


block 2

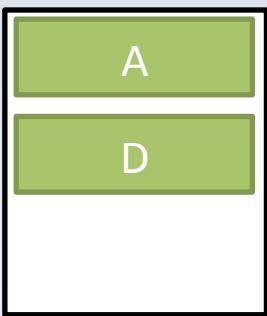


Example $R \cap S$

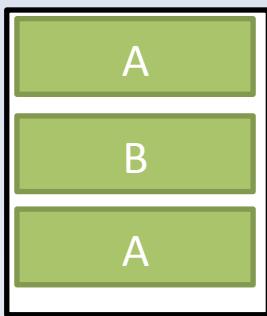
block 0



block 1



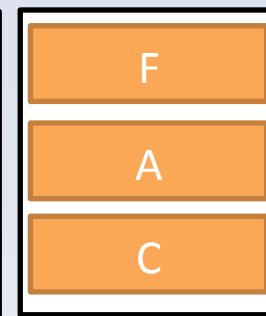
block 2



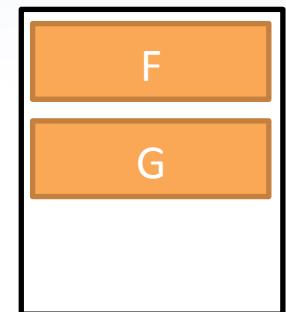
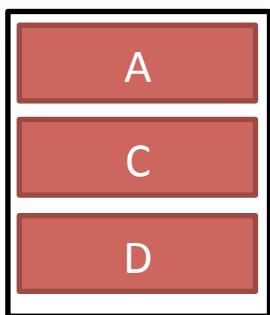
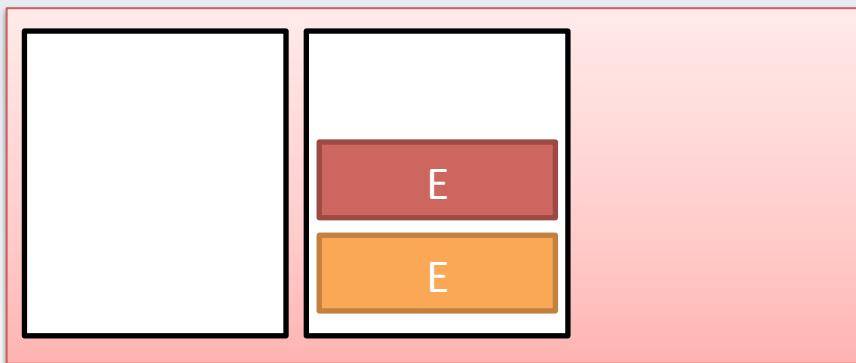
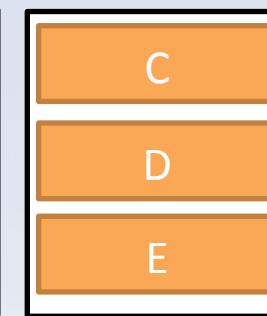
block 0



block 1

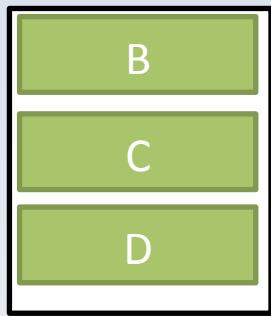


block 2

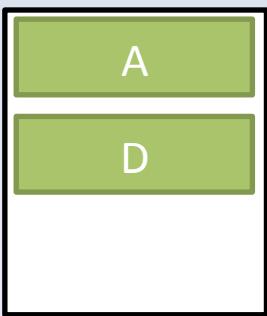


Example $R \cap S$

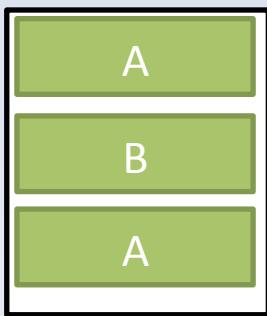
block 0



block 1



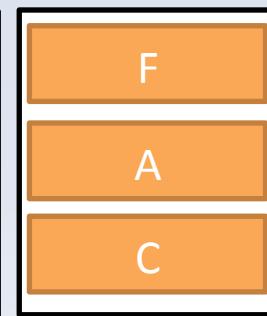
block 2



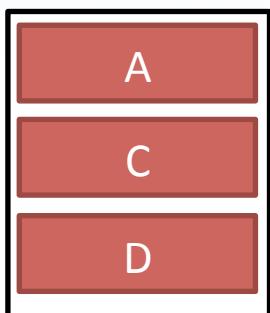
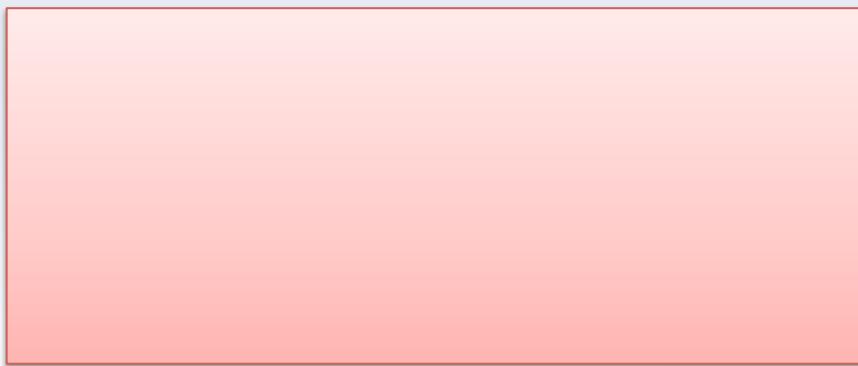
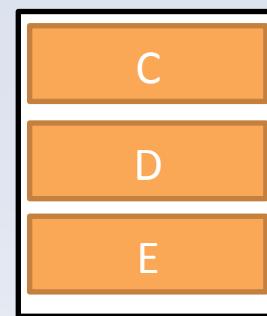
block 0



block 1



block 2



Sorting-based Set Difference

- We can implement $R - S$ similar to the way we implemented $R \cap S$
- Instead of only outputting when value appears in both relations, only output when it appears in R and not S



Sorting-Based Summary

Operator	M required	I/O Cost
δ, γ	\sqrt{B}	$3B$
$\cup, \cap, -$	$\sqrt{B(R) + B(S)}$	$3(B(R) + B(S))$



Next time...

- Use TPMMS to implement 2 pass algorithms for \bowtie operation
- Use *hashing* to implement 2 pass algorithms for operators
- Take advantage of indexes to implement algorithms for operators