

UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Query Optimization



illinois.edu

Announcements

- HW3 is due today
- MP3 is coming
- Piazza poll results:
 - data mining
 - distributed databases



Review

- What steps are involved in query processing?
- What is a parse tree?
- What is the role of the preprocessor?
- What is a logical query plan?
- What is a physical query plan?



Review

- Why do we care about rewriting our query plans?
- What are some heuristics for optimizing a query plan?
- Why do we care about estimating cost?



Cost Estimation

- We start with metadata (data dictionary)
 - $B(R)$ number of blocks holding R
 - $T(R)$ number of tuples in R
 - $V(R,a)$ number of distinct values for attribute a in relation R
- Computed periodically offline
- Possible to compute better statistics
 - e.g. histograms on page 805



Projection Estimation

- Can be computed exactly
 - $B(S) = (T(R) * \text{size}(\text{newtuple})) / \text{blocksize}$
 - $T(S) = T(R)$
 - $V(S, a) = V(R, a)$ (if a in S)



Selection Estimation

- For $S = \sigma_{a=c}(R)$
 - $T(S) \approx T(R)/V(R,A)$
- For $S = \sigma_{a < c}(R)$
 - $T(S) \approx T(R)/3$
- Treat AND as multiple nested selections
- Treat OR as independent (more on this later)
 - $T(S) \approx n(1 - (1 - m_1/n)(1 - m_2/n))$



Example

- Given $R(a,b,c)$ with:
 - $T(R)=10,000$
 - $V(R,a)=50, V(R,b)=20, V(R,c)=100$
- Let $S=\sigma_{a=5 \text{ AND } b<100} (R)$
 - $T(S)\approx(10,000/50)/3=67$



Selection with OR

- $S = \sigma_{c_1 \text{ OR } c_2}(R)$
 - Assume m_1 tuples satisfy c_1 , m_2 satisfy c_2
 - Assume conditions are independent
 - $(1 - m_1/n)$ fraction of tuples that don't satisfy c_1
 - $(1 - m_2/n)$ fraction of tuples that don't satisfy c_2
 - $(1 - m_1/n)(1 - m_2/n)$ fraction that satisfy neither
 - $1 - (1 - m_1/n)(1 - m_2/n)$ fraction that satisfy either
 - $T(S) \approx n(1 - (1 - m_1/n)(1 - m_2/n))$



Example

- Given $R(a,b,c)$ with:
 - $T(R)=10,000$
 - $V(R,a)=50, V(R,b)=20, V(R,c)=100$
- Let $S=\sigma_{a=5 \text{ OR } b<100} (R)$
 - Tuples satisfying $a=5 \approx 10k/50=200=m_1$
 - Tuples satisfying $b<100 \approx 3333=m_2$
 - $T(S) \approx (1-(1-200/10k)(1-3333/10k))=3466$



Join Estimation

- Arguably the most important cost
- We need to predict how many tuples will relate between R and S
 - Could be anywhere between 0 and $T(R)T(S)$
- We'll first work with natural join
 - $U = R(X, Y) \bowtie S(Y, Z)$
- Can easily generalize to other joins



Join Estimation

- Simplifying assumptions:
 1. If join attribute has more values in R than in S, **all** values in S occur in R
 - $V(R, Y) \geq V(S, Y)$
 - in other words, assume every tuple in S joins
 2. Non-join attributes do not lose values
 - $V(R \bowtie S, A) \approx V(R, A)$ for A not in Y
 - $V(R \bowtie S, B) \approx V(S, B)$ for B not in Y



Join Estimation

- Probability a pair of tuples join:
 $1/\max(V(R,Y),V(S,Y))$
- Total number of pairs of tuples: $T(R)T(S)$
- $T(R \bowtie S) \approx T(R)T(S)/\max(V(R,Y),V(S,Y))$



Example

- $R(a,b,c) \bowtie S(b,d,e)$
 - $T(R)=1000$, $V(R,a)=100$, $V(R,b)=20$, $V(R,c)=200$
 - $T(S)=2000$, $V(S,b)=50$, $V(S,d)=100$, $V(S,e)=400$
- $T(R \bowtie S) \approx 1000 * 2000 / \max(20, 50) = 40,000$



Other Operators

- Assume $T(S) < T(R)$
 - $T(R \cup S) \approx T(R) + T(S)/2$
 - $T(R \cap S) \approx T(S)/2$
 - $T(R - S) \approx T(R) - T(S)/2$
 - $T(\delta(R)) \approx \min(T(R)/2, V(R,a)V(R,b), \dots V(R,z))$
 - $T(\gamma(R)) \approx \min(T(R)/2, V(R,a)V(R,b), \dots V(R,z))$



Cost estimation

- Final result of cost estimation:
 - Ignore the input and the output
 - Sum up total number of I/Os
 - Proportional to total size of intermediate relations, so sum this up



Example

Given $R(a,b)$ and $S(b,c)$

$$T(R)=5000, V(R,a)=50, V(R,b)=100$$

$$T(S)=2000, V(S,b)=200, V(S,c)=100$$

$$U=\delta(\sigma_{a=10}(R \bowtie S))$$

$$\left. \begin{aligned} T(R \bowtie S) &= 5000 * 2000 / 200 = 50000 \\ T(\sigma_{a=10}(R \bowtie S)) &= 50000 / 50 = 1000 \end{aligned} \right\} \text{cost}=51,000$$

$$T(U)=\min(500, 20000)=500$$



Example

Given $R(a,b)$ and $S(b,c)$

$$T(R)=5000, V(R,a)=50, V(R,b)=100$$

$$T(S)=2000, V(S,b)=200, V(S,c)=100$$

$$U=\delta(\sigma_{a=10}(R)) \bowtie \delta(S)$$

$$T(\sigma_{a=10}(R))=100, T(\delta(\sigma_{a=10}(R)))=50$$

$$T(\delta(S))=1000$$

$$T(U)=1000*50/200=250$$



Example

Given $R(a,b)$ and $S(b,c)$

$$T(R)=5000, V(R,a)=50, V(R,b)=100$$

$$T(S)=2000, V(S,b)=200, V(S,c)=100$$

$$U = \delta(\sigma_{a=10}(R)) \bowtie \delta(S)$$

$$\left. \begin{array}{l} T(\sigma_{a=10}(R))=100, T(\delta(\sigma_{a=10}(R)))=50 \\ T(\delta(S))=1000 \end{array} \right\} \text{cost}=1,150$$

$$T(U)=1000*50/200=250$$



Query Optimization

- We need three things:
 1. ~~Rewrite rules~~
 2. An optimization algorithm
 3. ~~A cost estimator~~



Naïve Approach

- Try all possible rewrites of our query
- Estimate the cost of each one
- Choose the one with the lowest cost



Naïve Approach

$$\gamma_{a,b,\text{avg}(x)}(\sigma_C(R \bowtie S \bowtie T \bowtie U \bowtie V \bowtie W))$$

- Number of binary trees with n leaves = $(n-1)$ Catalan number = $[2(n-1)]! / [(n!)(n-1)!]$
- Number of arrangements of n leaves = $n!$
- Number of ways to reorder n joins = $[2(n-1)]! / (n-1)!$

Reordering JUST the joins in this problem: 30,240

For 10 joins: 17,643,225,600



Naïve Approach

- For realistic queries, naïve approach is impractical
 - Can't afford to enumerate all possible plans
 - Takes longer than the query itself!



Search Strategies

- Optimization on a large search space
 - Greedy
 - Branch and Bound
 - Hill Climbing
 - Simulated annealing
 - Genetic algorithms
 - Stochastic optimization
 - ... and many, many, many, many more




Join Order

- Choosing join order is very important
 - Why?
- We will focus entirely on this task:
 - given a series of n joins
 - assume we can compute a cost for each join
 - Find the best tree



Left Join Tree



- A tree topology with two advantages:
 1. Only $n!$ different orderings possible
 2. Can be easily pipelined with our algorithms
- On examples/exams, only use left join trees



Dynamic Programming

- Main idea: in large combinatorial problems, many of the partial results are helpful in later steps
- Example:
 $\text{Cost}((R \bowtie S) \bowtie T)$ result is useful for computing
 $\text{Cost}((R \bowtie S) \bowtie T) \bowtie U$
 $\text{Cost}(((R \bowtie S) \bowtie T) \bowtie V)$
 $\text{Cost}(((R \bowtie S) \bowtie T) \bowtie W)$



Dynamic Programming

- For each subset of $\{R_1, R_2, \dots, R_n\}$, compute best plan for that subset
 - Do this in increasing order of cardinality
 - Step 0: $\{R_1\}, \{R_2\}, \dots, \{R_n\}$
 - Step 1: $\{R_1, R_2\}, \{R_1, R_3\}, \dots, \{R_{n-1}, R_n\}$
 - ...
 - Step n: $\{R_1, R_2, \dots, R_n\}$



Dynamic Programming

- For each subset $S \subseteq \{R_1, R_2, \dots, R_n\}$
 - Compute size of $T(S)$
 - Find the best plan for S
 - Compute cost of S



Dynamic Programming

- Step 0: for each R_i
 - $\text{Size}(R_i)=0$
 - $\text{Plan}(R_i)=R_i$
 - $\text{Cost}(R_i)=0$



Dynamic Programming

- Step i : $S \subseteq \{R_1, R_2, \dots, R_n\}$ of cardinality $i+1$
 - For each pair of smaller subsets: S' and S'' such that $S = S' \cup S''$, find cost of best plan for $S' \bowtie S''$
 - $\text{Size}(S) = T(S)$
 - $\text{Plan}(S) = \text{best plan}$
 - $\text{Cost}(S) = \text{cost of best plan}$



Dynamic Programming

- We already have best plan and cost for smaller subsets S' and S''
- $\text{Cost}(S) =$
 $\text{Cost}(S') + \text{Cost}(S'') + \text{Size}(S') + \text{Size}(S'')$
 - $\text{Size}(A) = 0$ if A is a relation (it's input!)



Example

- $\text{Cost}(R \bowtie S) = \text{Cost}(R) + \text{Cost}(S) + 0 + 0 = 0$
- $\text{Cost}((R \bowtie S) \bowtie T) =$
 $\text{Cost}(R \bowtie S) + \text{Cost}(T) + \text{Size}(R \bowtie S) + \text{Size}(T) =$
 $\text{Size}(R \bowtie S)$



Example

- Given $R(a,b), S(b,c), T(c,d), U(d,a)$
 - $T(R)=T(S)=T(T)=T(U)=1k$
 - $V(R,a)=100, V(R,b)=200$
 - $V(S,b)=100, V(S,c)=500$
 - $V(T,c)=20, V(T,d)=500$
 - $V(U,a)=50, V(U,d)=1k$



Example

Subset	Size(S)	Cost	Plan
{R,S}			
{R,T}			
{R,U}			
{S,T}			
{S,U}			
{T,U}			
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



REMEMBER!

- We're ONLY considering left join trees!
 - $((R \bowtie S) \bowtie T) \bowtie U$ is valid
 - $(R \bowtie S) \bowtie (T \bowtie U)$ is ignored!
 - It **is** possible to do dynamic programming that considers all possible trees, though
- $R \bowtie S$ and $S \bowtie R$ could have different costs
 - estimation could depend on the algorithm



Example

Subset	Size(S)	Cost	Plan
{R,S}			$R \bowtie S$
{R,T}			
{R,U}			
{S,T}			
{S,U}			
{T,U}			
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	$T(R)T(S)/\max(V(R,b),V(S,b))$		$R \bowtie S$
{R,T}			
{R,U}			
{S,T}			
{S,U}			
{T,U}			
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	1k*1k/200		$R \bowtie S$
{R,T}			
{R,U}			
{S,T}			
{S,U}			
{T,U}			
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}			
{R,U}			
{S,T}			
{S,U}			
{T,U}			
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}			
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	$(1k * 1k * 1k) / (500 * 200) = 10k$	$5k + 0$	$(R \bowtie S) \bowtie T$
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	$(1k * 1k * 1k) / (500 * 200) = 10k$	$1,000k + 0$	$(R \bowtie T) \bowtie S$
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	$(1k * 1k * 1k) / (500 * 200) = 10k$	$2k + 0$	$(S \bowtie T) \bowtie R$
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}			
{R,T,U}			
{S,T,U}			
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}			



Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}	$(1k * 1k * 1k * 1k) / (1k * 500 * 200 * 100) = 10k$	$10k + 2k + 0 + 0 = 12k$	$((S \bowtie T) \bowtie R) \bowtie U$

Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}	$(1k * 1k * 1k * 1k) / (1k * 500 * 200 * 100) = 10k$	$50k + 5k + 0 + 0 = 55k$	$((R \bowtie S) \bowtie U) \bowtie T$

Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}	$(1k * 1k * 1k * 1k) /$ $(1k * 500 * 200 * 100) = 10k$	$10k + 1k + 0 + 0 = 11k$	$((T \bowtie U) \bowtie R) \bowtie S$

Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}	$((1k * 1k * 1k * 1k) / (1k * 500 * 200 * 100)) = 10k$	$2k + 1k + 0 + 0 = 3k$	$((T \bowtie U) \bowtie S) \bowtie R$

Example

Subset	Size(S)	Cost	Plan
{R,S}	5k	0	$R \bowtie S$
{R,T}	1,000k	0	$R \bowtie T$
{R,U}	10k	0	$R \bowtie U$
{S,T}	2k	0	$S \bowtie T$
{S,U}	1,000k	0	$S \bowtie U$
{T,U}	1k	0	$T \bowtie U$
{R,S,T}	10k	2k	$(S \bowtie T) \bowtie R$
{R,S,U}	50k	5k	$(R \bowtie S) \bowtie U$
{R,T,U}	10k	1k	$(T \bowtie U) \bowtie R$
{S,T,U}	2k	1k	$(T \bowtie U) \bowtie S$
{R,S,T,U}	10k	3k	$((T \bowtie U) \bowtie S) \bowtie R$



Physical Query Plan

- At each node of logical query plan:
 - Replace relational algebra operator with execution algorithms we discussed
 - How much memory is available?
 - Are relations indexed?
 - Are results sorted?
 - Should we pipeline or materialize intermediate results?




Pipelining vs. Materialization

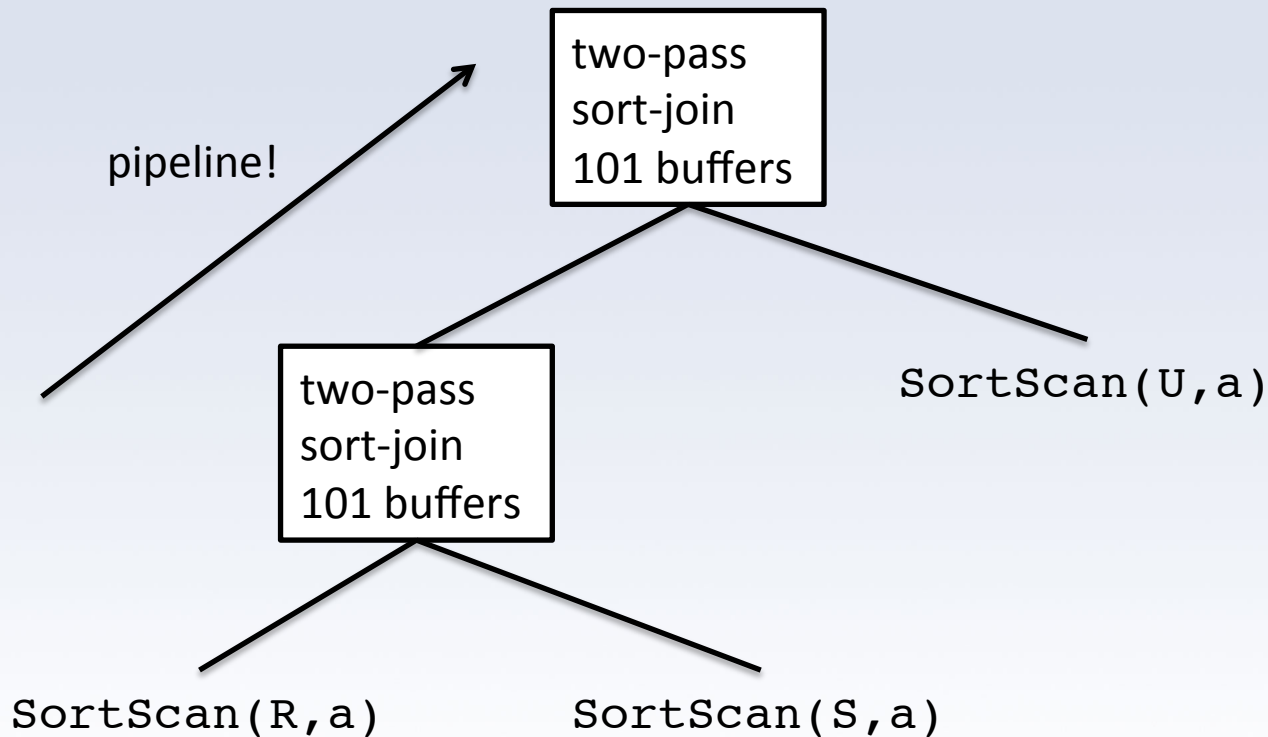
- Materialization - write the results of the operation back to the disk
 - results in 2 intermediate disk I/Os per block
- Pipelining - keep results in memory buffers
 - Results of one iterator passed to the next



Physical Query Plan



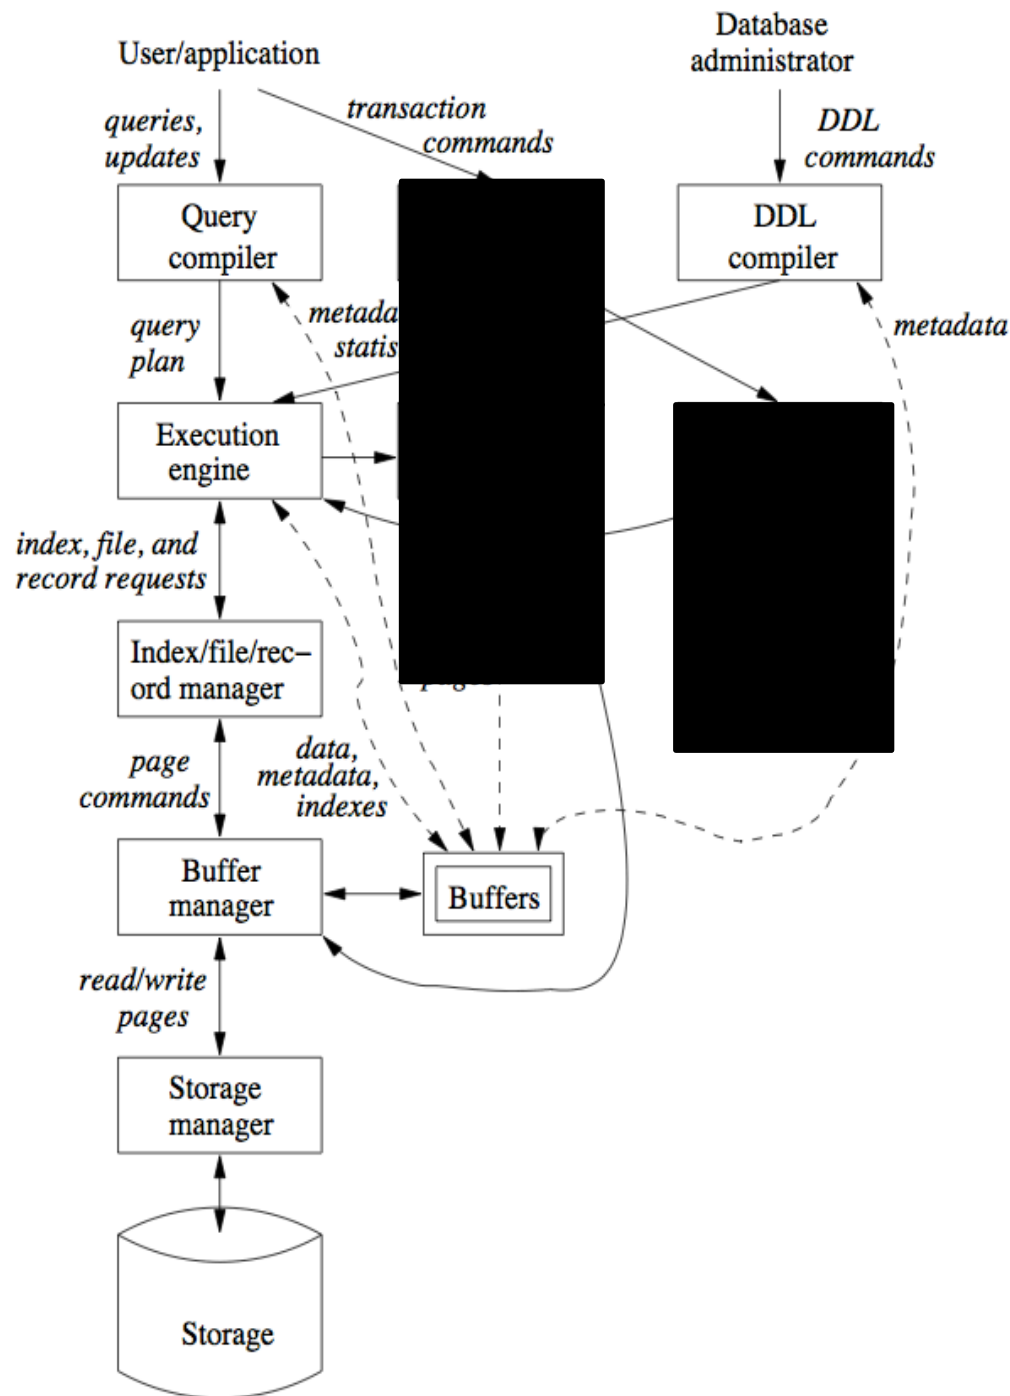
Physical Query Plan



MP3

- We'll give you
 - a series of query plans
 - a max buffer size (amount of memory)
 - a choice of implementation algorithms
- Your job:
 - improve the query plans





Next week...



